

Jon Spurring

Department of Computer Science,
University of Copenhagen

Learning to Program with F#

2022-09-01

Springer Nature

Preface

This book has been written as an introduction to programming for novice programmers. It is used in the first programming course at the University of Copenhagen's bachelor in computer science program. It has been typeset in \LaTeX , and all programs have been developed and tested in dotnet version 6.0.101

Jon Sporring
Professor, Ph.d.
Department of Computer Science,
University of Copenhagen
2022-09-01

Contents

1	Introduction	1
1.1	How to Learn to Solve Problems by Programming	1
1.2	How to Solve Problems	2
1.3	Approaches to Programming	3
1.4	Why Use F#	4
1.5	How to Read This Book	6
2	Solving problems by writing a program	7
2.1	Executing F# programs on a computer	9
2.2	Values have types and types reduce the risk of programming errors	11
2.3	Organizing often used code in functions	13
2.4	Asking the user for input	14
2.5	Conditionally execute code	15
2.6	Repeatedly execute code	16
2.7	Programming as a form of communication	18
2.8	Key concepts and terms in this chapter	20
3	Using F# as a Calculator	21

3.1	Literals and Basic Types	22
3.2	Operators on Basic Types	27
3.3	Boolean Arithmetic	31
3.4	Integer Arithmetic	32
3.5	Floating Point Arithmetic	35
3.6	Char and String Arithmetic	36
3.7	Programming Intermezzo: Hand Conversion Between Decimal and Binary Numbers	39
3.8	Key concepts and terms in this chapter	40
4	Values, Functions, and Statements	41
4.1	Value Bindings	45
4.2	Function Bindings	49
4.3	Do-Bindings	55
4.4	Conditional Expressions	55
4.5	Tracing code by hand	59
4.6	Key concepts and terms in this chapter	63
5	Programming with Types	65
5.1	Type products: Tuples	66
5.2	Type sums: Discriminated Unions	69
5.3	Records	72
5.4	Type Abbreviations	74
5.5	Variable Types	75
5.6	Key concepts and terms in this chapter	77
6	Making programs and documenting them	79

6.1	The 7-step guide to writing function	80
6.2	Programming as a communication activity	82
6.3	Key concepts and terms in this chapter	84
A	The Console in Windows, MacOS X, and Linux	87
A.1	The Basics	87
A.2	Windows	88
A.3	MacOS X and Linux	92
B	Number Systems on the Computer	97
B.1	Binary Numbers	97
B.2	IEEE 754 Floating Point Standard	98
C	Commonly Used Character Sets	103
C.1	ASCII	103
C.2	ISO/IEC 8859	104
C.3	Unicode	105
D	Common Language Infrastructure	109
	References	111
	Index	113

Chapter 1

Introduction

Programming is a creative process in which exciting problems may be solved and new tools and applications may be created. With programming skills, you can create high-level applications to run on a mobile device that interact with other users, databases, and artificial intelligence; you may create programs that run on supercomputers for simulating weather systems on alien planets or social phenomena in the internet economy; and you may create programs that run on small custom-made hardware for controlling your home appliances.

1.1 How to Learn to Solve Problems by Programming

In order to learn how to program, there are a couple of steps that are useful:

1. Choose a programming language: A programming language, such as F#, is a vocabulary and a set of grammatical rules for instructing a computer to perform a certain task. It is possible to program without a concrete language, but your ideas and thoughts must still be expressed in some fairly rigorous way. Theoretical computer scientists typically do not rely on computers nor programming languages but uses mathematics to prove properties of algorithms. However, most computer scientists program using a computer, and with a real language you have the added benefit of checking your algorithm, and hence your thoughts, rigorously on a real computer. This book teaches a subset of F#. The purpose is not to be a reference guide to this language but to use it as a vessel to teach you, the reader, how to convert your ideas into programs.
2. Learn the language: A computer language is a structure for thought, and it influences which thoughts you choose to express as a program, and how you

choose to do it. Any conversion requires you to acquire a sufficient level of fluency in order for you to be able to make programs. You do not need to be a master in F# nor to know every corner of the language, and you will expand your knowledge as you expose yourself to solving problems in the language, but you must invest an initial amount of time and energy in order to learn the basics of the language. This book aims at getting you started quickly, which is why we intentionally teach just a small subset of F#. On the internet and through other works you will be able to learn much more.

3. Practice: In order to be a good programmer, the most essential step is: practice, practice, practice! It has been estimated that to master anything, then you have to have spent at least 10000 hours practicing, so get started logging hours! It of course matters, how you practice. This book teaches a number of different programming themes. The point is that programming is thinking, and the scaffold you use shapes your thoughts. It is therefore important to recognize this scaffold and to have the ability to choose one which suits your ideas and your goals best. The best way to expand your abilities is to sharpen your present abilities, push yourself into new territory, and try something new. Do not be afraid to make errors or be frustrated at first. These are the experiences that make you grow.
4. Solve real problems: I have found that using my programming skills in real situations with customers demanding specific solutions, has forced me to put the programming tools and techniques that I use into perspective. Sometimes a task requires a cheap and fast solution, other times customers want a long-perspective solution with bug fixes, upgrades, and new features. Practicing solving real problems helps you strike a balance between the two when programming. It also allows makes you a more practical programmer, by allowing you to recognize its applications in your everyday experiences. Regardless, real problems create real programmers.

1.2 How to Solve Problems

Programming is the act of solving a problem by writing a program to be executed on a computer. A general method for solving problems, given by George Pólya [8] and adapted to programming, is:

Understand the problem: To solve any problem it is crucial that the problem formulation is understood. What is to be solved? Do you understand everything in the description of the problem? Is all information for finding the solution available or is something missing?

Design a plan: Good designs lead to programs are faster to implement, easier to find errors in, and easier to update in the future. Before you start typing a program consider things like: What are the requirements and constraints for the program? Which components should the program have? How are these components supposed to work together? Designing often involves drawing a diagram of the program and writing program sketches on paper.

Implement the plan: Implementation is the act of transforming a program design into code. A crucial part of any implementation is choosing which programming language to use. Furthermore, the solution to many problems will have a number of implementations which vary in how much code they require, to which degree they rely on external libraries, which programming style they are best suited for, what machine resources they require, and how long time they take to run on a computer. With a good design, the coding is usually easy, since the design will have uncovered the major issues and found solutions for these, but sometimes implementation reveals new problems, which require rethinking the design. Most often the implementation step also require a careful documentation of key aspects of the code, e.g., a user manual for the user, and internal notes for fellow programmers that are to maintain and update the code in the future.

Reflect on the result: A crucial part of any programming task is ensuring that the program solves the problem sufficiently. Ask yourself questions such as: What are the program's errors, is the documentation of the code sufficient and relevant for its intended use? Is the code easily maintainable and extendable by other programmers? Which parts of your method would you avoid or replicate in future programming sessions? Can you reuse some of the code you developed in other programs?

Programming is a very complicated process, and Pólya's list is a useful guide but not a fail-safe approach. Always approach problem-solving with an open mind.

1.3 Approaches to Programming

This book focuses on several fundamentally different approaches to programming:

Declarative programming emphasizes *what a program shall accomplish* but not *how*. We will consider Functional programming as an example of declarative programming. A *functional programming* language evaluates *functions* and avoids state changes. The program consists of *expressions* instead of statements. As an example, the function $f(x) = x^2$ takes a number x , evaluates the expression x^2 , and returns the resulting number. Everything about the function may be

characterized by the relation between the input and output values. Functional programming has its roots in lambda calculus [1]. The first language emphasizing functional programming was Lisp [6].

Imperative programming emphasizes *how a program shall accomplish a solution* and focusses less on *what the solution is*. A cooking recipe is an example of the spirit of imperative programming, where the recipe emphasizes what should be done in each step rather than describing the result. For example, a recipe for bread might tell you to first mix yeast and water, then add flour, etc. In imperative programming what should be done are called *statements* and in the recipe analogy, the steps are the statements. Statements influence the computer's *states*, in the same way that adding flour changes the state of our dough. Almost all computer hardware is designed to execute low-level programs written in imperative style. Imperative programming builds on the Turing machine [9]. As a historical note, the first major language was FORTRAN [5] which emphasized an imperative style of programming.

Structured programming emphasizes organization of programs in units of code and data. For example, a traffic light may consist of a state (red, yellow, green), and code for updating the state, i.e., switching from one color to the next. We will focus on Object-oriented programming as the example of structured programming. *Object-oriented programming* is a type of programming, where the code and data are structured into *objects*. E.g., a traffic light may be an object in a traffic-routing program. The first object-oriented programming language was Simula 67 developed by Dahl and Nygaard at the Norwegian Computing Center in Oslo [2].

Event-driven programming, which is often used when dynamically interacting with the real world. This is useful, for example, when programming graphical user interfaces, where programs will often need to react to a user clicking on the mouse or to text arriving from a web-server to be displayed on the screen. Event-driven programs are often programmed using *call-back functions*, which are small programs that are ready to run when events occur.

Most programs do not follow a single programming paradigm as, e.g., one of the above, but are a mix. Nevertheless, this book will treat each paradigm separately to emphasize its advantages and disadvantages.

1.4 Why Use F#

This book uses F#, also known as Fsharp, which is a functional first programming language, meaning that it is designed as a functional programming language that also

supports imperative and object-oriented programming. It was originally developed for Microsoft's .Net platform but is available as open source for many operating systems through Mono. As an introduction to programming, F# is a young programming language still under development, with syntax that at times is a bit complex. Still, it offers a number of advantages:

Interactive and compile mode: F# has an interactive and a compile mode of operation. In interactive mode you can write code that is executed immediately in a manner similar to working with a calculator, while in compile mode you combine many lines of code possibly in many files into a single application, which is easier to distribute to people who are not F# experts and is faster to execute.

Indentation for scope: F# uses indentation to indicate scope. Some lines of code belong together and should be executed in a certain order and may share data. Indentation helps in specifying this relationship.

Strongly typed: F# is strongly typed, reducing the number of runtime errors. That is, F# is picky, and will not allow the programmer to mix up types such as numbers and text. This is a great advantage for large programs.

Multi-platform: F# is available on Linux, Mac OS X, Android, iOS, Windows, GPUs, and browsers both via the public domain Mono platform and Microsoft's open source .Net system.

Free to use and open source: F# is supported by the Fsharp foundation (<http://fsharp.org>) and sponsored by Microsoft.

Assemblies: F# is designed to be able to easily communicate with other .Net and Mono programs through the language-independent, platform-independent byte-code called Common Intermediate Language (CIL) organized as assemblies. Thus, if you find that certain parts of a program are easy to express in F# and others in C++, then you will be able to combine these parts later into a single program.

Modern computing: F# supports all aspects of modern computing including Graphical User Interfaces, Web programming, Information rich programming, Parallel algorithms, . . .

Integrated development environments (IDE): F# is supported by IDEs such as Visual Studio (<https://www.visualstudio.com>) and Xamarin Studio (<https://www.xamarin.com>).

1.5 How to Read This Book

Learning to program requires mastering a programming language, however, most programming languages contain details that are rarely used or used in contexts far from a specific programming topic. Hence, this book only includes a subset of F# but focuses on language structures necessary to understand several common programming paradigms: Imperative programming mainly covered in 4–??, functional programming mainly covered in ??–??, object-oriented programming in ????, and event-driven programming in ??. A number of general topics are given in the appendix for reference. For further reading please consult <http://fsharp.org>.

Chapter 2

Solving problems by writing a program

Abstract In this chapter, you will find a quick introduction to several essential programming constructs with several examples that you can try on your computer using the `dotnet` command in your console. All constructs will be discussed in further detail in the following chapters. In this chapter, you will get a peek at:

- How to execute an F# program.
- How to perform simple arithmetic using F#.
- What types are and why they are important.
- How to write to and obtain written input from the user.
- How to perform conditional execution of code.
- How to define functions.
- How to repeat code without having to rewrite them.
- How to add textual comments to help yourself and other programmers understand your programs.

Programming is the art of solving problems by writing a program to be executed by a computer. For example, to solve the following problem,

Problem 2.1

What is the sum of 357 and 864?

we have written the program shown in Listing 2.1. In this book, we will show many

Listing 2.1 quickStartSum.fsx:

A script to add 2 numbers and print the result to the console.

```
1 let a = 357
2 let b = 864
3 let c = a + b
4 do printfn "%A" c

1 $ dotnet fsi quickStartSum.fsx
2 1221
```

programs, and for most, we will also show the result of executing the programs on a computer. Listing 2.1 shows both our program and how this program is executed on a computer. In the listing, we see our program was saved as a script in a file called `quickStartSum.fsx`, and in the console (also known as the terminal and the command-line) we executed the program by typing the command `dotnet fsi quickStartSum.fsx`. The result is then printed by the computer to the console as 1221. The colors are not part of the program but have been added to make it easier for us to identify different syntactical elements of the program.

The program consists of several lines. Our listing shows line numbers to the left. These are not part of the program but added for ease of discussion, since the order in which the lines appear the program matters. In this program, each line contains *expressions*, and this program has `let`-, `do`-expressions, and an addition. `let`-expressions defines aliases, and `do`-expressions defines computations. `let` and `do` are examples of *keywords*, and “+” is an example of an *operator*. Keywords, operators, and other sequences of characters, which F# recognizes are jointly called *lexemes*.

Reading the program from line 1, the first expression we encounter is `let a = 357`. This is known as a *let-binding* in F# and defines the equivalence between the name `a` and the value 357. F# does not accept a keyword as a name in a `let`-bindings. The consequence of this line is that in later lines there is no difference between writing the name `a` and the value 357. Similarly in line 2 the value 864 is bound to the name `b`. In contrast, line 3 contains an addition and a `let`-expression. It is at times useful to simulate the execution the computer does in a step-by-step manner by replacement:

`let c = a + b` \rightsquigarrow `let c = 357 + 864` \rightsquigarrow `let c = 1221`

Thus, since the expression on the right-hand side of the equal sign is evaluated, the result of line 3 is that the name `c` is bound to the value 1221.

Line 4 has a `do`-expression is also called a *do-binding* or a *statements*. In this `do`-binding, the *printfn* function `printfn` is called with 2 arguments, `"%A"` and `c`. All functions return values, and `printfn` the value 'nothing', which is denoted `()`. This function is very commonly used but also very special since it can take any number of arguments and produces output to the console. We say that "the output is printed to the screen". The first argument is called the *formatting string* and describes, what should be printed and how the remaining arguments if any, should be formatted. In this case, the value `c` is printed as an integer followed by a newline. Notice that in contrast to many other languages, F# does not use parentheses to frame the list of function arguments, nor does it use commas to separate them.

2.1 Executing F# programs on a computer

The main purpose of writing programs is to make computers execute or run them. F# has two modes of execution, *interactive* and *compiled*. Interactive mode allows the user to interact with F# as a dialogue: The user writes statements, and F# responds immediately. If a program has been saved as a file as in Listing 2.1 we do not need to rewrite the complete program every time we wish to execute it but can give the file as input to the F#'s interactive mode as demonstrated in Listing 2.1. Interactive mode is well suited for small experiments or back-of-an-envelope calculations, but not for programming in general, since each line is interpreted anew every time the program is run. In contrast, in compile mode, dotnet interprets the content of a source file once, and writes the result to disk, such that every when the user wishes to run the program, the interpretation step is not performed. For large programs, this can save considerable time. In the first chapters of this book, we will use interactive mode, and compile mode will be discussed in further detail in ??.

An interactive session is obtained by starting the console, typing the `fsharpi` command, typing the lines of the program, and ending the script-fragment with `;;`. The dialogue in Listing 2.2 demonstrates the workflow. What the user types has been highlighted by a box.

Listing 2.2: An interactive session.

```
1 $ dotnet fsi
2
3 Microsoft (R) F# Interactive version 12.0.0.0 for F# 6.0
4 Copyright (c) Microsoft Corporation. All Rights Reserved.
5
6 For help type #help;;
7
8 > let a = 3
9 - do printfn "%A" a;;
10 3
11 val a : int = 3
12 val it : unit = ()
13
14 > #quit;;
```

We see that after typing `fsharp`, the program starts by stating details about itself. Then F# writes `>` indicating that it is ready to receive commands. The user types `let a = 3` and presses `enter`, to which the interpreter responds with `-`. This indicates that the line has been received, that the script fragment is not yet completed, and that it is ready to receive more input. When the user types `do printfn "%A" a;;` followed by `enter`, then by `“;;”` the interpreter knows that the script-fragment is completed, it interprets the script-fragment, responds with `3` and some extra information about the entered code, and with `>` to indicate that it is ready for more script-fragments. The interpreter is stopped when the user types `#quit;;`. It is also possible to stop the interpreter by typing `ctrl-d`.

The interactive session results in extra output on the *type inference* performed. In Listing 2.2, F# states that the name `a` has *type* `int` and the value `3`. Likewise, the `do` statement F# refers to by the name `it`, and it has the type `unit` and value `“()”`. Types are very important to F# since they define how different program pieces fit together like lego bricks. They are a key ingredient for finding errors in programs, also known as *debugging*, and much of the rest of this book is concerned with types.

Instead of running `fsharp` interactively, we can write the script-fragment from Listing 2.2 into a file, here called `gettingStartedStump.fsx`. This file can be interpreted directly by `dotnet fsi` as shown in Listing 2.3.

Listing 2.3: Using the interpreter to execute a script.

```
1 $ dotnet fsi gettingStartedStump.fsx
2 3
```

Notice that in the file, `“;;”` is optional. In comparison to Listing 2.2, we see that the interpreter executes the code and prints the result on screen without the extra type information.

Files are important when programming, and F# and the console interprets files differently by the filename's suffix. A filename's suffix is the sequence of letters after the period in the filename. Generally, there are two types of files: *source code* and compiled programs. Until ??, we will concentrate on script files, which are source code, written in human-readable form using an editor, and has `.fsx` or `.fsscript` as suffix. In Table 2.1 is a complete list of possible suffixes used by F#.

Suffix	Human readable	Description
<code>.fs</code>	Yes	An <i>implementation file</i> , e.g., <code>myModule.fs</code>
<code>.fsi</code>	Yes	A <i>signature file</i> , e.g., <code>myModule.fsi</code>
<code>.fsx</code>	Yes	A <i>script file</i> , e.g., <code>gettingStartedStump.fsx</code>
<code>.fsscript</code>	Yes	Same as <code>.fsx</code> , e.g., <code>gettingStartedStump.fsscript</code>
<code>.dll</code>	No	A <i>library file</i> , e.g., <code>myModule.dll</code>
<code>.exe</code>	No	A stand-alone <i>executable file</i> , e.g., <code>gettingStartedStump.exe</code>

Table 2.1 Suffixes used, when programming F#.

2.2 Values have types and types reduce the risk of programming errors

Types are a central concept in F#. In the script 2.1 we bound values of integer type to names. There are several different integer types in F#, here we used the one called `[int]`. The values were not *declared* to have these types, instead the types were *inferred* by F#. Typing these bindings line by line in an interactive session, we see the inferred types as shown in Listing 2.4. The interactive session displays the type using the

Listing 2.4: Inferred types are given as part of the response from the interpreter.

```

1 > let a = 357;;
2 val a: int = 357
3
4 > let b = 864;;
5 val b: int = 864
6
7 > let c = a + b;;
8 val c: int = 1221
9
10 > do printfn "%A" c;;
11 1221
12 val it: unit = ()
```

`val` keyword followed by the name used in the binding, its type, and its value. Since the value is also returned, the last `printfn` statement is superfluous. Notice that `printfn` is automatically bound to the name `it` of type `unit` and value `()`. F# insists on binding all statements to values, and in lack of an explicit name, it will use `it`. Rumor has it that `it` is an abbreviation for "irrelevant".

In mathematics, types are also an important concept. For example, a number may belong to the set of natural \mathbb{N} , integer \mathbb{Z} , or real numbers, where all 3 sets are infinitely large and $\mathbb{N} \subset \mathbb{Z} \subset \mathbb{R}$ as illustrated in Figure 2.1. For many problems,




Fig. 2.1 In mathematics, the sets of natural, integer, and real numbers are each infinitely large, and real contains integers which in turn contains the set of natural numbers.

working with infinite sets is impractical, and instead, a lot of work in the early days of the computer's history was spent on designing finite sets of numbers, which have many of the properties of their mathematical equivalent, but which also are efficient for performing calculations on a computer. For example, the set of integers in F# is called `int` and is the set $\{-2\,147\,483\,648 \dots 2\,147\,483\,647\}$.

Types are also important when programming. For example, a the string `"863"` and the `int 863` may conceptually be identical but they are very different in the computer. F# is very picky about types, and generally does not allow types to be mixed, as demonstrated in the interactive session in Listing 2.5. In this ex-

Listing 2.5: Mixing types is often not allowed.

```

1 > let a = 357;;
2 val a: int = 357
3
4 > let b = "863";;
5 val b: string = "863"
6
7 > let c = a + b;;
8
9     let c = a + b;;
10    -----^
11
12 /Users/jrh630/repositories/fsharp-book/src/stdin(3,13): error
    FS0001: The type 'string' does not match the type 'int'
```

ample, we see that adding a string to an integer results in an error. The *error message* contains much information. First, it illustrates where dotnet could not understand the input by `-----^`. Then it repeats where the error is found as `/User/.../src/stdin(3,13)`, which means that dotnet was started in the directory `/User/.../src/`, the input was given on the standard-input meaning the keyboard, and the error was detected on line 3, column 13. Then F# gives the error number and a description of the error. Error numbers are an underdeveloped feature

in F# and should be ignored. However, the verbal description often contains useful information for correcting the program. Here, we are informed that there is a type mismatch in the expression. The reason for the mismatch is that since `a` is an integer, then the “+” operator must be integer addition, and thus for the expression to be executable, `b` can only be an integer.

2.3 Organizing often used code in functions

`printfn` is an example of a built-in function, and very often we wish to define our own. For example, in longer programs, some code needs to be used in several places and defining functions to *encapsulate* such code can be a great advantage for reducing the length of code, debugging, and writing code, which is easier to understand by other programmers. A function is defined using a `let`-binding. For example, to define a function, which takes two integers as input and returns their sum, we write

Listing 2.6: Defining the function `sum`

```
1 let sum x y =  
2   x + y
```

What this means is that we bind the name `sum` as a function, which takes two arguments and adds them. Further, in the function, the arguments are locally referred to by the names `x` and `y`. Indentation determines which lines should be evaluated when the function is called, and in this case, there is only one. The value of the last expression evaluated in a function is its return value. Here there is only one expression `x+y`, and thus, this function returns the value of the addition. This program does not do anything, since the function is neither called nor is its output used. However, we can modify Listing 2.1 to include it as shown in Listing 2.7. The output is the

Listing 2.7 `quickStartSumFct.fsx`:

Adding two integers with the use of a in-code defined function.

```
1 let sum x y =  
2   x + y  
3 let c = sum 357 864  
4 do printfn "%A" c  
  
-----  
1 $ dotnet fsi quickStartSumFct.fsx  
2 1221
```

same for the two programs, and the computation performed is almost the same. A step-by-step manner by replacement of the computation performed in line 3 is

```
let c = sum 357 864 ~> let c = 357 + 864 ~> let c = 1221
```

The main difference is that with the function `sum` we have an independent unit, which can be reused elsewhere in the code.

2.4 Asking the user for input

The `printfn` function allows us to write to the screen, which is useful, but sometimes we wish to start a dialogue with the user. One way to get user input is to ask the user to type something on the keyboard. Technically, input from the keyboard is called an *stdin stream*. This terminology is intended to remind us of characters streaming from the keyboard like the flow of water in a stream. Computer streams are different than water streams in that characters (or other items) only flow, when we ask for them. F# provides many libraries of prebuilt functions, and here we will use the `System.Console.ReadLine` function. The “.”-lexeme is read as `ReadLine` is a function which lies in `Console` which in turn lies in `System`. In the function documentation, we can read that `System.Console.ReadLine` takes a unit value as an argument and returns the *string* the user typed. A string is a built-in type, as is integers, and strings contains sequences of characters. The program will not advance until the user presses the newline. An example of a program that multiplies two integers supplied by a user is given in Listing 2.8. In this program, we find a user

Listing 2.8 quickStartSumInput.fsx:

Asking the user for input. The user entered 6, pressed the return button, 2, and pressed return again.

```
1 let sum x y = x + y
2 printfn "Adding a and b"
3 printf "Enter a: "
4 let a = int (System.Console.ReadLine ())
5 printf "Enter b: "
6 let b = int (System.Console.ReadLine ())
7 let c = sum a b
8 do printfn "%A" c
```

```
1 dotnet fsi quickStartSumInput.fsx
2 Adding a and b
3 Enter a: 6
4 Enter b: 2
5 8
```

dialogue, and we have designed it such that we assume that the user is unfamiliar with the inner workings of our program, and therefore helps the user understand the

purpose of the input and the expected result. This is good programming practice. Here, we will not discuss the program line-to-line, but it is advised to the novice programmer to match what is printed on the screen and from where in the code, the output comes from. However, let us focus on line 4 and 4, which introduce two new programming constructs. In each of these lines, 3 things happen: First the `System.Console.ReadLine` function is called with the “()” value as argument. This reads all the characters, the user types, up until the user presses the return key. The return value is a string of characters such as “6”. This value is different from the integer 6, and hence, to later be able to perform integer-addition, we *cast* the string value to `int`, meaning that we call the function `int` to convert the string-value to the corresponding integer value. Finally, the result is bound to the names `a` and `b` respectively.

2.5 Conditionally execute code

Often problem requires code evaluated based on conditions, which only can be decided at *runtime*, i.e., at the time, when the program is run. Consider a slight modification of our problem as

Problem 2.2

Ask for two integer values from the user, a and b , and print the result of the integer division a/b .

To solve this problem, we must decide what to do, if the user inputs $b = 0$, since division by zero is ill-defined. This is an example of a user input error, and later, we will investigate many different methods for handling such errors, but here, we will simply write an error message to the user, if the desired division is ill-defined. Thus, we need to decide at *runtime*, whether to divide a and b or to write an error message. For this we will use the `match-with` expression. In this program, the `match-with` expression covers line 7 to 12. When the computer executes these lines, it checks the value b against a list of patterns separated by “|”. The code belonging to each pattern follows the arrow, “ \rightarrow ”, and is called a *branch*, and which lines belongs to each branch is determined by *indentation*. Hence, the code belonging to the “`0.0`” branch is line 9 and to the “`_`” branch is line 11 to 12. The branches are checked one at a time from top to bottom. I.e., first b is compared with `0` which is equivalent to check whether $b = 0.0$. If this is true, then its branch is executed. Otherwise, the b is compared with the *wildcard* pattern, “`_`”. The wildcard pattern matches anything, and hence, if b is nonzero, then “`_`”-branch is executed. In most cases, F# will give an error, if the list of patterns do not cover the full domain of the type being matched. Here, b is an integer, and thus, we must write branches that takes *all* integer values into account. The wildcard pattern makes this easy and works as a catch-all-other cases, and is often placed as the last case, following the important cases. Assuming

Listing 2.9 quickStartDivisionInput.fsx:
Conditionally divide two user-given values.

```

1 let div x y = x / y
2 printfn "Dividing a by b"
3 printf "Enter a: "
4 let a = int (System.Console.ReadLine ())
5 printf "Enter b: "
6 let b = int (System.Console.ReadLine ())
7 match b with
8     0 ->
9         do printfn "Input error: Cannot divide by zero"
10    | _ ->
11        let c = div a b
12        do printfn "%A" c

```

```

1 % dotnet fsi quickStartDivisionInput.fsx
2 Dividing a by b
3 Enter a: 6
4 Enter b: 2
5 3
6 % dotnet fsi quickStartDivisionInput.fsx
7 Dividing a by b
8 Enter a: 6
9 Enter b: 0
10 Input error: Cannot divide by zero

```

that the user enters the value 0, then the step-by-step simplification of `match-with` expression is,

```

match b with 0 -> ... | _ -> ...
~> do printfn "Input error: Cannot divide by zero"

```

2.6 Repeatedly execute code

Often code needs to be evaluated many times or looped. For example, instead of stopping the program in Listing 2.9 if the user inputs $b = 0$, then we could repeat the question as many times as needed until the user inputs a non-zero value for b . This is called a loop, and there are several programming constructions for this purpose.

Let us first consider recursion. A recursive function is one, which calls itself, e.g., $f(f(f(\dots(x))))$ is an example of a function f which calls itself many times, possibly infinitely many. In the latter case, we say that the recursion has entered an infinite loop, and we will experience that either the program runs forever or that the execution stops due to a memory error. If we had infinite memory. To avoid this,

recursive functions must always have a stopping criterion. Thus, we can design a function for asking the user for a non-zero input value as shown in Listing 2.10. The function `readNonZeroValue` takes no input denoted by “()”, and repeatedly

Listing 2.10 `quickStartRecursiveInput.fsx`:
Recursively call `ReadLine` until a non-zero value is entered.

```

1 let rec readNonZeroValue () =
2     let a = int (System.Console.ReadLine ())
3     match a with
4     | 0 ->
5         printfn "Error: zero value entered. Try again"
6         readNonZeroValue ()
7     | _ ->
8         a
9 printfn "Please enter a non-zero value"
10 let b = readNonZeroValue ()
11 printfn "You typed: %A" b

```

```

1 $ dotnet fsi quickStartRecursiveInput.fsx
2 Please enter a non-zero value
3 0
4 Error: zero value entered. Try again
5 3
6 You typed: 3

```

calls itself until the $a \neq 0$ condition is met. It is recursive since its body contains a call to itself. For technical reasons, F# requires recursive functions to be declared by the `rec`-keyword as demonstrated. The function has been designed to stop if $a \neq 0$, and in F#, this is tested with the “<>” operator. Thus, if the stopping condition is satisfied, then the `then`-branch is executed, which does not call itself, and thus the recursion goes no deeper. If the condition is not met, then the `else`-branch is executed, and the function is eventually called anew. The example execution of the program demonstrates this for the case that the user first inputs the value 0 and then the value 3.

As an alternative to recursive functions, loops may also be implemented using the `while`-expression. In Listing 2.10 is an example of a solution where the recursive loop has been replaced with `while`-loop. As for other constructs, the lines to be repeated are indicated by indentation, in this case, lines 4 to 5, and in the end, the result of the `readNonZeroValueAlt` function is the last expression evaluated, which is the trivial expression `a` in line 6. In comparison with the recursive version of the program, the `while`-loop has a continuation conditions (line 3), i.e., the content of the loop is repeated as long as `a = 0` evaluates to `true`. Another difference is that in Listing 2.10 we could simplify our program to only using `let` value-bindings, here we need a new concept: *variables* also known as a `mutable` value. Mutable values allow us to update the value associated with a given name. Thus, the value associated with a name of mutable type depends on when it is accessed.

Listing 2.11 quickStartWhileInput.fsx:
Replacing recursion in Listing 2.10 with a `while`-loop.

```

1 let readNonZeroValueAlt () =
2     let mutable a = int (System.Console.ReadLine ())
3     while a = 0 do
4         printfn "Error: zero value entered. Try again"
5         a <- int (System.Console.ReadLine ())
6     a
7 printfn "Please enter a non-zero value"
8 let b = readNonZeroValueAlt ()
9 printfn "You typed: %A" b

```

```

1 $ dotnet fsi quickStartWhileInput.fsx
2 Please enter a non-zero value
3 0
4 Error: zero value entered. Try again
5 3
6 You typed: 3.0

```

This construction makes programs much more complicated and error-prone, and their use should be minimized. The syntax of mutable values is that first it should be defined with the `mutable`-keyword as shown in line 2, and when its value is to be updated then the “<-”-notation must be used as demonstrated in line 5. Note that the execution of the two programs Listing 2.10 and Listing 2.11, gives identical output, when presented with identical input. Hence, they solve the same problem by two quite different means. This is a common property of solutions to problems as a program: Often several different solutions exist, which are identical on the surface, but where the quality of the solution depends on how quality is defined and which programming constructions have been used. Here, the main difference is that the recursive solution avoids the use of mutable values, which turns out to be better for proving the correctness of programs and for adapting programs to super-computer architectures. However, recursive solutions may be very memory intensive, if the recursive call is anywhere but the last line of the function.

2.7 Programming as a form of communication

When programming it is important to consider the time dimension of a program. Some usually very small programs are only used for a short while, e.g., to test a programming construction or an idea to a solution. Others small as well as large may be used again and again over a long period, and possibly given to other programmers to use, maintain, and extend. In this case, programming is an act of communication, where what is being communicated is the solution to a problem as well as

the thoughts behind the chosen solution. Common experiences among programmers are that it is difficult to fully understand the thoughts behind a program written by a fellow programmer from its source code alone, and for code written perhaps just weeks earlier by the same programmer, said programmer can find it difficult to remember the reasons for specific programming choices. To support this communication, programmers use *code-comments*. As a general concept, this is also called in-code documentation. Documentation may also be an accompanying manual or report. Documentation serves several purposes:

1. Communicate what the code should be doing, e.g., describe functions in terms of their input-output relation.
2. Highlight big insights essential for the code.
3. Highlight possible conflicts and/or areas where the code could be changed later.

F# has two different syntaxes for comments. A block comment is everything bracketed by `(* *)`, and a line comment, which is everything between `//` and the end of the line. For example, adding comments to Listing 2.10 could look like Listing 2.12. Comments are ignored by the computer and serve solely as programmer-

Listing 2.12 `quickStartRecursiveInputComments.fsx`:
Adding comments to Listing 2.10.

```

1  (*
2     Demonstration of recursion for keyboard input.
3     Author: Jon Spurring
4     Date: 2022/7/28
5  *)
6
7  // Description: Repeatedly ask the user for a non-zero number
8  // until a non-zero value is entered.
9  // Arguments: None
10 // Result: the non-zero value entered
11 let rec readNonZeroValue () =
12     // Note that the value of a is different for every
13     // recursive call.
14     let a = int (System.Console.ReadLine ())
15     match a with
16     | 0 ->
17         printfn "Error: zero value entered. Try again"
18         readNonZeroValue ()
19     | _ ->
20         a
21 printfn "Please enter a non-zero value"
22 let b = readNonZeroValue ()
23 printfn "You typed: %A" b

```

to-programmer communication, there are no or few rules for specifying, what is good and bad documentation of a program. The essential point is that coding is a

journey in problem-solving, and proper documentation is an aid in understanding the solution and the journey that lead to it.

2.8 Key concepts and terms in this chapter

- F# has two modes of operation: **Interactive** and **compile** mode. The first chapters of this book will focus on the interactive mode.
- F# is accessed through the **console/terminal/command-line**, which is another program, in which text commands can be given such as starting the dotnet program in interactive mode.
- Programs are written in a human-readable form called the **source-code**.
- Source code consists of several syntactical elements such as **operators** such as "*" and "<-", **keywords** such as "let" and "while", **values** such as 1.2 and the string "hello world", and **user-defined names** such as "a" and "str". All words, which F# recognizes are called **lexemes**.
- A program consists of a sequence of **expressions**, which comes in two types: **let** and **do**.
- Values have **types** such as **int** and **string**. When performing calculations, the type defines which calculations can be done.
- **Functions** are a type of value and defined using a let-binding. They are used to encapsulate code to make the code easier to read and understand and to make code reusable.
- The **conditional match-with** expression is used to control what code is to be executed at **runtime**. Each piece of conditional code is called a **branch**.
- **Recursion** and **while**-loops are programming structure to execute the same code several times.
- **Mutable values** are in contrast to **immutable values** may change value over time, and makes programmer harder to understand.
- **Comments** are **in-code documentation** and are ignored by the computer but serve as an important tool for communication between programmers.

Chapter 3

Using F# as a Calculator

Abstract In the previous chapter, we introduced some key F# programming tools and concepts without going into depth on opportunities and limitations. In the following chapters, we will dive deeper into methods for solving problems by writing programs, and what facilities are available in F# to express these solutions. As a first step, we must acquaint ourselves with the basic building blocks of basic types, constants, and operators, and this chapter includes

- An introduction to the basic types and how to write constants of those types.
- Arithmetic of basic operations.

with these tools, you will be able to evaluate expressions as if F# were a simple calculator. Examples of problems, you will be able to solve after reading this chapter, is:

- What is the result of $3 \cos(4 * \pi/180) + 4 \sin(4 * \pi/180)$.
- Calculate how many characters are there in the text string "Hello World!".
- What is the ASCII value of the character 'J'.
- How to convert between whole and binary numbers.

3.1 Literals and Basic Types

All programs rely on processing of data, and an essential property of data is its *type*. A *literal* is a fixed value like the number 3, and if we type the number 3 in an interactive session at the input prompt, then F# responds as shown in Listing 3.1. What this means is that F# has inferred the type to be *int* and bound it to the

Listing 3.1: Typing the number 3.

```
1 > 3;;  
2 val it: int = 3
```

identifier *it*. For more on binding and identifiers see Chapter 4. Types matter, since the operations that can be performed on integers, are quite different from those that can be performed on, e.g., strings. Therefore, the number 3 has many different representations as shown in Listing 3.2. Each literal represents the number 3, but

Listing 3.2: Many representations of the number 3 but using different types.

```
1  
2 > 3;;  
3 val it: int = 3  
4  
5 > 3.0;;  
6 val it: float = 3.0  
7  
8 > '3';;  
9 val it: char = '3'  
10  
11 > "3";;  
12 val it: string = "3"
```

their types are different, and hence they are quite different values. The types *int* for integer numbers, *float* for floating point numbers, *bool* for Boolean values, *char* for characters, and *string* for strings of characters are the most common types of literals. A table of all *basic types* predefined in F# is given in Table 3.1. In addition to these built-in types, F# is designed such that it is easy to define new types.

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base 10*, which means that a value is represented as two sequences of decimal digits separated by a *decimal point*, where each *digit* *d* has a position and a value $d \in \{0, 1, 2, \dots, 9\}$. The part before the decimal point is called the *whole part* and the part after is called the *fractional part* of the number. An *integer* is a number with only a whole part and neither a decimal point nor a fractional part. As an example 35.7 is a decimal number, whose value is $3 \cdot 10^1 + 5 \cdot 10^0 + 7 \cdot 10^{-1}$, and 128 is an integer, whose value is $1 \cdot 10^2 + 2 \cdot 10^1 + 8 \cdot 10^0$. In F#, a decimal number is called a *floating point number*. Floating point numbers may alternatively be given

Metatype	Type name	Description
Boolean	<u>bool</u>	Boolean values true or false
Integer	<u>int</u>	Integer values from -2,147,483,648 to 2,147,483,647
	byte	Integer values from 0 to 255
	sbyte	Integer values from -128 to 127
	int8	Synonymous with sbyte
	uint8	Synonymous with byte
	int16	Integer values from -32768 to 32767
	uint16	Integer values from 0 to 65535
	int32	Synonymous with int
	uint32	Integer values from 0 to 4,294,967,295
	int64	Integer values from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
	uint64	Integer values from 0 to 18,446,744,073,709,551,615
Real	<u>float</u>	64-bit IEEE 754 floating point value from $-\infty$ to ∞
	double	Synonymous with float
	single	A 32-bit floating point type
	float32	Synonymous with single
	decimal	A floating point data type that has at least 28 significant digits
Character	<u>char</u>	Unicode character
	<u>string</u>	Unicode sequence of characters
None	<u>unit</u>	The value ()
Object	<u>obj</u>	An object
Exception	<u>exn</u>	An exception

Table 3.1 List of some of the basic types. The most commonly used types are underlined. For a description of integer see Appendix B.1, for floating point numbers see Appendix B.2, for ASCII and Unicode characters see Appendix C, for objects see ??, and for exceptions see ??.

using *scientific notation*, such as 3.5×10^{-4} and 4×10^2 , where the e-notation is translated to a value as $3.5 \times 10^{-4} = 3.5 \cdot 10^{-4} = 0.00035$, and $4 \times 10^2 = 4 \cdot 10^2 = 400$.

The basic unit of information in almost all computers is the binary digit or *bit* for short. Internally, programs and data are all represented as bits, hence F# has a strong support for binary numbers. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. E.g., the binary number $101.01_2 = 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 5.25$. Subscripts are often used to indicate the base of a number, e.g., 101.01_2 and 101.01_{10} are different numbers. Since base 10 is so common, the subscript for base 10 numbers is often omitted.

Binary numbers are closely related to *octal* and *hexadecimal numbers*. Octals uses 8 as basis and hexadecimals use 16 as basis. Each octal digit can be represented by exactly three bits, and each hexadecimal digit can be represented by exactly four bits. The hexadecimal digits use 0–9 to represent the values 0–9 and a–f in lower or alternatively upper case to represent the values 10–15. Thus, Octals and hexadecimals conveniently serve as shorthand for the much longer binary representation. As examples, the octal number 37_8 is $3 \cdot 8^1 + 7 \cdot 8^0 = 31$, and the hexadecimal number $f3_{16}$ is $15 \cdot 16^1 + 3 \cdot 16^0 = 243$.

To denote integers in bases different than 10, F# uses the prefix '0b' for binary, '0o' for octal, and '0x' for hexadecimal numbers. For example, the value 367_{10} may be written as an integer 367, as a binary number `0b101101111`, as a octal number `0o557`, and as a hexadecimal number `0x16f`. In F#, the character sequences `0b12` and `ff` are not recognized as numbers.

A *character* is a *Unicode code point*, and character literals are enclosed in single quotation marks. Appendix C.3 contains more details on code points. The character type in F# is denoted *char*. Examples of characters are 'a', 'D', '3', and examples of non-characters are '23' and 'abc'. Some characters, such as the tabulation character, do not have a visual representation. These can still be represented as a character using *escape sequences*. A character escape sequence starts with “\” followed by letter for simple escapes such as `\t` for tabulation and `\n` for newline. Escape sequences can also be a numerical representation of a code point, and three versions exist: The trigraph `\DDD`, where D is a decimal digit, is used to specify the first 256 code points, the hexadecimal escape codes `\uXXXX`, where X is a hexadecimal digit, is used to specify the first 65536 code points, and `\UXXXXXXXX` is used to specify any of the approximately $4.3 \cdot 10^9$ possible code points. All escape sequences are shown in Table 3.2. Examples of *char* representations of the letter 'a' are: 'a', '\097',

Character	Escape sequence	Description
BS	<code>\b</code>	Backspace
LF	<code>\n</code>	Line feed
CR	<code>\r</code>	Carriage return
HT	<code>\t</code>	Horizontal tabulation
\	<code>\\</code>	Backslash
"	<code>\"</code>	Quotation mark
'	<code>\'</code>	Apostrophe
BEL	<code>\a</code>	Bell
FF	<code>\f</code>	Form feed
VT	<code>\v</code>	Vertical tabulation
	<code>\uXXXX</code> , <code>\UXXXXXXXX</code> , <code>\DDD</code>	Unicode character ('X' is any hexadecimal digit, and 'D' is any decimal digit)

Table 3.2 Escape characters. The escapecode `\DDD` is sometimes called a tricode.

'\u0061', '\U00000061'.

A *string* is a sequence of characters enclosed in double quotation marks. Examples are "a", "this is a string", and "-&#@". Note that the string "a" and the character 'a' are not the same. Some strings are so common that they are given special names: One or more spaces " " is called *whitespace*, and both "`\n`" and "`\r\n`" are called *newline*. The escape-character “\” may be used to break a line in two. This and other examples are shown in Listing 3.3. Note that the response from `fsharp` is shown in double quotation marks, but this is not part of the string.

F# supports *literal types*, where the type of a literal is indicated as a prefix or suffix as shown in Table 3.3.

Listing 3.3: Examples of string literals.

```

1 > "abcde";;
2 val it: string = "abcde"
3
4 > "abc
5   de";;
6 val it: string = "abc
7   de"
8
9 > "abc\
10  de";;
11 val it: string = "abcde"
12
13 > "abc\nde";;
14 val it: string = "abc
15 de"

```

Type	syntax	Examples	Value
int, int32	<int xint> <int xint>l	3, 0x3 3l, 0x3l	3
uint32	<int xint>u <int xint>ul	3u 3ul	3
byte, uint8	<int xint>uy '<char>'B	97uy 'a'B	97
byte[]	"<string>"B @"<string>"B	"a\n"B @"a\n"B	[[97uy; 10uy]] [[97uy; 92uy; 110uy]]
sbyte, int8	<int xint>y	3y	3
int16	<int xint>s	3s	3
uint16	<int xint>us	3us	3
int64	<int xint>L	3L	3
uint64	<int xint>UL <int xint>uL	3UL 3uL	3
float, double	<float> <xint>LF	3.0 0x013LF	3.0 9.387247271e-323
single, float32	<float>F <float>f <xint>lf	3.0F 3.0f 0x013lf	3.0 3.0 4.4701421e-43f
decimal	<float int>M <float int>m	3.0M, 3M 3.0m, 3m	3.0
string	"<string>" @"<string>" ""<string>""	"\"quote\".\n" @"\"quote\".\n" ""\"quote\".\n""	"quote".<newline> "quote\".\n. "quote\".\n

Table 3.3 List of literal types. The syntax notation <> means that the programmer replaces the brackets and content with a value of the appropriate form. The <xint> is one of the integers on hexadecimal, octal, or binary forms such as 0x17, 0o21, and 0b10001. The [|] brackets means that the value is an array, see ?? for details.

The literal type is closely connected to how the values are represented internally. For example, a value of type `int32` use 32 bits and can be both positive and negative, while a `uint32` value also use 32 bits, but is unsigned. A `byte` is an 8-bit number, and `sbyte` is a signed 8-bit number. Values of type `float` use 64 bits, while `float32` only uses 32 bits. The number of bits used to represent numbers directly relates to the

range and precision these types can represent. This is summarized in Table 3.1 and discussed in more detail in Appendix B. String literals may be *verbatim* by the `@`-notation or triple double quotation marks, meaning that the escape sequences are not converted to their code point. The two types of string verbatim treat quotation marks differently, as illustrated in the table. Further examples are shown in Listing 3.4.

Listing 3.4: Named and implied literals.

```
1 > 3;;
2 val it: int = 3
3
4 > 4u;;
5 val it: uint32 = 4u
6
7 > 5.6;;
8 val it: float = 5.6
9
10 > 7.9f;;
11 val it: float32 = 7.9000000095f
12
13 > 'A';;
14 val it: char = 'A'
15
16 > 'B'B;;
17 val it: byte = 66uy
18
19 > "ABC";;
20 val it: string = "ABC"
21
22 > @"abc\nde";;
23 val it: string = "abc\nde"
```

Many basic types are compatible, and the type of a literal may be changed by *typecasting*. An example of casting to a float is shown in Listing 3.5. When

Listing 3.5: Casting an integer to a floating point number.

```
1 > float 3;;
2 val it: float = 3.0
```

`float` is given an argument, then it acts as a function rather than a type, and for the integer 3 it returns the floating point number `3.0`. For more on functions see Chapter 4. Boolean values are often treated as the integer values 0 and 1, but no short-hand function names exist for their conversions. Instead, use functions from the `System.Convert` family of functions, as demonstrated in Listing 3.6. Here `System.Convert.ToBoolean` is the identifier of a function `ToBoolean`, which is a *member* of the *class* `Convert` that is included in the *namespace* `System`. Namespaces, classes, and members will be discussed in ??.

Typecasting is often a destructive operation, e.g., typecasting a `float` to `int` removes the fractional part without rounding as shown in Listing 3.7. Here we typecasted to

Listing 3.6: Casting booleans.

```

1 > System.Convert.ToBoolean 1;;
2 val it: bool = true
3
4 > System.Convert.ToBoolean 0;;
5 val it: bool = false
6
7 > System.Convert.ToInt32 true;;
8 val it: int = 1
9
10 > System.Convert.ToInt32 false;;
11 val it: int = 0

```

Listing 3.7: Fractional part is removed by downcasting.

```

1 > int 357.6;;
2 val it: int = 357

```

a lesser type, in the sense that the set of integers is a subset of floating point numbers, and this is called *downcasting*. The opposite is called *upcasting* and is often non-destructive, as Listing 3.5 showed. Since floating point numbers are a superset of integers, the value is retained. As a side note, *rounding* a number $y.x$, where y is the *whole part* and x is the *fractional part*, is the operation of mapping numbers in the interval $y.x \in [y.0, y.5)$ to y , and those in $y.x \in [y.5, y + 1)$ to $y + 1$. This can be performed by downcasting, as shown in Listing 3.8. I.e., $357.6 + 0.5 = 358.1$

Listing 3.8: Rounding by modified downcasting.

```

1 > int (357.6 + 0.5);;
2 val it: int = 358

```

and removing the fractional part by downcasting results in 358, which is the correct answer.

3.2 Operators on Basic Types

Expressions are the basic building block of all F# programs, and this section will discuss operator expressions on basic types. A typical calculation, such used in Listing 3.8, is

$$\underbrace{357.6}_{\text{operand}} \quad \underbrace{+}_{\text{operator}} \quad \underbrace{0.5}_{\text{operand}} \quad (3.1)$$

is an example of an arithmetic *expression*, and the above expression consists of two *operands* and an *operator*. Since this operator takes two operands, it is called a

binary operator. The expression is written using *infix notation*, since the operands appear on each side of the operator.

In order to discuss general programming structures, we will use a simplified language to describe valid syntactical structures. In this simplified language, the syntax of basic binary operators is shown in the following.

Listing 3.9: Syntax for a binary expression.

```
1 <expr><op><expr>
```

Here `<expr>` is any expression supplied by the programmer, and `<op>` is a binary infix operator. F# supports a range of arithmetic binary infix operators on its built-in types, such as addition, subtraction, multiplication, division, and exponentiation, using the “+”, “-”, “*”, “/”, “**” lexemes, respectively. Not all operators are defined for all types, e.g., addition is defined for integer and float types as well as for characters and strings, but multiplication is only defined for integer and floating-point types. A complete list of built-in operators on basic types is shown in Table 3.4 and 3.5, and a range of mathematical functions is shown in Table 3.6. Note that

Operator	bool	ints	floats	char	string	Example	Result	Description
+		✓	✓	✓	✓	5 + 2	7	Addition
-		✓	✓			5.0 - 2.0	3.0	Subtraction
*		✓	✓			5 * 2	10	Multiplication
/		✓	✓			5.0 / 2.0	2.5	Division
%		✓	✓			5 % 2	1	Remainder
**			✓			5.0 ** 2.0	25.0	Exponentiation
+		✓	✓			+3	3	identity
-		✓	✓			-3.0	-3.0	negation
&&	✓					true && false	false	boolean and
	✓					true false	true	boolean or
not	✓					not true	false	boolean negation
&&&		✓				0b101 &&& 0b110	0b100	bitwise boolean and
		✓				0b101 0b110	0b111	bitwise boolean or
^^^		✓				0b101 ^^^ 0b110	0b011	bitwise boolean exclusive or
<<<		✓				0b110uy <<< 2	0b11000uy	bitwise left shift
>>>		✓				0b110uy >>> 2	0b1uy	bitwise right shift
~~~		✓				~~~0b110uy	0b11111001uy	bitwise boolean negation

**Table 3.4** Arithmetic operators on basic types. Ints and floats means all built-in integer and float types. Note that for the bitwise operations, digits 0 and 1 are taken to be `true` and `false`.

expressions can themselves be arguments to expressions, and thus, 4+5+6 is also a legal statement. Technically, F# interprets the expression as (4+5)+6 meaning that first 4+5 is evaluated according to the `<expr><op><expr>` syntax. Then the result replaces the parenthesis to yield 9+6, which, once again, is evaluated according to the `<expr><op><expr>` syntax to give 15. This is called *recursion*, which is the name for a type of rule or function that uses itself in its definition. See ?? for more on recursive functions.

Operator	bool	ints	floats	char	string	Example	Result	Description
<	✓	✓	✓	✓	✓	<code>true &lt; false</code>	<code>false</code>	Less than
>	✓	✓	✓	✓	✓	<code>5 &gt; 2</code>	<code>true</code>	Greater than
=	✓	✓	✓	✓	✓	<code>5.0 = 2.0</code>	<code>false</code>	Equal
<=	✓	✓	✓	✓	✓	<code>'a' &lt;= 'b'</code>	<code>true</code>	Less than or equal
>=	✓	✓	✓	✓	✓	<code>"ab" &gt;= "cd"</code>	<code>false</code>	Greater than or equal
<>	✓	✓	✓	✓	✓	<code>5 &lt;&gt; 2</code>	<code>true</code>	Not equal

**Table 3.5** Comparison operators on basic types. Types cannot be mixed, e.g., `3 < 'a'` is a syntax error.

Operator	bool	ints	floats	char	string	Example	Result	Description
abs		✓	✓			<code>abs -3</code>	3	Absolute value
acos			✓			<code>acos 0.8</code>	0.644	Inverse cosine
asin			✓			<code>asin 0.8</code>	0.927	Inverse sinus
atan			✓			<code>atan 0.8</code>	0.675	Inverse tangent
atan2			✓			<code>atan2 0.8 2.3</code>	0.335	Inverse tangentvariant
ceil			✓			<code>ceil 0.8</code>	1.0	Ceiling
cos			✓			<code>cos 0.8</code>	0.697	Cosine
exp			✓			<code>exp 0.8</code>	2.23	Natural exponent
floor			✓			<code>floor 0.8</code>	0.0	Floor
log			✓			<code>log 0.8</code>	-0.223	Natural logarithm
log10			✓			<code>log10 0.8</code>	-0.0969	Base-10 logarithm
max	✓	✓	✓	✓		<code>max 3.0 4.0</code>	4.0	Maximum
min	✓	✓	✓	✓		<code>min 3.0 4.0</code>	3.0	Minimum
pown		✓				<code>pown 3 2</code>	9	Integer exponent
round			✓			<code>round 0.8</code>	1.0	Rounding
sign		✓	✓			<code>sign -3</code>	-1	Sign
sin			✓			<code>sin 0.8</code>	0.717	Sinus
sqrt			✓			<code>sqrt 0.8</code>	0.894	Square root
tan			✓			<code>tan 0.8</code>	1.03	Tangent

**Table 3.6** Predefined functions for arithmetic operations.

Unary operators take only one argument and have the syntax:

#### Listing 3.10: A unary expressions.

```
1 <op><expr>
```

An example of a unary operator is `-3`, where `-` here is used to negate a positive integer. Since the operator appears before the operand, it is a *prefix operator*.

The concept of *precedence* is an important concept in arithmetic expressions. If parentheses are omitted in Listing 3.8, then F# will interpret the expression as `(int 357.6) + 0.5`, which is erroneous since the addition of an integer with a float is undefined. This is an example of precedence, i.e., function evaluation takes precedence over addition which means that function evaluation is performed first and

addition second. Consider the arithmetic expression shown in Listing 3.11. Here, the

**Listing 3.11: A simple arithmetic expression.**

```
1 > 3 + 4 * 5;;
2 val it: int = 23
```

addition and multiplication functions are shown in infix notation with the *operator* lexemes “+” and “*”. To arrive at the resulting value 23, F# has to decide in which order to perform the calculation. There are 2 possible orders,  $3 + (4 * 5)$  and  $(3 + 4) * 5$  that gives different results. For integer arithmetic, the correct order is, of course, multiplication before addition, and we say that multiplication takes *precedence* over addition. Every atomic operation that F# can perform is ordered in terms of its precedence, and for some common built-in operators shown in Table 3.7, the precedence is shown by the order they are given in the table.

Operator	Associativity	Description
+<expr> -<expr> ~~~<expr>	Left	Unary identity, negation, and bitwise negation operators
f <expr>	Left	Function application
<expr> ** <expr>	Right	Exponentiation
<expr> * <expr> <expr> / <expr> <expr> % <expr>	Left	Multiplication, division and remainder
<expr> + <expr> <expr> - <expr>	Left	Addition and subtraction binary operators
<expr> ^^^ <expr>	Right	Bitwise exclusive or
<expr> < <expr> <expr> <= <expr> <expr> > <expr> <expr> >= <expr> <expr> = <expr> <expr> <> <expr> <expr> <<< <expr> <expr> >>> <expr> <expr> &&& <expr> <expr>     <expr>	Left	Comparison operators, bitwise shift, and bitwise 'and' and 'or'.
<expr> && <expr>	Left	Boolean and
<expr>    <expr>	Left	Boolean or

**Table 3.7** Some common operators, their precedence, and their associativity. Rows are ordered from highest to lowest precedences, such that <expr> * <expr> has higher precedence than <expr> + <expr>. Operators in the same row have the same precedence..

Associativity describes the order in which calculations are performed for binary operators of the same precedence. Some operator's associativity are given in Table 3.7. In the table we see that “*” is left associative, which means that  $3.0 * 4.0 * 5.0$  is evaluated as  $(3.0 * 4.0) * 5.0$ . Conversely, “**” is right associative, so  $4.0 ** 3.0 ** 2.0$  is evaluated as  $4.0 ** (3.0 ** 2.0)$ . For some operators, like multiplication, association matters little, e.g.,  $4 * 3 * 2 = 4 * (3 * 2) = (4 * 3) * 2$ , and for other operators, like exponentiation, association makes a huge difference, e.g.,

$4^{(3^2)} \neq (4^3)^2$ . Examples of this is shown in Listing 3.12. **Whenever in doubt of** ★

**Listing 3.12: Precedence rules define implicit parentheses.**

```

1 > 4.0 * 3.0 * 2.0;;
2 val it: float = 24.0
3
4 > (4.0 * 3.0) * 2.0;;
5 val it: float = 24.0
6
7 > 4.0 * (3.0 * 2.0);;
8 val it: float = 24.0
9
10 > 4.0 ** 3.0 ** 2.0;;
11 val it: float = 262144.0
12
13 > (4.0 ** 3.0) ** 2.0;;
14 val it: float = 4096.0
15
16 > 4.0 ** (3.0 ** 2.0);;
17 val it: float = 262144.0

```

association or any other basic semantic rules, it is a good idea to use parentheses. It is also a good idea to test your understanding of the syntax and semantic rules by making a simple script.

### 3.3 Boolean Arithmetic

Boolean arithmetic is the basis of almost all computers and particularly important for controlling program flow, which will be discussed in ???. Boolean values are one of 2 possible values, true or false, which is also sometimes written as 1 and 0. Basic operations on Boolean values are 'and', 'or', and 'not', which in F# are written respectively as the binary operators &&, ||, and the function not. Since the domain of Boolean values is so small, all possible combination of input on these values can be written on the tabular form, known as a *truth table*, and the truth tables for the basic Boolean operators and functions are shown in Table 3.8. A good mnemonic

a	b	a && b	a    b	not a
false	false	false	false	true
false	true	false	true	true
true	false	false	true	false
true	true	true	true	false

**Table 3.8** Truth table for boolean 'and', 'or', and 'not' operators. Value 0 is false and 1 is true.

for remembering the result of the 'and' and 'or' operators is to use 1 for true, 0 for false, multiplication for the Boolean 'and' operator, and addition for the Boolean 'or' operator, e.g., true and false in this mnemonic translates to  $1 \cdot 0 = 0$ , and the result

translates back to the Boolean value `false`. In F#, the truth table for the basic Boolean operators can be produced by a program, as shown in Listing 3.13. Here, we used

**Listing 3.13: Boolean operators and truth tables.**

```

1
2 > printfn "a b a*b a+b not a"
3 printfn "%A %A %A %A %A"
4     false false (false && false) (false || false) (not false)
5 printfn "%A %A %A %A %A"
6     false true (false && true) (false || true) (not false)
7 printfn "%A %A %A %A %A"
8     true false (true && false) (true || false) (not true)
9 printfn "%A %A %A %A %A"
10    true true (true && true) (true || true) (not true);;
11 a b a*b a+b not a
12 false false false false true
13 false true false true true
14 true false false true false
15 true true true true false
16 val it: unit = ()

```

the `printfn` function to present the results of many expressions on something that resembles a tabular form. The spacing produced using the `printfn` function is not elegant, and in ?? we will discuss better options for producing more beautiful output. Notice that the arguments for `printfn` was given on the next line with indentation. The indentation is an important part of telling F# which part of what you write belong together. This is an example of the so-called lightweight syntax. Generally, F# ignores newlines and whitespaces except when using the lightweight syntax. The difference between verbose and lightweight syntax is discussed in Chapter 4.

### 3.4 Integer Arithmetic

The set of integers is infinitely large, but since all computers have limited resources, it is not possible to represent it in its entirety. The various integer types listed in Table 3.1 are finite subsets reduced by limiting their ranges. An in-depth description of integer implementation can be found in Appendix B. The type `int` is the most common type.

Table 3.4–3.6 give examples of operators and functions pre-defined for integer types. Notice that fewer functions are available for integers than for floating point numbers. For most addition, subtraction, multiplication, and negation, the result is straightforward. However, performing arithmetic operations on integers requires extra care, since the result may cause *overflow* and *underflow*. For example, an `sbyte` is specified using the “y”-literal and can hold values `[-128 .. 127]`. This causes problems



in the example in Listing 3.14. Here  $100 + 30 = 130$ , which is larger than the biggest

**Listing 3.14: Adding integers may cause overflow.**

```
1 > 100y;;
2 val it: sbyte = 100y
3
4 > 30y;;
5 val it: sbyte = 30y
6
7 > 100y + 30y;;
8 val it: sbyte = -126y
```

sbyte, and the result is an overflow. Similarly, we get an underflow, when the arithmetic result falls below the smallest value storable in an sbyte, as demonstrated in Listing 3.15. I.e., we were expecting a negative number but got a positive number

**Listing 3.15: Subtracting integers may cause underflow.**

```
1 > -100y - 30y;;
2 val it: sbyte = 126y
```

instead.

The overflow error in Listing 3.14 can be understood in terms of the binary representation of integers: In binary,  $130 = 10000010_2$ , and this binary pattern is interpreted differently as byte and sbyte, see Listing 3.16. That is, for signed bytes, the left-

**Listing 3.16: The leftmost bit is interpreted differently for signed and unsigned integers, which gives rise to potential overflow errors.**

```
1 > 0b10000010uy;;
2 val it: byte = 130uy
3
4 > 0b10000010y;;
5 val it: sbyte = -126y
```

most bit is used to represent the sign, and since the addition of  $100 = 01100100_2$  and  $30 = 00011110_2$  is  $130 = 10000010_2$ , which causes the left-most bit to be used, this is wrongly interpreted as a negative number when stored in an sbyte. Similar arguments can be made explaining underflows.

The `operator` discards the fractional part after division, and the *integer remainder* operator calculates the remainder after integer division, as demonstrated in Listing 3.17. Together, the integer division and remainder can form a lossless representation of the original number, see Listing 3.18. Here we see that integer division of 7 by 3 followed by multiplication by 3 is less than 7, and that the difference is  $7 \% 3$ .

Notice that neither overflow nor underflow error gave rise to an error message, which is why such bugs are difficult to find. Dividing any non-zero number by 0 is infinite, which is also outside the domain of any of the integer types, but in this case, F#

**Listing 3.17: Integer division and remainder operators.**

```

1 > 7 / 3;;
2 val it: int = 2
3
4 > 7 % 3;;
5 val it: int = 1

```

**Listing 3.18: Integer division and remainder is a lossless representation of an integer, compare with Listing 3.17.**

```

1 > (7 / 3) * 3;;
2 val it: int = 6
3
4 > (7 / 3) * 3 + (7 % 3);;
5 val it: int = 7

```

casts an *exception*, as shown in Listing 3.19. The output looks daunting at first sight,

**Listing 3.19: Integer division by zero causes an exception runtime error.**

```

1 > 3/0;;
2 System.DivideByZeroException: Attempted to divide by zero.
3 at <StartupCode$FSI_0002>.$FSI_0002.main@()

```

but the first and last lines of the error message are the most important parts, which tell us what exception was cast and why the program stopped. The middle contains technical details concerning which part of the program caused the error and can be ignored for the time being. Exceptions are a type of *runtime error*, and are discussed in ??

Integer exponentiation is not defined as an operator but is available as the built-in function `pown`. This function is demonstrated in Listing 3.20 for calculating  $2^5$ .

**Listing 3.20: Integer exponent function.**

```

1 > pown 2 5;;
2 val it: int = 32

```

For binary arithmetic on integers, the following operators are available: `<leftExpr> <<< <rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the left while inserting 0's to right; `<leftExpr> >>> <rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the right while inserting 0's to left; `~~~ <expr>` returns a new integer, where all 0 bits are changed to 1 bits and vice-versa; `<expr> &&& <expr>` returns the result of taking the Boolean 'and' operator position-wise; `<expr> ||| <expr>` returns the result of taking the Boolean 'or' operator position-wise; and `<expr> ^^^ <expr>` returns the result of the Boolean 'xor' operator defined by the truth table in Table 3.9.

a	b	a ^^^ b
false	false	false
false	true	true
true	false	true
true	true	false

**Table 3.9** Boolean exclusive or truth table.

### 3.5 Floating Point Arithmetic

Like integers, the set of reals is also infinitely large, hence, floating point types are finite subsets reduced by sampling the space of reals. An in-depth description of floating point implementations can be found in Appendix B. The type `float` is the most common type.

Table 3.4–3.6 give examples of operators and functions pre-defined for floating point types. Note that the remainder operator for floats calculates the remainder after division and discards the fractional part, see Listing 3.21. The remainder for

#### Listing 3.21: Floating point division and remainder operators.

```
1 > 7.0 / 2.5;;
2 val it: float = 2.8
3
4 > 7.0 % 2.5;;
5 val it: float = 2.0
```

floating point numbers can be fractional, but division, discarding fractional part, and the remainder is still a lossless representation of the original number, as demonstrated in Listing 3.22.

#### Listing 3.22: Floating point division, downcasting, and remainder is a lossless representation of a number.

```
1 > float (int (7.0 / 2.5));;
2 val it: float = 2.0
3
4 > (float (int (7.0 / 2.5))) * 2.5;;
5 val it: float = 5.0
6
7 > (float (int (7.0 / 2.5))) * 2.5 + 7.0 % 2.5;;
8 val it: float = 7.0
```

Arithmetic using `float` will not cause over- and underflow problems, since the IEEE 754 standard includes the special numbers  $\pm\infty$  and NaN. As shown in Listing 3.23, no exception is thrown. However, the `float` type has limited precision, since there is only a finite number of numbers that can be stored in a float. E.g., addition and subtraction can give surprising results, as demonstrated in Listing 3.24. That is, addition and subtraction associates to the left, hence the expression is interpreted as  $(357.8 + 0.1) - 357.9$  and we see that we do not get the expected 0. The

**Listing 3.23: Floating point numbers include infinity and Not-a-Number.**

```
1 > 1.0/0.0;;  
2 val it: float = infinity  
3  
4 > 0.0/0.0;;  
5 val it: float = nan
```

**Listing 3.24: Floating point arithmetic has finite precision.**

```
1 > 357.8 + 0.1 - 357.9;;  
2 val it: float = 5.684341886e-14
```

reason is that the calculation is done stepwise, and in the process, the numbers are represented using the imprecise floating point standard. Thus,  $357.8 + 0.1$  is represented as a number close to but not identical to what  $357.9$  is represented as, and thus, when subtracting these two representations, we get a very small nonzero number. Such errors tend to accumulate, and comparing the result of expressions of floating point values should, therefore, be treated with care. Thus, **equivalence of two floating point expressions should only be considered up to sufficient precision, e.g., comparing  $357.8 + 0.1$  and  $357.9$  up to  $1e-10$  precision should be tested as, `abs ((357.8 + 0.1) - 357.9) < 1e-10`.**

### 3.6 Char and String Arithmetic

Addition is the only operator defined for characters. Nevertheless, character arithmetic is often done by casting to an integer. A typical example is the conversion of character case, e.g., to convert the lowercase character 'z' to uppercase. Here, we use the *ASCIIbetical order*, add the difference between any Basic Latin Block letters in upper- and lowercase as integers, and cast back to `char`, see Listing 3.25. I.e., the

**Listing 3.25: Converting case by casting and integer arithmetic.**

```
1 > char (int 'z' - int 'a' + int 'A');;  
2 val it: char = 'Z'
```

code point difference between upper and lower case for any alphabetical character 'a' to 'z' is constant, hence we can change case by adding or subtracting the difference between any corresponding character. Unfortunately, this does not generalize to characters from other languages.

A large collection of operators and functions exist for `string`. The simplest is concatenation using the “+” operator, as demonstrated in Listing 3.26. Characters and strings cannot be concatenated, which is why the above example used the string of a space “ ” instead of the space character ' '. The characters of a string may

**Listing 3.26: Example of string concatenation.**

```

1 > "hello" + " " + "world";
2 val it: string = "hello world"

```

be indexed as using the `[]` notation. This is demonstrated in Listing 3.27. Notice

**Listing 3.27: String indexing using square brackets.**

```

1
2 > "abcdefg"[0];;
3 val it: char = 'a'
4
5 > "abcdefg"[3];;
6 val it: char = 'd'
7
8 > "abcdefg"[3..];;
9 val it: string = "defg"
10
11 > "abcdefg"[..3];;
12 val it: string = "abcd"
13
14 > "abcdefg"[1..3];;
15 val it: string = "bcd"
16
17 > "abcdefg"[*];;
18 val it: string = "abcdefg"

```

that the first character has index 0, and to get the last character in a string, we use the string's `length` property. This is done as shown in Listing 3.28. Since index

**Listing 3.28: String length attribute and string indexing.**

```

1 > "abcdefg".Length;;
2 val it: int = 7
3
4 > "abcdefg"[7-1];;
5 val it: char = 'g'

```

counting starts at 0, and since the string length is 7, the index of the last character is 6. There is a long list of built-in functions in `System.String` for working with strings, some of which will be discussed in ??.

The *dot notation* is an example of Structured programming, where technically speaking, the string `"abcdefg"` is an immutable *object* of *class* `string`, `[]` is an object *method*, and `Length` is a property. For more on objects, classes, and methods, see ??.

Strings are compared letter by letter. For two strings to be equal, they must have the same length and all the letters must be identical. E.g., `"abs" = "absalon"` is false, while `"abs" = "abs"` is true. The `"<>"` operator is the boolean negation of the `"="` operator, e.g., `"abs" <> "absalon"` is true, while `"abs" <> "abs"` is false.


For the “<”, “<=”, “>”, and “>=” operators, the strings are ordered alphabetically, such that `"abs" < "absalon" && "absalon" < "milk"` is true, that is, the “<” operator on two strings is true if the left operand should come before the right when sorting alphabetically. The algorithm for deciding the boolean value of `leftOp < rightOp` is as follows: we start by examining the first character, and if `leftOp[0]` and `rightOp[0]` are different, then `leftOp < rightOp` is equal to `leftOp[0] < rightOp[0]`. E.g., `"milk" < "abs"` is the same as `'m' < 'a'`, which is false, since the letter 'm' does not come before the letter 'a' in the alphabet, or more precisely, the codepoint of 'm' is not less than the codepoint of 'a'. If `leftOp[0]` and `rightOp[0]` are equal, then we move on to the next letter and repeat the investigation, e.g., `"abe" < "abs"` is true, since `"ab" = "ab"` is true and `'e' < 's'` is true. If we reach the end of either of the two strings, then the shorter word is smaller than the longer word, e.g., `"abs" < "absalon"` is true, while `"abs" < "abs"` is false. The “<=”, “>”, and “>=” operators are defined in a similar manner.

### 3.7 Programming Intermezzo: Hand Conversion Between Decimal and Binary Numbers

Conversion of integers between decimal and binary form is a key concept one must grasp in order to understand some of the basic properties of calculations on the computer. Converting from binary to decimal is straightforward if using the power-of-two algorithm, i.e., given a sequence of  $n + 1$  binary digits  $b_i$  written as  $b_n b_{n-1} \dots b_0$ , and where  $b_n$  and  $b_0$  are the most and least significant bits respectively, then the decimal value is calculated as,

$$v = \sum_{i=0}^n b_i 2^i \quad (3.2)$$

For example,  $10011_2 = 1 + 2 + 16 = 19$ . Converting from decimal to binary is a little more complex, but a simple divide-by-two algorithm exists. The key to understanding the divide-by-two algorithm is to realize that dividing a number by two is equivalent to shifting its binary representation one position to the right. E.g.,  $10 = 1010_2$  and  $10/2 = 5 = 101_2$ . Odd numbers have  $b_0 = 1$ , e.g.,  $11_{10} = 1011_2$  and  $11_{10}/2 = 5.5 = 101.1_2$ . Hence, if we divide any number by two and get a non-integer number, then its least significant bit was 1. Another way to express this is to say that the least significant bit is the remainder after integer division by two. Sequential application of this idea leads directly to the divide-by-two algorithm. E.g., if we were to convert the number  $11_{10}$  in decimal form to binary form, we would perform the following steps:



Here we used `div` and `rem` to signify the integer division and remainder operators. The algorithm stops when the result of integer division is zero. Reading off the remainder from below and up, we find the sequence  $1011_2$ , which is the binary form of the decimal number  $11_{10}$ . Using the interactive mode, we can perform the same calculation, as shown in Listing 3.29.

**Listing 3.29: Converting the number  $11_{10}$  to binary form.**

```

> printfn "(%d, %d)" (11 / 2) (11 % 2);;
(5, 1)
val it : unit = ()
> printfn "(%d, %d)" (5 / 2) (5 % 2);;
(2, 1)
val it : unit = ()
> printfn "(%d, %d)" (2 / 2) (2 % 2);;
(1, 0)
val it : unit = ()
> printfn "(%d, %d)" (1 / 2) (1 % 2);;
(0, 1)
val it : unit = ()

```

Thus, by reading the second integer-response from `printfn` from below and up, we again obtain the binary form of  $11_{10}$  to be  $1011_2$ . For integers with a fractional part, the divide-by-two algorithm may be used on the whole part, while multiply-by-two may be used in a similar manner on the fractional part.

### 3.8 Key concepts and terms in this chapter

In this chapter you have learned about:

- the **basic types**: **int** of various kinds which are all a subset of integers, **float** of various kinds which are all subsets of reals, **bool** which captures the notion of true and false, and **char** and **string** which holds characters and sequences of characters;
- how to write constants of the basic types, which are called **literals**;
- **operators** such as “+” and “-” and whose arguments are called **operands**;
- **unary** and **binary** operators;
- **escape sequences** for characters and strings;
- how to get the **length** of a string using the **dot** notation, and how to extract substrings using the **slice** notation.



## Chapter 4

# Values, Functions, and Statements

**Abstract** In the previous chapter, you got an introduction to the fundamental concepts of booleans, numbers, characters, and strings, how to perform calculations with them, and some of the limitations they have on the computer. This allows us to perform simple calculations as if the computer was an advanced pocket calculator. In this chapter, we will look at how we can:

- Make assign values to names to make programs that are easier to read and write.
- Organise lines of code in functions, to make the same line reusable, and to make programs shorter, and easier to understand.
- Use operators as functions.
- Conditionally execute code.
- How we can simulate the computer's execution of code by tracing-by-hand, to improve our understanding of, how it interprets the code, we write, and to find errors.

Examples of problems, you can solve, after having read this chapter is:

- Write a function that given the parameters  $a$  and  $b$  in  $f(x) = ax + b$ , find the value of  $x$  when  $f(x) = 0$ .
-

In this chapter, we will see how we can bind expressions to identifiers either as new constants, functions, or operators, how this saves time when building large programs, and how this makes programs easier to read and debug. As an example, consider the following problem,

#### Problem 4.1

or given set constants  $a$ ,  $b$ , and  $c$ , solve for  $x$  in

$$ax^2 + bx + c = 0 \quad (4.1)$$

To solve for  $x$  we use the quadratic formula from elementary algebra,

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}, \quad (4.2)$$

which gives the general solution for any values of the coefficients. Here, we will assume a positive discriminant,  $b^2 - 4ac > 0$ . In order to write a program where the code may be reused later, we define a function

```
discriminant : float -> float -> float -> float
```

that is, a function that takes 3 arguments,  $a$ ,  $b$ , and  $c$ , and calculates the discriminant. Likewise, we will define

```
positiveSolution : float -> float -> float -> float
```

and

```
negativeSolution : float -> float -> float -> float
```

that also take the polynomial's coefficients as arguments and calculate the solution corresponding to choosing the positive and negative sign for  $\pm$  in the equation. Details on function definition is given in Section 4.2. Our solution thus looks like Listing 4.1. Here, we have further defined names of values  $a$ ,  $b$ , and  $c$  which are used as inputs to our functions, and the results of function application are bound to the names  $d$ ,  $xn$ , and  $xp$ . The names of functions and values given here are examples of identifiers, and with these, we may reuse the quadratic formulas and calculated values later, while avoiding possible typing mistakes and reducing the amount of code which needs to be debugged.

The use of identifiers is central in programming. For F#, not to be confused with built-in functionality, identifiers must follow a specific set of rules:

Identifier

**Listing 4.1** identifiersExample.fsx:

Finding roots for quadratic equations using function name binding.

```

1 let discriminant a b c = b ** 2.0 - 4.0 * a * c
2 let positiveSolution a b c = (-b + sqrt (discriminant a b c))
  / (2.0 * a)
3 let negativeSolution a b c = (-b - sqrt (discriminant a b c))
  / (2.0 * a)
4
5 let a = 1.0
6 let b = 0.0
7 let c = -1.0
8 let d = discriminant a b c
9 let xp = positiveSolution a b c
10 let xn = negativeSolution a b c
11 do printfn "%0 = %A * x ** 2.0 + %A * x + %A" a b c
12 do printfn "    has discriminant %A and solutions %A and %A" d
   xn xp

```

---

```

1 $ dotnet fsi identifiersExample.fsx
2 0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
3 has discriminant 4.0 and solutions -1.0 and 1.0

```

- Identifiers are used as names for values, functions, types etc.
- They must start with a Unicode letter or underscore '_', but can be followed by zero or more of letters, digits, and a range of special characters except for SP, LF, and CR (space, line feed, and carriage return). See Appendix C.3 for more on codepoints that represents letters.
- They can also be a sequence of identifiers separated by a period.
- They cannot be keywords, see Table 4.1.

Examples of identifiers are: `a`, `theCharacter9`, `Next_Word`, `_tok`, and `f.sharp.rocks`. Since programmers often work in multilingual environment dominated by the English language it is advisable to **restrict identifiers to use letters from the English alphabet, numbers, period, and '_'**. However, the number of possible identifiers is enormous. The full definition refers to the Unicode general categories described in Appendix C.3, and there are currently 19,345 possible Unicode code points in the letter category and 2,245 possible Unicode code points in the special character category.

Identifiers may be used to carry information about their intended content and use, and careful selection of identifiers can aid programmers to communicate thoughts about the code. Thus, identifiers are often a word or several concatenated words conveying some relevant meaning. For example, in the function definition `let discriminant a b c = b ** 2.0 - 4.0 * a * c`, the function identifier has

Type	Keyword
Regular	<code>abstract</code> , <code>and</code> , <code>as</code> , <code>assert</code> , <code>base</code> , <code>begin</code> , <code>class</code> , <code>default</code> , <code>delegate</code> , <code>do</code> , <code>done</code> , <code>downcast</code> , <code>downto</code> , <code>elif</code> , <code>else</code> , <code>end</code> , <code>exception</code> , <code>extern</code> , <code>false</code> , <code>finally</code> , <code>for</code> , <code>fun</code> , <code>function</code> , <code>global</code> , <code>if</code> , <code>in</code> , <code>inherit</code> , <code>inline</code> , <code>interface</code> , <code>internal</code> , <code>lazy</code> , <code>let</code> , <code>match</code> , <code>member</code> , <code>module</code> , <code>mutable</code> , <code>namespace</code> , <code>new</code> , <code>null</code> , <code>of</code> , <code>open</code> , <code>or</code> , <code>override</code> , <code>private</code> , <code>public</code> , <code>rec</code> , <code>return</code> , <code>sig</code> , <code>static</code> , <code>struct</code> , <code>then</code> , <code>to</code> , <code>true</code> , <code>try</code> , <code>type</code> , <code>upcast</code> , <code>use</code> , <code>val</code> , <code>void</code> , <code>when</code> , <code>while</code> , <code>with</code> , and <code>yield</code> .
Reserved	<code>atomic</code> , <code>break</code> , <code>checked</code> , <code>component</code> , <code>const</code> , <code>constraint</code> , <code>constructor</code> , <code>continue</code> , <code>eager</code> , <code>fixed</code> , <code>fori</code> , <code>functor</code> , <code>include</code> , <code>measure</code> , <code>method</code> , <code>mixin</code> , <code>object</code> , <code>parallel</code> , <code>params</code> , <code>process</code> , <code>protected</code> , <code>pure</code> , <code>recursive</code> , <code>sealed</code> , <code>tailcall</code> , <code>trait</code> , <code>virtual</code> , and <code>volatile</code> .
Symbolic	<code>let!</code> , <code>use!</code> , <code>do!</code> , <code>yield!</code> , <code>return!</code> , <code> </code> , <code>-&gt;</code> , <code>&lt;-</code> , <code>..</code> , <code>:</code> , <code>(</code> , <code>)</code> , <code>[</code> , <code>]</code> , <code>[&lt;</code> , <code>&gt;]</code> , <code>[ </code> , <code> ]</code> , <code>{</code> , <code>}</code> , <code>'</code> , <code>#</code> , <code>:?&gt;</code> , <code>:?</code> , <code>:&gt;</code> , <code>...</code> , <code>::</code> , <code>:=</code> , <code>;;</code> , <code>;</code> , <code>:=</code> , <code>?</code> , <code>??</code> , <code>(*)</code> , <code>&lt;@</code> , <code>@&gt;</code> , <code>&lt;@@</code> , and <code>@&gt;</code> .
Reserved symbolic	<code>~</code> and <code>`</code>

**Table 4.1** Table of (possibly future) *keywords* and symbolic keywords in F#.

been chosen to be discriminant. F# places no special significance to the word 'discriminant', and the program would work exactly the same had the function been called `let f a b c = b ** 2.0 - 4.0 * a * c`. However, to programmers, the word 'discriminant' informs us of the intended role of the function and thus

- ★ is much preferred. This is a general principle: **identifier names should be chosen to reflect their semantic value.** The arguments `a`, `b`, and `c` are short, but adheres to a textbook tradition of elementary algebra. Again, we might as well have used, `let discriminant c a b = a ** 2.0 - 4.0 * c * b`, which is semantically identical to the original expression, but due to tradition, this would confuse most readers of the code. Thus, **identifier names should be chosen consistently with the readers' traditions.** Finally, identifiers are often concatenations of words, as `positiveSolution` in Listing 4.1. Concatenations can be difficult to read. Without the capitalization of the second word, we would have had `positivesolution`. This is readable at most times, but takes longer time to understand in general. Typical solutions are to use a separator, such as `positive_solution`, *lower camel case* also known as *mixed case* as in the example `positiveSolution`, and *upper camel case* also known as *pascal case* as `PositiveSolution`. In this book, we use lower camel case except where F# requires a capital first letter. Again, the choice does not influence what a program does, only how readable it is to a fellow programmer.
- ★ The important part is that **identifier names consisting of concatenated words are often preferred over names with few character, and concatenation should be emphasized, e.g., by camel casing.** Choosing the length of identifier names is a balancing act, since when working with large programs, very long identifier names can be tiresome to write, and a common practice is that the length of identifier names is proportional to the complexity of the program. I.e., complex programs use long names, simple programs use short names. What is complex and what is simple is naturally in the eye of the beholder, but when you program, remember that a future reader of the program most likely has not had time to work with the problem as long

as the programmer, thus **choose identifier names as if you were to explain the meaning of a program to a knowledgeable outsider.** ★

Another key concept in F# is expressions. An expression can be a mathematical expression, such as  $3 * 5$ , a function application, such as  $f3$ , and many other things. Central in this chapter is the binding of values and functions to identifiers, which is done with the keyword `let`, e.g., `let a = 1.0`.

Expressions are the main workhorse of F# and have an enormous variety in how they may be written. We will in this book gradually work through some of the more important facets.

### Expressions

- An Expression is a computation such as  $3 * 5$ .
- They can be value bindings between identifiers and expressions that evaluate to a value or a function, see Sections 4.1 and 4.2.
- They can be `do`-bindings that produce side-effects and whose result are ignored, see Section 4.2
- They can be assignments to variables, see Section 4.1.
- They can be a sequence of expressions separated with the “;” lexeme.
- They can be annotated with a type by using the “:” lexeme.

Before we begin a deeper discussion on bindings, note that F# adheres to two different syntaxes: *verbose* and *lightweight*. In the verbose syntax, newlines and whitespaces are generally ignored, while in lightweight syntax, certain keywords and lexemes may be replaced by newlines and whitespaces. The lightweight syntax is the most common and will be used in this book.

## 4.1 Value Bindings

Binding identifiers to literals, or expressions that are evaluated to be values, is called *value-binding*, and examples are `let a = 3.0` and `let b = cos 0.9`. Value bindings have the following syntax:

**Listing 4.2: Value binding expression.**

```
1 let <valueIdent> = <bodyExpr>
```

The `let` keyword binds a value-identifier with an expression. The above notation means that `<valueIdent>` is to be replaced with a name and `<bodyExpr>` with an expression that evaluates to a value. The binding *is* available in later lines until the end of the scope it is defined in.

The value identifier annotated with a type by using the “:” lexeme followed by the name of a type, e.g., `int`. The “_” lexeme may be used as a value-identifier. This lexeme is called the *wildcard pattern*, and for value-bindings it means that `<bodyExpr>` is evaluated, but the result is discarded. See ?? for more details on patterns.

For example, letting the identifier `p` be bound to the value `2.0` and using it in an expression is done as shown in Listing 4.3. Note that the expression `3.0 ** p` must

**Listing 4.3 letValueLightWeight.fsx:**

Lightweight syntax does not require the `in` keyword, but the expression must be aligned with the `let` keyword.

```
1 let p = 2.0
2 do printfn "%A" (3.0 ** p)

-----

1 $ dotnet fsi letValueLightWeight.fsx
2 9.0
```

be put in parentheses here, to force F# to *first* calculate the result of the expression and *then* give it to `printfn` to be printed on the screen. The same expression in interactive mode will also show with the inferred types, as shown in Listing 4.4. By

**Listing 4.4: Interactive mode also outputs inferred types.**

```
1 > let p = 2.0
2 do printfn "%A" (3.0 ** p);;
3 9.0
4 val p: float = 2.0
5 val it: unit = ()
```

the `val` keyword in the line `val p : float = 2.0`, we see that `p` is inferred to be of type `float` and bound to the value `2.0`. The inference is based on the type of the right-hand-side which is `float`. Identifiers may be defined to have a type using the “:” lexeme, but the types on the left-hand-side and right-hand-side of the “=” lexeme must be identical. Mixing types gives an error, as shown in Listing 4.5. Here, the left-hand-side is defined to be an identifier of type `float`, while the right-hand-side is a literal of type `integer`.

**Listing 4.5 letValueTypeError.fsx:**  
Binding error due to type mismatch.

```

1 let p = 2.0
2 do printfn "%A" (3 ** p)
-----
1 $ dotnet fsi letValueTypeError.fsx
2
3
4 letValueTypeError.fsx(2,18): error FS0001: The type 'int'
   does not support the operator 'Pow'
```

A key concept of programming is *scope*. When F# seeks the value bound to a name, it looks left and upward in the program text for its `let`-binding in the present or higher scopes. This is called *lexical scope*. Some special bindings are mutable, that is, they can change over time in which case F# uses the *dynamic scope*. This will be discussed in ??.

Scopes are given levels, and scopes may be nested, where the nested scope has a level one lower than its parent. F# distinguishes between the top and lower levels, and at the top level in the lightweight syntax, redefining values is not allowed, as shown in Listing 4.6. However, using parentheses, we create a *code block*, i.e., a

**Listing 4.6 letValueScopeLowerError.fsx:**  
Redefining identifiers is not allowed in lightweight syntax at top level.

```

1 let p = 3
2 let p = 4
3 do printfn "%A" p;
-----
1 $ fsharp --nologo -a letValueScopeLowerError.fsx
2
3 letValueScopeLowerError.fsx(2,5): error FS0037: Duplicate
   definition of value 'p'
```

*nested scope*, as demonstrated in Listing 4.7. In lower scope-levels, redefining is allowed but **avoid reusing names unless it's in a deeper scope**. Inside the block in Listing 4.7 we used indentation, which is good practice, but not required here. ★

Defining blocks is used for controlling the extent of a lexical scope of bindings. For example, bindings inside a nested scope are not available outside, as shown in Listing 4.8. Nesting is a natural part of structuring code, e.g., through function definitions to be discussed in Section 4.2 and flow control structures to be discussed in ??. Blocking code by nesting is a key concept for making robust code that is easy to use by others, without the user necessarily needing to know the details of the inner workings of a block of code.

**Listing 4.7** `letValueScopeBlockAlternative3.fsx`:  
A block may be created using parentheses.

```
1 let p = 3
2 (
3   let p = 4
4   do printfn "%A" p
5 )
```

---

```
1 $ dotnet fsi letValueScopeBlockAlternative3.fsx
2 4
```

**Listing 4.8** `letValueScopeNestedScope.fsx`:  
Bindings inside a scope are not available outside.

```
1 let p = 3
2 (
3   let q = 4
4   do printfn "%A" q
5 )
6 do printfn "%A %A" p q
```


---

```
1 $ fsharp --nologo -a letValueScopeNestedScope.fsx
2
3 letValueScopeNestedScope.fsx(6,22): error FS0039: The value
   or constructor 'q' is not defined.
```



## 4.2 Function Bindings

A function is a mapping between an input and output domain, as illustrated in Figure Figure 4.1. A key advantage of using functions when programming is that



**Fig. 4.1** A function  $f : X \rightarrow Y$  is a mapping between two domains, such that  $f(x) = y$ ,  $x \in X$ ,  $y \in Y$ .

they encapsulate code into smaller units, that may be reused and are easier to find errors in. F# is a functional first programming language and offers a number of alternative methods for specifying parameters, which will be discussed in this section. Binding identifiers to functions follows a syntax similar to value-binding,

### Listing 4.9: Function binding expression

```
1 let <funcIdent> (<arg> {<arg>} | () = <bodyExpr>)
```

Here **<funcIdent>** is an identifier and is the name of the function, **<arg>** is zero or more identifiers, that bind to the value used when calling the function, and which is to be used in the *body* of the function, the expression **<bodyExpr>**. The **|** notation denotes a choice, i.e., either that on the left-hand-side or that on the right-hand-side. Thus **let f x = x * x** and **let f () = 3** are valid function bindings, but **let f = 3** would be a value binding, not a function binding. The arguments and the function may be annotated with a type, in which case for arguments we write

### Listing 4.10: Function binding expression

```
1 let <funcIdent> ((<arg> : <type>) {(<arg> : <type>)} : <type> |  
  () : <type>) = <bodyExpr>
```

where **<type>** is a name of an existing type. The argument types are given in parentheses, and the return type is given last.

Functions are a key concept in F#, and in this chapter we will discuss the very basics. Recursive functions will be discussed in ?? and higher-order functions in ??.

The function's body can be placed either on the same line as the **let**-keyword or on the following lines using indentation. An example of defining a function and using it in interactive mode is shown in Listing 4.11. Here we see that the function is interpreted to have the type **val sum : x:float -> y:float -> float**. The “->” lexeme means a mapping between sets, in this case, floats. The function is also

**Listing 4.11: An example of a binding of an identifier and a function.**

```

1
2 > let sum (x : float) (y : float) : float = x + y
3 let c = sum 357.6 863.4
4 do printfn "%A" c;;
5 1221.0
6 val sum: x: float -> y: float -> float
7 val c: float = 1221.0
8 val it: unit = ()

```

a higher-order function, to be discussed in detail below, and here it suffices to think of `sum` as a function that takes 2 floats as argument and returns a float.

Not all types need to be declared, just a sufficient number for F# to be able to infer the types for the full statement. For the example, one is sufficient, and we could just have declare the type of the result, as in Listing 4.12. Or even just one of the

**Listing 4.12 letFunctionAlternative.fsx:  
Not every type needs to be declared.**

```

1 let sum x y : float = x + y

```

arguments, as in Listing 4.13. In both cases, since the `+` operator is only defined for

**Listing 4.13 letFunctionAlternative2.fsx:  
Just one type is often enough for F# to infer the rest.**

```

1 let sum (x : float) y = x + y

```

*operands* of the same type, declaring the type of either arguments or result implies the type of the remainder.

Arguments need not always be inferred to types, but may be of the generic type when *type safety* is ensured, as shown in Listing 4.14. Here, the function `second`

**Listing 4.14: Type safety implies that a function will work for any type.**

```

1 > let second x y = y
2 let a = second 3 5
3 do printfn "%A" a
4 let b = second "horse" 5.0
5 do printfn "%A" b;;
6 5
7 5.0
8 val second: x: 'a -> y: 'b -> 'b
9 val a: int = 5
10 val b: float = 5.0
11 val it: unit = ()

```

does not use the first argument `x`, which therefore can be of any type, and which F#, therefore, calls `'a`. The type of the second element, `y`, can also be of any type and not necessarily the same as `x`, so it is called `'b`. Finally, the result is the same type

as *y*, whatever it is. This is an example of a *generic function*, since it will work on any type.

A function may contain a sequence of expressions but must return a value. E.g., the quadratic formula may be written as shown in Listing 4.15. Here, we use a

**Listing 4.15** `identifiersExampleAdvance.fsx`:  
A function may contain sequences of expressions.

```

1 let solution a b c sgn =
2     let discriminant a b c =
3         b ** 2.0 - 4.0 * a * c
4     let d = discriminant a b c
5     (-b + sgn * sqrt d) / (2.0 * a)
6
7 let a = 1.0
8 let b = 0.0
9 let c = -1.0
10 let xp = solution a b c +1.0
11 let xn = solution a b c -1.0
12 do printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
13 do printfn "  has solutions %A and %A" xn xp

```

---

```

1 $ dotnet fsi identifiersExampleAdvance.fsx
2 0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
3   has solutions -1.0 and 1.0

```

set of nested scopes. In the scope of the function `solution` there are further two functions defined, each with their own scope, which F# identifies by the level of indentation. The amount of space used for indentation does not matter, but all lines in the same scope must use the same amount. The scope ends before the first line with the previous indentation or none. Notice how the last expression is not bound to an identifier, but is the result of the function, i.e., in contrast to many other languages, F# does not have an explicit keyword for returning values, but requires a final expression, which will be returned to the caller of the function. Note also that since the function `discriminant` is defined in the nested scope of `solution`, and because the scope ends before `let a = 1.0`, the function `discriminant` cannot be called outside `solution`.

*Lexical scope* and function definitions can be a cause of confusion, as the following example in Listing 4.16 shows. Here, the value-binding for `a` is redefined after it has been used to define a helper function `f`. So which value of `a` is used when we later apply `f` to an argument? To resolve the confusion, remember that value-binding is lexically defined, i.e., the binding `let f z = a * z` uses the value of `a` as it is defined by the ordering of the lines in the script, not dynamically by when `f` was called. Hence, **think of lexical scope as substitution of an identifier with its value or function immediately at the place of definition.** Since `a` and `3.0` are ★

**Listing 4.16 lexicalScopeNFunction.fsx:**Lexical scope means that  $f(z) = 3x$  and not  $4x$  at the time of calling.

```

1 let testScope x =
2   let a = 3.0
3   let f z = a * z
4   let a = 4.0
5   f x
6 do printfn "%A" (testScope 2.0)

```

---

```

1 $ dotnet fsi lexicalScopeNFunction.fsx
2 6.0

```

synonymous in the first lines of the program, the function `f` is really defined as `let f z = 3.0 * z`.

Functions do not need a name, but may be declared as an *anonymous function* using the `fun` keyword and the “`->`” lexeme, as shown in Listing 4.17. Here, a name

**Listing 4.17 functionDeclarationAnonymous.fsx:**

Anonymous functions are functions as values.

```

1 let first = fun x y -> x
2 do printfn "%d" (first 5 3)

```

---

```

1 $ dotnet fsi functionDeclarationAnonymous.fsx
2 5

```

is bound to an anonymous function which returns the first of two arguments. The difference to `let first x y = x` is that anonymous functions may be treated as values, meaning that they may be used as arguments to other functions and the new values may be reassigned to their identifiers when mutable, as will be discussed in ???. A common use of anonymous functions is as arguments to other functions, as demonstrated in Listing 4.18. Note that here `apply` is given 3 arguments: the

**Listing 4.18 functionDeclarationAnonymousAdvanced.fsx:**

Anonymous functions are often used as arguments for other functions.

```

1 let apply f x y = f x y
2 let mul = fun a b -> a * b
3 do printfn "%d" (apply mul 3 6)

```

---

```

1 $ dotnet fsi functionDeclarationAnonymousAdvanced.fsx
2 18

```

function `mul` and 2 integers. It is not given the result of `mul 3 6`, since that would not match the definition of `apply`. **Anonymous functions and functions as arguments**

**are powerful concepts, but tend to make programs harder to read, and their use should be limited.**

The result of one function is often used as an argument of another. This is function composition, and an example is shown in Listing 4.19. In the example we combine

**Listing 4.19 functionComposition.fsx:**  
Composing functions using intermediate bindings.

```
1 let f x = x + 1
2 let g x = x * x
3
4 let a = f 2
5 let b = g a
6 let c = g (f 2)
7 do printfn "a = %A, b = %A, c = %A" a b c
```

---

```
1 $ dotnet fsi functionComposition.fsx
2 a = 3, b = 9, c = 9
```

two functions *f* and *g* by storing the result of *f* 2 in *a* and using that as argument of *g*. This is the same as *g* (*f* 2), and in the later case, the compiler creates a temporary value for *f* 2. Such compositions are so common in F# that a special set of operators has been invented, called the *pipng* operators: “|>” and “<|”. They are used as demonstrated in Listing 4.20. The example shows regular composition,

**Listing 4.20 functionPiping.fsx:**  
Composing functions by piping.

```
1 let f x = x + 1
2 let g x = x * x
3
4 let a = g (f 2)
5 let b = 2 |> f |> g
6 let c = g <| (f <| 2)
7 do printfn "a = %A, b = %A, c = %A" a b c
```

---

```
1 $ dotnet fsi functionPiping.fsx
2 a = 9, b = 9, c = 9
```

left-to-right, and right-to-left piping. The word piping is a pictorial description of data as if it were flowing through pipes, where functions are connection points of pipes distributing data in a network. The three expressions in Listing 4.20 perform the same calculation. The left-to-right piping in line 5 corresponds to the left-to-right reading direction, i.e., the value 2 is used as argument to *f*, and the result is used as argument to *g*. In contrast, right-to-left piping in line 6 has the order of arithmetic composition as line 4. Unfortunately, since the piping operators are left-associative, without the parenthesis in line 6 *g* <| *f* <| 2, F# would read the expression as (*g*

<| f) <| 2. That would have been an error, since `g` takes an integer as argument, not a function. F# can also define composition on a function level. Further discussion on this is deferred to ??.

A *procedure* is a generalization of the concept of functions, and in contrast to functions, procedures need not return values. This is demonstrated in Listing 4.21. In F#, this is automatically given the unit type as the return value. Procedural

#### Listing 4.21 procedure.fsx:

A procedure is a function that has no return value, and in F# returns “()”.

```
1 let printIt a = printfn "This is '%A'" a
2 do printIt 3
3 do printIt 3.0

-----

1 $ dotnet fsi procedure.fsx
2 This is '3'
3 This is '3.0'
```

thinking is useful for *encapsulation* of scripts, but is prone to *side-effects*. For this reason, it is advised to **prefer functions over procedures**. More on side-effects in ??.

In F#, functions (and procedures) are *first-class citizens*, which means that functions are values: They may be passed as arguments, returned from a function, and bound to a name. For first-class citizens, the name it is bound to does not carry significance to the language, as, e.g., illustrated with the use of anonymous functions. Technically, a function is stored as a *closure*. A closure is a description of the function, its arguments, its expression, and the environment at the time it was created, i.e., the triple (*args*, *exp*, *env*). Consider the listing in Listing 4.22. It defines two functions

#### Listing 4.22 functionFirstClass.fsx:

The function `ApplyFactor` has a non-trivial closure.

```
1 let mul x y = x * y
2 let factor = 2.0
3 let applyFactor fct x =
4     let a = fct factor x
5     string a
6
7 do printfn "%g" (mul 5.0 3.0)
8 do printfn "%s" (applyFactor mul 3.0)

-----

1 $ dotnet fsi functionFirstClass.fsx
2 15
3 6
```

`mul` and `applyFactor`, where the latter is a higher-order function taking another

function as an argument and uses part of the environment to produce its result. The two closures are:

$$\text{mul} : (\text{args}, \text{exp}, \text{env}) = ((x, y), (x * y), ()) \quad (4.3)$$

$$\text{applyFactor} : (\text{args}, \text{exp}, \text{env}) = ((x, \text{fct}), (\text{body}), (\text{factor} \rightarrow 2.0)) \quad (4.4)$$

where lazily write `body` instead of the whole function's body. The function `mul` does not use its environment, and everything needed to evaluate its expression are values for its arguments. The function `applyFactor` also takes two arguments, a function and a value. It uses `factor` from the environment, thus this is stored in its closure. When `mul` is given as an argument in Listing 4.22 line 8, then it is its closure which is given to `applyFactor`, and the closure contains everything that `applyFactor` requires to use `mul`. Likewise, if `applyFactor` is given as argument to yet another function, then its closure includes the relevant part of its environment at the time of definition, `factor`, such that when `applyFactor` is applied to two arguments, then its closure contains everything needed to evaluate its expression.

### 4.3 Do-Bindings

Aside from `let`-bindings that binds names with values or functions, sometimes we just need to execute code. This is called a `do`-binding or, alternatively, a *statement*. The syntax is as follows:

Listing 4.23: Syntax for `do`-bindings.

```
1 [do ]<expr>
```

The expression `<expr>` must return `unit`. The keyword `do` is optional in most cases, but using it emphasizes that the expression is not a function that returns a useful value. Procedures are examples of such expressions, and a very useful family of procedures are the `printf` family described below. In the remainder of this book, we will refrain from using the `do` keyword.

### 4.4 Conditional Expressions

Programs often contain code which should only be executed under certain conditions. The `match – with` expressions allows us to write such expressions, and its syntax is as follows:

Listing 4.24: Syntax for `match`-expressions.

```

1 match <inputExpr> with
2   [| ]<pat> [when <guardExpr>] -> <caseExpr>
3   | <pat> [when <guardExpr>] -> <caseExpr>
4   | <pat> [when <guardExpr>] -> <caseExpr>
5   ...

```

where `<inputExpr>` is the *input pattern* to find matches of, `<pat>` is a pattern to match with, `<guardExpr>` is an optional guard expression, and `<caseExpr>` is the resulting expression. Each set starting with `<pat>` is called a case. The indentation must be equal to or higher than the indentation of `match`. All cases must return a value of the same type, and F# reports an error when the complete domain of the input pattern is not covered by cases in `match`-expressions.

For example, a calendar program may need to highlight weekends, and thus highlighting code should be called, if and only if the day of the week is either Saturday or Sunday, as shown in the following. Remember that `match – with` is the last

Listing 4.25 `weekdayPatternSimple.fsx`:  
Using `match – with` to print discriminated unions.

```

1 let whatToDoToday (d : int) : string =
2     match d with
3     | 0 -> "go to work"
4     | 1 -> "go to work"
5     | 2 -> "go to work"
6     | 3 -> "go to work"
7     | 4 -> "go to work"
8     | 5 -> "take time off"
9     | 6 -> "take time off"
10    | _ -> "unknown day of the week"
11
12 let dayOfWeek = 3
13 let task = whatToDoToday dayOfWeek
14 printfn "Day %A of the week you should %A" dayOfWeek task

```

---

```

1 $ dotnet fsi weekdayPatternSimple.fsx
2 Day 3 of the week you should "go to work"

```

expression in the function `whatToDoToday`, hence the resulting string of `match – with` is also the resulting string of the function. F# examines each case from top to bottom, and as soon as a matching case is found, then the code following the “->”-symbol is executed. This code is called the case’s *branch*. The “_” pattern is a *wildcard* and matches everything.

In the above code, each day has its own return string associated to it, which is great, if we want to return different messages for different days. However, in this code, we only return two different strings, and if, e.g., choose the one for weekdays, then we



need to make 5 identical changes to the code. This is wasteful and risks introducing new errors. So instead, we may compact this code by concatenating the “|”s as shown in Listing 4.26. It still seems unnecessarily clumsy to have to mention

**Listing 4.26 weekdayPatternSum.fsx:**  
Concatenating cases. Compare with Listing 4.25.

```
1 let whatToDoToday (d : int) : string =
2     match d with
3         0 | 1 | 2 | 3 | 4 -> "go to work"
4         | 5 | 6 -> "take time off"
5         | _ -> "unknown day of the week"
6
7 let dayOfWeek = 3
8 let task = whatToDoToday dayOfWeek
9 printfn "Day %A of the week you should %A" dayOfWeek task
```

---

```
1 $ dotnet fsi weekdayPatternSum.fsx
2 Day 3 of the week you should "go to work"
```

the number of each weekday explicitly, and to avoid this, we may use *guards* as illustrated in Listing 4.27. In the guard expression, we match *d* with the name *n*

**Listing 4.27 weekdayPatternGuard.fsx:**  
Using guards. Compare with Listing 4.26.

```
1 let whatToDoToday (d : int) : string =
2     match d with
3         n when 0 <= n && n < 5 -> "go to work"
4         | n when 5 <= n && n <= 6 -> "take time off"
5         | _ -> "unknown day of the week"
6
7 let dayOfWeek = 3
8 let task = whatToDoToday dayOfWeek
9 printfn "Day %A of the week you should %A" dayOfWeek task
```

---

```
1 $ dotnet fsi weekdayPatternGuard.fsx
2 Day 3 of the week you should "go to work"
```

as long as the condition is fulfilled, i.e., as long  $0 \leq n < 5$  or as it is written in F#, `0 <= n && n < 5`. The logical and-operator (“&&”) is needed in F#, since both `0 <= n` and `n < 5` needs to be true, and F# can only evaluate them individually.

Note that `match – with` expressions will give a warning, if there are not cases for the full domain of what is being match. I.e., if we don’t end with the wildcard pattern, when matching for integers, we would be missing cases for all other patterns

An alternative to `match – with` is the `if – then` expressions.

**Listing 4.28** weekdayPatternError.fsx:  
Using `match` – `with` to print discriminated unions.

```

1 let whatToDoToday (d : int) : string =
2     match d with
3     | n when 0 <= n && n < 5 -> "go to work"
4     | n when 5 <= n && n <= 6 -> "take time off"
5
6 let dayOfWeek = 3
7 let task = whatToDoToday dayOfWeek
8 printfn "Day %A of the week you should %A" dayOfWeek task

```

---

```

1 $ dotnet fsi weekdayPatternError.fsx
2
3
4 weekdayPatternError.fsx(2,9): warning FS0025: Incomplete
   pattern matches on this expression.
5
6 Day 3 of the week you should "go to work"

```

**Listing 4.29:** Conditional expressions.

```

1 if <cond> then <expr> {elif <cond> then <expr>} [else <expr>]

```

The condition `<con>` is an expression resulting in a Boolean value, and there can be zero or more `elif` conditions, as indicated by `{}`. Each expression `<expr>` is called a *branch*, and all branches must have the same type, such that regardless of which branch is chosen, the type of the result of the conditional expression is the same. Then the expression of the first if-branch, whose condition is true, is evaluate. If all conditions are false then the `else`-branch is evaluated. If no `else` expression is present, then `"()`" will be returned. See Listing 4.30 for a simple example. The branches are often several lines of code, in which case it is more

**Listing 4.30** weekdayIfThenElse.fsx:  
Conditional computation with `if` – `then`. Compare with Listing 4.27.

```

1 let whatToDoToday (d : int) : string =
2     if 0 <= d && d < 5 then "go to work"
3     elif 5 <= d && d <= 6 then "take time off"
4     else "unknown day of the week"
5
6 let dayOfWeek = 3
7 let task = whatToDoToday dayOfWeek
8 printfn "Day %A of the week you should %A" dayOfWeek task

```

---

```

1 $ dotnet fsi weekdayIfThenElse.fsx
2 Day 3 of the week you should "go to work"

```

useful to write them indented, below the keywords, which may also make the code

easier to read. The identical `whatToDoToday` function with indented branches is shown in Listing 4.31. In contrast to `match – with`, the `if – then` expression is

**Listing 4.31** `weekdayIfThenElseIndentation.fsx`:  
Conditional computation with `if – then`. Compare with Listing 4.27.

```
1 let whatToDoToday (d : int) : string =  
2     if 0 <= d && d < 5 then  
3         "go to work"  
4     elif 5 <= d && d <= 6 then  
5         "take time off"  
6     else  
7         "unknown day of the week"
```

not as thorough in investigating whether cases of the domain in question have been covered. For example, in both

```
let a = if true then 3
```

and

```
let a = if true then 3 elif false then 4
```

the value of `a` are uniquely determined, but F# finds them to be erroneous, since F# looks for the `else` to ensure all cases have been covered, and that the branches have the identical type. Hence,

```
let a = if true then 3 else 4
```

is the only valid expression of the 3. In practice, F# assumes that the omitted branch returns `()`, and thus it is fine to say `let a = if true then ()` and `if true then printfn "hej"`. Nevertheless, it is good practice in F# to always include an `else` branch.

## 4.5 Tracing code by hand

The concept of Tracing by hand, will be developed throughout this book. Here we will concentrate in the basics, and as we introduce more complicated programming structures, we will develop the Tracing by hand accordingly. Tracing may seem tedious in the beginning, but in conjunction with strategically placed debugging `printfn` statements, it is a very valuable tool for debugging.

**Listing 4.32 lexicalScopeTracing.fsx:**  
**Example of lexical scope and closure environment.**

```

1 let testScope x =
2   let a = 3.0
3   let f z = a * z
4   let a = 4.0
5   f x
6 printfn "%A" (testScope 2.0)

```

---

```

1 $ dotnet fsi lexicalScopeTracing.fsx
2 6.0

```

Consider the program in Listing 4.32. The program calls `testScope 2.0`, and by running the program, we see that the return-value is `6.0` and not `8.0`, as we had expected. Hence, we will use tracing to understand the result.

Tracing a program by hand means that we simulate its execution and, as part of that, keep track of the bindings, assignments closures, scopes, and input and output of the program. To do this, we need to consider the concept of *environments*.

Environments describe bindings available to the program at the present scope and at a particular time and place in the code. There is always an outer environment, called  $E_0$ , and each time we call a function or create a scope, we create a new environment. Only one environment can be active at a time, and it is updated as we simulate the execution of code with new bindings and temporary evaluations of expressions. Once a scope is closed, then its environment is deleted and a return-value is transported to its enclosing environment. In tracing, we note return-values explicitly. Likewise, output from, e.g., `printfn` is reported with a special notation.

To trace code, we make a table with 4 columns: Step, Line, Environment, and Bindings and evaluations. The Step column enumerates the steps performed. The Line column contains the program-line treated in a step *where* the present environment is updated. The Environment contains the name of the present environment, and Bindings ... shows *what* in the environment is updated.

The code in Listing 4.32 contains a function definition and a call, hence, the first lines of our table looks like,

Step	Line	Env.	Bindings and evaluations
0	-	$E_0$	()
1	1	$E_0$	testScope = ((x), testScope-body, ())
2	6	$E_0$	testScope 2.0 = ?

The elements of the table is to be understood as follows. Step 0 initializes the outer environment. In order for us to remember that the environment is empty, we write the unit value "`()`". Reading the code from top to bottom, the first nonempty and

non-comment line we meet is line 1, hence, in Step 1, we update the environment with the binding of a function to the name `testScope`. Since functions are values in F#, we note their bindings by their closures: a list of argument names, the function-body, and the values lexically available at the place of binding. See Section 4.2 for more information on closures. Following the function-binding, the `printfn` statement is called in line 6 to print the result `testScope 2.0`. However, before we can produce any output, we must first evaluate `testScope 2.0`. Since we do not yet know what this function evaluates to, in Step 2 we simply write the call with a question mark. The call causes the creation of a new environment, and we continue our table as follows,

Step	Line	Env.	Bindings and evaluations
3	1	$E_1$	$((x = 2.0), \text{testScope-body}, ())$

This means that we are going to execute the code in `testScope-body`. The function was called with 2.0 as argument, causing  $x = 2.0$ . Hence, the only binding available at the start of this environment is to the name `x`. In the `testScope-body`, we make 3 further bindings and a function call. First to `a`, then to `f`, then to another `a`, which will overshadow the previous binding, and finally we call `f`. Thus, our table is updated as follows,

Step	Line	Env.	Bindings and evaluations
4	2	$E_1$	$a = 3.0$
5	3	$E_1$	$f = ((z), a * z, (a = 3.0, x = 2.0))$
6	4	$E_1$	$a = 4.0$
7	5	$E_1$	$f\ x = ?$

Note that by lexical scope, the closure of `f` includes everything above its binding in  $E_1$ , and therefore we add  $a = 3.0$  and  $x = 2.0$  to the environment element in its closure. This has consequences for the following call to `f` in line 5, which creates a new environment based on `f`'s closure and the value of its arguments. The value of `x` in Step 7 is found by looking in the previous steps for the last binding to the name `x` in  $E_1$ , which occurs in Step 3. Note that the binding to a name `x` in Step 5 is an internal binding in the closure of `f` and is irrelevant here. Hence, we continue the table as,

Step	Line	Env.	Bindings and evaluations
8	3	$E_2$	$((z = 2.0), a * z, (a = 3.0, x = 2.0))$

Executing the body of `f`, we initially have 3 bindings available:  $z = 2.0$ ,  $a = 3.0$ , and  $x = 2.0$ . Thus, to evaluate the expression  $a * z$ , we use these bindings and write,

Step	Line	Env.	Bindings and evaluations
9	3	$E_2$	$a * z = 6.0$
10	3	$E_2$	$\text{return} = 6.0$

The 'return'-word is used to remind us that this is the value to replace the question mark with in Step 7. Here we will make a mental note and not physically replace the question mark with the calculated value. If you are ever in doubt which call is connected with which return value, seek upwards in the table from the return statement for the first question mark. Now we delete  $E_2$  and return to the enclosing environment,  $E_1$ . Here the function call was the last expression, hence the return-value from `testScope` will be equal to the return-value from `f`, and we write,

Step	Line	Env.	Bindings and evaluations
11	3	$E_1$	return = 6.0

Similarly, we delete  $E_1$  and return to the question mark in Step 2, which is replaced by the value 6.0. We can now finish the `println` statement and produce the output,

Step	Line	Env.	Bindings and evaluations
12	6	$E_0$	output = "6.0\n"

The return-value of a `println` statement is `()`, and since this line is the last of our program, we return `()` and end the program:

Step	Line	Env.	Bindings and evaluations
13	6	$E_0$	return = <code>()</code>

The full table is shown for completeness in Table 4.2. Hence, we conclude that the

Step	Line	Env.	Bindings and evaluations
0	-	$E_0$	<code>()</code>
1	1	$E_0$	<code>testScope = ((x), testScope-body, ())</code>
2	6	$E_0$	<code>testScope 2.0 = ?</code>
3	1	$E_1$	<code>((x = 2.0), testScope-body, ())</code>
4	2	$E_1$	<code>a = 3.0</code>
5	3	$E_1$	<code>f = ((z), a * z, (a = 3.0, x = 2.0))</code>
6	4	$E_1$	<code>a = 4.0</code>
7	5	$E_1$	<code>f x = ?</code>
8	3	$E_2$	<code>((z = 2.0), a * z, (a = 3.0, x = 2.0))</code>
9	3	$E_2$	<code>a * z = 6.0</code>
10	3	$E_2$	return = 6.0
11	3	$E_1$	return = 6.0
12	6	$E_0$	output = "6.0\n"
13	6	$E_0$	return = <code>()</code>

**Table 4.2** The complete table produced while tracing the program in Listing 4.32 by hand.

program outputs the value 6.0, since the function `f` uses the first binding of `a = 3.0`, and this is because the binding of `f` to the expression `a * z` creates a closure with a lexical scope. Thus, in spite that there is an overshadowing value of `a`, when `f` is called, this binding is ignored in the body of `f`. To correct this, we update the code as shown in Listing 4.33.

**Listing 4.33 lexicalScopeTracingCorrected.fsx:**

Tracing the code in Listing 4.32 by hand produced the table in Table 4.2, and to get the desired output, we correct the code as shown here.

```
1 let testScope x =  
2   let a = 4.0  
3   let f z = a * z  
4   f x  
5 printfn "%A" (testScope 2.0)  
  
1 $ dotnet fsi lexicalScopeTracingCorrected.fsx  
2 8.0
```

## 4.6 Key concepts and terms in this chapter

In this chapter, we have taken the first look at organizing code. Key concepts have been:

- **Binding** values and functions using the **let** statements.
- Function **closures**, which are the values bound in **let** statements of functions.
- Certain names are forbidden in particular **keywords**, which are reserved for other uses.
- In contrast to **let** expressions, **do** statements do something, such as **printfn** prints to the screen. They return the “()” value.
- Conditional expressions using **match-with** and some **patterns** and **guards**. Alternatively, **if-then** expressions can be used to a similar effect.
- **Trace by hand** as a method for simulating execution and in particular keeping track of the **environments** defined by the code structure.





## Chapter 5

# Programming with Types

**Abstract** In the previous chapter, we took the first step in organizing code such that it better reflects the solutions we seek, is reusable, and is easier to find possible errors in. A fundamental structure in all of this is types, which much like sets in mathematics, form the basic building blocks of what we express in code. In this chapter, we will focus on making new types to better express our solutions. We will examine how to combine types to express higher-order information such as

- defining new types that simultaneously combine existing types.
- define alternatives, i.e., a type that can be one of several other types.

After you have read this chapter, you will be able to model values that contain

- a combination of types such as an address consisting of both street names as a string and zip codes as an integer.
- an alternative list of types such as a shape being either a circle parametrized by its center and radius or square parametrized by its four corners.

F# is a strongly typed language, meaning that types are known or inferred at compile time. In the previous chapters, we have used *primitive types* such as `float` and `bool`, function types, and compound types implicitly defined by tuples. These types are used for simple programming tasks, and everything that can be programmed can be accomplished using these types. However, larger programs are often easier to read and write when using more complicated type structures. In this chapter, we will discuss type abbreviations, enumerated types, discriminated unions, records, and structs. Class types are discussed in depth in ??.

## 5.1 Type products: Tuples

*Tuples* are a direct extension of constants. They are immutable and have neither concatenations nor indexing operations. Tuples are unions of immutable types and have the following syntax:

**Listing 5.1:** Tuples are list of expressions separated by commas.

```
1 <expr>{, <expr>}
```

Tuples are identified by the “,” lexeme and often enclosed in parentheses, but that is not required. An example is a triple, also known as a 3-tuple, `(2, true, "hello")`. In interactive mode, the type of tuples is demonstrated in Listing 5.2. The values

**Listing 5.2:** Tuple types are products of sets.

```
1
2 > let tp = (2, true, "hello")
3 printfn "%A" tp;;
4 (2, true, "hello")
5 val tp: int * bool * string = (2, true, "hello")
6 val it: unit = ()
```

`2`, `true`, and `"hello"` are *members*, and the number of elements of a tuple is its *length*. From the response of F#, we see that the tuple is inferred to have the type `int * bool * string`. The “*” denotes the Cartesian product between sets. Tuples can be products of any types and follow the lexical scope rules like value and function bindings. Notice also that a tuple may be printed as a single entity by the `%A` placeholder. In the example we bound `tp` to the tuple. The opposite is also possible, as demonstrated in Listing 5.3. In this example, a function is defined that takes 1 argument, a 3-tuple. If we wanted a function with 3 arguments, then the function binding should have been `let deconstructNPrint a b c = ....`. The value binding `let (a, b, c) = tp`, binds a tuple with 3 named members to a value, thus deconstructing it in terms of its members. This is called pattern matching and will be discussed in further details in ??. Since we used the `\%A` placeholder in

**Listing 5.3: Definition of a tuple.**

```

1 > let deconstructNPrint tp =
2   let (a, b, c) = tp
3   printfn "tp = (%A, %A, %A)" a b c
4 deconstructNPrint (2, true, "hello")
5 deconstructNPrint (3.14, "Pi", 'p');;
6 tp = (2, true, "hello")
7 tp = (3.14, "Pi", 'p')
8 val deconstructNPrint: 'a * 'b * 'c -> unit
9 val it: unit = ()

```

the `printfn` function, the function can be called with 3-tuples of different types. F# informs us that the tuple type is variable by writing `'a * 'b * 'c`. The “`''`” notation means that the type can be decided at run-time, see Section 5.5 for more on variable types.

Pairs or 2-tuples are so common that F# includes two built-in functions, `fst` and `snd`, to extract the first and second element of a pair. This is demonstrated in Listing 5.4.

**Listing 5.4 pair.fsx:****Deconstruction of pairs with the built-in functions `fst` and `snd`.**

```

1 let pair = ("first", "second")
2 printfn "fst(pair) = %s, snd(pair) = %s" (fst pair) (snd pair)

```

---

```

1 $ fsharp --nologo pair.fsx && mono pair.exe
2 fst(pair) = first, snd(pair) = second

```

Tuples of equal lengths can be compared, and the comparison is defined similarly to string comparison. Tuples of equal length are compared element by element. E.g.,  $(1, 2) = (1, 3)$  is false, while  $(1, 2) = (1, 2)$  is true. The “`<>`” operator is the boolean negation of the “`=`” operator. For the “`<`”, “`<=`”, “`>`”, and “`>=`” operators, the strings are ordered lexicographically, such that  $(\text{'a'}, \text{'b'}, \text{'c'}) < (\text{'a'}, \text{'b'}, \text{'s'})$  &&  $(\text{'a'}, \text{'b'}, \text{'s'}) < (\text{'c'}, \text{'o'}, \text{'s'})$  is true, that is, the “`<`” operator on two tuples is true if and only if the left operand should come before the right when sorting alphabetically. See Listing 5.5 for an example. The algorithm for deciding the boolean value of  $(a_1, a_2) < (b_1, b_2)$  is as follows: we start by examining the first elements, and if  $a_1$  and  $b_1$  are different, then the result of  $(a_1, a_2) < (b_1, b_2)$  is equal to the result of  $a_1 < b_1$ . If  $a_1$  and  $b_1$  are equal, then we move on to the next letter and repeat the investigation. The “`<=`”, “`>`”, and “`>=`” operators are defined similarly.

Binding tuples to mutables does not make the tuple mutable. This is demonstrated in Listing 5.6. However, it is possible to define a mutable variable of type tuple such that new tuple values can be assigned to it, as shown in Listing 5.7. Mutable tuples are value types, meaning that binding to new names makes copies, not aliases, as

**Listing 5.5 tupleCompare.fsx:**

Tuples comparison is similar to string comparison.

```

1 let lessThan (a, b, c) (d, e, f) =
2     if a <> d then a < d
3     elif b <> e then b < d
4     elif c <> f then c < f
5     else false
6
7 let printTest x y =
8     printfn "%A < %A is %b" x y (lessThan x y)
9
10 let a = ('a', 'b', 'c');
11 let b = ('d', 'e', 'f');
12 let c = ('a', 'b', 'b');
13 let d = ('a', 'b', 'd');
14 printTest a b
15 printTest a c
16 printTest a d

```

---

```

1 $ fsharp --nologo tupleCompare.fsx && mono tupleCompare.exe
2 ('a', 'b', 'c') < ('d', 'e', 'f') is true
3 ('a', 'b', 'c') < ('a', 'b', 'b') is false
4 ('a', 'b', 'c') < ('a', 'b', 'd') is true

```

**Listing 5.6 tupleOfMutables.fsx:**

A mutable changes value, but the tuple defined by it does not refer to the new value.

```

1 let mutable a = 1
2 let mutable b = 2
3 let c = (a, b)
4 printfn "%A, %A, %A" a b c
5 a <- 3
6 printfn "%A, %A, %A" a b c

```

---

```

1 $ fsharp --nologo tupleOfMutables.fsx && mono
   tupleOfMutables.exe
2 1, 2, (1, 2)
3 3, 2, (1, 2)

```

demonstrated in Listing 5.8. The use of tuples shortens code and highlights semantic content at a higher level, e.g., instead of focusing on the elements, tuples focus on their union. While this may look elegant and short there is the risk of *obfuscation*, i.e., writing compact code that is difficult to read, where an unprepared reader of the code may not easily understand the computation nor appreciate its elegance without

★ an accompanying explanation. Hence, **always keep an eye out for compact and concise ways to write code, but never at the expense of readability.**

**Listing 5.7 mutableTuple.fsx:**

A mutable tuple can be assigned a new value.

```

1 let mutable pair = 1,2
2 printfn "%A" pair
3 pair <- (3,4)
4 printfn "%A" pair

```

---

```

1 $ fsharp --nologo mutableTuple.fsx && mono mutableTuple.exe
2 (1, 2)
3 (3, 4)

```

**Listing 5.8 mutableTupleValue.fsx:**

A mutable tuple is a value type.

```

1 let mutable pair = 1,2
2 let mutable aCopy = pair
3 pair <- (3,4)
4 printfn "%A %A" pair aCopy

```

---

```

1 $ fsharp --nologo mutableTupleValue.fsx && mono
   mutableTupleValue.exe
2 (3, 4) (1, 2)

```

**5.2 Type sums: Discriminated Unions**

A discriminated union is a union of a set of named cases. These cases can further be of specified types. The syntax for defining a discriminated union is as follows:

**Listing 5.9: Syntax for type abbreviation.**

```

1 [<attributes>]
2 type <ident> =
3   [| ]<ident> [of [<ident> :] <aType> [* [<ident> :] <aType>
   ...]]
4   | <ident> [of [<ident> :] <aType> [* [<ident> :] <aType>
   ...]]
5   ...

```

Discriminated unions are reference types, i.e., their content is stored on *The Heap*, see ?? for a discussion on reference types. Since they are immutable, there is no risk of side-effects. As reference types, they only pass a reference when used as arguments to and returned from a function. This is in contrast to value types, which transport a complete copy of the data structure. Discriminated unions are thus effective for large data structures. Discriminated unions can also be represented as structures using the [<Struct>] attribute, in which case they are value types. See ?? for a discussion on structs.

An example just using the named cases but no further specification of types is given in Listing 5.10. Here, we define a discriminated union as three named cases

**Listing 5.10 discriminatedUnions.fsx:**

**A discriminated union of medals. Compare with ??.**

```

1 type medal =
2   Gold
3   | Silver
4   | Bronze
5
6 let aMedal = medal.Gold
7 printfn "%A" aMedal

```

---

```

1 $ fsharp --nologo discriminatedUnions.fsx && mono
   discriminatedUnions.exe
2 Gold

```

signifying three different types of medals. Comparing with the enumerated type in ??, we see that the only difference is that the cases of the discriminated unions have no value. A commonly used discriminated union is the *option type*, see ?? for more detail.

Discriminated unions may also be used to store data. Where the names in enumerated types are aliases of single values, the names used in discriminated unions can hold any value specified at the time of creation. An example is given in Listing 5.11. In this case, we define a discriminated union of two and three-dimensional vectors.

**Listing 5.11 discriminatedUnionsOf.fsx:**

**A discriminated union using explicit subtypes.**

```

1 type vector =
2   Vec2D of float * float
3   | Vec3D of x : float * y : float * z : float
4
5 let v2 = Vec2D (1.0, -1.2)
6 let v3 = Vec3D (x = 1.0, z = -1.2, y = 0.9)
7 printfn "%A and %A" v2 v3

```

---

```

1 $ fsharp --nologo discriminatedUnionsOf.fsx && mono
   discriminatedUnionsOf.exe
2 Vec2D (1.0, -1.2) and Vec3D (1.0, 0.9, -1.2)

```

Values of these types are created using their names followed by a tuple of their arguments. The names are also called field names. The field names may be used when creating discriminated union values, as shown in Line 6. When used, then the arguments may be given in arbitrary order, nevertheless, values for all fields must be given.


Discriminated unions can be defined recursively. This feature is demonstrated in Listing 5.12. In this example we define a tree as depicted in Figure 5.1.

**Listing 5.12** discriminatedUnionTree.fsx:  
A discriminated union modelling binary trees.

```
1 type Tree =
2   Leaf of int
3   | Node of Tree * Tree
4
5 let one = Leaf 1
6 let two = Leaf 2
7 let three = Leaf 3
8 let tree = Node (Node (one, two), three)
9 printfn "%A" tree
```

---

```
1 $ fsharp --nologo discriminatedUnionTree.fsx && mono
   discriminatedUnionTree.exe
2 Node (Node (Leaf 1, Leaf 2), Leaf 3)
```



**Fig. 5.1** The tree with 3 leaves.

Pattern matching must be used in order to define functions on values of a discriminated union. E.g., in Listing 5.13 we define a function that traverses a tree and prints the content of the nodes.

**Listing 5.13** discriminatedUnionPatternMatching.fsx:  
A discriminated union modelling binary trees.

```
1 type Tree = Leaf of int | Node of Tree * Tree
2 let rec traverse (t : Tree) : string =
3   match t with
4     Leaf(v) -> string v
5     | Node(left, right) -> (traverse left) + ", " +
6       (traverse right)
7
8 let tree = Node (Node (Leaf 1, Leaf 2), Leaf 3)
9 printfn "%A: %s" tree (traverse tree)
```

---

```
1 $ fsharp --nologo discriminatedUnionPatternMatching.fsx &&
   mono discriminatedUnionPatternMatching.exe
2 Node (Node (Leaf 1, Leaf 2), Leaf 3): 1, 2, 3
```

Discriminated unions are very powerful and can often be used instead of class hierarchies. Class hierarchies are discussed in ??.

### 5.3 Records

A record is a compound of named values, and a record type is defined as follows:

**Listing 5.14: Syntax for defining record types.**

```
1 [ <attributes> ]
2 type <ident> = {
3     [ mutable ] <label1> : <type1>
4     [ mutable ] <label2> : <type2>
5     ...
6 }
```

Records are collections of named variables and values of possibly different types. They are reference types, and thus their content is stored on *The Heap*, see ?? for a discussion on reference types. Records can also be *struct records* using the [`<Struct>`] attribute, in which case they are value types. See ?? for a discussion on structs. An example of using records is given in Listing 5.15. The values of individual members of a record are obtained using the “.” notation. This example illustrates a

**Listing 5.15 records.fsx:**

A record is defined for holding information about a person.

```
1 type person = {
2     name : string
3     age : int
4     height : float
5 }
6
7 let author = {name = "Jon"; age = 50; height = 1.75}
8 printfn "%A\nname = %s" author author.name
9
10 -----
11 $ fsharp --nologo records.fsx && mono records.exe
12 { name = "Jon"
13   age = 50
14   height = 1.75 }
15 name = Jon
```

how record type is used to store varied data about a person.

If two record types are defined with the same label set, then the latter dominates the former. This is demonstrated in Listing 5.16. In the example, two identical record types are defined, and we use the built-in `GetType()` method to inspect the type



**Listing 5.16** recordsDominance.fsx:

Redefined types dominate old record types, but earlier definitions are still accessible using explicit or implicit specification for bindings.

```

1 type person = { name : string; age : int; height : float }
2 type teacher = { name : string; age : int; height : float }
3
4 let lecturer = {name = "Jon"; age = 50; height = 1.75}
5 printfn "%A : %A" lecturer (lecturer.GetType())
6 let author : person = {name = "Jon"; age = 50; height = 1.75}
7 printfn "%A : %A" author (author.GetType())
8 let father = {person.name = "Jon"; age = 50; height = 1.75}
9 printfn "%A : %A" author (author.GetType())
-----
1 $ fsharpc --nologo recordsDominance.fsx && mono
   recordsDominance.exe
2 { name = "Jon"
3   age = 50
4   height = 1.75 } : RecordsDominance+teacher
5 { name = "Jon"
6   age = 50
7   height = 1.75 } : RecordsDominance+person
8 { name = "Jon"
9   age = 50
10  height = 1.75 } : RecordsDominance+person

```

of bindings. We see that `lecturer` is of `RecordsDominance+teacher` type, since `teacher` dominates the identical `person` type definition. However, we may enforce the `person` type by either specifying it for the name, as in `let author : person = ...`, or by fully or partially specifying it in the record expression following the “=” sign. In both cases, they are of `RecordsDominance+person` type. The built-in `GetType()` method is inherited from the base class for all types, see ?? for a discussion on classes and inheritance.

Note that when creating a record you must supply a value to all fields, and you cannot refer to other fields of the same record, i.e., `{name = "Jon"; age = height * 3; height = 1.75}` is illegal.

Since records are per default reference types, binding creates aliases, not copies. This matters for mutable members, in which case when copying, we must explicitly create a new record with the old data. Copying can be done either by using referencing to the individual members of the source or using the short-hand *with* notation. This is demonstrated in Listing 5.17. Here, `age` is defined as a mutable value and can be changed using the usual “<-” assignment operator. The example demonstrates two different ways to create records. Note that when the mutable value `author.age` is changed in line 10, then `authorAlias` also changes, since it is an alias of `author`, but neither `authorCopy` nor `authorCopyAlt` changes, since they are copies. As

**Listing 5.17 recordCopy.fsx:**

Bindings are references. To copy and not make an alias, explicit copying must be performed.

```

1 type person = {
2     name : string;
3     mutable age : int;
4 }
5
6 let author = {name = "Jon"; age = 50}
7 let authorAlias = author
8 let authorCopy = {name = author.name; age = author.age}
9 let authorCopyAlt = {author with name = "Noj"}
10 author.age <- 51
11 printfn "author : %A" author
12 printfn "authorAlias : %A" authorAlias
13 printfn "authorCopy : %A" authorCopy
14 printfn "authorCopyAlt : %A" authorCopyAlt

```

---

```

1 $ fsharp -nologo recordCopy.fsx && mono recordCopy.exe
2 author : { name = "Jon"
3     age = 51 }
4 authorAlias : { name = "Jon"
5     age = 51 }
6 authorCopy : { name = "Jon"
7     age = 50 }
8 authorCopyAlt : { name = "Noj"
9     age = 50 }

```

illustrated, copying using `with` allows for easy copying and partial updates of another record value.

## 5.4 Type Abbreviations

F# allows for renaming of types, which is called *type abbreviation* or *type aliasing*. The syntax is:

**Listing 5.18: Syntax for type abbreviation.**

```
1 type <ident> = <aType>
```

where the identifier is a new name, and the type-name is an existing type or a compound of existing types. Listing 5.19 shows examples of the definition of several type abbreviations. Here we define the abbreviations `size`, `position`, `person`, and `intToFloat`, and later make bindings enforcing the usage of these abbreviations.

**Listing 5.19 typeAbbreviation.fsx:**

Defining four type abbreviations, three of which are compound types.

```

1 type size = int
2 type position = float * float
3 type person = string * int
4 type intToFloat = int -> float
5
6 let sz : size = 3
7 let pos : position = (2.5, -3.2)
8 let pers : person = ("Jon", 50)
9 let conv : intToFloat = fun a -> float a
10 printfn "%A, %A, %A, %A" sz pos pers (conv 2)

```

---

```

1 $ fsharpc --nologo typeAbbreviation.fsx && mono
   typeAbbreviation.exe
2 3, (2.5, -3.2), ("Jon", 50), 2.0

```

Type abbreviations are used as short abbreviations of longer types, and they add semantic content to the program text, thus making programs shorter and easier to read. Type abbreviations allow the programmer to focus on the intended structure at a higher level by, e.g., programming in terms of a type `position` rather than `float * float`. Thus, they often result in programs with fewer errors. Type abbreviations also make maintenance easier. For instance, if we at a later stage decide that positions can only have integer values, then we only need to change the definition of the type abbreviation, not every place a value of type `position` is used.

## 5.5 Variable Types

An advanced topic in F# is *variable types*. There are three different versions of variable types in F#: *runtime resolved*, which have the syntax '`<ident>`', *anonymous*, which are written as “_”, and *statically resolved*, which have the syntax '^<ident>'. Variable types are particularly useful for functions that work for many types. An example of a generic function and its use is given in Listing 5.20. In this example, the function `apply` has runtime resolved variable type, and it accepts three parameters: `f`, `x`, and `y`. The function will work as long as the parameters for `f` is a function of two parameters of identical type, and `x` and `y` are values of the same type. Thus, in the `printfn` statement we are able to use `apply` for both an integer and a float variant.

The example in Listing 5.20 illustrates a very complicated way to add two numbers. The “+” operator works for both types out of the box, so why not something simpler like relying on the F# type inference system by not explicitly specifying types, as

**Listing 5.20 variableType.fsx:**

A function apply with runtime resolved types.

```

1 let apply (f : 'a -> 'a -> 'a) (x : 'a) (y : 'a) : 'a = f x y
2 let intPlus (x : int) (y : int) : int = x + y
3 let floatPlus (x : float) (y : float) : float = x + y
4
5 printfn "%A %A" (apply intPlus 1 2) (apply floatPlus 1.0 2.0)

```

---

```

1 $ fsharp --nologo variableType.fsx && mono variableType.exe
2 3 3.0

```

attempted in Listing 5.21? Unfortunately, the example fails to compile, since the type

**Listing 5.21 variableTypeError.fsx:**Even though the “+” operator is defined for both integers and floats, the type inference is static and infers `plus : int -> int`.

```

1 let plus x y = x + y
2
3 printfn "%A %A" (plus 1 2) (plus 1.0 2.0)

```

---

```

1 $ fsharp --nologo variableTypeError.fsx && mono
  variableTypeError.exe
2
3 variableTypeError.fsx(3,34): error FS0001: This expression
  was expected to have type
4     'int'
5 but here has type
6     'float'
7
8 variableTypeError.fsx(3,38): error FS0001: This expression
  was expected to have type
9     'int'
10 but here has type
11     'float'

```

inference is performed at compile time, and by `plus 1 2`, it is inferred that `plus : int -> int`. Hence, calling `plus 1.0 2.0` is a type error. Function bindings allow for the use of the `inline` keyword, and adding this successfully reuses the definition of `plus` for both types, as shown in Listing 5.22. In the example, adding the `inline` does two things: Firstly, it copies the code to be performed to each place the function is used, and secondly, it forces statically resolved variable type checking independently in each place. The type annotations inferred as a result of the `inline`-keyword may be written explicitly, as shown in Listing 5.23. The example in Listing 5.23 demonstrates the statically resolved variable type syntax, `^<ident>`, as well as the use of *type constraints*, using the keyword `when`. Type constraints have a rich syntax, but will not be discussed further in this book. In the example, the type constraint `when ^a : (static member ( + ) : ^a * ^a -> ^a)`

**Listing 5.22 variableTypeInline.fsx:**

The keyword `inline` forces static and independent inference each place the function is used. Compare to the error case in Listing 5.21.

```
1 let inline plus x y = x + y
2
3 printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
-----
1 $ fsharpc --nologo variableTypeInline.fsx && mono
  variableTypeInline.exe
2 3 3.0
```

**Listing 5.23 compiletimeVariableType.fsx:**

Explicitly spelling out of the statically resolved type variables from Listing 5.21.

```
1 let inline plus (x : ^a) (y : ^a) : ^a when ^a : (static
  member ( + ) : ^a * ^a -> ^a) = x + y
2
3 printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
-----
1 $ fsharpc --nologo compiletimeVariableType.fsx && mono
  compiletimeVariableType.exe
2 3 3.0
```

is given using the object-oriented properties of the type variable `^a`, meaning that the only acceptable type values are those which have a member function `(+)` taking a tuple and giving a value all of identical type, and where the type can be inferred at compile time. See ?? for details on member functions.

The `inline` construction is useful when generating generic functions and still profiting from static type checking. However, explicit copying of functions is often something better left to the compiler to optimize over. An alternative seems to be using runtime resolved variable types with the `'<ident>` syntax. Unfortunately, this is not possible in case of most operators, since they have been defined in the `FSharp.Core` namespace to be statically resolved variable types. E.g., the `“+”` operator has type `( + ) : ^T1 -> ^T2 -> ^T3 (requires ^T1 with static member (+) and ^T2 with static member (+))`.

## 5.6 Key concepts and terms in this chapter

In this chapter you have learned about:

- the **product type** also known as a **tuple**, which is equivalent to a set product;
- the **sum type** also known as a **discriminate union**;
- the **records** which are similar to tuples, but allows you to name the entries
- and as an advanced topic, you have seen F# has flexibility in specifying types either at compile or runtime.

## Chapter 6

# Making programs and documenting them

**Abstract** Programs are more than a set of instructions, which when executed produces the desired result. Programming is an activity, and a program is the result of a process, in which a problem has been expressed, analyzed, subdivided, implemented, tested, and possibly rephrased. And often the process and its result are to be wrapped and documented for it to be useful by the programmer or others. In this chapter, we will zoom out, and focus on some of these surrounding processes. The chapter will describe:

- How to design functions.
- How and why to document programs using in-code documentation.

## 6.1 The 7-step guide to writing function

Pólya's problem solving technique described in Section 1.2 is a useful starting point for solving problems, and for the object-oriented programming paradigm to be discussed in later chapters ??, approaches such as Pólya's have been put into systems. Regardless of the origin, there is always a point, where a programmer has to focus on small scale problems such as: Which functions, should be used, and how should the functions be designed? This is not an area heavily investigated in the literature, but often a skill that programmers pick up by actively engaging in programming alone or with other programmers. However, here I will venture a recipes for designing functions, which me and my colleagues call The 7-step guide to writing functions. It is not meant as the ultimate guide or as required steps, but it is our experience that these steps contain essential elements that consciously or perhaps unconsciously always take part in designing useful and reusable functions.

To decide which functions to write and how to write them:

1. Note: Write a short note on what a function should do.
2. Name: Invent a name for the function. Semantically meaningful names should be preferred.
3. External test: Write a small test-program, which uses the yet to be written function.
4. Type: Decide what type, the function should have, e.g., by how you used it in your test-program.
5. Implement: Write the function and possibly its helper functions.
6. Internal test: Extend your test-program with more examples, where you use the function and based on its implementation.
7. Run: Run the test-program
8. Document: Write a brief in-code documentation of the function, see Section 6.2.

As an example, let us revisit the problem of solving a quadratic equation: The task is to find the zero-crossings of a second degree polynomial, i.e.,  $f(x) = ax^2 + bx + c = 0$ . The process could be as follows:

1. Note: We decide to stick to the mathematical description:

Given parameters  $a$ ,  $b$ , and  $c$ , the function should return the 0, 1, or possibly 2 locations  $x$ , where  $f(x) = 0$ .



- Name: This function may be used together with solvers for other equations, so we decide to give it the rather long name

```
solveQuadraticEquation.
```

- External test: We decide to a single tests to get a feeling of how the function is to be used:

#### Listing 6.1: Defining the function sum

```
1 let p = solveQuadraticEquation 1.0 0.3 -1.0
2 printfn "0=1.0x^2+0.3x-1.0 => x = %1" p;;
```

- Type: Since there may be 0-2 points  $x$ , where  $f(x) = 0$ , the output answers could be a tuple. Furhter, since  $a, b, c, x \in \mathbb{R}$ , we will use floats. Hence,

```
solveQuadraticEquation -> float -> float -> float -> float*float.
```

- Implement: Thinking about how to write `solveQuadraticEquation`, we decide that since the calculation of the discriminant is done twice, we will add it as a helper function. Our resulting code is:

#### Listing 6.2: Defining the function sum

```
1 let discriminant a b c = b ** 2.0 - 4.0 * a * c
2
3 let solveQuadraticEquation a b c =
4   let d = discriminant a b c
5   ((-b + sqrt d) / (2.0 * a),
6    (-b - sqrt d) / (2.0 * a))
```

- Internal test: Working with the code, we realize that it is unclear, what happens, when there are 0 og 1 solutions, so we update the external test and add more tests:

#### Listing 6.3: Defining the function sum

```
1 let p1 = solveQuadraticEquation 1.0 0.3 -1.0
2 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p1
3 let p2 = solveQuadraticEquation 1.0 0.0 0.0
4 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p2
5 let p3 = solveQuadraticEquation 1.0 0.0 1.0
6 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p3
```

- Run: The complete code with examples and its output, when executed is shown in Listing 6.4

**Listing 6.4 solveQuadraticEquation.fsx:**  
Solving quadratic equations

```

1 let discriminant a b c = b ** 2.0 - 4.0 * a * c
2
3 let solveQuadraticEquation a b c =
4     let d = discriminant a b c
5     ((-b + sqrt d) / (2.0 * a),
6      (-b - sqrt d) / (2.0 * a))
7
8 let p1 = solveQuadraticEquation 1.0 0.3 -1.0
9 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p1
10 let p2 = solveQuadraticEquation 1.0 0.0 0.0
11 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p2
12 let p3 = solveQuadraticEquation 1.0 0.0 1.0
13 printfn "0=1.0x^2+0.3x-1.0 => x = %A" p3

```

---

```

1 $ dotnet fsi solveQuadraticEquation.fsx
2 0=1.0x^2+0.3x-1.0 => x = (0.8611874208, -1.161187421)
3 0=1.0x^2+0.3x-1.0 => x = (0.0, -0.0)
4 0=1.0x^2+0.3x-1.0 => x = (nan, nan)

```

8. Document: The following section will discuss how to perform in-code documentation and use Listing 6.4 as an example.

## 6.2 Programming as a communication activity

Documentation is a very important part of writing programs, since it is most unlikely that you will be writing really obvious code. Moreover, what seems obvious at the point of writing may be mystifying months later to the author and to others. Documentation serves several purposes:

1. Communicate to the user of the code, what it does and how to use it. In this book, we will emphasize the XML-standard for this purpose.
2. Highlight big insights essential for the code, which is important for other programmers to understand and maintain the code.
3. Highlight possible conflicts and/or areas where the code could be changed later, which is also targeted programmers rather than users of the code.

The essential point is that coding is a journey in problem-solving, and proper documentation is an aid in understanding the solution and the journey that lead to it. Documentation is most often a mixture of in-code documentation and accompanying

documents. Here, we will focus on in-code documentation which arguably causes problems in multi-language environments and run the risk of bloating code. Since documentation is about human-to-human communication, there is no correct documentation. However, as in all things the documentation can both be too little and too much, and the ability to produce documentation is best learned by example and by doing.

F# has two different syntaxes for comments. Comments can be block comments:

**Listing 6.5: Block comments.**

```
1 (*<any text>*)
```

The comment text (<any text>) can be any text and is still parsed by F# as keywords and basic types, implying that `(* a comment (* in a comment *) *)` and `(* " " *)` are valid comments, while `(* " *)` is invalid.

Alternatively, comments may also be line comments,

**Listing 6.6: Line comments.**

```
1 //<any text>
```

where the comment text ends after the first newline.

The block and line comments are used principally for communicating insights and comments into the code between programmers which want to understand and/or maintain the code.

Users of the code, are most likely also programmers but have an outside perspective. They are more interested in what the code does, and how it is to be used. For this we recommend the *Extensible Markup Language* documentation standard (*XML-standard*)¹. All lines of the XML-standard starts with a triple-slash `///`. Thus, it is a line-comment, where an extra slash has been added for visual flair. XML consists of tags which always appear in pairs, e.g., the tag “tag” would look like `<tag> . . . </tag>`. A subset of tags are listed in Table 6.1. If no tags are used, then it is automatically assumed to be a `<summary>`. An example of a documented script is shown in Listing 6.7. is:

Several tools exists which extracts the comments from source code and reorder the comments into manual type structures, such as Doxygen. Popular output from such tools are both HTML and  $\text{\LaTeX}$ . However, for this text, the usage of the XML-standard as a way to standardize comments will suffice.

---

¹ For specification of C# documentations comments see ECMA-334: <http://www.ecma-international.org/publications/files/ECMA-ST/ECma-334.pdf>

Tag	Description
<c>	Set text in a code-font.
<code>	Set one or more lines in code-font.
<example>	Set as an example.
<exception>	Describe the exceptions a function can throw.
<list>	Create a list or table.
<para>	Set text as a paragraph.
<param>	Describe a parameter for a function or constructor.
<paramref>	Identify that a word is a parameter name.
<permission>	Document the accessibility of a member.
<remarks>	Further describe a function.
<returns>	Describe the return value of a function.
<see>	Set as link to other functions.
<seealso>	Generate a See Also entry.
<summary>	Main description of a function or value.
<typeparam>	Describe a type parameter for a generic type or method.
<typeparamref>	Identify that a word is a type parameter name.
<value>	Describe a value.

**Table 6.1** Recommended XML tags for documentation comments, from ECMA-334 3rd Edition, Annex E, Section 2.

### 6.3 Key concepts and terms in this chapter

This chapter has considered elements that are an important part of the activity of programming, but to some extent complement the specific act of writing source code. You have seen:

- How to use the **7-step guide** to design functions, which emphasizes writing examples of function usage before implementing the function itself.
- Write **in-code** documentation to support the understanding of the code.
- Documentation is written for programmers and there are at least two different types: **users** and **maintainers**.
- The **XML standard** uses `///`, is for both types of programmers, and documents what a program does and how it is to be used.
- The **line** and **block** comments are for implementation-specific details and intended to be read by programmers who seek to understand and maintain the code.
- There is no such thing as the correct documentation, but you are well advised to follow the XML standard and to improve your skill by writing documentation and sharing it with others.

**Listing 6.7** commentExample.fsx:  
Code with XML comments.

```

1  /// The discriminant of a quadratic equation with
2  /// parameters a, b, and c
3  let discriminant a b c = b ** 2.0 - 4.0 * a * c
4
5  /// <summary>Find x when 0 = ax^2+bx+c.</summary>
6  /// <remarks>Negative discriminants are not checked.</remarks>
7  /// <example>
8  ///     The following code:
9  ///     <code>
10     ///         let p = solveQuadraticEquation 1.0 0.3 -1.0
11     ///         printfn "0=1.0x^2+0.3x-1.0 => x = %A" p
12     ///     </code>
13     ///     prints <c>0=1.0x^2+0.3x-1.0 => x = (0.9, -1.2)</c>.
14     /// </example>
15     /// <param name="a">Quadratic coefficient.</param>
16     /// <param name="b">Linear coefficient.</param>
17     /// <param name="c">Constant coefficient.</param>
18     /// <returns>The solution to x as a tuple.</returns>
19     let solveQuadraticEquation a b c =
20         let d = discriminant a b c
21         ((-b + sqrt d) / (2.0 * a),
22          (-b - sqrt d) / (2.0 * a))
23
24     let p1 = solveQuadraticEquation 1.0 0.3 -1.0
25     printfn "0=1.0x^2+0.3x-1.0 => x = %A" p1
26     let p2 = solveQuadraticEquation 1.0 0.0 0.0
27     printfn "0=1.0x^2+0.3x-1.0 => x = %A" p2
28     let p3 = solveQuadraticEquation 1.0 0.0 1.0
29     printfn "0=1.0x^2+0.3x-1.0 => x = %A" p3

```



## Appendix A


### The Console in Windows, MacOS X, and Linux

Almost all popular operating systems are accessed through a user-friendly *graphical user interface (GUI)* that is designed to make typical tasks easy to learn to solve. As a computer programmer, you often need to access some of the functionalities of the computer, which, unfortunately, are sometimes complicated by this particular graphical user interface. The *console*, also called the *terminal* and the *Windows command line*, is the right hand of a programmer. The console is a simple program that allows you to complete text commands. Almost all the tasks that can be done with the graphical user interface can be done in the console and vice versa. Using the console, you will benefit from its direct control of the programs we write, and in your education, you will benefit from the fast and raw information you get through the console.

#### A.1 The Basics

When you open a *directory* or *folder* in your preferred operating system, the directory will have a location in the file system, whether from the console or through the operating system's graphical user interface. The console will almost always be associated with a particular directory or folder in the file system, and it is said that it is the directory that the console is in. The exact structure of file systems varies between Linux, MacOS X, and Windows, but common is that it is a hierarchical structure. This is illustrated in Figure A.1.

There are many predefined console commands, available in the console, and you can also make your own. In the following sections, we will review the most important commands in the three different operating systems. These are summarized in Table A.1.



**Fig. A.1** The top file hierarchy levels of common operating systems.

Windows	MacOS X/Linux	Description
<code>dir</code>	<code>ls</code>	Show content of present directory.
<code>cd &lt;d&gt;</code>	<code>cd &lt;d&gt;</code>	Change present directory to <d>.
<code>mkdir &lt;d&gt;</code>	<code>mkdir &lt;d&gt;</code>	Create directory <d>.
<code>rmdir &lt;d&gt;</code>	<code>rmdir &lt;d&gt;</code>	Delete <d> (Warning: cannot be reverted).
<code>move &lt;f&gt; &lt;f   d&gt;</code>	<code>mv &lt;f&gt; &lt;f   d&gt;</code>	Move <f1> to <f   d>.
<code>copy &lt;f1&gt; &lt;f2&gt;</code>	<code>cp &lt;f1&gt; &lt;f2&gt;</code>	Create a new file called <f2> as a copy of <f1>.
<code>del &lt;f&gt;</code>	<code>rm &lt;f&gt;</code>	delete <f> (Warning: cannot be reverted).
<code>echo &lt;s   v&gt;</code>	<code>echo &lt;s   v&gt;</code>	Write a string or content of a variable to screen.

**Table A.1** The most important console commands for Windows, MacOS X, and Linux. Here <f*> is shorthand for any filename, <d> for any directory name, <s> for any string, and <v> for any shell-variable.

## A.2 Windows

In this section we will discuss the commands summarized in Table A.1. Windows 7 and earlier versions: To open the console, press **Start** -> **Run** in the lower left corner, and then type `cmd` in the box. In Windows 8 and 10, you right-click on the windows icon, choose **Run** or equivalent in your local language, and type `cmd`. Alternatively, you can type **Windows-key** + **R**. Now you should open a console window with a prompt showing something like Listing A.1.



**Listing A.1: The Windows console.**

```
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation. All rights
reserved.

C:\Users\sporrington>
```

To see which files are in the directory, use *dir*, as shown in Listing A.2.

**Listing A.2: Directory listing with dir.**

```
C:\Users\sporrington>dir
Volume in drive C has no label.
Volume Serial Number is 94F0-31BD

Directory of C:\Users\sporrington

30-07-2015  15:23    <DIR>          .
30-07-2015  15:23    <DIR>          ..
30-07-2015  14:27    <DIR>          Contacts
30-07-2015  14:27    <DIR>          Desktop
30-07-2015  17:40    <DIR>          Documents
30-07-2015  15:11    <DIR>          Downloads
30-07-2015  14:28    <DIR>          Favorites
30-07-2015  14:27    <DIR>          Links
30-07-2015  14:27    <DIR>          Music
30-07-2015  14:27    <DIR>          Pictures
30-07-2015  14:27    <DIR>          Saved Games
30-07-2015  17:27    <DIR>          Searches
30-07-2015  14:27    <DIR>          Videos
                0 File(s)                0 bytes
                13 Dir(s)  95.004.622.848 bytes free

C:\Users\sporrington>
```

We see that there are no files and thirteen directories (DIR). The columns tell from left to right: the date and time of their creation, the file size or if it is a folder, and the name file or directory name. The first two folders “.” and “..” are found in each folder and refer to this folder as well as the one above in the hierarchy. In this case, the folder “.” is an alias for C:\Users\sporrington and “..” for C:\Users.

Use *cd* to change directory, e.g., to Documents, as in Listing A.3.

**Listing A.3: Change directory with cd.**

```
C:\Users\sporrington>cd Documents

C:\Users\sporrington\Documents>
```

Note that some systems translate default filenames, so their names may be given different names in different languages in the graphical user interface as compared to the console.

You can use *mkdir* to create a new directory called, e.g., *myFolder*, as illustrated in Listing A.4.

**Listing A.4: Creating a directory with *mkdir*.**

```
C:\Users\sporrington\Documents>mkdir myFolder

C:\Users\sporrington\Documents>dir
Volume in drive C has no label.
Volume Serial Number is 94F0-31BD

Directory of C:\Users\sporrington\Documents

30-07-2015  19:17    <DIR>          .
30-07-2015  19:17    <DIR>          ..
30-07-2015  19:17    <DIR>          myFolder
               0 File(s)                0 bytes
               3 Dir(s)  94.656.638.976 bytes free

C:\Users\sporrington\Documents>
```

By using *dir* we inspect the result.

Files can be created by, e.g., *echo* and *redirection*, as demonstrated in Listing A.5.

**Listing A.5: Creating a file with *echo* and *redirection*.**

```
C:\Users\sporrington\Documents>echo "Hi" > hi.txt

C:\Users\sporrington\Documents>dir
Volume in drive C has no label.
Volume Serial Number is 94F0-31BD

Directory of C:\Users\sporrington\Documents

30-07-2015  19:18    <DIR>          .
30-07-2015  19:18    <DIR>          ..
30-07-2015  19:17    <DIR>          myFolder
30-07-2015  19:18                8 hi.txt
               1 File(s)                8 bytes
               3 Dir(s)  94.656.634.880 bytes free

C:\Users\sporrington\Documents>
```

To move the file *hi.txt* to the directory *myFolder*, use *move*, as shown in Listing A.6.

**Listing A.6: Move a file with move.**

```
C:\Users\sporrington\Documents>move hi.txt myFolder
1 file(s) moved.

C:\Users\sporrington\Documents>
```

Finally, use *del* to delete a file and *rmdir* to delete a directory, as shown in Listing A.7.

**Listing A.7: Delete files and directories with del and rmdir.**

```
C:\Users\sporrington\Documents>cd myFolder

C:\Users\sporrington\Documents\myFolder>del hi.txt

C:\Users\sporrington\Documents\myFolder>cd ..

C:\Users\sporrington\Documents>rmdir myFolder

C:\Users\sporrington\Documents>dir
Volume in drive C has no label.
Volume Serial Number is 94F0-31BD

Directory of C:\Users\sporrington\Documents

30-07-2015 19:20 <DIR>          .
30-07-2015 19:20 <DIR>          ..
               0 File(s)              0 bytes
               2 Dir(s)  94.651.142.144 bytes free

C:\Users\sporrington\Documents>
```

The commands available from the console must be in its *search path*. The search path can be seen using *echo*, as shown in Listing A.8.

**Listing A.8: Displaying the search path.**

```
C:\Users\sporrington\Documents>echo %Path%
C:\Windows\system32;C:\Windows;C:\Windows\System32\Wbem;
C:\Windows\System32\WindowsPowerShell\v1.0\;"\Program
Files\emacs-24.5\bin\"

C:\Users\sporrington\Documents>
```

The path can be changed using the Control panel in the graphical user interface. In Windows 7, choose the Control panel, choose System and Security → System → Advanced system settings → Environment Variables. In Windows 10, you can find this window by searching for “Environment” in the Control panel. In

the window's **System variables** box, double-click on **Path** and add or remove a path from the list. The search path is a list of paths separated by “;”. Beware, Windows uses the search path for many different tasks, so remove only paths that you are certain are not used for anything.

A useful feature of the console is that you can use the **tab**-key to cycle through filenames. E.g., if you write `cd` followed by a space and **tab** a couple of times, then the console will suggest to you the available directories.

### A.3 MacOS X and Linux

MacOS X (OSX) and Linux are very similar, and both have the option of using *bash* as console. It is in the standard console on MacOS X and on many Linux distributions. A summary of the most important *bash* commands is shown in Table A.1. In MacOS X, you find the console by opening **Finder** and navigating to **Applications** → **Utilities** → **Terminal**. In Linux, the console can be started by typing **Ctrl + Alt + T**. Some Linux distributions have other key-combinations such as **Super + T**.

Once opened, the console is shown in a window with content, as shown in Listing A.9.

#### Listing A.9: The MacOS console.

```
Last login: Thu Jul 30 11:52:07 on ttys000
FN11194:~ sporring$
```

“FN11194” is the name of the computer, the character `~` is used as an alias for the user's home directory, and “sporring” is the username for the user presently logged onto the system. Use *ls* to see which files are present, as shown in Listing A.10.

#### Listing A.10: Display a directory content with *ls*.

```
FN11194:~ sporring$ ls
Applications  Documents    Library      Music
Public
Desktop       Downloads    Movies        Pictures
FN11194:~ sporring$
```

More details about the files are available by using flags to *ls* as demonstrated in Listing A.11.

**Listing A.11: Display extra information about files using flags to `ls`.**

```

FN11194:~ sporring$ ls -l
drwx----- 6 sporring  staff   204 Jul 30 14:07
    Applications
drwx-----+ 32 sporring  staff  1088 Jul 30 14:34 Desktop
drwx-----+ 76 sporring  staff  2584 Jul  2 15:53 Documents
drwx-----+  4 sporring  staff   136 Jul 30 14:35 Downloads
drwx-----@ 63 sporring  staff  2142 Jul 30 14:07 Library
drwx-----+  3 sporring  staff   102 Jun 29 21:48 Movies
drwx-----+  4 sporring  staff   136 Jul  4 17:40 Music
drwx-----+  3 sporring  staff   102 Jun 29 21:48 Pictures
drwxr-xr-x+  5 sporring  staff   170 Jun 29 21:48 Public
FN11194:~ sporring$

```

The flag `-l` means long, and many other flags can be found by querying the built-in manual with `man ls`. The output is divided into columns, where the left column shows a number of codes: “d” stands for directory, and the set of three of optional “rwx” denote whether respectively the owner, the associated group of users, and anyone can respectively “r” - read, “w” - write, and “x” - execute the file. In all directories but the Public directory, only the owner can do any of the three. For directories, “x” means permission to enter. The second column can often be ignored, but shows how many links there are to the file or directory. Then follows the username of the owner, which in this case is `sporring`. The files are also associated with a group of users, and in this case, they all are associated with the group called `staff`. Then follows the file or directory size, the date of last change, and the file or directory name. There are always two hidden directories: “.” and “..”, where “.” is an alias for the present directory, and “..” for the directory above. Hidden files will be shown with the `-a` flag.

Use `cd` to change to the directory, for example to Documents as shown in Listing A.12.

**Listing A.12: Change directory with `cd`.**

```

FN11194:~ sporring$ cd Documents/
FN11194:Documents sporring$

```

Note that some graphical user interfaces translate standard filenames and directories to the local language, such that navigating using the graphical user interface will reveal other files and directories, which, however, are aliases.

You can create a new directory using `mkdir`, as demonstrated in Listing A.13.

**Listing A.13: Creating a directory using `mkdir`.**

```
FN11194:Documents sporring$ mkdir myFolder
FN11194:Documents sporring$ ls
myFolder
FN11194:tmp sporring$
```

A file can be created using `echo` and with *redirection*, as shown in Listing A.14.

**Listing A.14: Creating a file with `echo` and redirection.**

```
FN11194:Documents sporring$ echo "hi" > hi.txt
FN11194:Documents sporring$ ls
hi.txt          myFolder
```

To move the file `hi.txt` into `myFolder`, use `mv`. This is demonstrated in Listing A.15.

**Listing A.15: Moving files with `mv`.**

```
FN11194:Documents sporring$ echo mv hi.txt myFolder/
FN11194:Documents sporring$
```

To delete the file and the directory, use `rm` and `rmdir`, as shown in Listing A.16.

**Listing A.16: Deleting files and directories.**

```
FN11194:Documents sporring$ cd myFolder/
FN11194:myFolder sporring$ rm hi.txt
FN11194:myFolder sporring$ cd ..
FN11194:Documents sporring$ rmdir myFolder/
FN11194:Documents sporring$ ls
FN11194:Documents sporring$
```

Only commands found on the *search path* are available in the console. The content of the search path is seen using the `echo` command, as demonstrated in Listing A.17.

**Listing A.17: The content of the search path.**

```
FN11194:Documents sporring$ echo $PATH
/Applications/Maple
17:/Applications/PackageManager.app/Contents/MacOS/:
/Applications/MATLAB_R2014b.app/bin:/opt/local/bin:
/opt/local/sbin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:
/sbin:/opt/X11/bin:/Library/TeX/texbin
FN11194:Documents sporring$
```

The search path can be changed by editing the setup file for Bash. On MacOS X it is called `~/.profile`, and on Linux it is either `~/.bash_profile` or `~/.bashrc`.

Here new paths can be added by adding the following line: `export PATH=<new path>:<another new path>:$PATH`.

A useful feature of Bash is that the console can help you write commands. E.g., if you write `fs` followed by pressing the `tab`-key, and if `Mono` is in the search path, then Bash will typically respond by completing the line as `fsharp`, and by further pressing the `tab`-key some times, Bash will show the list of options, typically `fshpari` and `fsharpc`. Also, most commands have an extensive manual which can be accessed using the `man` command. E.g., the manual for `rm` is retrieved by `man rm`.





## Appendix B

# Number Systems on the Computer

### B.1 Binary Numbers

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base 10* meaning that a decimal number consists of a sequence of digits separated by a *decimal point*, where each *digit* can have values  $d \in \{0, 1, 2, \dots, 9\}$  and the weight of each digit is proportional to its place in the sequence of digits with respect to the decimal point, i.e., the number  $357.6 = 3 \cdot 10^2 + 5 \cdot 10^1 + 7 \cdot 10^0 + 6 \cdot 10^{-1}$ , or in general, for a number consisting of digits  $d_i$  with  $n + 1$  and  $m$  digits to the left and right of the decimal point, the value  $v$  is calculated as:

$$v = \sum_{i=-m}^n d_i 10^i. \quad (\text{B.1})$$

The basic unit of information in almost all computers is the binary digit, or *bit* for short. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values  $b \in \{0, 1\}$ , and the base is 2. The general equation is,

$$v = \sum_{i=-m}^n b_i 2^i, \quad (\text{B.2})$$

and examples are  $1011.1_2 = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 + 1 \cdot 2^{-1} = 11.5$ . Notice that we use subscript 2 to denote a binary number, while no subscript is used for decimal numbers. The left-most bit is called the *most significant bit*, and the right-most bit is called the *least significant bit*. Due to typical organisation of computer memory, 8 binary digits is called a *byte*, and the term *word* is not universally defined but typically related to the computer architecture, a program is running on, such as 32 or 64 bits.

Other number systems are often used, e.g., *octal numbers*, which are base 8 numbers and have digits  $o \in \{0, 1, \dots, 7\}$ . Octals are useful short-hand for binary, since 3 binary digits map to the set of octal digits. Likewise, *hexadecimal numbers* are base 16 with digits  $h \in \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, b, c, d, e, f\}$ , such that  $a_{16} = 10$ ,  $b_{16} = 11$  and so on. Hexadecimals are convenient, since 4 binary digits map directly to the set of hexadecimal digits. Thus  $367 = 101101111_2 = 557_8 = 16f_{16}$ . A list of the integers 0–63 in various bases is given in Table B.1.

Dec	Bin	Oct	Hex	Dec	Bin	Oct	Hex
0	0	0	0	32	100000	40	20
1	1	1	1	33	100001	41	21
2	10	2	2	34	100010	42	22
3	11	3	3	35	100011	43	23
4	100	4	4	36	100100	44	24
5	101	5	5	37	100101	45	25
6	110	6	6	38	100110	46	26
7	111	7	7	39	100111	47	27
8	1000	10	8	40	101000	50	28
9	1001	11	9	41	101001	51	29
10	1010	12	a	42	101010	52	2a
11	1011	13	b	43	101011	53	2b
12	1100	14	c	44	101100	54	2c
13	1101	15	d	45	101101	55	2d
14	1110	16	e	46	101110	56	2e
15	1111	17	f	47	101111	57	2f
16	10000	20	10	48	110000	60	30
17	10001	21	11	49	110001	61	31
18	10010	22	12	50	110010	62	32
19	10011	23	13	51	110011	63	33
20	10100	24	14	52	110100	64	34
21	10101	25	15	53	110101	65	35
22	10110	26	16	54	110110	66	36
23	10111	27	17	55	110111	67	37
24	11000	30	18	56	111000	70	38
25	11001	31	19	57	111001	71	39
26	11010	32	1a	58	111010	72	3a
27	11011	33	1b	59	111011	73	3b
28	11100	34	1c	60	111100	74	3c
29	11101	35	1d	61	111101	75	3d
30	11110	36	1e	62	111110	76	3e
31	11111	37	1f	63	111111	77	3f

**Table B.1** A list of the integers 0–63 in decimal, binary, octal, and hexadecimal.

## B.2 IEEE 754 Floating Point Standard

The set of real numbers, also called *reals*, includes all fractions and irrational numbers. It is infinite in size both in the sense that there is no largest nor smallest number,

and that between any 2 given numbers there are infinitely many numbers. Reals are widely used for calculation, but since any computer only has finite memory, there are infinitely many numbers which cannot be represent on a computer. Hence, any computation performed on a computer with reals must rely on approximations. *IEEE 754 double precision floating-point format (binary64)*, known as a *double*, is a standard for representing an approximation of reals using 64 bits. These bits are divided into 3 parts: sign, exponent and fraction,

$$s e_1 e_2 \dots e_{11} m_1 m_2 \dots m_{52},$$

where  $s$ ,  $e_i$ , and  $m_j$  are binary digits. The bits are converted to a number using the equation by first calculating the exponent  $e$  and the mantissa  $m$ ,

$$e = \sum_{i=1}^{11} e_i 2^{11-i}, \quad (\text{B.3})$$

$$m = \sum_{j=1}^{52} m_j 2^{-j}. \quad (\text{B.4})$$

I.e., the exponent is an integer, where  $0 \leq e < 2^{11}$ , and the mantissa is a rational, where  $0 \leq m < 1$ . For most combinations of  $e$  and  $m$ , the real number  $v$  is calculated as,

$$v = (-1)^s (1 + m) 2^{e-1023} \quad (\text{B.5})$$

with the exceptions that

	$m = 0$	$m \neq 0$
$e = 0$	$v = (-1)^s 0$ (signed zero)	$v = (-1)^s m 2^{1-1023}$ (subnormals)
$e = 2^{11} - 1$	$v = (-1)^s \infty$	$v = (-1)^s \text{NaN}$ (not-a-number)

where  $e = 2^{11} - 1 = 11111111111_2 = 2047$ . The largest and smallest number that is not infinity is thus

$$e = 2^{11} - 2 = 2046, \quad (\text{B.6})$$

$$m = \sum_{j=1}^{52} 2^{-j} = 1 - 2^{-52} \simeq 1, \quad (\text{B.7})$$

$$v_{\max} = \pm \left(2 - 2^{-52}\right) 2^{1023} \simeq \pm 2^{1024} \simeq \pm 10^{308}. \quad (\text{B.8})$$

The density of numbers varies in such a way that when  $e - 1023 = 52$ , then

$$v = (-1)^s \left( 1 + \sum_{j=1}^{52} m_j 2^{-j} \right) 2^{52} \quad (\text{B.9})$$

$$= \pm \left( 2^{52} + \sum_{j=1}^{52} m_j 2^{-j} 2^{52} \right) \quad (\text{B.10})$$

$$= \pm \left( 2^{52} + \sum_{j=1}^{52} m_j 2^{52-j} \right) \quad (\text{B.11})$$

$$\stackrel{k=52-j}{=} \pm \left( 2^{52} + \sum_{k=51}^0 m_{52-k} 2^k \right), \quad (\text{B.12})$$

which are all integers in the range  $2^{52} \leq |v| < 2^{53}$ . When  $e - 1023 = 53$ , then the same calculation gives

$$v \stackrel{k=53-j}{=} \pm \left( 2^{53} + \sum_{k=52}^1 m_{53-k} 2^k \right), \quad (\text{B.13})$$

which are every second integer in the range  $2^{53} \leq |v| < 2^{54}$ , and so on for larger values of  $e$ . When  $e - 1023 = 51$ , the same calculation gives,

$$v \stackrel{k=51-j}{=} \pm \left( 2^{51} + \sum_{k=50}^{-1} m_{51-k} 2^k \right), \quad (\text{B.14})$$

which is a distance between numbers of  $1/2$  in the range  $2^{51} \leq |v| < 2^{52}$ , and so on for smaller values of  $e$ . Thus we may conclude that the distance between numbers in the interval  $2^n \leq |v| < 2^{n+1}$  is  $2^{n-52}$ , for  $-1022 = 1 - 1023 \leq n < 2046 - 1023 = 1023$ . For subnormals, the distance between numbers is

$$v = (-1)^s \left( \sum_{j=1}^{52} m_j 2^{-j} \right) 2^{-1022} \quad (\text{B.15})$$

$$= \pm \left( \sum_{j=1}^{52} m_j 2^{-j} 2^{-1022} \right) \quad (\text{B.16})$$

$$= \pm \left( \sum_{j=1}^{52} m_j 2^{-j-1022} \right) \quad (\text{B.17})$$

$$\stackrel{k=-j-1022}{=} \pm \left( \sum_{j=-1023}^{-1074} m_{-k-1022} 2^k \right), \quad (\text{B.18})$$

which gives a distance between numbers of  $2^{-1074} \simeq 10^{-323}$  in the range  $0 < |v| < 2^{-1022} \simeq 10^{-308}$ .



## Appendix C

### Commonly Used Character Sets

Letters, digits, symbols, and space are the core of how we store data, write programs, and communicate with computers and each other. These symbols are in short called characters and represent a mapping between numbers, also known as codes, and a pictorial representation of the character. E.g., the ASCII code for the letter 'A' is 65. These mappings are for short called character sets, and due to differences in natural languages and symbols used across the globe, many different character sets are in use. E.g., the English alphabet contains the letters 'a' to 'z'. These letters are common to many other European languages which in addition use even more symbols and accents. For example, Danish has further the letters 'æ', 'ø', and 'å'. Many non-European languages have completely different symbols, where the Chinese character set is probably the most extreme, and some definitions contain 106,230 different characters, albeit only 2,600 are included in the official Chinese language test at the highest level.

Presently, the most common character set used is Unicode Transformation Format (UTF), whose most popular encoding schemes are 8-bit (UTF-8) and 16-bit (UTF-16). Many other character sets exist, and many of the later build on the American Standard Code for Information Interchange (ASCII). The ISO-8859 codes were an intermediate set of character sets that are still in use, but which is greatly inferior to UTF. Here we will briefly give an overview of ASCII, ISO-8859-1 (Latin1), and UTF.

#### C.1 ASCII

The *American Standard Code for Information Interchange (ASCII)* [7], is a 7 bit code tuned for the letters of the English language, numbers, punctuation symbols, control codes and space, see Tables C.1 and C.2. The first 32 codes are reserved for

x0+0x	00	10	20	30	40	50	60	70
00	NUL	DLE	SP	0	@	P	'	p
01	SOH	DC1	!	1	A	Q	a	q
02	STX	DC2	"	2	B	R	b	r
03	ETX	DC3	#	3	C	S	c	s
04	EOT	DC4	\$	4	D	T	d	t
05	ENQ	NAK	%	5	E	U	e	u
06	ACK	SYN	&	6	F	V	f	v
07	BEL	ETB	'	7	G	W	g	w
08	BS	CAN	(	8	H	X	h	x
09	HT	EM	)	9	I	Y	i	y
0A	LF	SUB	*	:	J	Z	j	z
0B	VT	ESC	+	;	K	[	k	{
0C	FF	FS	,	<	L	\	l	
0D	CR	GS	-	=	M	]	m	}
0E	SO	RS	.	>	N	^	n	~
0F	SI	US	/	?	O	_	o	DEL

**Table C.1** ASCII

non-printable control characters to control printers and similar devices or to provide meta-information. The meaning of each control character is not universally agreed upon.

The code order is known as *ASCIIbetical order*, and it is sometimes used to perform arithmetic on codes, e.g., an uppercase letter with code *c* may be converted to lower case by adding 32 to its code. The ASCIIbetical order also has a consequence for sorting, i.e., when sorting characters according to their ASCII code, 'A' comes before 'a', which comes before the symbol '{'.

## C.2 ISO/IEC 8859

The ISO/IEC 8859 report [http://www.iso.org/iso/catalogue_detail?csnumber=28245](http://www.iso.org/iso/catalogue_detail?csnumber=28245) defines 10 sets of codes specifying up to 191 codes and graphics characters using 8 bits. Set 1, also known as ISO/IEC 8859-1, Latin alphabet No. 1, or *Latin1*, covers many European languages and is designed to be compatible with ASCII, such that code for the printable characters in ASCII is the same in ISO 8859-1. Table C.3 shows the characters above 7e. Codes 00-1f and 7f-9f are undefined in ISO 8859-1.



Code	Description
NUL	Null
SOH	Start of heading
STX	Start of text
ETX	End of text
EOT	End of transmission
ENQ	Enquiry
ACK	Acknowledge
BEL	Bell
BS	Backspace
HT	Horizontal tabulation
LF	Line feed
VT	Vertical tabulation
FF	Form feed
CR	Carriage return
SO	Shift out
SI	Shift in
DLE	Data link escape
DC1	Device control one
DC2	Device control two
DC3	Device control three
DC4	Device control four
NAK	Negative acknowledge
SYN	Synchronous idle
ETB	End of transmission block
CAN	Cancel
EM	End of medium
SUB	Substitute
ESC	Escape
FS	File separator
GS	Group separator
RS	Record separator
US	Unit separator
SP	Space
DEL	Delete

Table C.2 ASCII symbols.

### C.3 Unicode

Unicode is a character standard defined by the Unicode Consortium, <http://unicode.org>, as the *Unicode Standard*. Unicode allows for 1,114,112 different codes. Each code is called a *code point* which represents an abstract character. Code points are divided into 17 planes, each with  $2^{16} = 65,536$  code points. Planes are further subdivided into named *blocks*. The first plane is called the *Basic Multilingual plane* and its block of the first 128 code points is called the *Basic Latin block* and is identical to ASCII, see Table C.1, and code points 128-255 are called the *Latin-1 Supplement block*, and are identical to the upper range of ISO 8859-1, see Table C.3. Each code-point has a number of attributes such as the *Unicode general category*. Presently more than 128,000 code points are defined as covering 135 modern and historical writing systems, and obtained

x0+0x	80	90	A0	B0	C0	D0	E0	F0
00			NBSP	°	À	Đ	à	đ
01			¡	±	Á	Ñ	á	ñ
02			¢	²	Â	Ò	â	ò
03			£	³	Ã	Ó	ã	ó
04			¤	´	Ä	Ö	ä	ö
05			¥	µ	Å	Õ	å	õ
06			¦	¶	Æ	Ö	æ	ö
07			§	·	Ç	×	ç	÷
08			¨	¸	È	Ø	è	ø
09			©	¹	É	Ù	é	ù
0a			ª	º	Ê	Ú	ê	ú
0b			«	»	Ë	Û	ë	û
0c			¬	$\frac{1}{4}$	Ì	Ü	ì	ü
0d			SHY	$\frac{1}{2}$	Í	Ý	í	ý
0e			®	$\frac{3}{4}$	Î	Þ	î	þ
0f			¯		Ï	ß	ï	ÿ

**Table C.3** ISO-8859-1 (latin1) non-ASCII part. Note that the codes 7f – 9f are undefined.

Code	Description
NBSP	Non-breakable space
SHY	Soft hyphen

**Table C.4** ISO-8859-1 special symbols.

at <http://www.unicode.org/Public/UNIDATA/UnicodeData.txt>, which includes the code point, name, and general category.

A Unicode code point is an abstraction from the encoding and the graphical representation of a character. A code point is written as “U+” followed by its hexadecimal number, and for the Basic Multilingual plane, 4 digits are used, e.g., the code point with the unique name LATIN CAPITAL LETTER A has the Unicode code point “U+0041”, and is in this text visualized as ‘A’. More digits are used for code points of the remaining planes.

The general category is used to specify valid characters that do not necessarily have a visual representation but possibly transform text. Some categories and their letters in the first 256 code points are shown in Table C.5.

To store and retrieve code points, they must be encoded and decoded. A common encoding is *UTF-8*, which encodes code points as 1 to 4 bytes, and which is backward-compatible with ASCII and ISO 8859-1. Hence, in all 3 coding systems, the character with code 65 represents the character ‘A’. Another popular encoding scheme is *UTF-16*, which encodes characters as 2 or 4 bytes, but which is not backward-compatible with ASCII or ISO 8859-1. UTF-16 is used internally in many compilers, interpreters, and operating systems.

General category	Code points	Name
Lu	U+0041–U+005A, U+00C0–U+00D6, U+00D8–U+00DE	Upper case letter
Ll	U+0061–U+007A, U+00B5, U+00DF–U+00F6, U+00F8–U+00FF	Lower case letter
Lt	None	Digraphic letter, with first part uppercase
Lm	None	Modifier letter
Lo	U+00AA, U+00BA	Gender ordinal indicator
Nl	None	Letterlike numeric character
Pc	U+005F	Low line
Mn	None	Nonspacing combining mark
Mc	None	Spacing combining mark
Cf	U+00AD	Soft Hyphen

**Table C.5** Some general categories for the first 256 code points.



## Appendix D

### Common Language Infrastructure

The *Common Language Infrastructure (CLI)*, not to be confused with *Command Line Interface* with the same acronym, is a technical standard developed by Microsoft [4, 3]. The standard specifies a language, its format, and a runtime environment that can execute the code. The main feature is that it provides a common interface between many languages and many platforms, such that programs can collaborate in a language-agnostic manner and can be executed on different platforms without having to be recompiled. Main features of the standard are:

Common Type System (CTS) which defines a common set of types that can be used across different languages as if it were their own.

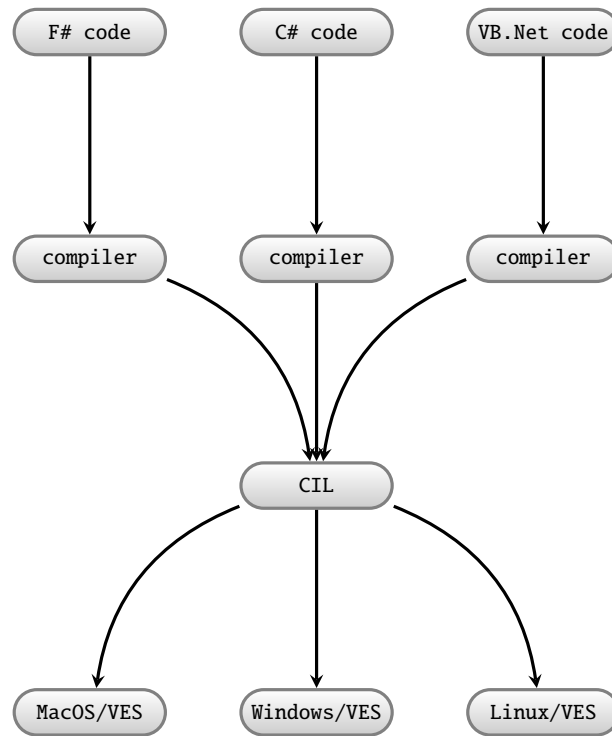
Metadata which defines a common method for referencing programming structures such as values and functions in a language-independent manner.

Common Intermediate Language (CIL) which is a platform-independent, stack-based, object-oriented assembly language that can be executed by the Virtual Execution System.

Virtual Execution System (VES) which is a platform dependent, virtual machine, which combines the above into code that can be executed at runtime. Microsoft's implementation of VES is called *Common Language Runtime (CLR)* and uses *just-in-time* compilation. In this book, we have been using the `mono` command.

The process of running an F# program is shown in Figure D.1. First the F# code is compiled or interpreted to CIL. This code possibly combined with other CIL code is then converted to a machine-readable code, and the result is then executed on the platform.

CLI defines a *module* as a single file containing executable code by VES. Hence, CLI's notion of a module is somewhat related to F#'s notion of module, but the



**Fig. D.1** The relation between some .NET/Mono languages with the Common intermediate language (CIL), and the Virtual execution systems (VES) on some operating system (mono).

two should not be confused. A collection of modules, a *manifest*, and possibly other resources, which jointly define a complete program is called an *assembly*. The manifest is the description of which files are included in the assembly together with its version, name, security information, and other bookkeeping information.

## References

1. Alonzo Church. An unsolvable problem of elementary number theory. *American Journal of Mathematics*, 58:345—363, 1936.
2. Ole-Johan Dahl and Kristen Nygaard. SIMULA a language for programming and description of discrete event systems. introduction and user's manual. Technical report, Norwegian Computing Center, 1967.
3. European Computer Manufacturers Association (ECMA). Standard ecma-335, common language infrastructure (cli). <http://www.ecma-international.org/publications/standards/Ecma-335.htm>.
4. International Organization for Standardization. Iso/iec 23271:2012, common language infrastructure (cli). <https://www.iso.org/standard/58046.html>.
5. Programming Research Group. Specifications for the ibm mathematical formula translating system, fortran. Technical report, Applied Science Division, International Business Machines Corporation, 1954.
6. John McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, 1960.
7. X3: ASA Sectional Committee on Computers and Information Processing. American standard code for information interchange. Technical Report ASA X3.4-1963, American Standards Association (ASA), 1963. <http://worldpowersystems.com/projects/codes/X3.4-1963/>.
8. George Pólya. *How to solve it*. Princeton University Press, 1945.
9. Alan M. Turing. On computable numbers, with an application to the entscheidungsproblem. *Proceedings of the London Mathematical Society*, s2-42(1):230–265, 1936.





# Index

- `(**)`, 83
- `->`, 52
- `.`, 72
- `//`, 83
- `:`, 45, 46
- `;`, 45
- `_`, 46
- `abs`, 28
- `acos`, 28
- `asin`, 28
- `atan2`, 28
- `atan`, 28
- `ceil`, 28
- `cosh`, 28
- `cos`, 28
- `exp`, 28
- `floor`, 28
- `log10`, 28
- `log`, 28
- `max`, 28
- `min`, 28
- `pown`, 28
- `round`, 28
- `sign`, 28
- `sinh`, 28
- `sin`, 28
- `sqrt`, 28
- `tanh`, 28
- `tan`, 28
- `_`, 75
- `()`, 9, 11
- `[]`, 37
- American Standard Code for Information Interchange, 103
- `and`, 31
- anonymous functions, 52
- anonymous variable type, 75
- ASCII, 103
- ASCIIbetical order, 36, 104
- assembly, 110
- base, 22, 97
- `bash`, 92
- Basic Latin block, 105
- Basic Multilingual plane, 105
- basic types, 22
- binary number, 23, 97
- binary operator, 28
- `binary64`, 99
- bit, 23, 97
- `bool`, 22
- branch, 15, 56, 58
- byte, 97
- `byte[]`, 24
- byte, 24
- call-back function, 4
- cast, 15
- `cd`, 89, 93
- `char`, 22, 24
- character, 24
- CIL, 109
- class, 26, 37
- CLI, 109
- closure, 54
- CLR, 109
- code block, 47
- code point, 24, 105
- Command Line Interface, 109
- comments, 19
- Common Intermediate Language, 109
- Common Language Infrastructure, 109
- Common Language Runtime, 109

- Common Type System, 109
- compile mode, 9
- console, 87
- CTS, 109
  
- debugging, 10
- decimal, 24
- decimal number, 22, 97
- decimal point, 22, 97
- declarative programming, 3
- del, 91
- digit, 22, 97
- dir, 89
- directory, 87
- do, 55
- do-binding, 9, 55
- dot notation, 37
- double, 99
- double, 24
- downcast, 27
- dynamic scope, 47
  
- echo, 90, 94
- elif, 57
- else, 57
- encapsulate, 13
- encapsulation, 49, 54
- environment, 60
- error message, 12
- escape sequences, 24
- event-driven programming, 4
- exception, 34
- exclusive or, 34
- executable file, 11
- exn, 22
- expression, 3, 8, 27, 45
- Extensible Markup Language, 83
  
- first-class citizenship, 54
- float, 22
- float32, 24
- floating point number, 22
- folder, 87
- formatting string, 9
- fractional part, 22, 27
- fst, 67
- fun, 52
- function, 3, 9
- function body, 49
- functional programming, 3
  
- generic function, 51
- graphical user interface, 87
- guard, 57
  
- GUI, 87
  
- hexadecimal number, 23, 98
  
- identifier, 42
- IEEE 754 double precision floating-point  
format, 99
- if, 57
- imperative programming, 4
- implementation file, 11
- indentation, 15
- infix notation, 28
- inline, 76
- input pattern, 56
- int, 22
- int16, 24
- int32, 24
- int64, 24
- int8, 24
- integer, 22
- integer division, 33
- integer remainder, 33
- interactive mode, 9
- it, 11, 22
  
- just-in-time, 109
  
- keyword, 8, 44
  
- Latin-1 Supplement block, 105
- Latin1, 104
- least significant bit, 97
- length, 66
- let, 46
- let-binding, 8
- lexeme, 8
- lexical scope, 47, 51
- library file, 11
- lightweight syntax, 45
- literal, 22
- literal type, 24
- lower camel case, 44
- ls, 92
  
- manifest, 110
- member, 26, 66
- Metadata, 109
- method, 37
- mixed case, 44
- mkdir, 90, 93
- module, 109
- most significant bit, 97
- move, 90
- mutable, 17

- mv, 94
- namespace, 26
- NaN, 99
- nested scope, 47
- newline, 24
- not, 31
- not-a-number, 99
- obfuscation, 68
- obj, 22
- object, 4, 37
- object-oriented programming, 4
- octal number, 23, 98
- operand, 27, 50
- operator, 8, 27, 30, 50
- option type, 70
- or, 31
- overflow, 32
- pascal case, 44
- pipng, 53
- precedence, 29, 30
- prefix operator, 29
- primitive types, 66
- printfn, 9
- procedure, 54
- reals, 98
- rec, 17
- recursion, 28
- redirection, 90, 94
- rm, 94
- rmdir, 91, 94
- rounding, 27
- runtime, 15
- runtime error, 34
- runtime resolved variable type, 75
- sbyte, 24
- scientific notation, 23
- scope, 47
- script file, 11
- search path, 91, 94
- side-effect, 54
- signature file, 11
- single, 24
- snd, 67
- source code, 11
- state, 4
- statement, 4, 9, 55
- statically resolved variable type, 75
- stdin, 14
- stream, 14
- string, 14, 24
- string, 22
- struct records, 72
- structured programming, 4
- subnormals, 99
- terminal, 87
- The Heap, 69, 72
- then, 57
- truth table, 31
- tuple, 66
- type, 10, 22
- type abbreviation, 74
- type aliasing, 74
- type constraints, 76
- type declaration, 11
- type inference, 10, 11
- type safety, 50
- typecasting, 26
- uint16, 24
- uint32, 24
- uint64, 24
- uint8, 24
- underflow, 32
- Unicode, 24
- unicode block, 105
- Unicode general category, 105
- Unicode Standard, 105
- unit, 22
- upcast, 27
- upper camel case, 44
- UTF-16, 106
- UTF-8, 106
- value-binding, 45
- variable, 17
- variable types, 75
- verbatim, 26
- verbose syntax, 45
- VES, 109
- Virtual Execution System, 109
- when, 76
- while, 17
- whitespace, 24
- whole part, 22, 27
- wildcard, 15, 56
- wildcard pattern, 46
- Windows command line, 87
- with, 73
- word, 97
- XML-standard, 83
- xor, 34