

Learning to program with F#

Jon Spurring

September 14, 2016

Contents

1	Preface	5
2	Introduction	6
2.1	How to learn to program	6
2.2	How to solve problems	7
2.3	Approaches to programming	7
2.4	Why use F#	8
2.5	How to read this book	9
I	F# basics	10
3	Executing F# code	11
3.1	Source code	11
3.2	Executing programs	11
4	Quick-start guide	14
5	Using F# as a calculator	19
5.1	Literals and basic types	19
5.2	Operators on basic types	24
5.3	Boolean arithmetic	26
5.4	Integer arithmetic	27
5.5	Floating point arithmetic	30
5.6	Char and string arithmetic	31
5.7	Programming intermezzo	33

6	Constants, functions, and variables	34
6.1	Values	37
6.2	Non-recursive functions	41
6.3	User-defined operators	45
6.4	The Printf function	47
6.5	Variables	49
7	In-code documentation	54
8	Controlling program flow	59
8.1	For and while loops	59
8.2	Conditional expressions	63
8.3	Programming intermezzo	65
8.4	Recursive functions	65
9	Ordered series of data	68
9.1	Tuples	69
9.2	Lists	71
9.3	Arrays	73
10	Testing programs	79
10.1	White-box testing	81
10.2	Back-box testing	84
10.3	Debugging by tracing	87
11	Exceptions	96
12	Input and Output	102
12.1	Interacting with the console	103
12.2	Storing and retrieving data from a file	104
12.3	Working with files and directories.	107
12.4	Programming intermezzo	108

II	Imperative programming	110
13	Graphical User Interfaces	112
14	Imperative programming	113
14.1	Introduction	113
14.2	Generating random texts	115
14.2.1	0'th order statistics	115
14.2.2	1'th order statistics	117
III	Declarative programming	118
15	Sequences and computation expressions	119
15.1	Sequences	119
16	Patterns	124
16.1	Pattern matching	124
17	Types and measures	127
17.1	Unit of Measure	127
18	Functional programming	131
IV	Structured programming	134
19	Namespaces and Modules	135
20	Object-oriented programming	137
V	Appendix	138
A	Number systems on the computer	139
A.1	Binary numbers	139
A.2	IEEE 754 floating point standard	139

B	Commonly used character sets	143
B.1	ASCII	143
B.2	ISO/IEC 8859	144
B.3	Unicode	144
3	A brief introduction to Extended Backus-Naur Form	148
D	F_b	151
E	Language Details	156
E.1	Arithmetic operators on basic types	156
E.2	Basic arithmetic functions	159
E.3	Precedence and associativity	160
E.4	Lightweight Syntax	162
F	The Some Basic Libraries	163
F.1	System.String	163
F.2	List, arrays, and sequences	170
F.3	Mutable Collections	171
F.3.1	Mutable lists	171
F.3.2	Stacks	171
F.3.3	Queues	171
F.3.4	Sets and dictionaries	171
	Bibliography	172
	Index	173

Chapter 6

Constants, functions, and variables

In the previous chapter, we saw how to use F# as a calculator working with literals, operators and built-in functions. To save time and make programs easier to read and debug, it is useful to bind expressions to identifiers either as new constants, functions, or operators. As an example, consider the problem,

Problem 6.1:

For given set constants a , b , and c , solve for x in

$$ax^2 + bx + c = 0 \tag{6.1}$$

To solve for x we use the quadratic formula from elementary algebra,

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}, \tag{6.2}$$

which gives the general solution for any values of the coefficients. Here, we will assume a positive discriminant, $b^2 - 4ac > 0$. In order to write a program, where the code may be reused later, we define a function `discriminant : float -> float -> float`, that is, a function that takes 3 arguments, `a`, `b`, and `c`, and calculates the discriminant. Details on function definition is given in Section 6.2. Likewise, we will define functions `positiveSolution : float -> float -> float` and `negativeSolution : float -> float -> float`, that also takes the polynomial's coefficients as arguments and calculates the solution corresponding to choosing the positive and negative sign for \pm in the equation. Our solution thus looks like Listing 6.1.

Listing 6.1, identifiersExample.fsx:**Finding roots for quadratic equations using function name binding.**

```
let discriminant a b c = b ** 2.0 - 4.0 * a * c
let positiveSolution a b c = (-b + sqrt (discriminant a b c)) / (2.0 * a)
let negativeSolution a b c = (-b - sqrt (discriminant a b c)) / (2.0 * a)

let a = 1.0
let b = 0.0
let c = -1.0
let d = discriminant a b c
let xp = positiveSolution a b c
let xn = negativeSolution a b c
printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
printfn "    has discriminant %A and solutions %A and %A" d xn xp

-----
0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
    has discriminant 4.0 and solutions -1.0 and 1.0
```

Here, we have further defined names of values `a`, `b`, and `c` used as input to our functions, and the results of function application is bound to the names `d`, `xn`, and `xp`. The names of functions and values given here are examples of identifiers, and with these, we may reuse the quadratic formulas and calculated values later, while avoiding possible typing mistakes and reducing amount of code, which needs to be debugged.

Before we begin a deeper discussion note that F# adheres to two different syntax: regular and *lightweight*. In the regular syntax, newlines and whitespaces are generally ignored, while in lightweight syntax, certain keywords and lexemes may be replaced by specific use of newlines and whitespaces. Lightweight syntax is the most common, but the syntaxes may be mixed, and we will highlight the options, when relevant.

· lightweight
syntax

The use of identifiers is central in programming. For F# not to be confused by built-in functionality, identifiers must follow a specific grammar: An identifier must start with a letter, but can be followed by zero or more of letters, digits, and a range of special characters except SP, LF, and CR (space, line feed, and carriage return). An identifier must not be a keyword or a reserved-keyword listed in Figures 6.1. An identifier is a name for a constant, an expression, or a type, and it is defined by the following EBNF:

Keywords:

`abstract`, `and`, `as`, `assert`, `base`, `begin`, `class`, `default`, `delegate`, `do`, `done`, `downcast`, `downto`, `elif`, `else`, `end`, `exception`, `extern`, `false`, `finally`, `for`, `fun`, `function`, `global`, `if`, `in`, `inherit`, `inline`, `interface`, `internal`, `lazy`, `let`, `match`, `member`, `module`, `mutable`, `namespace`, `new`, `null`, `of`, `open`, `or`, `override`, `private`, `public`, `rec`, `return`, `sig`, `static`, `struct`, `then`, `to`, `true`, `try`, `type`, `upcast`, `use`, `val`, `void`, `when`, `while`, `with`, and `yield`.

Reserved keywords for possible future use:

`atomic`, `break`, `checked`, `component`, `const`, `constraint`, `constructor`, `continue`, `eager`, `fixed`, `fori`, `functor`, `include`, `measure`, `method`, `mixin`, `object`, `parallel`, `params`, `process`, `protected`, `pure`, `recursive`, `sealed`, `tailcall`, `trait`, `virtual`, and `volatile`.

Symbolic keywords:

`let!`, `use!`, `do!`, `yield!`, `return!`, `|`, `->`, `<-`, `..`, `:`, `(,)`, `[,]`, `[<, >]`, `[|, |]`, `{, }`, `'`, `#`, `:?>`, `:?`, `:>`, `..`, `::`, `:=`, `;;`, `;`, `=`, `_`, `?`, `??`, `(*)`, `<@`, `@>`, `<@@`, and `@@>`.

Reserved symbolic keywords for possible future:

`~` and ```.

Figure 6.1: List of (possibly future) keywords and symbolic keywords in F#.

Listing 6.2: to do

```
ident = (letter | "_" ) {letter | dDigit | specialChar};
longIdent = ident | ident "." longIdent; (*no space around "."*)

longIdentOrOp = [longIdent "."] identOrOp; (*no space around "."*)
identOrOp =
    ident
    | "(" infixOp | prefixOp ")"
    | "(*)";

dDigit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9";
letter = Lu | Ll | Lt | Lm | Lo | Nl; (*e.g. "A", "B" ... and "a", "b", ...*)
specialChar = Pc | Mn | Mc | Cf; (*e.g., "_"*)

codePoint = ?Any unicode codepoint?;
Lu = ?Upper case letters?;
Ll = ?Lower case letters?;
Lt = ?Digraphic letters, with first part uppercase?;
Lm = ?Modifier letters?;
Lo = ?Gender ordinal indicators?;
Nl = ?Letterlike numeric characters?;
Pc = ?Low lines?;
Mn = ?Nonspacing combining marks?;
Mc = ?Spacing combining marks?;
Cf = ?Soft Hyphens?;
```

Thus, examples of identifiers are `a`, `theCharacter9`, `Next_Word`, `_tok`. Typically, only letters from the english alphabet are used as `letter`, and only `_` is used for `specialChar`, but the full definition refers to the Unicode general categories described in Appendix B.3, and there are currently 19.345 possible Unicode code points in the `letter` category and 2.245 possible Unicode code points in the `specialChar` category.

Expressions are a central concept in F#. An expression can be a mathematical expression, such as `3*5`,

a function application, such as $f3$, and many other things. Central in this chapter is the binding of values and functions to identifiers, which is done with the keyword `let`, using the following simplified syntax, e.g., `let a = 1.0`.

Expressions has an enormous variety in how they may be written, we will in this book gradually work through some of the more important facets. For this we will extend the EBNF notation with ellipses: `...`, to denote that what is shown is part of the complete EBNF production rule. E.g., the part of expressions, we will discuss in this chapter is specified in EBNF by,

Listing 6.3: to do

```
expr = ...
| expr ":" type (*type annotation*)
| expr ";" expr (*sequence of expressions*)
| "let" valueDefn "in" expr (*binding a value or variable*)
| "let" ["rec"] functionDefn "in" expr (*binding a function or operator*)
| "fun" argumentPats "->" expr (*anonymous function*)
| expr "<-" expr (*assingment*)

type = ...
| longIdent (*named such as "int"*)

valueDefn = ["mutable"] pat "=" expr;

pat = ...
| "_" (*wildcard*)
| ident (*named*)
| pat ":" type (*type constraint*)
| "(" pat ")" (*paranthesized*)

functionDefn = identOrOp argumentPats [":" type] "=" expr;
argumentPats = pat | pat argumentPats;
```

In the following sections, we will work through this bit by bit.

6.1 Values

Binding identifiers to literals or expressions that are evaluated to be values, is called value-binding, and examples are `let a = 3.0` and `let b = cos 0.9`. On EBNF the simplified syntax,

Listing 6.4: to do

```
expr = ...
| "let" valueDefn "in" expr (*binding a value or variable*)
```

The `let` bindings defines relations between patterns `pat` and expressions `expr` for many different purposes. Most often the pattern is an identifier `ident`, which `let` defines to be an alias of the expression `expr`. The pattern may also be defined to have specific type using the `:` lexeme and a named type. The `_` pattern is called the *wild card* pattern and, when it is in the value-binding, then the expression is evaluated but the result is discarded. The binding may be mutable as indicated by the keyword `mutable`, which will be discussed in Section 6.5, and the binding holds *lexically* for the last expression as indicated by the `in` keyword. For example, letting the identifier `p` be bound to the value 2.0 and using it in an expression is done as follows,

- `let`
- `:`
- wild card
- `mutable`
- lexically
- `in`

Listing 6.5, letValue.fsx:

The identifier `p` is used in the expression following the `in` keyword.

```
let p = 2.0 in printfn "%A" (3.0 ** p)

9.0
```

F# will ignore most newlines between lexemes, i.e., the above is equivalent to writing,

Listing 6.6, letValueLF.fsx:

Newlines after `in` make the program easier to read.

```
let p = 2.0 in
printfn "%A" (3.0 ** p)

9.0
```

F# also allows for an alternative notation called *lightweight syntax*, where e.g., the `in` keyword is replaced with a newline, and the expression starts on the next line at the same column as `let` starts in, i.e., the above is equivalent to

· lightweight
syntax

Listing 6.7, letValueLightWeight.fsx:

Lightweight syntax does not require the `in` keyword, but expression must be aligned with the `let` keyword.

```
let p = 2.0
printfn "%A" (3.0 ** p)

9.0
```

The same expression in interactive mode will also respond the inferred types, e.g.,

Listing 6.8, letValueLightWeightTypes.fsx:

Interactive mode also responds inferred types.

```
> let p = 2.0
- printfn "%A" (3.0 ** p);;
9.0

val p : float = 2.0
val it : unit = ()
```

By the `val` keyword in the line `val p : float = 2.0` we see that `p` is inferred to be of type `float` and bound to the value `2.0`. The inference is based on the type of the right-hand-side, which is of type `float`. Identifiers may be defined to have a type using the `:` lexeme, but the types on the left-hand-side and right-hand-side of the `=` lexeme must be identical. I.e., mixing types gives an error,

Listing 6.9, letValueTypeError.fsx:
Binding error due to type mismatch.

```
let p : float = 3
printfn "%A" (3.0 ** p)

/Users/sporring/repositories/fsharpNotes/src/letValueTypeError.fsx(1,17):
error FS0001: This expression was expected to have type
float
but here has type
int
```

Here, the left-hand-side is defined to be an identifier of type float, while the right-hand-side is a literal of type integer.

An expression can be a sequence of expressions separated by the lexeme `;`, e.g.,

Listing 6.10, letValueSequence.fsx:
A value-binding for a sequence of expressions.

```
let p = 2.0 in printfn "%A" p; printfn "%A" (3.0 ** p)

2.0
9.0
```

The lightweight syntax automatically inserts the `;` lexeme at newlines, hence using the lightweight syntax the above is the same as,

Listing 6.11, letValueSequenceLightWeight.fsx:
A value-binding for a sequence using lightweight syntax.

```
let p = 2.0
printfn "%A" p
printfn "%A" (3.0 ** p)

2.0
9.0
```

A key concept of programming is *scope*. In F#, the scope of a value-binding is lexically meaning that when F# determines the value bound to a name, it looks left and upward in the program text for the `let` statement defining it, e.g., · scope

Listing 6.12, letValueScopeLower.fsx:
Redefining identifiers is allowed in lower scopes.

```
let p = 3 in let p = 4 in printfn " %A" p;

4
```

F# also has to option of using dynamic scope, where the value of a binding is defined by when it is used, and this will be discussed in Section 6.5.

Scopes are given levels, and scopes may be nested, where the nested scope has a level one lower than its parent.¹ F# distinguishes between the top and lower levels, and at the top level in the lightweight syntax, redefining values is not allowed, e.g.,

Listing 6.13, letValueScopeLowerError.fsx:

Redefining identifiers is not allowed in lightweight syntax at top level.

```
let p = 3
let p = 4
printfn "%A" p;
```

```
/Users/sporring/repositories/fsharpNotes/src/letValueScopeLowerError.fsx
(2,5): error FS0037: Duplicate definition of value 'p'
```

But using parentheses, we create a *block*, i.e., a *nested scope*, and then redefining is allowed, e.g.,

· block
· nested scope

Listing 6.14, letValueScopeBlockAlternative3.fsx:

A block may be created using parentheses.

```
(
  let p = 3
  let p = 4
  printfn "%A" p
)
```

```
4
```

In both cases we used indentation, which is good practice, but not required here. Bindings inside are not available outside the nested scope, e.g.,

Listing 6.15, letValueScopeNestedScope.fsx:

Bindings inside a scope are not available outside.

```
let p = 3
(
  let q = 4
  printfn "%A" q
)
printfn "%A %A" p q
```

```
/Users/sporring/repositories/fsharpNotes/src/letValueScopeNestedScope.fsx
(6,19): error FS0039: The value or constructor 'q' is not defined
```

Nesting is a natural part of structuring code, e.g., through function definitions to be discussed in Section 6.2 and flow control structures to be discussed in Chapter 8. Blocking code by nesting is a key concept for making robust code that is easy to use by others without the user necessarily needing to know the details of the inner workings of a block of code.

Defining blocks is useful for controlling the extend of a lexical scope of bindings. For example, adding a second `printfn` statement,

¹Todo: Drawings would be good to describe scope

Listing 6.16, letValueScopeBlockProblem.fsx:
Overshadowing hides the first binding.

```
let p = 3 in let p = 4 in printfn "%A" p; printfn "%A" p
```

```
4
4
```

will print the value 4 last bound to the identifier `p`, since F# interprets the above as `let p = 3 in let p = 4 in (printfn "%A" p; printfn "%A" p)`. Had we intended to print the two different values of `p`, then we should have created a block as,

Listing 6.17, letValueScopeBlock.fsx:
Blocks allow for the return to the previous scope.

```
let p = 3 in (let p = 4 in printfn "%A" p); printfn "%A" p;
```

```
4
3
```

Here, the lexical scope of `let p = 4 in ...` is for the nested scope, which ends at `)`, returning to the lexical scope of `let p = 3 in`

6.2 Non-recursive functions

A function is a mapping between an input and output domain. A key advantage of using functions, when programming, is that they *encapsulate code* into smaller units, that are easier to debug and may be reused. F# is a functional first programming language, and offers a number of alternative methods for specifying parameters, which will be discussed in this section. Binding identifiers to functions follows a syntax similar to value-binding,

· encapsulate
code

Listing 6.18: to do

```
expr = ...
| "let" functionDefn "in" expr (*binding a function or operator*)
```

Functions may also be recursive, which will be discussed in Chapter 8. An example in interactive mode is,

Listing 6.19, letFunction.fsx:
An example of a binding of an identifier and a function.

```
> let sum (x : float) (y : float) : float = x + y in
- let c = sum 357.6 863.4 in
- printfn "%A" c;;
1221.0

val sum : x:float -> y:float -> float
val c : float = 1221.0
val it : unit = ()
```

and we see that the function is interpreted to have the type `val sum : x:float -> y:float ->`

`float`. The `->` lexeme means a mapping between sets, in this case floats. The function is also a higher order function, to be discussed in detail below, and here it suffices to think of `sum` as a function that takes 2 floats as argument and returns a float.

Not all types need to be declared, just sufficient for F# to be able to infer the types for the full statement. In the example, one specification is sufficient, and we could just have specified the type of the result,

Listing 6.20: All types need most often not be specified.

```
let sum x y : float = x + y
```

or even just one of the arguments,

Listing 6.21: Just one type is often enough for F# to infer the rest.

```
let sum (x : float) y = x + y
```

In both cases, since the `+` operator is only defined for *operands* of the same type, then when the type of either the result, any or both operands are declared, then the type of the remaining follows directly. As for values, lightweight syntax automatically inserts the keyword `in` and the lexeme `;`,

Listing 6.22, letFunctionLightWeight.fsx:
Lightweight syntax for function definitions.

```
let sum x y : float = x + y
let c = sum 357.6 863.4
printfn "%A" c

1221.0
```

Arguments need not always be inferred to types, but may be of generic type, which F# prefers, when *type safety* is ensured, e.g.,

Listing 6.23, functionDeclarationGeneric.fsx:
Typesafety implies that a function will work for any type, and hence it is generic.

```
> let second x y = y
- let a = second 3 5
- printfn "%A" a
- let b = second "horse" 5.0
- printfn "%A" b;;
5
5.0

val second : x:'a -> y:'b -> 'b
val a : int = 5
val b : float = 5.0
val it : unit = ()
```

Here, the function `second` does not use the first argument `x`, which therefore can be of any type, and which F# therefore calls `'a`, and the type of the second element, `y`, can also be of any type and not necessarily the same as `x`, so it is called `'b`. Finally the result is the same type as `y`, whatever it is. This is an example of a *generic function*,

A function may contain a sequence of expressions, but must return a value. E.g., the quadratic formula may be written as,

Listing 6.24, identifiersExampleAdvance.fsx:
A function may contain sequences of expressions.

```
let solution a b c sgn =
    let discriminant a b c =
        b ** 2.0 - 2.0 * a * c
    let d = discriminant a b c
    (-b + sgn * sqrt d) / (2.0 * a)

let a = 1.0
let b = 0.0
let c = -1.0
let xp = solution a b c +1.0
let xn = solution a b c -1.0
printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
printfn "  has solutions %A and %A" xn xp
```

```
0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
  has solutions -0.7071067812 and 0.7071067812
```

Here, we used the lightweight syntax, where the `=` identifies the start of a nested scope, and `F#` identifies the scope by indentation. The amount of space used for indentation does not matter, but all lines following the first must use the same. The scope ends before the first line with the previous indentation or none. Notice how the last expression is not bound to an identifier, but is the result of the function, i.e., in contrast to many other languages, `F#` does not have an explicit keyword for returning values, but requires a final expression, which will be returned to the caller of the function. Note also that since the function `discriminant` is defined in the nested scope of `solution`, then `discriminant` cannot be called outside `solution`, since the scope ends before `let a = 1.0`.

Lexical scope and function definitions can be a cause of confusion as the following example shows,²

· lexical scope

Listing 6.25, lexicalScopeNFunction.fsx:
Lexical scope means that $f(z) = 3x$ and not $4x$ at the time of calling.

```
let testScope x =
    let a = 3.0
    let f z = a * z
    let a = 4.0
    f x
printfn "%A" (testScope 2.0)
```

```
6.0
```

Here, the value-binding for `a` is redefined, after it has been used to define a helper function `f`. So which value of `a` is used, when we later apply `f` to an argument? To resolve the confusion, remember that value-binding is lexically defined, i.e., the binding `let f z = a * x` uses the value of `a`, it has by the ordering of the lines in the script, not dynamically by when `f` was called. Hence, **think of lexical scope as substitution of an identifier with its value or function immediately at the place of definition**. I.e., since `a` and `3.0` are synonymous in the first lines of the program, then the function `f` is really defined as, `let f z = 3.0 * x`.

Advice

²Todo: Add a drawing or possibly a spell-out of lexical scope here.

Functions do not need a name, but may be declared as an *anonymous function* using the `fun` keyword and the `->` lexeme,

· anonymous function

Listing 6.26, functionDeclarationAnonymous.fsx:
Anonymous functions are functions as values.

```
let first = fun x y -> x
printfn "%d" (first 5 3)
```

5

Here, a name is bound to an anonymous function, which returns the first of two arguments. The difference to `let first x y = x` is that anonymous functions may be treated as values, meaning that they may be used as arguments to other functions, and new values may be reassigned to their identifiers, when mutable, as will be discussed in Section 6.5. A common use of anonymous functions is as arguments to other functions, e.g.,

Listing 6.27, functionDeclarationAnonymousAdvanced.fsx:
Anonymous functions are often used as arguments for other functions.

```
let apply f x y = f x y
let mul = fun a b -> a * b
printfn "%d" (apply mul 3 6)
```

18

Note that here `apply` is given 3 arguments, the function `mul` and 2 integers. It is not given the result of `mul 3 6`, since that would not match the definition of `apply`. **Anonymous functions and functions as arguments are powerful concepts, but tend to make programs harder to read, and their use should be limited.**

Advice

Functions may be declared from other functions

Listing 6.28, functionDeclarationCurrying.fsx:

```
let mul x y = x*y
let timesTwo = mul 2.0
printfn "%g" (mul 5.0 3.0)
printfn "%g" (timesTwo 3.0)
```

15

6

Here, `mul 2.0` is a partial specification of the function `mul x y`, where the first argument is fixed, and hence, `timesTwo` is a function of 1 argument being the second argument of `mul`. This notation is called *currying* in tribute of Haskell Curry, and Currying is often used in functional programming, but generally **currying should be used carefully, since currying may seriously reduce readability of code.**

· currying
Advice

A *procedure* is a generalisation of the concept of functions, and in contrast to functions procedures need not return values,

· procedure

Listing 6.29, procedure.fsx:

A procedure is a function that has no return value, which in F# implies() as return value.

```
let printIt a = printfn "This is '%A'" a
printIt 3
printIt 3.0
```

```
-----
This is '3'
This is '3.0'
```

In F# this is automatically given the unit type as return value. Procedural thinking is useful for *encapsulation* of scripts, but is prone to *side-effects* and should be minimized by being replaced by functional thinking. More on side-effects in Section 6.5. **Procedural thinking is useful for encapsulation, but is prone to side-effects and should be minimized by being replaced by functional thinking.**

· encapsulation
· side-effects
Advice

6.3 User-defined operators

Operators are functions, and in F#, the infix multiplication operator `*` is equivalent to the function `(*)`, e.g.,

Listing 6.30, addOperatorNFunction.fsx:

```
let a = 3.0
let b = 4.0
let c = a + b
let d = (+) a b
printfn "%A plus %A is %A and %A" a b c d
```

```
-----
3.0 plus 4.0 is 7.0 and 7.0
```

All operator has this option, and you may redefine them and define your own operators, but in F# names of user-defined operators are limited by the following simplified EBNF:

Listing 6.31: Grammar for infix and prefix lexemes

```
infixOrPrefixOp = "+" | "-" | "+." | "-." | "%" | "&" | "&&";
prefixOp = infixOrPrefixOp | "~" {"~"} | "!" {opChar} - "!=";
infixOp =
  {"."} (
    infixOrPrefixOp
    | "-" {opChar}
    | "+" {opChar}
    | "||"
    | "<" {opChar}
    | ">" {opChar}
    | "="
    | " |" {opChar}
    | "&" {opChar}
    | "^" {opChar}
    | "*" {opChar}
    | "/" {opChar}
    | "%" {opChar}
    | "!=" )
  | ":@" | "::" | "$" | "?";
opChar =
  "!" | "%" | "&" | "*" | "+" | "-" | "." | "/"
  | "<" | "=" | ">" | "@" | "^" | "|" | "~";
```

The precedence rules and associativity of user-defined operators follows the rules for which they share prefixes with built-in rules, see Table E.6. E.g., `.*`, `+++`, and `<+` are valid operator names for infix operators, they have precedence as ordered, and their associativity are all left. Using `~` as the first character in the definition of an operator makes the operator unary and will not be part of the name. Examples of definitions and use of operators are,

Listing 6.32, operatorDefinitions.fsx:

```
let (.*) x y = x * y + 1
printfn "%A" (3 .* 4)
let (+++) x y = x * y + y
printfn "%A" (3 +++ 4)
let (<+) x y = x < y + 2.0
printfn "%A" (3.0 <+ 4.0)
let (~+.) x = x+1
printfn "%A" (+.1)
```

```
13
16
true
2
```

Operators beginning with `*` must use a space in its definition, (`*` in order for it not to be confused with the beginning of a comment (`*` , see Chapter 7 for more on comments in code).

Beware, redefining existing operators lexically redefines all future uses of the operators for all types, hence **it is not a good idea to redefine operators, but better to define new**. In Chapter 20 we will discuss how to define type specific operators including prefix operators.

Advice

Specifier	Type	Description
%b	bool	Replaces with boolean value
%s	string	
%c	char	
%d, %i	basic integer	
%u	basic unsigned integers	
%x	basic integer	formatted as unsigned hexadecimal with lower case letters
%X	basic integer	formatted as unsigned hexadecimal with upper case letters
%o	basic integer	formatted as unsigned octal integer
%f, %F,	basic floats	formatted on decimal form
%e, %E,	basic floats	formatted on scientific form. Lower case uses "e" while upper case uses "E" in the formatting.
%g, %G,	basic floats	formatted on the shortest of the corresponding decimal or scientific form.
%M	decimal	
%O	Objects ToString method	
%A	any built-in types	Formatted as a literal type
%a	Printf.TextWriterFormat -> 'a -> ()	
%t	(Printf.TextWriterFormat -> ()	

Table 6.1: Printf placeholder string

6.4 The Printf function

A common way to output information to the console is to use one of the family of *printf* commands. These functions are special, since they take a variable number of arguments, and the number is decided by the first - the format string,

· printf

Listing 6.33: to do

```
"printf" formatString {ident}
```

where a formatString is a string (simple or verbatim) with placeholders. The function printf prints formatString to the console, where all placeholder has been replaced by the value of the corresponding argument formatted as specified, e.g., in **printfn "1 2 %d"3** the formatString is "1 2 %d", and the placeholder is %d, and the printf replaced the placeholder with the value of the corresponding argument, and the result is printed to the console, in this case 1 2 3. Possible formats for the placeholder are,

Listing 6.34: to do

```
placeholder = "%%" | ("% " [flags] [width] [ "." precision] specifier) (* No
spaces between rules *)
flags = ["0"] ["+"] [SP] (* No spaces between rules *)
width = ["-"] ("*" | [dInt]) (* No spaces between rules *)
specifier = "b" | "d" | "i" | "u" | "x" | "X" | "o" | "e" | "E" | "f" | "F" |
" g" | "G" | "M" | "O" | "A" | "a" | "t"
```

There are specifiers for all the basic types and more as elaborated in Table 6.1. The placeholder can

be given a specified with, either by setting a specific integer, or using the * character, indicating that the width is given as an argument prior to the replacement value. Default is for the value to be right justified in the field, but left justification can be specified by the - character. For number types, you can specify their format by: "0" for padding the number with zeros to the left, when right justifying the number; "+" to explicitly show a plus sign for positive numbers; SP to enforce a space, where there otherwise would be a plus sign for positive numbers. For floating point numbers, the precision integer specifies the number of digits displayed of the fractional part. Examples of some of these combinations are,

Listing 6.35, printfExample.fsx:
Examples of printf and some of its formatting options.

```
let pi = 3.1415192
let hello = "hello"
printf "An integer: %d\n" (int pi)
printf "A float %f on decimal form and on %e scientific form, and a char
      '%c'\n" pi pi
printf "A char '%c' and a string \"%s\"\n" hello.[0] hello
printf "Float using width 8 and 1 number after the decimal:\n"
printf "  \"%8.1f\" \"%-8.1f\"\n" pi -pi
printf "  \"%08.1f\" \"%08.1f\"\n" pi -pi
printf "  \"% 8.1f\" \"% 8.1f\"\n" pi -pi
printf "  \"%-8.1f\" \"%-8.1f\"\n" pi -pi
printf "  \"%+8.1f\" \"%+8.1f\"\n" pi -pi
printf "  \"%8s\" \"%-8s\"\n" "hello" "hello"
```

```
An integer: 3
A char 'h' and a string "hello"
Float using width 8 and 1 number after the decimal:
  "    3.1" "   -3.1"
  "000003.1" "-00003.1"
  "    3.1" "   -3.1"
  "3.1" "   -3.1"
  "    +3.1" "   -3.1"
  "    hello"
"hello"
```

Not all combinations of flags and identifier types are supported, e.g., strings cannot have number of integers after the decimal specified. The placeholder types "%A", "%a", and "%t" are special for F#, examples of their use are,

Listing 6.36, printfExampleAdvance.fsx:

```
let noArgument writer = printf "I will not print anything"
let customFormatter writer arg = printf "Custom formatter got: \"%A\"" arg
printf "Print examples: %A, %A, %A\n" 3.0m 3uy "a string"
printf "Print function with no arguments: %t\n" noArgument
printf "Print function with 1 argument: %a\n" customFormatter 3.0
```

```
Print examples: 3.0M, 3uy, "a string"
Print function with no arguments: I will not print anything
Print function with 1 argument: Custom formatter got: "3.0"
```

The %A is special in that all built-in types including tuples, lists, and arrays to be discussed in Chapter 9 can be printed using this formatting string, but notice that the formatting performed includes the

Function	Example	Description
<code>printf</code> <code>printfn</code>	<code>printf "%d apples" 3</code>	Prints to the console, i.e., <code>stdout</code> as <code>printf</code> and adds a newline.
<code>fprintf</code> <code>fprintfn</code>	<code>fprintf stream "%d apples" 3</code>	Prints to a stream, e.g., <code>stderr</code> and <code>stdout</code> , which would be the same as <code>printf</code> and <code>eprintf</code> . as <code>fprintf</code> but with added newline.
<code>eprintf</code> <code>eprintfn</code>	<code>eprintf "%d apples" 3</code>	Print to <code>stderr</code> as <code>eprintf</code> but with added newline.
<code>sprintf</code>	<code>printf "%d apples" 3</code>	Return printed string
<code>failwithf</code>	<code>failwithf "%d failed apples" 3</code>	prints to a string and used for raising an exception.

Table 6.2: The family of printf functions.

named literal string. The two formatting strings `%t` and `%a` are options for user-customizing the formatting, and will not be discussed further.

Beware, `formatString` is not a `string` but a `Printf.TextWriterFormat`, so to predefine a `formatString` as, e.g., `let str = "hello %s" in printf str "world"` will be a type error.

The family of `printf` is shown in Table 6.2. The function `fprintf` prints to a stream, e.g., `stderr` and `stdout`, of type `System.IO.TextWriter`. Streams will be discussed in further detail in Chapter 12. The function `failwithf` is used with exceptions, see Chapter 11 for more details. The function has a number of possible return value types, and for testing the `ignore` function ignores it all, e.g., `ignore (failwithf "%d failed apples" 3)`

6.5 Variables

The `mutable` in `let` bindings means that the identifier may be rebound to a new value using the `<-` lexeme, e.g.,³

Listing 6.37: to do

```
expr = ...
| expr "<-" expr (*assignment*)
```

Mutable data is synonymous with the term *variable*. A variable is an area in the computers working memory associated with an identifier and a type, and this area may be read from and written to during program execution. For example,

³Todo: Discussion on heap and stack should be added here.

Listing 6.38, mutableAssignReassingShort.fsx:
A variable is defined and later reassigned a new value.

```
let mutable x = 5
printfn "%d" x
x <- -3
printfn "%d" x
```

```
5
-3
```

Here, an area in memory was denoted `x`, initially assigned the integer value 5, hence the type was inferred to be `int`. Later, this value of `x` was replaced with another integer using the `<-` lexeme. The `<-` lexeme is used to distinguish the assignment from the comparison operator, i.e., if we by mistake had written,

Listing 6.39, mutableEqual.fsx:
Common error - mistaking `=` and `<-` lexemes for mutable variables. The former is the test operator, while the latter is the assignment expression.

```
> let mutable a = 0
- a = 3;;

val mutable a : int = 0
val it : bool = false
```

then we instead would have obtained the default assignment of the result of the comparison of the content of `a` with the integer 3, which is false. However, it's important to note, that when the variable is initially defined, then the `'='` operator must be used, while later reassignments must use the `<-` expression.

Assignment type mismatches will result in an error,

Listing 6.40, mutableAssignReassingTypeError.fsx:
Assignment type mismatching causes a compile time error.

```
let mutable x = 5
printfn "%d" x
x <- -3.0
printfn "%d" x
```

```
/Users/sporring/repositories/fsharpNotes/src/
  mutableAssignReassingTypeError.fsx(3,6): error FS0001: This expression
    was expected to have type
      int
but here has type
      float
```

I.e., once the type of an identifier has been declared or inferred, then it cannot be changed.

A typical variable is a counter of type integer, and a typical use of counters is to increment them, i.e., erasing a new value to be one more than its previous value. For example,

Listing 6.41, mutableAssignIncrement.fsx:
Variable increment is a common use of variables.

```
let mutable x = 5 // Declare a variable x and assign the value 5 to it
printfn "%d" x
x <- x + 1 // Assign a new value -3 to x
printfn "%d" x
```

5
6

A function that elegantly implements the incrementation operation may be constructed as,

Listing 6.42, mutableAssignIncrementEncapsulation.fsx:

```
let incr =
    let mutable counter = 0
    fun () ->
        counter <- counter + 1
        counter
printfn "%d" (incr ())
printfn "%d" (incr ())
printfn "%d" (incr ())
```

1
2
3

⁴ Here, the output of `incr` is an anonymous function, that takes no argument, increments the variable of `incr` and returns the new value of the counter. This construction is called *encapsulation*, since the variable `counter` is hidden by the function `incr` from the user, i.e., the user need not be concerned with how the increment operator is implemented and the variable name used by `incr` does not clutter the scope where it is used.

· encapsulation

Variables implement dynamic scope, e.g., in comparison with the lexical scope, where the value of an identifier depends on which line in the program, an identifier is defined, dynamic scope depends on, when it is used. E.g., the script in Listing 6.25 defines a function using lexical scope and returns the number 6.0, however, if `a` is made `mutable`, then the behaviour is different:

Listing 6.43, dynamicScopeNFunction.fsx:
Mutual variables implement dynamics scope rules. Compare with Listing 6.25.

```
let testScope x =
    let mutable a = 3.0
    let f z = a * x
    a <- 4.0
    f x
printfn "%A" (testScope 2.0)
```

8.0

Here, the response is 8.0, since the value of `a` changed before the function `f` was called.

⁴Todo: Explain why this works!

Variables cannot be returned from functions, that is,

Listing 6.44, mutableAssignReturnValue.fsx:

```
let g () =  
    let x = 0  
    x  
printfn "%d" (g ())  
  
0
```

declares a function that has no arguments and returns the value 0, while the same for a variable is invalid,

Listing 6.45, mutableAssignReturnVariable.fsx:

```
let g () =  
    let mutual x = 0  
    x  
printfn "%d" (g ())  
  
/Users/sporring/repositories/fsharpNotes/src/mutableAssignReturnVariable.  
fsx(3,3): error FS0039: The value or constructor 'x' is not defined
```

There is a workaround for this by using *reference cells* by the build-in function `ref` and operators `!` and `:=`, · reference cells

Listing 6.46, mutableAssignReturnRefCell.fsx:

```
let g () =  
    let x = ref 0  
    x  
let y = g ()  
printfn "%d" !y  
y := 3  
printfn "%d" !y  
  
0  
3
```

That is, the `ref` function creates a reference variable, the `!` and the `:=` operators reads and writes its value. Reference cells are in some language called pointers, and their use is strongly discouraged, since they may cause *side-effects*, which is the effect that one function changes the state of another, such as the following example demonstrates,⁵ · side-effects

⁵Todo: Discuss side-effects!

Listing 6.47, mutableAssignReturnSideEffect.fsx:

```
let updateFactor factor =  
    factor := 2  
  
let multiplyWithFactor x =  
    let a = ref 1  
    updateFactor a  
    !a * x  
  
printfn "%d" (multiplyWithFactor 3)  
  
6
```

In the example, the function `updateFactor` changes a variable in the scope of `multiplyWithFactor`, which is prone to errors, since the style of programming does not follow the usual assignment syntax. Better style of programming is,

Listing 6.48, mutableAssignReturnWithoutSideEffect.fsx:

```
let updateFactor () =  
    2  
  
let multiplyWithFactor x =  
    let a = ref 1  
    a := updateFactor ()  
    !a * x  
  
printfn "%d" (multiplyWithFactor 3)  
  
6
```

Here, there can be no doubt in `multiplyWithFactor` that the value of 'a' is changing. Side-effects do have their use, but should in general be avoided at almost all costs, and in general it is advised to refrain from using ref cells.

6

⁶Todo: Add something about mutable functions

Chapter 7

In-code documentation

Documentation is a very important part of writing programs, since it is most unlikely, that you will be writing really obvious code. And what seems obvious at the point of writing may be mystifying months later to the author and to others. The documentation serves several purposes:

1. Communicate what the code should be doing
2. Highlight big insights essential for the code
3. Highlight possible conflicts and/or areas, where the code could be changed later

The essential point is that coding is a journey in problem solving, and proper documentation is an aid in understanding the solution and the journey leading to it. Documentation is most often a mixture between in-code documentation and accompanying documents. Here we will focus on in-code documentation, but arguably this does cause problems in multi-language environments, and run the risk of bloating code.

F# has the following simplified syntax for in-code documentation,

```
blockComment = "(*" {codePoint} "*)" ;  
lineComment = "//" {codePoint - newline} newline;
```

That is, text framed as a `blockComment` is still parsed by F# as keywords and basic types implying that `(* a comment (* in a comment *))` and `(* "*)"` are valid comments, while `(* " *)` is invalid.¹

The F# compiler has an option for generating *Extensible Markup Language (XML)* files from scripts using the C# documentation comments tags². The XML documentation starts with a triple-slash `///`, i.e., a `lineComment` and a slash, which serves as comments for the code construct, that follows immediately after. XML consists of tags which always appears in pairs, e.g., the tag “tag” would look like `<tag> ... </tag>`. The F# accept any tags, but recommends those listed in Table 7.1. If no tags are used, then it is automatically assumed to be a `<summary>`. An example of a documented script is,

· Extensible
Markup
Language
· XML

¹Todo: **lstlisting** colors is bad.

²For specification of C# documentations comments see ECMA-334 3rd Edition, Annex E, Section 2: <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-334.pdf>

Tag	Description
<c>	Set text in a code-font.
<code>	Set one or more lines in code-font.
<example>	Set as an example.
<exception>	Describe the exceptions a function can throw.
<list>	Create a list or table.
<para>	Set text as a paragraph.
<param>	Describe a parameter for a function or constructor.
<paramref>	Identify that a word is a parameter name.
<permission>	Document the accessibility of a member.
<remarks>	Further describe a function.
<returns>	Describe the return value of a function.
<see>	Set as link to other functions.
<seealso>	Generate a See Also entry.
<summary>	Main description of a function or value.
<typeparam>	Describe a type parameter for a generic type or method.
<typeparamref>	Identify that a word is a type parameter name.
<value>	Describe a value.

Table 7.1: Recommended XML tags for documentation comments, from ECMA-334 3rd Edition, Annex E, Section 2.

Listing 7.1, commentExample.fsx:
Code with XML comments.

```

/// Calculate the determinant of a quadratic equation with parameters a, b
/// , and c
let determinant a b c =
    b ** 2.0 - 2.0 * a * c

/// <summary>Find x when 0 = ax^2+bx+c.</summary>
/// <remarks>Negative determinants are not checked.</remarks>
/// <example>
///     The following code:
///     <code>
///         let a = 1.0
///         let b = 0.0
///         let c = -1.0
///         let xp = (solution a b c +1.0)
///         printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
///     </code>
///     results in <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> printed to
///     the console.
/// </example>
/// <param name="a">Quadratic coefficient.</param>
/// <param name="b">Linear coefficient.</param>
/// <param name="c">Constant coefficient.</param>
/// <param name="sgn">+1 or -1 indicating which solution is to be
///     calculated.</param>
/// <returns>The solution to x.</returns>
let solution a b c sgn =
    let d = determinant a b c
    (-b + sgn * sqrt d) / (2.0 * a)

let a = 1.0
let b = 0.0
let c = -1.0
let xp = (solution a b c +1.0)
printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp

```

```

0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7

```

Mono's `fsharpc` command may be used to extract the comments into an XML file,

Listing 7.2, Converting in-code comments to XML.

```
$ fsharpc --doc:commentExample.xml commentExample.fsx
F# Compiler for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License
```

This results in an XML file with the following content,

Listing 7.3, An XML file generated by `fsharpc`.

```
<?xml version="1.0" encoding="utf-8"?>
<doc>
<assembly><name>commentExample</name></assembly>
<members>
<member name="M:CommentExample.solution(System.Double,System.Double,System
    Double,System.Double)">
    <summary>Find x when 0 = ax^2+bx+c.</summary>
    <remarks>Negative determinants are not checked.</remarks>
    <example>
        The following code:
        <code>
            let a = 1.0
            let b = 0.0
            let c = -1.0
            let xp = (solution a b c +1.0)
            printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
        </code>
        results in <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> printed to the
            console.
    </example>
    <param name="a">Quadratic coefficient.</param>
    <param name="b">Linear coefficient.</param>
    <param name="c">Constant coefficient.</param>
    <param name="sgn">+1 or -1 indicating which solution is to be calculated.
        </param>
    <returns>The solution to x.</returns>
</member>
<member name="M:CommentExample.determinant(System.Double,System.Double,
    System.Double)">
    <summary>
        Calculate the determinant of a quadratic equation with parameters a, b,
            and c
    </summary>
</member>
</members>
</doc>
```

The extracted XML is written in C# type by convention, since F# is part of the Mono and .Net framework that may be used by any of the languages using Assemblies. Besides the XML inserted in the script, the XML has added `<?xml ...>` header, `<doc>`, `<assembly>`, `<members>`, and `<member>` tags. The header and the `<doc>` tag are standards for XML. The extracted XML is geared towards documenting big libraries of codes and thus highlights the structured programming organization, see Part IV, and `<assembly>`, `<members>`, and `<member>` are indications for where the functions belong in the hierarchy. As an example, the prefix `M:CommentExample.` means that it is a

method in the namespace `CommentExample`, which in this case is the name of the file. Further, the function type `val solution : a:float -> b:float -> c:float -> sgn:float -> float` is in the XML documentation `M:CommentExample.solution(System.Double,System.Double,System.Double,System.Double)`, which is the C# equivalent.

An accompanying program in the Mono suite is `mdoc`, which primary use is to perform a syntax analysis of an assembly and generate a scaffold XML structure for an accompanying document. With the `-i` flag, it is further possible to include the in-code comments as initial descriptions in the XML. The XML may be updated gracefully by `mdoc` as the code develops, without destroying manually entered documentation in the accompanying documentation. Finally, the XML may be exported to HTML

The primary use of the `mdoc` command is to analyze compiled code and generate an empty XML structure with placeholders to describe functions, values, and variables. This structure can then be updated and edited as the program develops. The edited XML files can then be exported to *Hyper Text Markup Language* (HTML) files, which can be viewed in any browser. Using the console, all of this is accomplished by,

· Hyper Text
Markup
Language
· HTML

Listing 7.4, Converting an XML file to HTML.

```
$ mdoc update -o commentExample -i commentExample.xml commentExample.exe
New Type: CommentExample
Member Added: public static double determinant (double a, double b, double
c);
Member Added: public static double solution (double a, double b, double c,
double sgn);
Member Added: public static double a { get; }
Member Added: public static double b { get; }
Member Added: public static double c { get; }
Member Added: public static double xp { get; }
Namespace Directory Created:
New Namespace File:
Members Added: 6, Members Deleted: 0
$ mdoc export-html -out commentExampleHTML commentExample
.CommentExample
```

The primary use of the `mdoc` command is to analyze compiled code and generate an empty XML structure with placeholders to describe functions, values, and variables. This structure can then be updated and edited as the program develops. The edited XML files can then be exported to HTML files, which can be viewed in any browser, an example of which is shown in Figure 7.1. A full description of how to use `mdoc` is found here³.

³<http://www.mono-project.com/docs/tools+libraries/tools/monodoc/generating-documentation/>

solution Method

Find x when $0 = ax^2 + bx + c$.

Syntax

```
[Microsoft.FSharp.Core.CompilationArgumentCounts(Mono.Cecil.CustomAttributeArgument[])]  
public static double solution (double a, double b, double c, double sgn)
```

Parameters

a
Quadratic coefficient.

b
Linear coefficient.

c
Constant coefficient.

sgn
+1 or -1 indicating which solution is to be calculated.

Returns

The solution to x .

Remarks

Negative determinants are not checked.

Example

The following code:

```
Example  
let a = 1.0  
let b = 0.0  
let c = -1.0  
let xp = (solution a b c +1.0)  
printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
```

results in $0 = 1.0x^2 + 0.0x + -1.0 \Rightarrow x_+ = 0.7$ printed to the console.

Requirements

Namespace:
Assembly: commentExample (in commentExample.dll)
Assembly Versions: 0.0.0.0

Figure 7.1: Part of the HTML documentation as produce by `mdoc` and viewed in a browser.

Chapter 8

Controlling program flow

Non-recursive functions encapsulates code and allows for some control of flow, that is, if there is a piece of code, which we need to to have executed many times, then we can encapsulate it in the body of a function, and then call the function several times. In this chapter, we will look at more general control of flow via loops, conditional execution, and recursion, and therefore we look at further extension of the `expr` rule,

Listing 8.1: to do

```
expr = ...
| "if" expr "then" expr {"elif" expr "then" expr} ["else" expr] (*
  conditional*)
| "while" expr "do" expr ["done"] (*while*)
| "for" ident "=" expr "to" expr "do" expr ["done"] (* simple for expression
  *)
| "let" ["rec"] functionDefn "in" expr (*binding a function or operator*)
```

8.1 For and while loops

Many programming constructs need to be repeated. The most basic example is counting, e.g., from 1 to 10 with a `for`-loop,¹

· `for`

Listing 8.2, count.fsx:
Counting from 1 to 10 using a `for`-loop.

```
> for i = 1 to 10 do
-   printf "%d " i
-   printfn " ";
1 2 3 4 5 6 7 8 9 10

val it : unit = ()
```

As this interactive script demonstrates, the identifier `i` takes all the values between 1 and 10, but in spite of its changing state, it is not mutable. Note also that the return value of the `for` expression is

¹Todo: Is it clear enough that the body of the loop is repeated?




Figure 8.1: The Fibonacci spiral is an approximation of the golden spiral. Each square has side lengths of successive Fibonacci numbers, and the curve in each square is the circular arc with radius of the square it is drawn in. Figure by Dicklyon <https://commons.wikimedia.org/w/index.php?curid=3730979>

() like the `printf` functions. The `for` and `while` loops follow the syntax,

Listing 8.3: U

```
expr = ...
| "while" expr "do" expr ["done"] (*while*)
| "for" ident "=" expr "to" expr "do" expr ["done"] (* simple for expression
*)
```

using lightweight syntax the script block between the `do` and `done` keywords may be replaced by a newline and indentation, e.g.,

· `do`
· `done`

Listing 8.4, `countLightweight.fsx`: Counting from 1 to 10 using a `for`-loop.

```
for i = 1 to 10 do
  printf "%d " i
  printfn ""
1 2 3 4 5 6 7 8 9 10
```

A more complicated example is,

Problem 8.1:

Write a program that prints the n 'th Fibonacci number.

The Fibonacci numbers is the series of numbers 1, 1, 2, 3, 5, 8, 13..., where the $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$, and they are related to Golden spirals shown in Figure 8.1.² We could solve this problem with a `for`-loop as follows,

²Todo: Should add to the figure, quadratic paper squares and area annotations to strengthen the relation to Fibonacci series.

Listing 8.5, fibFor.fsx:

The n 'th Fibonacci number as the sum of the previous 2 numbers, which are sequentially updated from 3 to n .

```
let fib n =
    let mutable a = 1
    let mutable b = 1
    let mutable f = 0
    for i = 3 to n do
        f <- a + b
        a <- b
        b <- f
    f

printfn "fib(1) = 1"
printfn "fib(2) = 1"
for i = 3 to 10 do
    printfn "fib(%d) = %d" i (fib i)
```

```
fib(1) = 1
fib(2) = 1
fib(3) = 2
fib(4) = 3
fib(5) = 5
fib(6) = 8
fib(7) = 13
fib(8) = 21
fib(9) = 34
fib(10) = 55
```

The basic idea of the solution is that if we are given the $(n - 1)$ 'th and $(n - 2)$ 'th numbers, then the n 'th number is trivial to compute. And assume that $\text{fib}(1)$ and $\text{fib}(2)$ are given, then it is trivial to calculate the $\text{fib}(3)$. Now we have the first 3 numbers, so we disregard $\text{fib}(1)$ and calculate $\text{fib}(4)$ from $\text{fib}(2)$ and $\text{fib}(3)$, and this process continues until we have reached the desired $\text{fib}(n)$.

For the alternative `for`-loop, consider the problem,

Problem 8.2:

Write a program that identifies prime factors of a given integer n .

Prime numbers are integers divisible only by 1 and themselves with zero remainder. Let's assume that we already have identified a list of primes from 2 to n , then we could write a program that checks the remainder as follows,

Listing 8.6, primeCheck.fsx:

Checking whether a given number has remainder zero after division by some low prime numbers.

```
let primeFactorCheck n =
    printfn "%d %% i = 0?" n
    for i in [2; 3; 5; 7; 11; 13; 17] do
        printfn "i = %d? %b" i (n%i = 0)
    ()

primeFactorCheck 10

10 % i = 0?
i = 2? true
i = 3? false
i = 5? true
i = 7? false
i = 11? false
i = 13? false
i = 17? false
```

In this example, the variable `i` runs through the elements of a list, which will be discussed in further detail in Chapter 9.

The `while`-loop is simpler than the `for`-loop and does not contain a builtin counter structure. Hence, if we are to repeat the count-to-10 program from Listing 8.2 example, it would look somewhat like, `while`

Listing 8.7, countWhile.fsx:

Count to 10 with a counter variable.

```
let mutable i = 1
while i <= 10 do
    printf "%d " i
    i <- i + 1
printf "\n"

1 2 3 4 5 6 7 8 9 10
```

In this case, the `for`-loop is to be preferred, since more lines of code typically means more chances of making a mistake. But the `while`-loop allows for other logical structures. E.g., let's find the biggest Fibonacci number less than 100,

Listing 8.8, fibWhile.fsx:

Search for the largest Fibonacci number less than a specified number.

```

let largestFibLeq n =
    let mutable a = 1
    let mutable b = 1
    let mutable f = 0
    while f <= n do
        f <- a + b
        a <- b
        b <- f
    a

printfn "largestFibLeq(1) = 1"
printfn "largestFibLeq(2) = 1"
for i = 3 to 10 do
    printfn "largestFibLeq(%d) = %d" i (largestFibLeq i)

```

```

largestFibLeq(1) = 1
largestFibLeq(2) = 1
largestFibLeq(3) = 3
largestFibLeq(4) = 3
largestFibLeq(5) = 5
largestFibLeq(6) = 5
largestFibLeq(7) = 5
largestFibLeq(8) = 8
largestFibLeq(9) = 8
largestFibLeq(10) = 8

```

Thus, **while**-loops are most often used, when the number of iteration cannot easily be decided, when entering the loop.

Both **for**- and **while**-loops are often associated with variables, i.e., values that change while looping. If one mistakenly used values and rebinding, then the result would in most cases be of little use, e.g.,

Listing 8.9, forScopeError.fsx:

Lexical scope error. While rebinding is valid F# syntax, has little effect due to lexical scope.

```

let a = 1
for i = 1 to 10 do
    let a = a + 1
    printf "(%d, %d) " a i
printf "\n"

```

```

(2, 1) (2, 2) (2, 3) (2, 4) (2, 5) (2, 6) (2, 7) (2, 8) (2, 9) (2, 10)

```

I.e., the **let** expression rebinds **a** every iteration of the loop, but the value on the right-hand-side is taken lexically from above, where **a** has the value 1, so every time the result is the value 2.

8.2 Conditional expressions

Consider the task,

Problem 8.3:

Write a function that given n writes the sentence, “I have n apple(s)”, where the plural ‘s’ is added appropriately.

For this we need to test on n ’s size, and one option is to use conditional expressions like,

Listing 8.10: Using conditional expression to generate different strings.

```
let applesIHave n =
  if n < 0 then "I owe " + (string -n) + " apples"
  elif n < 1 then "I have no apples"
  elif n < 2 then "I have 1 apple"
  else "I have " + (string n) + " apples"

printfn "%A" (applesIHave -3)
printfn "%A" (applesIHave -1)
printfn "%A" (applesIHave 0)
printfn "%A" (applesIHave 1)
printfn "%A" (applesIHave 2)
printfn "%A" (applesIHave 10)
```

The grammar for conditional expressions is,

Listing 8.11: w

```
expr = ...
| "if" expr "then" expr {"elif" expr "then" expr} ["else" expr] (*
  conditional*)
```

here the `expr` following `if` and `elif` are *conditions*, i.e., expressions that evaluate to a boolean value. The `expr` following `then` and `else` are called *branches*, and all branches must have same type. The result of the conditional expression is the first branch, for which its condition was true. The lightweight syntax allows for the visually more simple expression of scope by use of indentation

- `if`
- `elif`
- conditions
- `then`
- `else`
- branches

Listing 8.12: Lightweight syntax allows for making blocks of code by indentation in order to make code more for easy to read.

```
let applesIHave n =
  if n < 0 then
    "I owe " + (string -n) + " apples"
  elif n < 1 then
    "I have no apples"
  elif n < 2 then
    "I have 1 apple"
  else
    "I have " + (string n) + " apples"
```

Note that both `elif` and `else` branches are optional, which may cause problems, e.g., both `let a = if true then 3` and `let a = if true then 3 elif false then 4` will be invalid, since F# is not smart enough to realize that the type of the expression is uniquely determined. Instead F# looks for the `else` to ensure all cases have been covered, and that `a` always will be given a unique value of the same type regardless of the branches taken in the conditional statement, hence, `let a = if true then 3 else 4` is the only valid expression of the 3. However, the omitted branches are assumed to return `()`, and thus it is fine to say `let a = if true then ()` and `if true then printfn "hej"`

8.3 Programming intermezzo

Using loops and conditional expressions we are now able to solve the following problem

Problem 8.4:

Given an integer on decimal form, write its equivalent value on binary form

To solve this problem, consider odd numbers: They all have the property, that the least significant bit is 1, e.g., $1_2 = 1$, $101_2 = 5$, $110_2 = 6$, and that division by 2 is equal to right-shifting by 1, e.g., $1_2/2 = 0.1_2 = 0.5$, $101_2/2 = 10.1_2 = 2.5$, $110_2/2 = 11_2 = 3$. Thus by integer division by 2 and checking the remainder, we may sequentially read off the least significant bit. This leads to the following algorithm,

Listing 8.13, dec2bin.fsx:

Using integer division and remainder to convert any positive integer to binary form.

```
let dec2bin n =
    if n < 0 then
        "Illegal value"
    elif n = 0 then
        "0b0"
    else
        let mutable v = n
        let mutable str = ""
        while v > 0 do
            str <- (string (v % 2)) + str
            v <- v / 2
        "0b"+str
printfn "%d -> %s" -1 (dec2bin -1)
printfn "%d -> %s" 0 (dec2bin 0)
printfn "%d -> %s" 1 (dec2bin 1)
printfn "%d -> %s" 2 (dec2bin 2)
printfn "%d -> %s" 3 (dec2bin 3)
printfn "%d -> %s" 10 (dec2bin 10)
printfn "%d -> %s" 1023 (dec2bin 1023)

-----
-1 -> Illegal value
0 -> 0b0
1 -> 0b1
2 -> 0b10
3 -> 0b11
10 -> 0b1010
1023 -> 0b1111111111
```

8.4 Recursive functions

Recursion is a central concept in F#. A *recursive function* is a function, which calls itself. From a compiler point of view, this is challenging, since the function is used before the compiler has completed its analysis. However, this there is a technical solution for, and we will just concern ourselves with the logics of using recursion for programming. An example of a recursive function that counts from 1 to

· recursive
function

10 similarly to Listing 8.2 is,³

Listing 8.14, countRecursive.fsx:
Counting to 10 using recursion.

```
let rec prt a b =  
  if a > b then  
    printf "\n"  
  else  
    printf "%d " a  
    prt (a + 1) b  
  
prt 1 10  
  
1 2 3 4 5 6 7 8 9 10
```

Here the `prt` calls itself repeatedly, such that the first call is `prt 1 10`, which calls `prt 2 10`, and so on until the last call `prt 10 10`. Calling `prt 11 10` would not result in recursive calls, since when `a` is higher than 10 then the *stopping criterium* is met and a newline is printed. For values of `a` smaller than or equal `b` then the recursive branch is executed. Since `prt` calls itself as the last all but the stopping condition, then this is a *tail-recursive* function. Most compilers achieve high efficiency in terms of speed and memory, so **prefer tail-recursion whenever possible**.

· stopping
 criterium
· tail-recursive
Advice

Listing 8.15: U

```
expr = ...  
| "let" "rec" functionDefn "in" expr (*binding a function or operator*)
```

sing recursion to calculate the Fibonacci number as Listing 8.5.

³Todo: A drawing showing the stack for the example would be good.

Listing 8.16, fibRecursive.fsx:
The n 'th Fibonacci number using recursive.

```
let rec fib n =
  if n < 1 then
    0
  elif n = 1 then
    1
  else
    fib (n - 1) + fib (n - 2)

for i = 0 to 10 do
  printfn "fib(%d) = %d" i (fib i)
```

```
fib(0) = 0
fib(1) = 1
fib(2) = 1
fib(3) = 2
fib(4) = 3
fib(5) = 5
fib(6) = 8
fib(7) = 13
fib(8) = 21
fib(9) = 34
fib(10) = 55
```

Here we used the fact that including $\text{fib}(0) = 0$ in the Fibonacci series also produces it using the rule $\text{fib}(n) = \text{fib}(n-2) + \text{fib}(n-1)$, $n \geq 0$, which allowed us to define a function that is well defined for the complete set of integers. I.e., a negative argument returns 0. This is a general advice: **make functions that fails gracefully**. The recursive definition allows for recursive value definitions and defining several values and functions in one expression. Recursive values is particularly useful for defining infinite sequences, see Section 15.1.

Advice

4 5

⁴Todo: Add short-cut if-then-else with and and or logical operators.

⁵Todo: We need the `and` notation to define mutually recursive functions.

Bibliography

- [1] Alonzo Church. A set of postulates for the foundation of logic. *Annals of Mathematics*, 33(2):346–366, 1932.
- [2] Programming Research Group. Specifications for the ibm mathematical formula translating system, fortran. Technical report, Applied Science Division, International Business Machines Corporation, 1954.
- [3] John McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, 1960.
- [4] X3: ASA Sectional Committee on Computers and Information Processing. American standard code for information interchange. Technical Report ASA X3.4-1963, American Standards Association (ASA), 1963. <http://worldpowersystems.com/projects/codes/X3.4-1963/>.
- [5] George Pólya. *How to solve it*. Princeton University Press, 1945.

Index

. [], 31
ReadKey, 103
ReadLine, 103
Read, 103
System.Console.ReadKey, 103
System.Console.ReadLine, 103
System.Console.Read, 103
System.Console.WriteLine, 103
System.Console.Write, 103
WriteLine, 103
Write, 103
abs, 159
acos, 159
asin, 159
atan2, 159
atan, 159
bignum, 22
bool, 19
byte[], 22
byte, 22
ceil, 159
char, 19
cosh, 159
cos, 159
decimal, 22
double, 22
eprintfn, 48
eprintf, 48
exn, 19
exp, 159
failwithf, 48
float32, 22
float, 19
floor, 159
fprintfn, 48
fprintf, 48
ignore, 48
int16, 22
int32, 22
int64, 22
int8, 22
int, 19
it, 19
log10, 159
log, 159

max, 159
min, 159
nativeint, 22
obj, 19
pown, 159
printfn, 48
printf, 46, 48
round, 159
sbyte, 22
sign, 159
single, 22
sinh, 159
sin, 159
sprintf, 48
sqrt, 159
stderr, 48, 103
stdin, 103
stdout, 48, 103
string, 19
tanh, 159
tan, 159
uint16, 22
uint32, 22
uint64, 22
uint8, 22
unativeint, 22
unit, 19

American Standard Code for Information Inter-
change, 143

and, 26
anonymous function, 42
array sequence expressions, 123
Array.toArray, 76
Array.toList, 76
ASCII, 143
ASCIIbetical order, 31, 143

base, 19, 139
Basic Latin block, 144
Basic Multilingual plane, 144
basic types, 19
binary, 139
binary number, 21
binary operator, 25
binary64, 141

- binding, 14
- bit, 21, 139
- black-box testing, 80
- block, 38
- blocks, 144
- boolean and, 160
- boolean or, 160
- branches, 64
- branching coverage, 81
- bug, 79
- byte, 139

- character, 21
- class, 24, 32
- code point, 21, 144
- compiled, 11
- computation expressions, 71, 73
- conditions, 64
- Cons, 73
- console, 11
- coverage, 81
- currying, 43

- debugging, 13, 80, 87
- decimal number, 19, 139
- decimal point, 20, 139
- Declarative programming, 8
- digit, 20, 139
- dot notation, 32
- double, 141
- downcasting, 24

- EBNF, 20, 148
- efficiency, 80
- encapsulate code, 40
- encapsulation, 44, 50
- environment, 87
- exception, 29
- exclusive or, 29
- executable file, 11
- expression, 14, 24
- expressions, 8
- Extended Backus-Naur Form, 20, 148
- Extensible Markup Language, 54

- file, 102
- floating point number, 20
- format string, 14
- fractional part, 20, 24
- function, 17
- function coverage, 81
- Functional programming, 8, 113
- functional programming, 8
- functionality, 79
- functions, 8

- generic function, 41

- hand tracing, 87
- Head, 73
- hexadecimal, 139
- hexadecimal number, 21
- HTML, 57
- Hyper Text Markup Language, 57

- IEEE 754 double precision floating-point format, 141
- Imperativ programming, 113
- Imperative programming, 8
- implementation file, 11
- infix notation, 25
- infix operator, 24
- integer, 20
- integer division, 28
- interactive, 11
- IsEmpty, 73
- Item, 73

- jagged arrays, 76

- keyword, 14

- Latin-1 Supplement block, 144
- Latin1, 144
- least significant bit, 139
- Length, 73
- length, 69
- lexeme, 17
- lexical scope, 16, 42
- lexically, 36
- lightweight syntax, 33, 37
- list, 71
- list sequence expression, 123
- List.Empty, 73
- List.toArray, 73
- List.toList, 73
- literal, 19
- literal type, 22

- machine code, 113
- maintainability, 80
- member, 24, 69
- method, 32
- mockup code, 87
- module elements, 135
- modules, 11
- most significant bit, 139
- Mutable data, 49

- namespace, 24
- namespace pollution, 130
- NaN, 141

- nested scope, 38
- newline, 22
- not, 26
- not a number, 141

- obfuscation, 71
- object, 32
- Object oriented programming, 113
- Object-oriented programming, 8
- objects, 8
- octal, 139
- octal number, 21
- operand, 41
- operands, 25
- operator, 25, 41
- or, 26
- overflow, 27
- overshadows, 39

- pattern matching, 124, 131
- portability, 80
- precedence, 25
- prefix operator, 25
- Procedural programming, 113
- procedure, 44
- production rules, 148

- ragged multidimensional list, 73
- raise an exception, 96
- range expression, 72
- reals, 139
- recursive function, 65
- reference cells, 51
- reliability, 79
- remainder, 28
- rounding, 24
- run-time error, 29

- scientific notation, 20
- scope, 38
- script file, 11
- script-fragment, 17
- script-fragments, 11
- Seq.initInfinite, 122
- Seq.item, 120
- Seq.take, 120
- Seq.toArray, 122
- Seq.toList, 122
- side-effect, 75
- side-effects, 44, 52
- signature file, 11
- slicing, 75
- software testing, 80
- state, 8
- statement, 14
- statement coverage, 81
- statements, 8, 113
- states, 113
- stopping criterium, 66
- stream, 102
- string, 14, 22
- Structured programming, 8
- subnormals, 141

- Tail, 73
- tail-recursive, 66
- terminal symbols, 148
- tracing, 87
- truth table, 26
- tuple, 69
- type, 15, 19
- type declaration, 15
- type inference, 13, 15
- type safety, 41
- typecasting, 23

- unary operator, 25
- underflow, 27
- Unicode, 21
- unicode general category, 144
- Unicode Standard, 144
- unit of measure, 127
- unit testing, 80
- unit-less, 128
- unit-testing, 13
- upcasting, 24
- usability, 80
- UTF-16, 146
- UTF-8, 146

- variable, 49
- verbatim, 23

- white-box testing, 80, 81
- whitespace, 22
- whole part, 20, 24
- wild card, 36
- word, 139

- XML, 54
- xor, 29

- yield bang, 120