

Learning to program with F#

Jon Sparring

July 13, 2016

Contents

1	Executing F# code	3
1.1	Source code	3
1.2	Executing programs	3
2	Quick-start guide	6
3	Numbers, Characters, and Strings	13
3.1	Mutable Data	13
4	Functions and procedures	18
4.1	Procedures	20
5	Controlling program flow	21
5.0.1	Conditional expressions	21
5.0.2	For and while loops	22
6	Tuples, Lists, Arrays, and Sequences	26
6.1	Tuples	26
6.2	Lists	26
6.3	Arrays	26
6.3.1	1 dimensional arrays	26
6.3.2	Multidimensional Arrays	28
6.4	Sequences	31
7	Exceptions	32
7.1	Exception Handling	32
8	Testing programs	33
9	The Collection	34
9.1	Mutable Collections	34
9.1.1	Mutable lists	34
9.1.2	Stacks	34
9.1.3	Queues	34
9.1.4	Sets and dictionaries	34
10	Input/Output	35
10.1	Console I/O	35
10.2	File I/O	35
11	Graphical User Interfaces	38

12 Imperative programming	39
12.1 Introduction	39
12.2 Generating random texts	39
12.2.1 0'th order statistics	39
12.2.2 1'th order statistics	41
13 Functional programming	44
14 Object-oriented programming	45
A Number systems on the computer	46
A.1 Binary numbers	46
A.2 IEEE 754 floating point standard	46
B Commonly used character sets	50
B.1 ASCII	50
B.2 ISO/IEC 8859	50
B.3 Unicode	51
C A brief introduction to Extended Backus-Naur Form	54
Bibliography	57
Index	58

Chapter 1

Executing F# code

1.1 Source code

F# is a functional first programming language that also supports imperative and object oriented programming. It also has strong support for parallel programming and information rich programs. It was originally developed for Microsoft's .Net platform, but is available as open source for many operating systems through Mono. In this text we consider F# 4.0 and its Mono implementation, which is different from .Net mainly in terms of the number of libraries accessible. The complete language specification is described in <http://fsharp.org/specs/language-spec/4.0/FSharpSpec-4.0-latest.pdf>.

F# has 2 modes of execution, *interactive* and *compiled*. Interactive mode is well suited for small experiments or back-of-an-envelope calculations, but not for programming in general. In Mono, the interactive system is started by calling **fsharpi** from the *console*, while compilation is performed with **fsharpirc** and execution of the compiled code is performed using the **mono** command. The various forms of fsharp programs are identified by suffixes:

interactive
compiled
console

.fs An *implementation file*

implementation
file

.fsi A *signature file*

signature file

.fsx A *script file*

script file

.fscript Same as **.fsx**

.exe An *executable file*

executable file

The implementation, signature, and script files are all typically compiled to produce an executable file, but syntactically correct code can also be entered into the interactive system, in which case these are called *script-fragments*. The implementation and signature files are special kinds of script files used for building *modules*.

script-
fragments
modules

1.2 Executing programs

Programs may either be executed by the interpreter or by compiling and executing the compiled code.

In Mono the interpreter is called **fsharpi** and can be used in 2 ways: interactively, where a user enters 1 or more script-fragments separated by the `";;"` token, or to execute a script file treated as a single script-fragment. To illustrate the difference, consider the following program, which declares a value **a** to be the decimal value 3.0 and finally print it to the console:

```
gettingStartedStump.fsx
```

```
let a = 3.0
printfn "%g" a
```

An interactive session is obtained by starting the console, typing the `fsharpi` command, typing the lines of the program, and ending the script-fragment with the `;;` token:

```
$ fsharpi

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let a = 3.0
- printfn "%g" a;;
3

val a : float = 3.0
val it : unit = ()

> #quit;;
```

The interpreter is stopped by pressing `ctrl-d` or typing `"#quit;;"`. Conversely, executing the file with the interpreter as follows,

```
$ fsharpi gettingStartedStump.fsx
3
```

Finally, compiling and executing the code is performed as,

```
$ fsharpc gettingStartedStump.fsx
F# Compiler for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License
$ mono gettingStartedStump.exe
3
```

Both the interpreter and the compiler translates the source code into a format, which can be executed by the computer. While the compiler performs this translation once and stores the result in the executable file, the interpreter translates the code every time the code is executed. Thus, to run the program again with the interpreter, then it must be retranslated as `"$fsharpi gettingStartedStump.fsx"`, but since the program has been compiled, then the compile-execute only needs to be re-executed `"$ mono gettingStartedStump.exe"`. On a Macbook Pro, with a 2.9 Ghz Intel Core i5, the time the various stages takes for this script are.

Command	Time
<code>fsharpi gettingStartedStump.fsx</code>	1.88s
<code>fsharpc gettingStartedStump.fsx</code>	1.90s
<code>mono randomTextOrder0.exe</code>	0.05s

I.e., executing the script with `fsharpi` is slightly faster than by first compiling it with `fsharpc` and then executing the result with `mono`, $1.88s < 0.05s + 1.90s$, if the script were to be executed only once,

but every future execution of the script using the compiled version requires only the use of `mono`, which is much faster than `fsharp`, $1.88s \gg 0.05s$.

The interactive session results in extra output on the *type inference* performed, which is very useful for debugging and development of code-fragments, but both executing programs with the interpreted directly from a file and compiling and executing the program is much preferred for programming complete programs, since the starting state is well defined, and since this better supports *unit-testing* a tool for debugging programs consisting of many code-fragments.

Chapter 2

Quick-start guide

Programming is the art of solving problems by writing a program to be executed by a computer. For example, to solve the following problem,

What is the sum of 357 and 864?

we could write the following program in F#,

quickStartSum.fsx

```
let a = 357
let b = 864
let c = a + b
printfn "%A" c
```

```
> fsharpi quickStartSum.fsx

1221
```

Listing 2.1: A script to add 2 numbers and print the result to the console.

In box the above, we see our program was saved as a script in a file called `quickStartSum.fsx`, and in the console we executed the program by typing the command `fsharpi quickStartSum.fsx`. The result is then printed in the console to be 1221.

To solve the program we made program consisting of several lines, where each line was a *statement*. The first statement `let a = 357` used the `let` keyword to bind the value 357 to the name `a`. Likewise, we bound the value 864 to the name `b`, but to the name `c` we bound the result of evaluating the *expression* `a + b`. That is, first the value `a + b` was calculated by substituting the names of `a` and `b` with their values to give the expression, `357 + 864`, then this expression was evaluated by adding the values to give, 1221, and this value was finally bound to the name `c`. The last line printed the value of `c` to the console followed by a LF (line feed, see Appendix B.1) with the `printfn` function. Here `printfn` is a function of 2 arguments: `"%A"` and `c`. Notice, that in contrast to many other languages, F# does not use parentheses to frame the list of arguments, nor does it use commas to separate them. In general, the `printfn` function always has 1 or more arguments, and the first is a *format string*. A *string* is a sequence of characters starting and ending with double quotation marks. E.g., `let s = "this is a string of characters"` binds the string `"this is..."` to the name `s`. For the `printfn` function, the format string may be any string, but if it contains format character sequences, such as `%A`, then the values following the format string are substituted. The format string must match the value *type*, that is, here `c` is of type integer, whereas the format string `%A` matches any type.

Types are a central concept in F#. The language contains a number of built-in types, and it designed such that it is easy to define new types. The simplest types are called *primitive types*, and a table of some of the most commonly used primitive types are shown in Table 2.1. In the script 2.1 we

statement
keyword

expression

format string
string

type

primitive types

Type	Description
bool	Boolean values are true and false.
int	Integer values from -2,147,483,648 to 2,147,483,647.
float, double	64-bit IEEE 754 floating point value from $-\infty$ to ∞ , see Appendix A.2.
char	Unicode character values, see Appendix B.
string	Unicode sequence of characters.
unit	No value denoted ().

Table 2.1: Some primitive types.

bound values of types `int` and `string` to names. The values were not declared to have these types, instead the types were inferred by F#. Had we typed these statements line by line in an interactive session, then we would have seen the inferred types:

```
$ fsharp

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let a = 357;;

val a : int = 357

> let b = 864;;

val b : int = 864

> let c = a + b;;

val c : int = 1221

> printfn "%A" c;;
1221
val it : unit = ()

> #quit;;
```

The an interactive session displays the type using the `val` keyword. Using Extended Backus-Naur form (see Appendix C), the response from the the interpreter follows the syntax:

```
"val" name ":" type "=" value
```

Since the value is also responded, then the last `printfn` statement is superfluous. However, it is ill advised to design programs to be run in an interactive session, since the scripts needs to be manually copied every time it is to be run, and since the starting state may be unclear.

Advice!

Were we to solve a slightly different problem,

What is the sum of 357.6 and 863.4?

then we would have to use floating point arithmetic instead of integers. A subtle point here is that the above problem is stated in using base 10 numbers, while almost all computers perform their calculations

in base 2 numbers. E.g., the number 357.6 in base 10 can be considered the sum,

$$357.6 = 3 \cdot 10^2 + 5 \cdot 10^1 + 7 \cdot 10^0 + 6 \cdot 10^{-1}, \quad (2.1)$$

and in general any base 10 number, also known as *decimal numbers*, has the representation $a_n a_{n-1} a_{n-2} \dots a_0 . a_{-1} a_{-2} \dots a_{-m}$ for $n + 1$ digits to the left and m digits to the right of the period, and where $0 \leq a_i \leq 9$, and which translate to the equation

$$v = \sum_{i=-m}^n a_i 10^i. \quad (2.2)$$

Base 2 numbers, also known as *binary numbers*, has the general form

$$v = \sum_{i=-p}^q b_i 2^i, \quad (2.3)$$

where $0 \leq b_i \leq 1$ are binary digits, and where $357.6_{10} = 101100101.\overline{10011}_2$, using subscript to denote base. and where the bar notation means that the digits are repeated infinitely many times. I.e.,

$$357.6 = 1 \cdot 10^8 + 1 \cdot 10^6 + 1 \cdot 10^5 + 1 \cdot 10^2 + 1 \cdot 10^0 + 1 \cdot 10^{-1} + 1 \cdot 10^{-4} + 1 \cdot 10^{-5} \dots \quad (2.4)$$

where all terms with coefficient 0 have been removed for brevity. The example illustrates that while the number 357.6 is short to write in decimal, it has infinite number of digits in binary, and since any computer only has finite memory, then we must use an approximation. Floating point numbers is a popular approximation, whose type is `float` which is synonymous with `double` in F#, and which implements the IEEE 754 floating point standard for doubles. Thus the program looks like,

quickStartSumFloat.fsx

```
let a = 357.6
let b = 863.4
let c = a + b
printfn "%A" c

> fsharp quickStartSumFloat.fsx

1221.0
```

Listing 2.2: Floating point types and arithmetic.

Note that although the response is an integer, it has type `float` which is indicated by `1221.0` instead of `1221`.

F# is very picky about types, and generally does not allow types to be mixed. E.g., in an interactive session,

```
$ fsharp

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let a = 357;;

val a : int = 357

> let b = 863.4;;
```

```

val b : float = 863.4

> let c = a + b;;

    let c = a + b;;
    -----^

/Users/sporring/Desktop/fsharpNotes/stdin(3,13): error FS0001: The
    type 'float' does not match the type 'int'
> #quit;;

```

we see that binding a name to a number without a decimal point is inferred to be integer, while when binding to a number with a decimal point, then the type is inferred to be a float, and when trying to add values of integer and floating point, then we get an error. The result is floating point, so to perform the sum, we must convert the integer to floating point, and we could either write the floating point equivalent to the integer 357,

```
let a = 357.0
```

or *type cast* the integer to the float type as

```
let a = float 357
```

type cast

Type casting is a potential destructive operation, e.g., type casting a float to int removes the part after the decimal point without rounding,

quickStartDownCast.fsx

```

let a = 357.6
let b = int a
printfn "%A -> %A" a b

```

```
> fsharpi quickStartDownCast.fsx
```

```
357.6 -> 357
```

Listing 2.3: Type casting is potentially a destructive action, here the fractional part is removed.

The type casting in the example is sometimes called *down casting*, since the floating time is casted to a lesser type, in the sense that floating point numbers can represents more numbers than integers.

down casting

F# is a functional first programming language, and one implication is that names are constant and cannot be changed. If attempted, then F# will return an error as, e.g.,

quickStartRebindError.fsx

```

let a = 357
let a = 864

```

```
> fsharpi quickStartRebindError.fsx
```

```
/Users/sporring/repositories/fsharpNotes/quickStartRebindError.fsx
```

```
(2,5): error FS0037: Duplicate definition of value 'a'
```

Listing 2.4: A name cannot be rebound.

However, if the same was performed in an interactive session,

```
$ fsharp

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let a = 357;;

val a : int = 357

> let a = 864;;

val a : int = 864

> #quit;;
```

then apparently rebinding is legal. The difference is that the `;;` token defines a new nested *scope*. A scope is an area in a program, where a binding is valid. Scopes can be nested, and in F# a binding may reuse names in a nested scope, in which case the previous value is overshadowed. I.e., attempting the same without `;;` between the two `let` statements results in an error, e.g.,

```
$ fsharp

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let a = 357
- let a = 864;;

    let a = 864;;
    ----^

/Users/sporring/Desktop/fsharpNotes/stdin(2,5): error FS0037:
    Duplicate definition of value 'a'
> #quit;;
```

This is yet another reason why the interactive session should be avoided except for back-of-an-envelope experiment. Scopes can be nested squares as shown in Figure 2.1.

In F# functions are also values, and to define a function, we use the syntax

```
"let" name arg { arg } "=" expr
```

where `arg` is a list of argument names, and `expr` is an expression. Defining a function `sum` as part of




Figure 2.1: Binding of the the same name in the same scope is illegal in F# 2, but legal in a different scopes. In (a) the two bindings are in the same scope, which is illegal, while in (b) the bindings are in separate scopes by the extra `;;` token, which is legal.

the solution to the above program gives,

quickStartSumFct.fsx

```
let sum x y = x + y
let c = sum 357 864
printfn "%A" c
```

```
> fsharp quickStartSumFct.fsx
1221
```

Listing 2.5: A script to add 2 numbers using a user defined function.

Entering the function into an interactive session will illustrate the inferred type, the function `sum` has: `val sum : x:int * y:int -> int`, by which is meant that `sum` is a mapping from the set product of the set of integers with integers into integers. Type inference in F# may cause problems, since, e.g., the type of a function is based on the context in which it is defined. E.g., in an interactive session, defining the `sum` in one scope on a single line will default the types to integers, F#'s favorite type, which will give an error, if it in a nested scope is to be used for floats. E.g.,

```
$ fsharp

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let sum x y = x + y;;

val sum : x:int -> y:int -> int

> let c = sum 357.6 863.4;;

    let c = sum 357.6 863.4;;
    ~~~~~~

/Users/sporring/Desktop/fsharpNotes/stdin(2,13): error FS0001: This
expression was expected to have type
    int
but here has type
    float
> #quit;;
```

A remedy is to either define the function in the same scope as its use,

```
$ fsharpi

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let sum x y = x + y
- let c = sum 357.6 863.4;;

val sum : x:float -> y:float -> float
val c : float = 1221.0
> #quit;;
```

or specify the type of the function at point of definition using the notation,

```
"let" name argWType { argWType } [ ":" type ] "=" expr
argWType = arg | "(" arg ":" type ")"
```

where not all types need to be declared, just sufficient for F# to be able to infer the types for the full statement. In the example, one sufficient specification is,

```
$ fsharpi

F# Interactive for F# 4.0 (Open Source Edition)
Freely distributed under the Apache 2.0 Open Source License

For help type #help;;

> let sum (x : float) (y : float) = x + y;;

val sum : x:float -> y:float -> float

> let c = sum 357.6 863.4;;

val c : float = 1221.0

> #quit;;
```

but alternatively we could have specified the type of the result,

```
let sum x y : float = x + y
```

or even just one of the arguments,

```
let sum (x : float) y = x + y
```

In both cases, since the *+* operator is only defined for *operands* of the same type, then when the type of either the result, any or both operands are declared, then the type of the remaining follows directly.

operator
operands

Chapter 3

Numbers, Characters, and Strings

...

3.1 Mutable Data

The most common syntax for a value definition is

```
"let" [ "mutable" ] ident [ ":" type ] "=" expr "in" expr
```

or alternatively

```
"let" ["mutable"] ident [ ":" type ] "=" expr ["in"]  
expr
```

In the above, `ident` may be replaced with a more complicated pattern, but this is outside the scope of this text. If a value has been defined as mutable, then its value may be changed using the following syntax,

```
expr "<-" expr
```

Mutable data is synonymous with the term *variable*. A variable is an area in the computer's working memory associated with a name and a type, and this area may be read from and written to during program execution. For example,

Mutable data
variable

mutableAssignReassign.fsx

```
let mutable x = Unchecked.defaultof<int> // Declare a variable x of  
    type int and assign the corresponding default value to it.  
printfn "%d" x  
x <- 5 // Assign a new value 5 to x  
printfn "%d" x  
x <- -3 // Assign a new value -3 to x  
printfn "%d" x
```

```
> fsharpi mutableAssignReassign.fsx
```

```
0
```

```
5
-3
```

Here a area in memory was denoted x, declared as type integer and assigned a default value. Later, this value of x was replaced with another integer and yet another integer. The operator '<-' is used to distinguish the statement from the mathematical concept of equality. A short-hand for the above is available as,

mutableAssignReassingShort.fsx

```
let mutable x = 5 // Declare a variable x and assign the value 5 to it
printfn "%d" x
x <- -3 // Assign a new value -3 to x
printfn "%d" x
```

```
> fsharp mutableAssignReassingShort.fsx
5
-3
```

where the assignment of the default value was skipped, and the type was inferred from the assignment operation. However, it's important to note, that when the variable is declared, then the '=' operator must be used, while later reassignments must use the '<-' operator. Type mismatches will result in an error,

mutableAssignReassingTypeError.fsx

```
let mutable x = 5 // Declare a variable x and assign the value 5 to it
printfn "%d" x
x <- -3.0 // This is illegal, since -3.0 is a float, while x is of type int
printfn "%d" x
```

```
> fsharp mutableAssignReassingTypeError.fsx

/Users/sporring/repositories/fsharpNotes/
  mutableAssignReassingTypeError.fsx(3,6): error FS0001: This
    expression was expected to have type
      int
but here has type
  float
```

A typical variable is a counter of type integer, and a typical use of counters is to increment them, i.e., erasing a new value to be one more than its previous value. For example,

mutableAssignIncrement.fsx

```
let mutable x = 5 // Declare a variable x and assign the value 5 to it
printfn "%d" x
```

```
x <- x + 1 // Assign a new value -3 to x
printfn "%d" x
```

```
> fsharp Interactive mutableAssignIncrement.fsx
```

```
5
6
```

An function that elegantly implements the incrementation operation may be constructed as,

mutableAssignIncrementEncapsulation.fsx

```
let incr =
    let mutable counter = 0
    fun () ->
        counter <- counter + 1
        counter
printfn "%d" (incr ())
printfn "%d" (incr ())
printfn "%d" (incr ())
```

```
> fsharp Interactive mutableAssignIncrementEncapsulation.fsx
```

```
1
2
3
```

Here the output of `incr` is an anonymous function, that takes no argument, increments the variable of `incr` and returns the new value of the counter. This construction is called *encapsulation*, since the variable `counter` is hidden by the function `incr` from the user, i.e., the user need not be concerned with how the increment operator is implemented and the variable name used by `incr` does not clutter the scope where it is used.

Variables cannot be returned from functions, that is,

mutableAssignReturnValue.fsx

```
let g () =
    let x = 0
    x
printfn "%d" (g ())
```

```
> fsharp Interactive mutableAssignReturnValue.fsx
```

```
0
```

declares a function that has no arguments and returns the value 0, while the same for a variable is illegal,

mutableAssignReturnVariable.fsx

```
let g () =
    let mutual x = 0
    x
printfn "%d" (g ())
```

Explain why
this works!
encapsulation


```
> fsharp mutableAssignReturnVariable.fsx

/Users/sporring/repositories/fsharpNotes/
  mutableAssignReturnVariable.fsx(3,3): error FS0039: The value
    or constructor 'x' is not defined
```

There is a workaround for this by using *reference cells* by the build-in function `ref` and operators `!` and `:=`, reference cells

mutableAssignReturnRefCell.fsx

```
let g () =
    let x = ref 0
    x
let y = g ()
printfn "%d" !y
y := 3
printfn "%d" !y

> fsharp mutableAssignReturnRefCell.fsx

0
3
```

That is, the `ref` function creates a reference variable, the `!` and the `:=` operators reads and writes its value. Reference cells are in some language called pointers, and their use is strongly discouraged, since they may cause *side-effects*, which is the effect that one function changes the state of another, side-effects such as the following example demonstrates,

mutableAssignReturnSideEffect.fsx

```
let updateFactor factor =
    factor := 2

let multiplyWithFactor x =
    let a = ref 1
    updateFactor a
    !a * x

printfn "%d" (multiplyWithFactor 3)

> fsharp mutableAssignReturnSideEffect.fsx

6
```

In the example, the function `updateFactor` changes a variable in the scope of `multiplyWithFactor`, which is prone to errors, since the style of programming does not follow the usual assignment syntax. Better style of programming is,

mutableAssignReturnWithoutSideEffect.fsx

```
let updateFactor () =  
    2  
  
let multiplyWithFactor x =  
    let a = ref 1  
    a := updateFactor ()  
    !a * x  
  
printfn "%d" (multiplyWithFactor 3)  
  
-----  
> fsharp -i mutableAssignReturnWithoutSideEffect.fsx  
6
```

Here there can be no doubt in `multiplyWithFactor` that the value of 'a' is changing. Side-effects do have their use, but should in general be avoided at almost all costs, and in general it is advised to refrain from using ref cells.

Chapter 4

Functions and procedures

A function is a mapping between an input and output domain. F# is a functional first programming language, and offers a number of alternative methods for specifying parameters. A simple example is,

functionDeclarationMul.fsx

```
let mul (x, y) = x*y
let z = mul (3, 5)
printfn "%d" z
```

```
> fsharp -i functionDeclarationMul.fsx
15
```

which declares a function of a tuple and returns their multiplication. The types are inferred from its first use in the second line, i.e., `mul` is `val mul : x:int * y:int -> int`. An argument may be of generic type for input, which need not be inferred without sacrificing type safety, e.g.,

functionDeclarationGeneric.fsx

```
let second (x, y) = y
let a = second (3, 5)
printfn "%A" a
let b = second ("horse", 5.0)
printfn "%A" b
```

```
> fsharp -i functionDeclarationGeneric.fsx
5
5.0
```

Here the function `second` does not use the first element in the tuple, `x`, and the type of the second element, `y`, can safely be anything.

Functions may be anonymously declared using the `fun` keyword,

functionDeclarationAnonymous.fsx

```
let first = fun (x, y) -> x
printfn "%d" (first (5, 3))
```

```
> fsharp -i functionDeclarationAnonymous.fsx
5
```

Anonymous functions are often used as arguments to other functions, e.g.,

```
functionDeclarationAnonymousAdvanced.fsx

let apply (f, x, y) = f (x, y)
let z = apply ((fun (a, b) -> a * b), 3, 6)
printfn "%d" z

> fsharp -i functionDeclarationAnonymousAdvanced.fsx
18
```

This is a powerful concept, but can make programs hard to read, and overuse is not recommended.

Functions may be declared using pattern matching, which is a flexible method for declaring output depending on conditions on the input value. The most common pattern matching method is by use of the `match with` syntax,

```
functionDeclarationMatchWith.fsx

let rec factorial n =
    match n with
    | 0 -> 1
    | 1 -> 1
    | _ -> n * (factorial (n - 1))

printfn "%d" (factorial 5)

> fsharp -i functionDeclarationMatchWith.fsx
120
```

A short-hand only for functions of 1 parameter is the `function` syntax,

```
functionDeclarationFunction.fsx

let rec factorial = function
    | 0 -> 1
    | 1 -> 1
    | n -> n * (factorial (n - 1))

printfn "%d" (factorial 5)

> fsharp -i functionDeclarationFunction.fsx
120
```

Note that the name given in the match, here `n`, is not used in the first line, and is arbitrary at the line of pattern matching, and may even be different on each line. For these reasons this syntax is discouraged.

Functions may be declared from other functions

functionDeclarationTupleCurrying.fsx

```
let mul (x, y) = x*y
let double y = mul (2.0, y)
printfn "%g" (mul (5.0, 3.0))
printfn "%g" (double 3.0)
```

```
> fsharp -i functionDeclarationTupleCurrying.fsx
15
6
```

For functions of more than 1 argument, there exists a short notation, which is called *currying* in tribute of Haskell Curry, currying

functionDeclarationCurrying.fsx

```
let mul x y = x*y
let double = mul 2.0
printfn "%g" (mul 5.0 3.0)
printfn "%g" (double 3.0)
```

```
> fsharp -i functionDeclarationCurrying.fsx
15
6
```

Here `mul 2.0` is a partial specification of the function `mul x y`, where the first argument is fixed, and hence, `double` is a function of 1 argument being the second argument of `mul`. Currying is often used in functional programming, but generally currying should be used carefully, since currying may seriously reduce readability of code.

4.1 Procedures

A *procedure* is a generalisation of the concept of functions, and in contrast to functions procedures procedure need not return values. An example, we've already seen is the `printfn`, which is used to print text on the console, but does not return a value. Coincidentally, since the console is a state, printing to it is a side-effect. Above we examined

mutableAssignReturnSideEffectStump.fsx

```
let updateFactor factor =
    factor := 2
```

which also does not have a return value. Procedural thinking is useful for encapsulation, but is prone to side-effects and should be minimized by being replaced by functional thinking.

Chapter 5

Controlling program flow

5.0.1 Conditional expressions

```
"if" expr "then" expr  
[{"elif" expr "then" expr}  
"else" expr]
```

A basic flow control mechanism used both for functional and imperative programming is the **if-then-else** construction, e.g.,

flowIfThen.fsx

```
let printOnlyPostiveValues x =  
    if x > 0 then  
        printfn "%d" x  
printOnlyPostiveValues 3  
printOnlyPostiveValues -3  
  
> fsharpit flowIfThen.fsx  
3
```

I.e., if and only if the value of the argument is postive, then it will be printed on screen. More common is to include the **else**

flowIfThenElse.fsx

```
let abs x =  
    if x < 0 then  
        -x  
    else  
        x  
printfn "%d" (abs 3)  
printfn "%d" (abs -3)  
  
> fsharpit flowIfThenElse.fsx  
3  
3
```

A common construction is a nested list of **if-then-else**,

flowIfThenElseNested.fsx

```
let digitToString x =  
    if x < 1 then  
        '0'  
    else  
        if x < 2 then  
            '1'  
        else  
            '2'  
  
printfn "%c" (digitToString 1)  
printfn "%c" (digitToString 3)  
printfn "%c" (digitToString -3)
```

```
> fsharpi flowIfThenElseNested.fsx
```

```
1  
2  
0
```

where the integers 0-2 are converted to characters, and integers outside this domain is converted to the nearest equivalent number. This construction is so common that a short-hand notation exists, and we may equivalently have written,

flowIfThenElseNestedShort.fsx

```
let digitToString x =  
    if x < 1 then  
        '0'  
    elif x < 2 then  
        '1'  
    else  
        '2'  
  
printfn "%c" (digitToString 1)  
printfn "%c" (digitToString 3)  
printfn "%c" (digitToString -3)
```

```
> fsharpi flowIfThenElseNestedShort.fsx
```

```
1  
2  
0
```

5.0.2 For and while loops

A major difference between functional and imperative programming is how loops are expressed. Consider the problem of printing the numbers 1 to 5 on the console with a **while** loop can be done as follows,

flowWhile.fsx

```

let mutable i = 1
while i <= 5 do
    printf "%d " i
    i <- i + 1
printf "\n"

> fsharpi flowWhile.fsx

1 2 3 4 5

```

where the same result by recursion as

flowWhileRecursion.fsx

```

let rec prt a b =
    if a <= b then
        printf "%d " a
        prt (a + 1) b
    else
        printf "\n"
prt 1 5

> fsharpi flowWhileRecursion.fsx

1 2 3 4 5

```

The counting example is so often used that a special notation is available, the **for** loop, where the above could be implemented as

flowFor.fsx

```

for i = 1 to 5 do
    printf "%d " i
printf "\n"

> fsharpi flowFor.fsx

1 2 3 4 5

```

Note that *i* is a value and not a variable here. For a more complicated example, consider the problem of calculating average grades from a list of courses and grades. Using the above construction, this could be performed as,

flowForListsIndex.fsx

```

let courseGrades =
    ["Introduction to programming", 95;
     "Linear algebra", 80;
     "User Interaction", 85;]

let mutable sum = 0;
let mutable n = 0;
for i = 0 to (List.length courseGrades) - 1 do
    let (title, grade) = courseGrades.[i]
    printfn "Course: %s, Grade: %d" title grade

```



```

    sum <- sum + grade;
    n <- n + 1;
let avg = (float sum) / (float n)
printfn "Average grade: %g" avg

```

```

> fsharpi flowForListsIndex.fsx

Course: Introduction to programming, Grade: 95
Course: Linear algebra, Grade: 80
Course: User Interaction, Grade: 85
Average grade: 86.6667

```

However, an elegant alternative is available as

flowForLists.fsx

```

let courseGrades =
    ["Introduction to programming", 95;
     "Linear algebra", 80;
     "User Interaction", 85;]

let mutable sum = 0;
let mutable n = 0;
for (title, grade) in courseGrades do
    printfn "Course: %s, Grade: %d" title grade
    sum <- sum + grade;
    n <- n + 1;
let avg = (float sum) / (float n)
printfn "Average grade: %g" avg

```

```

> fsharpi flowForLists.fsx

Course: Introduction to programming, Grade: 95
Course: Linear algebra, Grade: 80
Course: User Interaction, Grade: 85
Average grade: 86.6667

```

This to be preferred, since we completely can ignore list boundary conditions and hence avoid out of range indexing. For comparison see a recursive implementation of the same,

flowForListsRecursive.fsx

```

let courseGrades =
    ["Introduction to programming", 95;
     "Linear algebra", 80;
     "User Interaction", 85;]

let rec printAndSum lst =
    match lst with
    | (title, grade)::rest ->
        printfn "Course: %s, Grade: %d" title grade
        let (sum, n) = printAndSum rest
        (sum + grade, n + 1)
    | _ -> (0, 0)

```

```
let (sum, n) = printAndSum courseGrades
let avg = (float sum) / (float n)
printfn "Average grade: %g" avg
```

```
> fsharp flowForListsRecursive.fsx

Course: Introduction to programming, Grade: 95
Course: Linear algebra, Grade: 80
Course: User Interaction, Grade: 85
Average grade: 86.6667
```

Note how this implementation avoids the use of variables in contrast to the previous examples.

Chapter 6

Tuples, Lists, Arrays, and Sequences

6.1 Tuples

6.2 Lists

6.3 Arrays

6.3.1 1 dimensional arrays

Roughly speaking, arrays are mutable lists, and may be created and indexed in the same manner, e.g.,

arrayCreation.fsx

```
let A = [| 1; 2; 3; 4; 5 |]
let B = [| 1 .. 5 |]
let C = [| for a in 1 ..5 do yield a |]

let printArray (a : int array) =
    for i = 0 to a.Length - 1 do
        printf "%d " a.[i]
    printf "\n"

printArray A
printArray B
printArray C
```

```
> fsharp arrayCreation.fsx
```

```
1 2 3 4 5
1 2 3 4 5
1 2 3 4 5
```

Notice that as for lists, arrays are indexed starting with 0, and that in this particular case it was necessary to specify the type of the argument for `printArray` as an array of integers with the `array` keyword. The `array` keyword is synonymous with `'[]'`. Arrays do not support pattern matching, cannot be resized, but are mutable,

arrayReassign.fsx

```
let A = [| 1; 2; 3; 4; 5 |]

let printArray (a : int array) =
    for i = 0 to a.Length - 1 do
        printf "%d " a.[i]
    printf "\n"

let square (a : int array) =
    for i = 0 to a.Length - 1 do
        a.[i] <- a.[i] * a.[i]

printArray A
square A
printArray A
```

```
> fsharpi arrayReassign.fsx
```

```
1 2 3 4 5
1 4 9 16 25
```

Notice that in spite of the missing `mutable` keyword, the function `square` still had the side-effect of squaring all entries in `A`. Arrays support *slicing*, that is, indexing an array with a range results in a slicing copy of array with values corresponding to the range, e.g.,

arraySlicing.fsx

```
let A = [| 1; 2; 3; 4; 5 |]
let B = A.[1..3]
let C = A[..2]
let D = A.[3..]
let E = A.[*]

let printArray (a : int array) =
    for i = 0 to a.Length - 1 do
        printf "%d " a.[i]
    printf "\n"

printArray A
printArray B
printArray C
printArray D
printArray E
```

```
> fsharpi arraySlicing.fsx
```

```
1 2 3 4 5
2 3 4
1 2 3
4 5
1 2 3 4 5
```

As illustrated, the missing start or end index implies from the first or to the last element.

There are quite a number of built-in procedures for all arrays some of which we summarize in Table 6.1. Thus, the `arrayReassign.fsx` program can be written using arrays as,

`arrayReassignModule.fsx`

```
let A = [| 1 .. 5 |]

let printArray (a : int array) =
    Array.iter (fun x -> printf "%d " x) a
    printf "\n"

let square a = a * a

printArray A
let B = Array.map square A
printArray A
printArray B
```

```
> fsharpi arrayReassignModule.fsx
1 2 3 4 5
1 2 3 4 5
1 4 9 16 25
```

and the `flowForListsIndex.fsx` program can be written using arrays as,

`flowForListsIndexModule.fsx`

```
let courseGrades =
    ["Introduction to programming", 95;
     "Linear algebra", 80;
     "User Interaction", 85;]

let A = Array.ofList courseGrades
let printCourseNGrade (title, grade) =
    printfn "Course: %s, Grade: %d" title grade
Array.iter printCourseNGrade A
let (titles,grades) = Array.unzip A
let avg = (float (Array.sum grades)) / (float grades.Length)
printfn "Average grade: %g" avg
```

```
> fsharpi flowForListsIndexModule.fsx
Course: Introduction to programming, Grade: 95
Course: Linear algebra, Grade: 80
Course: User Interaction, Grade: 85
Average grade: 86.6667
```

Both cases avoid the use of variables and side-effects which is a big advantage for code safety.

6.3.2 Multidimensional Arrays

Higher dimensional arrays can be created as arrays of arrays (of arrays ...). These are known as *jagged arrays*, since there is no inherent control of that all sub-arrays are of similar size. E.g., the following jagged arrays is a jagged array of increasing width,

append	Creates an array that contains the elements of one array followed by the elements of another array.
average	Returns the average of the elements in an array.
blit	Reads a range of elements from one array and writes them into another.
choose	Applies a supplied function to each element of an array. Returns an array that contains the results <i>x</i> for each element for which the function returns <i>Some(x)</i> .
collect	Applies the supplied function to each element of an array, concatenates the results, and returns the combined array.
concat	Creates an array that contains the elements of each of the supplied sequence of arrays.
copy	Creates an array that contains the elements of the supplied array.
create	Creates an array whose elements are all initially the supplied value.
empty	Returns an empty array of the given type.
exists	Tests whether any element of an array satisfies the supplied predicate.
fill	Fills a range of elements of an array with the supplied value.
filter	Returns a collection that contains only the elements of the supplied array for which the supplied condition returns true.
find	Returns the first element for which the supplied function returns true. Raises <i>System.Collections.Generic.KeyNotFoundException</i> if no such element exists.
findIndex	Returns the index of the first element in an array that satisfies the supplied condition. Raises <i>System.Collections.Generic.KeyNotFoundException</i> if none of the elements satisfy the condition.
fold	Applies a function to each element of an array, threading an accumulator argument through the computation. If the input function is <i>f</i> and the array elements are <i>i0...iN</i> , this function computes <i>f (...(f s i0)...) iN</i> .
foldBack	Applies a function to each element of an array, threading an accumulator argument through the computation. If the input function is <i>f</i> and the array elements are <i>i0...iN</i> , this function computes <i>f i0 (...(f iN s))</i> .
forall	Tests whether all elements of an array satisfy the supplied condition.
isEmpty	Tests whether an array has any elements.
iter	Applies the supplied function to each element of an array.
init	...
length	Returns the length of an array. The <i>System.Array.Length</i> property does the same thing.
map	Creates an array whose elements are the results of applying the supplied function to each of the elements of a supplied array.
mapI	
max	Returns the largest of all elements of an array. <i>Operators.max</i> is used to compare the elements.
min	Returns the smallest of all elements of an array. <i>Operators.min</i> is used to compare the elements.
ofList	Creates an array from the supplied list.
ofSeq	Creates an array from the supplied enumerable object.
partition	Splits an array into two arrays, one containing the elements for which the supplied condition returns true, and the other containing those for which it returns false.
rev	Reverses the order of the elements in a supplied array.
sort	Sorts the elements of an array and returns a new array. <i>Operators.compare</i> is used to compare the elements.
sub	Creates an array that contains the supplied subrange, which is specified by starting index and length.
sum	Returns the sum of the elements in the array.
toList	Converts the supplied array to a list.
toSeq	Views the supplied array as a sequence.
unzip	Splits an array of tuple pairs into a tuple of two arrays.
zeroCreate	Creates an array whose elements are all initially zero.
zip	Combines two arrays into an array of tuples that have two elements. The two arrays must have equal lengths; otherwise, <i>System.ArgumentException</i> is raised.

Table 6.1: Some built-in procedures in the Array module for arrays (from <https://msdn.microsoft.com/en-us/visualfsharpdocs/conceptual/fsharp-core-library-reference>)

arrayJagged.fsx

```
let A = [| for n in 1..3 do yield [| 1 .. n |] |]  
  
let printArrayOfArrays (a : int array array) =  
    for i = 0 to a.Length - 1 do  
        for j = 0 to a.[i].Length - 1 do  
            printf "%d " a.[i].[j]  
        printf "\n"  
  
printArrayOfArrays A  
  
-----  
> fsharpi arrayJagged.fsx  
  
1  
1 2  
1 2 3
```

Indexing arrays of arrays is done sequentially, in the sense that in the above example, the number of outer arrays is `a.Length`, `a.[i]` is the *i*'th array, the length of the *i*'th array is `a.[i].Length`, and the *j*'th element of the *i*'th array is thus `a.[i].[j]`. Often 2 dimensional square arrays are used, which can be implemented as a jagged array as,

arrayJaggedSquare.fsx

```
let pownArray (a : int array) p =  
    for i = 0 to a.Length - 1 do  
        a.[i] <- pown a.[i] p  
    a  
  
let A = [| for n in 1..3 do yield (pownArray [| 1 .. 4 |] n ) |]  
  
let printArrayOfArrays (a : int array array) =  
    for i = 0 to a.Length - 1 do  
        for j = 0 to a.[i].Length - 1 do  
            printf "%2d " a.[i].[j]  
        printf "\n"  
  
printArrayOfArrays A  
  
-----  
> fsharpi arrayJaggedSquare.fsx  
  
1  2  3  4  
1  4  9 16  
1  8 27 64
```

In fact, square arrays of dimensions 2 to 4 are so common that fsharp has built-in modules for their support. In the following describe `Array2D`. The workings of `Array3D` and `Array4D` are very similar. An example of creating the same 2 dimensional array as above but as an `Array2D` is,

array2D.fsx

```
let A = Array2D.create 3 4 0
```

blit	Reads a range of elements from one array and writes them into another.
copy	Creates an array that contains the elements of the supplied array.
create	Creates an array whose elements are all initially the supplied value.
iter	Applies the supplied function to each element of an array.
length1	Returns the length of an array in the first dimension.
length2	Returns the length of an array in the second dimension.
map	Creates an array whose elements are the results of applying the supplied function to each of the elements of a supplied array.
mapi	
zeroCreate	Creates an array whose elements are all initially zero.

Table 6.2: Some built-in procedures in the Array2D module for arrays (from <https://msdn.microsoft.com/en-us/visualfsharpdocs/conceptual/fsharp-core-library-reference>)

```

for i = 0 to (Array2D.length1 A) - 1 do
  for j = 0 to (Array2D.length2 A) - 1 do
    A.[i,j] <- pown (j + 1) (i + 1)

let printArray2D (a : int [,]) =
  for i = 0 to (Array2D.length1 a) - 1 do
    for j = 0 to (Array2D.length2 a) - 1 do
      printf "%2d " a.[i, j]
    printf "\n"

printArray2D A

```

```

> fsharpi array2D.fsx

1  2  3  4
1  4  9 16
1  8 27 64

```

Notice that the indexing uses a slightly different notation '[,]' and the length functions are also slightly different. The statement `A.Length` would return the total number of elements in the array, in this case 12.

There are a bit few built-in procedures for 2 dimensional array types, some of which are summarized in Table 6.2

note that
`A.[1,*]` is
a `Array` but
`A.[1..1,*]` is
an `Array2D`.

6.4 Sequences

Chapter 7

Exceptions

7.1 Exception Handling

Exception handling allows programmers to catch and handle errors whenever an application enters an invalid state.

...

Chapter 8

Testing programs

...

Chapter 9

The Collection

9.1 Mutable Collections

`System.Collections.Generic`

9.1.1 Mutable lists

`List`, `LinkedList`

9.1.2 Stacks

`Stack`

9.1.3 Queues

`Queue`

9.1.4 Sets and dictionaries

`HashSet`, and `Dictionary` from

Chapter 10

Input/Output

Reading and writing to files and the console window.

10.1 Console I/O

...

10.2 File I/O

filenamediialogue.fsx

```
let getAFileName () =
    let mutable filename = Unchecked.defaultof<string>
    let mutable fileExists = false
    while not(fileExists) do
        System.Console.WriteLine("Enter Filename: ")
        filename <- System.Console.ReadLine()
        fileExists <- System.IO.File.Exists filename
    filename

let listOfFiles = System.IO.Directory.GetFiles(".")
printfn "Directory contains: %A" listOfFiles
let filename = getAFileName ()
printfn "You typed: %s" filename
```

readFile.fsx

```
let rec printFile (reader : System.IO.StreamReader) =
    if not(reader.EndOfStream) then
        let line = reader.ReadLine ()
        printfn "%s" line
        printFile reader

let filename = "readFile.fsx"
let reader = System.IO.File.OpenText filename
printFile reader
```

```
> fsharpi readFile.fsx
```

```

let rec printFile (reader : System.IO.StreamReader) =
    if not(reader.EndOfStream) then
        let line = reader.ReadLine ()
        printfn "%s" line
        printFile reader

let filename = "readFile.fsx"
let reader = System.IO.File.OpenText filename
printFile reader

```

reverseFile.fsx

```

let rec readFile (stream : System.IO.StreamReader) =
    if not(stream.EndOfStream) then
        (stream.ReadLine ()) :: (readFile stream)
    else
        []

let rec writeFile (stream : System.IO.StreamWriter) text =
    match text with
    | (l : string) :: ls ->
        stream.WriteLine l
        writeFile stream ls
    | _ -> ()

let reverseString (s : string) =
    System.String(Array.rev (s.ToCharArray()))

let inputStream = System.IO.File.OpenText "reverseFile.fsx"
let text = readFile inputStream
let reverseText = List.map reverseString (List.rev text)
let outputStream = System.IO.File.CreateText "xsf.eliFesrever"
writeFile outputStream reverseText
outputStream.Close ()
printfn "%A" reverseText

```

> fsharp reverseFile.fsx

```

["txeTesrever "A%" nftnirp"; ")( esolC.maertStuptuo";
"txeTesrever maertStuptuo eliFetirw";
""reverseFile.fsx" txeTetaerC.eliF.OI.metsyS = maertStuptuo tel";
")txet ver.tsiL( gnirtSesrever pam.tsiL = txeTesrever tel";
"maertStupni eliFdaer = txet tel";
""xsf.eliFesrever" txeTnepO.eliF.OI.metsyS = maertStupni tel"; "";
"))(yarrArahCoT.s( ver.yarrA(gnirtS.metsyS ";
"= )gnirts : s( gnirtSesrever tel"; ""; ")( >- _ | ";
"sl maerts eliFetirw "; "l eniLetirW.maerts ";
">- sl :: )gnirts : l( | "; "htiw txet hctam ";
"= txet )retirWmaertS.OI.metsyS : maerts( eliFetirw cer tel"; "";
"][ ";
"esle "; ")maerts eliFdaer( :: ))( eniLdaeR.maerts( ";
"neht )maertSf0dnE.maerts(ton fi ";
"= )redaeRmaertS.OI.metsyS : maerts( eliFdaer cer tel"]

```



Chapter 11

Graphical User Interfaces

...

Chapter 12

Imperative programming

12.1 Introduction

Imperativ programming focusses on how a problem is to be solved as a list of *statements* and a set of *states*, where states may change over time. An example is a baking recipe, e.g.,

1. Mix yeast with water
2. Stir in salt, oil, and flour
3. Knead until the dough has a smooth surface
4. Let the dough rise until it has double size
5. Shape dough into a loaf
6. Let the loaf rise until double size
7. Bake in oven until the bread is golden brown

Each line in this example consists of one or more statements that are to be executed, and while executing them states such as size of the dough, color of the bread changes, and some execution will halt execution until certain conditions of these states are fulfilled, e.g., the bread will not be put into the oven for baking before it has risen sufficiently.

Statements may be grouped into procedures, and structuring imperative programs heavily into procedures is called *Procedural programming*, which is sometimes considered as a separate paradigm from imperative programming. *Object oriented programming* is an extension of imperative programming, where statements and states are grouped into classes and will be treated elsewhere.

Almost all computer hardware is designed for *machine code*, which is a common term used for many low-level computer programming language, and almost all machine languages follow the imperative programming paradigm.

Functional programming may be considered a subset of imperative programming, in the sense that functional programming does not include the concept of a state, or one may think of functional programming as only have one unchanging state. Functional programming has also a bigger focus on what should be solved, by declaring rules but not explicitly listing statements describing how these rules should be combined and executed in order to solve a given problem. Functional programming will be treated elsewhere.

Imperativ pro-
gramming
statements
states

Procedural
programming
Object ori-
ented pro-
gramming
machine code
Functional
programming

12.2 Generating random texts

12.2.1 0'th order statistics

randomTextOrder0.fsx

```
let histToCumulativeProbability hist =
    let appendSum (acc : float array) (elem : int) =
        let sum =
            if acc.Length = 0 then
                float elem
            else
                acc.[acc.Length-1] + (float elem)
        Array.append acc [| sum |]

    let normalizeProbability k v = v/k

    let cumSum = Array.fold appendSum Array.empty<float> hist
    if cumSum.[cumSum.Length - 1] > 0.0 then
        Array.map (normalizeProbability cumSum.[cumSum.Length - 1])
            cumSum
    else
        Array.create cumSum.Length (1.0 / (float cumSum.Length))

let lookup (hist : float array) (v : float) =
    Array.findIndex (fun h -> h > v) hist

let countEqual A v =
    Array.fold (fun acc elem -> if elem = v then acc+1 else acc) 0 A

let intToIdx i = i - (int ' ')

let idxToInt i = i + (int ' ')

let legalIndex size idx =
    (0 <= idx) && (idx <= size - 1)

let analyzeFile (reader : System.IO.StreamReader) size pushForward
    =
    let hist = Array.create size 0
    let mutable c = Unchecked.defaultof<int>
    while not(reader.EndOfStream) do
        c <- pushForward (reader.Read ())
        if legalIndex size c then
            hist.[c] <- hist.[c] + 1
    hist

let sampleFromCumulativeProbability cumulative noSamples =
    let rnd = System.Random ()
    let rndArray = Array.init noSamples (fun _ -> rnd.NextDouble ())
    Array.map (lookup cumulative) rndArray

let filename = "randomTextOrder0.fsx"
let noSamples = 200
let histSize = 126 - 32 + 1

let reader = System.IO.File.OpenText filename
```

```

let hist = analyzeFile reader histSize intToIdx
reader.Close ()
let idxValue = Array.mapi (fun i v -> (idxToInt i, v)) hist
printfn "%A" idxValue
printfn "%d zeros out of %d elements" (countEqual hist 0) hist.
    Length
let cumulative = histToCumulativeProbability hist
let rndIdx = sampleFromCumulativeProbability cumulative noSamples
let rndInt = Array.map idxToInt rndIdx
let rndChar = Array.map (fun v -> char v) rndInt
let text = System.String.Concat rndChar // System.String is not the
    same as String !!!
printfn "%A" text

```

> fsharpi randomTextOrder0.fsx

```

[(32, 339); (33, 3); (34, 8); (35, 0); (36, 0); (37, 4); (38, 2);
(39, 4);
(40, 27); (41, 27); (42, 0); (43, 5); (44, 1); (45, 13); (46, 42);
; (47, 4);
(48, 11); (49, 9); (50, 3); (51, 1); (52, 0); (53, 0); (54, 1);
(55, 0);
(56, 0); (57, 0); (58, 5); (59, 0); (60, 6); (61, 31); (62, 9);
(63, 0);
(64, 0); (65, 19); (66, 0); (67, 8); (68, 1); (69, 3); (70, 7);
(71, 0);
(72, 0); (73, 14); (74, 0); (75, 0); (76, 7); (77, 0); (78, 1);
(79, 5);
(80, 6); (81, 0); (82, 3); (83, 26); (84, 9); (85, 1); (86, 2);
(87, 0);
(88, 0); (89, 0); (90, 0); (91, 6); (92, 0); (93, 6); (94, 0);
(95, 1);
(96, 0); (97, 100); (98, 14); (99, 41); (100, 50); (101, 143);
(102, 29);
(103, 12); (104, 38); (105, 87); (106, 0); (107, 5); (108, 89);
(109, 63);
(110, 75); (111, 50); (112, 26); (113, 2); (114, 88); (115, 47);
(116, 121);
(117, 55); (118, 18); (119, 3); (120, 21); (121, 31); (122, 12);
(123, 0);
(124, 2); (125, 0); (126, 0)]
29 zeros out of 95 elements
"e=simm yAo"ten ihynail Ohond etsvtrmxo o= .lmi ue.flyatr e
uxhtrtLuitSlrty iceiulotru.tg idn pm drame (ionzA>ait >)l
sidahtn( r(aeym.. nn.(et1nlxflrh i'uPu damtrtrauls eusl g s.
olcyi]daeCt)g0sei [t."

```

12.2.2 1'th order statistics

randomTextOrder0.fsx

```

let histToCumulativeProbability hist =
    let appendSum (acc : float array) (elem : int) =

```

```

    let sum =
        if acc.Length = 0 then
            float elem
        else
            acc.[acc.Length-1] + (float elem)
    Array.append acc [| sum |]

let normalizeProbability k v = v/k

let cumSum = Array.fold appendSum Array.empty<float> hist
if cumSum.[cumSum.Length - 1] > 0.0 then
    Array.map (normalizeProbability cumSum.[cumSum.Length - 1])
        cumSum
else
    Array.create cumSum.Length (1.0 / (float cumSum.Length))

let lookup (hist : float array) (v : float) =
    Array.findIndex (fun h -> h > v) hist

let countEqual A v =
    Array.fold (fun acc elem -> if elem = v then acc+1 else acc) 0 A

let intToIdx i = i - (int ' ')

let idxToInt i = i + (int ' ')

let legalIndex size idx =
    (0 <= idx) && (idx <= size - 1)

let analyzeFile (reader : System.IO.StreamReader) size pushForward
    =
    let hist = Array.create size 0
    let mutable c = Unchecked.defaultof<int>
    while not(reader.EndOfStream) do
        c <- pushForward (reader.Read ())
        if legalIndex size c then
            hist.[c] <- hist.[c] + 1
    hist

let sampleFromCumulativeProbability cumulative noSamples =
    let rnd = System.Random ()
    let rndArray = Array.init noSamples (fun _ -> rnd.NextDouble ())
    Array.map (lookup cumulative) rndArray

let filename = "randomTextOrder0.fsx"
let noSamples = 200
let histSize = 126 - 32 + 1

let reader = System.IO.File.OpenText filename
let hist = analyzeFile reader histSize intToIdx
reader.Close ()
let idxValue = Array.mapi (fun i v -> (idxToInt i, v)) hist
printfn "%A" idxValue

```

```

printfn "%d zeros out of %d elements" (countEqual hist 0) hist.
    Length
let cumulative = histToCumulativeProbability hist
let rndIdx = sampleFromCumulativeProbability cumulative noSamples
let rndInt = Array.map idxToInt rndIdx
let rndChar = Array.map (fun v -> char v) rndInt
let text = System.String.Concat rndChar // System.String is not the
    same as String !!!
printfn "%A" text

```

```

> fsharpi randomTextOrder0.fsx

```

```

[|(32, 339); (33, 3); (34, 8); (35, 0); (36, 0); (37, 4); (38, 2);
  (39, 4);
  (40, 27); (41, 27); (42, 0); (43, 5); (44, 1); (45, 13); (46, 42)
    ; (47, 4);
  (48, 11); (49, 9); (50, 3); (51, 1); (52, 0); (53, 0); (54, 1);
    (55, 0);
  (56, 0); (57, 0); (58, 5); (59, 0); (60, 6); (61, 31); (62, 9);
    (63, 0);
  (64, 0); (65, 19); (66, 0); (67, 8); (68, 1); (69, 3); (70, 7);
    (71, 0);
  (72, 0); (73, 14); (74, 0); (75, 0); (76, 7); (77, 0); (78, 1);
    (79, 5);
  (80, 6); (81, 0); (82, 3); (83, 26); (84, 9); (85, 1); (86, 2);
    (87, 0);
  (88, 0); (89, 0); (90, 0); (91, 6); (92, 0); (93, 6); (94, 0);
    (95, 1);
  (96, 0); (97, 100); (98, 14); (99, 41); (100, 50); (101, 143);
    (102, 29);
  (103, 12); (104, 38); (105, 87); (106, 0); (107, 5); (108, 89);
    (109, 63);
  (110, 75); (111, 50); (112, 26); (113, 2); (114, 88); (115, 47);
    (116, 121);
  (117, 55); (118, 18); (119, 3); (120, 21); (121, 31); (122, 12);
    (123, 0);
  (124, 2); (125, 0); (126, 0)|]
29 zeros out of 95 elements
"e=simm yAo"ten ihynail Ohond etsvtrmxo o= .lmi ue.flyatr e
  uxhtrtLuitSlrty iceiulotru.tg idn pm drame (ionzA>ait >)l
  sidahtn( r(aeym.. nn.(et1nlxflrh i'uPu damtrtrauls eusl g s.
  olcyi]daeCt)g0sei [t."

```

Chapter 13

Functional programming

Chapter 14

Object-oriented programming

Appendix A

Number systems on the computer

A.1 Binary numbers

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base 10* means that for a number consisting of a sequence of digits separated by a *decimal point*, where each *digit* can have values $d \in \{0, 1, 2, \dots, 9\}$ and the weight of each digit is proportional to its place in the sequence of digits w.r.t. the decimal point, i.e., the number $357.6 = 3 \cdot 10^2 + 5 \cdot 10^1 + 7 \cdot 10^0 + 6 \cdot 10^{-1}$ or in general:

decimal num-
ber
base
decimal point
digit

$$v = \sum_{i=-m}^n d_i 10^i \quad (\text{A.1})$$

The basic unit of information in almost all computers is the binary digit or *bit* for short. A *binary* number consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. The general equation is,

bit
binary

$$v = \sum_{i=-m}^n b_i 2^i \quad (\text{A.2})$$

and examples are $1011.1_2 = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 + 1 \cdot 2^{-1} = 11.5$. Notice that we use subscript 2 to denote a binary number, while no subscript is used for decimal numbers. Due to typical organization of computer memory, 8 binary digits is called a *byte*, and 32 digits a *word*.

byte

Other number systems are often used, e.g., *octal* numbers, which are base 8 numbers, where each digit is $o \in \{0, 1, \dots, 7\}$. Octals are useful short-hand for binary, since 3 binary digits maps to the set of octal digits. Likewise, *hexadecimal* numbers are base 16 with digits $h \in \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, b, c, d, e, f\}$ such that $a_{16} = 10$, $b_{16} = 11$ and so on. Hexadecimals are convenient since 4 binary digits map directly to the set of octal digits. Thus $367 = 101101111_2 = 557_8 = 16f_{16}$. A list of the intergers 0–63 is various bases is given in Table A.1.

word
octal
hexadecimal

A.2 IEEE 754 floating point standard

The set of real numbers also called *reals* includes all fractions and irrational numbers. It is infinite in size both in the sense that there is no largest nor smallest number and between any 2 given numbers there are infinitely many numbers. Reals are widely used for calculation, but since any computer only has finite memory, it is impossible to represent all possible reals. Hence, any computation performed on a computer with reals must rely on approximations. *IEEE 754 double precision floating-point format* (*binary64*), known as a *double*, is a standard for representing an approximation of reals using 64 bits. These bits are divided into 3 parts: sign, exponent and fraction,

reals

IEEE 754 dou-
ble precision
floating-point
format
binary64
double

$$s e_1 e_2 \dots e_{11} m_1 m_2 \dots m_{52},$$

Dec	Bin	Oct	Hex	Dec	Bin	Oct	Hex
0	0	0	0	32	100000	40	20
1	1	1	1	33	100001	41	21
2	10	2	2	34	100010	42	22
3	11	3	3	35	100011	43	23
4	100	4	4	36	100100	44	24
5	101	5	5	37	100101	45	25
6	110	6	6	38	100110	46	26
7	111	7	7	39	100111	47	27
8	1000	10	8	40	101000	50	28
9	1001	11	9	41	101001	51	29
10	1010	12	a	42	101010	52	2a
11	1011	13	b	43	101011	53	2b
12	1100	14	c	44	101100	54	2c
13	1101	15	d	45	101101	55	2d
14	1110	16	e	46	101110	56	2e
15	1111	17	f	47	101111	57	2f
16	10000	20	10	48	110000	60	30
17	10001	21	11	49	110001	61	31
18	10010	22	12	50	110010	62	32
19	10011	23	13	51	110011	63	33
20	10100	24	14	52	110100	64	34
21	10101	25	15	53	110101	65	35
22	10110	26	16	54	110110	66	36
23	10111	27	17	55	110111	67	37
24	11000	30	18	56	111000	70	38
25	11001	31	19	57	111001	71	39
26	11010	32	1a	58	111010	72	3a
27	11011	33	1b	59	111011	73	3b
28	11100	34	1c	60	111100	74	3c
29	11101	35	1d	61	111101	75	3d
30	11110	36	1e	62	111110	76	3e
31	11111	37	1f	63	111111	77	3f

Table A.1: A list of the intergers 0–63 in decimal, binary, octal, and hexadecimal.

where s , e_i , and m_j are binary digits. The bits are converted to a number using the equation by first calculating the exponent e and the mantissa m ,

$$e = \sum_{i=1}^{11} e_i 2^{11-i}, \quad (\text{A.3})$$

$$m = \sum_{j=1}^{52} m_j 2^{-j}. \quad (\text{A.4})$$

I.e., the exponent is an integer, where $0 \leq e < 2^{11}$, and the mantissa is a rational, where $0 \leq m < 1$. For most combinations of e and m the real number v is calculated as,

$$v = (-1)^s (1 + m) 2^{e-1023} \quad (\text{A.5})$$

with the exception that

	$m = 0$	$m \neq 0$
$e = 0$	$v = (-1)^s 0$ (signed zero)	$v = (-1)^s m 2^{1-1023}$ (subnormals)
$e = 2^{11} - 1$	$v = (-1)^s \infty$	$v = (-1)^s \text{NaN}$ (not a number)

subnormals
NaN
not a number

where $e = 2^{11} - 1 = 11111111111_2 = 2047$. The largest and smallest number that is not infinity is thus

$$e = 2^{11} - 2 = 2046 \quad (\text{A.6})$$

$$m = \sum_{j=1}^{52} 2^{-j} = 1 - 2^{-52} \simeq 1. \quad (\text{A.7})$$

$$v_{\max} = \pm (2 - 2^{-52}) 2^{1023} \simeq \pm 2^{1024} \simeq \pm 10^{308} \quad (\text{A.8})$$

The density of numbers varies in such a way that when $e - 1023 = 52$, then

$$v = (-1)^s \left(1 + \sum_{j=1}^{52} m_j 2^{-j} \right) 2^{52} \quad (\text{A.9})$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{-j} 2^{52} \right) \quad (\text{A.10})$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{52-j} \right) \quad (\text{A.11})$$

$$\stackrel{k=52-j}{=} \pm \left(2^{52} + \sum_{k=51}^0 m_{52-k} 2^k \right) \quad (\text{A.12})$$

which are all integers in the range $2^{52} \leq |v| < 2^{53}$. When $e - 1023 = 53$, then the same calculation gives

$$v \stackrel{k=53-j}{=} \pm \left(2^{53} + \sum_{k=52}^1 m_{53-k} 2^k \right) \quad (\text{A.13})$$

which are every second integer in the range $2^{53} \leq |v| < 2^{54}$, and so on for larger e . When $e - 1023 = 51$, then the same calculation gives,

$$v \stackrel{k=51-j}{=} \pm \left(2^{51} + \sum_{k=50}^{-1} m_{51-k} 2^k \right) \quad (\text{A.14})$$

which gives a distance between numbers of $1/2$ in the range $2^{51} \leq |v| < 2^{52}$, and so on for smaller e . Thus we may conclude that the distance between numbers in the interval $2^n \leq |v| < 2^{n+1}$ is 2^{n-52} , for $-1022 = 1 - 1023 \leq n < 2046 - 1023 = 1023$. For subnormals the distance between numbers are

$$v = (-1)^s \left(\sum_{j=1}^{52} m_j 2^{-j} \right) 2^{-1022} \quad (\text{A.15})$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j} 2^{-1022} \right) \quad (\text{A.16})$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j-1022} \right) \quad (\text{A.17})$$

$${}^{k=-j-1022}_{=} \pm \left(\sum_{j=-1023}^{-1074} m_{-k-1022} 2^k \right) \quad (\text{A.18})$$

which gives a distance between numbers of $2^{-1074} \simeq 10^{-323}$ in the range $0 < |v| < 2^{-1022} \simeq 10^{-308}$.

Appendix B

Commonly used character sets

Letters, digits, symbols and space are the core of how we store data, write programs, and communicate with computers and each others. These symbols are in short called characters, and represents a mapping between numbers, also known as codes, and a pictorial representation of the character. E.g., the ASCII code for the letter 'A' is 65. These mappings are for short called character sets, and due to differences in natural languages and symbols used across the globe, many different character sets are in use. E.g., the English alphabet contains the letters 'a' to 'z', which is shared by many other European languages, but which have other symbols and accents for example, Danish has further the letters 'æ', 'ø', and 'å'. Many non-european languages have completely different symbols, where Chinese character set is probably the most extreme, where some definitions contains 106,230 different characters albeit only 2,600 are included in the official Chinese language test at highest level.

Presently, the most common character set used is Unicode Transformation Format (UTF), whose most popular encoding schemes are 8-bit (UTF-8) and 16-bit (UTF-16). Many other character sets exists, and many of the later builds on the American Standard Code for Information Interchange (ASCII). The ISO-8859 codes were an intermediate set of character sets that are still in use, but which is greatly inferior to UTF. Here we will briefly give an overview of ASCII, ISO-8859-1 (Latin1), and UTF.

B.1 ASCII

The *American Standard Code for Information Interchange* (ASCII) [1], is a 7 bit code tuned for the letters of the english language, numbers, punctuation symbols, control codes and space, see Tables B.1 and B.2. The first 32 codes are reserved for non-printable control characters to control printers and similar devices or to provide meta-information. The meaning of each control characters is not universally agreed upon.

The code order is known as *ASCIIbetical order*, and it is sometimes used to perform arithmetic on codes, e.g., an upper case letter with code c may be converted to lower case by adding 32 to its code. The ASCIIbetical order also has consequence for sorting, i.e., when sorting characters according to their ASCII code, then 'A' comes before 'a', which comes before the symbol '{'.

American
Standard
Code for In-
formation
Interchange
ASCII
ASCIIbetical
order

B.2 ISO/IEC 8859

The ISO/IEC 8859 report http://www.iso.org/iso/catalogue_detail?csnumber=28245 defines 10 sets of codes specifying up to 191 codes and graphic characters using 8 bits. Set 1 also known as ISO/IEC 8859-1, Latin alphabet No. 1, or *Latin1* covers many European languages and is designed to be compatible with ASCII, such that code for the printable characters in ASCII are the same in ISO 8859-1. In Table B.3 is shown the characters above 7e. Codes 00-1f and 7f-9f are undefined in ISO 8859-1.

Latin1

x0+0x	00	10	20	30	40	50	60	70
00	NUL	DLE	SP	0	@	P	'	p
01	SOH	DC1	!	1	A	Q	a	q
02	STX	DC2	"	2	B	R	b	r
03	ETX	DC3	#	3	C	S	c	s
04	EOT	DC4	\$	4	D	T	d	t
05	ENQ	NAK	%	5	E	U	e	u
06	ACK	SYN	&	6	F	V	f	v
07	BEL	ETB	,	7	G	W	g	w
08	BS	CAN	(8	H	X	h	x
09	HT	EM)	9	I	Y	i	y
0A	LF	SUB	*	:	J	Z	j	z
0B	VT	ESC	+	;	K	[k	{
0C	FF	FS	,	<	L	\	l	
0D	CR	GS	-	=	M]	m	}
0E	SO	RS	.	>	N	^	n	~
0F	SI	US	/	?	O	_	o	DEL

Table B.1: ASCII

B.3 Unicode

Unicode is a character standard defined by the Unicode Consortium, <http://unicode.org> as the *Unicode Standard*. Unicode allows for 1,114,112 different codes. Each code is called a *code point*, which represents an abstract character. However, not all abstract characters requires a unit of several code points to be specified. Code points are divided into 17 planes each with $2^{16} = 65,536$ code points. Planes are further subdivided into named *blocks*. The first plane is called the *Basic Multilingual plane* and it are the first 128 code points is called the *Basic Latin block* and are identical to ASCII, see Table B.1, and code points 128-255 is called the *Latin-1 Supplement block*, and are identical to the upper range of ISO 8859-1, see Table B.3. Presently more than 128,000 code points covering 135 modern and historic writing systems.

A unicode code point is an abstraction from the encoding and the graphical representation of a character. A code point is written as “U+” followed by its hexadecimal number, and for the Basic Multilingual plane 4 digits are used, e.g., the code point with the unique name LATIN CAPITAL LETTER A has the unicode code point is “U+0041”, and in this text it is visualized as ‘A’. More digits are used for code points of the remaining planes.

To store and retrieve code points, they must be encoded and decoded. A common encoding is *UTF-8*, which encodes code points as 1 to 4 bytes, and which is backward-compatible with ASCII and ISO 8859-1. Hence, in all 3 coding systems the character with code 65 represents the character ‘A’. Another popular encoding scheme is *UTF-16*, which encodes characters as 2 or 4 bytes, but which is not backward-compatible with ASCII or ISO 8859-1. UTF-16 is used internally in many compiles, interpreters and operating systems.

Unicode Stan-
dard
code point
blocks
Basic Multilin-
gual plane
Basic Latin
block
Latin-1 Sup-
plement block

UTF-8

UTF-16

Code	Description
NUL	Null
SOH	Start of heading
STX	Start of text
ETX	End of text
EOT	End of transmission
ENQ	Enquiry
ACK	Acknowledge
BEL	Bell
BS	Backspace
HT	Horizontal tabulation
LF	Line feed
VT	Vertical tabulation
FF	Form feed
CR	Carriage return
SO	Shift out
SI	Shift in
DLE	Data link escape
DC1	Device control one
DC2	Device control two
DC3	Device control three
DC4	Device control four
NAK	Negative acknowledge
SYN	Synchronous idle
ETB	End of transmission block
CAN	Cancel
EM	End of medium
SUB	Substitute
ESC	Escape
FS	File separator
GS	Group separator
RS	Record separator
US	Unit separator
DEL	Delete

Table B.2: ASCII symbols.

x0+0x	80	90	A0	B0	C0	D0	E0	F0
00			NBSP	°	À	Ð	à	ð
01			¡	±	Á	Ñ	á	ñ
02			¢	²	Â	Ò	â	ò
03			£	³	Ã	Ó	ã	ó
04			¤	´	Ä	Ô	ä	ô
05			¥	µ	Å	Õ	å	õ
06			¦	¶	Æ	Ö	æ	ö
07			§	·	Ç	×	ç	÷
08			¨	, ¸	È	Ø	è	ø
09			©	¹	É	Ù	é	ù
0a			ª	º	Ê	Ú	ê	ú
0b			«	»	Ë	Û	ë	û
0c			¬	$\frac{1}{4}$	Ì	Ü	ì	ü
0d			SHY	$\frac{1}{2}$	Í	Ý	í	ý
0e			®	$\frac{3}{4}$	Î	Þ	î	þ
0f			¯	¸	Ï	ß	ï	ÿ

Table B.3: ISO-8859-1 (latin1) non-ASCII part. Note that the codes 7f – 9f are undefined.

Code	Description
NBSP	Non-breakable space
SHY	Soft hyphen

Table B.4: ISO-8859-1 special symbols.

Appendix C

A brief introduction to Extended Backus-Naur Form

Extended Backus-Naur Form (EBNF) is a language to specify programming languages in. The name is a tribute to John Backus who used it to describe the syntax of ALGOL58 and Peter Naur for his work on ALGOL 60.

An EBNF consists of *terminal symbols* and *production rules*. Examples of typical terminal symbol are characters, numbers, punctuation marks, and whitespaces, e.g.,

```
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
```

Extended
Backus-Naur
Form
EBNF
terminal sym-
bols
production
rules

A production rule specifies a method of combining other production rules and terminal symbols, e.g.,

```
number = { digit } ;
```

A proposed standard for EBNF (proposal ISO/IEC 14977, <http://www.cl.cam.ac.uk/~mgk25/iso-14977.pdf>) is,

'=' definition, e.g.,

```
zero = "0" ;
```

here **zero** is the terminal symbol 0.

',' concatenation, e.g.,

```
one = "1" ;  
eleven = one , one ;
```

here **eleven** is the terminal symbol 11.

',' termination of line

'|' alternative options, e.g.,

```
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
```

here `digit` is the single character terminal symbol, such as 3.

'[...]' optional, e.g.,

```
zero = "0" ;
nonZeroDigit = "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
nonZero = [ zero ], nonZeroDigit
```

here `nonZero` is a non-zero digit possibly preceded by zero, such as 02.

'{ ... }' repetition zero or more times, e.g.,

```
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
number = digit, { digit }
```

here `number` is a word consisting of 1 or more digits, such as 12.

'(...)' grouping, e.g.,

```
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
number = digit, { digit }
expression = number, { ( "+" | "-" ), number };
```

here `expression` is a number or a sum of numbers such as 3 + 5.

'" ... "' a terminal string, e.g.,

```
string = "abc" ;
```

'" ... "' a terminal string, e.g.,

```
string = 'abc' ;
```

'(* ... *)' a comment (* ... *)

```
(* a binary digit *) digit = "0" | "1" (* from this all numbers
    may be constructed *) ;
```

Everything inside the comments are not part of the formal definition.

'? ... ?' special sequence, a notation reserved for future extensions of EBNF.

'-' exception, e.g.,


```

letter = "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H"
        | "I" | "J" | "K" | "L" | "M" | "N" | "O" | "P" | "Q"
        | "R" | "S" | "T" | "U" | "V" | "W" | "X" | "Y" | "Z" ;
vowel = "A" | "E" | "I" | "O" | "U" ;
consonant = letter - vowel ;

```

here `consonant` are all letters except vowels.

The proposal allows for identifiers that includes space, but often a reduced form is used, where identifiers are single words, in which case the concatenation symbol `,` is omitted. Likewise, the termination symbol `;` is often replaced with the new-line character, and if long lines must be broken, then indentation is used to signify continuation.

In this relaxed EBNF, the EBNF syntax itself can be expressed in EBNF as,

```

letter = "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H"
        | "I" | "J" | "K" | "L" | "M" | "N" | "O" | "P" | "Q"
        | "R" | "S" | "T" | "U" | "V" | "W" | "X" | "Y" | "Z"
        | "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h"
        | "i" | "j" | "k" | "l" | "m" | "n" | "o" | "p" | "q"
        | "r" | "s" | "t" | "u" | "v" | "w" | "x" | "y" | "z"
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
symbol = "[" | "]" | "{" | "}" | "(" | ")" | "<" | ">"
        | "'" | '"' | "=" | "|" | "." | "," | ";"
underscore = "_"
identifier = letter { letter | digit | underscore }
character = letter | digit | symbol | underscore
string = character { character }
terminal = '"' string '"' | "'" string "'"
rhs = identifier
      | terminal
      | "[" rhs "]"
      | "{" rhs "}"
      | "(" rhs ")"
      | rhs "|" rhs
(* | rhs "," rhs *)
rule = identifier "=" rhs (* ";" *)
grammar = rule { rule }

```

Here the comments demonstrate, the relaxed modification.

Bibliography

- [1] X3: ASA Sectional Committee on Computers and Information Processing. American standard code for information interchange. Technical Report ASA X3.4-1963, American Standards Association (ASA), 1963. <http://worldpowersystems.com/projects/codes/X3.4-1963/>.

Index

- American Standard Code for Information Inter-
change, 50
- ASCII, 50
- ASCIIbetical order, 50
- base, 46
- Basic Latin block, 51
- Basic Multilingual plane, 51
- binary, 46
- binary numbers, 8
- binary64, 46
- bit, 46
- blocks, 51
- byte, 46
- code point, 51
- compiled, 3
- console, 3
- currying, 20
- decimal number, 46
- decimal numbers, 8
- decimal point, 46
- digit, 46
- double, 46
- down casting, 9
- EBNF, 54
- encapsulation, 15
- executable file, 3
- expression, 6
- Extended Backus-Naur Form, 54
- format string, 6
- Functional programming, 39
- hexadecimal, 46
- IEEE 754 double precision floating-point format,
46
- Imperativ programming, 39
- implementation file, 3
- interactive, 3
- jagged arrays, 28
- keyword, 6
- Latin-1 Supplement block, 51
- Latin1, 50
- machine code, 39
- modules, 3
- Mutable data, 13
- NaN, 48
- not a number, 48
- Object oriented programming, 39
- octal, 46
- operands, 12
- operator, 12
- primitive types, 6
- Procedural programming, 39
- procedure, 20
- production rules, 54
- reals, 46
- reference cells, 16
- scope, 10
- script file, 3
- script-fragments, 3
- side-effects, 16
- signature file, 3
- slicing, 27
- statement, 6
- statements, 39
- states, 39
- string, 6
- subnormals, 48
- terminal symbols, 54
- type, 6
- type cast, 9
- type inference, 5
- Unicode Standard, 51
- unit-testing, 5
- UTF-16, 51
- UTF-8, 51
- variable, 13
- word, 46