Learning to Program with F#

Jon Sporring

Department of Computer Science,
University of Copenhagen

2018-10-04 14:37:52+02:00

Contents

Contents

Contents

1 | Preface

This book has been written as an introduction to programming for novice programmers. It is used in the first programming course at the University of Copenhagen's bachelor in computer science program. It has been typeset in LaTeX, and all programs have been developed and tested in Mono version 5.10.1.57.

Jon Sporring
Associate Professor, Ph.d.
Department of Computer Science,
University of Copenhagen
2018-10-04

2 | Introduction

Programming is a creative process in which exciting problems may be solved and new tools and applications may be created. With programming skills, you can create high-level applications to run on a mobile device that interact with other users, databases, and artificial intelligence; you may create programs that run on supercomputers for simulating weather systems on alien planets or social phenomena in the internet economy; and you may create programs that run on small custom-made hardware for controlling your home appliances.

2.1. How to Learn to Solve Problems by Programming

In order to learn how to program, there are a couple of steps that are useful:

1. Choose a programming language: A programming language, such as F#, is a vocabulary and a set of grammatical rules for instructing a computer to perform a certain task. It is possible to program without a concrete language, but your ideas and thoughts must still be expressed in some fairly rigorous way. Theoretical computer scientists typically do not rely on computers nor programming languages but uses mathematics to prove properties of algorithms. However, most computer scientists program using a computer, and with a real language you have the added benefit of checking your algorithm, and hence your thoughts, rigorously on a real computer. This book teaches a subset of F#. The purpose is not to be a reference guide to this language but to use it as a vessel to teach you, the reader, how to convert your ideas into programs.

2. Learn the language: A computer language is a structure for thought, and it influences which thoughts you choose to express as a program, and how you choose to do it. Any conversion requires you to acquire a sufficient level of fluency in order for you to be able to make programs. You do not need to be a master in F# nor to know every corner of the language, and you will expand your knowledge as you expose yourself to solving problems in the language, but you must invest an initial amount of time and energy in order to learn the basics of the language. This book aims at getting you started quickly, which is why we intentionally teach just a small subset of F#. On the internet and through other works you will be able to learn much more.

3. Practice: In order to be a good programmer, the most essential step is: practice, practice, practice! It has been estimated that to master anything, then you have to have spent at least 10000 hours practicing, so get started logging hours! It of course matters, how you practice. This book teaches a number of different programming themes. The point is that programming is thinking, and the scaffold you use shapes

your thoughts. It is therefore important to recognize this scaffold and to have the ability to choose one which suits your ideas and your goals best. The best way to expand your abilities is to sharpen your present abilities, push yourself into new territory, and try something new. Do not be afraid to make errors or be frustrated at first. These are the experiences that make you grow.

4. Solve real problems: I have found that using my programming skills in real situations with customers demanding specific solutions, has forced me to put the programming tools and techniques that I use into perspective. Sometimes a task requires a cheap and fast solution, other times customers want a long-perspective solution with bug fixes, upgrades, and new features. Practicing solving real problems helps you strike a balance between the two when programming. It also allows makes you a more practical programmer, by allowing you to recognize its applications in your everyday experiences. Regardless, real problems create real programmers.

2.2. How to Solve Problems

Programming is the act of solving a problem by writing a program to be executed on a computer. A general method for solving problems, given by George Pólya [9] and adapted to programming, is:

Understand the problem: To solve any problem it is crucial that the problem formulation is understood. What is to be solved? Do you understand everything in the description of the problem? Is all information for finding the solution available or is something missing?

Design a plan: Good designs lead to programs are faster to implement, easier to find errors in, and easier to update in the future. Before you start typing a program consider things like: What are the requirements and constraints for the program? Which components should the program have? How are these components supposed to work together? Designing often involves drawing a diagram of the program and writing program sketches on paper.

Implement the plan: Implementation is the act of transforming a program design into code. A crucial part of any implementation is choosing which programming language to use. Furthermore, the solution to many problems will have a number of implementations which vary in how much code they require, to which degree they rely on external libraries, which programming style they are best suited for, what machine resources they require, and how long time they take to run on a computer. With a good design, the coding is usually easy, since the design will have uncovered the major issues and found solutions for these, but sometimes implementation reveals new problems, which require rethinking the design. Most often the implementation step also require a careful documentation of key aspects of the code, e.g., a user manual for the user, and internal notes for fellow programmers that are to maintain and update the code in the future.

Reflect on the result: A crucial part of any programming task is ensuring that the program solves the problem sufficiently. Ask yourself questions such as: What are the program's errors, is the documentation of the code sufficient and relevant for its intended use? Is the code easily maintainable and extendable by other programmers? Which parts of your method would you avoid or replicate in future programming sessions? Can you reuse some of the code you developed in other programs?

Programming is a very complicated process, and Pólya's list is a useful guide but not a fail-safe approach. Always approach problem-solving with an open mind.

2.3. Approaches to Programming

This book focuses on several fundamentally different approaches to programming:

Imperative programming emphasizes *how a program shall accomplish a solution* and focusses less on *what the solution is*. A cooking recipe is an example of the spirit of imperative programming, where the recipe emphasizes what should be done in each step rather than describing the result. For example, a recipe for bread might tell you to first mix yeast and water, then add flour, etc. In imperative programming what should be done are called *statements* and in the recipe analogy, the steps are the statements. Statements influence the computer's *states*, in the same way that adding flour changes the state of our dough. Almost all computer hardware is designed to execute low-level programs written in imperative style. Imperative programming builds on the Turing machine [10]. As a historical note, the first major language was FORTRAN [6] which emphasized an imperative style of programming.

· imperative programming

· statement
· state

Declarative programming emphasizes *what a program shall accomplish* but not *how*. We will consider Functional programming as an example of declarative programming. A *functional programming* language evaluates *functions* and avoids state changes. The program consists of *expressions* instead of statements. As an example, the function $f(x) = x^2$ takes a number x, evaluates the expression x^2, and returns the resulting number. Everything about the function may be characterized by the relation between the input and output values. Functional programming has its roots in lambda calculus [1]. The first language emphasizing functional programming was Lisp [7].

· declarative programming

· functional programming
· function
· expression

Structured programming emphasizes organization of programs in units of code and data. For example, a traffic light may consist of a state (red, yellow, green), and code for updating the state, i.e., switching from one color to the next. We will focus on Object-oriented programming as the example of structured programming. *Object-oriented programming* is a type of programming, where the code and data are structured into *objects*. E.g., a traffic light may be an object in a traffic-routing program. The first object-oriented programming language was Simula 67 developed by Dahl and Nygaard at the Norwegian Computing Center in Oslo [2].

· structured programming

· object-oriented programming
· object

Event-driven programming, which is often used when dynamically interacting with the real world. This is useful, for example, when programming graphical user interfaces, where programs will often need to react to a user clicking on the mouse or to text arriving from a web-server to be displayed on the screen. Event-driven programs are often programmed using *call-back functions*, which are small programs that are ready to run when events occur.

· event-driven programming

· call-back functions

Most programs do not follow a single programming paradigm as, e.g., one of the above, but are a mix. Nevertheless, this book will treat each paradigm separately to emphasize its advantages and disadvantages.

2.4. Why Use F#

This book uses F#, also known as Fsharp, which is a functional first programming language, meaning that it is designed as a functional programming language that also supports imperative and object-oriented programming. It was originally developed for Microsoft's .Net platform but is available as open source for many operating systems through Mono. As an introduction to programming, F# is a young programming language still under development, with syntax that at times is a bit complex. Still, it offers a number of advantages:

Interactive and compile mode: F# has an interactive and a compile mode of operation. In interactive mode you can write code that is executed immediately in a manner similar to working with a calculator, while in compile mode you combine many lines of code possibly in many files into a single application, which is easier to distribute to people who are not F# experts and is faster to execute.

Indentation for scope: F# uses indentation to indicate scope. Some lines of code belong together and should be executed in a certain order and may share data. Indentation helps in specifying this relationship.

Strongly typed: F# is strongly typed, reducing the number of runtime errors. That is, F# is picky, and will not allow the programmer to mix up types such as numbers and text. This is a great advantage for large programs.

Multi-platform: F# is available on Linux, Mac OS X, Android, iOS, Windows, GPUs, and browsers via the Mono platform.

Free to use and open source: F# is supported by the Fsharp foundation (`http://fsharp.org`) and sponsored by Microsoft.

Assemblies: F# is designed to be able to easily communicate with other .Net and Mono programs through the language-independent, platform-independent bytecode called Common Intermediate Language (CIL) organized as assemblies. Thus, if you find that certain parts of a program are easy to express in F# and others in C++, then you will be able to combine these parts later into a single program.

Modern computing: F# supports all aspects of modern computing including Graphical User Interfaces, Web programming, Information rich programming, Parallel algorithms, . . .

Integrated development environments (IDE): F# is supported by major IDEs such as Visual Studio (`https://www.visualstudio.com`) and Xamarin Studio (`https://www.xamarin.com`).

2.5. How to Read This Book

Learning to program requires mastering a programming language, however, most programming languages contain details that are rarely used or used in contexts far from a specific programming topic. Hence, this book only includes a subset of F# but focuses on language structures necessary to understand several common programming paradigms: Imperative programming mainly covered in Chapters 6 to 11, functional programming mainly covered

in Chapters 13 to 16, object-oriented programming in Chapters 20 and 22, and event-driven programming in Chapter 23. A number of general topics are given in the appendix for reference. For further reading please consult `http://fsharp.org`.

3 | Executing F# Code

3.1. Source Code

F# is a functional first programming language, meaning that it has strong support for functional programming, but also supports imperative and object-oriented programming.

F# has two modes of execution, *interactive* and *compiled*. Interactive mode allows the user to interact with F# as a dialogue: The user writes statements, and F# responds immediately. Interactive mode is well suited for small experiments or back-of-an-envelope calculations, but not for programming in general. In compile mode, the user writes a complete program, which is translated or compiled using the F# compiler into a compiled file. The compiled file can be run or executed as a stand-alone program using a virtual machine called `mono`. In both the interactive and compile mode, F# statements are translated into something that can be executed on the computer. A major difference is that in interactive mode, the translation is performed everytime the program is executed, while in compiled mode the translation is performed only once.

· interactive mode
· compile mode

Both interactive and compile modes can be accessed via the *console*, see Appendix A for more information on the console. The interactive system is started by calling `fsharpi` at the command prompt in the console, while compilation is performed with `fsharpc`, and execution of the compiled code is performed using the `mono` command.

· console

F# programs come in many forms, which are identified by suffixes. The *source code* is an F# program written in human-readable form using an editor. F# recognizes the following types of source code files:

· source code

`.fs` An *implementation file*, e.g., `myModule.fs`

· implementation file

`.fsi` A *signature file*, e.g., `myModule.fsi`

· signature file

`.fsx` A *script file*, e.g., `gettingStartedStump.fsx`

· script file

`.fsscript` Same as `.fsx`, e.g., `gettingStartedStump.fsscript`

Compiled code is source code translated into a machine-readable language, which can be executed by a machine. Compiled F# code is either:

`.dll` A *library file*, e.g., `myModule.dll`

· library file

`.exe` A stand-alone *executable file*, e.g., `gettingStartedStump.exe`

· executable file

The implementation, signature, and script files are all typically compiled to produce an executable file, in which case they are called *scripts*, but can also be entered into the

· scripts

interactive system, in which case these are called *script-fragments*. The implementation · script-fragment
and signature files are special kinds of script files used for building libraries. Libraries
in F# are called modules, and they are collections of smaller programs used by other
programs, which will be discussed in detail in Chapter 9.

3.2. Executing Programs

Programs may either be executed by the interpreter or by compiling and executing the
compiled code. In Mono the interpreter is called `fsharpi` and can be used in two ways:
interactively, where a user enters one or more script-fragments separated by the "`;;`" char-
acters, or to execute a script file treated as a single script-fragment. [1]

To illustrate the difference between interactive and compile mode, consider the program
in Listing 3.1.

Listing 3.1 gettingStartedStump.fsx:
A simple demonstration script.

```
1  let a = 3.0
2  do printfn "%g" a
```

The code declares a value `a` to be the decimal value 3.0 and finally prints it to the console.
The `do printfn` is a statement for displaying the content of a value to the screen, and
`"%g"` is a special notation to control how the value is printed. In this case, it is printed
as a decimal number. This and more will be discussed at length in the following chapters.
For now, we will concentrate on how to interact with the F# interpreter and compiler.

An interactive session is obtained by starting the console, typing the `fsharpi` command,
typing the lines of the program, and ending the script-fragment with "`;;`". The dialogue
in Listing 3.2 demonstrates the workflow. What the user types has been highlighted by a
box.

---

[1] Jon: **Too early to introduce lexeme: "F# uses many characters which at times are given
special meanings, e.g., the characters ";;" is compound character denoting the end of a
script-fragment. Such possibly compound characters are called lexemes."**

Listing 3.2: An interactive session.

```
1  $ fsharpi
2
3  F# Interactive for F# 4.1 (Open Source Edition)
4  Freely distributed under the Apache 2.0 Open Source License
5
6  For help type #help;;
7
8  > let a = 3.0
9  - do printfn "%g" a;;
10 3
11
12 val a : float = 3.0
13 val it : unit = ()
14
15 > #quit;;
```

We see that after typing `fsharpi`, the program starts by stating details about itself. Then F# writes > indicating that it is ready to receive commands. The user types `let a = 3.0` and presses `enter`, to which the interpreter responds with `-`. This indicates that the line has been received, that the script-fragment is not yet completed, and that it is ready to receive more input. When the user types `do printfn "%g" a;;` followed by `enter`, then by ";;" the interpreter knows that the script-fragment is completed, it interprets the script-fragment, responds with `3` and some extra information about the entered code, and with > to indicate that it is ready for more script-fragments. The interpreter is stopped when the user types `#quit;;`. It is also possible to stop the interpreter by typing `ctrl-d`.

The interactive session results in extra output on the *type inference* performed. In Listing 3.2 F# states that the name `a` has *type* `float` and the value `3.0`. Likewise, the `do` statement F# refers to by the name `it`, and it has the type `unit` and value `()`. Types are very important to F# since they define how different program pieces fit together like lego bricks. They are a key ingrediens for finding errors in programs, also known as *debugging*, and much of the rest of this book is concerned with types.

· type inference

· type

· debugging

Instead of running `fsharpi` interactively, we can write the script-fragment from Listing 3.1 into a file, here called `gettingStartedStump.fsx`. This file can be interpreted directly by `fsharpi` as shown in Listing 3.3.

Listing 3.3: Using the interpreter to execute a script.

```
1  $ fsharpi gettingStartedStump.fsx
2  3
```

Notice that in the file, ";;" is optional. We see that the interpreter executes the code and prints the result on screen without the extra type information as compared to Listing 3.2.

Finally, the file containing Listing 3.1 may be compiled into an executable file with the program `fsharpc`, and run using the program `mono` from the console. This is demonstrated in Listing 3.4.

Listing 3.4: Compiling and executing a script.

```
1  $ fsharpc gettingStartedStump.fsx
2  F# Compiler for F# 4.1 (Open Source Edition)
3  Freely distributed under the Apache 2.0 Open Source License
4  $ mono gettingStartedStump.exe
5  3
```

The compiler takes `gettingStartedStump.fsx` and produces `gettingStarted.exe`, which can be run using `mono`.

Both the interpreter and the compiler translates the source code into a format which can be executed by the computer. While the compiler performs this translation once and stores the result in the executable file, the interpreter translates the code every time the code is executed. Thus, to run the program again with the interpreter, it must be retranslated as `"$fsharpi gettingStartedStump.fsx`". In contrast, compiled code does not need to be recompiled to be run again, only re-executed using `"$ mono gettingStartedStump.exe"`. On a MacBook Pro, with a 2.9 GHz Intel Core i5, the time the various stages take for this script are:

Command	Time
`fsharpi gettingStartedStump.fsx`	1.88s
`fsharpc gettingStartedStump.fsx`	1.90s
`mono gettingStartedStump.exe`	0.05s

I.e., executing the script with `fsharpi` is slightly faster than by first compiling it with `fsharpc` and then executing the result with `mono` (1.88s < 0.05s + 1.90s), if the script were to be executed only once, but every future execution of the script using the compiled version requires only the use of `mono`, which is much faster than `fsharpi` (1.88s ≫ 0.05s).

Executing programs with the interpreted directly from a file and compiling and executing the program is much preferred for programming complete programs, since the starting state is well defined, and since this better supports *unit testing*, which is a method for debugging programs. Thus, **prefer compiling over interpretation.**

· unit testing

Advice

4 | Quick-start Guide

Programming is the art of solving problems by writing a program to be executed by a computer. For example, to solve the following problem,

> **Problem 4.1**
>
> What is the sum of 357 and 864?

we have written the programshown in Listing 4.1.

> **Listing 4.1 quickStartSum.fsx:**
> **A script to add 2 numbers and print the result to the console.**
>
> ```fsharp
> 1  let a = 357
> 2  let b = 864
> 3  let c = a + b
> 4  do printfn "%A" c
> ```
>
> ```
> 1  $ fsharpc --nologo quickStartSum.fsx && mono quickStartSum.exe
> 2  1221
> ```

In the box the above, we see our program was saved as a script in a file called `quickStartSum.fsx`, and in the console we executed the program by typing the command `fsharpc --nologo quickStartSum.fsx && mono quickStartSum.exe`. The result is then printed in the console to be `1221`. Here, as in the rest of this book, we have used the optional flag `--nologo`, which informs `fsharpc` not to print information about its version etc., thus making the output shorter. The `&&` notation tells the console to first run the command on the left, and if that did not report any errors, then run the command on the right. This could also have been performed as two separate commands to the console, and throughout this book we will use the above shorthand when convenient.

To solve the problem, we made program consisting of several lines, where each line was an *expressions*. The first expression, `let a = 357`, in line 1 used the *let keyword* to *bind* the value 357 to the name `a`. This is called a *let binding* and makes the name synonymous with the value. Another notable point is that F# identifies 357 as an *integer number*, which is F#'s preferred number type, since computations on integers are very efficient, and since integers are very easy to communicate to other programs. In line 2 we bound the value 864 to the name `b`, and to the name `c` we bound the result of evaluating the sum `a + b` in line 3. Line 4 is a *do binding*, as noted by the keyword *do*. Do bindings are also sometimes called *statements*, and the `do` keyword is optional in F#. Here the value of `c` was printed to the console followed by a newline with the *printfn function*. A function

· expression
· `let`
· keyword
· binding
· let binding
· integer number
· do binding
· `do`
· statement
· `printfn`
· function

15

in F# is an entity that takes zero or more arguments and returns a value. The function `printfn` is very special, since it can take any number of arguments and returns the special value *"()"* which has type ***unit***. The `do` tells F# to ignore this value. Here `printfn` has been used with 2 arguments: `"%A"` and `c`. Notice that in contrast to many other languages, F# does not use parentheses to frame the list of arguments, nor does it use commas to separate them. In general, the `printfn` function always has 1 or more arguments, and the first is a *format string*. A *string* is a sequence of characters starting and ending with double quotation marks. E.g., `let s = "this is a string of characters"` binds the string `"this is..."` to the name `s`. For the `printfn` function, the format string may be any string, but if it contains format character sequences, such as `%A`, then format character sequence are replaced by the arguments to `printfn` which follows the format string. The format string must match the value *type*, that is, here `c` is of type integer, whereas the format string `%A` matches many types.

· ()
· unit

· format string
· string

· type

Types are a central concept in F#. In the script 4.1 we bound values of integer type to names. There are several different integer types in F#, here we used the one called `int`. The values were not *declared* to have these types, instead the types were *inferred* by F#. Typing these bindings line by line in an interactive session, we see the inferred types as shown in Listing 4.2.

· type declaration
· type inference

> **Listing 4.2: Inferred types are given as part of the response from the interpreter.**

```
1   > let a = 357;;
2   val a : int = 357
3
4   > let b = 864;;
5   val b : int = 864
6
7   > let c = a + b;;
8   val c : int = 1221
9
10  > do printfn "%A" c;;
11  1221
12  val it : unit = ()
```

The interactive session displays the type using the ***val*** keyword followed by the name used in the binding, its type, and its value. Since the value is also responded, the last `printfn` statement is superfluous. However, **it is ill-advised to design programs to be run in an interactive session, since the scripts need to be manually copied every time it is to be run, and since the starting state may be unclear.** Notice that `printfn` is automatically bound to the name ***it*** of type `unit` and value *"()"*. F# insists on binding all statements to values, and in lack of an explicit name, it will use `it`. Rumor has it that `it` is an abbreviation for "irrelevant".

· val

Advice

· it
· ()

The following problem,

> **Problem 4.2**
>
> What is the sum of 357.6 and 863.4?

uses *decimal point* numbers instead of integers. These numbers are called *floating point numbers*, and their internal representation is quite different to integer numbers used previously. Likewise, algorithms used to perform arithmetic are quite different from integers.

· decimal point
· floating point number

Now the program would look like Listing 4.3.

Listing 4.3 quickStartSumFloat.fsx:
Floating point types and arithmetic.

```
1  let a = 357.6
2  let b = 863.4
3  let c = a + b
4  do printfn "%A" c
```

```
1  $ fsharpc --nologo quickStartSumFloat.fsx && mono
     quickStartSumFloat.exe
2  1221.0
```

On the surface, this could appear as an almost negligible change, but the set of integers and the set of real numbers (floats) require quite different representations in order to be effective on a computer, and as a consequence, the implementation of their operations, such as addition, are very different. Thus, although the response is an integer, it has type `float` which is indicated by `1221.0` and which is not the same as `1221`. F# is very picky about types, and generally does not allow types to be mixed, as demonstrated in the interactive session in Listing 4.4.

Listing 4.4: Mixing types is often not allowed.

```
1   > let a = 357;;
2   val a : int = 357
3
4   > let b = 863.4;;
5   val b : float = 863.4
6
7   > let c = a + b;;
8     let c = a + b;;
9     ------------^
10
11  stdin(4,13): error FS0001: The type 'float' does not match the
      type 'int'
```

We see that binding a name to a number without a decimal point is inferred to be an integer, while when binding to a number with a decimal point the type is inferred to be a float, and that our attempt of adding an integer and floating point value gives an error. The *error message* contains much information. First, it states that the error is in `stdin(4,13)`, · error message which means that the error was found on standard-input at line 4 and column 13. Since the program was executed using `fsharpi quickStartSumFloat.fsx`, here standard input means the file `quickStartSumFloat.fsx` shown in Listing 4.3. The corresponding line and column are also shown in Listing 4.4. After the file, line, and column number, F# informs us of the error number and a description of the error. Error numbers are an underdeveloped feature in Mono and should be ignored. However, the verbal description often contains useful information for *debugging*. In the example we are informed that there · debugging is a type mismatch in the expression, i.e., since `a` is an integer, F# expected `b` to be one too. Debugging is the process of solving errors in programs, and here we can solve the error by either making `a` into a float or `b` into an int. The right solution depends on the application.

F# is a functional first programming language, and one implication of this is that names have a *lexical scope*. A scope is the lines in a program where a binding is valid, and lexical scope means that to find the value of a name, F# looks for the value in the above lines. Furthermore, at the outermost level, rebinding is not allowed. If attempted, then F# will return an error as shown in Listing 4.5.

· lexical scope

**Listing 4.5 quickStartRebindError.fsx:
A name cannot be rebound.**

```
1  let a = 357
2  let a = 864
```

```
1  $ fsharpc --nologo -a quickStartRebindError.fsx
2
3  quickStartRebindError.fsx(2,5): error FS0037: Duplicate
     definition of value 'a'
```

However, if the same code is executed in an interactive session, then rebinding does not cause an error, as shown in Listing 4.6.

Listing 4.6: Names may be reused when separated by the lexeme ";;".

```
1  > let a = 357;;
2  val a : int = 357
3
4  > let a = 864;;
5  val a : int = 864
```

The difference is that the *";;" lexeme* is used to specify the end of a *script-fragment*. A lexeme is a letter or word, which F# considers as an atomic unit. Script-fragments may be defined both in scripts and in interactive mode, and rebinding is not allowed at the outermost level in script-fragments. Even with the ";;" lexeme, rebinding is not allowed in compile-mode. In general, **avoid rebinding of names.**

· ;;
· lexeme
· script-fragment

Advice

In F#, *functions* are also values, and we may define a function `sum` as part of the solution to the above program, as shown in Listing 4.7.

· function

**Listing 4.7 quickStartSumFct.fsx:
A script to add 2 numbers using a user-defined function.**

```
1  let sum x y = x + y
2  let c = sum 357 864
3  do printfn "%A" c
```

```
1  $ fsharpc --nologo quickStartSumFct.fsx && mono
     quickStartSumFct.exe
2  1221
```

Functions are useful to *encapsulate* code, such that we can focus on the transformation of data by a function while ignoring the details on how this is done. Functions are also

· encapsulation

useful for code reuse, i.e., instead of repeating a piece of code in several places, such code can be encapsulated in a function and replaced with function calls. This makes debugging and maintenance considerably simpler. Entering the function into an interactive session will illustrate the inferred type the function `sum` has: `val sum : x:int -> y:int -> int`. The "`->`" is the mapping operator in the sense that functions are mappings between sets. The type of the function `sum`, should be read as `val sum : x:int -> (y:int -> int)`, that is, `sum` takes an integer and returns a function, which takes an integer and returns an integer. This is an example of a higher-order function.

Type inference in F# may cause problems, since the type of a function is inferred based on the context in which it is defined. E.g., in an interactive session, defining the sum in one scope on a single line will default the types to integers, F#'s favorite type. Thus, if the next script-fragment uses the function with floats, then we will get an error message as shown in Listing 4.8.

Listing 4.8: Types are inferred in blocks, and F# tends to prefer integers.

```
1   val sum : x:int -> y:int -> int
2
3   > let c = sum 357.6 863.4;;
4     let c = sum 357.6 863.4;;
5     ------------^^^^^
6
7   stdin(3,13): error FS0001: This expression was expected to
      have type
8       'int'
9   but here has type
10      'float'
```

A remedy is to define the function in the same script-fragment as it is used, such as shown in Listing 4.9.

Listing 4.9: Type inference is per script-fragment.

```
1   > let sum x y = x + y
2   - let c = sum 357.6 863.4;;
3   val sum : x:float -> y:float -> float
4   val c : float = 1221.0
```

Alternatively, the types may be explicitly stated as shown in Listing 4.10.

Listing 4.10: Function argument and return types may be stated explicitly.

```
1   > let sum (x : float) (y : float) : float = x + y;;
2   val sum : x:float -> y:float -> float
3
4   > let c = sum 357.6 863.4;;
5   val c : float = 1221.0
```

The function `sum` has two arguments and a return type, and in Listing 4.10 we have specified all three. This is done using the "`:`" lexeme, and to resolve confusion, we must use parentheses around the arguments, such as `(y : float)`, otherwise F# would not be

able to understand whether the type annotation was for the argument or the return value. Often it is sufficient to specify just some of the types, since type inference will enforce the remaining types. E.g., in this example, the "+" operator is defined for identical types, so specifying the return value of `sum` to be a float implies that the result of the "+" operator is a float, and therefore that its arguments must be floats, and finally then that the arguments for `sum` must be floats. However, in this book we advocate the following advice: **specify**   Advice **types unless explicitly working with generic functions.**

In this chapter, we have scratched the surface of learning how to program by concentrating on a number of key programming concepts and how they are expressed in the F# language. In the following chapters, we will expand the description of F# with features used in all programming approaches.

5 | Using F# as a Calculator

In this chapter, we will exclusively use the interactive mode to illustrate basic types and operations in F#.

5.1. Literals and Basic Types

All programs rely on processing of data, and an essential property of data is its *type*. A *literal* is a fixed value like the number 3, and if we type the number 3 in an interactive session at the input prompt, then F# responds as shown in Listing 5.1.

· type
· literal

> Listing 5.1: Typing the number 3.

```
1  > 3;;
2  val it : int = 3
```

What this means is that F# has inferred the type to be *int* and bound it to the identifier *it*. For more on binding and identifiers see Chapter 6. Types matter, since the operations that can be performed on integers, are quite different from those that can be performed on, e.g., strings. Therefore, the number 3 has many different representations as shown in Listing 5.2.

· int
· it

> Listing 5.2: Many representations of the number 3 but using different types.

```
1   > 3;;
2   val it : int = 3
3
4   > 3.0;;
5   val it : float = 3.0
6
7   > '3';;
8   val it : char = '3'
9
10  > "3";;
11  val it : string = "3"
```

Each literal represents the number 3, but their types are different, and hence they are quite different values. The types int for integer numbers, *float* for floating point numbers, *bool* for Boolean values, *char* for characters, and *string* for strings of characters are the most

· float
· bool
· char
· string

Metatype	Type name	Description
Boolean	bool	Boolean values true or false
Integer	int	Integer values from -2,147,483,648 to 2,147,483,647
	byte	Integer values from 0 to 255
	sbyte	Integer values from -128 to 127
	int8	Synonymous with sbyte
	uint8	Synonymous with byte
	int16	Integer values from -32768 to 32767
	uint16	Integer values from 0 to 65535
	int32	Synonymous with int
	uint32	Integer values from 0 to 4,294,967,295
	int64	Integer values from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
	uint64	Integer values from 0 to 18,446,744,073,709,551,615
Real	float	64-bit IEEE 754 floating point value from $-\infty$ to ∞
	double	Synonymous with float
	single	A 32-bit floating point type
	float32	Synonymous with single
	decimal	A floating point data type that has at least 28 significant digits
Character	char	Unicode character
	string	Unicode sequence of characters
None	unit	The value ()
Object	obj	An object
Exception	exn	An exception

Table 5.1.: List of some of the basic types. The most commonly used types are underlined. For a description of integer see Appendix B.1, for floating point numbers see Appendix B.2, for ASCII and Unicode characters see Appendix C, for objects see Chapter 20, and for exceptions see Chapter 18.

common types of literals. A table of all *basic types* predefined in F# is given in Table 5.1. In addition to these built-in types, F# is designed such that it is easy to define new types. · basic types

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base* 10, which means that a value is represented as two sequences of decimal digits separated by a *decimal point*, where each *digit d* has a position and a value $d \in \{0, 1, 2, \ldots, 9\}$. The part before the decimal point is called the *whole part* and the part after is called the *fractional part* of the number. An *integer* is a number with only a whole part and neither a decimal point nor a fractional part. As an example `35.7` is a decimal number, whose value is $3 \cdot 10^1 + 5 \cdot 10^0 + 7 \cdot 10^{-1}$, and `128` is an integer, whose value is $1 \cdot 10^2 + 2 \cdot 10^1 + 8 \cdot 10^0$. In F#, a decimal number is called a *floating point number*. Floating point numbers may alternatively be given using *scientific notation*, such as `3.5e-4` and `4e2`, where the `e`-notation is translated to a value as `3.5e-4` $= 3.5 \cdot 10^{-4} = 0.00035$, and `4e2` $= 4 \cdot 10^2 = 400$.

· decimal number
· base
· decimal point
· digit
· whole part
· fractional part
· integer
· floating point number
· scientific notation

The basic unit of information in almost all computers is the binary digit or *bit* for short. Internally, programs and data are all represented as bits, hence F# has a strong support for binary numbers. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. E.g., the binary number $101.01_2 = 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 5.25$. Subscripts are often used to indicate the base of a number, e.g., 101.01_2 and 101.01_{10} are different numbers. Since base 10 is so common, the subscript for base 10 numbers is often omitted.

· bit

· binary number

Character	Escape sequence	Description
BS	\b	Backspace
LF	\n	Line feed
CR	\r	Carriage return
HT	\t	Horizontal tabulation
\	\\	Backslash
"	\"	Quotation mark
'	\'	Apostrophe
BEL	\a	Bell
FF	\f	Form feed
VT	\v	Vertical tabulation
	\uXXXX, \UXXXXXXXX, \DDD	Unicode character ('X' is any hexadecimal digit, and 'D' is any decimal digit)

Table 5.2.: Escape characters. The escapecode `\DDD` is sometimes called a tricode.

Binary numbers are closely related to *octal* and *hexadecimal numbers*. Octals uses 8 as basis and hexadecimals use 16 as basis. Each octal digit can be represented by exactly three bits, and each hexadecimal digit can represented by exactly four bits. The hexadecimal digits use `0–9` to represent the values 0–9 and `a–f` in lower or alternatively upper case to represent the values 10-15. Thus, Octals and hexadecimals conveniently serve as shorthand for the much longer binary representation. As examples, the octal number 37_8 is $3\cdot 8^1 + 7\cdot 8^0 = 31$, and the hexadecimal number $\mathtt{f3}_{16}$ is $15 \cdot 16^1 + 3 \cdot 16^0 = 243$.

· octal number
· hexadecimal number

To denote integers in bases different than 10, F# uses the prefix '0b' for binary, '0o' for octal, and '0x' for hexadecimal numbers. For example, the value 367_{10} may be written as an integer `367`, as a binary number `0b101101111`, as a octal number `0o557`, and as a hexadecimal number `0x16f`. In F#, the character sequences `0b12` and `ff` are not recognized as numbers.

A *character* is a *Unicode code point*, and character literals are enclosed in single quotation marks. Appendix C.3 contains more details on code points. The character type in F# is denoted **char**. Examples of characters are 'a', 'D', '3', and examples of non-characters are '23' and 'abc'. Some characters, such as the tabulation character, do not have a visual representation. These can still be represented as a character using *escape sequences*. A character escape sequence starts with "\" followed by letter for simple escapes such as `\t` for tabulation and `\n` for newline. Escape sequences can also be a numerical representation of a code point, and three versions exist: The trigraph `\DDD`, where `D` is a decimal digit, is used to specify the first 256 code points, the hexadecimal escape codes `\uXXXX`, where `X` is a hexadecimal digit, is used to specify the first 65536 code points, and `\UXXXXXXXX` is used to specify any of the approximately $4.3 \cdot 10^9$ possible code points. All escape sequences are shown in Table 5.2. Examples of **char** representations of the letter 'a' are: `'a'`, `'\097'`, `'\u0061'`, `'\U00000061'`.

· character
· Unicode
· code point
· **char**
· escape sequences

A *string* is a sequence of characters enclosed in double quotation marks. Examples are `"a"`, `"this is a string"`, and `"-&#\@"`. Note that the string `"a"` and the character `'a'` are not the same. Some strings are so common that they are given special names: One or more spaces `" "` is called *whitespace*, and both `"\n"` and `"\r\n"` are called *newline*. The escape-character "\" may be used to break a line in two. This and other examples are shown in Listing 5.3.

· string

· whitespace
· newline

Type	syntax	Examples	Value
int, int32	`<int or hex>`	3, 0x3	3
	`<int or hex>l`	3l, 0x3l	
uint32	`<int or hex>u`	3u	3
	`<int or hex>ul`	3ul	
byte, uint8	`<int or hex>uy`	97uy	97
	`'<char>'B`	`'a'B`	
byte[]	`"<string>"B`	`"a\n"B`	[|97uy; 10uy|]
	`@"<string>"B`	`@"a\n"B`	[|97uy; 92uy; 110uy|]
sbyte, int8	`<int or hex>y`	3y	3
int16	`<int or hex>s`	3s	3
uint16	`<int or hex>us`	3us	3
int64	`<int or hex>L`	3L	3
uint64	`<int or hex>UL`	3UL	3
	`<int or hex>uL`	3uL	
float, double	`<float>`	3.0	3.0
	`<hex>LF`	0x013fLF	9.387247271e-323
single, float32	`<float>F`	3.0F	3.0
	`<float>f`	3.0f	3.0
	`<hex>lf`	0x013flf	4.4701421e-43f
decimal	`<float or int>M`	3.0M,3M	3.0
	`<float or int>m`	3.0m,3m	
string	`"<string>"`	`"a \"quote\".\n"`	a "quote".<newline>
	`@"<string>"`	`@"a ""quote"".\n"`	a "quote".\n.
	`"""<string>"""`	`"""a "quote".\n"""`	a "quote".\n

Table 5.3.: List of literal types. The syntax notation `<>` means that the programmer replaces the brackets and content with a value of the appropriate form. The `[||]` notation means that the value is an array, see Section 11.3 for details.

Listing 5.3: Examples of string literals.

```
1  > "abcde";;
2  val it : string = "abcde"
3
4  > "abc
5  -    de";;
6  val it : string = "abc
7     de"
8
9  > "abc\
10 -    de";;
11 val it : string = "abcde"
12
13 > "abc\nde";;
14 val it : string = "abc
15 de"
```

Note that the response from `fsharpi` is shown in double quotation marks, but this is not part of the string.

F# supports *literal types*, where the type of a literal is indicated as a prefix or suffix as   · literal type shown in Table 5.3.

The literal type is closely connected to how the values are represented internally. For example, a value of type `int32` use 32 bits and can be both positive and negative, while a `uint32` value also use 32 bits, but is unsigned. A `byte` is an 8-bit number, and `sbyte` is a signed 8-bit number. Values of type `float` use 64 bits, while `float32` only uses 32 bits. The number of bits used to represent numbers directly relates to the range and precession these types can represent. This is summarized in Table 5.1 and discussed in more detail in Appendix B. String literals may be *verbatim* by the @-notation or triple double quotation · verbatim marks, meaning that the escape sequences are not converted to their code point. The two types of string verbatim treat quotation marks differently, as illustrated in the table. Further examples are shown in Listing 5.4.

Listing 5.4: Named and implied literals.

```
1  > 3;;
2  val it : int = 3
3
4  > 4u;;
5  val it : uint32 = 4u
6
7  > 5.6;;
8  val it : float = 5.6
9
10 > 7.9f;;
11 val it : float32 = 7.9000001f
12
13 > 'A';;
14 val it : char = 'A'
15
16 > 'B'B;;
17 val it : byte = 66uy
18
19 > "ABC";;
20 val it : string = "ABC"
21
22 > @"abc\nde";;
23 val it : string = "abc\nde"
```

Many basic types are compatible, and the type of a literal may be changed by *typecasting*. · typecasting An example of casting to a `float` is shown in Listing 5.5.

Listing 5.5: Casting an integer to a floating point number.

```
1  > float 3;;
2  val it : float = 3.0
```

When `float` is given an argument, then it acts as a function rather than a type, and for the integer `3` it returns the floating point number `3.0`. For more on functions see Chapter 6. Boolean values are often treated as the integer values 0 and 1, but no short-hand function names exist for their conversions. Instead, use functions from the `System.Convert` family of functions, as demonstrated in Listing 5.6.

Listing 5.6: Casting booleans.

```
1  > System.Convert.ToBoolean 1;;
2  val it : bool = true
3
4  > System.Convert.ToBoolean 0;;
5  val it : bool = false
6
7  > System.Convert.ToInt32 true;;
8  val it : int = 1
9
10 > System.Convert.ToInt32 false;;
11 val it : int = 0
```

Here `System.Convert.ToBoolean` is the identifier of a function `ToBoolean`, which is a *member* of the *class* `Convert` that is included in the *namespace* `System`. Namespaces, classes, and members will be discussed in Chapter 9.

· member
· class
· namespace

Typecasting is often a destructive operation, e.g., typecasting a `float` to `int` removes the fractional part without rounding as shown in Listing 5.7.

Listing 5.7: Fractional part is removed by downcasting.

```
1  > int 357.6;;
2  val it : int = 357
```

Here we typecasted to a lesser type, in the sense that the set of integers is a subset of floating point numbers, and this is called *downcasting*. The opposite is called *upcasting* and is often non-destructive, as Listing 5.5 showed. Since floating point numbers are a superset of integers, the value is retained. As a side note, *rounding* a number $y.x$, where y is the *whole part* and x is the *fractional part*, is the operation of mapping numbers in the interval $y.x \in [y.0, y.5)$ to y, and those in $y.x \in [y.5, y+1)$ to $y+1$. This can be performed by downcasting, as shown in Listing 5.8.

· downcasting
· upcasting
· rounding
· whole part
· fractional part

Listing 5.8: Rounding by modified downcasting.

```
1  > int (357.6 + 0.5);;
2  val it : int = 358
```

I.e., $357.6 + 0.5 = 358.1$ and removing the fractional part by downcasting results in 358, which is the correct answer.

5.2. Operators on Basic Types

Expressions are the basic building block of all F# programs, and this section will discuss operator expressions on basic types. A typical calculation, such used in Listing 5.8, is

$$\underbrace{357.6}_{\text{operand}} \quad \underbrace{+}_{\text{operator}} \quad \underbrace{0.5}_{\text{operand}} \tag{5.1}$$

is an example of an arithmetic *expression*, and the above expression consists of two *operands*

and an *operator*. Since this operator takes two operands, it is called a *binary operator*. The expression is written using *infix notation*, since the operands appear on each side of the operator.

· operator

· binary operator

· infix notation

In order to discuss general programming structures, we will use a simplified language to describe valid syntactical structures. In this simplified language, the syntax of basic binary operators is shown in the following.

Listing 5.9: **Syntax for a binary expression.**

```
<expr><op><expr>
```

Here `<expr>` is any expression supplied by the programmer, and `<op>` is a binary infix operator. F# supports a range of arithmetic binary infix operators on its built-in types, such as addition, subtraction, multiplication, division, and exponentiation, using the "+", "-", "*", "/", "**" lexemes, respectively. Not all operators are defined for all types, e.g., addition is defined for integer and float types as well as for characters and strings, but multiplication is only defined for integer and floating-point types. A complete list of built-in operators on basic types is shown in Table E.1 and E.2, and a range of mathematical functions is shown in Table E.3. Note that expressions can themselves be arguments to expressions, and thus, `4+5+6` is also a legal statement. This is called *recursion*, which means that a rule or a function is used by the rule or function itself in its definition. See Chapter 13 for more on recursive functions.

· recursion

Unary operators take only one argument and have the syntax:

Listing 5.10: **A unary expressions.**

```
<op><expr>
```

An example of a unary operator is `-3`, where `-` here is used to negate a positive integer. Since the operator appears before the operand, it is a *prefix operator*.

· prefix operator

The concept of *precedence* is an important concept in arithmetic expressions.[1] If parentheses are omitted in Listing 5.8, then F# will interpret the expression as `(int 357.6) + 0.5`, which is erroneous since the addition of an integer with a float is undefined. This is an example of precedence, i.e., function evaluation takes precedence over addition which means that function evaluation is performed first and addition second. Consider the arithmetic expression shown in Listing 5.11.

· precedence

Listing 5.11: A simple arithmetic expression.

```
> 3 + 4 * 5;;
val it : int = 23
```

Here, the addition and multiplication functions are shown in infix notation with the *operator* lexemes "+" and "*". To arrive at the resulting value 23, F# has to decide in which order to perform the calculation. There are 2 possible orders, `3 + (4 * 5)` and `(3 + 4) * 5` that gives different results. For integer arithmetic, the correct order is, of course,

· operator

---

[1]Jon: **minor comment on indexing and slice-ranges.**

Operator	Associativity	Description
`+<expr>`, `-<expr>`, `~~~<expr>`	Left	Unary identity, negation, and bitwise negation operators
`f <expr>`	Left	Function application
`<expr> ** <expr>`	Right	Exponentiation
`<expr> * <expr>`, `<expr> / <expr>`, `<expr> % <expr>`	Left	Multiplication, division and remainder
`<expr> + <expr>`, `<expr> - <expr>`	Left	Addition and subtraction binary operators
`<expr> ^^^ <expr>`	Right	Bitwise exclusive or
`<expr> < <expr>`, `<expr> <= <expr>`, `<expr> > <expr>`, `<expr> >= <expr>`, `<expr> = <expr>`, `<expr> <> <expr>`, `<expr> <<< <expr>`, `<expr> >>> <expr>`, `<expr> &&& <expr>`, `<expr> ||| <expr>`,	Left	Comparison operators, bitwise shift, and bitwise 'and' and 'or'.
`<expr> && <expr>`	Left	Boolean and
`<expr> || <expr>`	Left	Boolean or

Table 5.4.: Some common operators, their precedence, and their associativity. Rows are ordered from highest to lowest precedences, such that `<*expr*> * <*expr*>` has higher precedence than `<*expr*> + <*expr*>`. Operators in the same row have the same precedence. Full table is given in Table E.5.

multiplication before addition, and we say that multiplication takes *precedence* over addition. Every atomic operation that F# can perform is ordered in terms of its precedence, and for some common built-in operators shown in Table 5.4, the precedence is shown by the order they are given in the table.

· precedence

Associativity describes the order in which calculations are performed for binary operators of the same precedence. Some operator's associativity are given in Table 5.4. In the table we see that "`*`" is left associative, which means that `3.0 * 4.0 * 5.0` is evaluated as `(3.0 * 4.0) * 5.0`. Conversely, `**` is right associative, so `4.0**3.0**2.0` is evaluated as `4.0**(3.0**2.0)`. For some operators, like multiplication, association matters little, e.g., $4*3*2 = 4*(3*2) = (4*3)*2$, and for other operators, like exponentiation, association makes a huge difference, e.g., $4^{(3^2)} \neq (4^3)^2$. Examples of this is shown in Listing 5.12.

a	b	a && b	a || b	not a
false	false	false	false	true
false	true	false	true	true
true	false	false	true	false
true	true	true	true	false

Table 5.5.: Truth table for boolean 'and', 'or', and 'not' operators. Value 0 is false and 1 is true.

Listing 5.12: Precedence rules define implicit parentheses.

```
1   > 4.0 * 3.0 * 2.0;;
2   val it : float = 24.0
3
4   > (4.0 * 3.0) * 2.0;;
5   val it : float = 24.0
6
7   > 4.0 * (3.0 * 2.0);;
8   val it : float = 24.0
9
10  > 4.0 ** 3.0 ** 2.0;;
11  val it : float = 262144.0
12
13  > (4.0 ** 3.0) ** 2.0;;
14  val it : float = 4096.0
15
16  > 4.0 ** (3.0 ** 2.0);;
17  val it : float = 262144.0
```

Advice

Whenever in doubt of association or any other basic semantic rules, it is a good idea to use parentheses. It is also a good idea to test your understanding of the syntax and semantic rules by making a simple script.

5.3. Boolean Arithmetic

Boolean arithmetic is the basis of almost all computers and particularly important for controlling program flow, which will be discussed in Chapter 8. Boolean values are one of 2 possible values, true or false, which is also sometimes written as 1 and 0. Basic operations on Boolean values are 'and', 'or', and 'not', which in F# are written respectively as the binary operators &&, ||, and the function not. Since the domain of Boolean values is so small, all possible combination of input on these values can be written on the tabular form, known as a *truth table*, and the truth tables for the basic Boolean operators and functions are shown in Table 5.5. A good mnemonic for remembering the result of the 'and' and 'or' operators is to use 1 for true, 0 for false, multiplication for the Boolean 'and' operator, and addition for the Boolean 'or' operator, e.g., true and false in this mnemonic translates to $1 \cdot 0 = 0$, and the result translates back to the Boolean value false. In F#, the truth table for the basic Boolean operators can be produced by a program, as shown in Listing 5.13.

· and
· or
· not
· truth table

Listing 5.13: Boolean operators and truth tables.

```
1  > printfn "a b a*b a+b not a"
2  - printfn "%A %A %A %A %A"
3  -    false false (false && false) (false || false) (not false)
4  - printfn "%A %A %A %A %A"
5  -    false true (false && true) (false || true) (not false)
6  - printfn "%A %A %A %A %A"
7  -    true false (true && false) (true || false) (not true)
8  - printfn "%A %A %A %A %A"
9  -    true true (true && true) (true || true) (not true);;
10 a b a*b a+b not a
11 false false false false true
12 false true false true true
13 true false false true false
14 true true true true false
15 val it : unit = ()
```

Here, we used the `printfn` function to present the results of many expressions on something that resembles a tabular form. The spacing produced using the `printfn` function is not elegant, and in Section 6.5 we will discuss better options for producing more beautiful output. Notice that the arguments for `printfn` was given on the next line with indentation. The indentation is an important part of telling F# which part of what you write belong together. This is an example of the so-called lightweight syntax. Generally, F# ignores newlines and whitespaces except when using the lightweight syntax. The difference between verbose and lightweight syntax is discussed in Chapter 6.

5.4. Integer Arithmetic

The set of integers is infinitely large, but since all computers have limited resources, it is not possible to represent it in its entirety. The various integer types listed in Table 5.1 are finite subsets reduced by limiting their ranges. An in-depth description of integer implementation can be found in Appendix B. The type `int` is the most common type.

Table E.1–E.3 give examples of operators and functions pre-defined for integer types. Notice that fewer functions are available for integers than for floating point numbers. For most addition, subtraction, multiplication, and negation, the result is straightforward. However, performing arithmetic operations on integers requires extra care, since the result may cause *overflow* and *underflow*. For example, an `sbyte` is specified using the "y"-literal and can · overflow
hold values $[-128 \ldots 127]$. This causes problems in the example in Listing 5.14. · underflow

---

Listing 5.14: Adding integers may cause overflow.

```
1  > 100y;;
2  val it : sbyte = 100y
3
4  > 30y;;
5  val it : sbyte = 30y
6
7  > 100y + 30y;;
8  val it : sbyte = -126y
```

Here $100 + 30 = 130$, which is larger than the biggest `sbyte`, and the result is an overflow. Similarly, we get an underflow, when the arithmetic result falls below the smallest value storable in an `sbyte`, as demonstrated in Listing 5.15.

Listing 5.15: Subtracting integers may cause underflow.

```
1  > -100y - 30y;;
2  val it : sbyte = 126y
```

I.e., we were expecting a negative number but got a positive number instead.

The overflow error in Listing 5.14 can be understood in terms of the binary representation of integers: In binary, $130 = 10000010_2$, and this binary pattern is interpreted differently as `byte` and `sbyte`, see Listing 5.16.

Listing 5.16: The leftmost bit is interpreted differently for signed and unsigned integers, which gives rise to potential overflow errors.

```
1  > 0b10000010uy;;
2  val it : byte = 130uy
3
4  > 0b10000010y;;
5  val it : sbyte = -126y
```

That is, for signed bytes, the left-most bit is used to represent the sign, and since the addition of $100 = 01100100_2$ and $30 = 00011110_b$ is $130 = 10000010_2$, which causes the left-most bit to be used, this is wrongly interpreted as a negative number when stored in an `sbyte`. Similar arguments can be made explaining underflows.

The operator discards the fractional part after division, and the *integer remainder* operator calculates the remainder after integer division, as demonstrated in Listing 5.17.

· integer division

· integer remainder

Listing 5.17: Integer division and remainder operators.

```
1  > 7 / 3;;
2  val it : int = 2
3
4  > 7 % 3;;
5  val it : int = 1
```

Together, the integer division and remainder can form a lossless representation of the original number, see Listing 5.18.

Listing 5.18: Integer division and remainder is a lossless representation of an integer, compare with Listing 5.17.

```
1  > (7 / 3) * 3;;
2  val it : int = 6
3
4  > (7 / 3) * 3 + (7 % 3);;
5  val it : int = 7
```

Here we see that integer division of 7 by 3 followed by multiplication by 3 is less than 7, and that the difference is `7 % 3`.

Notice that neither overflow nor underflow error gave rise to an error message, which is why such bugs are difficult to find. Dividing any non-zero number by 0 is infinite, which is also outside the domain of any of the integer types, but in this case, F# casts an *exception*, · exception as shown in Listing 5.19.

Listing 5.19: Integer division by zero causes an exception runtime error.

```
1  > 3/0;;
2  System.DivideByZeroException: Attempted to divide by zero.
3    at <StartupCode$FSI_0002>.$FSI_0002.main@ () [0x00001] in
       <da48886c466e4b80be4566752c596fa8>:0
4    at (wrapper managed-to-native)
       System.Reflection.MonoMethod:InternalInvoke
       (System.Reflection.MonoMethod,object,object[],System.Exception&)
5    at System.Reflection.MonoMethod.Invoke (System.Object obj,
       System.Reflection.BindingFlags invokeAttr,
       System.Reflection.Binder binder, System.Object[] parameters,
       System.Globalization.CultureInfo culture) [0x00032] in
       <c9f8153c41de4f8cbafd0e32f9bf6b28>:0
6  Stopped due to error
```

The output looks daunting at first sight, but the first and last lines of the error message are the most important parts, which tell us what exception was cast and why the program stopped. The middle contains technical details concerning which part of the program caused the error and can be ignored for the time being. Exceptions are a type of *runtime error*, and are discussed in Chapter 18   · runtime error

Integer exponentiation is not defined as an operator but is available as the built-in function `pown`. This function is demonstrated in Listing 5.20 for calculating 2^5.

Listing 5.20: Integer exponent function.

```
1  > pown 2 5;;
2  val it : int = 32
```

For binary arithmetic on integers, the following operators are available: `<leftExpr>` `<<<` `<rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the left while inserting 0's to right; `<leftExpr>` `>>>` `<rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the right while inserting 0's to left; `~~~` `<expr>` returns a new integer, where all 0 bits are changed to 1 bits and vice-versa; `<expr>` `&&&` `<expr>` returns the result of taking the Boolean 'and' operator position-wise; `<expr>` `|||` `<expr>` returns the result of taking the Boolean 'or' operator position-wise; and `<expr>` `^^^` `<expr>` returns the result of the Boolean 'xor' operator defined by the truth table in Table 5.6.   · xor
                                                                    · exclusive or

a	b	a ^^^ b
false	false	false
false	true	true
true	false	true
true	true	false

Table 5.6.: Boolean exclusive or truth table.

5.5. Floating Point Arithmetic

Like integers, the set of reals is also infinitely large, hence, floating point types are finite subsets reduced by sampling the space of reals. An in-depth description of floating point implementations can be found in Appendix B. The type `float` is the most common type.

Table E.1–E.3 give examples of operators and functions pre-defined for floating point types. Note that the remainder operator for floats calculates the remainder after division and discards the fractional part, see Listing 5.21.

Listing 5.21: Floating point division and remainder operators.

```
1  > 7.0 / 2.5;;
2  val it : float = 2.8
3
4  > 7.0 % 2.5;;
5  val it : float = 2.0
```

The remainder for floating point numbers can be fractional, but division, discarding fractional part, and the remainder is still a lossless representation of the original number, as demonstrated in Listing 5.22.

Listing 5.22: Floating point division, downcasting, and remainder is a lossless representation of a number.

```
1  > float (int (7.0 / 2.5));;
2  val it : float = 2.0
3
4  > (float (int (7.0 / 2.5))) * 2.5;;
5  val it : float = 5.0
6
7  > (float (int (7.0 / 2.5))) * 2.5 + 7.0 % 2.5;;
8  val it : float = 7.0
```

Arithmetic using `float` will not cause over- and underflow problems, since the IEEE 754 standard includes the special numbers $\pm\infty$ and NaN. As shown in Listing 5.23, no exception is thrown.

Listing 5.23: Floating point numbers include infinity and Not-a-Number.

```
1  > 1.0/0.0;;
2  val it : float = infinity
3
4  > 0.0/0.0;;
5  val it : float = nan
```

However, the `float` type has limited precision, since there is only a finite number of numbers that can be stored in a float. E.g., addition and subtraction can give surprising results, as demonstrated in Listing 5.24.

Listing 5.24: Floating point arithmetic has finite precision.

```
1  > 357.8 + 0.1 - 357.9;;
2  val it : float = 5.684341886e-14
```

That is, addition and subtraction associates to the left, hence the expression is interpreted as `(357.8 + 0.1) - 357.9` and we see that we do not get the expected 0. The reason is that the calculation is done stepwise, and in the process, the numbers are represented using the imprecise floating point standard. Thus, `357.8 + 0.1` is represented as a number close to but not identical to what `357.9` is represented as, and thus, when subtracting these two representations, we get a very small nonzero number. Such errors tend to accumulate, and comparing the result of expressions of floating point values should, therefore, be treated with care. Thus, **equivalence of two floating point expressions should only be**   Advice **considered up to sufficient precision, e.g., comparing** `357.8 + 0.1` **and** `357.9` **up to** `1e-10` **precision should be tested as,** `abs ((357.8 + 0.1) - 357.9) < 1e-10`**.**

5.6. Char and String Arithmetic

Addition is the only operator defined for characters. Nevertheless, character arithmetic is often done by casting to an integer. A typical example is the conversion of character case, e.g., to convert the lowercase character 'z' to uppercase. Here, we use the *ASCIIbetical order*, add the difference between any Basic Latin Block letters in upper- and lowercase as   · ASCIIbetical order **integers**, and cast back to `char`, see Listing 5.25.

Listing 5.25: Converting case by casting and integer arithmetic.

```
1  > char (int 'z' - int 'a' + int 'A');;
2  val it : char = 'Z'
```

I.e., the code point difference between upper and lower case for any alphabetical character 'a' to 'z' is constant, hence we can change case by adding or subtracting the difference between any corresponding character. Unfortunately, this does not generalize to characters from other languages.

A large collection of operators and functions exist for `string`. The simplest is concatenation using the "+" operator, as demonstrated in Listing 5.26.

Listing 5.26: Example of string concatenation.

```
1  > "hello" + " " + "world";;
2  val it : string = "hello world"
```

Characters and strings cannot be concatenated, which is why the above example used the string of a space `" "` instead of the space character `' '`. The characters of a string may be indexed as using the `.[]` notation. This is demonstrated in Listing 5.27.

· `.[]`

Listing 5.27: String indexing using square brackets.

```
1   > "abcdefg".[0];;
2   val it : char = 'a'
3
4   > "abcdefg".[3];;
5   val it : char = 'd'
6
7   > "abcdefg".[3..];;
8   val it : string = "defg"
9
10  > "abcdefg".[..3];;
11  val it : string = "abcd"
12
13  > "abcdefg".[1..3];;
14  val it : string = "bcd"
15
16  > "abcdefg".[*];;
17  val it : string = "abcdefg"
```

Notice that the first character has index 0, and to get the last character in a string, we use the string's `length` property. This is done as shown in Listing 5.28.

Listing 5.28: String length attribute and string indexing.

```
1  > "abcdefg".Length;;
2  val it : int = 7
3
4  > "abcdefg".[7-1];;
5  val it : char = 'g'
```

Since index counting starts at 0, and since the string length is 7, the index of the last character is 6. There is a long list of built-in functions in `System.String` for working with strings, some of which will be discussed in Section 11.1.

The *dot notation* is an example of Structured programming, where technically speaking, the string `"abcdefg"` is an immutable *object* of *class* `string`, `[]` is an object *method*, and `Length` is a property. For more on objects, classes, and methods, see Chapter 20.

· dot notation
· object
· class
· method

Strings are compared letter by letter. For two strings to be equal, they must have the same length and all the letters must be identical. E.g., `"abs" = "absalon"` is false, while `"abs" = "abs"` is true. The "`<>`" operator is the boolean negation of the "`=`" operator, e.g., `"abs" <> "absalon"` is true, while `"abs" <> "abs"` is false. For the "<" , "<=", ">",

and ">=" operators, the strings are ordered alphabetically, such that `"abs" < "absalon"` `&& "absalon" < "milk"` is true, that is, the "<" operator on two strings is true if the left operand should come before the right when sorting alphabetically. The algorithm for deciding the boolean value of `leftOp < rightOp` is as follows: we start by examining the first character, and if `leftOp.[0]` and `rightOp.[0]` are different, then `leftOp < rightOp` is equal to `leftOp.[0] < rightOp.[0]`. E.g., `"milk" < "abs"` is the same as `'m' < 'a'`, which is false, since the letter 'm' does not come before the letter 'a' in the alphabet, or more precisely, the codepoint of 'm' is not less than the codepoint of 'a'. If `leftOp.[0]` and `rightOp.[0]` are equal, then we move on to the next letter and repeat the investigation, e.g., `"abe" < "abs"` is true, since `"ab" = "ab"` is true and `'e' < 's'` is true. If we reach the end of either of the two strings, then the shorter word is smaller than the longer word, e.g., `"abs" < "absalon"` is true, while `"abs" < "abs"` is false. The "<=", ">", and ">=" operators are defined in a similar manner.

5.7. Programming Intermezzo: Hand Conversion Between Decimal and Binary Numbers

Conversion of integers between decimal and binary form is a key concept one must grasp in order to understand some of the basic properties of calculations on the computer. Converting from binary to decimal is straightforward if using the power-of-two algorithm, i.e., given a sequence of $n + 1$ binary digits b_i written as $b_n b_{n-1} \ldots b_0$, and where b_n and b_0 are the most and least significant bits respectively, then the decimal value is calculated as,

$$v = \sum_{i=0}^{n} b_i 2^i \tag{5.2}$$

For example, $10011_2 = 1 + 2 + 16 = 19$. Converting from decimal to binary is a little more complex, but a simple divide-by-two algorithm exists. The key to understanding the divide-by-two algorithm is to realize that dividing a number by two is equivalent to shifting its binary representation one position to the right. E.g., $10 = 1010_2$ and $10/2 = 5 = 101_2$. Odd numbers have $b_0 = 1$, e.g., $11_{10} = 1011_2$ and $11_{10}/2 = 5.5 = 101.1_2$. Hence, if we divide any number by two and get a non-integer number, then its least significant bit was 1. Another way to express this is to say that the least significant bit is the remainder after integer division by two. Sequential application of this idea leads directly to the divide-by-two algorithm. E.g., if we were to convert the number 11_{10} in decimal form to binary form, we would perform the following steps:

$$11 \text{ div } 2 = 5, \; 11 \text{ rem } 2 = \boxed{1}$$
$$5 \text{ div } 2 = 2, \; 5 \text{ rem } 2 = \boxed{1}$$
$$2 \text{ div } 2 = 1, \; 2 \text{ rem } 2 = \boxed{0}$$
$$1 \text{ div } 2 = 0, \; 1 \text{ rem } 2 = \boxed{1}$$

Here we used div and rem to signify the integer division and remainder operators. The algorithm stops when the result of integer division is zero. Reading off the remainder from below and up, we find the sequence 1011_2, which is the binary form of the decimal

number 11_{10}. Using the interactive mode, we can perform the same calculation, as shown in Listing 5.29.

Listing 5.29: Converting the number 11_{10} to binary form.

```
 1  > printfn "(%d, %d)" (11 / 2) (11 % 2);;
 2  (5, 1)
 3  val it : unit = ()
 4  > printfn "(%d, %d)" (5 / 2) (5 % 2);;
 5  (2, 1)
 6  val it : unit = ()
 7  > printfn "(%d, %d)" (2 / 2) (2 % 2);;
 8  (1, 0)
 9  val it : unit = ()
10  > printfn "(%d, %d)" (1 / 2) (1 % 2);;
11  (0, 1)
12  val it : unit = ()
```

Thus, by reading the second integer-response from `printfn` from below and up, we again obtain the binary form of 11_{10} to be 1011_2. For integers with a fractional part, the divide-by-two algorithm may be used on the whole part, while multiply-by-two may be used in a similar manner on the fractional part.

6 | Values and Functions

In this chapter, we will see how we can bind expressions to identifiers either as new constants, functions, or operators, how this saves time when building large programs, and how this makes programs easier to read and debug. As an example, consider the following problem,

> **Problem 6.1**
>
> For given set constants a, b, and c, solve for x in
>
> $$ax^2 + bx + c = 0 \tag{6.1}$$

To solve for x we use the quadratic formula from elementary algebra,

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}, \tag{6.2}$$

which gives the general solution for any values of the coefficients. Here, we will assume a positive discriminant, $b^2 - 4ac > 0$. In order to write a program where the code may be reused later, we define a function `discriminant : float -> float -> float -> float`, that is, a function that takes 3 arguments, `a`, `b`, and `c`, and calculates the discriminant. Likewise, we will define `positiveSolution : float -> float -> float -> float` and `negativeSolution : float -> float -> float -> float`, that also take the polynomial's coefficients as arguments and calculate the solution corresponding to choosing the positive and negative sign for $\pm$ in the equation. Details on function definition is given in Section 6.2. Our solution thus looks like Listing 6.1.

> **Listing 6.1 identifiersExample.fsx:**
> **Finding roots for quadratic equations using function name binding.**
>
> ```
> 1  let discriminant a b c = b ** 2.0 - 4.0 * a * c
> 2  let positiveSolution a b c = (-b + sqrt (discriminant a b c))
>      / (2.0 * a)
> 3  let negativeSolution a b c = (-b - sqrt (discriminant a b c))
>      / (2.0 * a)
> 4
> 5  let a = 1.0
> 6  let b = 0.0
> 7  let c = -1.0
> 8  let d = discriminant a b c
> 9  let xp = positiveSolution a b c
> 10 let xn = negativeSolution a b c
> 11 do printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
> 12 do printfn "  has discriminant %A and solutions %A and %A" d
>      xn xp
> ```
>
> --
>
> ```
> 1  $ fsharpc --nologo identifiersExample.fsx && mono
>      identifiersExample.exe
> 2  0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
> 3    has discriminant 4.0 and solutions -1.0 and 1.0
> ```

Here, we have further defined names of values `a`, `b`, and `c` which are used as inputs to our functions, and the results of function application are bound to the names `d`, `xn`, and `xp`. The names of functions and values given here are examples of identifiers, and with these, we may reuse the quadratic formulas and calculated values later, while avoiding possible typing mistakes and reducing the amount of code which needs to be debugged.

The use of identifiers is central in programming. For F#, not to be confused with built-in functionality, identifiers must follow a specific set of rules:

· identifier

Identifier

- Identifiers are used as names for values, functions, types etc.

- They must start with a Unicode letter or underscore '_', but can be followed by zero or more of letters, digits, and a range of special characters except for SP, LF, and CR (space, line feed, and carriage return). See Appendix C.3 for more on codepoints that represents letters.

- They can also be a sequence of identifiers separated by a period.

- They cannot be keywords, see Table 6.1.

Examples of identifiers are: `a`, `theCharacter9`, `Next_Word`, `_tok`, and `f.sharp.rocks`. Since programmers often work in multilingual environment dominated by the English language it is advicable to **restrict identifiers to use letters from the English alphabet,** Advice **numbers, period, and '_'.** However, the number of possible identifiers is enormous. The full definition refers to the Unicode general categories described in Appendix C.3, and there are currently 19.345 possible Unicode code points in the letter category and 2.245 possible Unicode code points in the special character category.

Type	Keyword
Regular	abstract, and, as, assert, base, begin, class, default, delegate, do, done, downcast, downto, elif, else, end, exception, extern, false, finally, for, fun, function, global, if, in, inherit, inline, interface, internal, lazy, let, match, member, module, mutable, namespace, new, null, of, open, or, override, private, public, rec, return, sig, static, struct, then, to, true, try, type, upcast, use, val, void, when, while, with, and yield.
Reserved	atomic, break, checked, component, const, constraint, constructor, continue, eager, fixed, fori, functor, include, measure, method, mixin, object, parallel, params, process, protected, pure, recursive, sealed, tailcall, trait, virtual, and volatile.
Symbolic	let!, use!, do!, yield!, return!, |, ->, <-, ., :, (,), [,], [<, >], [|, |], {, }, ', #, :?>, :?, :>, .., ::, :=, ;;, ;, =, _, ?, ??, (*), <@, @>, <@@, and @@>.
Reserved symbolic	~ and `

Table 6.1.: Table of (possibly future) *keywords* and symbolic keywords in F#.

Identifiers may be used to carry information about their intended content and use, and careful selection of identifiers can aid programmers to communicate thoughts about the code. Thus, identifiers are often a word or several concatenated words conveying some relevant meaning. For example, in the function definition let discriminant a b c = b ** 2.0 - 4.0 * a * c, the function identifier has been chosen to be discriminant. F# places no special significance to the word 'discriminant', and the program would work exactly the same had the function been called let f a b c = b ** 2.0 - 4.0 * a * c. However, to programmers, the word 'discriminant' informs us of the intended role of the function and thus is much preferred. This is a general principle: **identifier names** **should be chosen to reflect their semantic value.**. The arguments a, b, and c are short, but adheres to a textbook tradition of elementary algebra. Again, we might as well have used, let discriminant c a b = a ** 2.0 - 4.0 * c * b, which is semantically identical to the original expression, but due to tradition, this would confuse most readers of the code. Thus, **identifier names should be chosen consistently with the readers'** **traditions.** Finally, identifiers are often concatenations of words, as positiveSolution in Listing 6.1. Concatenations can be difficult to read. Without the capitalization of the second word, we would have had positivesolution. This is readable at most times, but takes longer time to understand in general. Typical solutions are to use a separator, such as positive_solution, *lower camel case* also known as *mixed case* as in the example positiveSolution, and *upper camel case* also known as *pascal case* as PositiveSolution. In this book, we use lower camel case except where F# requires a capital first letter. Again, the choice does not influence what a program does, only how readable it is to a fellow programmer. The important part is that **identifier names consisting of concatenated** **words are often preferred over names with few character, and concatenation** **should be emphasized, e.g., by camel casing.** Choosing the length of identifier names is a balancing act, since when working with large programs, very long identifier names can be tiresome to write, and a common practice is that the length of identifier names is proportional to the complexity of the program. I.e., complex programs use long names, simple programs use short names. What is complex and what is simple is naturally in the eye of the beholder, but when you program, remember that a future reader of the program most likely has not had time to work with the problem as long as the programmer, thus

Advice

Advice

· lower camel case
· mixed case
· upper camel case
· pascal case

Advice

choose identifier names as if you were to explain the meaning of a program to a knowledgeable outsider. · Advice

Another key concept in F# is expressions. An expression can be a mathematical expression, such as $3 * 5$, a function application, such as $f3$, and many other things. Central in this chapter is the binding of values and functions to identifiers, which is done with the keyword `let`, e.g., `let a = 1.0`.

Expressions are the main workhorse of F# and have an enormous variety in how they may be written. We will in this book gradually work through some of the more important facets.

· expression

Expressions

- An Expression is a computation such as `3 * 5`.

- They can be value bindings between identifiers and expressions that evaluate to a value or a function, see Sections 6.1 and 6.2.

- They can be `do`-bindings that produce side-effects and whose result are ignored, see Section 6.2

- They can be assignments to variables, see Section 6.1.

- They can be a sequence of expressions separated with the *";"* lexeme. · ;

- They can be annotated with a type by using the *":"* lexeme. · :

Before we begin a deeper discussion on bindings, note that F# adheres to two different syntaxes: *verbose* and *lightweight*. In the verbose syntax, newlines and whitespaces are generally ignored, while in lightweight syntax, certain keywords and lexemes may be replaced by newlines and whitespaces. The lightweight syntax is the most common, but the syntaxes may be mixed, and we will highlight the options, when relevant. · verbose syntax · lightweight syntax

6.1. Value Bindings

Binding identifiers to literals, or expressions that are evaluated to be values, is called *value-binding*, and examples are `let a = 3.0` and `let b = cos 0.9`. Value bindings have the following syntax: · value-binding

```
Listing 6.2:   Value binding expression.

1   let <valueIdent> = <bodyExpr> [in <expr>]
```

The `let` keyword binds a value-identifier `<valueIdent>` with an expression `<bodyExpr>` that evaluates to a value. The following square bracket notation `[]` means that the enclosed is optional, and F# is able to identify whether or not the optional part is used by whether or not the `in` keyword is present in the binding expression. If the `in` keyword is used, then the value-identifier is a local definition in the following `<expr>` expression but in the following lines. For lightweight syntax, the `in` keyword is replaced with a newline, and the binding *is* valid in the following lines at the level of scope of the value-binding, or deeper, *lexically*. · let · in · lexically

The value identifier annotated with a type by using the *":"* lexeme followed by the name of   · :
a type, e.g., `int`. The *"_"* lexeme may be used as a value-identifier. This lexeme is called   · _
the *wildcard* pattern, and for value-bindings it means that `<bodyExpr>` is evaluated, but   · wildcard
the result is discarded. See Chapter 15 for more details on patterns.

For example, letting the identifier `p` be bound to the value `2.0` and using it in an expression
is done as shown in Listing 6.3.

Listing 6.3 letValue.fsx:
The identifier p is used in the expression following the `in` keyword.

```
1  let p = 2.0 in do printfn "%A" (3.0 ** p)
```
```
1  $ fsharpc --nologo letValue.fsx && mono letValue.exe
2  9.0
```

F# will ignore most newlines between lexemes, i.e., the above is equivalent to writing as
shown in Listing 6.4.

Listing 6.4 letValueLF.fsx:
Newlines after `in` make the program easier to read.

```
1  let p = 2.0 in
2  do printfn "%A" (3.0 ** p)
```
```
1  $ fsharpc --nologo letValueLF.fsx && mono letValueLF.exe
2  9.0
```

F# also allows for an alternative notation called *lightweight syntax*, where e.g., the `in`   · lightweight syntax
keyword is replaced with a newline, and the expression starts on the next line at the same
column as `let` starts in, i.e., the above is equivalent to Listing 6.5.

Listing 6.5 letValueLightWeight.fsx:
Lightweight syntax does not require the `in` keyword, but the expression
must be aligned with the `let` keyword.

```
1  let p = 2.0
2  do printfn "%A" (3.0 ** p)
```
```
1  $ fsharpc --nologo letValueLightWeight.fsx
2  $ mono letValueLightWeight.exe
3  9.0
```

The same expression in interactive mode will also show with the inferred types, as shown
in Listing 6.6.

Listing 6.6: Interactive mode also outputs inferred types.

```
1  > let p = 2.0
2  - do printfn "%A" (3.0 ** p);;
3  9.0
4  val p : float = 2.0
5  val it : unit = ()
```

By the `val` keyword in the line `val p : float = 2.0`, we see that `p` is inferred to be of type `float` and bound to the value `2.0`. The inference is based on the type of the right-hand-side which is `float`. Identifiers may be defined to have a type using the ":" lexeme, but the types on the left-hand-side and right-hand-side of the "=" lexeme must be identical. Mixing types gives an error, as shown in Listing 6.7.

Listing 6.7 letValueTypeError.fsx:
Binding error due to type mismatch.

```
1  let p : float = 3
2  do printfn "%A" (3.0 ** p)
```

```
1  $ fsharpc --nologo letValueTypeError.fsx && mono
     letValueTypeError.exe
2
3  letValueTypeError.fsx(1,17): error FS0001: This expression was
     expected to have type
4      'float'
5  but here has type
6      'int'
```

Here, the left-hand-side is defined to be an identifier of type float, while the right-hand-side is a literal of type integer.

An expression can be a sequence of expressions separated by the lexeme ";", see Listing 6.8.

Listing 6.8 letValueSequence.fsx:
A value-binding for a sequence of expressions.

```
1  let p = 2.0 in do printfn "%A" p; do printfn "%A" (3.0 ** p)
```

```
1  $ fsharpc --nologo letValueSequence.fsx && mono
     letValueSequence.exe
2  2.0
3  9.0
```

The lightweight syntax automatically inserts the ";" lexeme at newlines, hence using the lightweight syntax, the above is the same as shown in Listing 6.9.

Listing 6.9 letValueSequenceLightWeight.fsx:
A value-binding for a sequence using lightweight syntax.

```
1  let p = 2.0
2  do printfn "%A" p
3  do printfn "%A" (3.0 ** p)
```

```
1  $ fsharpc --nologo letValueSequenceLightWeight.fsx
2  $ mono letValueSequenceLightWeight.exe
3  2.0
4  9.0
```

A key concept of programming is *scope*. In F#, the scope of a value-binding is lexically. · scope
This means that when F# seeks the value bound to a name, it looks left and upward in
the program text for its `let`-binding in the present or higher scopes, see Listing 6.10 for
an example.

Listing 6.10 letValueScopeLower.fsx:
Redefining identifiers is allowed in lower scopes.

```
1  let p = 3 in let p = 4 in do printfn " %A" p;
```

```
1  $ fsharpc --nologo letValueScopeLower.fsx && mono
     letValueScopeLower.exe
2    4
```

Some special bindings are mutable, in which case F# uses the dynamic scope, that is, the
value of a binding is defined by when it is used. This will be discussed in Section 6.7.

Scopes are given levels, and scopes may be nested, where the nested scope has a level one
lower than its parent.[1] F# distinguishes between the top and lower levels, and at the top
level in the lightweight syntax, redefining values is not allowed, as shown in Listing 6.11.

Listing 6.11 letValueScopeLowerError.fsx:
Redefining identifiers is not allowed in lightweight syntax at top level.

```
1  let p = 3
2  let p = 4
3  do printfn "%A" p;
```

```
1  $ fsharpc --nologo -a letValueScopeLowerError.fsx
2
3  letValueScopeLowerError.fsx(2,5): error FS0037: Duplicate
     definition of value 'p'
```

However, using parentheses, we create a *code block*, i.e., a *nested scope*, and then redefining · code block
is allowed, as demonstrated in Listing 6.12. · nested scope

---

[1]Jon: **Drawings would be good to describe scope**

Listing 6.12 letValueScopeBlockAlternative3.fsx:
A block may be created using parentheses.

```
1  (
2     let p = 3
3     let p = 4
4     do printfn "%A" p
5  )
```

```
1  $ fsharpc --nologo letValueScopeBlockAlternative3.fsx
2  $ mono letValueScopeBlockAlternative3.exe
3  4
```

Nevertheless, **avoid reusing names unless it's in a deeper scope.**   Advice

Inside the block in Listing 6.12 we used indentation, which is good practice, but not required here.

Bindings inside a nested scope are not available outside, as shown in Listing 6.13.

Listing 6.13 letValueScopeNestedScope.fsx:
Bindings inside a scope are not available outside.

```
1  let p = 3
2  (
3     let q = 4
4     do printfn "%A" q
5  )
6  do printfn "%A %A" p q
```

```
1  $ fsharpc --nologo -a letValueScopeNestedScope.fsx
2
3  letValueScopeNestedScope.fsx(6,22): error FS0039: The value or
      constructor 'q' is not defined. Maybe you want one of the
      following:
4      p
```

Nesting is a natural part of structuring code, e.g., through function definitions to be discussed in Section 6.2 and flow control structures to be discussed in Chapter 8. Blocking code by nesting is a key concept for making robust code that is easy to use by others, without the user necessarily needing to know the details of the inner workings of a block of code.

Defining blocks is used for controlling the extent of a lexical scope of bindings. For example, adding a second `printfn` statement, as in Listing 6.14, will print the value 4, last bound to the identifier p, since F# interprets the above as `let p = 3 in let p = 4 in (printfn "%A" p; printfn "%A" p)`.

Listing 6.14 letValueScopeBlockProblem.fsx:
Overshadowing hides the first binding.

```
1  let p = 3 in let p = 4 in do printfn "%A" p; do printfn "%A" p
```

```
1  $ fsharpc --nologo letValueScopeBlockProblem.fsx
2  $ mono letValueScopeBlockProblem.exe
3  4
4  4
```

Had we intended to print the two different values of p, then we should have created a block as in Listing 6.15.

Listing 6.15 letValueScopeBlock.fsx:
Blocks allow for the return to the previous scope.

```
1  let p = 3 in (let p = 4 in do printfn "%A" p); do printfn "%A"
     p;
```

```
1  $ fsharpc --nologo letValueScopeBlock.fsx && mono
     letValueScopeBlock.exe
2  4
3  3
```

6.2. Function Bindings

A function is a mapping between an input and output domain. A key advantage of using functions when programming is that they encapsulate code into smaller units, that are easier to debug and may be reused. F# is a functional first programming language and offers a number of alternative methods for specifying parameters, which will be discussed in this section. Binding identifiers to functions follows a syntax similar to value-binding,

Listing 6.16: Function binding expression

```
1  let <funcIdent> <arg> {<arg>} | () = <bodyExpr> [in <expr>]
```

Here `<funcIdent>` is an identifier and is the name of the function, `<arg>` is zero or more identifiers, that bind to the value used when calling the function, and which is to be used in the body of the function, the expression `<bodyExpr>`. The | notation denotes a choice, i.e., either that on the left-hand-side or that on the right-hand-side. Thus `let f x = x * x` and `let f () = 3` are valid function bindings, but `let f = 3` would be a value binding, not a function binding. The arguments and the function may be annotated with a type, in which case for arguments we write

Listing 6.17:   Function binding expression

```
let <funcIdent> (<arg> : <type>) {(<arg> : <type>)} : <type> |
    () : <type> = <bodyExpr> [in <expr>]
```

where `<type>` is a name of an existing type. The argument types are given in parentheses, and the return type is given last.

Functions are a key concept in F#, and in this chapter we will discuss the very basics. Recursive functions will be discussed in Chapter 8 and higher-order functions in Chapter 16.

An example of defining a function and using it in interactive mode is shown in Listing 6.18.

Listing 6.18: An example of a binding of an identifier and a function.

```
> let sum (x : float) (y : float) : float = x + y in
- let c = sum 357.6 863.4 in
- do printfn "%A" c;;
1221.0
val sum : x:float -> y:float -> float
val c : float = 1221.0
val it : unit = ()
```

Here we see that the function is interpreted to have the type `val sum : x:float -> y:float -> float`. The "`->`" lexeme means a mapping between sets, in this case, floats. The function is also a higher order function, to be discussed in detail below, and here it suffices to think of `sum` as a function that takes 2 floats as argument and returns a float.

Not all types need to be declared, just a sufficient number for F# to be able to infer the types for the full statement. For the example, one is sufficient, and we could just have declare the type of the result, as in Listing 6.19.

Listing 6.19 letFunctionAlterantive.fsx:
Not every type needs to be declared.

```
let sum x y : float = x + y
```

Or even just one of the arguments, as in Listing 6.20.

Listing 6.20 letFunctionAlterantive2.fsx:
Just one type is often enough for F# to infer the rest.

```
let sum (x : float) y = x + y
```

In both cases, since the + *operator* is only defined for *operands* of the same type, declaring the type of either arguments or result implies the type of the remainder. As for values, lightweight syntax automatically inserts the keyword `in` and the lexeme "`;`", as shown in Listing 6.21.

· operator

· operand

Listing 6.21 letFunctionLightWeight.fsx:
Lightweight syntax for function definitions.

```
1  let sum x y : float = x + y
2  let c = sum 357.6 863.4
3  do printfn "%A" c
```

```
1  $ fsharpc --nologo letFunctionLightWeight.fsx
2  $ mono letFunctionLightWeight.exe
3  1221.0
```

Arguments need not always be inferred to types, but may be of the generic type when *type safety* is ensured, as shown in Listing 6.22.   · type safety

---

Listing 6.22: Type safety implies that a function will work for any type.

```
1   > let second x y = y
2   - let a = second 3 5
3   - do printfn "%A" a
4   - let b = second "horse" 5.0
5   - do printfn "%A" b;;
6   5
7   5.0
8   val second : x:'a -> y:'b -> 'b
9   val a : int = 5
10  val b : float = 5.0
11  val it : unit = ()
```

Here, the function `second` does not use the first argument `x`, which therefore can be of any type, and which F#, therefore, calls `'a`. The type of the second element, `y`, can also be of any type and not necessarily the same as `x`, so it is called `'b`. Finally, the result is the same type as `y`, whatever it is. This is an example of a *generic function*, since it will work   · generic function
on any type.

A function may contain a sequence of expressions but must return a value. E.g., the quadratic formula may be written as shown in Listing 6.23.

> **Listing 6.23 identifiersExampleAdvance.fsx:**
> **A function may contain sequences of expressions.**
>
> ```
> 1  let solution a b c sgn =
> 2    let discriminant a b c =
> 3      b ** 2.0 - 4.0 * a * c
> 4    let d = discriminant a b c
> 5    (-b + sgn * sqrt d) / (2.0 * a)
> 6
> 7  let a = 1.0
> 8  let b = 0.0
> 9  let c = -1.0
> 10 let xp = solution a b c +1.0
> 11 let xn = solution a b c -1.0
> 12 do printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
> 13 do printfn "  has solutions %A and %A" xn xp
> ```
>
> ```
> 1  $ fsharpc --nologo identifiersExampleAdvance.fsx
> 2  $ mono identifiersExampleAdvance.exe
> 3  0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
> 4    has solutions -1.0 and 1.0
> ```

Here, we used the lightweight syntax, where the "=" identifies the start of a nested scope, and F# identifies the scope by indentation. The amount of space used for indentation does not matter, but all lines in the same scope must use the same amount. The scope ends before the first line with the previous indentation or none. Notice how the last expression is not bound to an identifier, but is the result of the function, i.e., in contrast to many other languages, F# does not have an explicit keyword for returning values, but requires a final expression, which will be returned to the caller of the function. Note also that since the function `discriminant` is defined in the nested scope of `solution`, and becausethe scope ends before `let a = 1.0`, `discriminant` cannot be called outside `solution`.

Lexical scope and function definitions can be a cause of confusion, as the following example · lexical scope
in Listing 6.24 shows.[2]

> **Listing 6.24 lexicalScopeNFunction.fsx:**
> **Lexical scope means that $f(z) = 3x$ and not $4x$ at the time of calling.**
>
> ```
> 1  let testScope x =
> 2    let a = 3.0
> 3    let f z = a * z
> 4    let a = 4.0
> 5    f x
> 6  do printfn "%A" (testScope 2.0)
> ```
>
> ```
> 1  $ fsharpc --nologo lexicalScopeNFunction.fsx
> 2  $ mono lexicalScopeNFunction.exe
> 3  6.0
> ```

Here, the value-binding for `a` is redefined after it has been used to define a helper function `f`. So which value of `a` is used when we later apply `f` to an argument? To resolve the

---

[2]Jon: **Add a drawing or possibly a spell-out of lexical scope here.**

confusion, remember that value-binding is lexically defined, i.e., the binding `let f z = a * z` uses the value of `a` as it is defined by the ordering of the lines in the script, not dynamically by when `f` was called. Hence, **think of lexical scope as substitution of an identifier with its value or function immediately at the place of definition.** Since `a` and `3.0` are synonymous in the first lines of the program, the function `f` is really defined as `let f z = 3.0 * z`.

Advice

Functions do not need a name, but may be declared as an *anonymous function* using the *fun* keyword and the "`->`" lexeme, as shown in Listing 6.25.

· anonymous function
· `fun`
· `->`

Listing 6.25 functionDeclarationAnonymous.fsx: Anonymous functions are functions as values.

```
1  let first = fun x y -> x
2  do printfn "%d" (first 5 3)
```

```
1  $ fsharpc --nologo functionDeclarationAnonymous.fsx
2  $ mono functionDeclarationAnonymous.exe
3  5
```

Here, a name is bound to an anonymous function which returns the first of two arguments. The difference to `let first x y = x` is that anonymous functions may be treated as values, meaning that they may be used as arguments to other functions and the new values may be reassigned to their identifiers when mutable, as will be discussed in Section 6.7. A common use of anonymous functions is as arguments to other functions, as demonstrated in Listing 6.26.

Listing 6.26 functionDeclarationAnonymousAdvanced.fsx: Anonymous functions are often used as arguments for other functions.

```
1  let apply f x y  = f x y
2  let mul = fun a b -> a * b
3  do printfn "%d" (apply mul 3 6)
```

```
1  $ fsharpc --nologo functionDeclarationAnonymousAdvanced.fsx
2  $ mono functionDeclarationAnonymousAdvanced.exe
3  18
```

Note that here `apply` is given 3 arguments: the function `mul` and 2 integers. It is not given the result of `mul 3 6`, since that would not match the definition of `apply`. **Anonymous functions and functions as arguments are powerful concepts, but tend to make programs harder to read, and their use should be limited.**

Advice

The result of one function is often used as an argument of another. This is function composition, and an example is shown in Listing 6.27.

Listing 6.27 functionComposition.fsx:
Composing functions using intermediate bindings.

```
1  let f x = x + 1
2  let g x = x * x
3
4  let a = f 2
5  let b = g a
6  let c = g (f 2)
7  do printfn "a = %A, b = %A, c = %A" a b c
```

```
1  $ fsharpc --nologo functionComposition.fsx
2  $ mono functionComposition.exe
3  a = 3, b = 9, c = 9
```

In the example we combine two functions `f` and `g` by storing the result of `f 2` in `a` and using that as argument of `g`. This is the same as `g (f 2)`, and in the later case, the compile creates a temporary value for `f 2`. Such compositions are so common in F# that a special set of operators has been invented, called the *piping* operators: *"/>"* and *"</"*. They are used as demonstrated in Listing 6.28.

· piping
· |>
· <|

Listing 6.28 functionPiping.fsx:
Composing functions by piping.

```
1  let f x = x + 1
2  let g x = x * x
3
4  let a = g (f 2)
5  let b = 2 |> f |> g
6  let c = g <| (f <| 2)
7  do printfn "a = %A, b = %A, c = %A" a b c
```

```
1  $ fsharpc --nologo functionPiping.fsx && mono
     functionPiping.exe
2  a = 9, b = 9, c = 9
```

The example shows regular composition, left-to-right, and right-to-left piping. The word piping is a picturial description of data as if it were flowing through pipes, where functions are connection points of pipes distributing data in a network. The three expressions in Listing 6.28 perform the same calculation. The left-to-right piping in line 5 corresponds to the left-to-right reading direction, i.e., the value 2 is used as argument to `f`, and the result is used as argument to `g`. In contrast, right-to-left piping in line 6 has the order of arithmetic composition as line 4. Unfortunately, since the piping operators are left-associatitve, without the parenthesis in line 6 `g <| f <| 2`, F# would read the expression as `(g <| f) <| 2`. That would have been an error, since `g` takes an integer as argument, not a function. F# can also define composition on a function level. Further discussion on this is deferred to Chapter 16. The piping operator comes in four variants: *"||>"*, *"<||"*, *"|||>"*, and *"<|||"*. These allow for piping between pairs and triples to functions of 2 and 3 arguments, see Listing 6.29 for an example.

> **Listing 6.29 functionTuplePiping.fsx:**
> **Tuples can be piped to functions of more than one argument.**
>
> ```
> 1  let f x = printfn "%A" x
> 2  let g x y = printfn "%A %A" x y
> 3  let h x y z = printfn "%A %A %A" x y z
> 4
> 5  1 |> f
> 6  (1, 2) ||> g
> 7  (1, 2, 3) |||> h
> ```
> --
> ```
> 1  $ fsharpc --nologo functionTuplePiping.fsx
> 2  $ mono functionTuplePiping.exe
> 3  1
> 4  1 2
> 5  1 2 3
> ```

The example demonstrates right-to-left piping, left-to-right works analogously.[3]

A *procedure* is a generalization of the concept of functions, and in contrast to functions, procedures need not return values. This is demonstrated in Listing 6.30.

· procedure

> **Listing 6.30 procedure.fsx:**
> **A procedure is a function that has no return value, and in F# returns "()".**
>
> ```
> 1  let printIt a = printfn "This is '%A'" a
> 2  do printIt 3
> 3  do printIt 3.0
> ```
> --
> ```
> 1  $ fsharpc --nologo procedure.fsx && mono procedure.exe
> 2  This is '3'
> 3  This is '3.0'
> ```

In F#, this is automatically given the unit type as the return value. Procedural thinking is useful for *encapsulation* of scripts, but is prone to *side-effects*. For this reason, it is adviced to **prefer functions over procedures.** More on side-effects in Section 6.7.

· encapsulation

· side-effect

Advice

In F#, functions (and procedures) are *first-class citizens*, which means that functions are values: They may be passed as arguments, returned from a function, and bound to a name. For first-class citizens, the name it is bound to does not carry significance to the language, as, e.g., illustrated with the use of anonymous functions. Technically, a function is stored as a *closure*. A closure is a description of the function, its arguments, its expression, and the environment at the time it was created, i.e., the triple $(args, exp, env)$. Consider the listing in Listing 6.31.

· first-class citizens

· closure

---

[3]Jon: **Tuples have not yet been introduced!**

Listing 6.31 functionFirstClass.fsx:
The function `ApplyFactor` has a non-trivial closure.

```
1  let mul x y = x * y
2  let factor = 2.0
3  let applyFactor fct x =
4    let a = fct factor x
5    string a
6
7  do printfn "%g" (mul 5.0 3.0)
8  do printfn "%s" (applyFactor mul 3.0)
```

```
1  $ fsharpc --nologo functionFirstClass.fsx && mono
     functionFirstClass.exe
2  15
3  6
```

It defines two functions `mul` and `applyFactor`, where the latter is a higher order function taking another function as an argument and uses part of the environment to produce its result. The two closures are:

$$\texttt{mul} : (\text{args}, \text{exp}, \text{env}) = \big((\texttt{x, y}), (\texttt{x * y}), ()\big) \tag{6.3}$$

$$\texttt{applyFactor} : (\text{args}, \text{exp}, \text{env}) = \Big((\texttt{x, fct}), \big(\begin{smallmatrix}\texttt{let a = fct factor x}\\ \texttt{string a}\end{smallmatrix}\big), (\texttt{factor} \rightarrow 2.0)\Big) \tag{6.4}$$

The function `mul` does not use its environment, and everything needed to evaluate its expression are values for its arguments. The function `applyFactor` also takes two arguments, a function and a value. It uses `factor` from the environment, thus this is stored in its closure. When `mul` is given as an argument in Listing 6.31 line 8, then it is its closure which is given to `applyFactor`, and the closure contains everything that `applyFactor` requires to use `mul`. Likewise, if `applyFactor` is given as argument to yet another function, then its closure includes the relevant part of its environment at the time of definition, `factor`, such that when `applyFactor` is applied to two arguments, then its closure contains everything needed to evaluate its expression.

6.3. Operators

Operators are functions, and in F#, the infix multiplication operator + is equivalent to the function (+), as shown in Listing 6.32.

Listing 6.32 addOperatorNFunction.fsx:
Operators have function equivalents.

```
1  let a = 3.0
2  let b = 4.0
3  let c = a + b
4  let d = (+) a b
5  do printfn "%A plus %A is %A and %A" a b c d
```

```
1  $ fsharpc --nologo addOperatorNFunction.fsx
2  $ mono addOperatorNFunction.exe
3  3.0 plus 4.0 is 7.0 and 7.0
```

All operators have this option, and you may redefine them and define your own operators, but in F# names of user-defined operators are limited:

- A *unary operator* name can be: "+", "-", "+.", "-.", "%", "&", "&&", "~~", "~~~", · unary operator "~~~~", ..., apostropheOp. Here apostropheOp is an operator name starting with "!" and followed by one or more of either "!", "%", "&", "*", "+", "-", ".", "/", "<", "=", ">", "@", "^", "|", "~", but apostropheOp cannot be "!=".

- An *binary operator* name can be: "+", "-", "+.", "-.", "%", "&", "&&", ":=", "::", "$", · binary operator "?", dotOp. Here dotOp is an operator name starting with "." and followed by "+", "-", "+.", "-.", "%", "&", "&&", "-", "+", "||", "<", ">", "=", "|", "&", "^", "*", "/", "%", "!=". Only "?" and "?<-" may start with "?".

The precedence rules and associativity of user-defined operators follow the rules for which they share prefixes with built-in rules, see Table E.5. E.g., .*, +++, and <+ are valid operator names for infix operators, they have precedence as ordered, and their associativities are all left. Using ~ as the first character in the definition of an operator makes the operator unary and will not be part of the name. Examples of definitions and use of operators are,

---

Listing 6.33 operatorDefinitions.fsx:
Operators may be (re)defined by their function equivalent.

```
1  let (.*) x y = x * y + 1
2  printfn "%A" (3 .* 4)
3  let (+++) x y = x * y + y
4  printfn "%A" (3 +++ 4)
5  let (<+) x y = x < y + 2.0
6  printfn "%A" (3.0 <+ 4.0)
7  let (~+.) x = x+1
8  printfn "%A" (+.1)
```

```
1  $ fsharpc --nologo operatorDefinitions.fsx
2  $ mono operatorDefinitions.exe
3  13
4  16
5  true
6  2
```

Operators beginning with * must use a space in its definition. For example, without a

space (`*` would be confused with the beginning of a comment (`*`, see Chapter 7 for more on comments in the code.

Beware, redefining existing operators lexically redefines all future uses of the operators for all types, hence **it is not a good idea to redefine operators, but better to define** **new ones.** In Chapter 20 we will discuss how to define type-specific operators, including prefix operators.

<div style="text-align: right">Advice</div>

6.4. Do Bindings

Aside from `let` bindings that binds names with values or functions, sometimes we just need to execute code. This is called a *do* binding or, alternatively, a *statement*. The syntax is as follows:

<div style="text-align: right">· do
· do binding
· statement</div>

> **Listing 6.34:**   Syntax for `do` bindings.
>
> ```
> [do ]<expr>
> ```

The expression `<expr>` must return `unit`. The keyword `do` is optional in most cases, but using it emphasizes that the expression is not a function that returns a useful value. Procedures are examples of such expressions, and a very useful family of procedures are the `printf` family described below. In the remainder of this book, we will refrain from using the `do` keyword.

6.5. The Printf Function

A common way to output information to the console is to use one of the family of *printf* commands. These functions are special, since they take a variable number of arguments, and the number is decided by the first argument - the format string. The syntax for the `printf` commands are as follows:

<div style="text-align: right">· printf</div>

> **Listing 6.35:**   `printf` statement.
>
> ```
> printf <format-string> {<ident>}
> ```

The `formatString` is a string (simple or verbatim) with placeholders. The function `printf` prints `formatString` to the console, where all `placeholder` have been replaced by the values of the corresponding arguments formatted as specified. For example, in `printfn "1 2 %d" 3`, the `formatString` is `"1 2 %d"` and the placeholder is `%d`. When executed, `printf` will replace `%d` with the following argument, `3`, and print the result to the console: `1 2 3`. There are specifiers for all the basic types, and more, as elaborated in Table 6.2. The placeholder can be parametrized, e.g., the placeholder string `%8s` will print a right-aligned string which that is eight characters wide and padded with spaces, as needed. For floating point numbers, `%8f` will print a number that is exactly seven digits and a decimal point, making eight characters in total. Zeros are added after the decimal point, as needed. Alternatively, we may specify the number of digits after the decimal point, such that `%8.1f` will print a floating point number, aligned to the right, with one

Specifier	Type	Description
`%b`	`bool`	replaces with boolean value
`%s`	`string`	
`%c`	`char`	
`%d`, `%i`	basic integer	
`%u`	basic unsigned integers	
`%x`	basic integer	formatted as unsigned hexadecimal with lower case letters
`%X`	basic integer	formatted as unsigned hexadecimal with upper case letters
`%o`	basic integer	formatted as unsigned octal integer
`%f`, `%F`,	basic floats	formatted on decimal form
`%e`, `%E`,	basic floats	formatted on scientific form. Lower case uses "e" while upper case uses "E" in the formatting.
`%g`, `%G`,	basic floats	formatted on the shortest of the corresponding decimal or scientific form.
`%M`	decimal	
`%O`	Objects `ToString` method	
`%A`	any built-in types	formatted as a literal type
`%a`	`Printf.TextWriterFormat ->'a -> ()`	
`%t`	`(Printf.TextWriterFormat -> ()`	

Table 6.2.: Printf placeholder string

digit after the decimal point padded with spaces, as needed. The default is for the value to be right justified in the field, but left justification can be specified by the - character. For number types, you can specify their format by `"0"` for padding the number with zeros to the left when right justifying the number; `"+"` to explicitly show a plus sign for positive numbers; `SP` to enforce a space, where there otherwise would be a plus sign for positive numbers. The placeholder parameter may also be given as an argument to `printf` which case the placeholder should use the `*` character instead of an integer.

Examples of placeholder parametrization are shown in Listing 6.36.

> **Listing 6.36 printfExample.fsx:**
> **Examples of printf and some of its formatting options.**
>
> ```
> 1  let pi = 3.1415192
> 2  let hello = "hello"
> 3  printf "An integer: %d\n" (int pi)
> 4  printf "A float %f on decimal form and on %e scientific
>       form\n" pi pi
> 5  printf "A char '%c' and a string \"%s\"\n" hello.[0] hello
> 6  printf "Float using width 8 and 1 number after the decimal:\n"
> 7  printf "   \"%8.1f\" \"%8.1f\"\n" pi -pi
> 8  printf "   \"%08.1f\" \"%08.1f\"\n" pi -pi
> 9  printf "   \"% 8.1f\" \"% 8.1f\"\n" pi -pi
> 10 printf "   \"%-8.1f\" \"%-8.1f\"\n" pi -pi
> 11 printf "   \"%+8.1f\" \"%+8.1f\"\n" pi -pi
> 12 printf "   \"%8s\"\n\"%-8s\"\n" "hello" "hello"
> ```
>
> ```
> 1  $ fsharpc --nologo printfExample.fsx && mono printfExample.exe
> 2  An integer: 3
> 3  A float 3.141519 on decimal form and on 3.141519e+000
>       scientific form
> 4  A char 'h' and a string "hello"
> 5  Float using width 8 and 1 number after the decimal:
> 6     "     3.1" "    -3.1"
> 7     "000003.1" "-00003.1"
> 8     "     3.1" "    -3.1"
> 9     "3.1     " "-3.1    "
> 10    "    +3.1" "    -3.1"
> 11    "   hello"
> 12 "hello   "
> ```

Not all combinations of flags and identifier types are supported, e.g., strings cannot have the number of integers after the decimal point specified. The placeholder types `"%A"`, `"%a"`, and `"%t"` are special for F#, examples of their use are shown in Listing 6.37.

> **Listing 6.37 printfExampleAdvance.fsx:**
> **Custom format functions may be used to specialise output.**
>
> ```
> 1  let noArgument writer = printf "I will not print anything"
> 2  let customFormatter writer arg = printf "Custom formatter got:
>       \"%A\"" arg
> 3  printf "Print examples: %A, %A, %A\n" 3.0m 3uy "a string"
> 4  printf "Print function with no arguments: %t\n" noArgument
> 5  printf "Print function with 1 argument: %a\n" customFormatter
>       3.0
> ```
>
> ```
> 1  $ fsharpc --nologo printfExampleAdvance.fsx
> 2  $ mono printfExampleAdvance.exe
> 3  Print examples: 3.0M, 3uy, "a string"
> 4  Print function with no arguments: I will not print anything
> 5  Print function with 1 argument: Custom formatter got: "3.0"
> ```

The `%A` is special in that all built-in types, including tuples, lists, and arrays to be discussed

Function	Example	Description
`printf`	`printf "%d apples" 3`	Prints to the console, i.e., `stdout`
`printfn`		As `printf` and adds a newline.
`fprintf`	`fprintf stream "%d apples" 3`	Prints to a stream, e.g., `stderr` and `stdout`, which would be the same as `printf` and `eprintf`.
`fprintfn`		As `fprintf` but with added newline.
`eprintf`	`eprintf "%d apples" 3`	Prints to `stderr`
`eprintfn`		As `eprintf` but with added newline.
`sprintf`	`printf "%d apples" 3`	Return printed string
`failwithf`	`failwithf "%d failed apples" 3`	Prints to a string and used for raising an exception.

Table 6.3.: The family of printf functions.

in Chapter 11, can be printed using this formatting string, but notice that the formatting performed includes the named literal string. The two formatting strings `%t` and `%a` are options for user-customizing the formatting and will not be discussed further.

Beware, `formatString` is not a `string` but a `Printf.TextWriterFormat`, so to predefine a `formatString` as, e.g., `let str = "hello %s" in printf str "world"`, will be a type error.

The family of `printf` is shown in Table 6.3. The function `fprintf` prints to a stream, e.g., `stderr` and `stdout`, of type `System.IO.TextWriter`. For the moment it is sufficient to think of both `stderr` and `stdout` to be the console. Streams will be discussed in further detail in Chapter 19. The function `failwithf` is used with exceptions, see Chapter 18 for more details. The function has a number of possible return value types, and for testing, the *ignore* function ignores it all, e.g., `ignore (failwithf "%d failed apples" 3)`. · ignore

6.6. Reading from the Console

The `printf` and `printfn` functions allow us to write text on the screen. A program often needs to ask a user to input data, e.g., by typing text on a keyboard. Text typed on the keyboard is accessible through the `stdin` stream, and F# provides several library functions for capturing text typed on the keyboard. In the following section, we will briefly discuss the `System.Console.ReadLine` function. For more details and other methods of input see Chapters 19 and 23.

The function `System.Console.ReadLine` takes a unit value as an argument and returns the string the user typed. The program will not advance until the user presses newline. An example of a program that multiplies two floating point numbers supplied by a user is given in Listing 6.38,

Listing 6.38 userDialoguePrintf.fsx:
Interacting with a user using `ReadLine`.

```
1  printfn "To perform the multiplication of a and b"
2  printf "Enter a: "
3  let a = float (System.Console.ReadLine ())
4  printf "Enter b: "
5  let b = float (System.Console.ReadLine ())
6  printfn "a * b = %A" (a * b)
```

and an example dialogue is shown in Listing 6.39.

Listing 6.39: An example dialogue of running Listing 6.38. What the user typed has been framed in red boxes.

```
1  $ fsharpc --nologo userDialoguePrintf.fsx && mono
     userDialoguePrintf.exe
2  To perform the multiplication of a and b
3  Enter a: 3.5
4  Enter b: 7.4
5  a * b = 25.9
```

Note that the string is immediately cast to floats such that we can multiply the input using the float multiplication operator. This also implies that if the user inputs a non-number, then `mono` will halt with an error message.

6.7. Variables

Identifiers may be mutable, which means that the it may be rebound to a new value. Mutable identifiers are specified using the *mutable* keyword with the following syntax: · `mutable`

Listing 6.40:   Syntax for defining mutable values with an initial value.

```
1  let mutable <ident> = <expr> [in <expr>]
```

Changing the value of an identifier is called *assignment* and is done using the *"<-"* lexeme. · assignment
Assignments have the following syntax:[4] · `<-`

Listing 6.41:   Value reassignment for mutable variables.

```
1  <ident> <- <ident>
```

Mutable values is synonymous with the term *variable*. A variable is an area in the computer's working memory associated with an identifier and a type, and this area may be read from and written to during program execution, see Listing 6.42 for an example. · mutable value
 · variable

---

[4]Jon: **Discussion on heap and stack should be added here.**

Listing 6.42 mutableAssignReassingShort.fsx:
A variable is defined and later reassigned a new value.

```
1  let mutable x = 5
2  printfn "%d" x
3  x <- -3
4  printfn "%d" x
```

```
1  $ fsharpc --nologo mutableAssignReassingShort.fsx
2  $ mono mutableAssignReassingShort.exe
3  5
4  -3
```

Here, an area in memory was denoted `x`, initially assigned the integer value 5, hence the type was inferred to be `int`. Later, this value of `x` was replaced with another integer using the "`<-`" lexeme. The "`<-`" lexeme is used to distinguish the assignment from the comparison operator. For example, the statement `a = 3` in Listing 6.43 is not an assignment but a comparison which is evaluated to be false.

Listing 6.43: It is a common error to mistake "=" and "<-" lexemes for mutable variables.

```
1  > let mutable a = 0
2  - a = 3;;
3  val mutable a : int = 0
4  val it : bool = false
```

However, it is important to note that when the variable is initially defined, then the "="' operator must be used, while later reassignments must use the "`<-`" expression.

Assignment type mismatches will result in an error, as demonstrated in Listing 6.44.

Listing 6.44 mutableAssignReassingTypeError.fsx:
Assignment type mismatching causes a compile-time error.

```
1  let mutable x = 5
2  printfn "%d" x
3  x <- -3.0
4  printfn "%d" x
```

```
1  $ fsharpc --nologo mutableAssignReassingTypeError.fsx
2
3  mutableAssignReassingTypeError.fsx(3,6): error FS0001: This
     expression was expected to have type
4     'int'
5  but here has type
6     'float'
7  $ mono mutableAssignReassingTypeError.exe
8  Cannot open assembly 'mutableAssignReassingTypeError.exe': No
     such file or directory.
```

I.e., once the type of an identifier has been declared or inferred, it cannot be changed.

A typical variable is a counter of type integer, and a typical use of counters is to increment them, see Listing 6.45 for an example.

Listing 6.45 mutableAssignIncrement.fsx:
Variable increment is a common use of variables.

```
1  let mutable x = 5 // Declare a variable x and assign the value
      5 to it
2  printfn "%d" x
3  x <- x + 1 // Increment the value of x
4  printfn "%d" x
```

```
1  $ fsharpc --nologo mutableAssignIncrement.fsx
2  $ mono mutableAssignIncrement.exe
3  5
4  6
```

Using variables in expressions, as opposed to the left-hand-side of an assignment operation, reads the value of the variable. Thus, when using a variable as the return value of a function, then the value is copied from the local scope of the function to the scope from which it is called. This is demonstrated in Listing 6.46.

Listing 6.46: Returning a mutable variable returns its value.

```
1  > let g () =
2  -    let mutable y = 0
3  -    y
4  - printfn "%d" (g ());;
5  0
6  val g : unit -> int
7  val it : unit = ()
```

In the example we see that the type is a value, and not mutable.

Variables implement dynamic scope, that is, the value of an identifier depends on *when* it is used. This is in contrast to lexical scope, where the value of an identifier depends on *where* it is defined. As an example, consider the script in Listing 6.24 which defines a function using lexical scope and returns the number `6.0`, however, if `a` is made `mutable`, then the behavior is different, as shown in Listing 6.47.

Listing 6.47 dynamicScopeNFunction.fsx:
Mutual variables implement dynamic scope rules. Compare with Listing 6.24.

```
1  let testScope x =
2    let mutable a = 3.0
3    let f z = a * z
4    a <- 4.0
5    f x
6  printfn "%A" (testScope 2.0)
```

```
1  $ fsharpc --nologo dynamicScopeNFunction.fsx
2  $ mono dynamicScopeNFunction.exe
3  8.0
```

Here, the response is `8.0`, since the value of `a` changed before the function `f` was called.

6.8. Reference Cells

F# has a variation of mutable variables called *reference cells*. Reference cells have the built-in function *ref* and the operators "*!*" and "*:=*", where `ref` creates a reference variable, and the "*!*" and the "*:=*" operators respectively reads and writes its value. An example of using reference cells is given in Listing 6.48.

· reference cells
· `ref`
· `:=`

Listing 6.48 refCell.fsx:
Reference cells are variants of mutable variables.

```
1  let x = ref 0
2  printfn "%d" !x
3  x := !x + 1
4  printfn "%d" !x
```

```
1  $ fsharpc --nologo refCell.fsx && mono refCell.exe
2  0
3  1
```

Reference cells are different from mutable variables, since their content is allocated on *The Heap*. The Heap is a global data storage that is not destroyed when a function returns, which is in contrast to the *call stack*, also known as *The Stack*. The Stack maintains all the local data for a specific instance of a function call, see Section 13.2 for more details. As a consequence, when a reference cell is returned from a function, then it is the reference to the location on The Heap, which is returned as a value. Since this points outside the local data area of the function, this location is still valid after the function returns, and the variable stored there is accessible to the caller. This is illustrated in Figure 6.1

· The Heap
· call stack
· The Stack

Reference cells may cause *side-effects*, where variable changes are performed across independent scopes. Some side-effects are useful, e.g., the `printf` family changes the content of the screen, and the screen is outside the scope of the caller. Another example of a useful side-effect is a counter shown in Listing 6.49.

· side-effect

Figure 6.1.: A reference cell is a pointer to The Heap, and the content is not destroyed when its reference falls out of scope.

Listing 6.49 refEncapsulation.fsx:
An increment function with a local state using a reference cell.

```
1  let incr =
2    let counter = ref 0
3    fun () ->
4      counter := !counter + 1
5      !counter
6  printfn "%d" (incr ())
7  printfn "%d" (incr ())
8  printfn "%d" (incr ())
```

```
1  $ fsharpc --nologo refEncapsulation.fsx && mono
     refEncapsulation.exe
2  1
3  2
4  3
```

Here `incr` is an anonymous function with an internal state `counter`. At first glance, it may be surprising that `incr ()` does not return the value `1` at every call. The reason is that the value of the `incr` is the closure of the anonymous function `fun () -> counter := ...`, which is

$$\texttt{incr} : (\text{args}, \text{exp}, \text{env}) = \left((), \left(\begin{smallmatrix}\texttt{counter := !counter + 1}\\ \texttt{!counter}\end{smallmatrix}\right), (\texttt{counter} \rightarrow \texttt{ref 0})\right). \quad (6.5)$$

Thus, `counter` is only initiated once at the initial binding, while every call of `incr ()` updates its value on The Heap. Such a programming structure is called *encapsulation*,   · encapsulation since the `counter` state has been encapsulated in the anonymous function, and the only way to access it is by calling the same anonymous function. In general, it is advisable to **use encapsulation to hide implementation details irrelevant to the user of the**   Advice **code.**

The `incr` example in Listing 6.49 is an example of a useful side-effect. An example to be avoided is shown in Listing 6.50.

Listing 6.50 refSideEffect.fsx:
Intertwining independent scopes is typically a bad idea.

```
1  let updateFactor factor =
2    factor := 2
3
4  let multiplyWithFactor x =
5    let a = ref 1
6    updateFactor a
7    !a * x
8
9  printfn "%d" (multiplyWithFactor 3)
```

```
1  $ fsharpc --nologo refSideEffect.fsx && mono refSideEffect.exe
2  6
```

In the example, the function `updateFactor` changes a variable in the scope of the function `multiplyWithFactor`. The code style is prone to errors, since the computations are not local at the place of writing, i.e., in `multiplyWithFactor`, and if `updateFactor` were defined in a library, then the source code may not be available. Better style of programming is shown in Listing 6.51.

Listing 6.51 refWithoutSideEffect.fsx:
A solution similar to Listing 6.50 without side-effects.

```
1  let updateFactor () =
2    2
3
4  let multiplyWithFactor x =
5    let a = ref 1
6    a := updateFactor ()
7    !a * x
8
9  printfn "%d" (multiplyWithFactor 3)
```

```
1  $ fsharpc --nologo refWithoutSideEffect.fsx
2  $ mono refWithoutSideEffect.exe
3  6
```

Here, there can be no doubt in `multiplyWithFactor` that the value of `a` is changing. Side-effects do have their use, but should, in general, be avoided at almost all costs, and it is advised to **minimize the use of side effects**.

Advice

Reference cells give rise to an effect called *aliasing*, where two or more identifiers refer to the same data, as illustrated in Listing 6.52.

· aliasing

Listing 6.52 refCellAliasing.fsx:
Aliasing can cause surprising results and should be avoided.

```
1  let a = ref 1
2  let b = a
3  printfn "%d, %d" !a !b
4  b := 2
5  printfn "%d, %d" !a !b
```

```
1  $ fsharpc --nologo refCellAliasing.fsx && mono
     refCellAliasing.exe
2  1, 1
3  2, 2
```

Here, `a` is defined as a reference cell, and by defining `b` to be equal to `a`, we have created an alias. This can be very confusing since as the example shows, changing the value of `b` causes `a` to change as well. Aliasing is a variant of side-effects, and **aliasing should be avoided at all costs**. ·Advice

Since F# version 4.0, the compiler has automatically converted mutable variables to reference cells, where needed. E.g., Listing 6.49 can be rewritten using a mutable variable, as shown in Listing 6.53.

Listing 6.53 mutableEncapsulation.fsx:
Local mutable content can be indirectly accessed outside its scope.

```
1  let incr =
2    let mutable counter = 0
3    fun () ->
4      counter <- counter + 1
5      counter
6  printfn "%d" (incr ())
7  printfn "%d" (incr ())
8  printfn "%d" (incr ())
```

```
1  $ fsharpc --nologo mutableEncapsulation.fsx
2  $ mono mutableEncapsulation.exe
3  1
4  2
5  3
```

Reference cells are preferred over mutable variables for encapsulation, in order to avoid confusion.

6.9. Tuples

Tuples are a direct extension of constants. They are immutable and have neither concatenations nor indexing operations. Tuples are unions of immutable types and have the following syntax: ·tuple

> **Listing 6.54:   Tuples are list of expressions separated by commas.**

```
1  <expr>{, <expr>}
```

Tuples are identified by the ",," lexeme and often enclosed in parentheses, but that is not required. An example is a triple, also known as a 3-tuple, (2,true,"hello"). In interactive mode, the type of tuples is demonstrated in Listing 6.55.

> **Listing 6.55: Tuple types are products of sets.**

```
1  > let tp = (2, true, "hello")
2  - printfn "%A" tp;;
3  (2, true, "hello")
4  val tp : int * bool * string = (2, true, "hello")
5  val it : unit = ()
```

The values 2, true, and "hello" are *members*, and the number of elements of a tuple is · member
its *length*. From the response of F#, we see that the tuple is inferred to have the type · length
int * bool * string. The "*" denotes the Cartesian product between sets. Tuples can
be products of any types and follow the lexical scope rules like value and function bindings.
Notice also that a tuple may be printed as a single entity by the %A placeholder. In the
example we bound tp to the tuple. The opposite is also possible, as demonstrated in
Listing 6.56.

> **Listing 6.56: Definition of a tuple.**

```
1  > let deconstructNPrint tp =
2  -    let (a, b, c) = tp
3  -    printfn "tp = (%A, %A, %A)" a b c
4  -
5  - deconstructNPrint (2, true, "hello")
6  - deconstructNPrint (3.14, "Pi", 'p');;
7  tp = (2, true, "hello")
8  tp = (3.14, "Pi", 'p')
9  val deconstructNPrint : 'a * 'b * 'c -> unit
10 val it : unit = ()
```

In this example, a function is defined that takes 1 argument, a 3-tuple. If we wanted a function with 3 arguments, then the function binding should have been let deconstructNPrint a b c = ... The value binding let (a, b, c) = tp, binds a tuple with 3 named members to a value, thus deconstructing it in terms of its members. This is called pattern matching and will be discussed in further details in Chapter 15. Since we used the \%A placeholder in the printfn function, the function can be called with 3-tuples of different types. F# informs us that the tuple type is variable by writing 'a * 'b * 'c. The "'" notation means that the type can be decided at run-time, see Section 14.6 for more on variable types.

Pairs or 2-tuples are so common that F# includes two built-in functions, *fst* and *snd*, to · fst
extract the first and second element of a pair. This is demonstrated in Listing 6.57. · snd

Listing 6.57 pair.fsx:
Deconstruction of pairs with the built-in functions `fst` and `snd`.

```
1  let pair = ("first", "second")
2  printfn "fst(pair) = %s, snd(pair) = %s" (fst pair) (snd pair)
```

```
1  $ fsharpc --nologo pair.fsx && mono pair.exe
2  fst(pair) = first, snd(pair) = second
```

Tuples of equal lengths can be compared, and the comparison is defined similarly to string comparison. Tuples of equal length are compared element by element. E.g., `(1,2) = (1,3)` is false, while `(1,2) = (1,2)` is true. The "`<>`" operator is the boolean negation of the "`=`" operator. For the "`<`" , "`<=`", "`>`", and "`>=`" operators, the strings are ordered lexicographically, such that `('a', 'b', 'c') < ('a', 'b', 's') && ('a', 'b', 's') < ('c', 'o', 's')` is true, that is, the "`<`" operator on two tuples is true if and only if the left operand should come before the right when sorting alphabetically. See Listing 6.58 for an example.

Listing 6.58 tupleCompare.fsx:
Tuples comparison is similar to string comparison.

```
1  let lessThan (a, b, c) (d, e, f) =
2    if a <> d then a < d
3    elif b <> e then b < d
4    elif c <> f then c < f
5    else false
6
7  let printTest x y =
8    printfn "%A < %A is %b" x y (lessThan x y)
9
10  let a = ('a', 'b', 'c');
11  let b = ('d', 'e', 'f');
12  let c = ('a', 'b', 'b');
13  let d = ('a', 'b', 'd');
14  printTest a b
15  printTest a c
16  printTest a d
```

```
1  $ fsharpc --nologo tupleCompare.fsx && mono tupleCompare.exe
2  ('a', 'b', 'c') < ('d', 'e', 'f') is true
3  ('a', 'b', 'c') < ('a', 'b', 'b') is false
4  ('a', 'b', 'c') < ('a', 'b', 'd') is true
```

The algorithm for deciding the boolean value of `(a1, a2) < (b1, b2)` is as follows: we start by examining the first elements, and if `a1` and `b1` are different, then the result of `(a1, a2) < (b1, b2)` is equal to the result of `a1 < b1`. If `a1` and `b1` are equal, then we move on to the next letter and repeat the investigation. The "`<=`", "`>`", and "`>=`" operators are defined similarly.

Binding tuples to mutables does not make the tuple mutable. This is demonstrated in Listing 6.59.

> **Listing 6.59 tupleOfMutables.fsx:**
> **A mutable changes value, but the tuple defined by it does not refer to the new value.**
>
> ```
> let mutable a = 1
> let mutable b = 2
> let c = (a, b)
> printfn "%A, %A, %A" a b c
> a <- 3
> printfn "%A, %A, %A" a b c
> ```
> ```
> $ fsharpc --nologo tupleOfMutables.fsx && mono
>     tupleOfMutables.exe
> 1, 2, (1, 2)
> 3, 2, (1, 2)
> ```

However, it is possible to define a mutable variable of type tuple such that new tuple values can be assigned to it, as shown in Listing 6.60.

> **Listing 6.60 mutableTuple.fsx:**
> **A mutable tuple can be assigned a new value.**
>
> ```
> let mutable pair = 1,2
> printfn "%A" pair
> pair <- (3,4)
> printfn "%A" pair
> ```
> ```
> $ fsharpc --nologo mutableTuple.fsx && mono mutableTuple.exe
> (1, 2)
> (3, 4)
> ```

Mutable tuples are value types, meaning that binding to new names makes copies, not aliases, as demonstrated in Listing 6.61.

> **Listing 6.61 mutableTupleValue.fsx:**
> **A mutable tuple is a value type.**
>
> ```
> let mutable pair = 1,2
> let mutable aCopy = pair
> pair <- (3,4)
> printfn "%A %A" pair aCopy
> ```
> ```
> $ fsharpc --nologo mutableTupleValue.fsx && mono
>     mutableTupleValue.exe
> (3, 4) (1, 2)
> ```

The use of tuples shortens code and highlights semantic content at a higher level, e.g., instead of focusing on the elements, tuples focus on their union. While this may look elegant and short there is the risk of *obfuscation*, i.e., writing compact code that is difficult to read, where an unprepared reader of the code may not easily understand the computation   · obfuscation

nor appreciate its elegance without an accompanying explanation. Hence, **always keep an eye out for compact and concise ways to write code, but never at the expense of readability.**   Advice

7 | In-code Documentation

Documentation is a very important part of writing programs, since it is most unlikely that you will be writing really obvious code. Moreover, what seems obvious at the point of writing may be mystifying months later to the author and to others. Documentation serves several purposes:

1. Communicate what the code should be doing.

2. Highlight big insights essential for the code.

3. Highlight possible conflicts and/or areas where the code could be changed later.

The essential point is that coding is a journey in problem-solving, and proper documentation is an aid in understanding the solution and the journey that lead to it. Documentation is most often a mixture of in-code documentation and accompanying documents. Here, we will focus on in-code documentation which arguably causes problems in multi-language environments and run the risk of bloating code.

F# has two different syntaxes for comments. Comments can be block comments: [1]

Listing 7.1: Block comments.

```
1  (*<any text>*)
```

The comment text (`<any text>`) can be any text and is stilled parsed by F# as keywords and basic types, implying that `(* a comment (* in a comment *) *)` and `(* "*)" *)` are valid comments, while `(* " *)` is invalid.

Alternatively, comments may also be line comments,

Listing 7.2: Line comments.

```
1  //<any text>
```

where the comment text ends after the first newline.

The F# compiler has an option for generating *Extensible Markup Language* (*XML*) files from scripts using the C# documentation comments tags[2]. The XML documentation starts

· Extensible Markup Language
· XML

---

[1] Jon: **color of '*' is wrong.**

[2] For specification of C# documentations comments see ECMA-334: `http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-334.pdf`

Tag	Description
`<c>`	Set text in a code-font.
`<code>`	Set one or more lines in code-font.
`<example>`	Set as an example.
`<exception>`	Describe the exceptions a function can throw.
`<list>`	Create a list or table.
`<para>`	Set text as a paragraph.
`<param>`	Describe a parameter for a function or constructor.
`<paramref>`	Identify that a word is a parameter name.
`<permission>`	Document the accessibility of a member.
`<remarks>`	Further describe a function.
`<returns>`	Describe the return value of a function.
`<see>`	Set as link to other functions.
`<seealso>`	Generate a See Also entry.
`<summary>`	Main description of a function or value.
`<typeparam>`	Describe a type parameter for a generic type or method.
`<typeparamref>`	Identify that a word is a type parameter name.
`<value>`	Describe a value.

Table 7.1.: Recommended XML tags for documentation comments, from ECMA-334 3rd Edition, Annex E, Section 2.

with a triple-slash `///`, i.e., a lineComment and a slash, which serve as comments for the code construct that follows immediately after. XML consists of tags which always appear in pairs, e.g., the tag "tag" would look like `<tag> ... </tag>`. F# accept any tags, but recommends those listed in Table 7.1. If no tags are used, then it is automatically assumed to be a `<summary>`. An example of a documented script is shown in Listing 7.3. is:

Listing 7.3 commentExample.fsx:
Code with XML comments.

```
 1  /// The discriminant of a quadratic equation with parameters
        a, b, and c
 2  let discriminant a b c = b ** 2.0 - 4.0 * a * c
 3
 4  /// <summary>Find x when 0 = ax^2+bx+c.</summary>
 5  /// <remarks>Negative discriminants are not checked.</remarks>
 6  /// <example>
 7  ///   The following code:
 8  ///   <code>
 9  ///      let a = 1.0
10  ///      let b = 0.0
11  ///      let c = -1.0
12  ///      let xp = (solution a b c +1.0)
13  ///      printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b
        c xp
14  ///   </code>
15  ///   prints <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> to
        the console.
16  /// </example>
17  /// <param name="a">Quadratic coefficient.</param>
18  /// <param name="b">Linear coefficient.</param>
19  /// <param name="c">Constant coefficient.</param>
20  /// <param name="sgn">+1 or -1 determines the solution.</param>
21  /// <returns>The solution to x.</returns>
22  let solution a b c sgn =
23    let d = discriminant a b c
24    (-b + sgn * sqrt d) / (2.0 * a)
25
26  let a = 1.0
27  let b = 0.0
28  let c = -1.0
29  let xp = (solution a b c +1.0)
30  printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
```

```
 1  $ fsharpc --nologo commentExample.fsx && mono
        commentExample.exe
 2  0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 1.0
```

Mono's `fsharpc` command may be used to extract the comments into an XML file, as demonstrated in Listing 7.4.

Listing 7.4, Converting in-code comments to XML.

```
 1  $ fsharpc --doc:commentExample.xml commentExample.fsx
 2  F# Compiler for F# 4.0 (Open Source Edition)
 3  Freely distributed under the Apache 2.0 Open Source License
```

This results in an XML file with the content shown in Listing 7.5.

Listing 7.5, An XML file generated by `fsharpc`.

```
1  <?xml version="1.0" encoding="utf -8"?>
2  <doc>
3  <assembly><name>commentExample</name></assembly>
4  <members>
5  <member name="M:CommentExample.solution(System.Double,System.
     Double,System.Double,System.Double)">
6   <summary>Find x when 0 = ax^2+bx+c.</summary>
7   <remarks>Negative discriminants are not checked.</remarks>
8   <example>
9     The following code:
10    <code>
11      let a = 1.0
12      let b = 0.0
13      let c = -1.0
14      let xp = (solution a b c +1.0)
15      printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c
     xp
16    </code>
17    prints <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> to the
     console.
18  </example>
19  <param name="a">Quadratic coefficient.</param>
20  <param name="b">Linear coefficient.</param>
21  <param name="c">Constant coefficient.</param>
22  <param name="sgn">+1 or -1 determines the solution.</param>
23  <returns>The solution to x.</returns>
24 </member>
25 <member name="M:CommentExample.discriminant(System.Double,
     System.Double,System.Double)">
26 <summary>
27  The discriminant of a quadratic equation with parameters a, b,
      and c
28 </summary>
29 </member>
30 </members>
31 </doc>
```

The extracted XML is written in C# type by convention, since F# is part of the Mono and .Net framework that may be used by any of the languages using Assemblies. Besides the XML inserted in the script, the XML has added the `<?xml ...>` header, `<doc>`, `<assembly>`, `<members>`, and `<member>` tags. The header and the `<doc>` tag are standards for XML. The extracted XML is geared towards documenting big libraries of codes and thus highlights the structured programming organisation, see Chapters 9 and 20, and `<assembly>`, `<members>`, and `<member>` are indications for where the functions belong in the hierarchy. As an example, the prefix `M:CommentExample.` indicates that the method is in the namespace `commentExample`, which in this case is the name of the file. Furthermore, the function type `val` `solution : a:float -> b:float -> c:float -> sgn:float -> float` is in the XML documentation

M:CommentExample.solution(System.Double,System.Double,System.Double,System.Double),

which is the C# equivalent.

An accompanying program in the Mono suite is `mdoc`, whose primary use is to perform a

syntax analysis of an assembly and generate a scaffold XML structure for an accompanying document. With the `-i` flag, it is further possible to include the in-code comments as initial descriptions in the XML. The XML may be updated gracefully by `mdoc` as the code develops, without destroying manually entered documentation in the accompanying documentation. Finally, the XML may be exported to HTML.

The primary use of the `mdoc` command is to analyze compiled code and generate an empty XML structure with placeholders to describe functions, values, and variables. This structure can be updated and edited as the program develops, and the edited XML files can be exported to *Hyper Text Markup Language* (*HTML*) files and viewed in any browser. Using the console, all of this is accomplished by the procedure shown in Listing 7.6, and the result is shown in Figure 7.1.

· Hyper Text Markup Language
· HTML

Listing 7.6, Converting an XML file to HTML.

```
1  $ mdoc update -o commentExample -i commentExample.xml
     commentExample.exe
2  New Type: CommentExample
3  Member Added: public static double determinant (double a,
     double b, double c);
4  Member Added: public static double solution (double a, double
     b, double c, double sgn);
5  Member Added: public static double a { get; }
6  Member Added: public static double b { get; }
7  Member Added: public static double c { get; }
8  Member Added: public static double xp { get; }
9  Namespace Directory Created:
10 New Namespace File:
11 Members Added: 6, Members Deleted: 0
12 $ mdoc export-html -out commentExampleHTML commentExample
13 .CommentExample
```

A full description of how to use `mdoc` is found here[3].

---

[3]http://www.mono-project.com/docs/tools+libraries/tools/monodoc/
generating-documentation/

solution Method

Find x when 0 = ax^2+bx+c.

Syntax

[Microsoft.FSharp.Core.CompilationArgumentCounts(Mono.Cecil.CustomAttributeArgument[])]
public static double **solution** (double a, double b, double c, double sgn)

Parameters

a

Quadratic coefficient.

b

Linear coefficient.

c

Constant coefficient.

sgn

+1 or -1 determines the solution.

Returns

The solution to x.

Remarks

Negative discriminants are not checked.

Example

The following code:

```
Example
    let a = 1.0
    let b = 0.0
    let c = -1.0
    let xp = (solution a b c +1.0)
    printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
```

prints `0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7` to the console.

Requirements

Namespace:
Assembly: commentExample (in commentExample.dll)
Assembly Versions: 0.0.0.0

Figure 7.1.: Part of the HTML documentation as produced by `mdoc` and viewed in a browser.

8 | Controlling Program Flow

Non-recursive functions encapsulate code and allow for control of execution flow. That is, if a piece of code needs to be executed many times, then we can encapsulate it in the body of a function and call this function several times. In this chapter, we will look at more general control of flow via loops and conditional execution. Recursion is another mechanism for controlling flow, but this is deferred to Chapter 13.

8.1. While and For Loops

Many programming constructs need to be repeated, and F# contains many structures for repetition. A *while*-loop has the following syntax:

> **Listing 8.1: While loop.**
>
> ```
> 1  while <condition> do <expr> [done]
> ```

The *condition* `<condition>` is an expression that evaluates to true or false. A while-loop repeats the `<expr>` expression as long as the condition is true. Using lightweight syntax, the block following the *do* keyword up to and including the *done* keyword may be replaced by a newline and indentation.

The program in Listing 8.5 is an example of a while-loop which counts from 1 to 10.

> **Listing 8.2 countWhile.fsx:**
> **Count to 10 with a counter variable.**
>
> ```
> 1  let mutable i = 1 in while i <= 10 do printf "%d " i; i <- i +
>      1 done;
> 2  printf "\n"
> ```
> --
> ```
> 1  $ fsharpc --nologo countWhile.fsx && mono countWhile.exe
> 2  1 2 3 4 5 6 7 8 9 10
> ```

The variable `i` is customarily called the counter variable. The counting is done by performing the following computation: In line 1, the counter variable is first given an initial value of 1. Then execution enters the while-loop and examines the condition. Since $1 <= 10$, the condition is true, and execution enters the body of the loop. The body prints the value of the counter to the screen and increases the counter by 1. Then execution returns to the

top of the while-loop. Now the condition is $2 <= 10$, which is also true, and so execution enters the body and so on until the counter has reached the value 11, in which case the condition $11 <= 10$ is false, and execution continues in line 2.

In lightweight syntax, this would be as shown in Listing 8.3.

Listing 8.3 countWhileLightweight.fsx:
Count to 10 with a counter variable using lightweight syntax.

```
1  let mutable i = 1
2  while i <= 10 do
3    printf "%d " i
4    i <- i + 1
5  printf "\n"
```

```
1  $ fsharpc --nologo countWhileLightweight.fsx
2  $ mono countWhileLightweight.exe
3  1 2 3 4 5 6 7 8 9 10
```

Notice that although the expression following the condition is preceded with a `do` keyword, and `do <expr>` is a do binding, the keyword `do` is mandatory.

Counters are so common that a special syntax has been reserved for loops using counters. These are called *for*-loops. For-loops come in several variants, and here we will focus on the one using an explicit counter. Its syntax is:

· for

Listing 8.4:   For loop.

```
1  for <ident> = <firstExpr> to <lastExpr> do <bodyExpr> [done]
```

A for-loop initially binds the counter identifier `<ident>` to be the value `<firstExpr>`. Then execution enters the body, and `<bodyExpr>` is evaluated. Once done, the counter is increased, and execution evaluates `<bodyExpr>` once again. This is repeated as long as the counter is not greater than `<lastExpr>`. As for while-loops, when using lightweight syntax the block following the *do* keyword up to and including the *done* keyword may be replaced by a newline and indentation.

· do
· done

The counting example from Listing 8.2 using a `for`-loop is shown in Listing 8.5

Listing 8.5 count.fsx:
Counting from 1 to 10 using a `for`-loop.

```
1  for i = 1 to 10 do printf "%d " i done
2  printfn ""
```

```
1  $ fsharpc --nologo count.fsx && mono count.exe
2  1 2 3 4 5 6 7 8 9 10
```

As this interactive script demonstrates, the identifier `i` takes all the values between 1 and 10, but in spite of its changing state, it is not mutable. Note also that the return value of

Figure 8.1.: The Fibonacci spiral is an approximation of the golden spiral. Each square has side lengths of successive Fibonacci numbers, and the curve in each square is the circular arc with a radius of the square it is drawn in.

the `for` expression is "()", like the `printf` functions. The lightweight equivalent is shown in Listing 8.6.

Listing 8.6 countLightweight.fsx:
Counting from 1 to 10 using a for-loop using the lightweight syntax.

```
1  for i = 1 to 10 do
2    printf "%d " i
3  printfn ""
```

```
1  $ fsharpc --nologo countLightweight.fsx && mono
     countLightweight.exe
2  1 2 3 4 5 6 7 8 9 10
```

To further compare for- and while-loops, consider the following problem.

Problem 8.1

Write a program that calculates the n'th Fibonacci number.

Fibonacci numbers is a sequence of numbers starting with $1, 1$, and where the next number is calculated as the sum of the previous two. Hence the first ten numbers are: $1, 1, 2, 3, 5, 8, 13, 21, 34, 55$. Fibonacci numbers are related to Golden spirals shown in Figure 8.1. Often the sequence is extended with a preceding number 0, to be $0, 1, 1, 2, 3, \ldots$, which we will do here as well.

We could solve this problem with a `for`-loop, as shown in Listing 8.7.

Listing 8.7 fibFor.fsx:
The n'th Fibonacci number calculated using a for-loop.

```
1  let fib n =
2    let mutable pair = (0, 1)
3    for i = 2 to n do
4      pair <- (snd pair, (fst pair) + (snd pair))
5    snd pair
6
7  printfn "fib(1) = %d" (fib 1)
8  printfn "fib(2) = %d" (fib 2)
9  printfn "fib(3) = %d" (fib 3)
10 printfn "fib(10) = %d" (fib 10)
```

```
1  $ fsharpc --nologo fibFor.fsx && mono fibFor.exe
2  fib(1) = 1
3  fib(2) = 1
4  fib(3) = 2
5  fib(10) = 55
```

The basic idea of the solution is that if we are given the $(n-1)$'th and $(n-2)$'th numbers, the n'th number is trivial to compute. And assuming that fib(1) and fib(2) are given, then it is trivial to calculate fib(3). For fib(4), we only need fib(3) and fib(2), hence we may disregard fib(1). Thus, we realize that we can cyclicly update the previous, current, and next values by shifting values until we have reached the desired fib(n). This is implement in Listing 8.7 as the function `fib`, which takes an integer `n` as argument and returns the n'th Fibonacci number. The function does this iteratively using a `for`-loop, where `i` is the counter value, and `pair` is the pair of the $i-1$'th and i'th Fibonacci numbers. In the body of the loop, the i'th and $i+1$'th numbers are assigned to `pair`. The `for`-loop automatically updates `i` for next iteration. When $n < 2$ the body of the for-loop is not evaluated, and 1 is returned. This is of course wrong for $n < 1$, but we will ignore this for now.

Listing 8.8 shows a program similar to Listing 8.7 using a while-loop instead of for-loop.

> **Listing 8.8 fibWhile.fsx:**
> **The n'th Fibonacci number calculated using a while-loop.**
>
> ```fsharp
> let fib (n : int) : int =
>   let mutable pair = (0, 1)
>   let mutable i = 1
>   while i < n do
>     pair <- (snd pair, fst pair + snd pair)
>     i <- i + 1
>   snd pair
>
> printfn "fib(1) = %d" (fib 1)
> printfn "fib(2) = %d" (fib 2)
> printfn "fib(3) = %d" (fib 3)
> printfn "fib(10) = %d" (fib 10)
> ```
>
> ```
> $ fsharpc --nologo fibWhile.fsx && mono fibWhile.exe
> fib(1) = 1
> fib(2) = 1
> fib(3) = 2
> fib(10) = 55
> ```

The programs are almost identical. In this case, the `for`-loop is to be preferred, since more lines of code typically mean more chances of making a mistake. However, while-loops are somewhat easier to argue correctness about.

The correctness of `fib` in Listing 8.8 can be proven using a *loop invariant*. An *invariant* is a statement that is always true at a particular point in a program, and a loop invariant is a statement which is true at the beginning and end of a loop. In line 4 in Listing 8.8, we may state the invariant: The variable `pair` is the pair of the $i - 1$'th and i'th Fibonacci numbers. This is provable by induction:

· loop invariant
· invariant

Base case: Before entering the while loop, `i` is 1, `pair` is (0, 1). Thus, the invariant is true.

Induction step: Assuming that `pair` is the $i - 1$'th and i'th Fibonacci numbers, the body first assigns a new value to `pair` as the i'th and $i + 1$'th Fibonacci numbers, then increases i by one such that at the end of the loop the `pair` again contains the the $i - 1$'th and i'th Fibonacci numbers.

Thus, since our invariant is true for the first case, and any iteration following an iteration where the invariant is true, is also true, then it is true for all iterations.

Thus we know that the second value in `pair` holds the value of the i'th Fibonacci number, and since we further may prove that $i = n$ when line 7 is reached, then it is proven that `fib` returns the n'th Fibonacci number.

While-loops also allow for logical structures other than for-loops, such as the case when the number of iteration cannot easily be decided when entering the loop. As an example, consider a slight variation of the above problem, where we wish to find the largest Fibonacci number less or equal some number. A solution to this problem is shown in Listing 8.9.

Listing 8.9 fibWhileLargest.fsx:
Search for the largest Fibonacci number less than a specified number.

```
1  let largestFibLeq n =
2    let mutable pair = (0, 1)
3    while snd pair <= n do
4      pair <- (snd pair, fst pair + snd pair)
5    fst pair
6
7  for i = 1 to 10 do
8    printfn "largestFibLeq(%d) = %d" i (largestFibLeq i)
```

```
1  $ fsharpc --nologo fibWhileLargest.fsx && mono
     fibWhileLargest.exe
2  largestFibLeq(1) = 1
3  largestFibLeq(2) = 2
4  largestFibLeq(3) = 3
5  largestFibLeq(4) = 3
6  largestFibLeq(5) = 5
7  largestFibLeq(6) = 5
8  largestFibLeq(7) = 5
9  largestFibLeq(8) = 8
10 largestFibLeq(9) = 8
11 largestFibLeq(10) = 8
```

The strategy here is to iteratively calculate Fibonacci numbers until we've found one larger than the argument **n**, and then return the previous. This could not be calculated with a for-loop.

8.2. Conditional Expressions

Programs often contain code which should only be executed under certain conditions. This can be expressed with `if`-expressions, whose syntax is as follows.

· if
· then
· elif
· else

Listing 8.10:   Conditional expressions.

```
1  if <cond> then <expr> {elif <cond> then <expr>} [else <expr>]
```

The condition `<con>` is an expression resulting in a Boolean value, and there can be zero or more `elif` conditions, as indicated by `{}`. Each expression `<expr>` is called a *branch*, · branch and all branches must have the same type, such that regardless of which branch is chosen, the type of the result of the conditional expression is the same. Then the expression of the first if-branch, whose condition is true, is evaluate. If all conditions are false then the `else`-branch is evaluated. If no `else` expression is present, then "()" will be returned. See Listing 8.11 for a simple example.

> **Listing 8.11 condition.fsx:**
> **Conditions evaluate their branches depending on the value of the condition.**
>
> ```
> 1  if true then printfn "hi" else printfn "bye"
> 2  if false then printfn "hi" else printfn "bye"
> ```
> ```
> 1  $ fsharpc --nologo condition.fsx && mono condition.exe
> 2  hi
> 3  bye
> ```

The lightweight syntax allows for newlines entered everywhere, but indentation must be used to express scope.

To demonstrate conditional expressions, let us write a program which writes the sentence "I have n apple(s)", where the plural 's' is added appropriately for various n's. This is done in Listing 8.12, using the lightweight syntax.

> **Listing 8.12 conditionalLightweight.fsx:**
> **Using conditional expression to generate different strings.**
>
> ```
> 1   let applesIHave n =
> 2     if n < -1 then
> 3       "I owe " + (string -n) + " apples"
> 4     elif n < 0 then
> 5       "I owe " + (string -n) + " apple"
> 6     elif n < 1 then
> 7       "I have no apples"
> 8     elif n < 2 then
> 9       "I have 1 apple"
> 10    else
> 11      "I have " + (string n) + " apples"
> 12
> 13  printfn "%A" (applesIHave -3)
> 14  printfn "%A" (applesIHave -1)
> 15  printfn "%A" (applesIHave 0)
> 16  printfn "%A" (applesIHave 1)
> 17  printfn "%A" (applesIHave 2)
> 18  printfn "%A" (applesIHave 10)
> ```
> ```
> 1  $ fsharpc --nologo conditionalLightWeight.fsx
> 2  $ mono conditionalLightWeight.exe
> 3  "I owe 3 apples"
> 4  "I owe 1 apple"
> 5  "I have no apples"
> 6  "I have 1 apple"
> 7  "I have 2 apples"
> 8  "I have 10 apples"
> ```

The sentence structure and its variants give rise to a more compact solution, since the language to be returned to the user is a variant of "I have/owe no/number apple(s)", i.e., certain conditions determine whether the sentence should use "have" and "owe" and so forth. So, we could instead make decisions on each of these sentence parts, and then built the final sentence from its parts. This is accomplished in the following example:

Listing 8.13 conditionalLightweightAlt.fsx:
Using sentence parts to construct the final sentence.

```
1  let applesIHave n =
2    let haveOrOwe = if n < 0 then "owe" else "have"
3    let pluralS = if (n = 0) || (abs n) > 1 then "s" else ""
4    let number = if n = 0 then "no" else (string (abs n))
5
6    "I " + haveOrOwe + " " + number + " apple" + pluralS
7
8  printfn "%A" (applesIHave -3)
9  printfn "%A" (applesIHave -1)
10 printfn "%A" (applesIHave 0)
11 printfn "%A" (applesIHave 1)
12 printfn "%A" (applesIHave 2)
13 printfn "%A" (applesIHave 10)
```

```
1  $ fsharpc --nologo conditionalLightWeightAlt.fsx
2  $ mono conditionalLightWeightAlt.exe
3  "I owe 3 apples"
4  "I owe 1 apple"
5  "I have no apples"
6  "I have 1 apple"
7  "I have 2 apples"
8  "I have 10 apples"
```

While arguably shorter, this solution is also denser, and most likely more difficult to debug and maintain.

Note that both `elif` and `else` branches are optional, which may cause problems. For example, both `let a = if true then 3` and `let a = if true then 3 elif false then 4` will be invalid, since F# is not smart enough to realize that the type of the expression is uniquely determined. Instead, F# looks for the `else` to ensure all cases have been covered, and that `a` always will be given a unique value of the same type regardless of the branch taken in the conditional statement. Hence, `let a = if true then 3 else 4` is the only valid expression of the 3. In practice, F# assumes that the omitted branch returns "()", and thus it is fine to say `let a = if true then ()` and `if true then printfn "hej"`. Nevertheless, it is good practice in F# to always include an `else` branch.

8.3. Programming Intermezzo: Automatic Conversion of Decimal to Binary Numbers

Using loops and conditional expressions, we are now able to solve the following problem:

Problem 8.2

Given an integer on decimal form, write its equivalent value on the binary form.

To solve this problem, consider odd numbers: They all have the property that the least significant bit is 1, e.g., $1_2 = 1, 101_2 = 5$, in contrast to even numbers such as $110_2 = 6$. Division by 2 is equal to right-shifting by 1, e.g., $1_2/2 = 0.1_2 = 0.5, 101_2/2 = 10.1_2 =$

$2.5, 110_2/2 = 11_2 = 3$. Thus, through dividing by 2 and checking the remainder, we may sequentially read off the least significant bit. This leads to the algorithm shown in Listing 8.14.

Listing 8.14 dec2bin.fsx:
Using integer division and remainder to write any positive integer in binary form.

```
1  let dec2bin n =
2    if n < 0 then
3      "Illegal value"
4    elif n = 0 then
5      "0b0"
6    else
7      let mutable v = n
8      let mutable str = ""
9      while v > 0 do
10       str <- (string (v % 2)) + str
11       v <- v / 2
12     "0b" + str
13
14
15 printfn "%4d -> %s" -1 (dec2bin -1)
16 printfn "%4d -> %s" 0 (dec2bin 0)
17 for i = 0 to 3 do
18   printfn "%4d -> %s" (pown 10 i) (dec2bin (pown 10 i))
```

```
1  $ fsharpc --nologo dec2bin.fsx && mono dec2bin.exe
2    -1 -> Illegal value
3     0 -> 0b0
4     1 -> 0b1
5    10 -> 0b1010
6   100 -> 0b1100100
7  1000 -> 0b1111101000
```

In the code, the states `v` and `str` are iteratively updated until `str` finally contains the desired solution.

To prove that Listing 8.14 calculates the correct sequence, we use induction. First we realize that for $v < 1$, the while-loop is skipped, and the result is trivially true. We will concentrate on line 9 in Listing 8.14 and will prove the following loop invariant: The string `str` contains all the bits of `n` to the right of the bit pattern remaining in variable `v`.

Base case $n = 000\ldots000x$: If n only uses the lowest bit, then $n = 0$ or $n = 1$. If $n = 0$, then it is trivially correct. Considering the case $n = 1$: Before entering into the loop, `v` is 1, and `str` is the empty string, so the invariant is true. The condition of the while-loop is $1 > 0$, so execution enters the loop. Since integer division of 1 by 2 gives 0 with remainder 1, `str` is set to `"1"` and `v` to 0. Now we reexamine the while-loop's condition, $0 > 0$, which is false, so we exit the loop. At this point, `v` is 0 and `str` is `"1"`, so all bits have been shifted from `n` to `str`, and none are left in `v`. Thus the invariant is true. Finally, the program returns `"0b1"`.

Induction step: Consider the case of $n > 1$, and assume that the invariant is true when entering the loop, i.e., that m bits already have been shifted to `str` and that $n > 2^m$.

In this case, `v` contains the remaining bits of `n`, which is the integer division `v = n / 2**m`. Since $n > 2^m$, `v` is non-zero, and the loop conditions is true, so we enter the loop body. In the loop body we concatenate the rightmost bit of `v` to the left of `str` using `v % 2`, and right-shift `v` one bit to the right with `v <- v / 2`. Thus, when returning to the condition the invariant is true, since the right-most bit in `v` has been shifted to `str`. This continues until all bits have been shifted to `str` and `v = 0`, in which case the loop terminates, and `"0b"+str` is returned.

Thus we have proven that `dec2bin` correctly converts integers to strings representing binary numbers.

9 | Organising Code in Libraries and Application Programs

In this chapter, we will focus on a number of ways to make the code available as *library* · library
functions in F#. A library is a collection of types, values, and functions that an application
program can use. A library does not perform calculations on its own.

F# includes several programming structures to organize code in libraries: Modules, names-
paces, and classes. In this chapter, we will describe modules and namespaces. Classes will
be described in detail in Chapter 20.

9.1. Modules

An F# *module*, not to be confused with a Common Language Infrastructure module (see · module
Appendix D), is a programming structure used to organize type declarations, values, func-
tions, etc.

Every implementation and script file in F# implicitly defines a module, and the module
name is given by the filename. Consider the script file `Meta.fsx` shown in Listing 9.1.[1]

> **Listing 9.1 Meta.fsx:**
> **A script file defining the `apply` function.**
>
> ```
> 1  type floatFunction = float -> float -> float
> 2  let apply (f : floatFunction) (x : float) (y : float) : float
>      = f x y
> ```

Here, we have implicitly defined a module with the name `Meta`. Another script file may
now use this function, which is accessed using the ".". notation, i.e., `Meta.apply` will refer
to this function in other programs. An application program could be as the one shown in
Listing 9.3.

---

[1]Jon: **Type definitions have not been introduced at this point!**

Listing 9.2 MetaApp.fsx:
Defining a script calling the module.

```
1  let add : Meta.floatFunction = fun x y -> x + y
2  let result = Meta.apply add 3.0 4.0
3  printfn "3.0 + 4.0 = %A" result
```

In the exmple above, we have explicitly used the module's type definition for illustration purposes. A shorter and possibly simpler program would have been to define `add` as `let add x y = x + y`, since F#'s type system will infer the implied type. However, **explicit definitions of types is recommended for readability.** Hence, an alternative to the above example's use of anonymous functions is: `let add (x: float) (y: float) : float = x + y`. To compile the module and the application program, we write as demonstrated in Listing 9.3.

Advice

Listing 9.3: Compiling both the module and the application code. Note that file order matters when compiling several files.

```
1  $ fsharpc --nologo Meta.fsx MetaApp.fsx && mono MetaApp.exe
2  3.0 + 4.0 = 7.0
```

Since the F# compiler reads through the files once, the order of the filenames in the compile command is very important. Hence, the script containing the module and function definitions must be to the left of the script containing their use. Notice also that if not otherwise specified, the F# compiler produces an `.exe` file derived from the last filename in the list of filenames.

We may also explicitly define the module name using the *module* with the following syntax,   · `module`

Listing 9.4:   Outer module.

```
1  module <ident>
2  <script>
```

Here, the identifier `<ident>` is a name not necessarily related to the filename, and the script `<script>` is an expression. An example is given in Listing 9.20.

Listing 9.5 MetaExplicit.fsx:
Explicit definition of the outermost module.

```
1  module Meta
2  type floatFunction = float -> float -> float
3  let apply (f : floatFunction) (x : float) (y : float) : float
      = f x y
```

Since we have created a new file, where the module `Meta` is explicitly defined, we can use the same application program. This is demonstrated in Listing 9.6.

Listing 9.6: Changing the module definition to explicit naming has no effect on the application nor the compile command.

```
1  $ fsharpc --nologo MetaExplicit.fsx MetaApp.fsx && mono
     MetaApp.exe
2  3.0 + 4.0 = 7.0
```

Since `MetaExplicit.fsx` explicitly defines the module name, `apply` is not available to an application program as `MetaExplicit.apply`. **It is recommended that module names are defined explicitly, since filenames may change due to external conditions.** In other words, filenames are typically set from the perspective of the filesystem. The user may choose to change names to suit a filesystem structure, or different platforms may impose different file naming conventions. Thus, direct linking of filenames with the internal workings of a program is a needless complication of structure.

Advice

The definitions inside a module may be accessed directly from an application program, omitting the ".·"-notation, by use of the *open* keyword,

·open

Listing 9.7:   Open module.

```
1  open <ident>
```

We can modify `MetaApp.fsx`, as shown in Listing 9.9.

Listing 9.8 MetaAppWOpen.fsx:
Avoiding the ".·"-notation by the open keyword.

```
1  open Meta
2  let add : floatFunction = fun x y -> x + y
3  let result = apply add 3.0 4.0
4  printfn "3.0 + 4.0 = %A" result
```

In this case, the namespace of our previously defined module is included into the scope of the application functions, and its types, values, functions, etc. can be used directly, as shown in Listing 9.9.

Listing 9.9: How the application program opens the module has no effect on the module code nor compile command.

```
1  $ fsharpc --nologo MetaExplicit.fsx MetaAppWOpen.fsx && mono
     MetaAppWOpen.exe
2  3.0 + 4.0 = 7.0
```

The open-keyword should be used sparingly, since including a library's definitions into the application scope can cause surprising naming conflicts, because the user of a library typically has no knowledge of the inner workings of the library. E.g., the user may accidentally use code defined in the library, but with different type and functionality than intended, which the type system will use to deduce types in the application program, and therefore will either give syntax or runtime errors that are difficult to understand. This problem is known as *namespace pollution*, and for clarity, **it is recommended to use the open-**

·namespace pollution

Advice

keyword sparingly. Note that for historical reasons, the phrase 'namespace pollution' is used to cover pollution both due to modules and namespaces.

Modules may also be nested, in which case the nested definitions must use the "="-sign and must be appropriately indented.

Listing 9.10:   Nested modules.

```
module <ident> = <script>
```

In lightweight syntax, a newline may be entered before the script `<script>`, and the script must be indented. An example is shown in Listing 9.11.

Listing 9.11 nestedModules.fsx:
Modules may be nested.

```
module Utilities
let PI = 3.1415
module Meta =
  type floatFunction = float -> float -> float
  let apply (f : floatFunction) (x : float) (y : float) :
    float = f x y
module MathFcts =
  let add : Meta.floatFunction = fun x y -> x + y
```

In this case, `Meta` and `MathFcts` are defined at the same level and said to be siblings, while `Utilities` is defined at a higher level. In this relation, the former two are said to be the children of the latter. Note that the nesting respects the lexical scope rules, such that the constant `PI` is directly accessible in both modules `Meta` and `MathFcts`, as is the module `Meta` in `MathFcts`, but not `MathFcts` in `Meta`. The "."-notation is reused to index deeper into the module hierarchy, as the example in Listing 9.12 shows.

Listing 9.12 nestedModulesApp.fsx:
Applications using nested modules require additional usage of the "." notation to navigate the nesting tree.

```
let add : Utilities.Meta.floatFunction = fun x y -> x + y
let result = Utilities.Meta.apply Utilities.MathFcts.add 3.0
  Utilities.PI
printfn "3.0 + 4.0 = %A" result
```

Modules can be recursive using the `rec`-keyword, meaning that in our example we can make the outer module recursive, as demonstrated in Listing 9.13.[2]

·`rec`

---

[2]Jon: **Dependence on version 4.1 and higher.**

> **Listing 9.13 nestedRecModules.fsx:**
> **Mutual dependence on nested modules requires the `rec` keyword in the module definition.**
>
> ```
> 1  module rec Utilities
> 2    module Meta =
> 3      type floatFunction = float -> float -> float
> 4      let apply (f : floatFunction) (x : float) (y : float) :
>      float = f x y
> 5    module MathFcts =
> 6      let add : Meta.floatFunction = fun x y -> x + y
> ```

The consequence is that the modules `Meta` and `MathFcts` are accessible in both modules, but compilation will now give a warning since soundness of the code will first be checked at runtime. In general, it is advised to **avoid programming constructions whose** Advice **validity cannot be checked at compile-time.**

9.2. Namespaces

An alternative way to structure code in modules is to use a *namespace*, which can only · namespace hold modules and type declarations and only works in compiled mode. Namespaces are defined as explicitly defined outer modules, using the *namespace* keyword in accordance · `namespace` with the following syntax.

> **Listing 9.14: Namespace.**
>
> ```
> 1  namespace <ident>
> 2  <script>
> ```

An example is given in Listing 9.15.

> **Listing 9.15 namespace.fsx:**
> **Defining a namespace is similar to explicitly named modules.**
>
> ```
> 1  namespace Utilities
> 2  type floatFunction = float -> float -> float
> 3  module Meta =
> 4    let apply (f : floatFunction) (x : float) (y : float) :
>    float = f x y
> ```

Notice that when organizing code in a namespace, the first line of the file, other than comments and compiler directives, must be the one starting with `namespace`.

As for modules, the content of a namespace is accessed using the "`.`" notation, as demonstrated in Listing 9.17.

Listing 9.16 namespaceApp.fsx:
The ".".-notation lets the application program access functions and types in a namespace.

```
1  let add : Utilities.floatFunction = fun x y -> x + y
2  let result = Utilities.Meta.apply add 3.0 4.0
3  printfn "3.0 + 4.0 = %A" result
```

Likewise, the compilation is performed in the same way as for modules, see Listing 9.17.

Listing 9.17: Compilation of files including namespace definitions uses the same procedure as modules.

```
1  $ fsharpc --nologo namespace.fsx namespaceApp.fsx && mono
     namespaceApp.exe
2  3.0 + 4.0 = 7.0
```

Hence, from an application point of view, it is not immediately possible to see that `Utilities` is defined as a namespace and not a module. However, in contrast to modules, namespaces may span several files. E.g., we may add a third file extending the `Utilities` namespace with the `MathFcts` module, as demonstrated in Listing 9.18.

Listing 9.18 namespaceExtension.fsx:
Namespaces may span several files. Here is shown an extra file which extends the `Utilities` namespace.

```
1  namespace Utilities
2  module MathFcts =
3    let add : floatFunction = fun x y -> x + y
```

To compile, we now need to include all three files in the right order, see Listing 9.19.

Listing 9.19: Compilation of namespaces defined in several files requires careful consideration of order, since the compiler reads once and only once through the files in the order they are given.

```
1  $ fsharpc --nologo namespace.fsx namespaceExtension.fsx
     namespaceApp.fsx && mono namespaceApp.exe
2  3.0 + 4.0 = 7.0
```

The order matters, since `namespaceExtension.fsx` relies on the definition of `floatFunction` in the file `namespace.fsx`. You can use extensions to extend existing namespaces included with the F# compiler.[34]

Namespaces may also be nested. In contrast to modules, nesting is defined using the "." notation. That is, to create a child namespace `more` of `Utilities`, we must use initially write `namespace Utilities.more`. Indentation is ignored in the `namespace` line, thus

---

[3] Jon: **Something about intrinsic and optional extension** https://docs.microsoft.com/en-us/dotnet/fsharp/language-reference/type-extensions.

[4] Jon: **Perhaps something about the global namespace `global`.**

left-most indentation is almost always used. Namespaces follow lexical scope rules, and identically to modules, namespaces containing mutually dependent children can be declared using the `rec` keyword, e.g., `namespace rec Utilities`.

9.3. Compiled Libraries

Libraries may be distributed in compiled form as `.dll` files. This saves the user from having to recompile a possibly large library every time library functions needs to be compiled with an application program. In order to produce a library file from `MetaExplicitModuleDefinition.fsx` and then compile an application program, we first use the compiler's `-a` option to produce the `.dll`. A demonstration is given in Listing 9.20.

Listing 9.20: A stand-alone `.dll` file is created and used with special compile commands.

```
1   $ fsharpc --nologo -a MetaExplicit.fsx
```

This produces the file `MetaExplicit.dll`, which may be linked to an application by using the `-r` option during compilation, see Listing 9.21.[5]

Listing 9.21: The library is linked to an application during compilation to produce runnable code.

```
1   $ fsharpc --nologo -r MetaExplicit.dll MetaApp.fsx && mono
      MetaApp.exe
2   3.0 + 4.0 = 7.0
```

A library can be the result of compiling a number of files into a single `.dll` file. `.dll`-files may be loaded dynamically in script files (`.fsx`-files) by using the **#r** *directive*, as   · **#r** directive illustrated in Listing 9.23.

Listing 9.22 MetaHashApp.fsx:
The `.dll` file may be loaded dynamically in `.fsx` script files and in interactive mode. Nevertheless, this usage is not recommended.

```
1   #r "MetaExplicit.dll"
2   let add : Meta.floatFunction = fun x y -> x + y
3   let result = Meta.apply add 3.0 4.0
4   printfn "3.0 + 4.0 = %A" result
```

We may now omit the explicit mentioning of the library when compiling, as shown in Listing 9.23.

---

[5]Jon: **This is the MacOS option standard, Windows is slightly different.**

> **Listing 9.23: When using the #r directive, then the `.dll` file need not be explicitly included in the list of files to be compiled.**
>
> ```
> 1  $ fsharpc --nologo MetaHashApp.fsx && mono MetaHashApp.exe
> 2  3.0 + 4.0 = 7.0
> ```

The `#r` directive is also used to include a library in interactive mode. However, for the code to be compiled, the use of the `#r` directive requires that the filesystem path to the library is coded inside the script. As for module names, direct linking of filenames with the internal workings of a program is a needless complication of structure, and **it is recommended not to rely on the use of the #r directive.** ·Advice

In the above listings we have compiled *script files* into libraries. However, F# has reserved the `.fs` filename suffix for library files, and such files are called *implementation files*. In contrast to script files, implementation files do not support the `#r` directive. When compiling a list of implementation and script files, all but the last file must explicitly define a module or a namespace. ·script file ·implementation file

Both script and implementation files may be augmented with *signature files*. A signature file contains no implementation, only type definitions. Signature files offer three distinct features: ·signature file

1. Signature files can be used as part of the documentation of code, since type information is of paramount importance for an application programmer to use a library.

2. Signature files may be written before the implementation file. This allows for a higher-level programming design that focuses on *which* functions should be included and *how* they can be composed.

3. Signature files allow for access control. Most importantly, if a type definition is not available in the signature file, then it is not available to the application program. Such definitions are private and can only be used internally in the library code. More fine-grained control related to classes is available and will be discussed in Chapter 20.

Signature files can be generated automatically using the `--sig:<filename>` compiler directive. To demonstrate this feature, we will first move the definition of `add` to the implementation file, see Listing 9.28.

> **Listing 9.24 MetaWAdd.fs:**
> **An implementation file including the `add` function.**
>
> ```
> 1  module Meta
> 2  type floatFunction = float -> float -> float
> 3  let apply (f : floatFunction) (x : float) (y : float) : float
>      = f x y
> 4  let add (x : float) (y : float) : float = x + y
> ```

A signature file may be automatically generated, as shown in Listing 9.25.

Listing 9.25: Automatic generation of a signature file at compile time.

```
1  $ fsharpc --nologo --sig:MetaWAdd.fsi MetaWAdd.fs
2
3  MetaWAdd.fs(4,48): warning FS0988: Main module of program is
      empty: nothing will happen when it is run
```

The warning can safely be ignored, since at this point it is not our intention to produce runnable code. The above listing has generated the signature file in Listing 9.26.

Listing 9.26 MetaWAdd.fsi:
An automatically generated signature file from `MetaWAdd.fs`.

```
1  module Meta
2  type floatFunction = float -> float -> float
3  val apply : f:floatFunction -> x:float -> y:float -> float
4  val add : x:float -> y:float -> float
```

We can generate a library using the automatically generated signature file by writing `fsharpc -a MetaWAdd.fsi MetaWAdd.fs`, which is identical to compiling the `.dll` file without the signature file. However, if we remove, e.g., the type definition for `add` in the signature file, then this function becomes private to the module and cannot be accessed outside. Hence, using the signature file in Listing 9.27, and recompiling the `.dll` with Listing 9.24 does not generates errors.

Listing 9.27 MetaWAddRemoved.fsi:
Removing the type defintion for `add` **from** `MetaWAdd.fsi`.

```
1  module Meta
2  type floatFunction = float -> float -> float
3  val apply : f:floatFunction -> x:float -> y:float -> float
```

Listing 9.28: Automatic generation of a signature file at compile time.

```
1  $ fsharpc --nologo -a MetaWAddRemoved.fsi MetaWAdd.fs
```

However, when using the newly created `MetaWAdd.dll` with an application that does not itself supply a definition of `add`, we get a syntax error, since `add` now is inaccessible to the application program. This is demonstrated in Listing 9.29 and 9.30.

Listing 9.29 MetaWOAddApp.fsx:
A version of Listing 9.3 without a definition of `add`.

```
1  let result = Meta.apply add 3.0 4.0
2  printfn "3.0 + 4.0 = %A" result
```

Listing 9.30: Automatic generation of a signature file at compile time.

```
$ fsharpc --nologo -r MetaWAdd.dll MetaWOAddApp.fsx

MetaWOAddApp.fsx(1,25): error FS0039: The value or constructor
    'add' is not defined.
```

10 | Testing Programs

A software bug is an error in a computer program that causes it to produce an incorrect result or behave in an unintended manner. The term 'bug' was used by Thomas Edison in 1878[1][2], but made popular in computer science by Grace Hopper, who found a moth interfering with the electronic circuits of the Harward Mark II electromechanical computer and coined the term *bug* for errors in computer programs. The original bug is shown in Figure 10.1. Software is everywhere, and errors therein have a huge economic impact on our society and can threaten lives[3].

· bug

The ISO/IEC organizations have developed standards for software testing[4]. To illustrate basic concepts of software quality, consider a hypothetical route planning system. Essential factors of its quality are:

· functionality

Functionality: Does the software compile and run without internal errors. Does it solve the problem it was intended to solve? E.g., does the route planning software find a suitable route from point a to b?

· reliability

Reliability: Does the software work reliably over time? E.g., does the route planning software work when there are internet dropouts?

· usability

Usability: Is the software easy and intuitive to use by humans? E.g., is it easy to enter addresses and alternative routes in the software's interface?
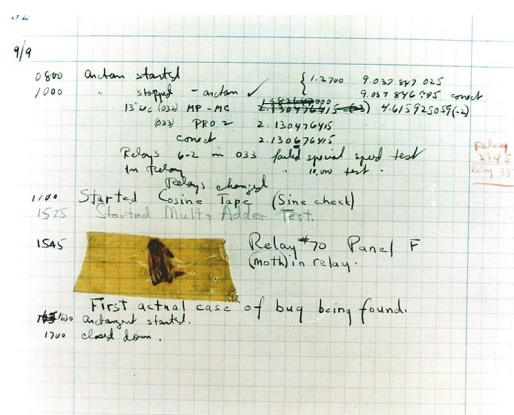
· efficiency

Figure 10.1.: The first computer bug, caught by Grace Hopper, U.S. Naval Historical Center Online Library Photograph NH 96566-KN.

---

[1] https://en.wikipedia.org/wiki/Software_bug
[2] http://edison.rutgers.edu/NamesSearch/DocImage.php3?DocId=LB003487
[3] https://en.wikipedia.org/wiki/List_of_software_bugs
[4] ISO/IEC 9126, International standard for the evaluation of software quality, December 19, 1991, later replaced by ISO/IEC 25010:2011

Efficiency: How many computer and human resources does the software require? E.g., does it take milliseconds or hours to find a requested route? Can the software run on a mobile platform with limited computer speed and memory?

Maintainability: In case of the discovery of new bugs, is it easy to test and correct the software? Is it easy to extend the software with new functionality? E.g., is it easy to update the map with updated roadmaps and new information? Can the system be improved to work both for car drivers and bicyclists?

· maintainability

Portability: Is it easy to port the software to new systems such as new server architecture and screen sizes? E.g., if the routing software originally was written for IOS devices, will it be easy to port to Android systems?

· portability

The above-mentioned concepts are ordered based on the requirements of the system. Functionality and reliability are perhaps the most important concepts, since if the software does not solve the specified problem, then the software design process has failed. However, many times the problem definition will evolve along with the software development process. But as a bare minimum, the software should run without internal errors and not crash under a well-defined set of circumstances. Furthermore, it is often the case that software designed for the general public requires a lot of attention to the usability of the software, since in many cases non-experts are expected to be able to use the software with little or no prior training. On the other hand, software used internally in companies will be used by a small number of people who become experts in using the software, and it is often less important that the software is easy to understand by non-experts. An example is text processing software like Microsoft Word versus Gnu Emacs and LaTeX. Word is designed to be used by non-experts for small documents such as letters and notes and relies heavily on interfacing with the system using click-interaction. On the other hand, Emacs and LaTeX are for experts for longer and professionally typeset documents and relies heavily on keyboard shortcuts and text-codes for typesetting document entities.

The purpose of *software testing* is to find bugs. When errors are found, then we engage in *debugging*, which is the process of diagnosing and correcting bugs. Once we have a failed software test, i.e., one that does not find any bugs, then we have strengthened our belief in the software, but it is important to note that software testing and debugging rarely removes all bugs, and with each correction or change of software there is a fair risk new bugs being introduced. It is not exceptional that the testing-software is as large as the software being tested.

· software testing
· debugging

In this chapter, we will focus on two approaches to software testing which emphasize functionality: *white-box* and *black-box testing*. An important concept in this context is *unit testing*, where the program is considered in smaller pieces, called units, and for which accompanying programs for testing can be made which test these units automatically. Black-box testing considers the problem formulation and the program interface, and can typically be written early in the software design phase. In contrast, white-box testing considers the program text, and thus requires the program to be available. Thus, there is a tendency for black-box test programs to be more stable, while white-box testing typically is developed incrementally alongside the software development.

· white-box testing
· black-box testing
· unit testing

To illustrate software testing, we'll start with a problem:

> **Problem 10.1**
>
> Given any date in the Gregorian calendar, calculate the day of the week.

Facts about dates in the Gregorian calendar are:

- Combinations of dates and weekdays repeat themselves every 400 years.

- The typical length of the months January, February, . . . follow the knuckle rule, i.e., January belongs to the index knuckle, February to the space between the index and the middle finger, and August restarts or starts on the other hand. All knuckle months have 31 days, all spacing months have 30 days except February, which has 29 days on leap years and 28 days all other years.

- A leap year is a multiple of 4, except if it is also a multiple of 100 but not of 400.

Many solutions to the problem have been discovered, and here we will base our program on Gauss' method, which is based on integer division and calculates the weekday of the 1st of January of a given year. For any other date, we will count our way through the weeks from the previous 1st of January. The algorithm relies on an enumeration of weekdays starting with Sunday = 0, Monday = 1, . . . , and Saturday = 6. Our proposed solution is shown in Listing 10.1.[5]

Listing 10.1 date2Day.fsx:
A function that can calculate day-of-week from any date in the Gregorian calendar.

```
1   let januaryFirstDay (y : int) =
2     let a = (y - 1) % 4
3     let b = (y - 1) % 100
4     let c = (y - 1) % 400
5     (1 + 5 * a + 4 * b + 6 * c) % 7
6
7   let rec sum (lst : int list) j =
8     if 0 <= j && j < lst.Length then
9       lst.[0] + sum lst.[1..] (j - 1)
10    else
11      0
12
13  let date2Day d m y =
14    let dayPrefix =
15      ["Sun"; "Mon"; "Tues"; "Wednes"; "Thurs"; "Fri"; "Satur"]
16    let feb = if (y % 4 = 0) && ((y % 100 <> 0) || (y % 400 =
      0)) then 29 else 28
17    let daysInMonth = [31; feb; 31; 30; 31; 30; 31; 31; 30; 31;
      30; 31]
18    let dayOne = januaryFirstDay y
19    let daysSince = (sum daysInMonth (m - 2)) + d - 1
20    let weekday = (dayOne + daysSince) % 7;
21    dayPrefix.[weekday] + "day"
```

10.1. White-box Testing

White-box testing considers the text of a program. The degree to which the text of the program is covered in the test is called the *coverage*. Since our program is small, we have the opportunity to ensure that all functions are called at least once, which is called *function*

· white-box testing

· coverage

---

[5]Jon: **This example relies on lists, which has not been introduced yet.**

coverage, and we will also be able to test every branching in the program, which is called *branching coverage*. If both are fulfilled, we say that we have *statement coverage*. The procedure is as follows:

· function coverage

· branching coverage

· statement coverage

1. Decide which units to test: The program shown in Listing 10.1 has 3 functions, and we will consider these each as a unit, but we might as well just have chosen `date2Day` as a single unit. The important part is that the union of units must cover the whole program text, and since `date2Day` calls both `januaryFirstDay` and `sum`, designing test cases for the latter two is superfluous. However, we may have to do this anyway when debugging, and we may choose at a later point to use these functions separately, and in both cases, we will be able to reuse the testing of the smaller units.

2. Identify branching points: The function `januaryFirstDay` has no branching function, `sum` has one, and depending on the input values, two paths through the code may be used, and `date2Day` has one where the number of days in February is decided. Note that in order to test this, our test-date must be March 1 or later. In this example, there are only examples of `if`-branch points, but they may as well be loops and pattern matching expressions. In the Listing 10.2 it is shown that the branch points have been given a comment and a number.

> **Listing 10.2 date2DayAnnotated.fsx:**
> **In white-box testing, the branch points are identified.**
>
> ```
> 1  // Unit: januaryFirstDay
> 2  let januaryFirstDay (y : int) =
> 3    let a = (y - 1) % 4
> 4    let b = (y - 1) % 100
> 5    let c = (y - 1) % 400
> 6    (1 + 5 * a + 4 * b + 6 * c) % 7
> 7
> 8  // Unit: sum
> 9  let rec sum (lst : int list) j =
> 10    (* WB: 1 *)
> 11    if 0 <= j && j < lst.Length then
> 12      lst.[0] + sum lst.[1..] (j - 1)
> 13    else
> 14      0
> 15
> 16  // Unit: date2Day
> 17  let date2Day d m y =
> 18    let dayPrefix =
> 19      ["Sun"; "Mon"; "Tues"; "Wednes"; "Thurs"; "Fri";
> 20    "Satur"]
> 20    (* WB: 1 *)
> 21    let feb = if (y % 4 = 0) && ((y % 100 <> 0) || (y % 400
> 21    = 0)) then 29 else 28
> 22    let daysInMonth = [31; feb; 31; 30; 31; 30; 31; 31; 30;
> 22    31; 30; 31]
> 23    let dayOne = januaryFirstDay y
> 24    let daysSince = (sum daysInMonth (m - 2)) + d - 1
> 25    let weekday = (dayOne + daysSince) % 7;
> 26    dayPrefix.[weekday] + "day"
> ```

3. For each unit, produce an input set that tests each branch: In our example, the branch points depend on a Boolean expression, and for good measure, we are going

to test each term that can lead to branching. Using 't' and 'f' for `true` and `false`, we thus write,

Unit	Branch	Condition	Input	Expected output
`januaryFirstDay`	0	-	2016	5
`sum`	1	`0 <= j &&` `  j < lst.Length`		
	1a	`t && t`	`[1; 2; 3] 1`	3
	1b	`f && t`	`[1; 2; 3] -1`	0
	1c	`t && f`	`[1; 2; 3] 10`	0
	1d	`f && f`	-	-

`date2Day`	1	`(y % 4 = 0) &&`
		`  ((y % 100 <> 0)`
		`  ||`
		`    (y % 400 = 0))`		

	-	`t && (t || t)`	-	-
	1a	`t && (t || f)`	`8 9 2016`	Thursday
	1b	`t && (f || t)`	`8 9 2000`	Friday
	1c	`t && (f || f)`	`8 9 2100`	Wednesday
	-	`f && (t || t)`	-	-
	1d	`f && (t || f)`	`8 9 2015`	Tuesday
	-	`f && (f || t)`	-	-
	-	`f && (f || f)`	-	-

The impossible cases have been intentionally blank, e.g., it is not possible for $j < 0$ and $j > n$ for some positive value n.

4. Write a program that tests all these cases and checks the output, see Listing 10.3.

> **Listing 10.3 date2DayWhiteTest.fsx:**
> **The tests identified by white-box analysis. The program from List-**
> **ing 10.2 has been omitted for brevity.**

```
1  printfn "White-box testing of date2Day.fsx"
2  printfn "  Unit: januaryFirstDay"
3  printfn "    Branch: 0 - %b" (januaryFirstDay 2016 = 5)
4
5  printfn "  Unit: sum"
6  printfn "    Branch: 1a - %b" (sum [1; 2; 3] 1 = 3)
7  printfn "    Branch: 1b - %b" (sum [1; 2; 3] -1 = 0)
8  printfn "    Branch: 1c - %b" (sum [1; 2; 3] 10 = 0)
9
10 printfn "  Unit: date2Day"
11 printfn "    Branch: 1a - %b" (date2Day 8 9 2016 =
     "Thursday")
12 printfn "    Branch: 1b - %b" (date2Day 8 9 2000 =
     "Friday")
13 printfn "    Branch: 1c - %b" (date2Day 8 9 2100 =
     "Wednesday")
14 printfn "    Branch: 1d - %b" (date2Day 8 9 2015 =
     "Tuesday")
```

```
1  $ fsharpc --nologo date2DayWhiteTest.fsx && mono
     date2DayWhiteTest.exe
2  White-box testing of date2Day.fsx
3    Unit: januaryFirstDay
4      Branch: 0 - true
5    Unit: sum
6      Branch: 1a - true
7      Branch: 1b - true
8      Branch: 1c - true
9    Unit: date2Day
10     Branch: 1a - true
11     Branch: 1b - true
12     Branch: 1c - true
13     Branch: 1d - true
```

Notice that the output of the tests is organized such that they are enumerated per unit, hence we can rearrange as we like and still uniquely refer to a unit's test. Also, the output of the test program produces a list of tests that should return true or success or a similar positively loaded word, but without further or only little detail, such that we at a glance can identify any test that produced unexpected results.

After the white-box testing has failed to find errors in the program, we have some confidence in the program, since we have run every line at least once. It is, however, in no way a guarantee that the program is error free, which is why white-box testing is often accompanied with black-box testing to be described next.

10.2. Black-box Testing

In black-box testing, the program is considered a black box, and no knowledge is required about how a particular problem is solved. In fact, it is often useful not to have that

knowledge at all. It is rarely possible to test all input to a program, so in black-box testing, the solution is tested for typical and extreme cases based on knowledge of the problem. The procedure is as follows:

1. Decide on the interface to use: It is useful to have an agreement with the software developers about what interface is to be used, e.g., in our case, the software developer has made a function `date2Day d m y` where `d`, `m`, and `y` are integers specifying the day, month, and year.

2. Make an overall description of the tests to be performed and their purpose:

   1 a consecutive week, to ensure that all weekdays are properly returned

   2 two set of consecutive days across boundaries that may cause problems: across a new year, and across a regular month boundary.

   3 a set of consecutive days across February-March boundaries for a leap and non-leap year

   4 four dates after February in a non-leap year, a non-multiple-of-100 leap year, a multiple-of-100-but-not-of-400 non-leap year, and a multiple-of-400 leap year.

   Given no information about the program's text, there are other dates that one could consider as likely candidates for errors, but the above is judged to be a fair coverage.

3. Choose a specific set of input and expected output relations on the tabular form:

Test number	Input	Expected output
1a	1 1 2016	Friday
1b	2 1 2016	Saturday
1c	3 1 2016	Sunday
1d	4 1 2016	Monday
1e	5 1 2016	Tuesday
1f	6 1 2016	Wednesday
1g	7 1 2016	Thursday
2a	31 12 2014	Wednesday
2b	1 1 2015	Thursday
2c	30 9 2017	Saturday
2d	1 10 2017	Sunday
3a	28 2 2016	Sunday
3b	29 2 2016	Monday
3c	1 3 2016	Tuesday
3d	28 2 2017	Tuesday
3e	1 3 2017	Wednesday
4a	1 3 2015	Sunday
4b	1 3 2012	Thursday
4c	1 3 2000	Wednesday
4d	1 3 2100	Monday

4. Write a program executing the tests, as shown in Listing 10.4 and 10.5.

Listing 10.4 date2DayBlackTest.fsx:
The tests identified by black-box analysis. The program from List-
ing 10.2 has been omitted for brevity.

```
28  let testCases = [
29    ("A complete week",
30     [(1,  1,  2016,  "Friday");
31      (2,  1,  2016,  "Saturday");
32      (3,  1,  2016,  "Sunday");
33      (4,  1,  2016,  "Monday");
34      (5,  1,  2016,  "Tuesday");
35      (6,  1,  2016,  "Wednesday");
36      (7,  1,  2016,  "Thursday");]);
37    ("Across boundaries",
38     [(31, 12, 2014,  "Wednesday");
39      (1,  1,  2015,  "Thursday");
40      (30, 9,  2017,  "Saturday");
41      (1,  10, 2017,  "Sunday")]);
42    ("Across Feburary boundary",
43     [(28, 2,  2016,  "Sunday");
44      (29, 2,  2016,  "Monday");
45      (1,  3,  2016,  "Tuesday");
46      (28, 2,  2017,  "Tuesday");
47      (1,  3,  2017,  "Wednesday")]);
48    ("Leap years",
49     [(1,  3,  2015,  "Sunday");
50      (1,  3,  2012,  "Thursday");
51      (1,  3,  2000,  "Wednesday");
52      (1,  3,  2100,  "Monday")]);
53    ]
54
55  printfn "Black-box testing of date2Day.fsx"
56  for i = 0 to testCases.Length - 1 do
57    let (setName, testSet) = testCases.[i]
58    printfn "  %d. %s" (i+1) setName
59    for j = 0 to testSet.Length - 1 do
60      let (d, m, y, expected) = testSet.[j]
61      let day = date2Day d m y
62      printfn "    test %d - %b" (j+1) (day = expected)
```

Listing 10.5: Output from Listing 10.4.

```
1  $ fsharpc --nologo date2DayBlackTest.fsx && mono
     date2DayBlackTest.exe
2  Black-box testing of date2Day.fsx
3    1. A complete week
4      test 1 - true
5      test 2 - true
6      test 3 - true
7      test 4 - true
8      test 5 - true
9      test 6 - true
10     test 7 - true
11   2. Across boundaries
12     test 1 - true
13     test 2 - true
14     test 3 - true
15     test 4 - true
16   3. Across Feburary boundary
17     test 1 - true
18     test 2 - true
19     test 3 - true
20     test 4 - true
21     test 5 - true
22   4. Leap years
23     test 1 - true
24     test 2 - true
25     test 3 - true
26     test 4 - true
```

Notice how the program has been made such that it is almost a direct copy of the table produced in the previous step.

A black-box test is a statement of what a solution should fulfill for a given problem. Hence, **it is a good idea to make a black-box test early in the software design phase, in order to clarify the requirements for the code to be developed and take an outside view of the code prior to developing it.** Advice

After the black-box testing has failed to find errors in the program, we have some confidence in the program, since, from a user's perspective, the program produces sensible output in many cases. It is, however, in no way a guarantee that the program is error free.

10.3. Debugging by Tracing

Once an error has been found by testing, the *debugging* phase starts. The cause of a bug · debugging can either be that the chosen algorithm is the wrong one for the job, or the implementation of it has an error. In the debugging process, we have to keep an open mind and not rely on assumptions, since assumptions tend to blind the reader of a text. A frequent source of errors is that the state of a program is different than expected, e.g., because the calculation performed is different than intended, or that the return of a library function is different than expected. The most important tool for debugging is *simplification*. This is similar to white-box testing, but where the units tested are very small. E.g., the suspected piece of code could be broken down into smaller functions or code snippets, which are given

well-defined input, and, e.g., use `printfn` statements to obtain the output of the code
snippet. Another related technique is to use *mockup code*, which replaces parts of the code   · mockup code
with code that produces safe and relevant results. If the bug is not obvious, then more
rigorous techniques must be used, such as *tracing*. Some development interfaces have a   · tracing
built-in tracing system, e.g., `fsharpi` will print inferred types and some binding values.
However, often the source of a bug is due to a misunderstanding of the flow of data through
a program execution, and we will in the following introduce *hand tracing* as a technique   · hand tracing
to simulate the execution of a program by hand.

Consider the program in Listing 10.6. [6]

Listing 10.6 gcd.fsx:
The greatest common divisor of 2 integers.

```
1  let rec gcd a b =
2    if a < b then
3      gcd b a
4    elif b > 0 then
5      gcd b (a % b)
6    else
7      a
8
9  let a = 10
10  let b = 15
11  printfn "gcd %d %d = %d" a b (gcd a b)
```
- -
```
1  $ fsharpc --nologo gcd.fsx && mono gcd.exe
2  gcd 10 15 = 5
```

The program includes a function for calculating the greatest common divisor of 2 integers,
and calls this function with the numbers 10 and 15. Hand tracing this program means that
we simulate its execution and, as part of that, keep track of the bindings, assignments and
input and output of the program. To do this, we need to consider code snippets' *environ-*
ments. E.g., to hand trace the above program, we start by noting the outer environment,   · environment
called E_0 for short. In line 1, the `gcd` identifier is bound to a function, hence we write:

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$

Function bindings like this one are noted as closures, which are triplets of the form
$\big(\text{arguments}, \text{expression}, \text{environment}\big)$. The symbol $\varnothing$ denotes the empty environment. The
closure is everything needed for the expression to be calculated. Here, we wrote gcd-body
to denote everything after the equal sign in the function binding. Next, F# executes line 9
and 10, and we update our environment to reflect the bindings as,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$

In line 11, the function is evaluated. This initiates a new environment E_1, and we update

---

[6]Jon: **This program uses recursion, which has not been introduced yet.**

our trace as,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
$$\text{line 11: gcd a b} \rightarrow ?$$
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$

where the new environment is noted to have gotten its argument names `a` and `b` bound to the values `10` and `15` respectively, and where the return of the function to environment E_0 is yet unknown, so it is noted as a question mark. In line 2, the comparison `a < b` is checked, and since we are in environment E_1, this is the same as checking `10 < 15`, which is true, so the program executes line 3. Hence, we initiate a new environment E_2 and update our trace as,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
$$\text{line 11: gcd a b} \rightarrow ?$$
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$
$$\text{line 3: gcd b a} \rightarrow ?$$
$$E_2 : \big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$$

where in the new environment, `a` and `b` are bound to the values `15` and `10` respectively. In E_2, `10 < 15` is false, so the program evaluates `b > 0`, which is true, hence line 5 is executed. This calls `gcd` once again, but with new arguments. Since `a % b` is parenthesized, it is evaluated before `gcd` is called.

Hence, we update our trace as,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
$$\text{line 11: gcd a b} \rightarrow ?$$
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$
$$\text{line 3: gcd b a} \rightarrow ?$$
$$E_2 : \big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$$
$$\text{line 5: a \% b} \rightarrow 5$$
$$\text{line 5: gcd b (a \% b)} \rightarrow ?$$
$$E_3 : \big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$$

Again we fall through to line 5, evaluate the remainder operator and initiate a new envi-

ronment,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
line 11: gcd a b → ?
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$
line 3: gcd b a → ?
$$E_2 : \big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$$
line 5: a % b → 5
line 5: gcd b (a % b) → ?
$$E_3 : \big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$$
line 5: a % b → 0
line 5: gcd b (a % b) → ?
$$E_4 : \big((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \varnothing\big)$$

This time both `a < b` and `b > 0` are false, so we fall through to line 7, and `gcd` from E_4 returns its value of `a`, which is `5`, so we scratch E_4 and change the question mark in E_3 to `5`,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
line 11: gcd a b → ?
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$
line 3: gcd b a → ?
$$E_2 : \big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$$
line 5: a % b → 5
line 5: gcd b (a % b) → ?
$$E_3 : \big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$$
line 5: a % b → 0
line 5: gcd b (a % b) → ~~?~~ 5
$$\cancel{E_4} : \big((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \varnothing\big)$$

Now, line 5 in E_3 is also a return point of `gcd`, hence we scratch E_3 and change the question

mark in E_2 to 5,

E_0 :

    gcd $\rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$

    $a \rightarrow 10$

    $b \rightarrow 15$

    line 11: gcd a b $\rightarrow$?

E_1 :$\big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$

    line 3: gcd b a $\rightarrow$?

E_2 :$\big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$

    line 5: a % b $\rightarrow$ 5

    line 5: gcd b (a % b) $\rightarrow$ ⁊ 5

E_3 :$\big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$

    line 5: a % b $\rightarrow$ 0

    line 5: gcd b (a % b) $\rightarrow$ ⁊ 5

E_4 :$\big((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \varnothing\big)$

and likewise, for E_2 and E_1,

E_0 :

    gcd $\rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$

    $a \rightarrow 10$

    $b \rightarrow 15$

    line 11: gcd a b $\rightarrow$ ⁊ 5

E_1 :$\big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$

    line 3: gcd b a $\rightarrow$ ⁊ 5

E_2 :$\big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$

    line 5: a % b $\rightarrow$ 5

    line 5: gcd b (a % b) $\rightarrow$ ⁊ 5

E_3 :$\big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$

    line 5: a % b $\rightarrow$ 0

    line 5: gcd b (a % b) $\rightarrow$ ⁊ 5

E_4 :$\big((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \varnothing\big)$

Now we are able to continue the program in environment E_0 with the `printfn` statement,

and we write,

$$E_0 :$$
$$\text{gcd} \rightarrow \big((a, b), \text{gcd-body}, \varnothing\big)$$
$$a \rightarrow 10$$
$$b \rightarrow 15$$
$$\text{line 11: gcd a b} \rightarrow \lightning 5$$
$$\text{line 11: stdout} \rightarrow \text{"gcd 10 15 = 5"}$$
$$E_1 : \big((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \varnothing\big)$$
$$\text{line 3: gcd b a} \rightarrow \lightning 5$$
$$E_2 : \big((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \varnothing\big)$$
$$\text{line 5: a \% b} \rightarrow 5$$
$$\text{line 5: gcd b (a \% b)} \rightarrow \lightning 5$$
$$E_3 : \big((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \varnothing\big)$$
$$\text{line 5: a \% b} \rightarrow 0$$
$$\text{line 5: gcd b (a \% b)} \rightarrow \lightning 5$$
$$E_4 : \big((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \varnothing\big)$$

which completes the hand tracing of `gcd.fsx`.

F# uses the lexical scope, which implies that besides function arguments, we also at times need to consider the environment at the place of writing. Consider the program in Listing 10.7.

Listing 10.7 lexicalScopeTracing.fsx:
Example of lexical scope and closure environment.

```
1  let testScope x =
2    let a = 3.0
3    let f z = a * z
4    let a = 4.0
5    f x
6  printfn "%A" (testScope 2.0)
```

```
1  $ fsharpc --nologo lexicalScopeTracing.fsx
2  $ mono lexicalScopeTracing.exe
3  6.0
```

To hand trace this, we start by creating the outer environment, define the closure for `testScope`, and reach line 6,

$$E_0 :$$
$$\text{testScope} \rightarrow \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope 2.0} \rightarrow ?$$

We create a new environment for `testScope` and note the bindings,

E_0 :

testScope $\rightarrow \big(x, \text{testScope-body}, \varnothing\big)$

line 6: testScope 2.0 $\rightarrow$?

E_1 : $\big(x \rightarrow 2.0, \text{testScope-body}, \varnothing\big)$

$a \rightarrow 3.0$

f $\rightarrow \big(z, \text{a * z}, (a \rightarrow 3.0, x \rightarrow 2.0)\big)$

$a \rightarrow 4.0$

Since we are working with lexical scope, `a` is noted twice, and its interpretation is by lexical order. Hence, the environment for the closure of `f` is everything above in E_1, so we add $a \rightarrow 3.0$ and $x \rightarrow 2.0$. In line 5, `f` is called, so we create an environment based on its closure,

E_0 :

testScope $\rightarrow \big(x, \text{testScope-body}, \varnothing\big)$

line 6: testScope 2.0 $\rightarrow$?

E_1 : $\big(x \rightarrow 2.0, \text{testScope-body}, \varnothing\big)$

$a \rightarrow 3.0$

f $\rightarrow \big(z, \text{a * z}, (a \rightarrow 3.0, x \rightarrow 2.0)\big)$

$a \rightarrow 4.0$

line 5: f x $\rightarrow$?

E_2 : $\big(z \rightarrow 2.0, \text{a * z}, (a \rightarrow 3.0, x \rightarrow 2.0)\big)$

The expression in the environment E_2 evaluates to `6.0`, and unraveling the scopes we get,

E_0 :

testScope $\rightarrow \big(x, \text{testScope-body}, \varnothing\big)$

line 6: testScope 2.0 $\rightarrow$ ⤳ 6.0

line 6: stdout $\rightarrow$ "6.0"

$\cancel{E_1}$: $\big(x \rightarrow 2.0, \text{testScope-body}, \varnothing\big)$

$a \rightarrow 3.0$

f $\rightarrow \big(z, \text{a * z}, (a \rightarrow 3.0, x \rightarrow 2.0)\big)$

$a \rightarrow 4.0$

line 5: f x $\rightarrow$ ⤳ 6.0

$\cancel{E_2}$: $\big(z \rightarrow 2.0, \text{a * z}, (a \rightarrow 3.0, x \rightarrow 2.0)\big)$

For mutable binding, the scope is dynamic. For this, we need the concept of storage. Consider the program in Listing 10.8.

Listing 10.8 dynamicScopeTracing.fsx:
Example of lexical scope and closure environment.

```
1  let testScope x =
2    let mutable a = 3.0
3    let f z = a * z
4    a <- 4.0
5    f x
6  printfn "%A" (testScope 2.0)
```

```
1  $ fsharpc --nologo dynamicScopeTracing.fsx
2  $ mono dynamicScopeTracing.exe
3  8.0
```

We add a storage area to our hand tracing, e.g., line 6,

$$\text{Store} :$$
$$E_0 :$$
$$\text{testScope} \rightarrow \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \rightarrow ?$$

So when we generate environment E_1, the mutable binding is to a storage location,

$$\text{Store} :$$
$$\alpha_1 \rightarrow 3.0$$
$$E_0 :$$
$$\text{testScope} \rightarrow \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \rightarrow ?$$
$$E_1 : \big(x \rightarrow 2.0, \text{testScope-body}, \varnothing\big)$$
$$a \rightarrow \alpha_1$$

which is assigned the value `3.0` at the definition of `a`. Now, the definition of `f` is using the storage location,

$$\text{Store} :$$
$$\alpha_1 \rightarrow 3.0$$
$$E_0 :$$
$$\text{testScope} \rightarrow \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \rightarrow ?$$
$$E_1 : \big(x \rightarrow 2.0, \text{testScope-body}, \varnothing\big)$$
$$a \rightarrow \alpha_1$$
$$f \rightarrow \big(z, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0)\big)$$

and in line 4, it is the value in the storage which is updated,

$$\text{Store}:$$
$$\alpha_1 \to \cancel{3.0}\ 4.0$$
$$E_0:$$
$$\text{testScope} \to \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \to ?$$
$$E_1:\big(x \to 2.0, \text{testScope-body}, \varnothing\big)$$
$$a \to \alpha_1$$
$$\text{f} \to \big(z, \text{a * z}, (a \to \alpha_1, x \to 2.0)\big)$$

Hence,

$$\text{Store}:$$
$$\alpha_1 \to \cancel{3.0}\ 4.0$$
$$E_0:$$
$$\text{testScope} \to \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \to ?$$
$$E_1:\big(x \to 2.0, \text{testScope-body}, \varnothing\big)$$
$$a \to \alpha_1$$
$$\text{f} \to \big(z, \text{a * z}, (a \to \alpha_1, x \to 2.0)\big)$$
$$\text{line 5: f x} \to ?$$
$$E_2:\big(z \to 2.0, \text{a * z}, (a \to \alpha_1, x \to 2.0)\big)$$

and the return value from `f` evaluated in environment E_2 now reads the value `4.0` for `a` and returns `8.0`. Hence,

$$\text{Store}:$$
$$\alpha_1 \to \cancel{3.0}\ 4.0$$
$$E_0:$$
$$\text{testScope} \to \big(x, \text{testScope-body}, \varnothing\big)$$
$$\text{line 6: testScope } 2.0 \to \cancel{?}\ 8.0$$
$$\text{line 6: stdout} \to \text{"8.0"}$$
$$\cancel{E_1}:\big(x \to 2.0, \text{testScope-body}, \varnothing\big)$$
$$a \to \alpha_1$$
$$\text{f} \to \big(z, \text{a * z}, (a \to \alpha_1, x \to 2.0)\big)$$
$$\text{line 5: f x} \to \cancel{?}\ 8.0$$
$$\cancel{E_2}:\big(z \to 2.0, \text{a * z}, (a \to \alpha_1, x \to 2.0)\big)$$

By the above examples, it is seen that hand tracing can be used to study the flow of data through a program in detail. It may seem tedious in the beginning, but the care illustrated above is useful to ensure rigor in the analysis. Most will find that once accustomed to the method, the analysis can be performed rigorously but with less paperwork, and in conjunction with strategically placed debugging `printfn` statements, it is a very valuable tool for debugging.

11 | Collections of Data

F# is tuned to work with collections of data, and there are several built-in types of collections with various properties making them useful for different tasks. Examples include strings, lists, and arrays. Strings were discussed in Chapter 5 and will be revisited here in more details.

The data structures discussed below all have operators, properties, methods, and modules to help you write elegant programs using them.

Properties and methods are common object-oriented terms used in conjunction with the discussed functionality. They are synonymous with values and functions and will be discussed in Chapter 20. Properties and methods for a value or variable are called using the *dot notation*, i.e., with the "`.`"-lexeme. For example, `"abcdefg".Length` is a property and · dot notation
is equal to the length of the string, and `"abcdefg".ToUpper()` is a method and creates a new string where all characters have been converted to upper case.

The data structures also have accompanying modules with a wealth of functions and where some are mentioned here. Further, the data structures are all implemented as classes offering even further functionality. The modules are optimized for functional programming, see Chapters 13 to 17, while classes are designed to support object-oriented programming, see Chapters 20 to 22.

In the following, a brief overview of many properties, methods, and functions is given by describing their name and type-definition, and by giving a short description and an example of their use. Several definitions are general and works with many different types. To describe this we will use the notation of generic types, see Section 6.2. The name of a generic type starts with the "`'`" lexeme, such as `'T`. The implication of the appearance of a generic type in, e.g., a function's type-definition, is that the function may be used with any real type such as `int` or `char`. If the same generic type name is used in several places in the type-definition, then the function must use a real type consistently. For example, The `List.fromArray` function has type `arr:'T [] -> 'T list`, meaning that it takes an array of some type and returns a list of the same type.

See the F# Language Reference at `https://docs.microsoft.com/en-us/dotnet/fsharp/` for a full description of all available functionality including variants of those included here.

11.1. Strings

Strings have been discussed in Chapter 5, the content of which will be briefly revisited here followed by a description of some of the many supporting built-in functions in F# on strings.

A *string* is a sequence of characters. Each character is represented using UTF-16, see   · string
Appendix C for further details on the unicode standard. The type `string` is an alias for
`System.string`. String literals are delimited by double quotation marks ""⟩" and inside   · `System.string`
the delimiters, character escape sequences are allowed (see Table 5.2), which are replaced
by the corresponding character code. Examples are `"This is a string"`, `"\tTabulated`
`string"`, `"A \"quoted\" string"`, and `""`. Strings may span several lines, and new lines
inside strings are part of the string unless the line is ended with a backslash. Strings
may be *verbatim* by preceding the string with "`@`", in which case escape sequences are not   · verbatim string
replaced, but two double quotation marks are an escape sequence which is replaced by a
one double quotation mark. Examples of "`@`"-verbatim strings are: `@"This is a string"`,
`@"\tNon-tabulated string"`, `@"A ""quoted"" string"`, and `@""`. Alternatively, a ver-
batim string may be delimited by three double quotation marks. Examples of "`"""`"-
verbatim strings are: `"""This is a string"""`, `"""\tNon-tabulated string"""`, `"""A`
`"quoted" string"""`, and `""""""`. Strings may be indexed using the `.[]` notation, as
demonstrated in Listing 5.27.

11.1.1. String Properties and Methods

Strings have a few properties which are values attached to each string and accessed using
the "`.`" notation. The only to be mentioned here is:

`IndexOf()`: `str:string -> int`. Returns the index of the first occurence of `s` or -1, if
   `str` does not appear in the string.

> **Listing 11.1:** `IndexOf()`
> ```
> 1  > "Hello World".IndexOf("World");;
> 2  val it : int = 6
> ```

`Length`: `int`. Returns the length of the string.

> **Listing 11.2:** `Length`
> ```
> 1  val it : int = 4
> ```

`ToLower()`: `unit -> string`. Returns a copy of the string where each letter has been
   converted to lower case.

> **Listing 11.3:** `ToLower()`
> ```
> 1  > "aBcD".ToLower();;
> 2  val it : string = "abcd"
> ```

`ToUpper()`: `unit -> string`. Returns a copy of the string where each letter has been
   converted to upper case.

Listing 11.4: ToUpper()

```
1  > "aBcD".ToUpper ();;
2  val it : string = "ABCD"
```

Trim(): `unit -> string`. Returns a copy of the string where leading and trailing whitespaces have been removed.

Listing 11.5: Trim()

```
1  > "  Hello World  ".Trim ();;
2  val it : string = "Hello World"
```

Split(): `unit -> string []`. Splits a string of words separated by spaces into an array of words. See Section 11.3 for more information about arrays.

Listing 11.6: Split()

```
1  > "Hello World".Split ();;
2  val it : string [] = [|"Hello"; "World"|]
```

11.1.2. The String Module

The `String` module offers many functions for working with strings. Some of the most powerful ones are listed below, and they are all higher order functions.

String.collect: `f:(char -> string) -> str:string -> string`. Creates a new string whose characters are the results of applying `f` to each of the characters of `str` and concatenating the resulting strings.

Listing 11.7: String.collect

```
1  > String.collect (fun c -> (string c) + ", ") "abc";;
2  val it : string = "a, b, c, "
```

String.exists: `f:(char -> bool) -> str:string -> bool`. Returns true if any character in `str` evaluates to true when using `f`.

Listing 11.8: String.exists

```
1  > String.exists (fun c -> c = 'd') "abc";;
2  val it : bool = false
```

String.forall: `f:(char -> bool) -> str:string -> bool`. Returns true if all characters in `str` evalutes to true when using `f`.

Listing 11.9: `String.forall`

```
1  > String.forall (fun c -> c < 'd') "abc";;
2  val it : bool = true
```

`String.init: n:int -> f:(int -> string) -> string`. Creates a new string with length `n` and whose characters are the result of applying `f` to each index of that string.

Listing 11.10: `String.init`

```
1  > String.init 5 (fun i -> (string i) + ", ");;
2  val it : string = "0, 1, 2, 3, 4, "
```

`String.iter: f:(char -> unit) -> str:string -> unit`. Applies `f` to each character in `str`.

Listing 11.11: `String.iter`

```
1  > String.iter (fun c -> printfn "%c" c) "abc";;
2  a
3  b
4  c
5  val it : unit = ()
```

`String.map: f:(char -> char) -> str:string -> string`. Creates a new string whose characters are the results of applying `f` to each of the characters of `str`.

Listing 11.12: `String.map`

```
1  > let toUpper c = c + char (int 'A' - int 'a')
2  - String.map toUpper "abcd";;
3  val toUpper : c:char -> char
4  val it : string = "ABCD"
```

11.2. Lists

Lists are unions of immutable values of the same type. A list can be expressed as a *sequence* · list
expression,   · sequence expression

Listing 11.13: The syntax for a list using the sequence expression.

```
1  [[<expr>{; <expr>}]]
```

Examples are a list of integers `[1; 2; 3]`, a list of strings `["This"; "is"; "a"; "list"]`, a list of anonymous functions `[(fun x -> x); (fun x -> x*x)]`, and an empty list `[]`. Lists may also be given as ranges,

Figure 11.1.: A list is a linked list: Here is illustrated the linked list of `["This"; "is"; "a"; "list"]`.

Listing 11.14: The syntax for a list using the range expressions.

```
1  [<expr> .. <expr> [.. <expr>]]
```

where `<expr>` in *range expressions* must be of integers, floats, or characters. Examples are `[1 .. 5]`, `[-3.0 .. 2.0]`, and `['a' .. 'z']`. Range expressions may include a step size, thus, `[1 .. 2 .. 10]` evaluates to `[1; 3; 5; 7; 9]`.

· range expressions

A list type is identified with the `list` keyword, such that a list of integers has the type `int list`. Like strings, lists may be indexed using the "`.[]`" notation, the lengths of lists is retrieved using the `Length` property, and we may test whether a list is empty by using the `isEmpty` property. These features are demonstrated in Listing 11.15.

· `list`
· `.[]`
· `Length`
· `isEmpty`

Listing 11.15 listIndexing.fsx:
Lists are indexed as strings and has a `Length` property.

```
1  let printList (lst : int list) : unit =
2    for i = 0 to lst.Length - 1 do
3      printf "%A " lst.[i]
4    printfn ""
5
6  let lst = [3; 4; 5]
7  printfn "lst = %A, lst.[1] = %A" lst lst.[1]
8  printfn "lst.Length = %A, lst.isEmpty = %A" lst.Length
     lst.IsEmpty
9  printList lst
```

```
1  $ fsharpc --nologo listIndexing.fsx && mono listIndexing.exe
2  lst = [3; 4; 5], lst.[1] = 4
3  lst.Length = 3, lst.isEmpty = false
4  3 4 5
```

F# implements lists as linked lists, as illustrated in Figure 11.1. As a consequence, indexing element i has *computational complexity* $\mathcal{O}(i)$. The computational complexity of an operation is a description of how long a computation will take without considering the hardware it is performed on. The notation is sometimes called *Big-O* notation or *Landau notation*. In the present case, the complexity is $\mathcal{O}(i)$, which means that the complexity is linear in i and indexing element $i+1$ takes 1 unit longer than indexing element i when i is very large. The size of the unit is on purpose unspecified and depends on implementation and hardware details. Nevertheless, Big-O notation is a useful tool for reasoning about the efficiency of an operation. F# has access to the list's elements only by traversing the list from its beginning. I.e., to obtain the value of element i, F# starts with element 0, follows the link to element 1 and so on, until element i is reached. To reach element $i+1$ instead, we would need to follow 1 more link, and assuming that following a single link takes some

· computational complexity
· Big-O
· Landau notation

constant amount of time we find that the computational complexity is $\mathcal{O}(i)$. Compared to arrays, to be discussed below, this is slow, which is why **indexing lists should be avoided.**   Advice

Notice especially that lists are zero-indexed, and thus, the last element in a list `lst` is `lst.Length -1`. This is a very common source of error! Therefore, indexing in lists using *for*-loops is supported using a special notation with the *in* keyword,   · for
· in

Listing 11.16:   For-in loop with in expression.

```
1   for <ident> in <list> do <bodyExpr> [done]
```

In `for-in` loops, the loop runs through each element of the `<list>`, and assigns it to the identifier `<ident>`. This is demonstrated in Listing 11.17.

Listing 11.17 listFor.fsx:
The `for`-`in` loops are preferred over `for`-`to` loops.

```
1   let printList (lst : int list) : unit =
2     for elm in lst do
3       printf "%A " elm
4     printfn ""
5
6   printList [3; 4; 5]
```

```
1   $ fsharpc --nologo listFor.fsx && mono listFor.exe
2   3 4 5
```

Using `for-in`-expressions remove the risk of off-by-one indexing errors, and thus, `for-in`   Advice
is to be preferred over `for-to`.

Lists support slicing identically to strings, as demonstrated in Listing 11.18.

Listing 11.18: Examples of list slicing. Compare with Listing 5.27.

```
1  > let lst = ['a' .. 'g'];;
2  val lst : char list = ['a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g']
3
4  > lst.[0];;
5  val it : char = 'a'
6
7  > lst.[3];;
8  val it : char = 'd'
9
10 > lst.[3..];;
11 val it : char list = ['d'; 'e'; 'f'; 'g']
12
13 > lst.[..3];;
14 val it : char list = ['a'; 'b'; 'c'; 'd']
15
16 > lst.[1..3];;
17 val it : char list = ['b'; 'c'; 'd']
18
19 > lst.[*];;
20 val it : char list = ['a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g']
```

Lists may be concatenated using either the *"@"*[1] *concatenation* operator or the *"::" cons*   · @
operators. The difference is that "@" concatenates two lists of identical types, while "::"   · list concatenation
concatenates an element and a list of identical types. This is demonstrated in Listing 11.19.   · ::
   · list cons

Listing 11.19: Examples of list concatenation.

```
1  > ([1] @ [2; 3]);;
2  val it : int list = [1; 2; 3]
3
4  > ([1; 2] @ [3; 4]);;
5  val it : int list = [1; 2; 3; 4]
6
7  > (1 :: [2; 3]);;
8  val it : int list = [1; 2; 3]
```

Since lists are represented as linked lists, the cons operator is very efficient and has computational complexity $\mathcal{O}(1)$, while concatenation has computational complexity $\mathcal{O}(n)$, where n is the length of the first list.

It is possible to make multidimensional lists as lists of lists, as shown in Listing 11.20.

---

[1]Jon: **why does the at-symbol not appear in the index?**

Figure 11.2.: A list is a ragged linked list: Here is illustrated the linked list of `[[1;2];[3;4;5]]`.

Listing 11.20 listMultidimensional.fsx:
A ragged multidimensional list, built as lists of lists, and its indexing.

```
1  let a = [[1;2];[3;4;5]]
2  let row = a.Item 0 in printfn "%A" row
3  let elm = row.Item 1 in printfn "%A" elm
4  let elm = (a.Item 0).Item 1 in printfn "%A" elm
```

```
1  $ fsharpc --nologo listMultidimensional.fsx
2  $ mono listMultidimensional.exe
3  [1; 2]
4  2
5  2
```

The example shows a *ragged multidimensional list*, since each row has a different number of elements. This is also illustrated in Figure 11.2. · ragged multidimensional list

The indexing of a particular element is slow due to the linked list implementation of lists, which is why arrays are often preferred for two- and higher-dimensional data structures, see Section 11.3.

11.2.1. List Properties

Lists support a number of properties, some of which are listed below.

`Head:` Returns the first element of a list.

Listing 11.21: `Head`

```
1  > [1; 2; 3].Head;;
2  val it : int = 1
```

`IsEmpty:` Returns true if the list is empty.

Listing 11.22: `Head`

```
1  > [1; 2; 3].IsEmpty;;
2  val it : bool = false
```

Length: Returns the number of elements in the list.

> **Listing 11.23:** Length

```
1  > [1; 2; 3].Length;;
2  val it : int = 3
```

Tail: Returns the list, except for its first element.

> **Listing 11.24:** Tail

```
1  > [1; 2; 3].Tail;;
2  val it : int list = [2; 3]
```

11.2.2. The List Module

The built-in `List` module contains a wealth of functions for lists, some of which are briefly summarized below:

List.collect: f:('T -> 'U list) -> lst:'T list -> 'U list. Applies `f` to each element in `lst` and return a concatenated list of the results.

> **Listing 11.25:** List.collect

```
1  > List.collect (fun elm -> [elm; elm; elm]) [1; 2; 3];;
2  val it : int list = [1; 1; 1; 2; 2; 2; 3; 3; 3]
```

List.contains: elm:'T -> lst:'T list -> bool. Returns true or false depending on whether or not `elm` is contained in `lst`.

> **Listing 11.26:** List.contains

```
1  > List.contains 3 [1; 2; 3];;
2  val it : bool = true
```

List.filter: f:('T -> bool) -> lst:'T list -> 'T list. Returns a new list with all the elements of `lst` for which `f` evaluates to true.

> **Listing 11.27:** List.filter

```
1  > List.filter (fun x -> x % 2 = 1) [0 .. 9];;
2  val it : int list = [1; 3; 5; 7; 9]
```

List.find: f:('T -> bool) -> lst:'T list -> 'T. Returns the first element of `lst` for which `f` is true.

Listing 11.28: `List.find`

```
1  > List.find (fun x -> x % 2 = 1) [0 .. 9];;
2  val it : int = 1
```

`List.findIndex: f:('T -> bool) -> lst:'T list -> int.` Returns the index of the first element of `lst` for which `f` is true.

Listing 11.29: `List.findIndex`

```
1  > List.findIndex (fun x -> x = 'k') ['a' .. 'z'];;
2  val it : int = 10
```

`List.fold: f:('State -> 'T -> 'State) -> elm:'State -> lst:'T list -> 'State.` Updates an accumulator iteratively by applying `f` to each element in `lst`. The initial value of the accumulator is `elm`. For example, when `lst` consists of `n+1` elements `List.fold` calculates:

$$f \ (\dots \ (f \ (f \ \text{elm} \ \text{lst.}[0]) \ \text{lst.}[1]) \ \dots) \ \text{lst.}[n].$$

Listing 11.30: `List.fold`

```
1  > let addSquares acc elm = acc + elm*elm
2  - List.fold addSquares 0 [0 .. 9];;
3  val addSquares : acc:int -> elm:int -> int
4  val it : int = 285
```

`List.foldBack: f:('T -> 'State -> 'State) -> lst:'T list -> elm:'State -> 'State.` Updates an accumulator iteratively backwards by applying `f` to each element in `lst`. The initial value of the accumulator is `elm`. For exampel, when `lst` consists of `n+1` elements `List.foldBack` calculates:

$$f \ \text{lst.}[0] \ (f \ \text{lst.}[1] \ (\dots(f \ \text{lst.}[n] \ \text{elm}) \ \dots)).$$

Listing 11.31: `List.foldBack`

```
1  > let addSquares elm acc = acc + elm*elm
2  - List.foldBack addSquares [0 .. 9] 0;;
3  val addSquares : elm:int -> acc:int -> int
4  val it : int = 285
```

`List.forall: f:('T -> bool) -> lst:'T list -> bool.` Returns true if all elements in `lst` are true when `f` is applied to them.

Listing 11.32: `List.forall`

```
1  > List.forall (fun x -> x % 2 = 1) [0 .. 9];;
2  val it : bool = false
```

`List.head: lst:'T list -> int.` Returns the first element in `lst`. An exception is raised

if `lst` is empty. See Section 18.1 for more on exceptions.

Listing 11.33: `List.head`

```
1  > List.head [1; -2; 0];;
2  val it : int = 1
```

`List.init: m:int -> f:(int -> 'T) -> 'T list`. Create a list with `m` elements and whose value is the result of applying `f` to the index of the element.

Listing 11.34: `List.init`

```
1  > List.init 10 (fun i -> i * i);;
2  val it : int list = [0; 1; 4; 9; 16; 25; 36; 49; 64; 81]
```

`List.isEmpty: lst:'T list -> int`. Returns true if `lst` is empty.

Listing 11.35: `List.isEmpty`

```
1  > List.isEmpty [1; 2; 3];;
2  val it : bool = false
```

`List.iter: f:('T -> unit) -> lst:'T list -> unit`. Applies `f` to every element in `lst`.

Listing 11.36: `List.iter`

```
1  > List.iter (fun x -> printfn "%A " x) [0; 1; 2];;
2  0
3  1
4  2
5  val it : unit = ()
```

`List.map: f:('T -> 'U) -> lst:'T list -> 'U list`. Returns a list as a concatenation of applying `f` to every element of `lst`.

Listing 11.37: `List.map`

```
1  > List.map (fun x -> x*x) [0 .. 9];;
2  val it : int list = [0; 1; 4; 9; 16; 25; 36; 49; 64; 81]
```

`List.ofArray: arr:'T [] -> 'T list`. Returns a list whose elements are the same as `arr`. See Section 11.3 for more on arrays.

Listing 11.38: `List.ofArray`

```
1  > List.ofArray [|1; 2; 3|];;
2  val it : int list = [1; 2; 3]
```

List.rev: `lst:'T list -> 'T list`. Returns a new list with the same elements as in `lst` but in reversed order.

> **Listing 11.39:** `List.rev`

```
1  > List.rev [1; 2; 3];;
2  val it : int list = [3; 2; 1]
```

List.sort: `lst:'T list -> 'T list`. Returns a new list with the same elements as in `lst` but where the elements are sorted.

> **Listing 11.40:** `List.sort`

```
1  > List.sort [3; 1; 2];;
2  val it : int list = [1; 2; 3]
```

List.tail: `'T list -> 'T list`. Returns a new list identical to `lst` but without its first element. An Exception is raised if `lst` is empty. See Section 18.1 for more on exceptions.

> **Listing 11.41:** `List.tail`

```
1  > List.tail [1; 2; 3];;
2  val it : int list = [2; 3]
3
4  > let a = [1; 2; 3] in List.tail a;;
5  val it : int list = [2; 3]
```

List.toArray: `lst:'T list -> 'T []`. Returns an array whose elements are the same as `lst`. See Section 11.3 for more on arrays.

> **Listing 11.42:** `List.toArray`

```
1  > List.toArray [1; 2; 3];;
2  val it : int [] = [|1; 2; 3|]
```

List.unzip: `lst:('T1 * 'T2) list -> 'T1 list * 'T2 list`. Returns a pair of lists of all the first elements and all the second elements of `lst`, respectively.

> **Listing 11.43:** `List.unzip`

```
1  > List.unzip [(1, 'a'); (2, 'b'); (3, 'c')];;
2  val it : int list * char list = ([1; 2; 3], ['a'; 'b';
     'c'])
3
4  >
```

List.zip: `lst1:'T1 list -> lst2:'T2 list -> ('T1 * 'T2) list`. Returns a list of pairs, where elements in `lst1` and `lst2` are iteratively paired.

> **Listing 11.44:** `List.zip`
>
> ```
> 1  > List.zip [1; 2; 3] ['a'; 'b'; 'c'];;
> 2  val it : (int * char) list = [(1, 'a'); (2, 'b'); (3, 'c')]
> ```

11.3. Arrays

One dimensional *arrays*, or just arrays for short, are mutable lists of the same type and follow a similar syntax as lists. Arrays can be stated as a *sequence expression*,

· arrays

· sequence expression

> **Listing 11.45:** The syntax for an array using the sequence expression.
>
> ```
> 1  [|[<expr>{; <expr>}]|]
> ```

Examples are arrays of integers `[|1; 2; 3|]`, of strings `[|"This"; "is"; "an"; "array"|]`, of functions `[|(fun x -> x); (fun x -> x*x)|]`, and an empty array `[||]`. Arrays may also be given as ranges,

> **Listing 11.46:** The syntax for an array using the range expression.
>
> ```
> 1  [|<expr> .. <expr> [.. <expr>]|]
> ```

but arrays of *range expressions* must be of `<expr>` integers, floats, or characters. Examples are `[|1 .. 5|]`, `[|-3.0 .. 2.0|]`, and `[|'a' .. 'z'|]`. Range expressions may include a step size, thus, `[|1 .. 2 .. 10|]` evaluates to `[|1; 3; 5; 7; 9|]`.

· range expressions

The array type is defined using the **array** keyword or alternatively the "`[]`" lexeme. Like strings and lists, arrays may be indexed using the "`.[]`" notation. Arrays cannot be resized, but are mutable, as shown in Listing 11.47.

· `.[]`

> **Listing 11.47 arrayReassign.fsx:**
> Arrays are mutable in spite of the missing `mutable` keyword.
>
> ```
> 1  let square (a : int array) =
> 2    for i = 0 to a.Length - 1 do
> 3      a.[i] <- a.[i] * a.[i]
> 4
> 5  let A = [| 1; 2; 3; 4; 5 |]
> 6  printfn "%A" A
> 7  square A
> 8  printfn "%A" A
> ```
> ```
> 1  $ fsharpc --nologo arrayReassign.fsx && mono arrayReassign.exe
> 2  [|1; 2; 3; 4; 5|]
> 3  [|1; 4; 9; 16; 25|]
> ```

Notice that in spite of the missing `mutable` keyword, the function `square` still has the

side-effect of squaring all entries in `A`. F# implements arrays as chunks of memory and · side-effect
indexes arrays via address arithmetic. I.e., element i in an array, whose first element is in
memory address α and whose elements fill β addresses each, is found at address $\alpha + i\beta$.[2]
Hence, indexing has computational complexity of $\mathcal{O}(1)$, but appending and prepending
values to arrays and array concatenation requires copying the new and existing values to
a fresh area in memory and thus has computational complexity $\mathcal{O}(n)$, where n is the total
number of elements. Thus, **indexing arrays is fast, but cons and concatenation is** Advice
slow and should be avoided.

Arrays support *slicing*, that is, indexing an array with a range result in a copy of the array · slicing
with values corresponding to the range. This is demonstrated in Listing 11.48.

Listing 11.48: Examples of array slicing. Compare with Listing 11.18 and Listing 5.27.

```
 1  > let arr = [|'a' .. 'g'|];;
 2  val arr : char [] = [|'a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g'|]
 3
 4  > arr.[0];;
 5  val it : char = 'a'
 6
 7  > arr.[3];;
 8  val it : char = 'd'
 9
10  > arr.[3..];;
11  val it : char [] = [|'d'; 'e'; 'f'; 'g'|]
12
13  > arr.[..3];;
14  val it : char [] = [|'a'; 'b'; 'c'; 'd'|]
15
16  > arr.[1..3];;
17  val it : char [] = [|'b'; 'c'; 'd'|]
18
19  > arr.[*];;
20  val it : char [] = [|'a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g'|]
```

As illustrated, the missing start or end index imply from the first or to the last element,
respectively.

Arrays do not have explicit operator support for appending and concatenation, instead the
`Array` namespace includes an `Array.append` function, as shown in Listing 11.49.

Listing 11.49 arrayAppend.fsx:
Two arrays are appended with `Array.append`.

```
 1  let a = [|1; 2;|]
 2  let b = [|3; 4; 5|]
 3  let c = Array.append a b
 4  printfn "%A, %A, %A" a b c
```
- -
```
 1  $ fsharpc --nologo arrayAppend.fsx && mono arrayAppend.exe
 2  [|1; 2|], [|3; 4; 5|], [|1; 2; 3; 4; 5|]
```

---

[2]Jon: **Add a figure illustrating address indexing.**

Arrays are *reference types*, meaning that identifiers are references and thus suffer from · reference types aliasing, as illustrated in Listing 11.50.

Listing 11.50 arrayAliasing.fsx:
Arrays are reference types and suffer from aliasing.

```
1  let a = [|1; 2; 3|];
2  let b = a
3  a.[0] <- 0
4  printfn "a = %A, b = %A" a b;;
```

```
1  $ fsharpc --nologo arrayAliasing.fsx && mono arrayAliasing.exe
2  a = [|0; 2; 3|], b = [|0; 2; 3|]
```

11.3.1. Array Properties and Methods

Some important properties and methods for arrays are:

`Clone(): 'T []`. Returns a copy of the array.

Listing 11.51: `Clone`

```
1  > let a = [|1; 2; 3|];
2  - let b = a.Clone()
3  - a.[0] <- 0
4  - printfn "a = %A, b = %A" a b;;
5  a = [|0; 2; 3|], b = [|1; 2; 3|]
6  val a : int [] = [|0; 2; 3|]
7  val b : obj = [|1; 2; 3|]
8  val it : unit = ()
```

`Length: int`. Returns the number of elements in the array.

Listing 11.52: `Length`

```
1  > [|1; 2; 3|].Length;;
2  val it : int = 3
```

11.3.2. The Array Module

There are quite a number of built-in procedures for arrays in the `Array` module, some of which are summarized below.

`Array.append: arr1:'T [] -> arr2:'T [] -> 'T []`. Creates an new array whose elements are a concatenated copy of `arr1` and `arr2`.

Listing 11.53: `Array.append`

```
1  > Array.append [|1; 2;|] [|3; 4; 5|];;
2  val it : int [] = [|1; 2; 3; 4; 5|]
```

`Array.contains: elm:'T -> arr:'T [] -> bool`. Returns true if `arr` contains `elm`.

Listing 11.54: `Array.contains`

```
1  > Array.contains 3 [|1; 2; 3|];;
2  val it : bool = true
```

`Array.exists: f:('T -> bool) -> arr:'T [] -> bool`. Returns true if any application of `f` evaluates to true when applied to the elements of `arr`.

Listing 11.55: `Array.exists`

```
1  > Array.exists (fun x -> x % 2 = 1) [|0 .. 2 .. 4|];;
2  val it : bool = false
```

`Array.filter: f:('T -> bool) -> arr:'T [] -> 'T []`. Returns an array of elements from `arr` who evaluate to true when `f` is applied to them.

Listing 11.56: `Array.filter`

```
1  > Array.filter (fun x -> x % 2 = 1) [|0 .. 9|];;
2  val it : int [] = [|1; 3; 5; 7; 9|]
```

`Array.find: f:('T -> bool) -> arr:'T [] -> 'T`. Returns the first element in `arr` for which `f` evaluates to true. The `System.Collections.Generic.KeyNotFoundException` exception is raised if no element is found. See Section 18.1 for more on exceptions.

Listing 11.57: `Array.find`

```
1  > Array.find (fun x -> x % 2 = 1) [|0 .. 9|];;
2  val it : int = 1
```

`Array.findIndex: f:('T -> bool) -> arr:'T [] -> int`. Returns the index of the first element in in `arr` for which `f` evaluates to true. If none are found, then the `System.Collections.Generic.KeyNotFoundException` exception is raised. See Section 18.1 for more on exceptions.

Listing 11.58: `Array.findIndex`

```
1  > Array.findIndex (fun x -> x = 'k') [|'a' .. 'z'|];;
2  val it : int = 10
```

`Array.fold: f:('State -> 'T -> 'State) -> elm:'State -> arr:'T [] -> 'State`. Up-

dates an accumulator iteratively by applying `f` to each element in `arr`. The initial value of the accumulator is `elm`. For example, when `arr` consists of `n+1` elements `Array.fold` calculates:

$$f\ (\dots (f\ (f\ elm\ arr.[0])\ arr.[1])\ \dots)\ arr.[n].$$

Listing 11.59: `Array.fold`

```
1  > let addSquares acc elm = acc + elm*elm
2  - Array.fold addSquares 0 [|0 .. 9|];;
3  val addSquares : acc:int -> elm:int -> int
4  val it : int = 285
```

`Array.foldBack: f:('T -> 'State -> 'State) -> arr:'T [] -> elm:'State -> 'State.` Updates an accumulator iteratively backwards by applying `f` to each element in `arr`. The initial value of the accumulator is `elm`. For exampel, when `arr` consists of `n+1` elements `List.foldBack` calculates:

$$f\ arr.[0]\ (f\ arr.[1]\ (\dots(f\ arr.[n]\ elm)\ \dots)).$$

Listing 11.60: `Array.foldBack`

```
1  > let addSquares elm acc = acc + elm*elm
2  - Array.foldBack addSquares [|0 .. 9|] 0;;
3  val addSquares : elm:int -> acc:int -> int
4  val it : int = 285
```

`Array.forall: f:('T -> bool) -> arr:'T [] -> bool.` Returns true if `f` evaluates to true for every element in `arr`.

Listing 11.61: `Array.forall`

```
1  > Array.forall (fun x -> (x % 2 = 1)) [|0 .. 9|];;
2  val it : bool = false
```

`Array.init: m:int -> f:(int -> 'T) -> 'T [].` Create an array with `m` elements and whose value is the result of applying `f` to the index of the element.

Listing 11.62: `Array.init`

```
1  > Array.init 10 (fun i -> i * i);;
2  val it : int [] = [|0; 1; 4; 9; 16; 25; 36; 49; 64; 81|]
```

`Array.isEmpty: arr:'T [] -> bool.` Returns true if `arr` is empty.

Listing 11.63: `Array.isEmpty`

```
1  > Array.isEmpty [||];;
2  val it : bool = true
```

`Array.iter: f:('T -> unit) -> arr:'T [] -> unit`. Applies `f` to each element of `arr`.

```
Listing 11.64: Array.iter

1  > Array.iter (fun x -> printfn "%A " x) [|0; 1; 2|];;
2  0
3  1
4  2
5  val it : unit = ()
```

`Array.map: f:('T -> 'U) -> arr:'T [] -> 'U []`. Creates an new array whose elements are the results of applying `f` to each of the elements of `arr`.

```
Listing 11.65: Array.map

1  > Array.map (fun x -> x * x) [|0 .. 9|];;
2  val it : int [] = [|0; 1; 4; 9; 16; 25; 36; 49; 64; 81|]
```

`Array.ofList: lst:'T list -> 'T []`. Creates an array whose elements are copied from `lst`.

```
Listing 11.66: Array.ofList

1  > Array.ofList [1; 2; 3];;
2  val it : int [] = [|1; 2; 3|]
```

`Array.rev: arr:'T [] -> 'T []`. Creates a new array whose elements are identical to `arr` but in reverse order.

```
Listing 11.67: Array.rev

1  > Array.rev [|1; 2; 3|];;
2  val it : int [] = [|3; 2; 1|]
```

`Array.sort: arr:'T[] -> 'T []`. Creates a new array with the same elements as in `arr` but in sorted order

```
Listing 11.68: Array.sort

1  > Array.sort [|3; 1; 2|];;
2  val it : int [] = [|1; 2; 3|]
```

`Array.toList: arr:'T [] -> 'T list`. Creates a new list whose elements are copied from `arr`.

Listing 11.69: `Array.toList`

```
1  > Array.toList [|1; 2; 3|];;
2  val it : int list = [1; 2; 3]
```

`Array.unzip: arr:('T1 * 'T2) [] -> 'T1 [] * 'T2 []`. Returns a pair of arrays of all the first elements and all the second elements of `arr`, respectively.

Listing 11.70: `Array.unzip`

```
1  > Array.unzip [|(1, 'a'); (2, 'b'); (3, 'c')|];;
2  val it : int [] * char [] = ([|1; 2; 3|], [|'a'; 'b';
     'c'|])
```

`Array.zip: arr1:'T1 [] -> arr2:'T2 [] -> ('T1 * 'T2) []`. Returns a list of pairs, where elements in `arr1` and `arr2` are iteratively paired.

Listing 11.71: `Array.zip`

```
1  > Array.zip [|1; 2; 3|] [|'a'; 'b'; 'c'|];;
2  val it : (int * char) [] = [|(1, 'a'); (2, 'b'); (3, 'c')|]
```

11.4. Multidimensional Arrays

Multidimensional arrays can be created as arrays of arrays (of arrays . . .). These are known as *jagged arrays*, since there is no inherent guarantee that all sub-arrays are of the same size. The example in Listing 11.72 is a jagged array of increasing width.

· multidimensional arrays

· jagged arrays

Listing 11.72 arrayJagged.fsx:
An array of arrays. When row lengths are of non-equal elements, then it is a jagged array.

```
1  let arr = [|[|1|]; [|1; 2|]; [|1; 2; 3|]|]
2
3  for row in arr do
4    for elm in row do
5      printf "%A " elm
6    printf "\n"
```
```
1  $ fsharpc --nologo arrayJagged.fsx && mono arrayJagged.exe
2  1
3  1 2
4  1 2 3
```

Indexing arrays of arrays is done sequentially, in the sense that in the above example, the number of outer arrays is `a.Length`, `a.[i]` is the i'th array, the length of the i'th array is `a.[i].Length`, and the j'th element of the i'th array is thus `a.[i].[j]`. Often

2-dimensional rectangular arrays are used, which can be implemented as a jagged array, as shown in Listing 11.73.

Listing 11.73 arrayJaggedSquare.fsx:
A rectangular array.

```
let pownArray  (arr : int array array) p =
  for i = 1 to arr.Length - 1 do
    for j = 1 to arr.[i].Length - 1 do
      arr.[i].[j] <- pown arr.[i].[j] p

let printArrayOfArrays (arr : int array array) =
  for row in arr do
    for elm in row do
      printf "%3d " elm
    printf "\n"

let A = [|[|1 .. 4|]; [|1 .. 2 .. 7|]; [|1 .. 3 .. 10|]|]
pownArray A 2
printArrayOfArrays A
```

```
$ fsharpc --nologo arrayJaggedSquare.fsx && mono
  arrayJaggedSquare.exe
  1   2   3   4
  1   9  25  49
  1  16  49 100
```

Note that the for-in cannot be used in `pownArray`, e.g.,

> for row in arr do for elm in row do elm <- pown elm p done done,

since the iterator value `elm` is not mutable, even though `arr` is an array.

Square arrays of dimensions 2 to 4 are so common that F# has built-in modules for their support. Here, we will describe *Array2D*. The workings of *Array3D* and *Array4D* are very similar. A generic `Array2D` has type `'T [,]`, and it is indexed also using the `[,]` notation. The `Array2D.length1` and `Array2D.length2` functions are supplied by the `Array2D` module for obtaining the size of an array along the first and second dimension. Rewriting the with jagged array example in Listing 11.73 to use `Array2D` gives a slightly simpler program, which is shown in Listing 11.74.

· Array2D
· Array3D
· Array4D

Listing 11.74 array2D.fsx:
Creating a 3 by 4 rectangular array of integers.

```
1  let arr = Array2D.create 3 4 0
2  for i = 0 to (Array2D.length1 arr) - 1 do
3    for j = 0 to (Array2D.length2 arr) - 1 do
4      arr.[i,j] <- j * Array2D.length1 arr + i
5  printfn "%A" arr
```

```
1  $ fsharpc --nologo array2D.fsx && mono array2D.exe
2  [[0; 3; 6; 9]
3   [1; 4; 7; 10]
4   [2; 5; 8; 11]]
```

Note that the `printf` supports direct printing of the 2-dimensional array. `Array2D` arrays support slicing. The "*" lexeme is particularly useful to obtain all values along a dimension. This is demonstrated in Listing 11.75.

Listing 11.75: Examples of Array2D slicing. Compare with Listing 11.74.

```
1  > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j);;
2  val arr : int [,] = [[0; 10; 20; 30]
3                       [1; 11; 21; 31]
4                       [2; 12; 22; 32]]
5
6  > arr.[2,3];;
7  val it : int = 32
8
9  > arr.[1..,3..];;
10 val it : int [,] = [[31]
11                     [32]]
12
13 > arr.[..1,*];;
14 val it : int [,] = [[0; 10; 20; 30]
15                     [1; 11; 21; 31]]
16
17 > arr.[1,*];;
18 val it : int [] = [|1; 11; 21; 31|]
19
20 > arr.[1..1,*];;
21 val it : int [,] = [[1; 11; 21; 31]]
```

Note that in almost all cases, slicing produces a sub-rectangular 2 dimensional array, except for `arr.[1,*]`, which is an array, as can be seen by the single "[". In contrast, `A.[1..1,*]` is an Array2D. Note also that `printfn` typesets 2 dimensional arrays as `[[ ... ]]` and not `[|[| ... |]|]`, which can cause confusion with lists of lists.

Multidimensional arrays have the same properties and methods as arrays, see Section 11.3.1.

11.4.1. The Array2D Module

There are quite a number of built-in procedures for arrays in the **Array2D** namespace, some of which are summarized below.[3]

copy: arr:'T [,] -> 'T [,]. Creates a new array whose elements are copied from **arr**.

```
Listing 11.76: Array2D.copy

1  > let a = Array2D.init 3 4 (fun i j -> i + 10 * j)
2  - let b = Array2D.copy a;;
3  val a : int [,] = [[0; 10; 20; 30]
4                     [1; 11; 21; 31]
5                     [2; 12; 22; 32]]
6  val b : int [,] = [[0; 10; 20; 30]
7                     [1; 11; 21; 31]
8                     [2; 12; 22; 32]]
```

create: m:int -> n:int -> v:'T -> 'T [,]. Creates an **m** by **n** array whose elements are set to **v**.

```
Listing 11.77: Array2D.create

1  > Array2D.create 2 3 3.14;;
2  val it : float [,] = [[3.14; 3.14; 3.14]
3                        [3.14; 3.14; 3.14]]
```

init: m:int -> n:int -> f:(int -> int -> 'T) -> 'T [,]. Creates an **m** by **n** array whose elements are the result of applying **f** to the index of an element.

```
Listing 11.78: Array2D.init

1  > Array2D.init 3 4 (fun i j -> i + 10 * j);;
2  val it : int [,] = [[0; 10; 20; 30]
3                      [1; 11; 21; 31]
4                      [2; 12; 22; 32]]
```

iter: f:('T -> unit) -> arr:'T [,] -> unit. Applies **f** to each element of **arr**.

```
Listing 11.79: Array2D.iter

1  > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
2  - Array2D.iter (fun elm -> printf "%A " elm) arr
3  - printfn "";;
4  0 10 20 30 1 11 21 31 2 12 22 32
5  val arr : int [,] = [[0; 10; 20; 30]
6                       [1; 11; 21; 31]
7                       [2; 12; 22; 32]]
8  val it : unit = ()
```

---

[3]Jon: **rewrite description**

`length1: arr:'T [,] -> int`. Returns the length the first dimension of `arr`.

> **Listing 11.80:** `Array2D.length1`
>
> ```
> 1  > let arr = Array2D.create 2 3 0.0 in Array2D.length1 arr;;
> 2  val it : int = 2
> ```

`length2: arr:'T [,] -> int`. Returns the length of the second dimension of `arr`.

> **Listing 11.81:** `Array2D.forall length2`
>
> ```
> 1  > let arr = Array2D.create 2 3 0.0 in Array2D.length2 arr;;
> 2  val it : int = 3
> ```

`map: f:('T -> 'U) -> arr:'T [,] -> 'U [,]`. Creates a new array whose elements are the results of applying `f` to each of the elements of `arr`.

> **Listing 11.82:** `Array2D.map`
>
> ```
> 1  > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
> 2  - Array2D.map (fun x -> x * x) arr;;
> 3  val arr : int [,] = [[0; 10; 20; 30]
> 4                       [1; 11; 21; 31]
> 5                       [2; 12; 22; 32]]
> 6  val it : int [,] = [[0; 100; 400; 900]
> 7                      [1; 121; 441; 961]
> 8                      [4; 144; 484; 1024]]
> ```

12 | The Imperative Programming paradigm

Imperative programming is a paradigm for programming states. In imperative programming, the focus is on how a problem is to be solved, as a list of *statements* that affects *states*. In F#, states are mutable and immutable values, and they are affected by functions and procedures. An imperative program is typically identified as using:

Mutable values
   Mutable values are holders of states, they may change over time, and thus have a dynamic scope.

Procedures
   Procedures are functions that returns "()", as opposed to functions that transform data. They are the embodiment of side-effects.

Side-effects
   Side-effects are changes of state that are not reflected in the arguments and return values of a function. The `printf` is an example of a procedure that uses side-effects to communicate with the terminal.

Loops
   The `for`- and `while`-loops typically use an iteration value to update some state, e.g., `for`-loops are often used to iterate through a list and summarize its contents.

Mono state or stateless programs, as *functional programming*, can be seen as a subset of imperative programming and is discussed in Chapter 17. *Object-oriented programming* is an extension of imperative programming, where statements and states are grouped into classes. For a discussion on object-oriented programming, see Chapter 22.

An imperative program is like a Turing machine, a theoretical machine introduced by Alan Turing in 1936 [10]. Almost all computer hardware is designed for *machine code*, which is a common term used for many low-level computer programming languages, and almost all machine languages follow the imperative programming paradigm.

A prototypical example is a baking recipe, e.g., to make a loaf of bread, do the following:

1. Mix yeast with water.

2. Stir in salt, oil, and flour.

3. Knead until the dough has a smooth surface.

4. Let the dough rise until it has doubled its size.

5. Shape dough into a loaf.

6. Let the loaf rise until double size.

7. Bake in the oven until the bread is golden brown.

Each line in this example consists of one or more statements that are to be executed, and while executing them, states such as the size of the dough and the color of the bread changes. Some execution will halt execution until certain conditions of these states are fulfilled, e.g., the bread will not be put into the oven for baking before it has risen sufficiently.

12.1. Imperative Design

Programming is the act of solving a problem by writing a program to be executed on a computer. The imperative programming paradigm focuses on states. To solve a problem, you could work through the following list of actions:

1. Understand the problem. As Pólya described it, see Chapter 2, the first step in any solution is to understand the problem. A good trick to check whether you understand the problem, is to briefly describe it in your own words.

2. Identify the main values, variables, functions, and procedures needed. If the list of procedures is large, then you most likely should organize them in modules.

3. For each function and procedure, write a precise description of what it should do. This can conveniently be performed as an in-code comment for the procedure, using the F# XML documentation standard.

4. Make mockup functions and procedures using the intended types, but do not necessarily compute anything sensible. Run through examples in your mind, using this mockup program to identify any obvious oversights.

5. Write a suite of unit tests that tests the basic requirements for your code. The unit tests should be runnable with your mockup code. Writing unit tests will also allow you to evaluate the usefulness of the code pieces as seen from an application point of view.

6. Replace the mockup functions in a prioritized order, i.e., write the must-have code before you write the nice-to-have code, while regularly running your unit tests to keep track of your progress.

7. Evaluate the code in relation to the desired goal, and reiterate earlier actions as needed until the task has been sufficiently completed.

8. Complete your documentation both in-code and outside to ensure that the intended user has sufficient knowledge to effectively use your program and to ensure that you or a fellow programmer will be able to maintain and extend the program in the future.

13 | Recursion

Recursion is a central concept in F# and is used to control flow in loops without the `for` and `while` constructions. Figure 13.1 illustrates the concept of an infinite loop with recursion.

Figure 13.1.: An infinitely long string of "forever forever forever...", conceptually calculated by `let rec` forever () = `"forever "` + (forever ()).

13.1. Recursive Functions

A *recursive function* is a function which calls itself, and the syntax for defining recursive functions is an extension of that for regular functions:

· recursive function

> **Listing 13.1:** Syntax for defining one or more mutually dependent recursive functions.
>
> ```
> 1  let rec <ident> = <expr> {and <ident> = <expr>} [in] <expr>
> ```

From a compiler point of view, the `rec` is necessary, since the function is used before the compiler has completed its analysis. If two functions are mutually recursive, then they must be defined jointly using the `and` keyword.

· rec

· and

An example of a recursive function that counts from 1 to 10 similarly to Listing 8.5 is given

in Listing 13.2.

Listing 13.2 countRecursive.fsx:
Counting to 10 using recursion.

```
1  let rec prt a b =
2    if a > b then
3        printf "\n"
4    else
5        printf "%d " a
6        prt (a + 1) b
7
8  prt 1 10
```

```
1  $ fsharpc --nologo countRecursive.fsx && mono
     countRecursive.exe
2  1 2 3 4 5 6 7 8 9 10
```

Here the `prt` function calls itself repeatedly, such that the first call is `prt 1 10`, which calls `prt 2 10`, and so on until the last call `prt 11 10`. Each time `prt` is called, new bindings named `a` and `b` are made to new values. This is illustrated in Figure 13.2. The old values are no longer accessible, as indicated by subscripts in the figure. E.g., in prt_3, the scope has access to a_3 but not a_2 and a_1. Thus, in this program, the process is similar to a `for` loop, where the counter is `a`, and in each loop its value is reduced.

The structure of the function is typical for recursive functions. They very often follow the following pattern.

Listing 13.3: Recursive functions consist of a stopping criterium, a stopping expression, and a recursive step.

```
1  let rec f a =
2    if <stopping condition>
3    then <stopping step>
4    else <recursion step>
```

The *match − with* are also very common conditional structures. In Listing 13.2, `a > b` is the *stopping condition*, `printfn "\n"` is the *stopping step*, and `printfn "\%d " a; prt (a + 1) b` is the *recursion step*.

· match
· with
· stopping condition
· stopping step
· recursion step

13.2. The Call Stack and Tail Recursion

Fibonacci's sequence of numbers is a recursive sequence of numbers with relations to the Golden ratio and structures in biology. The Fibonacci sequence is the sequence of numbers $1, 1, 2, 3, 5, 8, 13, \ldots$. The sequence starts with $1, 1$ and the next number is recursively given as the sum of the two previous ones. A direct implementation of this is given in Listing 8.7.

```
$ fsharpc countRecursive.fsx && mono countRecursive.exe

prt 1 10

  prt₁: a₁= 1, b₁= 10

    1 > 10: ✗
       printf "1 "
       prt 2 10

         prt₂: a₂= 2, b₂= 10

           2 > 10: ✗
              printf "2 "
              prt 3 10

                prt₃: a₃= 3, b₃= 10

                  3 > 10: ✗
                     printf "3 "
                     prt 4 10
                        ⋱
                          prt 11 10

                            prt₁₁: a₁₁= 1, b₁₁= 10

                              11 > 10 ✓
                                 printf "\\n"
                              - - - - - - - - - - - - - - - -
                              ()

                        ⋰
                  - - - - - - - - - - - - - - - - - - - -
                  ()

           - - - - - - - - - - - - - - - - - - - - - - - -
           ()

    - - - - - - - - - - - - - - - - - - - - - - - - - - -
    ()

- - - - - - - - - - - - - - - - - - - - - - - - - - - - -
()
```

Figure 13.2.: Illustration of the recursion used to write the sequence "1 2 3 ... 10" in line 8 in Listing 13.2. Each frame corresponds to a call to `prt`, where new values overshadow old ones. All calls return `unit`.


Listing 13.4 fibRecursive.fsx:
The *n*'th Fibonacci number using recursion.

```
1  let rec fib n =
2     if n < 1 then
3        0
4     elif n = 1 then
5        1
6     else
7        fib (n - 1) + fib (n - 2)
8
9  for i = 0 to 10 do
10    printfn "fib(%d) = %d" i (fib i)
```
```
1  $ fsharpc --nologo fibRecursive.fsx && mono fibRecursive.exe
2  fib(0) = 0
3  fib(1) = 1
4  fib(2) = 1
5  fib(3) = 2
6  fib(4) = 3
7  fib(5) = 5
8  fib(6) = 8
9  fib(7) = 13
10 fib(8) = 21
11 fib(9) = 34
12 fib(10) = 55
```

Here we extended the sequence to $0, 1, 1, 2, 3, 5, \ldots$ with the starting sequence $0, 1$, allowing us to define all fib$(n) = 0$, $n < 1$. Thus, our function is defined for all integers, and for the irrelevant negative arguments it fails gracefully by returning 0. This is a general piece of advice: **make functions that fail gracefully.**   Advice

A visualization of the calls and the scopes created by `fibRecursive` is shown in Figure 13.3. The figure illustrates that each recursive step results in two calls to the function, thus creating two new scopes. And it gets worse. Figure 13.4 illustrates the tree of calls for `fib 5`. Thus, a call to the function `fib` generates a tree of calls that is five levels deep and has `fib(5)` number of nodes. In general for the program in Listing 13.4, a call to `fib(n)` produces a tree with `fib(n)` $\leq c\alpha^n$ calls to the function for some positive constant c and $\alpha \geq \frac{1+\sqrt{5}}{2} \sim 1.6$[1]. Each call takes time and requires memory, and we have thus created a slow and somewhat memory-intensive function. This is a hugely ineffective implementation of calculating entries into Fibonacci's sequence, since many of the calls are identical. E.g., in Figure 13.4, `fib 1` is called five times. Before we examine a faster algorithm, we first need to discuss how F# executes function calls.

When a function is called, then memory is dynamically allocated internally for the function on what is known as the *call stack*. Stacks are used for many things in programming, but typically the call stack is considered special, since it is almost always implicitly part of any program execution. Hence, it is often just referred to as *The Stack*. When a function is called, a new *stack frame* is stacked (pushed) on the call stack, including its arguments, local storage such as mutable values, and where execution should return to when the function is finished. When the function finishes, the stack frame is unstacked (popped) and in its stead, the return value of the function is stacked. This return value is then unstacked

· call stack

· The Stack

· stack frame

---

[1]Jon: `https://math.stackexchange.com/questions/674533/prove-upper-bound-big-o-for-fibonaccis-sequence`

```
$ fsharpc fibRecursive.fsx && mono fibRecursive.exe
```

fib 3

fib$_1$: n$_1$= 3

3 < 1: ✗
  3 = 1: ✗
    fib 2 + fib 1

fib$_2$: n$_2$= 2

2 < 1: ✗
 2 = 1: ✗
  fib 1 + fib 0

fib$_4$: n$_4$= 1

1 < 1: ✗
 1 = 1:
✓

1

fib$_5$: n$_5$= 0

0 < 1: ✓

0

1

fib$_3$: n$_3$ = 1

1 < 1: ✗
 1 = 1: ✓

1

2

3

Figure 13.3.: Illustration of the recursion used to write the sequence "1 2 3 ... 10" in line 8 in Listing 13.2. Each frame corresponds to a call to `fib`, where new values overshadow old ones.

```
                                fib 5
                     fib 4              fib 3
                fib 3    fib 2      fib 2    fib 1
            fib 2  fib 1  fib 1 fib 0  fib 1 fib 0
        fib 1 fib 0
```

Figure 13.4.: The function calls involved in calling `fib 5`.

Figure 13.5.: A call to `fib 5` in Listing 13.4 starts a sequence of function calls and stack frames on the call stack.

and used by the caller. After unstacking the return value, the call stack is identical to its state prior to the call. Figure 13.5 shows snapshots of the call stack when calling `fib 5` in Listing 13.4. The call first stacks a frame onto the call stack with everything needed to execute the function body plus a reference to where the return to, when the execution is finished. Then the body of `fib` is executed, which includes calling `fib 4` and `fib 3` in turn. The call to `fib 4` stacks a frame onto the call stack, and its body is executed. Once execution is returned from the call to `fib 4`, the result of the function is on top of the stack. It is unstacked, saved and the call to `fib 3` is treated equally. When the end of `fib 5` is reached, its frame is unstacked, and its result is stacked. In this way, the call stack is returned to its original state except for the result of the function, and execution is returned to the point right after the original call to `fib 5`. Thus, for Listing 13.4 $\mathcal{O}\left(\alpha^n\right)$, $\alpha = \frac{1+\sqrt{5}}{2}$ stacking operations are performed for a call to `fib n`. The $\mathcal{O}\left(f(n)\right)$ is the *Landau symbol* used to denote the order of a function, such that if $g(n) = \mathcal{O}(f(n))$ then there exists two real numbers $M > 0$ and a n_0 such that for all $n \geq n_0$, $|g(n)| \leq M|f(n)|$.[2] As indicated by the tree in Figure 13.4, the call tree is at most n high, which corresponds to a maximum of n additional stack frames as compared to the starting point.

· Landau symbol

The implementation of Fibonacci's sequence in Listing 13.4 can be improved to run faster and use less memory. One such algorithm is given in Listing 13.5

Listing 13.5 fibRecursiveAlt.fsx:
A fast, recursive implementation of Fibonacci's numbers. Compare with Listing 13.4.

```
1  let fib n =
2    let rec fibPair n pair =
3      if n < 2 then pair
4      else fibPair (n - 1) (snd pair, fst pair + snd pair)
5    if n < 1 then 0
6    elif n = 1 then 1
7    else fibPair n (0, 1) |> snd
8
9  printfn "fib(10) = %d" (fib 10)
```

```
1  $ fsharpc --nologo fibRecursiveAlt.fsx && mono
     fibRecursiveAlt.exe
2  fib(10) = 55
```

Calculating the 45th Fibonacci number a MacBook Pro, with a 2.9 GHz Intel Core i5 using Listing 13.4 takes about 11.2s while using Listing 13.5 is about 224 times faster and only

---

[2]Jon: **Introduction of Landau notation needs to be moved earlier, since it used in Collections chapter.**

takes 0.050s. The reason is that `fib` in Listing 13.5 calculates every number in the sequence once and only once by processing the list recursively while maintaining the previous two values needed to calculate the next in the sequence. I.e., the function `helper` transforms the pair `(a,b)` to `(b,a+b)` such that, e.g., the 4th and 5th pair `(3,5)` is transformed into the 5th and the 6th pair `(5,8)` in the sequence. What complicates the algorithm is that besides the transformation, we must keep track of when to stop, which here is done using a counter variable, that is recursively reduced by 1 until our stopping criterium.

Listing 13.5 also uses much less memory than Listing 13.4, since its recursive call is the last expression in the function, and since the return value of two recursive calls to `helper` is the same as the return value of the last. In fact, the return value of any number of recursive calls to `helper` is the return value of the last. This structure is called *tail-recursion*. · tail-recursion
Compilers can easily optimize the call stack usage for tail recursion, since when in this example `helper` calls itself, then its frame is no longer needed, and may be replaced by the new `helper` with the slight modification, that the return point should be to `fib` and not the end of the previous `helper`. Once the recursion reaches the stopping criteria, then instead of popping a long list of calls of `helper` frames, then there is only one, and the return value is equal to the return value of the last call and the return point is to `fib`. Thus, many stack frames in tail recursion are replaced by one. Hence, **prefer tail-recursion** Advice
whenever possible.

13.3. Mutually Recursive Functions

Functions that recursively call each other are called *mutually recursive* functions. F# · mutually recursive
offers the `let – rec – and` notation for co-defining mutually recursive functions. As an · `let`
example, consider the function `even : int -> bool`, which returns true if its argument · `rec`
is even and false otherwise, and the opposite function `odd : int -> bool`. A mutually · `and`
recursive implementation of these functions can be developed from the following relations:
`even 0 = `true`, `odd 0 = `false`, and for $n > 0$, `even n = odd (n-1)`, which implies that
for $n > 0$, `odd n = even (n-1)`:

Listing 13.6 mutuallyRecursive.fsx:
Using mutual recursion to implement even and odd functions.

```
1  let rec even x =
2    if x = 0 then true
3    else odd (x - 1)
4  and odd x =
5    if x = 0 then false
6    else even (x - 1);;
7
8  let w = 5;
9  printfn "%*s %*s %*s" w "i" w "even" w "odd"
10 for i = 1 to w do
11   printfn "%*d %*b %*b" w i w (even i) w (odd i)
```

```
1  $ fsharpc --nologo mutuallyRecursive.fsx && mono
     mutuallyRecursive.exe
2      i  even   odd
3      1 false   true
4      2  true false
5      3 false   true
6      4  true false
7      5 false   true
```

Notice that in the lightweight notation the `and` must be on the same indentation level as the original `let`.

Without the `and` keyword, F# will issue a compile error at the definition of even. However, it is possible to implement mutual recursion by using functions as an argument, e.g.,

> **Listing 13.7 mutuallyRecursiveAlt.fsx:**
> **Mutual recursion without the `and` keyword requires a helper function.**
>
> ```
> let rec evenHelper (notEven: int -> bool) x =
>   if x = 0 then true
>   else notEven (x - 1)
>
> let rec odd x =
>   if x = 0 then false
>   else evenHelper odd (x - 1);;
>
> let even x = evenHelper odd x
>
> let w = 5;
> printfn "%*s %*s %*s" w "i" w "Even" w "Odd"
> for i = 1 to w do
>   printfn "%*d %*b %*b" w i w (even i) w (odd i)
> ```
>
> ```
> $ fsharpc --nologo mutuallyRecursiveAlt.fsx
> $ mono mutuallyRecursiveAlt.exe
>    i  Even   Odd
>    1 false  true
>    2  true false
>    3 false  true
>    4  true false
>    5 false  true
> ```

This being said, Listing 13.6 is clearly to be preferred over Listing 13.7.

In the example above, we used the **even** and **odd** function problems to demonstrate mutual recursion. There is, of course, a much simpler solution, which does not use recursion at all:

> **Listing 13.8 parity.fsx:**
> **A better way to test for parity without recursion.**
>
> ```
> let even x = (x % 2 = 0)
> let odd x = not (even x)
> ```

This is to be preferred anytime as the solution to the problem. [3]

---

[3]Jon: **Here it would be nice to have an intermezzo, giving examples of how to write a recursive program by thinking the problem has been solved.**

14 | Programming with Types

F# is a strongly typed language, meaning that types are known or inferred at compile time. In the previous chapters, we have used *primitive types* such as `float` and `bool`, · primitive types function types, and compound types implicitly defined by tuples. These types are used for simple programming tasks, and everything that can be programmed can be accomplished using these types. However, larger programs are often easier to read and write when using more complicated type structures. In this chapter, we will discuss type abbreviations, enumerated types, discriminated unions, records, and structs. Class types are discussed in depth in Chapter 20.

14.1. Type Abbreviations

F# allows for renaming of types, which is called *type abbreviation* or *type aliasing*. The · type abbreviation syntax is: · type aliasing

Listing 14.1: Syntax for type abbreviation.

```
type <ident> = <aType>
```

where the identifier is a new name, and the type-name is an existing type or a compound of existing types. Listing 14.2 shows examples of the defintion of several type abbreviations.

Listing 14.2 typeAbbreviation.fsx:
Defining three type abbreviations, two of which are compound types.

```
type size = int
type position = float * float
type person = string * int
type intToFloat = int -> float

let sz : size = 3
let pos : position = (2.5, -3.2)
let pers : person = ("Jon", 50)
let conv : intToFloat = fun a -> float a
printfn "%A, %A, %A, %A" sz pos pers (conv 2)
```
```
$ fsharpc --nologo typeAbbreviation.fsx && mono
    typeAbbreviation.exe
3, (2.5, -3.2), ("Jon", 50), 2.0
```

Here we define the abbreviations `size`, `position`, `person`, and `intToFloat`, and later make bindings enforcing the usage of these abbreviations.

Type abbreviations are used as short abbreviations of longer types, and they add semantic content to the program text, thus making programs shorter and easier to read. Type abbreviations allow the programmer to focus on the intended structure at a higher level by, e.g., programming in terms of a type `position` rather than `float * float`. Thus, they often result in programs with fewer errors. Type abbreviations also make maintenance easier. For instance, if we at a later stage decide that positions can only have integer values, then we only need to change the definition of the type abbreviation, not every place a value of type `position` is used.

14.2. Enumerations

Enumerations or *enums* for short are types with named values. Names in enums are assigned to a subset of integer or char values. Their syntax is as follows:

· enumerations
· enums

> **Listing 14.3: Syntax for enumerations.**
> ```
> 1  type <ident> =
> 2    [ | ] <ident> = <integerOrChar>
> 3    | <ident> = <integerOrChar>
> 4    | <ident> = <integerOrChar>
> 5    ...
> ```

An example of using enumerations is given in Listing 14.4.

> **Listing 14.4 enum.fsx:**
> **An enum type acts as a typed alias to a set of integers or chars.**
> ```
> 1  type medal =
> 2    Gold = 0
> 3    | Silver = 1
> 4    | Bronze = 2
> 5
> 6  let aMedal = medal.Gold
> 7  printfn "%A has value %d" aMedal (int aMedal)
> ```
> ```
> 1  $ fsharpc --nologo enum.fsx && mono enum.exe
> 2  Gold has value 0
> ```

In this example, we define an enumerated type for medals, which allows us to work with the names rather than the values. Since the values most often are arbitrary, we can program using semantically meaningful names instead. Being able to refer to an underlying integer type allows us to interface with other – typically low-level – programs that require integers, and to perform arithmetic. E.g., for the medal example, we can typecast the enumerated types to integers and calculate an average medal harvest.

14.3. Discriminated Unions

A discriminated union is a union of a set of named cases. These cases can further be of specified types. The syntax for defining a discriminated union is as follows:

Listing 14.5:   Syntax for type abbreviation.

```
1  [<attributes>]
2  type <ident> =
3    [| ]<ident> [of [<ident> :] <aType> [* [<ident> :] <aType>
     ...]]
4    | <ident> [of [<ident> :] <aType> [* [<ident> :] <aType> ...]]
5    ...
```

Discriminated unions are reference types, i.e., their content is stored on *The Heap*, see   · The Heap
Section 6.8 for a discussion on reference types. Since they are immutable, there is no risk of side-effects. As reference types, they only pass a reference when used as arguments to and returned from a function. This is in contrast to value types, which transport a complete copy of the data structure. Discriminated unions are thus effective for large data structures. Discriminated unions can also be represented as structures using the `[<Struct>]` attribute, in which case they are value types. See Section 14.5 for a discussion on structs.

An example just using the named cases but no further specification of types is given in Listing 14.6.

Listing 14.6 discriminatedUnions.fsx:
A discriminated union of medals. Compare with Listing 14.4.

```
1  type medal =
2    Gold
3    | Silver
4    | Bronze
5
6  let aMedal = medal.Gold
7  printfn "%A" aMedal
```

```
1  $ fsharpc --nologo discriminatedUnions.fsx
2  $ mono discriminatedUnions.exe
3  Gold
```

Here, we define a discriminated union as three named cases signifying three different types of medals. Comparing with the enumerated type in Listing 14.4, we see that the only difference is that the cases of the discriminated unions have no value. A commonly used discriminated union is the *option type*, see Section 18.2 for more detail.   · option type

Discriminated unions may also be used to store data. Where the names in enumerated types are aliases of single values, the names used in discriminated unions can hold any value specified at the time of creation. An example is given in Listing 14.7.

Figure 14.1.: The tree with 3 leaves.

Listing 14.7 discriminatedUnionsOf.fsx:
A discriminated union using explicit subtypes.

```
1  type vector =
2    Vec2D of float * float
3    | Vec3D of x : float * y : float * z : float
4
5  let v2 = Vec2D (1.0, -1.2)
6  let v3 = Vec3D (x = 1.0, z = -1.2, y = 0.9)
7  printfn "%A and %A" v2 v3
```

```
1  $ fsharpc --nologo discriminatedUnionsOf.fsx
2  $ mono discriminatedUnionsOf.exe
3  Vec2D (1.0,-1.2) and Vec3D (1.0,0.9,-1.2)
```

In this case, we define a discriminated union of two and three-dimensional vectors. Values of these types are created using their names followed by a tuple of their arguments. The names are also called field names. The field names may be used when creating discrimated union values, as shown in Line 6. When used, then the arguments may be given in arbitrary order, nevertheless, values for all fields must be given.

Discriminated unions can be defined recursively. This feature is demonstrated in Listing 14.8.

Listing 14.8 discriminatedUnionTree.fsx:
A discriminated union modelling binary trees.

```
1  type Tree =
2    Leaf of int
3    | Node of Tree * Tree
4
5  let one = Leaf 1
6  let two = Leaf 2
7  let three = Leaf 3
8  let tree = Node (Node (one, two), three)
9  printfn "%A" tree
```

```
1  $ fsharpc --nologo discriminatedUnionTree.fsx
2  $ mono discriminatedUnionTree.exe
3  Node (Node (Leaf 1,Leaf 2),Leaf 3)
```

In this example we define a tree as depicted in Figure 14.1.

Pattern matching must be used in order to define functions on values of a discriminated union. E.g., in Listing 14.9 we define a function that traverses a tree and prints the content of the nodes.[1]

Listing 14.9 discriminatedUnionPatternMatching.fsx: A discriminated union modelling binary trees.

```
type Tree = Leaf of int | Node of Tree * Tree
let rec traverse (t : Tree) : string =
    match t with
      Leaf(v) -> string v
      | Node(left, right) -> (traverse left) + ", " +
   (traverse right)

let tree = Node (Node (Leaf 1, Leaf 2), Leaf 3)
printfn "%A: %s" tree (traverse tree)
```

```
$ fsharpc --nologo discriminatedUnionPatternMatching.fsx
$ mono discriminatedUnionPatternMatching.exe
Node (Node (Leaf 1,Leaf 2),Leaf 3): 1, 2, 3
```

Discriminated unions are very powerful and can often be used instead of class hierarchies. Class hierarchies are discussed in Section 21.1.

14.4. Records

A record is a compound of named values, and a record type is defined as follows:

Listing 14.10:   Syntax for defining record types.

```
[ <attributes> ]
type <ident> = {
  [ mutable ] <label1> : <type1>
  [ mutable ] <label2> : <type2>
  ...
}
```

Records are collections of named variables and values of possibly different types. They are reference types, and thus their content is stored on *The Heap*, see Section 6.8 for a discussion on reference types. Records can also be *struct records* using the `[<Struct>]` attribute, in which case they are value types. See Section 14.5 for a discussion on structs. An example of using records is given in Listing 14.11. The values of individual members of a record are obtained using the *"."* notation

· The Heap

· struct records

· .

---

[1]Jon: **Example uses pattern matching, which has yet to be introduced.**

Listing 14.11 records.fsx:
A record is defined for holding information about a person.

```
1  type person = {
2      name : string
3      age : int
4      height : float
5  }
6
7  let author = {name = "Jon"; age = 50; height = 1.75}
8  printfn "%A\nname = %s" author author.name
```

```
1  $ fsharpc --nologo records.fsx && mono records.exe
2  {name = "Jon";
3   age = 50;
4   height = 1.75;}
5  name = Jon
```

This examples illustrate a how record type is defined to store varied data about a person, and how a value is created by a record expression defining its field values.

If two record types are defined with the same label set, then the latter dominates the former. This is demonstrated in Listing 14.12.

Listing 14.12 recordsDominance.fsx:
Redefined types dominate old record types, but earlier definitions are still accessible using explicit or implicit specification for bindings.

```
1  type person = { name : string; age : int; height : float }
2  type teacher = { name : string; age : int; height : float }
3
4  let lecturer = {name = "Jon"; age = 50; height = 1.75}
5  printfn "%A : %A" lecturer (lecturer.GetType())
6  let author : person = {name = "Jon"; age = 50; height = 1.75}
7  printfn "%A : %A" author (author.GetType())
8  let father = {person.name = "Jon"; age = 50; height = 1.75}
9  printfn "%A : %A" author (author.GetType())
```

```
1  $ fsharpc --nologo recordsDominance.fsx && mono
      recordsDominance.exe
2  {name = "Jon";
3   age = 50;
4   height = 1.75;} : RecordsDominance+teacher
5  {name = "Jon";
6   age = 50;
7   height = 1.75;} : RecordsDominance+person
8  {name = "Jon";
9   age = 50;
10  height = 1.75;} : RecordsDominance+person
```

In the example, two identical record types are defined, and we use the built-in `GetType()` method to inspect the type of bindings. We see that `lecturer` is of `RecordsDominance+teacher` type, since `teacher` dominates the identical `author` type definition. However, we may en-

force the `person` type by either specifying it for the name, as in `let author : person = ...`, or by fully or partially specifying it in the record expression following the "=" sign. In both cases, they are of `RecordsDominance+author` type. The built-in `GetType()` method is inherited from the base class for all types, see Chapter 20 for a discussion on classes and inheritance.

Note that when creating a record you must supply a value to all fields, and you cannot refer to other fields of the same record, i.e., `{name = "Jon"; age = height * 3; height = 1.75}` is illegal.

Since records are per default reference types, binding creates aliases, not copies. This matters for mutable members, in which case when copying, we must explicitly create a new record with the old data. Copying can be done either by using referencing to the individual members of the source or using the short-hand *with* notation. This is demonstrated in Listing 14.13.   · `with`

Listing 14.13 recordCopy.fsx:
Bindings are references. To copy and not make an alias, explicit copying must be performed.

```
1  type person = {
2    name : string;
3    mutable age : int;
4  }
5
6  let author = {name = "Jon"; age = 50}
7  let authorAlias = author
8  let authorCopy = {name = author.name; age = author.age}
9  let authorCopyAlt = {author with name = "Noj"}
10 author.age <- 51
11 printfn "author : %A" author
12 printfn "authorAlias : %A" authorAlias
13 printfn "authorCopy : %A" authorCopy
14 printfn "authorCopyAlt : %A" authorCopyAlt
```

```
1  $ fsharpc --nologo recordCopy.fsx && mono recordCopy.exe
2  author : {name = "Jon";
3   age = 51;}
4  authorAlias : {name = "Jon";
5   age = 51;}
6  authorCopy : {name = "Jon";
7   age = 50;}
8  authorCopyAlt : {name = "Noj";
9   age = 50;}
```

Here, `age` is defined as a mutable value and can be changed using the usual "`<-`" assignment operator. The example demonstrates two different ways to create records. Note that when the mutable value `author.age` is changed in line 10, then `authorAlias` also changes, since it is an alias of `author`, but neither `authorCopy` nor `authorCopyAlt` changes, since they are copies. As illustrated, copying using `with` allows for easy copying and partial updates of another record value.

14.5. Structures

Structures, or *structs* for short, have much in common with records. They specify a compound type with named fields, but they are value types, and they allow for some customization of what is to happen when a value of its type is created. Since they are value types, they are best used for small amounts of data. The syntax for defining struct types are:

· structures
· structs

```
Listing 14.14:  Syntax for type abbreviation.

[ <attributes> ]
[<Struct >]
type <ident> =
  val [ mutable ] <label1> : <type1>
  val [ mutable ] <label2> : <type2>
  ...
  [new (<arg1>, <arg2>, ...) = {<label1> = <arg1>; <label1> =
  <arg2>; ...}
  [new (<arg1>, <arg2>, ...) = {<label1> = <arg1>; <label1> =
  <arg2>; ...}
  ...
```

The syntax makes use of the `val` and `new` keywords. Like `let`, the keyword `val` binds a name to a value, but unlike `let`, the value is always the type's default value. The `new` keyword denotes the function used to fill values into the fields at time of creation. This function is called the *constructor*. No `let` or `do` bindings are allowed in structure definitions. Fields are accessed using the *"."* notation. An example is given in Listing 14.15.

· val
·
· new
· constructor
· ·

```
Listing 14.15 struct.fsx:
Defining a struct type and creating a value of it.

[<Struct >]
type position =
  val x : float
  val y : float
  new (a : float, b : float) = {x = a; y = b}

let p = position (3.0, 4.2)
printfn "%A: x = %A, y = %A" p p.x p.y
```
```
$ fsharpc --nologo struct.fsx && mono struct.exe
Struct+position: x = 3.0, y = 4.2
```

Structs are small versions of classes and allows, e.g., for overloading of the `new` constructor and for overriding of the inherited *ToString()* function. This is demonstrated in Listing 14.16.

· overload
· override
· ToString()

> **Listing 14.16 structOverloadNOverride.fsx:**
> **Overloading the `new` constructor and overriding the default `ToString()` function.**
>
> ```
> 1  [<Struct >]
> 2  type position =
> 3    val x : float
> 4    val y : float
> 5    new (a : float, b : float) = {x = a; y = b}
> 6    new (a : int, b : int) = {x = float a; y = float b}
> 7    override this.ToString () =
> 8      "(" + (string this.x) + ", " + (string this.y) + ")"
> 9
> 10  let pFloat = position (3.0, 4.2)
> 11  let pInt = position (3, 4)
> 12  printfn "%A and %A" pFloat pInt
> ```
> ```
> 1  $ fsharpc --nologo structOverloadNOverride.fsx
> 2  $ mono structOverloadNOverride.exe
> 3  (3, 4.2) and (3, 4)
> ```

We defer further discussion of these concepts to Chapter 20.

The use of structs are generally discouraged, and instead, it is recommended to use enums, records, and discriminated unions, possibly with the `[<Struct>]` attribute for the last two in order to make them value types.

14.6. Variable Types

An advanced topic in F# is *variable types*. There are three different versions of variable types in F#: *runtime resolved*, which have the syntax `'<ident>`, *anonymous*, which are written as "`_`", and *statically resolved*, which have the syntax `^<ident>`. Variable types are particularly useful for functions that work for many types. An example of a generic function and its use is given in Listing 14.17.

· variable types
· runtime resolved variable type
· anonymous variable type
· _
· statically resolved variable type

> **Listing 14.17 variableType.fsx:**
> **A function `apply` with runtime resolved types.**
>
> ```
> 1  let apply (f : 'a -> 'a -> 'a) (x : 'a) (y : 'a) : 'a = f x y
> 2  let intPlus (x : int) (y : int) : int = x + y
> 3  let floatPlus (x : float) (y : float) : float = x + y
> 4
> 5  printfn "%A %A" (apply intPlus 1 2) (apply floatPlus 1.0 2.0)
> ```
> ```
> 1  $ fsharpc --nologo variableType.fsx && mono variableType.exe
> 2  3 3.0
> ```

In this example, the function `apply` has runtime resolved variable type, and it accepts three parameters: `f`, `x`, and `y`. The function will work as long as the parameters for `f` is a function of two parameters of identical type, and `x` and `y` are values of the same type.

Thus, in the `printfn` statement we are able to use `apply` for both an integer and a float variant.

The example in Listing 14.17 illustrates a very complicated way to add two numbers. The "+" operator works for both types out of the box, so why not something simpler like relying on the F# type inference system by not explicitly specifying types, as attempted in Listing 14.18?

Listing 14.18 variableTypeError.fsx:
Even though the "+" operator is defined for both integers and floats, the type inference is static and infers `plus : int -> int`.

```
1  let plus x y = x + y
2
3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
```

```
1  $ fsharpc --nologo variableTypeError.fsx && mono
      variableTypeError.exe
2
3  variableTypeError.fsx(3,34): error FS0001: This expression was
      expected to have type
4      'int'
5  but here has type
6      'float'
7
8  variableTypeError.fsx(3,38): error FS0001: This expression was
      expected to have type
9      'int'
10 but here has type
11     'float'
```

Unfortunately, the example fails to compile, since the type inference is performed at compile time, and by `plus 1 2`, it is inferred that `plus : int -> int`. Hence, calling `plus 1.0 2.0` is a type error. Function bindings allow for the use of the *inline* keyword, and adding this successfully reuses the definition of `plus` for both types, as shown in Listing 14.19.

· `inline`

Listing 14.19 variableTypeInline.fsx:
The keyword `inline` forces static and independent inference each place the function is used. Compare to the error case in Listing 14.18.

```
1  let inline plus x y = x + y
2
3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
```

```
1  $ fsharpc --nologo variableTypeInline.fsx && mono
      variableTypeInline.exe
2  3 3.0
```

In the example, adding the `inline` does two things: Firstly, it copies the code to be performed to each place the function is used, and secondly, it forces statically resolved variable type checking independently in each place. The type annotations inferred as a result of the `inline`-keyword may be written explicitly, as shown in Listing 14.20.

> **Listing 14.20 compiletimeVariableType.fsx:**
> **Explicitly spelling out of the statically resolved type variables from Listing 14.18.**
>
> ```
> 1  let inline plus (x : ^a) (y : ^a) : ^a when ^a : (static
>       member ( + ) :  ^a * ^a -> ^a) = x + y
> 2
> 3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
> ```
> ```
> 1  $ fsharpc --nologo compiletimeVariableType.fsx
> 2  $ mono compiletimeVariableType.exe
> 3  3 3.0
> ```

The example in Listing 14.20 demonstrates the statically resolved variable type syntax, `^<ident>`, as well as the use of *type constraints*, using the keyword *when*. Type constraints   ·type constraints
have a rich syntax, but will not be discussed further in this book.[2] In the example, the   ·`when`
type constraint `when ^a : (static member ( + ) :  ^a * ^a -> ^a)` is given using the
object-oriented properties of the type variable `^a`, meaning that the only acceptable type
values are those which have a member function `(+)` taking a tuple and giving a value all
of identical type, and where the type can be inferred at compile time. See Chapter 20 for
details on member functions.

The `inline` construction is useful when generating generic functions and still profiting
from static type checking. However, explicit copying of functions is often something better
left to the compiler to optimize over. An alternative seems to be using runtime resolved
variable types with the `'<ident>` syntax. Unfortunately, this is not possible in case of most
operators, since they have been defined in the `FSharp.Core` namespace to be statically
resolved variable types. E.g., the "+" operator has type `( + ) : ^T1 -> ^T2 -> ^T3`
`(requires ^T1 with static member (+) and ^T2 with static member (+))`.

---

[2]Jon: **Should I extend on type constraints? Perhaps it is better left for a specialize chapter
on generic functions.**

15 | Pattern Matching

Pattern matching is used to transform values and variables into a syntactical structure. The simplest example is value-bindings. The *let*-keyword was introduced in Section 6.1, · `let` its extension with pattern matching is given as,

Listing 15.1: Syntax for `let`-expressions with pattern matching.

```
[[<Literal >]]
let [mutable] <pat> [: <returnType>] = <bodyExpr> [in <expr>]
```

A typical use of this is to extract elements of tuples, as demonstrated in Listing 15.2.

Listing 15.2 letPattern.fsx:
Patterns in `let` expressions may be used to extract elements of tuples.

```
let a = (3,4)
let (x,y) = a
let (alsoX,_) = a
printfn "%A: %d %d %d" a x y alsoX
```

```
$ fsharpc --nologo letPattern.fsx && mono letPattern.exe
(3, 4): 3 4 3
```

Here we extract the elements of a pair twice. First by binding to x and y, and second by binding to alsoX while using the wildcard pattern to ignore the second element. Thus, again the wildcard pattern in value-bindings is used to underline a disregarded value.

Another common use of patterns is as an alternative to `if` – `then` – `else` expressions, particularly when parsing input for a function. Consider the example in Listing 15.3.

Listing 15.3 switch.fsx:
Using *if − then − else* to print discriminated unions.

```
1   type Medal = Gold | Silver | Bronze
2   let statement (m : Medal) : string =
3     if m = Gold then "You won"
4     elif m = Silver then "You almost won"
5     else "Maybe you will win next time"
6
7   let m = Silver
8   printfn "%A : %s" m (statement m)
```

```
1   $ fsharpc --nologo switch.fsx && mono switch.exe
2   Silver : You almost won
```

In the example, a discriminated union and a function are defined. The function converts each case to a supporting statement, using an `if`-expression. The same can be done with the *match − with* expression and patterns, as demonstrated in Listing 15.4.

· match
· with

Listing 15.4 switchPattern.fsx:
Using *match − with* to print discriminated unions.

```
1   type Medal = Gold | Silver | Bronze
2   let statement (m : Medal) : string =
3     match m with
4       Gold -> "You won"
5       | Silver -> "You almost won"
6       | _ -> "Maybe you can win next time"
7
8   let m = Silver
9   printfn "%A : %s" m (statement m)
```

```
1   $ fsharpc --nologo switchPattern.fsx && mono switchPattern.exe
2   Silver : You almost won
```

Here we used a pattern for the discriminated union cases and a wildcard pattern as default. The lightweight syntax for `match`-expressions is,

Listing 15.5: Syntax for `match`-expressions.

```
1   match <inputExpr> with
2   [| ]<pat> [when <guardExpr>] -> <caseExpr>
3   | <pat> [when <guardExpr>] -> <caseExpr>
4   | <pat> [when <guardExpr>] -> <caseExpr>
5   ...
```

where `<inputExpr>` is the *input pattern* to find matches of, `<pat>` is a pattern to match with, `<guardExpr>` is an optional guard expression, and `<caseExpr>` is the resulting expression. Each set starting with `<pat>` is called a case. In lightweight syntax, the indentation must be equal to or higher than the indentation of `match`. All cases must return a value of the same type, and F# reports an error when the complete domain of the input pattern

· input pattern

is not covered by cases in `match`-expressions.

Patterns are also used in a version of *for*-loop expressions, and its lightweight syntax is given as,

> **Listing 15.6:** Syntax for `for`-expressions with pattern matching.
>
> ```
> for <pat> in <sourceExpr> do
>   <bodyExpr>
> ```

Typically, `<sourceExpr>` is a list or an array. An example is given in Listing 15.7.

> **Listing 15.7 forPattern.fsx:**
> **Patterns may be used in `for`-loops.**
>
> ```
> for (_,y) in [(1,3); (2,1)] do
>   printfn "%d" y
> ```
> ```
> $ fsharpc --nologo forPattern.fsx && mono forPattern.exe
> 3
> 1
> ```

The wildcard pattern is used to disregard the first element in a pair while iterating over the complete list. It is good practice to **use wildcard patterns to emphasize unused values.**

The final expression involving patterns to be discussed is the *anonymous functions*. Patterns for anonymous functions have the syntax,

> **Listing 15.8:** Syntax for anonymous functions with pattern matching.
>
> ```
> fun <pat> [<pat> ...] -> <bodyExpr>
> ```

This is an extension of the syntax discussed in Section 6.2. A typical use for patterns in *fun*-expressions is shown in Listing 15.9.

> **Listing 15.9 funPattern.fsx:**
> **Patterns may be used in `fun`-expressions.**
>
> ```
> let f = fun _ -> "hello"
> printfn "%s" (f 3)
> ```
> ```
> $ fsharpc --nologo funPattern.fsx && mono funPattern.exe
> hello
> ```

Here we use an anonymous function expression and bind it to `f`. The expression has one argument of any type, which it ignores through the wildcard pattern. Some limitations

apply to the patterns allowed in `fun`-expressions.[1] The wildcard pattern in `fun`-expressions are often used for *mockup functions*, where the code requires the said function, but its content has yet to be decided. Thus, mockup functions can be used as loose place-holders while experimenting with program design.

· mockup functions

Patterns are also used in exceptions to be discussed in Section 18.1, and in conjunction with the `function`-keyword, a keyword we discourage in this book. We will now demonstrate a list of important patterns in F#.

15.1. Wildcard Pattern

A *wildcard pattern* is denoted "`_`" and matches anything, see e.g., Listing 15.10.

· wildcard pattern
· `_`

> **Listing 15.10 wildcardPattern.fsx:**
> **Constant patterns match to constants.**
>
> ```
> 1  let whatEver (x : int) : string =
> 2    match x with
> 3      _ -> "If you say so"
> 4
> 5  printfn "%s" (whatEver 42)
> ```
> --
> ```
> 1  $ fsharpc --nologo wildcardPattern.fsx && mono
>      wildcardPattern.exe
> 2  If you say so
> ```

In this example, anything matches the wildcard pattern, so all cases are covered and the function always returns the same sentence. This is rarely a useful structure on its own, since this could be replaced by a value binding or by a function ignoring its input. However, wildcard patterns are extremely useful, since they act as the final `else` in `if`-expressions.

15.2. Constant and Literal Patterns

A *constant pattern* matches any input pattern with constants, see e.g., Listing 15.11.

· constant pattern

---

[1]Jon: **Remove or elaborate.**

Listing 15.11 constPattern.fsx:
Constant patterns match to constants.

```
1  type Medal = Gold | Silver | Bronze
2  let intToMedal (x : int) : Medal =
3    match x with
4      0 -> Gold
5      | 1 -> Silver
6      | _ -> Bronze
7
8  printfn "%A" (intToMedal 0)
```

```
1  $ fsharpc --nologo constPattern.fsx && mono constPattern.exe
2  Gold
```

In this example, the input pattern is queried for a match with 0, 1, or the wildcard pattern. Any simple literal type constants may be used in the constant pattern, such as 8, 23y, 1010u, 1.2, "hello world", 'c', and false. Here we also use the wildcard pattern. Note that matching is performed in a lazy manner and stops at the first matching case from the top. Thus, although the wildcard pattern matches everything, its case expression is only executed if none of the previous patterns match the input.

Constants can also be pre-bound by the [<Literal>] attribute for value-bindings. This is demonstrated in Listing 15.12.

Listing 15.12 literalPattern.fsx:
A variant of constant patterns is literal patterns.

```
1  [<Literal>]
2  let TheAnswer = 42
3  let whatIsTheQuestion (x : int) : string =
4    match x with
5      TheAnswer -> "We will need to build a bigger machine..."
6      | _ -> "Don't know that either"
7
8  printfn "%A" (whatIsTheQuestion 42)
```

```
1  $ fsharpc --nologo literalPattern.fsx && mono
     literalPattern.exe
2  "We will need to build a bigger machine..."
```

The attribute is used to identify the value-binding `TheAnswer` to be used, as if it were a simple literal type. Literal patterns must be either uppercase or module prefixed identifiers.

15.3. Variable Patterns

A *variable pattern* is a single lower-case letter identifier. Variable pattern identifiers are assigned the value and type of the input pattern. Combinations of constant and variable patterns are also allowed in conjunction with with records and arrays. This is demonstrated in Listing 15.13.

· variable pattern

Listing 15.13 variablePattern.fsx:
Variable patterns are useful for, e.g., extracting and naming fields

```
1  let (name, age) = ("Jon", 50)
2  let getAgeString (age : int) : string =
3    match age with
4      0 -> "a newborn"
5      | 1 -> "1 year old"
6      | n -> (string n) + " years old"
7
8  printfn "%s is %s" name (getAgeString age)
```

```
1  $ fsharpc --nologo variablePattern.fsx && mono
     variablePattern.exe
2  Jon is 50 years old
```

In this example, the value identifier **n** has the function of a named wildcard pattern. Hence, the case could as well have been | _ -> (string age) + "years old", since age is already defined in this scope. However, variable patterns syntactically act as an argument to an anonymous function and thus act to isolate the dependencies. They are also very useful together with guards, see Section 15.4.

15.4. Guards

A *guard* is a pattern used together with `match`-expressions including the *when*-keyword, as shown in Listing 15.5.

· guard
· when

Listing 15.14 guardPattern.fsx:
Guard expressions can be used with other patterns to restrict matches.

```
1  let getAgeString (age : int) : string =
2    match age with
3      n when n < 1 -> "infant"
4      | n when n < 13 -> "child"
5      | n when n < 20 -> "teen"
6      | _ -> "adult"
7
8  printfn "A person aged %d is a/an %s" 50 (getAgeString 50)
```

```
1  $ fsharpc --nologo guardPattern.fsx && mono guardPattern.exe
2  A person aged 50 is a/an adult
```

Here guards are used to iteratively carve out subset of integers to assign different strings to each set. The guard expression in `<pat> when <guardExpr> -> <caseExpr>` is any expression evaluating to a Boolean, and the case expression is only executed for the matching case.

15.5. List Patterns

Lists have a concatenation pattern associated with them. The *"::"* cons-operator is used to to match the head and the rest of a list, and *"[]"* is used to match an empty list, which is also sometimes called the nil-case. This is very useful when recursively processing lists, as shown in Listing 15.15

· list pattern
· ::
· []

Listing 15.15 listPattern.fsx:
Recursively parsing a list with list patterns.

```
1  let rec sumList (lst : int list) : int =
2    match lst with
3      n :: rest -> n + (sumList rest)
4      | [] -> 0
5
6  let rec sumThree (lst : int list) : int =
7    match lst with
8      [a; b; c] -> a + b + c
9      | _ -> sumList lst
10
11  let aList = [1; 2; 3]
12  printfn "The sum of %A is %d, %d" aList (sumList aList)
     (sumThree aList)
```

```
1  $ fsharpc --nologo listPattern.fsx && mono listPattern.exe
2  The sum of [1; 2; 3] is 6, 6
```

In the example, the function `sumList` uses the cons operator to match the head of the list with `n` and the tail with `rest`. The pattern `n :: tail` also matches `3 :: []`, and in that case `tail` would be assigned the value `[]`. When `lst` is empty, then it matches with "`[]`". List patterns can also be matched explicitly named elements, as demonstrated in the `sumThree` function. The elements to be matched can be any mix of constants and variables.

15.6. Array, Record, and Discriminated Union Patterns

Array, *record*, and *discriminated union patterns* are direct extensions on constant, variable, and wildcard patterns. Listing 15.16 gives examples of array patterns.

· array pattern
· record pattern
· discriminated union patterns

Listing 15.16 arrayPattern.fsx:
Using variable patterns to match on size and content of arrays.

```
1  let arrayToString (x : int []) : string =
2    match x with
3      [|1;_;_|] -> "3 elements, first of is 1"
4      | [|x;1;_|] -> "3 elements, first is " + (string x) + "
     Second 1"
5      | x -> "A general array"
6
7  printfn "%s" (arrayToString [|1; 1; 1|])
8  printfn "%s" (arrayToString [|3; 1; 1|])
9  printfn "%s" (arrayToString [|1|])
```

```
1  $ fsharpc --nologo arrayPattern.fsx && mono arrayPattern.exe
2  3 elements, first of is 1
3  3 elements, first is 3 Second 1
4  A general array
```

In the function `arrayToString`, the first case matches arrays of 3 elements where the first is the integer 1, the second case matches arrays of 3 elements where the second is a 1 and names the first `x`, and the final case matches all arrays and works as a default match case. As demonstrated, the cases are treated from first to last, and only the expression of the first case that matches is executed.

For record patterns, we use the field names to specify matching criteria. This is demonstrated in Listing 15.17.

Listing 15.17 recordPattern.fsx:
Variable patterns for records to match on field values.

```
1  type Address = {street : string; zip : int; country : string}
2  let contact : Address = {
3      street = "Universitetsparken 1";
4      zip = 2100;
5      country = "Denmark"}
6  let getZip (adr : Address) : int =
7    match adr with
8      {street = _; zip = z; country = _} -> z
9
10 printfn "The zip-code is: %d" (getZip contact)
```

```
1  $ fsharpc --nologo recordPattern.fsx && mono recordPattern.exe
2  The zip-code is: 2100
```

Here, the record type `Address` is created, and in the function `getZip`, a variable pattern `z` is created for naming zip values, and the remaining fields are ignored. Since the fields are named, the pattern match need not mention the ignored fields, and the example match is equivalent to `{zip = z} -> z`. The curly brackets are required for record patterns.

Discriminated union patterns are similar. For discriminated unions with arguments, the arguments can be matched as constants, variables, or wildcards. A demonstration is given

in Listing 15.18.

Listing 15.18 unionPattern.fsx:
Matching on discriminated union types.

```
1   type vector =
2     Vec2D of float * float
3     | Vec3D of float * float * float
4
5   let project (vec : vector) : vector =
6     match vec with
7       Vec3D (a, b, _) -> Vec2D (a, b)
8       | v -> v
9
10  let v = Vec3D (1.0, -1.2, 0.9)
11  printfn "%A -> %A" v (project v)
```

```
1   $ fsharpc --nologo unionPattern.fsx && mono unionPattern.exe
2   Vec3D (1.0,-1.2,0.9) -> Vec2D (1.0,-1.2)
```

In the `project`-function, three-dimensional vectors are projected to two dimensions by removing the third element. Two-dimensional vectors are unchanged. The example uses the wildcard pattern to emphasize that the third element of three-dimensional vectors is ignored. Named arguments can also be matched, in which case ";" is used instead of "," to delimit the fields in the match.

15.7. Disjunctive and Conjunctive Patterns

Patterns may be combined using the *"|"* and *"&"* lexemes. These patterns are called dis-   · |
junctive and conjunctive patterns, respectively, and work similarly to their logical operator   · &
counter parts, "`||`" and "`&&`".

Disjunctive patterns require at least one pattern to match, as illustrated in Listing 15.19.   · disjunctive pattern

---

Listing 15.19 disjunctivePattern.fsx:
Patterns can be combined logically as 'or' syntax structures.

```
1   let vowel (c : char) : bool =
2     match c with
3       'a' | 'e' | 'i' | 'o' | 'u' | 'y' -> true
4       | _ -> false
5
6   String.iter (fun c -> printf "%A " (vowel c)) "abcdefg"
```

```
1   $ fsharpc --nologo disjunctivePattern.fsx && mono
      disjunctivePattern.exe
2   true false false false true false false
```

Here one or more cases must match for the final case expression, and thus, any vowel results in the value `true`. Everything else is matched with the wildcard pattern.

For *conjunctive patterns*, all patterns must match, which is illustrated in Listing 15.20.   · conjunctive patterns

---

Listing 15.20 conjunctivePattern.fsx:
Patterns can be combined logically as 'and' syntax structures.

```
1  let is11 (v : int * int) : bool =
2    match v with
3      (1,_) & (_,1) -> true
4      | _ -> false
5
6  printfn "%A" (List.map is11 [(0,0); (0,1); (1,0); (1,1)])
```
- -
```
1  $ fsharpc --nologo conjunctivePattern.fsx && mono
     conjunctivePattern.exe
2  [false; false; false; true]
```

In this case, we separately check the elements of a pair for the constant value `1` and return true only when both elements are `1`. In many cases, conjunctive patters can be replaced by more elegant matches, e.g., using tuples, and in the above example a single case `(1,1)` `-> true` would have been simpler. Nevertheless, conjunctive patterns are used together with active patterns, to be discussed below.

15.8. Active Patterns

The concept of patterns is extendable to functions. Such functions are called *active patterns*, and active patterns come in two flavors: regular and option types. The active   · active patterns
pattern cases are constructed as function bindings, but using a special notation. They all take the pattern input as last argument, and may take further preceding arguments. The syntax for active patterns is one of,

---

Listing 15.21:   Syntax for binding active patterns to expressions.

```
1  let (|<caseName>|[_| ]) [ <arg> [<arg> ... ]] <inputArgument> =
     <expr>
2  let (|<caseName>|<caseName>|...|<caseName>|) <inputArgumet> =
     <expr>
```

When using the `(|<caseName>|_|])` variants, then the active pattern function must return an option type. The multi-case variant `(|<caseName>|<caseName>|...|<caseName>|)` must return a `Fsharp.Core.Choice` type. All other variants can return any type. There are no restrictions on arguments `<arg>`, and `<inputArgumetn>` is the input pattern to be matched. Notice in particular that the multi-case variant only takes one argument and cannot be combined with the option-type syntax. Below we will demonstrate by example how the various patterns are used.

The single case, `(|<caseName>|])`, matches all and is useful for extracting information from complex types, as demonstrated in Listing 15.22.

Listing 15.22 activePattern.fsx:
Single case active pattern for deconstructing complex types.

```
1  type vec = {x : float; y : float}
2  let (|Cartesian|) (v : vec) = (v.x, v.y)
3  let (|Polar|) (v : vec) = (sqrt(v.x*v.x + v.y * v.y), atan2
     v.y v.x)
4  let printCartesian (p : vec) : unit =
5      match p with
6         Cartesian (x, y) -> printfn "%A:\n Cartesian (%A, %A)" p
     x y
7  let printPolar (p : vec) : unit =
8      match p with
9         Polar (a, d) -> printfn "%A:\n Polar (%A, %A)" p a d
10
11  let v = {x = 2.0; y = 3.0}
12  printCartesian v
13  printPolar v
```

```
1  $ fsharpc --nologo activePattern.fsx && mono activePattern.exe
2  {x = 2.0;
3   y = 3.0;}:
4   Cartesian (2.0, 3.0)
5  {x = 2.0;
6   y = 3.0;}:
7   Polar (3.605551275, 0.9827937232)
```

Here we define a record to represent two-dimensional vectors and two different single case active patterns. Note that in the binding of the active pattern functions in line 2 and 3, the argument is the input expression `match <inputExpr> with` ..., see Listing 15.5. However, the argument for the cases in line 6 and 9 are names bound to the output of the active pattern function.

Both `Cartesian` and `Polar` match a vector record, but they dismantle the contents differently. For an alternative solution using Class types, see Section 20.1.

More complicated behavior is obtainable by supplying additional arguments to the single case. This is demonstrated in Listing 15.23.

> **Listing 15.23 activeArgumentsPattern.fsx:**
> **All but the multi-case active pattern may take additional arguments.**
>
> ```
> type vec = {x : float; y : float}
> let (|Polar|) (o : vec) (v : vec) =
>   let x = v.x - o.x
>   let y = v.y - o.y
>   (sqrt(x*x + y * y), atan2 y x)
> let printPolar (o : vec) (p : vec) : unit =
>     match p with
>       Polar o (a, d) -> printfn "%A:\n Cartesian (%A, %A)" p a
>   d
>
> let v = {x = 2.0; y = 3.0}
> let offset = {x = 1.0; y = 1.0}
> printPolar offset v
> ```
>
> ```
> $ fsharpc --nologo activeArgumentsPattern.fsx
> $ mono activeArgumentsPattern.exe
> {x = 2.0;
>  y = 3.0;}:
>  Cartesian (2.236067977, 1.107148718)
> ```

Here we supply an offset, which should be subtracted prior to calculating lengths and angles. Notice in line 8 that the argument is given prior to the result binding.

Active pattern functions return option types are called *partial pattern functions*. The option type allows for specifying mismatches, as illustrated in Listing 15.24.

· partial pattern functions

> **Listing 15.24 activeOptionPattern.fsx:**
> **Option type active patterns mismatch on `None` results.**
>
> ```
> let (|Div|_|) (e,d) = if d <> 0.0 then Some (e/d) else None
>
> let safeDiv (p : float * float) =
>     match p with
>       | (0.0, 0.0) -> printfn "Div %A = undefined" p
>       | Div res -> printfn "Div %A = %A" p res
>       | _ -> printfn "Div %A = infinity" p
>
> List.iter safeDiv [(1.0,1.0); (0.0,1.0); (1.0,0.0); (0.0,0.0)]
> ```
>
> ```
> $ fsharpc --nologo activeOptionPattern.fsx
> $ mono activeOptionPattern.exe
> Div (1.0, 1.0) = 1.0
> Div (0.0, 1.0) = 0.0
> Div (1.0, 0.0) = infinity
> Div (0.0, 0.0) = undefined
> ```

In the example, we use the (`|<caseName>|_|`) variant to indicate that the active pattern returns an option type. Nevertheless, the result binding `res` in line 6 uses the underlying value of `Some`. And in contrast to the two previous examples of single case patterns, the value `None` results in a mismatch. Thus in this case, if the denominator is `0.0`, then

`Div` `res` does not match but the wildcard pattern does.

Multicase active patterns work similarly to discriminated unions without arguments.[2] An   · multicase active
example is given in Listing 15.25.   patterns

Listing 15.25 activeMultiCasePattern.fsx:
Multi-case active patterns have a syntactical structure similar to discriminated unions.

```
1  let (|Gold|Silver|Bronze|) inp =
2     if inp = 0 then Gold
3     elif inp = 1 then Silver
4     else Bronze
5
6  let intToMedal (i : int) =
7      match i with
8        Gold -> printfn "%d: It's gold!" i
9        | Silver -> printfn "%d: It's silver." i
10       | Bronze -> printfn "%d: It's no more than bronze." i
11
12 List.iter intToMedal [0..3]
```

```
1  $ fsharpc --nologo activeMultiCasePattern.fsx
2  $ mono activeMultiCasePattern.exe
3  0: It's gold!
4  1: It's silver.
5  2: It's no more than bronze.
6  3: It's no more than bronze.
```

In this example, we define three cases in line 1. The result of the active pattern function
must be one of these cases. For the `match`-expression, the match is based on the output of
the active pattern function, hence in line 8, the case expression is executed when the result
of applying the active pattern function to the input expression `i` is `Gold`. In this case, a
solution based on discriminated unions would probably be clearer.

15.9. Static and Dynamic Type Pattern

Input patterns can also be matched on type. For *static type matching*, the matching is   · static type pattern
performed at compile time and indicated using the *":"* lexeme followed by the type name to   · :
be matched. Static type matching is further used as input to the type inference performed
at compile time to infer non-specified types, as illustrated in Listing 15.26.

---

[2]Jon: **This maybe too advanced for this book.**

Listing 15.26 staticTypePattern.fsx:
Static matching on type binds the type of other values by type inference.

```
1  let rec sum lst =
2      match lst with
3        (n : int) :: rest -> n + (sum rest)
4        | [] -> 0
5
6  printfn "The sum is %d" (sum [0..3])
```
```
1  $ fsharpc --nologo staticTypePattern.fsx && mono
     staticTypePattern.exe
2  The sum is 6
```

Here the head of the list `n` in the list pattern is explicitly matched as an integer, and the type inference system thus concludes that `lst` must be a list of integers.

In contrast to static type matching, *dynamic type matching* is performed at runtimes and indicated using the "`:?`" lexeme followed by a type name. Dynamic type patterns allow for matching generic values at runtime. This is an advanced topic, which is included here for completeness. An example is given in Listing 15.27.

· dynamic type
  pattern
· :?

Listing 15.27 dynamicTypePattern.fsx:
Dynamic matching on type binds the type of other values by type inference.

```
1  let isString (x : obj) : bool =
2      match x with
3        :? string -> true
4        | _ -> false
5
6  let a = "hej"
7  printfn "Is %A a string? %b" a (isString a)
8  let b = 3
9  printfn "Is %A a string? %b" b (isString b)
```
```
1  $ fsharpc --nologo dynamicTypePattern.fsx && mono
     dynamicTypePattern.exe
2  Is "hej" a string? true
3  Is 3 a string? false
```

In F#, all types are also objects whose type is denoted `obj`. Thus, the example uses the generic type when defining the argument to `isString`, and then dynamic type pattern matching for further processing. See Chapter 20 for more on objects. Dynamic type patterns are often used for analyzing exceptions, which is discussed in Section 18.1. While dynamic type patterns are useful, they imply runtime checking, and **it is almost always better to prefer compile time over runtime type checking.**

Advice

A | The Console in Windows, MacOS X, and Linux

Almost all popular operating systems are accessed through a user-friendly *graphical user interface (GUI)* that is designed to make typical tasks easy to learn to solve. As a computer programmer, you often need to access some of the functionalities of the computer, which, unfortunately, are sometimes complicated by this particular graphical user interface. The *console*, also called the *terminal* and the *Windows command line*, is the right hand of a programmer. The console is a simple program that allows you to complete text commands. Almost all the tasks that can be done with the graphical user interface can be done in the console and vice versa. Using the console, you will benefit from its direct control of the programs we write, and in your education, you will benefit from the fast and raw information you get through the console.

· graphical user interface

· GUI

· console

· terminal

· Windows command line

A.1. The Basics

When you open a *directory* or *folder* in your preferred operating system, the directory will have a location in the file system, whether from the console or through the operating system's graphical user interface. The console will almost always be associated with a particular directory or folder in the file system, and it is said that it is the directory that the console is in. The exact structure of file systems varies between Linux, MacOS X, and Windows, but common is that it is a hierarchical structure. This is illustrated in Figure A.1.

· directory

· folder

There are many predefined console commands, available in the console, and you can also make your own. In the following sections, we will review the most important commands in the three different operating systems. These are summarized in Table A.1.

A.2. Windows

In this section we will discuss the commands summarized in Table A.1. Windows 7 and earlier versions: To open the console, press `Start->Run` in the lower left corner, and then type `cmd` in the box. In Windows 8 and 10, you right-click on the windows icon, choose `Run` or equivalent in your local language, and type `cmd`. Alternatively, you can type `Windows-key + R`. Now you should open a console window with a prompt showing something like Listing A.1.

(a) Linux
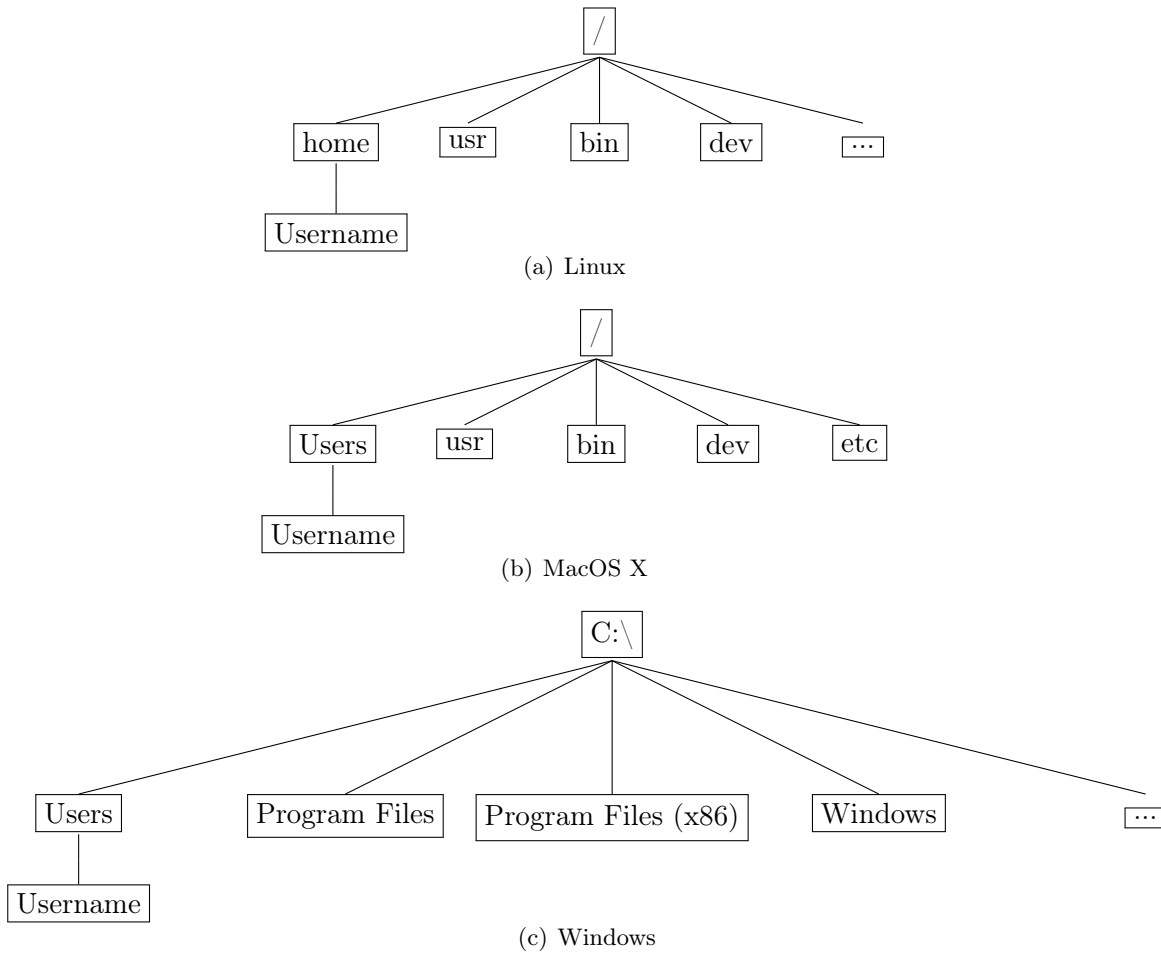
(b) MacOS X

(c) Windows

Figure A.1.: The top file hierarchy levels of common operating systems.

Windows	MacOS X/Linux	Description
`dir`	`ls`	Show content of present directory.
`cd <dir>`	`cd <dir>`	Change present directory to `<dir>`.
`mkdir <dir>`	`mkdir <dir>`	Create directory `<dir>`.
`rmdir <dir>`	`rmdir <dir>`	Delete `<dir>` (Warning: cannot be reverted).
`move <file> <file or dir>`	`mv <file> <file or dir>`	Move `<fil>` to `<file or dir>`.
`copy <file1> <file2>`	`cp <file1> <file2>`	Create a new file called `<file2>` as a copy of `<file1>`.
`del <file>`	`rm <file>`	delete `<file>` (Warning: cannot be reverted).
`echo <string or variable>`	`echo <string or variable>`	Write a string or content of a variable to screen.

Table A.1.: The most important console commands for Windows, MacOS X, and Linux.

Listing A.1: The Windows console.

```
1  Microsoft Windows [Version 6.1.7601]
2  Copyright (c) 2009 Microsoft Corporation.  All rights reserved.
3
4  C:\Users\sporring>
```

To see which files are in the directory, use *dir*, as shown in Listing A.2.   · dir

Listing A.2: Directory listing with dir.

```
1  C:\Users\sporring>dir
2   Volume in drive C has no label.
3   Volume Serial Number is 94F0-31BD
4
5   Directory of C:\Users\sporring
6
7  30-07-2015  15:23    <DIR>          .
8  30-07-2015  15:23    <DIR>          ..
9  30-07-2015  14:27    <DIR>          Contacts
10 30-07-2015  14:27    <DIR>          Desktop
11 30-07-2015  17:40    <DIR>          Documents
12 30-07-2015  15:11    <DIR>          Downloads
13 30-07-2015  14:28    <DIR>          Favorites
14 30-07-2015  14:27    <DIR>          Links
15 30-07-2015  14:27    <DIR>          Music
16 30-07-2015  14:27    <DIR>          Pictures
17 30-07-2015  14:27    <DIR>          Saved Games
18 30-07-2015  17:27    <DIR>          Searches
19 30-07-2015  14:27    <DIR>          Videos
20                0 File(s)              0 bytes
21               13 Dir(s)  95.004.622.848 bytes free
22
23 C:\Users\sporring>
```

We see that there are no files and thirteen directories (DIR). The columns tell from left to right: the date and time of their creation, the file size or if it is a folder, and the name file or directory name. The first two folders "." and ".." are found in each folder and refer to this folder as well as the one above in the hierarchy. In this case, the folder "." is an alias for `C:\Users\sporring` and ".." for `C:\Users`.

Use *cd* to change directory, e.g., to `Documents`, as in Listing A.3.   · cd

Listing A.3: Change directory with cd.

```
1  C:\Users\sporring>cd Documents
2
3  C:\Users\sporring\Documents>
```

Note that some systems translate default filenames, so their names may be given different names in different languages in the graphical user interface as compared to the console.

You can use *mkdir* to create a new directory called, e.g., `myFolder`, as illustrated in List-   · mkdir

ing A.4.

Listing A.4: Creating a directory with `mkdir`.

```
1   C:\Users\sporring\Documents>mkdir myFolder
2
3   C:\Users\sporring\Documents>dir
4    Volume in drive C has no label.
5    Volume Serial Number is 94F0-31BD
6
7    Directory of C:\Users\sporring\Documents
8
9   30-07-2015  19:17    <DIR>          .
10  30-07-2015  19:17    <DIR>          ..
11  30-07-2015  19:17    <DIR>          myFolder
12             0 File(s)              0 bytes
13             3 Dir(s)  94.656.638.976 bytes free
14
15  C:\Users\sporring\Documents>
```

By using `dir` we inspect the result.

Files can be created by, e.g., ***echo*** and *redirection*, as demonstrated in Listing A.5.   · `echo`
· redirection

Listing A.5: Creating a file with `echo` and redirection.

```
1   C:\Users\sporring\Documents>echo "Hi" > hi.txt
2
3   C:\Users\sporring\Documents>dir
4    Volume in drive C has no label.
5    Volume Serial Number is 94F0-31BD
6
7    Directory of C:\Users\sporring\Documents
8
9   30-07-2015  19:18    <DIR>          .
10  30-07-2015  19:18    <DIR>          ..
11  30-07-2015  19:17    <DIR>          myFolder
12  30-07-2015  19:18                 8 hi.txt
13             1 File(s)              8 bytes
14             3 Dir(s)  94.656.634.880 bytes free
15
16  C:\Users\sporring\Documents>
```

To move the file `hi.txt` to the directory `myFolder`, use ***move***, as shown in Listing A.6.   · `move`

Listing A.6: Move a file with `move`.

```
1   C:\Users\sporring\Documents>move hi.txt myFolder
2          1 file(s) moved.
3
4   C:\Users\sporring\Documents>
```

Finally, use ***del*** to delete a file and ***rmdir*** to delete a directory, as shown in Listing A.7.   · `del`
· `rmdir`

Listing A.7: Delete files and directories with `del` and `rmdir`.

```
1   C:\Users\sporring\Documents>cd myFolder
2
3   C:\Users\sporring\Documents\myFolder>del hi.txt
4
5   C:\Users\sporring\Documents\myFolder>cd ..
6
7   C:\Users\sporring\Documents>rmdir myFolder
8
9   C:\Users\sporring\Documents>dir
10   Volume in drive C has no label.
11   Volume Serial Number is 94F0-31BD
12
13   Directory of C:\Users\sporring\Documents
14
15  30-07-2015  19:20    <DIR>           .
16  30-07-2015  19:20    <DIR>           ..
17               0 File(s)           0 bytes
18               2 Dir(s)  94.651.142.144 bytes free
19
20  C:\Users\sporring\Documents>
```

The commands available from the console must be in its *search path*. The search path can   · search path
be seen using `echo`, as shown in Listing A.8.

Listing A.8: Displaying the search path.

```
1   C:\Users\sporring\Documents>echo %Path%
2   C:\Windows\system32;C:\Windows;C:\Windows\System32\Wbem;
      C:\Windows\System32\WindowsPowerShell\v1.0\;"\Program
      Files\emacs-24.5\bin\"
3
4   C:\Users\sporring\Documents>
```

The path can be changed using the Control panel in the graphical user interface. In Windows 7, choose the Control panel, choose `System and Security` → `System` → `Advanced system settings` → `Environment Variables`. In Windows 10, you can find this window by searching for "Environment" in the Control panel. In the window's `System variables` box, double-click on `Path` and add or remove a path from the list. The search path is a list of paths separated by ";". Beware, Windows uses the search path for many different tasks, so remove only paths that you are certain are not used for anything.

A useful feature of the console is that you can use the `tab`-key to cycle through filenames. E.g., if you write `cd` followed by a space and `tab` a couple of times, then the console will suggest to you the available directories.

A.3. MacOS X and Linux

MacOS X (OSX) and Linux are very similar, and both have the option of using *bash* as   · bash
console. It is in the standard console on MacOS X and on many Linux distributions. A summary of the most important `bash` commands is shown in Table A.1. In MacOS X,

you find the console by opening `Finder` and navigating to `Applications` → `Utilities` -> `Terminal`. In Linux, the console can be started by typing `Ctrl + Alt + T`. Some Linux distributions have other key-combinations such as `Super + T`.

Once opened, the console is shown in a window with content, as shown in Listing A.9.

Listing A.9: The MacOS console.

```
1  Last login: Thu Jul 30 11:52:07 on ttys000
2  FN11194:~ sporring$
```

"FN11194" is the name of the computer, the character ∼ is used as an alias for the user's home directory, and "sporring" is the username for the user presently logged onto the system. Use *ls* to see which files are present, as shown in Listing A.10.   · ls

---

Listing A.10: Display a directory content with `ls`.

```
1  FN11194:~ sporring$ ls
2  Applications    Documents    Library        Music         Public
3  Desktop         Downloads    Movies         Pictures
4  FN11194:~ sporring$
```

More details about the files are available by using flags to `ls` as demonstrated in Listing A.11.

Listing A.11: Display extra information about files using flags to `ls`.

```
1   FN11194:~ sporring$ ls -l
2   drwx------   6 sporring  staff   204 Jul 30 14:07 Applications
3   drwx------+ 32 sporring  staff  1088 Jul 30 14:34 Desktop
4   drwx------+ 76 sporring  staff  2584 Jul  2 15:53 Documents
5   drwx------+  4 sporring  staff   136 Jul 30 14:35 Downloads
6   drwx------@ 63 sporring  staff  2142 Jul 30 14:07 Library
7   drwx------+  3 sporring  staff   102 Jun 29 21:48 Movies
8   drwx------+  4 sporring  staff   136 Jul  4 17:40 Music
9   drwx------+  3 sporring  staff   102 Jun 29 21:48 Pictures
10  drwxr-xr-x+  5 sporring  staff   170 Jun 29 21:48 Public
11  FN11194:~ sporring$
```

The flag `-l` means long, and many other flags can be found by querying the built-in manual with `man ls`. The output is divided into columns, where the left column shows a number of codes: "d" stands for directory, and the set of three of optional "rwx" denote whether respectively the owner, the associated group of users, and anyone can respectively "r" - read, "w" - write, and "x" - execute the file. In all directories but the `Public` directory, only the owner can do any of the three. For directories, "x" means permission to enter. The second column can often be ignored, but shows how many links there are to the file or directory. Then follows the username of the owner, which in this case is `sporring`. The files are also associated with a group of users, and in this case, they all are associated with the group called `staff`. Then follows the file or directory size, the date of last change, and the file or directory name. There are always two hidden directories: "." and "..", where "." is an alias for the present directory, and ".." for the directory above. Hidden files will be shown with the `-a` flag.

Use `cd` to change to the directory, for example to `Documents` as shown in Listing A.12.   · `cd`

> **Listing A.12: Change directory with `cd`.**
>
> ```
> 1  FN11194:~ sporring$ cd Documents/
> 2  FN11194:Documents sporring$
> ```

Note that some graphical user interfaces translate standard filenames and directories to the local language, such that navigating using the graphical user interface will reveal other files and directories, which, however, are aliases.

You can create a new directory using *mkdir*, as demonstrated in Listing A.13.   · `mkdir`

> **Listing A.13: Creating a directory using `mkdir`.**
>
> ```
> 1  FN11194:Documents sporring$  mkdir myFolder
> 2  FN11194:Documents sporring$ ls
> 3  myFolder
> 4  FN11194:tmp sporring$
> ```

A file can be created using *echo* and with *redirection*, as shown in Listing A.14.   · `echo`
  · redirection

> **Listing A.14: Creating a file with `echo` and redirection.**
>
> ```
> 1  FN11194:Documents sporring$ echo "hi" > hi.txt
> 2  FN11194:Documents sporring$ ls
> 3  hi.txt        myFolder
> ```

To move the file `hi.txt` into `myFolder`, use *mv*. This is demonstrated in Listing A.15.   · `mv`

> **Listing A.15: Moving files with `mv`.**
>
> ```
> 1  FN11194:Documents sporring$ echo mv hi.txt myFolder/
> 2  FN11194:Documents sporring$
> ```

To delete the file and the directory, use *rm* and *rmdir*, as shown in Listing A.16.   · `rm`
  · `rmdir`

> **Listing A.16: Deleting files and directories.**
>
> ```
> 1  FN11194:Documents sporring$ cd myFolder/
> 2  FN11194:myFolder sporring$ rm hi.txt
> 3  FN11194:myFolder sporring$ cd ..
> 4  FN11194:Documents sporring$ rmdir myFolder/
> 5  FN11194:Documents sporring$ ls
> 6  FN11194:Documents sporring$
> ```

Only commands found on the *search path* are available in the console. The content of the   · search path
search path is seen using the `echo` command, as demonstrated in Listing A.17.

Listing A.17: The content of the search path.

```
1  FN11194:Documents sporring$ echo $PATH
2  /Applications/Maple
     17/:/Applications/PackageMaker.app/Contents/MacOS/:
     /Applications/MATLAB_R2014b.app/bin/:/opt/local/bin:
     /opt/local/sbin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:
     /sbin:/opt/X11/bin:/Library/TeX/texbin
3  FN11194:Documents sporring$
```

The search path can be changed by editing the setup file for Bash. On MacOS X it is called `~/.profile`, and on Linux it is either `~/.bash_profile` or `~/.bashrc`. Here new paths can be added by adding the following line: `export PATH="<new path>:<another new path>:$PATH"`.

A useful feature of Bash is that the console can help you write commands. E.g., if you write `fs` followed by pressing the `tab`-key, and if Mono is in the search path, then Bash will typically respond by completing the line as `fsharp`, and by further pressing the `tab`-key some times, Bash will show the list of options, typically `fshpari` and `fsharpc`. Also, most commands have an extensive manual which can be accessed using the `man` command. E.g., the manual for `rm` is retrieved by `man rm`.

B | Number Systems on the Computer

B.1. Binary Numbers

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base* 10 meaning that a decimal number consists of a sequence of digits separated by a *decimal point*, where each *digit* can have values $d \in \{0, 1, 2, \ldots, 9\}$ and the weight of each digit is proportional to its place in the sequence of digits with respect to the decimal point, i.e., the number $357.6 = 3 \cdot 10^2 + 5 \cdot 10^1 + 7 \cdot 10^0 + 6 \cdot 10^{-1}$, or in general, for a number consisting of digits d_i with $n + 1$ and m digits to the left and right of the decimal point, the value v is calculated as:

$$v = \sum_{i=-m}^{n} d_i 10^i. \tag{B.1}$$

· decimal number
· base
· decimal point
· digit

The basic unit of information in almost all computers is the binary digit, or *bit* for short. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. The general equation is,

$$v = \sum_{i=-m}^{n} b_i 2^i, \tag{B.2}$$

· bit
· binary number

and examples are $1011.1_2 = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 + 1 \cdot 2^{-1} = 11.5$. Notice that we use subscript 2 to denote a binary number, while no subscript is used for decimal numbers. The left-most bit is called the *most significant bit*, and the right-most bit is called the *least significant bit*. Due to typical organisation of computer memory, 8 binary digits is called a *byte*, and the term *word* is not universally defined but typically related to the computer architecture, a program is running on, such as 32 or 64 bits.

· most significant bit
· least significant bit
· byte
· word

Other number systems are often used, e.g., *octal numbers*, which are base 8 numbers and have digits $o \in \{0, 1, \ldots, 7\}$. Octals are useful short-hand for binary, since 3 binary digits map to the set of octal digits. Likewise, *hexadecimal numbers* are base 16 with digits $h \in \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, b, c, d, e, f\}$, such that $a_{16} = 10$, $b_{16} = 11$ and so on. Hexadecimals are convenient, since 4 binary digits map directly to the set of hexadecimal digits. Thus $367 = 101101111_2 = 557_8 = 16f_{16}$. A list of the integers 0–63 in various bases is given in Table B.1.

· octal number
· hexadecimal number

B.2. IEEE 754 Floating Point Standard

The set of real numbers, also called *reals*, includes all fractions and irrational numbers. It

Dec	Bin	Oct	Hex	Dec	Bin	Oct	Hex
0	0	0	0	32	100000	40	20
1	1	1	1	33	100001	41	21
2	10	2	2	34	100010	42	22
3	11	3	3	35	100011	43	23
4	100	4	4	36	100100	44	24
5	101	5	5	37	100101	45	25
6	110	6	6	38	100110	46	26
7	111	7	7	39	100111	47	27
8	1000	10	8	40	101000	50	28
9	1001	11	9	41	101001	51	29
10	1010	12	a	42	101010	52	2a
11	1011	13	b	43	101011	53	2b
12	1100	14	c	44	101100	54	2c
13	1101	15	d	45	101101	55	2d
14	1110	16	e	46	101110	56	2e
15	1111	17	f	47	101111	57	2f
16	10000	20	10	48	110000	60	30
17	10001	21	11	49	110001	61	31
18	10010	22	12	50	110010	62	32
19	10011	23	13	51	110011	63	33
20	10100	24	14	52	110100	64	34
21	10101	25	15	53	110101	65	35
22	10110	26	16	54	110110	66	36
23	10111	27	17	55	110111	67	37
24	11000	30	18	56	111000	70	38
25	11001	31	19	57	111001	71	39
26	11010	32	1a	58	111010	72	3a
27	11011	33	1b	59	111011	73	3b
28	11100	34	1c	60	111100	74	3c
29	11101	35	1d	61	111101	75	3d
30	11110	36	1e	62	111110	76	3e
31	11111	37	1f	63	111111	77	3f

Table B.1.: A list of the integers 0–63 in decimal, binary, octal, and hexadecimal.

is infinite in size both in the sense that there is no largest nor smallest number, and that between any 2 given numbers there are infinitely many numbers. Reals are widely used for calculation, but since any computer only has finite memory, there are infinitely many numbers which cannot be represent on a computer. Hence, any computation performed on a computer with reals must rely on approximations. *IEEE 754 double precision floating-point format* (*binary64*), known as a *double*, is a standard for representing an approximation of reals using 64 bits. These bits are divided into 3 parts: sign, exponent and fraction,

$$s\, e_1 e_2 \ldots e_{11}\, m_1 m_2 \ldots m_{52},$$

· IEEE 754 double precision floating-point format
· binary64
· double

where s, e_i, and m_j are binary digits. The bits are converted to a number using the equation by first calculating the exponent e and the mantissa m,

$$e = \sum_{i=1}^{11} e_i 2^{11-i}, \tag{B.3}$$

$$m = \sum_{j=1}^{52} m_j 2^{-j}. \tag{B.4}$$

I.e., the exponent is an integer, where $0 \le e < 2^{11}$, and the mantissa is a rational, where $0 \le m < 1$. For most combinations of e and m, the real number v is calculated as,

$$v = (-1)^s \, (1 + m) \, 2^{e-1023} \tag{B.5}$$

with the exceptions that

	$m = 0$	$m \ne 0$
$e = 0$	$v = (-1)^s \, 0$ (signed zero)	$v = (-1)^s \, m 2^{1-1023}$ (subnormals)
$e = 2^{11} - 1$	$v = (-1)^s \, \infty$	$v = (-1)^s \, \mathrm{NaN}$ (not-a-number)

where $e = 2^{11} - 1 = 11111111111_2 = 2047$. The largest and smallest number that is not infinity is thus

· subnormals
· NaN
· not-a-number

$$e = 2^{11} - 2 = 2046, \tag{B.6}$$

$$m = \sum_{j=1}^{52} 2^{-j} = 1 - 2^{-52} \simeq 1, \tag{B.7}$$

$$v_{\max} = \pm \left(2 - 2^{-52}\right) 2^{1023} \simeq \pm 2^{1024} \simeq \pm 10^{308}. \tag{B.8}$$

The density of numbers varies in such a way that when $e - 1023 = 52$, then

$$v = (-1)^s \left(1 + \sum_{j=1}^{52} m_j 2^{-j}\right) 2^{52} \tag{B.9}$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{-j} 2^{52}\right) \tag{B.10}$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{52-j}\right) \tag{B.11}$$

$$\overset{k=52-j}{=} \pm \left(2^{52} + \sum_{k=51}^{0} m_{52-k} 2^k\right), \tag{B.12}$$

which are all integers in the range $2^{52} \le |v| < 2^{53}$. When $e - 1023 = 53$, then the same calculation gives

$$v \overset{k=53-j}{=} \pm \left(2^{53} + \sum_{k=52}^{1} m_{53-k} 2^k\right), \tag{B.13}$$

which are every second integer in the range $2^{53} \le |v| < 2^{54}$, and so on for larger values of e. When $e - 1023 = 51$, the same calculation gives,

$$v \overset{k=51-j}{=} \pm \left(2^{51} + \sum_{k=50}^{-1} m_{51-k} 2^k\right), \tag{B.14}$$

which is a distance between numbers of $1/2$ in the range $2^{51} \leq |v| < 2^{52}$, and so on for smaller values of e. Thus we may conclude that the distance between numbers in the interval $2^n \leq |v| < 2^{n+1}$ is 2^{n-52}, for $-1022 = 1 - 1023 \leq n < 2046 - 1023 = 1023$. For subnormals, the distance between numbers is

$$v = (-1)^s \left(\sum_{j=1}^{52} m_j 2^{-j} \right) 2^{-1022} \tag{B.15}$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j} 2^{-1022} \right) \tag{B.16}$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j-1022} \right) \tag{B.17}$$

$$\overset{k=-j-1022}{=} \pm \left(\sum_{j=-1023}^{-1074} m_{-k-1022} 2^k \right), \tag{B.18}$$

which gives a distance between numbers of $2^{-1074} \simeq 10^{-323}$ in the range $0 < |v| < 2^{-1022} \simeq 10^{-308}$.

C | Commonly Used Character Sets

Letters, digits, symbols, and space are the core of how we store data, write programs, and communicate with computers and each other. These symbols are in short called characters and represent a mapping between numbers, also known as codes, and a pictorial representation of the character. E.g., the ASCII code for the letter 'A' is 65. These mappings are for short called character sets, and due to differences in natural languages and symbols used across the globe, many different character sets are in use. E.g., the English alphabet contains the letters 'a' to 'z'. These letters are common to many other European languages which in addition use even more symbols and accents. For example, Danish has further the letters 'æ', 'ø', and 'å'. Many non-European languages have completely different symbols, where the Chinese character set is probably the most extreme, and some definitions contain 106,230 different characters, albeit only 2,600 are included in the official Chinese language test at the highest level.

Presently, the most common character set used is Unicode Transformation Format (UTF), whose most popular encoding schemes are 8-bit (UTF-8) and 16-bit (UTF-16). Many other character sets exist, and many of the later build on the American Standard Code for Information Interchange (ASCII). The ISO-8859 codes were an intermediate set of character sets that are still in use, but which is greatly inferior to UTF. Here we will briefly give an overview of ASCII, ISO-8859-1 (Latin1), and UTF.

C.1. ASCII

The *American Standard Code for Information Interchange* (*ASCII*) [8], is a 7 bit code tuned for the letters of the English language, numbers, punctuation symbols, control codes and space, see Tables C.1 and C.2. The first 32 codes are reserved for non-printable control characters to control printers and similar devices or to provide meta-information. The meaning of each control character is not universally agreed upon.

· American Standard Code for Information Interchange
· ASCII

The code order is known as *ASCIIbetical order*, and it is sometimes used to perform arithmetic on codes, e.g., an uppercase letter with code c may be converted to lower case by adding 32 to its code. The ASCIIbetical order also has a consequence for sorting, i.e., when sorting characters according to their ASCII code, 'A' comes before 'a', which comes before the symbol '{'.

· ASCIIbetical order

x0+0x	00	10	20	30	40	50	60	70
00	NUL	DLE	SP	0	@	P	'	p
01	SOH	DC1	!	1	A	Q	a	q
02	STX	DC2	"	2	B	R	b	r
03	ETX	DC3	#	3	C	S	c	s
04	EOT	DC4	$	4	D	T	d	t
05	ENQ	NAK	%	5	E	U	e	u
06	ACK	SYN	&	6	F	V	f	v
07	BEL	ETB	'	7	G	W	g	w
08	BS	CAN	(	8	H	X	h	x
09	HT	EM	)	9	I	Y	i	y
0A	LF	SUB	*	:	J	Z	j	z
0B	VT	ESC	+	;	K	[	k	{
0C	FF	FS	,	<	L	\	l		
0D	CR	GS	−	=	M	]	m	}
0E	SO	RS	.	>	N	^	n	~
0F	SI	US	/	?	O	_	o	DEL

Table C.1.: ASCII

C.2. ISO/IEC 8859

The ISO/IEC 8859 report `http://www.iso.org/iso/catalogue_detail?csnumber=28245`
defines 10 sets of codes specifying up to 191 codes and graphics characters using 8 bits. Set
1, also known as ISO/IEC 8859-1, Latin alphabet No. 1, or *Latin1*, covers many European   · Latin1
languages and is designed to be compatible with ASCII, such that code for the printable
characters in ASCII is the same in ISO 8859-1. Table C.3 shows the characters above 7e.
Codes 00-1f and 7f-9f are undefined in ISO 8859-1.

C.3. Unicode

Unicode is a character standard defined by the Unicode Consortium, `http://unicode.org`,
as the *Unicode Standard*. Unicode allows for 1,114,112 different codes. Each code is called   · Unicode Standard
a *code point* which represents an abstract character. However, not all abstract characters   · code point
require a unit of several code points to be specified. Code points are divided into 17
planes, each with $2^{16} = 65,536$ code points. Planes are further subdivided into named
blocks. The first plane is called the *Basic Multilingual plane* and its block of the first 128   · unicode block
code points is called the *Basic Latin block* and is identical to ASCII, see Table C.1, and   · Basic Multilingual
code points 128-255 are called the *Latin-1 Supplement block*, and are identical to the upper   plane
range of ISO 8859-1, see Table C.3. Each code-point has a number of attributes such as the   · Basic Latin block
Unicode general category. Presently more than 128,000 code points are defined as covering   · Latin-1 Supplement
135 modern and historical writing systems, and obtained at `http://www.unicode.org/`   block
`Public/UNIDATA/UnicodeData.txt`, which includes the code point, name, and general   · Unicode general
category.   category

A Unicode code point is an abstraction from the encoding and the graphical representation
of a character. A code point is written as "U+" followed by its hexadecimal number, and
for the Basic Multilingual plane, 4 digits are used, e.g., the code point with the unique
name LATIN CAPITAL LETTER A has the Unicode code point "U+0041", and is in this

Code	Description
NUL	Null
SOH	Start of heading
STX	Start of text
ETX	End of text
EOT	End of transmission
ENQ	Enquiry
ACK	Acknowledge
BEL	Bell
BS	Backspace
HT	Horizontal tabulation
LF	Line feed
VT	Vertical tabulation
FF	Form feed
CR	Carriage return
SO	Shift out
SI	Shift in
DLE	Data link escape
DC1	Device control one
DC2	Device control two
DC3	Device control three
DC4	Device control four
NAK	Negative acknowledge
SYN	Synchronous idle
ETB	End of transmission block
CAN	Cancel
EM	End of medium
SUB	Substitute
ESC	Escape
FS	File separator
GS	Group separator
RS	Record separator
US	Unit separator
SP	Space
DEL	Delete

Table C.2.: ASCII symbols.

text visualized as 'A'. More digits are used for code points of the remaining planes.

The general category is used to specify valid characters that do not necessarily have a visual representation but possibly transform text. Some categories and their letters in the first 256 code points are shown in Table C.5.

To store and retrieve code points, they must be encoded and decoded. A common encoding is *UTF-8*, which encodes code points as 1 to 4 bytes, and which is backward-compatible · UTF-8 with ASCII and ISO 8859-1. Hence, in all 3 coding systems, the character with code 65 represents the character 'A'. Another popular encoding scheme is *UTF-16*, which encodes · UTF-16 characters as 2 or 4 bytes, but which is not backward-compatible with ASCII or ISO 8859-1. UTF-16 is used internally in many compilers, interpreters, and operating systems.

x0+0x	80	90	A0	B0	C0	D0	E0	F0
00			NBSP	°	À	Ð	à	ð
01			¡	±	Á	Ñ	á	ñ
02			¢	²	Â	Ò	â	ò
03			£	³	Ã	Ó	ã	ó
04			¤	´	Ä	Ô	ä	ô
05			¥	µ	Å	Õ	å	õ
06			¦	¶	Æ	Ö	æ	ö
07			§	·	Ç	×	ç	÷
08			¨	¸	È	Ø	è	ø
09			©	¹	É	Ù	é	ù
0a			ª	º	Ê	Ú	ê	ú
0b			«	»	Ë	Û	ë	û
0c			¬	¼	Ì	Ü	ì	ü
0d			SHY	½	Í	Ý	í	ý
0e			®	¾	Î	Þ	î	þ
0f			¯	¿	Ï	ß	ï	ÿ

Table C.3.: ISO-8859-1 (latin1) non-ASCII part. Note that the codes 7f – 9f are undefined.

Code	Description
NBSP	Non-breakable space
SHY	Soft hypen

Table C.4.: ISO-8859-1 special symbols.

General category	Code points	Name
Lu	U+0041–U+005A, U+00C0–U+00D6, U+00D8–U+00DE	Upper case letter
Ll	U+0061–U+007A, U+00B5, U+00DF–U+00F6, U+00F8–U+00FF	Lower case letter
Lt	None	Digraphic letter, with first part uppercase
Lm	None	Modifier letter
Lo	U+00AA, U+00BA	Gender ordinal indicator
Nl	None	Letterlike numeric character
Pc	U+005F	Low line
Mn	None	Nonspacing combining mark
Mc	None	Spacing combining mark
Cf	U+00AD	Soft Hyphen

Table C.5.: Some general categories for the first 256 code points.

Bibliography

[1] Alonzo Church. An unsolvable problem of elementary number theory. *American Journal of Mathematics*, 58:345—-363, 1936.

[2] Ole-Johan Dahl and Kristen Nygaard. SIMULA a language for programming and description of discrete event systems. introduction and user's manual. Technical report, Norwegian Computing Center, 1967.

[3] European Computer Manufacturers Association (ECMA). Standard ecma-335, common language infrastructure (cli). `http://www.ecma-international.org/publications/standards/Ecma-335.htm`.

[4] International Organization for Standardization. Iso/iec 23271:2012, common language infrastructure (cli). `https://www.iso.org/standard/58046.html`.

[5] Object Management Group. Uml version 2.0. `http://www.omg.org/spec/UML/2.0/`.

[6] Programming Research Group. Specifications for the ibm mathematical formula translating system, fortran. Technical report, Applied Science Division, International Business Machines Corporation, 1954.

[7] John McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, 1960.

[8] X3: ASA Sectional Committee on Computers and Information Processing. American standard code for information interchange. Technical Report ASA X3.4-1963, American Standards Association (ASA), 1963. `http://worldpowersystems.com/projects/codes/X3.4-1963/`.

[9] George Pólya. *How to solve it*. Princeton University Press, 1945.

[10] Alan M. Turing. On computable numbers, with an application to the entscheidungsproblem. *Proceedings of the London Mathematical Society*, s2-42(1):230–265, 1936.

Index