

Learning to program with F#

Jon Sparring

Department of Computer Science,
University of Copenhagen

June 2, 2018

Contents

1	Preface	7
2	Introduction	8
2.1	Learning how to solve problems by programming	8
2.2	How to solve problems	9
2.3	Approaches to programming	9
2.4	Why use F#	10
2.5	How to read this book	11
3	Executing F# code	12
3.1	Source code	12
3.2	Executing programs	13
4	Quick-start guide	15
5	Using F# as a calculator	20
5.1	Literals and basic types	20
5.2	Operators on basic types	25
5.3	Boolean arithmetic	28
5.4	Integer arithmetic	29
5.5	Floating point arithmetic	31
5.6	Char and string arithmetic	32
5.7	Programming intermezzo: Hand conversion between decimal and binary numbers	34
6	Values and functions	36
6.1	Value bindings	39

6.2	Function bindings	44
6.3	Operators	50
6.4	Do bindings	51
6.5	The Printf function	52
6.6	Reading from the console	55
6.7	Variables	55
6.8	Reference cells	58
6.9	Tuples	61
7	In-code documentation	65
8	Controlling program flow	71
8.1	While and for loops	71
8.2	Conditional expressions	76
8.3	Programming intermezzo: Automatic conversion of decimal to binary numbers	78
9	Organising code in libraries and application programs	81
9.1	Modules	81
9.2	Namespaces	84
9.3	Compiled libraries	86
10	Testing programs	90
10.1	White-box testing	92
10.2	Black-box testing	95
10.3	Debugging by tracing	98
11	Collections of data	107
11.1	Strings	107
11.1.1	String properties	107
11.1.2	String module	108
11.2	Lists	110
11.2.1	List properties	113
11.2.2	List module	113

11.3 Arrays	117
11.3.1 Array properties and methods	120
11.3.2 Array module	120
11.4 Multidimensional arrays	125
11.4.1 Array2D module	128
12 The imperative programming paradigm	130
12.1 Imperative design	131
13 Recursion	132
13.1 Recursive functions	132
13.2 The call stack and tail recursion	133
13.3 Mutual recursive functions	138
14 Programming with types	140
14.1 Type abbreviations	140
14.2 Enumerations	141
14.3 Discriminated Unions	142
14.4 Records	144
14.5 Structures	146
14.6 Variable types	148
15 Pattern matching	151
15.1 Wildcard pattern	154
15.2 Constant and literal patterns	154
15.3 Variable patterns	155
15.4 Guards	156
15.5 List patterns	156
15.6 Array, record, and discriminated union patterns	157
15.7 Disjunctive and conjunctive patterns	159
15.8 Active Pattern	159
15.9 Static and dynamic type pattern	162

16 Higher order functions	164
16.1 Function composition	166
16.2 Currying	166
17 The functional programming paradigm	169
17.1 Functional design	170
18 Handling Errors and Exceptions	172
18.1 Exceptions	172
18.2 Option types	181
18.3 Programming intermezzo: Sequential division of floats	182
19 Working with files	185
19.1 Command line arguments	186
19.2 Interacting with the console	187
19.3 Storing and retrieving data from a file	189
19.4 Working with files and directories.	192
19.5 Reading from the internet	193
19.6 Resource Management	194
19.7 Programming intermezzo: Ask user for existing file	195
20 Classes and objects	197
20.1 Constructors and members	197
20.2 Accessors	199
20.3 Objects are reference types	202
20.4 Static classes	203
20.5 Recursive members and classes	204
20.6 Function and operator overloading	205
20.7 Additional constructors	208
20.8 Interfacing with <code>printf</code> family	209
20.9 Programming intermezzo	211
21 Derived classes	215

21.1 Inheritance	215
21.2 Abstract class	218
21.3 Interfaces	220
21.4 Programming intermezzo: Chess	222
22 The object-oriented programming paradigm	232
22.1 Identification of objects, behaviors, and interactions by nouns-and-verbs	233
22.2 Class diagrams in the Unified Modelling Language	233
22.3 Programming intermezzo: designing a racing game	236
23 Graphical User Interfaces	241
23.1 Opening a window	241
23.2 Drawing geometric primitives	243
23.3 Programming intermezzo: Hilbert Curve	252
23.4 Handling events	255
23.5 Labels, buttons, and pop-up windows	257
23.6 Organising controls	263
24 The Event-driven programming paradigm	269
25 Where to go from here	270
A The console in Windows, MacOS X, and Linux	272
A.1 The basics	272
A.2 Windows	272
A.3 MacOS X and Linux	276
B Number systems on the computer	280
B.1 Binary numbers	280
B.2 IEEE 754 floating point standard	280
C Commonly used character sets	283
C.1 ASCII	283
C.2 ISO/IEC 8859	283

<i>CONTENTS</i>	6
C.3 Unicode	284
D Common Language Infrastructure	287
E Language Details	289
E.1 Arithmetic operators on basic types	289
E.2 Basic arithmetic functions	292
E.3 Precedence and associativity	293
F To Dos	301
Bibliography	303
Index	304

1 | Preface

This book has been written as an introduction to programming for novice programmers. It is used on the first programming course at the University of Copenhagen's bachelor in computer science program. It has been typeset in L^AT_EX, and all programs have been developed and tested in Mono version 5.2.0.

Jon Sparring
Associate Professor, Ph.d.
Department of Computer Science,
University of Copenhagen
June 2, 2018

2 | Introduction

Programming is a creative process in which exciting problems may be solved and new tools and applications may be created. With programming skills, you can create high-level applications to run on a mobile device that interacts with other users, databases, and artificial intelligence; you may create programs that run on supercomputers for simulating weather systems on alien planets or social phenomena in the internet economy; and you may create programs that run on small custom-made hardware for controlling your home appliances.

2.1 Learning how to solve problems by programming

In order to learn how to program, there are a couple of steps that are useful:

1. Choose a programming language: A programming language such as F# is a vocabulary and a set of grammatical rules for instructing a computer to perform a certain task. It is possible to program without a concrete language but your ideas and thoughts must be expressed in some fairly rigorous way. Actually, theoretical computer science typically does not rely on computers nor programming languages but uses mathematics to prove properties of algorithms. However, most computer scientists program, and with a real language, you have the added benefit of checking your algorithm, and hence your thoughts, rigorously on a real computer. This book teaches a subset of F#. The purpose is not to be a reference guide to this language but to use it as a vessel to teach you, the reader, how to convert your ideas into programs.
2. Learn the language: A computer language is a structure for thought, and it influences which thoughts you choose to express as a program, and how you choose to do it. Any conversion requires you to acquire a sufficient level of fluency, for you to be able to make programs. You do not need to be a master in F# nor to know every corner of the language, and you will expand your knowledge as you expose yourself to solving problems in the language but you must invest an initial amount of time and energy in order to learn the basics of the language. This book aims at getting you started quickly, which is why we intentionally teach just a small subset of F#. On the net and through other works, you will be able to learn much more.
3. Practice: If you want to be a good programmer, then there is only one way: practice, practice, practice! It has been estimated that to master anything, then you have to have spent at least 10000 hours of practice, so get started logging hours! It of course matters, what you practice. This book teaches 3 different programming themes. The point is that programming is thinking, and the scaffold, you use, shapes your thoughts. It is therefore important to recognize this scaffold and to have the ability to choose that which suits your ideas and your goals best. And the best way to expand your abilities is to sharpen your present abilities, push yourself into new territory, and trying something new. Do not be afraid to make errors or be frustrated at first. These are the experiences that make you grow.

4. Solve real problems: I have found that using my programming skills in real situations with customers demanding solutions, that work for them, has allowed me to put into perspective the programming tools and techniques that I use. Often customers want solutions that work, are secure, cheap, and delivered fast, which has pulled me as a programmer in the direction of “if it works, then sell it”. On the other hand, in the longer perspective customers also want bug fixes, upgrades, and new features, which require carefully designed code, well-written test-suites, and good documentation. And as always, the right solution is somewhere in between. Regardless, real problems create real programmers.

2.2 How to solve problems

Programming is the act of solving a problem by writing a program to be executed on a computer. A general method for solving problems was given by George Pólya [9] and adapted to programming is:

Understand the problem: To solve any problem it is crucial that the problem formulation is understood: What is to be solved? Do you understand everything in the description of the problem? Is all information for finding the solution available or is something missing?

Design a plan: Good designs mean that programs are faster to program, easier to find errors in and to update in the future. So, before you start typing a program consider things like: What are the requirements and constraints for the program? Which components should the program have? How are these components supposed to work together? Designing often involves drawing a diagram of the program and writing program sketches on paper.

Implement the plan: Implementation is the act of transforming a program design into code. A crucial part of any implementation is choosing which programming language to use. Also, the solution to many problems will have a number of implementations which vary in how much code they require, to which degree they rely on external libraries, which programming style they are best suited for, what machine resources they require, and what their running times are. With a good design, the coding is usually easy, since the design will have uncovered the major issues and found solutions for these but sometimes implementation reveals new problems, which requires rethinking the design. Most implementations also include writing documentation of the code.

Reflect on the result: A crucial part of any programming task is ensuring that the program solves the problem sufficiently. E.g., what are the program’s errors, is the documentation of the code sufficient and relevant for its intended use? Is the code easily maintainable and extendable by other programmers? Are there any general lessons to be learned from or general code developed by the programming experience, which may be used for future programming sessions?

Programming is a very complicated process, and Pólya’s list is a useful guide but not a fail-safe approach. Always approach problem-solving with an open mind.

2.3 Approaches to programming

This book focuses on 3 fundamentally different approaches to programming:

Imperative programming, emphasizes *how a program shall accomplish a solution* and less on *what the solution is*. A cooking recipe is an example of the spirit of imperative programming, where the recipe emphasizes what should be done in each step rather than describing the result. E.g., for making bread, you first mix yeast and water, then add flour, etc. In imperative programming

- imperative programming
- statement
- state

what should be done are called *statements* and they influence the computer's *states*, like adding flour changes the state of our dough. Almost all computer hardware is designed to execute low-level programs written in imperative style. Imperative programming builds on the Turing machine [10]. The first major language was FORTRAN [6] which emphasized an imperative style of programming.

Declarative programming, which emphasizes *what a program shall accomplish* but not *how*. We will consider Functional programming as an example of declarative programming. A *functional programming* language evaluates *functions* and avoids state changes. The program consists of *expressions* instead of statements. As an example the function $f(x) = x^2$ takes a number x and evaluates the expression x^2 , and returns the result. Functional programming has its roots in lambda calculus [1], and the first language emphasizing functional programming was Lisp [7].

- declarative programming
- functional programming
- function expression
- structured programming

Structured programming, which emphasizes organization of code in units with well-defined interfaces and isolation of internal states and code from other parts of the program. We will focus on Object-oriented programming as the example of structured programming. *Object-oriented programming* is a type of programming, where the states and programs are structured into *objects*. A typical object-oriented design takes a problem formulation and identifies key nouns as potential objects and verbs as potential actions to be taken on objects. The first object-oriented programming language was Simula 67 developed by Dahl and Nygaard at the Norwegian Computing Center in Oslo [2].

- Object-oriented programming
- object

Event-driven programming, is often used when dynamically interacting with the real world. E.g., when programming graphical user interfaces, programs will often need to react to a user clicking on the mouse or when text arrives from a web-server to be displayed on the screen. Event-driven programs are often programmed using *call-back functions*, which are small programs that are ready to run when events occur.

- event-driven programming
- call-back functions

Most programs do not follow a single programming paradigm as, e.g., one of the above but are a mix. Nevertheless, this book will treat each paradigm separately to emphasize their advantages and disadvantages.

2.4 Why use F#

This book uses F# also known as Fsharp, which is a functional first programming language that also supports imperative and object-oriented programming. It was originally developed for Microsoft's .Net platform but is available as open source for many operating systems through Mono. As an introduction to programming, F# is a young programming language still under development, with syntax that at times is a bit complex but it offers a number of advantages:

Interactive and compile mode: F# has an interactive and a compile mode of operation: In interactive mode you can write code that is executed immediately in a manner similar to working with a calculator, while in compile mode, you combine many lines of code possibly in many files into a single application, which is easier to distribute to non F# experts and is faster to execute.

Indentation for scope: F# uses indentation to indicate scope: Some lines of code belong together, e.g., should be executed in a certain order and may share data, and indentation helps in specifying this relationship.

Strongly typed: F# is strongly typed, reducing the number of runtime errors. That is, F# is picky, and will not allow the programmer to mix up types such as numbers and text. This is a great advantage for large programs.

Multi-platform: F# is available on Linux, Mac OS X, Android, iOS, Windows, GPUs, and browsers via the Mono platform.

Free to use and open source: F# is supported by the Fsharp foundation (<http://fsharp.org>) and sponsored by Microsoft.

Assemblies: F# is designed to be able to easily communicate with other .Net and Mono programs through the language-independent, platform-independent bytecode called Common Intermediate Language (CIL) organized as assemblies. Thus, if you find that certain parts of a program are easy to express in F# and others in C++, then you will be able to combine these parts later into a single program.

Modern computing: F# supports all aspects of modern computing including Graphical User Interfaces, Web programming, Information rich programming, Parallel algorithms, ...

Integrated development environments (IDE): F# is supported by major IDEs such as Visual Studio (<https://www.visualstudio.com>) and Xamarin Studio (<https://www.xamarin.com>).

2.5 How to read this book

Learning to program requires mastering a programming language, however, most programming languages contain details that are rarely used or used in contexts far from a specific programming topic. Hence, this book only includes a subset of F# but focuses on language structures necessary to understand 4 common programming paradigms: Imperative programming mainly covered in Chapters 6 to 11, functional programming mainly covered in Chapters 13 to 16, object-oriented programming in Chapters 20 and 22, and event-driven programming in Chapter 23. A number of general topics are given in the appendix for reference. The disadvantage of this approach is that no single part contains a reference guide to F#, and F# topics are revisited and expanded across the book. For further reading please consult <http://fsharp.org>.

3 | Executing F# code

3.1 Source code

F# is a functional first programming language, meaning that it has strong support for functional programming, but F# also supports imperative and object-oriented programming. It also has strong support for parallel programming and information-rich programs. It was originally developed for Microsoft's .Net platform, but is available as open source for many operating systems through Mono. In this text, we consider F# 4.1 and its Mono implementation, which is different from .Net mainly in terms of the number of libraries accessible. The complete language specification is described in <http://fsharp.org/specs/language-spec/>.

F# has 2 modes of execution, *interactive* and *compiled*. Interactive mode is well suited for small experiments or back-of-an-envelope calculations, but not for programming in general. Both modes can be accessed via the *console*, see Appendix A for more information on the console. The interactive system is started by calling `fsharpi` at the command prompt in the console, while compilation is performed with `fsharpc`, and execution of the compiled code is performed using the `mono` command.

F# programs come in many forms, which are identified by suffixes. The *source code* is an F# program written in human-readable form using an editor. F# recognizes the following types of source code files:

<code>.fs</code> An <i>implementation file</i> , e.g., <code>myModule.fs</code>	· implementation file
<code>.fsi</code> A <i>signature file</i> , e.g., <code>myModule.fsi</code>	· signature file
<code>.fsx</code> A <i>script file</i> , e.g., <code>gettingStartedStump.fsx</code>	· script file
<code>.fsscript</code> Same as <code>.fsx</code> , e.g., <code>gettingStartedStump.fsscript</code>	

Compiled code is source code translated into a machine-readable language, which can be executed by a machine. Compiled F# code is either:

<code>.dll</code> A <i>library file</i> , e.g., <code>myModule.dll</code>	· library file
<code>.exe</code> A stand-alone <i>executable file</i> , e.g., <code>gettingStartedStump.exe</code>	· executable file

The implementation, signature, and script files are all typically compiled to produce an executable file, in which case they are called *scripts*, but can also be entered into the interactive system, in which case these are called *script-fragments*. The implementation and signature files are special kinds of script files used for building libraries. Libraries in F# are called modules, and they are collections of smaller programs used by other programs, which will be discussed in detail in Chapter 9.

3.2 Executing programs

Programs may either be executed by the interpreter or by compiling and executing the compiled code. In Mono the interpreter is called `fsharpi` and can be used in two ways: interactively, where a user enters one or more script-fragments separated by the “`;;`” characters, or to execute a script file treated as a single script-fragment.¹

To illustrate the difference between interactive and compile mode, consider the program in Listing 3.1.

Listing 3.1 `getStartedStump.fsx`:
A simple demonstration script.

```
1 let a = 3.0
2 do printfn "%g" a
```

The code declares a value `a` to be the decimal value 3.0 and finally prints it to the console. The `do printfn` is a statement for displaying the content of a value to the screen, and `%g` is a special notation to control how the value is printed. In this case, it is printed as a decimal number. This and more will be discussed at length in the following chapters. For now, we will concentrate on how to interact with the F# interpreter and compiler.

An interactive session is obtained by starting the console, typing the `fsharpi` command, typing the lines of the program, and ending the script-fragment with “`;;`”. The dialogue in Listing 3.2 demonstrates the workflow. What the user types has been highlighted by a box.

Listing 3.2: An interactive session.

```
1 $ fsharpi
2
3 F# Interactive for F# 4.1 (Open Source Edition)
4 Freely distributed under the Apache 2.0 Open Source License
5
6 For help type #help;;
7
8 > let a = 3.0
9 - do printfn "%g" a;;
10 3
11
12 val a : float = 3.0
13 val it : unit = ()
14
15 > #quit;;
```

We see that after typing `fsharpi`, then the program starts by stating details about itself followed by `>` indicating that it is ready to receive commands. The user then types `let a = 3.0` and presses `enter`, to which the interpreter responds with `-`. This indicates that the line has been received, that the script-fragment is not yet completed, and that it is ready to receive more input. When the user types `do printfn "%g" a;;` followed by `enter`, then by “`;;`” the interpreter knows that the script-fragment is completed, it interprets the script-fragment, responds with `3` and extra type information about the entered code, and with `>` to indicate, that it is ready for more script-fragments. The interpreter is

¹Jon: Too early to introduce lexeme: “F# uses many characters which at times are given special meanings, e.g., the characters “`;;`” is compound character denoting the end of a script-fragment. Such possibly compound characters are called lexemes.”

stopped, when the user types `#quit;;`. It is also possible to stop the interpreter by typing `ctrl-d`.

Instead of running `fsharpi` interactively, we can write the script-fragment from Listing 3.1 into a file, here called `gettingStartedStump.fsx`. This file can be interpreted directly by `fsharpi` as shown in Listing 3.3.

Listing 3.3: Using the interpreter to execute a script.

```
1 $ fsharpi gettingStartedStump.fsx
2 3
```

Notice that in the file, “`;;`” is optional. We see that the interpreter executes the code and prints the result on screen without the extra type information.

Finally, the file containing Listing 3.1 may be compiled into an executable file with the program `fsharpc`, and run using the program `mono` from the console. This is demonstrated in Listing 3.4.

Listing 3.4: Compiling and executing a script.

```
1 $ fsharpc gettingStartedStump.fsx
2 F# Compiler for F# 4.1 (Open Source Edition)
3 Freely distributed under the Apache 2.0 Open Source License
4 $ mono gettingStartedStump.exe
5 3
```

The compiler takes `gettingStartedStump.fsx` and produces `gettingStarted.exe`, which can be run using `mono`.

Both the interpreter and the compiler translates the source code into a format, which can be executed by the computer. While the compiler performs this translation once and stores the result in the executable file, the interpreter translates the code every time the code is executed. Thus, to run the program again with the interpreter, it must be retranslated as “`$fsharpi gettingStartedStump.fsx`”. In contrast, compiled code does not need to be recompiled to be run again, only re-executed using “`$mono gettingStartedStump.exe`”. On a MacBook Pro, with a 2.9 GHz Intel Core i5, the time the various stages take for this script are:

Command	Time
<code>fsharpi gettingStartedStump.fsx</code>	1.88s
<code>fsharpc gettingStartedStump.fsx</code>	1.90s
<code>mono gettingStartedStump.exe</code>	0.05s

I.e., executing the script with `fsharpi` is slightly faster than by first compiling it with `fsharpc` and then executing the result with `mono`, $1.88s < 0.05s + 1.90s$, if the script were to be executed only once, but every future execution of the script using the compiled version requires only the use of `mono`, which is much faster than `fsharpi`, $1.88s \gg 0.05s$.

The interactive session results in extra output on the *type inference* performed, which is very useful for *debugging* and development of code-fragments, but both executing programs with the interpreted directly from a file and compiling and executing the program is much preferred for programming complete programs, since the starting state is well defined, and since this better supports *unit-testing*, which is a method for debugging programs. Thus, **prefer compiling over interpretation**.

- type inference
- debugging
- unit-testing
- Advice

4 | Quick-start guide

Programming is the art of solving problems by writing a program to be executed by a computer. For example, to solve the following problem,

Problem 4.1

What is the sum of 357 and 864?

we have written the program in F# shown in Listing 4.1.

Listing 4.1 quickStartSum.fsx:

A script to add 2 numbers and print the result to the console.

```
1 let a = 357
2 let b = 864
3 let c = a + b
4 do printfn "%A" c

1 $ fsharp --nologo quickStartSum.fsx && mono quickStartSum.exe
2 1221
```

In the box the above, we see our program was saved as a script in a file called `quickStartSum.fsx`, and in the console we executed the program by typing the command `fsharp --nologo quickStartSum.fsx && mono quickStartSum.exe`. The result is then printed in the console to be 1221. Here, as in the rest of this book, we have used the optional flag `--nologo`, which informs `fsharp` not to print information about its version etc., thus making the output shorter. The `&&` notation tells the console to first run the command on the left, and if that did not report any errors, then run that on the right. This could as well have been performed as two separate commands to the console, and throughout this book, we will use the above shorthand, when convenient.

To solve the problem, we made program consisting of several lines, where each line was a *expressions*. The first expression `let a = 357` in line 1 used the `let` keyword to *bind* the value 357 to the name `a`. This is called a *let-binding*, and a let-binding makes the name synonymous with the value. Another point to be noted is that F# identifies 357 as an *integer number*, which is F#'s preferred number type, since computations on integers are very efficient, and since integers are very easy to communicate to other programs. In line 2 we bound the value 864 to the name `b`, and to the name `c`, we bound the result of evaluating the sum `a + b` in line 3. Line 4 is a *do-binding*, as noted by the keyword `do`. Do-bindings are also sometimes called *statements*, and the `do` keyword is optional in F#. Here the value of `c` was printed to the console followed by a newline (LF possibly preceded by CR, see Appendix C.1) with the *printfn function*. A function in F# is an entity that takes zero or more arguments and returns a value. The function `printfn` is very special, since it can take any number of arguments. It need not return any value, but F# insists that every function must return a value, wherefore `printfn`

- expression
- `let`
- keyword
- binding
- let-binding
- integer number
- do-binding
- `do`
- statements
- `printfn`
- function

returns a special type of value called *unit* and written as “()”. The `do` tells F# to ignore this value. Here `printfn` has been used with 2 arguments: “%A” and `c`. Notice that in contrast to many other languages, F# does not use parentheses to frame the list of arguments, nor does it use commas to separate them. In general, the `printfn` function always has 1 or more arguments, and the first is a *format string*. A *string* is a sequence of characters starting and ending with double quotation marks. E.g., `let s = "this is a string of characters"` binds the string “this is...” to the name `s`. For the `printfn` function, the format string may be any string, but if it contains format character sequences, such as %A, then the values following the format string are substituted. The format string must match the value *type*, that is, here `c` is of type integer, whereas the format string %A matches many types.

· unit
· ()
· format string
· string
· type

Types are a central concept in F#. In the script 4.1 we bound values of integer type to names. There are several different integer types in F#, here we used the one called `int`. The values were not *declared* to have these types, instead the types were *inferred* by F#. Typing these bindings line by line in an interactive session, then we see the inferred types as shown in Listing 4.2.

· type declaration
· type inference

Listing 4.2: Inferred types are given as part of the response from the interpreter.

```

1  > let a = 357;;
2  val a : int = 357
3
4  > let b = 864;;
5  val b : int = 864
6
7  > let c = a + b;;
8  val c : int = 1221
9
10 > do printfn "%A" c;;
11 1221
12 val it : unit = ()
```

The interactive session displays the type using the `val` keyword followed by the name used in the binding, its type, and its value. Since the value is also responded, then the last `printfn` statement is superfluous. However, **it is ill-advised to design programs to be run in an interactive session, since the scripts need to be manually copied every time it is to be run, and since the starting state may be unclear.** Notice that `printfn` is automatically bound to the name `it` of type `unit` and value “()”. F# insists on binding all statements to values, and in lack of an explicit name, then it will use `it`. Rumor has it that `it` is an abbreviation for “irrelevant”.

· val
· Advice
· it
· ()

Were we to solve a slightly different problem,

Problem 4.2

What is the sum of 357.6 and 863.4?

where the only difference is that the numbers now use a *decimal point*. These are called *floating point numbers*, and the internal representation is quite different to integer numbers used previously, and the algorithms used to perform arithmetic are also quite different from integers. Now the program would look like Listing 4.3.

· decimal point
· floating point numbers

Listing 4.3 quickStartSumFloat.fsx:
Floating point types and arithmetic.

```

1  let a = 357.6
2  let b = 863.4
3  let c = a + b
4  do printfn "%A" c

1  $ fsharpc --nologo quickStartSumFloat.fsx && mono quickStartSumFloat.exe
2  1221.0

```

On the surface, this could appear as an almost negligible change, but the set of integers and the set of real numbers (floats) require quite different representations in order to be effective on a computer, and as a consequence, the implementation of their operations such as addition are very different. Thus, although the response is an integer, it has type `float`, which is indicated by `1221.0`, and which is not the same as `1221`. F# is very picky about types, and generally does not allow types to be mixed, as demonstrated in the interactive session in Listing 4.4.

Listing 4.4: Mixing types is often not allowed.

```

1  > let a = 357;;
2  val a : int = 357
3
4  > let b = 863.4;;
5  val b : float = 863.4
6
7  > let c = a + b;;
8     let c = a + b;;
9     -----^
10
11 stdin(4,13): error FS0001: The type 'float' does not match the type 'int'

```

We see that binding a name to a number without a decimal point is inferred to be an integer, while when binding to a number with a decimal point, then the type is inferred to be a float, and when trying to add values of integer and floating point, we get an error. The *error message* contains much information. First, it states that the error is in `stdin(4,13)`, which means that the error was found on standard-input at line 4 and column 13. Since the program was executed using `fsharp` `quickStartSumFloat.fsx`, then here standard input means the file `quickStartSumFloat.fsx` shown in Listing 4.3. The corresponding line and column are also shown in Listing 4.4. After the file, line, and column number, F# informs us of the error number and a description of the error. Error numbers are an underdeveloped feature in Mono and should be ignored. However, the verbal description often contains useful information for *debugging*. In the example we are informed that there is a type mismatch in the expression, i.e., since `a` is an integer, then F# had expected `b` to be one too. Debugging is the process of solving errors in programs, and here we can solve the error by either making `a` into a float or `b` into an int. The right solution depends on the application.

· error message

· debugging

F# is a functional first programming language, and one implication of this is that names have a *lexical scope*. A scope is the lines in a program, where a binding is valid, and lexical scope means that to find the value of a name F# looks for the value in the above lines. Further, at the outer most level, rebinding is not allowed. If attempted, then F# will return an error as shown in Listing 4.5.

· lexical scope

Listing 4.5 quickStartRebindError.fsx:
A name cannot be rebound.

```

1  let a = 357
2  let a = 864
-----
1  $ fsharp --nologo -a quickStartRebindError.fsx
2
3  quickStartRebindError.fsx(2,5): error FS0037: Duplicate definition of
    value 'a'

```

However, if the same is performed in an interactive session, then rebinding does not cause an error as shown in Listing 4.6.

Listing 4.6: Names may be reused when separated by the lexeme “;;”.

```

1  > let a = 357;;
2  val a : int = 357
3
4  > let a = 864;;
5  val a : int = 864

```

The difference is that the “;;” *lexeme* is used to specifies the end of a *script-fragment*. A lexeme is a letter or a word, which the F# considers as an atomic unit. Script-fragments may be defined both in scripts and in interactive mode, and rebinding is not allowed at the outermost level in script-fragments. Even with the “;;” lexeme, rebinding is not allowed in compile-mode. In general, **avoid rebinding of names**. Advice

In F#, *functions* are also values, and we may define a function `sum` as part of the solution to the above program as shown in Listing 4.7. · function

Listing 4.7 quickStartSumFct.fsx:
A script to add 2 numbers using a user-defined function.

```

1  let sum x y = x + y
2  let c = sum 357 864
3  do printfn "%A" c
-----
1  $ fsharp --nologo quickStartSumFct.fsx && mono quickStartSumFct.exe
2  1221

```

Functions are useful to *encapsulate* code, such that we can focus on the transformation of data by a function while ignoring the details on how this is done. Functions are also useful for code reuse, i.e., instead of repeating a piece of code in several places, such code can be encapsulated in a function and replaced with function calls. This makes debugging and maintenance considerably simpler. Entering the function into an interactive session will illustrate the inferred type, the function `sum` has: `val sum : x:int -> y:int -> int`. The “->” is the mapping operator in the sense that functions are mappings between sets. The type of the function `sum`, should be read as `val sum : x:int -> (y:int -> int)`, that is, `sum` takes an integer and returns a function, which takes an integer and returns an integer. This is an example of a higher-order function. · encapsulate

Type inference in F# may cause problems since the type of a function is inferred in the context, in which it is defined. E.g., in an interactive session, defining the `sum` in one scope on a single line will default the types to integers, F#'s favorite type. Thus, if the next script-fragment uses the function with floats, then we will get an error message as shown in Listing 4.8.

Listing 4.8: Types are inferred in blocks, and F# tends to prefer integers.

```

1  val sum : x:int -> y:int -> int
2
3  > let c = sum 357.6 863.4;;
4      let c = sum 357.6 863.4;;
5      -----^^^
6
7  stdin(3,13): error FS0001: This expression was expected to have type
8      'int'
9  but here has type
10     'float'
```

A remedy is to define the function in the same script-fragment as it is used such as shown in Listing 4.9.

Listing 4.9: Type inference is per script-fragment.

```

1  > let sum x y = x + y
2  - let c = sum 357.6 863.4;;
3  val sum : x:float -> y:float -> float
4  val c : float = 1221.0
```

Alternatively, the types may be explicitly stated as shown in Listing 4.10.

Listing 4.10: Function argument and return types may be stated explicitly.

```

1  > let sum (x : float) (y : float) : float = x + y;;
2  val sum : x:float -> y:float -> float
3
4  > let c = sum 357.6 863.4;;
5  val c : float = 1221.0
```

The function `sum` has two arguments and a return type and in Listing 4.10 we have specified all three. This is done using the “:” lexeme, and to resolve confusion, we must use parentheses around the arguments such as `(y : float)`, otherwise F# would not be able to understand, whether the type annotation was for the argument or the return value. Often it is sufficient to specify some of the types since type inference will enforce the remaining types. E.g., in this example, the “+” operator is defined for identical types, so specifying the return value of `sum` to be a float, implies that the result of the “+” operator is a float, and therefore its arguments must be floats, and finally then the arguments for `sum` must be floats. However, in this book we advocate the following advice: **specify types unless explicitly working with generic functions.** Advice

In this chapter, we have scratched the surface of learning how to program by concentrating on a number of key programming concepts and how they are expressed in the F# language. In the following chapters, we will expand the description of F# with features used in all programming approaches.

5 | Using F# as a calculator

In this chapter, we will exclusively use the interactive mode to illustrate basic types and operations in F#.

5.1 Literals and basic types

All programs rely on processing of data, and an essential property of data is its *type*. A *literal* is a fixed value like the number 3, and if we type the number 3 in an interactive session at the input prompt, then F# responds as shown in Listing 5.1.

Listing 5.1: Typing the number 3.

```
1 > 3;;
2 val it : int = 3
```

What this means is that F# has inferred the type to be *int* and bound it to the identifier *it*. For more on binding and identifiers see Chapter 6. Types matter, since the operations that can be performed on integers, are quite different from those that can be performed on, e.g., strings. E.g., the number 3 has many different representations as shown in Listing 5.2.

Listing 5.2: Many representations of the number 3 but using different types.

```
1 > 3;;
2 val it : int = 3
3
4 > 3.0;;
5 val it : float = 3.0
6
7 > '3';;
8 val it : char = '3'
9
10 > "3";;
11 val it : string = "3"
```

Each literal represents the number 3, but their types are different, and hence they are quite different values. The types *int* for integer numbers, *float* for floating point numbers, *bool* for Boolean values, *char* for characters, and *string* for strings of characters are the most common types of literals. A table of all *basic types* predefined in F# is given in Table 5.1. Besides these built-in types, F# is designed such that it is easy to define new types.

· *float*
· *bool*
· *char*
· *string*
· basic types

Metatype	Type name	Description
Boolean	<u>bool</u>	Boolean values true or false
Integer	<u>int</u>	Integer values from -2,147,483,648 to 2,147,483,647
	byte	Integer values from 0 to 255
	sbyte	Integer values from -128 to 127
	int8	Synonymous with sbyte
	uint8	Synonymous with byte
	int16	Integer values from -32768 to 32767
	uint16	Integer values from 0 to 65535
	int32	Synonymous with int
	uint32	Integer values from 0 to 4,294,967,295
Real	int64	Integer values from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
	uint64	Integer values from 0 to 18,446,744,073,709,551,615
	<u>float</u>	64-bit IEEE 754 floating point value from $-\infty$ to ∞
	double	Synonymous with float
	single	A 32-bit floating point type
Character	float32	Synonymous with single
	decimal	A floating point data type that has at least 28 significant digits
None	<u>char</u>	Unicode character
	<u>string</u>	Unicode sequence of characters
Object	<u>unit</u>	The value ()
Exception	<u>obj</u>	An object
	<u>exn</u>	An exception

Table 5.1: List of some of the basic types. The most commonly used types are underlined. For a description of integer see Appendix B.1, for floating point numbers see Appendix B.2, for ASCII and Unicode characters see Appendix C, for objects see Chapter 20, and for exceptions see Chapter 18.

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base 10*, which means that a value is represented as two sequences of decimal digits separated by a *decimal point*, where each *digit* d has a position and a value $d \in \{0, 1, 2, \dots, 9\}$. The part before the decimal point is called the *whole part* and the part after is called the *fractional part* of the number. The whole part with neither a decimal point nor a fractional part is called an *integer*. As an example 35.7 is a decimal number, whose value is $3 \cdot 10^1 + 5 \cdot 10^0 + 7 \cdot 10^{-1}$, and 128 is an integer, whose value is $1 \cdot 10^2 + 2 \cdot 10^1 + 8 \cdot 10^0$. In F# a decimal number is called a *floating point number*. Floating point numbers may alternatively be given using *scientific notation*, such as 3.5e-4 and 4e2, where the e-notation is translated to a value as $3.5e-4 = 3.5 \cdot 10^{-4} = 0.00035$, and $4e2 = 4 \cdot 10^2 = 400$.

The basic unit of information in almost all computers is the binary digit or *bit* for short. Internally, programs and data are all represented as bits, hence F# has a strong support for binary numbers. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. E.g., the binary number $101.01_2 = 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2} = 5.25$. Subscripts are often used to indicate the base of a number, e.g., 101.01_2 and 101.01_{10} are different numbers. Since base 10 is so common, the subscript for base 10 numbers is often omitted.

Binary numbers are closely related to *octal* and *hexadecimal numbers*, where octals uses 8 as basis, and where each octal digit can be represented by exactly 3 bits, while hexadecimal numbers use 16 as basis, and where each hexadecimal digit can be written in binary using exactly 4 bits. The hexadecimal digits use 0–9 to represent the values 0–9 and a–f in lower or alternatively upper case to represent the values 10–15. Thus, Octals and hexadecimals conveniently serve as shorthand for the much longer binary representation. As examples, the octal number 37_8 is $3 \cdot 8^1 + 7 \cdot 8^0 = 31$, and the hexadecimal number $f3_{16}$ is $15 \cdot 16^1 + 3 \cdot 16^0 = 243$.

- decimal number
- base
- decimal point
- digit
- whole part
- fractional part
- integer
- floating point number
- scientific notation
- bit
- binary number
- octal number
- hexadecimal number

Character	Escape sequence	Description
BS	<code>\b</code>	Backspace
LF	<code>\n</code>	Line feed
CR	<code>\r</code>	Carriage return
HT	<code>\t</code>	Horizontal tabulation
<code>\</code>	<code>\\</code>	Backslash
<code>"</code>	<code>\"</code>	Quotation mark
<code>'</code>	<code>\'</code>	Apostrophe
BEL	<code>\a</code>	Bell
FF	<code>\f</code>	Form feed
VT	<code>\v</code>	Vertical tabulation
	<code>\uXXXX</code> , <code>\UXXXXXXXX</code> , <code>\DDD</code>	Unicode character

Table 5.2: Escape characters. For the unicode characters 'X' are hexadecimal digits, while for tricode characters 'D' is a decimal character.

To denote integers on bases different than 10, F# uses the prefix '0b' for binary, '0o' for octal, and '0x' for hexadecimal numbers. For example, the value 367_{10} may be written as an integer `367`, as a binary number `0b101101111`, as a octal number `0o557`, and as a hexadecimal number `0x16f`. The character sequences `0b12` and `ff` are not numbers.

A *character* is a *Unicode code point*, and character literals are enclosed in single quotation marks. Appendix C.3 contains more details on code points. The character type in F# is denoted *char*. Examples of characters are 'a', 'D', '3'. However '23' and 'abc' are not characters. Some characters do not have a visual representation such as the tabulation character. These can still be represented as a character using *escape sequences*. A character escape sequence starts with “\” followed by either a letter for simple escapes such as `\t` for tabulation and `\n` for newline. Escape sequences can also be a numerical representation of a code point, and three versions exist: The trigraph `\DDD`, where D is a decimal digit, is used to specify the first 256 code points, the hexadecimal escape codes `\uXXXX`, where X is a hexadecimal digit, is used to specify the first 65536 code points, and `\UXXXXXXXX` is used to specify any of the approximately $4.3 \cdot 10^9$ possible code points. All escape sequences are shown in Table 5.2. Examples of *char* representations of the letter 'a' are: `'a'`, `'\097'`, `'\u0061'`, `'\U00000061'`.

- character
- Unicode
- code point
- *char*
- escape sequences

A *string* is a sequence of characters enclosed in double quotation marks. Examples are `"a"`, `"this is a string"`, and `"-&\#@"`. Note that the string `"a"` and the character `'a'` are not the same. Some strings are so common that they are given special names: One or more spaces, `" "` is called *whitespace*, and both `"\n"` and `"\r\n"` are called *newline*. The escape-character “\” may be used to break a line in two. This and other examples are shown in Listing 5.3.

- string
- whitespace
- newline

Type	syntax	Examples	Value
int, int32	<int or hex> <int or hex>l	3, 0x3 3l, 0x3l	3
uint32	<int or hex>u <int or hex>ul	3u 3ul	3
byte, uint8	<int or hex>uy '<char>'B	97uy 'a'B	97
byte[]	"<string>"B @"<string>"B	"a\n"B @"a\n"B	[97uy; 10uy] [97uy; 92uy; 110uy]
sbyte, int8	<int or hex>y	3y	3
int16	<int or hex>s	3s	3
uint16	<int or hex>us	3us	3
int64	<int or hex>L	3L	3
uint64	<int or hex>UL <int or hex>uL	3UL 3uL	3
float, double	<float> <hex>LF	3.0 0x013fLF	3.0 9.387247271e-323
single, float32	<float>F <float>f <hex>lf	3.0F 3.0f 0x013flf	3.0 3.0 4.4701421e-43f
decimal	<float or int>M <float or int>m	3.0M,3M 3.0m,3m	3.0
string	"<string>" @"<string>" "<string>"	"a \"quote\".\n" @"a \"quote\".\n" ""a \"quote\".\n""	a "quote".<newline> a "quote".\n. a "quote".\n

Table 5.3: List of literal type. Syntax notation is used such that, e.g., <> means that the programmer replaces the brackets and content with a value on appropriate form. The [| |] notation means that the value is an array, see Section 11.3 for details.

Listing 5.3: Examples of string literals.

```

1 > "abcde";;
2 val it : string = "abcde"
3
4 > "abc
5 -   de";;
6 val it : string = "abc
7   de"
8
9 > "abc\
10 -   de";;
11 val it : string = "abcde"
12
13 > "abc\nde";;
14 val it : string = "abc
15 de"

```

Note that the response from `fsharp` is shown in double quotation marks, but this is not part of the string.

F# supports *literal types*, where the type of a literal is indicated as a prefix or suffix as shown in the Table 5.3. The table uses a simple syntax notation such that <integer or hexadecimal>UL means that the user supplies an integer or a hexadecimal number followed by the characters 'UL'.

The literal type is closely connected to how the values are represented internally. E.g., a value of type `int32` use 32 bits and can be both positive and negative, while a `uint32` value also use 32 bits, but is unsigned. A `byte` is an 8-bit number, and `sbyte` is a signed 8-bit number. Values of type `float` uses 64 bits, while `float32` only uses 32 bits. The number of bits used to represent numbers directly relates to the range and precision these types can represent. This is summarized in Table 5.1 and discussed in more detail in Appendix B. String literals may be *verbatim* by the `@`-notation or triple double quotation marks, meaning that the escape sequences are not converted to their code point. The two types of string verbatim treat quotation marks differently as illustrated in the table. Further examples are shown in Listing 5.4.

Listing 5.4: Named and implied literals.

```

1  > 3;;
2  val it : int = 3
3
4  > 4u;;
5  val it : uint32 = 4u
6
7  > 5.6;;
8  val it : float = 5.6
9
10 > 7.9f;;
11 val it : float32 = 7.9000001f
12
13 > 'A';;
14 val it : char = 'A'
15
16 > 'B'B;;
17 val it : byte = 66uy
18
19 > "ABC";;
20 val it : string = "ABC"
21
22 > @"abc\nde";;
23 val it : string = "abc\nde"

```

Many basic types are compatible, and the type of a literal may be changed by *typecasting*. An example of casting to a `float` is shown in Listing 5.5.

Listing 5.5: Casting an integer to a floating point number.

```

1  > float 3;;
2  val it : float = 3.0

```

When `float` is given an argument, then it acts as a function rather than a type, and for the integer 3 it returns the floating point number 3.0. For more on functions see Chapter 6. Boolean values are often treated as the integer values 0 and 1, but no short-hand function names exist for their conversions. Instead, use functions from the `System.Convert` family of functions, as demonstrated in Listing 5.6.

Listing 5.6: Casting booleans.

```

1 > System.Convert.ToBoolean 1;;
2 val it : bool = true
3
4 > System.Convert.ToBoolean 0;;
5 val it : bool = false
6
7 > System.Convert.ToInt32 true;;
8 val it : int = 1
9
10 > System.Convert.ToInt32 false;;
11 val it : int = 0

```

Here `System.Convert.ToBoolean` is the identifier of a function `ToBoolean`, which is a *member* of the *class* `Convert` that is included in the *namespace* `System`. Namespaces, classes, and members will be discussed in Chapter 9.

- member
- class
- namespace

Typecasting is often a destructive operation, e.g., typecasting a `float` to `int` removes the fractional part without rounding as shown in Listing 5.7.

Listing 5.7: Fractional part is removed by downcasting.

```

1 > int 357.6;;
2 val it : int = 357

```

Here we typecasted to a lesser type, in the sense that the set of integers is a subset of floating point numbers, and this is called *downcasting*. The opposite is called *upcasting* and is often non-destructive, as Listing 5.5 showed, where an integer was cast to a float while retaining its value. As a side note, *rounding* a number $y.x$, where y is the *whole part* and x is the *fractional part*, is the operation of mapping numbers in the interval $y.x \in [y.0, y.5)$ to y and $y.x \in [y.5, y + 1)$ to $y + 1$. This can be performed by downcasting as shown in Listing 5.8.

- downcasting
- upcasting
- rounding
- whole part
- fractional part

Listing 5.8: Fractional part is removed by downcasting.

```

1 > int (357.6 + 0.5);;
2 val it : int = 358

```

As the example shows, for floating points whose fractional part is equal to or larger than $y.5$ adding 0.5 will make them above $(y + 1).0$, and downcasting will thus downcase to $(y + 1).0$. Conversely, fractional parts below will downcast to $y.0$. Thus, rounding is achieved by downcasting.

5.2 Operators on basic types

Listing 5.8 is an example of an arithmetic *expression* using an *binary operator* written using *infix notation*, since the operator appears in between the *operands*. The “+” operator is binary since it takes two arguments, and since it is written between its arguments, then it uses infix notation. Expressions are the basic building block of all F# programs and this section will discuss operator expressions on basic types.

- expression
- binary operator
- infix notation
- operands

The syntax of basic binary operators is shown in the following:

Listing 5.9 Syntax for a binary expression.

```
1 <expr><op><expr>
```

Here `<expr>` is any expression supplied by the programmer, and `<op>` is a binary, infix operator. F# supports a range of arithmetic binary infix operators on its built-in types such as addition, subtraction, multiplication, division, and exponentiation using the “+”, “-”, “*”, “/”, “**” lexemes. Not all operators are defined for all types, e.g., addition is defined for integer and float types as well as for characters and strings, but multiplication is only defined for integer and floating-point types. A complete list of built-in operators on basic types is shown in Table E.1 and E.2 and a range of mathematical functions shown in Table E.3. An example is `3+4`. Note that expressions can themselves be arguments to expressions, and thus, `4+5+6` is also a legal statement. This is called *recursion*, which means that a rule or a function is used by the rule or function itself in its definition. See Chapter 13 for more on recursive functions.

Unary operators take only one argument and have the syntax:

Listing 5.10 A unary expressions.

```
1 <op><expr>
```

An example of a unary operator is `-3`, where `-` here is used to negate a positive integer. Since the operator appears before the operand it is a *prefix operator*.

The concept of *precedence* is an important concept in arithmetic expressions.¹ If parentheses are omitted in Listing 5.8, then F# will interpret the expression as `(int 357.6) + 0.5`, which is erroneous since the addition of an integer with a float is undefined. This is an example of precedence, i.e., function evaluation takes precedence over addition meaning that it is performed before addition. Consider the arithmetic expression,

Listing 5.11: A simple arithmetic expression.

```
1 > 3 + 4 * 5;;
2 val it : int = 23
```

Here, the addition and multiplication functions are shown in infix notation with the *operator* lexemes “+” and “*”. To arrive at the resulting value 23, F# has to decide in which order to perform the calculation. There are 2 possible orders, `3 + (4 * 5)` or `(3 + 4) * 5`, which gives different results. For integer arithmetic, the correct order is, of course, to multiply before addition, and we say that multiplication takes *precedence* over addition. Every atomic operation that F# can perform is ordered in terms of its precedence, and for some common built-in operators shown in Table 5.4, the precedence is shown by the order they are given in the table.

Associativity implies the order in which calculations are performed for operators of the same precedence. For some operators and type combinations association matters little, e.g., multiplication associates to the left and exponentiation associates to the right, as demonstrated in Listing 5.12.

¹Jon: minor comment on indexing and slice-ranges.

Operator	Associativity	Description
+<expr>, -<expr>, ~~~<expr>	Left	Unary identity, negation, and bitwise negation operator
f <expr>	Left	Function application
<expr> ** <expr>	Right	Exponent
<expr> * <expr>, <expr> / <expr>, <expr> % <expr>	Left	Multiplication, division and remainder
<expr> + <expr>, <expr> - <expr>	Left	Addition and subtraction binary operators
<expr> ^^^ <expr>	Right	bitwise exclusive or
<expr> < <expr>, <expr> <= <expr>, <expr> > <expr>, <expr> >= <expr>, <expr> = <expr>, <expr> <> <expr>, <expr> <<< <expr>, <expr> >>> <expr>, <expr> &&& <expr>, <expr> <expr> ,	Left	Comparison operators, bitwise shift, and bitwise 'and' and 'or'.
<expr> && <expr>	Left	Boolean and
<expr> <expr>	Left	Boolean or

Table 5.4: Some common operators, their precedence, and their associativity. Rows are ordered from highest to lowest precedences, such that <expr> * <expr> has higher precedence than <expr> + <expr>. Operators in the same row have the same precedence. Full table is given in Table E.5.

Listing 5.12: Precedence rules define implicit parentheses.

```

1  > 3.0*4.0*5.0;;
2  val it : float = 60.0
3
4  > (3.0*4.0)*5.0;;
5  val it : float = 60.0
6
7  > 3.0*(4.0*5.0);;
8  val it : float = 60.0
9
10 > 4.0 ** 3.0 ** 2.0;;
11 val it : float = 262144.0
12
13 > (4.0 ** 3.0) ** 2.0;;
14 val it : float = 4096.0
15
16 > 4.0 ** (3.0 ** 2.0);;
17 val it : float = 262144.0

```

The expression for $3.0 * 4.0 * 5.0$ associates to the left, and thus is interpreted as $(3.0 * 4.0) * 5.0$, but gives the same results as $3.0 * (4.0 * 5.0)$, since association does not matter for multiplication of numbers. However, the expression for $4.0 ** 3.0 ** 2.0$ associates to the right, and thus is interpreted as $4.0 ** (3.0 ** 2.0)$, which is quite different from $(4.0 ** 3.0) ** 2.0$. **Whenever in doubt of association or any other basic semantic rules, it is a good idea to use parentheses as here. It is also a good idea to test your understanding of the syntax and semantic** Advice

a	b	a && b	a b	not a
false	false	false	false	true
false	true	false	true	true
true	false	false	true	false
true	true	true	true	false

Table 5.5: Truth table for boolean 'and', 'or', and 'not' operators. Value 0 is false and 1 is true.

rules by making a simple script.

5.3 Boolean arithmetic

Boolean arithmetic is the basis of almost all computers and particularly important for controlling program flow, which will be discussed in Chapter 8. Boolean values are one of 2 possible values, true or false, which is also sometimes written as 1 and 0. Basic operations on Boolean values are 'and', 'or', and 'not', which in F# are written respectively as the binary operators `&&`, `||`, and the function `not`. Since the domain of Boolean values is so small, all possible combination of input on these values can be written on the tabular form, known as a *truth table*, and the truth tables for the basic Boolean operators and functions are shown in Table 5.5. A good mnemonic for remembering the result of the 'and' and 'or' operators is to use 1 for true, 0 for false, multiplication for the Boolean 'and' operator, and addition for Boolean 'or' operator, e.g., true and false in this mnemonic translates to $1 \cdot 0 = 0$, and the result translates back to the Boolean value false. In F# the truth table for the basic Boolean operators can be produced by a program as shown in Listing 5.13.

Listing 5.13: Boolean operators and truth tables.

```
1 > printfn "a b a*b a+b not a"
2 - printfn "%A %A %A %A %A"
3 -   false false (false && false) (false || false) (not false)
4 - printfn "%A %A %A %A %A"
5 -   false true (false && true) (false || true) (not false)
6 - printfn "%A %A %A %A %A"
7 -   true false (true && false) (true || false) (not true)
8 - printfn "%A %A %A %A %A"
9 -   true true (true && true) (true || true) (not true);;
10 a b a*b a+b not a
11 false false false false true
12 false true false true true
13 true false false true false
14 true true true true false
15 val it : unit = ()
```

Here, we used the `printfn` function, to present the results of many expressions on something that resembles a tabular form. The spacing produced using the `printfn` function is not elegant, and in Section 6.5 we will discuss better options for producing more beautiful output. Notice that the arguments for `printfn` was given on the next line with indentation. The indentation is an important part of telling F# which part of what you write belongs together. This is an example of the so-called lightweight syntax. Generally, F# ignores newlines and whitespaces except when using the lightweight syntax. The difference between verbose and lightweight syntax is discussed in Chapter 6.

5.4 Integer arithmetic

The set of integers is infinitely large, but since all computers have limited resources, it is not possible to represent it in its entirety. The various integer types listed in Table 5.1 are finite subsets reduced by limiting their ranges. An in-depth description of integer implementation can be found in Appendix B. The type `int` is the most common type.

Table E.1, E.2, and E.3 give examples operators and functions pre-defined for integer types. Notice that fewer functions are available for integers than for floating point numbers. For most addition, subtraction, multiplication, and negation, the result is straightforward. However, performing arithmetic operations on integers requires extra care, since the result may cause *overflow* and *underflow*. E.g., the range of the integer type `sbyte` is $[-128 \dots 127]$, which causes problems in the example in Listing 5.14. · overflow
· underflow

Listing 5.14: Adding integers may cause overflow.

```
1 > 100y;;
2 val it : sbyte = 100y
3
4 > 30y;;
5 val it : sbyte = 30y
6
7 > 100y + 30y;;
8 val it : sbyte = -126y
```

Here $100 + 30 = 130$, which is larger than the biggest `sbyte`, and the result is an overflow. Similarly, we get an underflow, when the arithmetic result falls below the smallest value storable in an `sbyte` as demonstrated in Listing 5.15.

Listing 5.15: Subtracting integers may cause underflow.

```
1 > -100y - 30y;;
2 val it : sbyte = 126y
```

I.e., we were expecting a negative number but got a positive number instead.

The overflow error in Listing 5.14 can be understood in terms of the binary representation of integers: In binary, $130 = 10000010_2$, and this binary pattern is interpreted differently as `byte` and `sbyte`, see Listing 5.16.

Listing 5.16: The leftmost bit is interpreted differently for signed and unsigned integers, which gives rise to potential overflow errors.

```
1 > 0b10000010uy;;
2 val it : byte = 130uy
3
4 > 0b10000010y;;
5 val it : sbyte = -126y
```

That is, for signed bytes, the left-most bit is used to represent the sign, and since the addition of $100 = 01100100_2$ and $30 = 00011110_2$ is $130 = 10000010_2$ causes the left-most bit to be used, then this is wrongly interpreted as a negative number when stored in an `sbyte`. Similar arguments can be made explaining underflows.

The division and remainder operators, which discards the fractional part after division, and the *remainder* operator calculates the remainder after integer division, as demonstrated in Listing 5.17. · integer division · remainder

Listing 5.17: Integer division and remainder operators.

```
1 > 7 / 3;;
2 val it : int = 2
3
4 > 7 % 3;;
5 val it : int = 1
```

Together, integer division and the remainder is a lossless representation of the original number, see Listing 5.18.

Listing 5.18: Integer division and remainder is a lossless representation of an integer, compare with Listing 5.17.

```
1 > (7 / 3) * 3;;
2 val it : int = 6
3
4 > (7 / 3) * 3 + (7 % 3);;
5 val it : int = 7
```

And we see that integer division of 7 by 3 followed by multiplication by 3 is less than 7, and the difference is $7 \% 3$.

Notice that neither overflow nor underflow error gave rise to an error message, which is why such bugs are difficult to find. Dividing any non-zero number with 0 is infinite, which is also outside the domain of any of the integer types, but in this case, F# casts an *exception* as shown in Listing 5.19. · exception

Listing 5.19: Integer division by zero causes an exception runtime error.

```
1 > 3/0;;
2 System.DivideByZeroException: Attempted to divide by zero.
3   at <StartupCode$FSI_0002>.$FSI_0002.main@ () [0x00001] in
   <da48886c466e4b80be4566752c596fa8>:0
4   at (wrapper managed-to-native)
   System.Reflection.MonoMethod:InternalInvoke
   (System.Reflection.MonoMethod,object[],System.Exception&)
5   at System.Reflection.MonoMethod.Invoke (System.Object obj,
   System.Reflection.BindingFlags invokeAttr, System.Reflection.Binder
   binder, System.Object[] parameters, System.Globalization.CultureInfo
   culture) [0x00032] in <c9f8153c41de4f8cbafd0e32f9bf6b28>:0
6 Stopped due to error
```

The output looks daunting at first sight, but the first and last lines of the error message are the most important parts, which tells us what exception was cast and why the program stopped. The middle is technical details concerning which part of the program caused this and can be ignored for the time being. Exceptions are a type of *runtime error*, and are treated in Chapter 18 · runtime error

Integer exponentiation is not defined as an operator, but this is available as the built-in function `pown`. This function is demonstrated in Listing 5.20 for calculating 2^5 .

a	b	a ~~~ b
false	false	false
false	true	true
true	false	true
true	true	false

Table 5.6: Boolean exclusive or truth table.

Listing 5.20: Integer exponent function.

```

1 > pown 2 5;;
2 val it : int = 32

```

For binary arithmetic on integers, the following operators are available: `<leftExpr> <<< <rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the left while inserting 0's to right; `<leftExpr> >>> <rightExpr>`, which shifts the bit pattern of `<leftExpr>` `<rightExpr>` positions to the right while inserting 0's to left; `<expr> &&& <expr>`, bitwise 'and', returns the result of taking the Boolean 'and' operator position-wise; `<expr> ||| <expr>`, bitwise 'or', as 'and' but using the Boolean 'or' operator; and `<expr> ~~~ <expr>`, bitwise xor, which is returns the result of the Boolean 'xor' operator defined by the truth table in Table 5.6.

· xor
· exclusive or

5.5 Floating point arithmetic

Like integers, the set of reals is also infinitely large, hence, floating point types are finite subsets reduced by sampling the space of reals. An in-depth description of floating point implementations can be found in Appendix B. The type `float` is the most common type.

Table E.1, E.2, and E.3 give examples operators and functions pre-defined for floating point types. Note that the remainder operator for floats calculates the remainder after division and discarding the fractional part, see Listing 5.21.

Listing 5.21: Floating point division and remainder operators.

```

1 > 7.0 / 2.5;;
2 val it : float = 2.8
3
4 > 7.0 % 2.5;;
5 val it : float = 2.0

```

The remainder for floating point numbers can be fractional, but division, discarding fractional part, and the remainder is still a lossless representation of the original number as demonstrated in Listing 5.22.

Listing 5.22: Floating point division, truncation, and remainder is a lossless representation of a number.

```

1 > float (int (7.0 / 2.5));;
2 val it : float = 2.0
3
4 > (float (int (7.0 / 2.5))) * 2.5;;
5 val it : float = 5.0
6
7 > (float (int (7.0 / 2.5))) * 2.5 + 7.0 % 2.5;;
8 val it : float = 7.0

```

Arithmetic using `float` will not cause over- and underflow problems, since the IEEE 754 standard includes the special numbers $\pm\infty$ and NaN. As shown in Listing 5.23, no exception is thrown.

Listing 5.23: Floating point numbers include infinity and Not-a-Number.

```

1 > 1.0/0.0;;
2 val it : float = infinity
3
4 > 0.0/0.0;;
5 val it : float = nan

```

However, the `float` type has limited precision, since there is only a finite number of numbers that can be stored in a float. E.g., addition and subtraction can give surprising results as demonstrated in Listing 5.24.

Listing 5.24: Floating point arithmetic has finite precision.

```

1 > 357.8 + 0.1 - 357.9;;
2 val it : float = 5.684341886e-14

```

That is, addition and subtraction associates to the left, hence the expression is interpreted as $(357.8 + 0.1) - 357.9$ and we see that we do not get the expected 0. The reason is that the calculation is done stepwise, and in the process, the numbers are represented using the imprecise floating point standard. Thus, $357.8 + 0.1$ is represented as a number close to but not identical to what 357.9 is represented as, and thus, when subtracting these two representations, we get a very small number but not 0. Such errors tend to accumulate and comparing the result of expressions of floating point values should, therefore, be treated with care. Thus, **equivalence of two floating point expressions should only be considered up to sufficient precision, e.g., comparing $357.8 + 0.1$ and 357.9 up to $1e-10$ precision should be tested as, `abs ((357.8 + 0.1) - 357.9) < 1e-10`.** Advice

5.6 Char and string arithmetic

Addition is the only operator defined for characters. Nevertheless, character arithmetic is often done by casting to an integer. A typical example is the conversion of character case, e.g., to convert the lowercase character 'z' to uppercase, we use the *ASCIIbetical order* and add the difference between any Basic Latin Block letters in upper- and lowercase as `integers` and cast back to `char`, see Listing 5.25. · ASCIIbetical order

Listing 5.25: Converting case by casting and integer arithmetic.

```

1 > char (int 'z' - int 'a' + int 'A');;
2 val it : char = 'Z'

```

I.e., the code point difference between upper and lower case for any alphabetical character 'a' to 'z' is constant, hence we can change case by adding or subtracting the difference between any corresponding character. Unfortunately, this does not generalize to characters from other languages.

A large collection of operators and functions exist for `string`. The simplest is concatenation using the “+” operator as demonstrated in Listing 5.26.

Listing 5.26: Example of string concatenation.

```

1 > "hello" + " " + "world";;
2 val it : string = "hello world"

```

Characters and strings cannot be concatenated, which is why the above example used the string of a space " " instead of the space character ' '. The characters of a string may be indexed as using the `.[]` notation. This is demonstrated in Listing 5.27.

Listing 5.27: String indexing using square brackets.

```

1 > "abcdefg".[0];;
2 val it : char = 'a'
3
4 > "abcdefg".[3];;
5 val it : char = 'd'
6
7 > "abcdefg".[3..];;
8 val it : string = "defg"
9
10 > "abcdefg".[..3];;
11 val it : string = "abcd"
12
13 > "abcdefg".[1..3];;
14 val it : string = "bcd"
15
16 > "abcdefg".[*];;
17 val it : string = "abcdefg"

```

Notice that the first character has index 0, and to get the last character in a string, we use the string's `length` property. This is done as shown in Listing 5.28.

Listing 5.28: String length attribute and string indexing.

```

1 > "abcdefg".Length;;
2 val it : int = 7
3
4 > "abcdefg".[7-1];;
5 val it : char = 'g'

```

Since index counting starts at 0, and the string length is 7, then the index of the last character is 6. There is a long list of built-in functions in `System.String` for working with strings, some of which will be discussed in Section 11.1.

The *dot notation* is an example of Structured programming, where technically speaking, the string `"abcdefg"` is an immutable *object* of *class* `string`, `[]` is an object *method*, and `Length` is a property. For more on object, classes, and methods see Chapter 20.

· dot notation
· object
· class
· method

Strings are compared letter by letter. For two strings to be equal, they must have the same length and all the letters must be identical. E.g., `"abs" = "absalon"` is false, while `"abs" = "abs"` is true. The `<>` operator is the boolean negation of the `=` operator, e.g., `"abs" <> "absalon"` is true, while `"abs" <> "abs"` is false. For the `<`, `<=`, `>`, and `>=` operators, the strings are ordered alphabetically, such that `"abs" < "absalon" && "absalon" < "milk"` is true, that is, the `<` operator on two strings is true, if the left operand should come before the right, when sorting alphabetically. The algorithm for deciding the boolean value of `leftOp < rightOp` is as follows: we start by examining the first character, and if `leftOp.[0]` and `rightOp.[0]` are different, then the `leftOp < rightOp` is equal to `leftOp.[0] < rightOp.[0]`. E.g., `"milk" < "abs"` is the same as `'m' < 'a'`, which is false, since the letter 'm' does not come before the letter 'a' in the alphabet, or more precisely, the codepoint of 'm' is not less than the codepoint of 'a'. If `leftOp.[0]` and `rightOp.[0]` are equal, then we move onto the next letter and repeat the investigation, e.g., `"abe" < "abs"` is true, since `"ab" = "ab"` is true and `'e' < 's'` is true. If we reach the end of either of the two strings, then the short is smaller than the larger, e.g., `"abs" < "absalon"` is true, while `"abs" < "abs"` is false. The `<=`, `>`, and `>=` operators are defined similarly.


5.7 Programming intermezzo: Hand conversion between decimal and binary numbers

Conversion of integers between decimal and binary form is a key concept in order to understand some of the basic properties of calculations on the computer. From binary to decimal is straightforward using the power-of-two algorithm, i.e., given a sequence of $n + 1$ bits that represent an integer $b_n b_{n-1} \dots b_0$, where b_n and b_0 are the most and least significant bits, then the decimal value is calculated as,

$$v = \sum_{i=0}^n b_i 2^i \quad (5.1)$$

For example, $10011_2 = 1 + 2 + 16 = 19$. From decimal to binary is a little more complex, but a simple divide-by-two algorithm exists. The key to understanding the divide-by-two algorithm is to realize that when you divide a number by two, then that is equivalent to shifting its binary representation 1 to the right. E.g., $10 = 1010_2$ and $10/2 = 5 = 101_2$. Odd numbers have $b_0 = 1$, e.g., $11_{10} = 1011_2$ and $11_{10}/2 = 5.5 = 101.1_2$. Hence, if we divide any number by two and get a non-integer number, then its least significant bit was 1. Another way to express this is that the least significant bit is the remainder after integer division by two. Sequential application of this idea leads directly to the divide-by-two algorithm. E.g., if we were to convert the number 11_{10} on decimal form to binary form we would

perform the following steps:



Here we used `div` and `rem` to signify the integer division and remainder operators. The algorithm stops when the result of integer division is zero. Reading off the remainder from below and up we find the sequence 1011_2 , which is the binary form of the decimal number 11_{10} . Using interactive mode, we can calculate the same as shown in Listing 5.29.

Listing 5.29: Converting the number 11_{10} to binary form.

```

1 > printfn "%d, %d" (11 / 2) (11 % 2);;
2 (5, 1)
3 val it : unit = ()
4 > printfn "%d, %d" (5 / 2) (5 % 2);;
5 (2, 1)
6 val it : unit = ()
7 > printfn "%d, %d" (2 / 2) (2 % 2);;
8 (1, 0)
9 val it : unit = ()
10 > printfn "%d, %d" (1 / 2) (1 % 2);;
11 (0, 1)
12 val it : unit = ()

```

Thus, by reading the second integer-response from `printfn` from below and up, we again obtain the binary form of 11_{10} to be 1011_2 . For integers with a fractional part, the divide-by-two may be used on the whole part, while multiply may be used in a similar manner on the fractional part.

6 | Values and functions

In the previous chapter, we saw how to use F# as a calculator working with literals, operators and built-in functions. To save time and make programs easier to read and debug, it is useful to bind expressions to identifiers either as new constants, functions, or operators. As an example, consider the problem,

Problem 6.1

For given set constants a , b , and c , solve for x in

$$ax^2 + bx + c = 0 \tag{6.1}$$

To solve for x we use the quadratic formula from elementary algebra,

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}, \tag{6.2}$$

which gives the general solution for any values of the coefficients. Here, we will assume a positive discriminant, $b^2 - 4ac > 0$. In order to write a program, where the code may be reused later, we define a function `discriminant : float -> float -> float -> float`, that is, a function that takes 3 arguments, `a`, `b`, and `c`, and calculates the discriminant. Details on function definition is given in Section 6.2. Likewise, we will define functions `positiveSolution : float -> float -> float -> float` and `negativeSolution : float -> float -> float -> float`, that also takes the polynomial's coefficients as arguments and calculates the solution corresponding to choosing the positive and negative sign for \pm in the equation. Our solution thus looks like Listing 6.1.

Listing 6.1 identifiersExample.fsx:

Finding roots for quadratic equations using function name binding.

```

1  let discriminant a b c = b ** 2.0 - 4.0 * a * c
2  let positiveSolution a b c = (-b + sqrt (discriminant a b c)) / (2.0 * a)
3  let negativeSolution a b c = (-b - sqrt (discriminant a b c)) / (2.0 * a)
4
5  let a = 1.0
6  let b = 0.0
7  let c = -1.0
8  let d = discriminant a b c
9  let xp = positiveSolution a b c
10 let xn = negativeSolution a b c
11 do printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
12 do printfn "    has discriminant %A and solutions %A and %A" d xn xp

```

```

1  $ fsharp --nologo identifiersExample.fsx && mono identifiersExample.exe
2  0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
3    has discriminant 4.0 and solutions -1.0 and 1.0

```

Here, we have further defined names of values `a`, `b`, and `c` used as input to our functions, and the results of function application are bound to the names `d`, `xn`, and `xp`. The names of functions and values given here are examples of identifiers, and with these, we may reuse the quadratic formulas and calculated values later, while avoiding possible typing mistakes and reducing the amount of code, which needs to be debugged.

The use of identifiers is central in programming. For F# not to be confused by built-in functionality, identifiers must follow a specific set of rules:

Identifier

· identifier

- Identifiers are used as names for values, functions, types etc.
- They must start with a letter or underscore '`_`', but can be followed by zero or more of letters, digits, and a range of special characters except for SP, LF, and CR (space, line feed, and carriage return). Letters are most Unicode codepoints that produce regular letters, see Appendix C.3 for more on codepoints.
- They can also be a sequence of identifiers separated by a period.
- They cannot be keywords, see Table 6.1.

Examples of identifiers are: `a`, `theCharacter9`, `Next_Word`, `_tok`, and `f.sharp.rocks`. Since programmers often work in multilingual environment dominated by the English language **Restrict identifiers to use letters from the English alphabet, numbers, period, and '`_`'**. However, the number of possible identifiers is enormous: the full definition refers to the Unicode general categories described in Appendix C.3, and there are currently 19,345 possible Unicode code points in the letter category and 2,245 possible Unicode code points in the special character category.

Advice

Identifiers may be used to carry information about their intended content and use, and careful selection of identifiers can aid programmers to communicate thoughts about the code. Thus, identifiers are often a word or several concatenated words from the human language. For example in the function definition `let discriminant a b c = b ** 2.0 - 4.0 * a * c`, the function identifier has been chosen to be `discriminant`. F# places no special significance to the word 'discriminant', and the program would work exactly the same, had the function been called `let f a b c = b ** 2.0 - 4.0 * a * c`. However, to programmers, the word 'discriminant' informs us of the intended role of the function and

Type	Keyword
Regular	<code>abstract</code> , <code>and</code> , <code>as</code> , <code>assert</code> , <code>base</code> , <code>begin</code> , <code>class</code> , <code>default</code> , <code>delegate</code> , <code>do</code> , <code>done</code> , <code>downcast</code> , <code>downto</code> , <code>elif</code> , <code>else</code> , <code>end</code> , <code>exception</code> , <code>extern</code> , <code>false</code> , <code>finally</code> , <code>for</code> , <code>fun</code> , <code>function</code> , <code>global</code> , <code>if</code> , <code>in</code> , <code>inherit</code> , <code>inline</code> , <code>interface</code> , <code>internal</code> , <code>lazy</code> , <code>let</code> , <code>match</code> , <code>member</code> , <code>module</code> , <code>mutable</code> , <code>namespace</code> , <code>new</code> , <code>null</code> , <code>of</code> , <code>open</code> , <code>or</code> , <code>override</code> , <code>private</code> , <code>public</code> , <code>rec</code> , <code>return</code> , <code>sig</code> , <code>static</code> , <code>struct</code> , <code>then</code> , <code>to</code> , <code>true</code> , <code>try</code> , <code>type</code> , <code>upcast</code> , <code>use</code> , <code>val</code> , <code>void</code> , <code>when</code> , <code>while</code> , <code>with</code> , and <code>yield</code> .
Reserved	<code>atomic</code> , <code>break</code> , <code>checked</code> , <code>component</code> , <code>const</code> , <code>constraint</code> , <code>constructor</code> , <code>continue</code> , <code>eager</code> , <code>fixed</code> , <code>fori</code> , <code>functor</code> , <code>include</code> , <code>measure</code> , <code>method</code> , <code>mixin</code> , <code>object</code> , <code>parallel</code> , <code>params</code> , <code>process</code> , <code>protected</code> , <code>pure</code> , <code>recursive</code> , <code>sealed</code> , <code>tailcall</code> , <code>trait</code> , <code>virtual</code> , and <code>volatile</code> .
Symbolic	<code>let!</code> , <code>use!</code> , <code>do!</code> , <code>yield!</code> , <code>return!</code> , <code> </code> , <code>-></code> , <code><-</code> , <code>..</code> , <code>:</code> , <code>(</code> , <code>)</code> , <code>[</code> , <code>]</code> , <code>[<</code> , <code>>]</code> , <code>[</code> , <code>]</code> , <code>{</code> , <code>}</code> , <code>'</code> , <code>#</code> , <code>:?></code> , <code>:?</code> , <code>:</code> , <code>></code> , <code>..</code> , <code>::</code> , <code>:=</code> , <code>:::</code> , <code>:</code> , <code>=</code> , <code>_</code> , <code>?</code> , <code>??</code> , <code>(*)</code> , <code><@</code> , <code>@></code> , <code><@@</code> , and <code>@@></code> .
Reserved symbolic	<code>~</code> and <code>`</code>

Table 6.1: Table of (possibly future) *keywords* and symbolic keywords in F#.

thus is much preferred. This is a general principle, **identifier names should be chosen to reflect their semantic value**.. The arguments `a`, `b`, and `c` are short, but adheres to a textbook tradition of elementary algebra. Again, we might as well have used, `let discriminant c a b = a ** 2.0 - 4.0 * c * b`, which is semantically identical to the original expression, but due to tradition, this would confuse most readers of the code. Thus, **identifier names should be chosen consistently with readers tradition**. Finally, identifiers are often concatenations of words, as `positiveSolution` in Listing 6.1. Concatenations can be difficult to read. Without the capitalized second word, we would have had `positivesolution`. This is readable at most times but takes longer time for humans to parse in general. Typical solutions are to use a separator such as `positive_solution`, *lower camel case* also known as *mixed case* as in the example `positiveSolution`, and *upper camel case* also known as *pascal case* as `PositiveSolution`. In this book, we use lower camel case except where F# requires a capital first letter. Again, the choice does not influence what a program does, only how readable it is to a fellow programmer. The important part is that **identifier names consisting of concatenated words are often preferred over few character names, and concatenation should be emphasized, e.g., by camel casing**. The length of identifier names is a balancing act, since when working with large programs, very long identifier names can be tiresome to write, and a common practice is that the length of identifier names is proportional to the complexity of the program. I.e., complex programs use long names, simple use short. What is complex and what is simple is naturally in the eye of the beholder, but when we program, remember that a future reader of the program most likely has not had time to work the problem as long as the programmer, thus **choose identifier names as if you were to explain the meaning of a program to a knowledgeable outsider**.

- Advice
- lower camel case
 - mixed case
 - upper camel case
 - pascal case
- Advice
- Advice

Another key concept in F# is expressions. An expression can be a mathematical expression, such as `3 * 5`, a function application, such as `f3`, and many other things. Central in this chapter is the binding of values and functions to identifiers, which is done with the keyword `let`, e.g., `let a = 1.0`.

Expressions are the main workhorse of F# and have an enormous variety in how they may be written, we will in this book gradually work through some of the more important facets.

Expressions

- expression

- An Expression is a computation such as `3 * 5`.
- They can be value bindings between identifiers and expressions that evaluate to a value or a function, see Sections 6.1 and 6.2.
- They can be do bindings that produce side-effects and whose result is ignored, see Section 6.2
- They can be assignments to variables, see Section 6.1.

- They can be a sequence of expressions separated with “;” lexeme.
- They can be annotated with a type using the “:” lexeme.

· :

Before we begin a deeper discussion on bindings, note that F# adheres to two different syntaxes: *verbose* and *lightweight*. In the verbose syntax, newlines and whitespaces are generally ignored, while in lightweight syntax, certain keywords and lexemes may be replaced by newlines and whitespaces. The lightweight syntax is the most common, but the syntaxes may be mixed, and we will highlight the options, when relevant.

· verbose syntax

· lightweight syntax

6.1 Value bindings

Binding identifiers to literals or expressions that are evaluated to be values, is called *value-binding*, and examples are `let a = 3.0` and `let b = cos 0.9`. Value bindings have the following syntax,

· value-binding

Listing 6.2 Value binding expression.

```
1 let <valueIdent> = <bodyExpr> [in <expr>]
```

The `let` keyword binds a value-identifier `<valueIdent>` with an expression `<bodyExpr>` that evaluates to a value. If the `in` keyword is used, then the value-identifier becomes synonymous with the evaluated value in `<expr>` only. The square bracket notation `[]` means that the enclosed is optional. Here the meaning is that for lightweight syntax, the `in` keyword is replaced with a newline, and the binding is valid in the following lines at the level of scope of the value-binding or deeper *lexically*.

· let

· in

· lexically

The value identifier annotated with a type using the “:” lexeme followed by the name of a type, e.g., `int`. The “_” lexeme may be used as a value-identifier. This lexeme is called the *wildcard* pattern, and for value-bindings it means that the `<bodyExpr>` is evaluated, but the result is discarded. See Chapter 15 for more details on patterns.

· :

· _

· wildcard

For example, letting the identifier `p` be bound to the value `2.0` and using it in an expression is done as shown in Listing 6.3.

Listing 6.3 letValue.fsx:

The identifier `p` is used in the expression following the `in` keyword.

```
1 let p = 2.0 in do printfn "%A" (3.0 ** p)

-----

1 $ fsharp --nologo letValue.fsx && mono letValue.exe
2 9.0
```

F# will ignore most newlines between lexemes, i.e., the above is equivalent to writing as shown in Listing 6.4.

Listing 6.4 letValueLF.fsx:
Newlines after `in` make the program easier to read.

```
1 let p = 2.0 in
2 do printfn "%A" (3.0 ** p)

-----

1 $ fsharpc --nologo letValueLF.fsx && mono letValueLF.exe
2 9.0
```

F# also allows for an alternative notation called *lightweight syntax*, where e.g., the `in` keyword is replaced with a newline, and the expression starts on the next line at the same column as `let` starts in, i.e., the above is equivalent to Listing 6.5.

Listing 6.5 letValueLightWeight.fsx:
Lightweight syntax does not require the `in` keyword, but expression must be aligned with the `let` keyword.

```
1 let p = 2.0
2 do printfn "%A" (3.0 ** p)

-----

1 $ fsharpc --nologo letValueLightWeight.fsx
2 $ mono letValueLightWeight.exe
3 9.0
```

The same expression in interactive mode will also respond the inferred types as, e.g., shown in Listing 6.6.

Listing 6.6: Interactive mode also responds inferred types.

```
1 > let p = 2.0
2 - do printfn "%A" (3.0 ** p);;
3 9.0
4 val p : float = 2.0
5 val it : unit = ()
```

By the `val` keyword in the line `val p : float = 2.0` we see that `p` is inferred to be of type `float` and bound to the value 2.0. The inference is based on the type of the right-hand-side, which is of type `float`. Identifiers may be defined to have a type using the `:` lexeme, but the types on the left-hand-side and right-hand-side of the `=` lexeme must be identical. I.e., mixing types gives an error as shown in Listing 6.7.

Listing 6.7 `letValueTypeError.fsx`:
Binding error due to type mismatch.

```

1  let p : float = 3
2  do printfn "%A" (3.0 ** p)

```

```

1  $ fsharp --nologo letValueTypeError.fsx && mono letValueTypeError.exe
2
3  letValueTypeError.fsx(1,17): error FS0001: This expression was expected
   to have type
4    'float'
5  but here has type
6    'int'

```

Here, the left-hand-side is defined to be an identifier of type `float`, while the right-hand-side is a literal of type integer.

An expression can be a sequence of expressions separated by the lexeme “;”, see Listing 6.8.

Listing 6.8 `letValueSequence.fsx`:
A value-binding for a sequence of expressions.

```

1  let p = 2.0 in do printfn "%A" p; do printfn "%A" (3.0 ** p)

```

```

1  $ fsharp --nologo letValueSequence.fsx && mono letValueSequence.exe
2  2.0
3  9.0

```

The lightweight syntax automatically inserts the “;” lexeme at newlines, hence using the lightweight syntax the above is the same as shown in Listing 6.9.

Listing 6.9 `letValueSequenceLightWeight.fsx`:
A value-binding for a sequence using lightweight syntax.

```

1  let p = 2.0
2  do printfn "%A" p
3  do printfn "%A" (3.0 ** p)

```

```

1  $ fsharp --nologo letValueSequenceLightWeight.fsx
2  $ mono letValueSequenceLightWeight.exe
3  2.0
4  9.0

```

A key concept of programming is *scope*. In F#, the scope of a value-binding is lexically meaning that when F# determines the value bound to a name, it looks left and upward in the program text for the `let` statement defining it, see, e.g., Listing 6.10.

Listing 6.10 `letValueScopeLower.fsx`:
Redefining identifiers is allowed in lower scopes.

```
1 let p = 3 in let p = 4 in do printfn "%A" p;

-----

1 $ fsharp --nologo letValueScopeLower.fsx && mono letValueScopeLower.exe
2 4
```

Some special bindings are mutable, in which case F# uses the dynamic scope, that is, the value of a binding is defined by when it is used, and this will be discussed in Section 6.7.

Scopes are given levels, and scopes may be nested, where the nested scope has a level one lower than its parent.¹ F# distinguishes between the top and lower levels, and at the top level in the lightweight syntax, redefining values is not allowed, as shown in Listing 6.11.

Listing 6.11 `letValueScopeLowerError.fsx`:
Redefining identifiers is not allowed in lightweight syntax at top level.

```
1 let p = 3
2 let p = 4
3 do printfn "%A" p;

-----

1 $ fsharp --nologo -a letValueScopeLowerError.fsx
2
3 letValueScopeLowerError.fsx(2,5): error FS0037: Duplicate definition of
  value 'p'
```

But using parentheses, we create a *block*, i.e., a *nested scope*, and then redefining is allowed, as demonstrated in Listing 6.12.

· block
· nested scope

Listing 6.12 `letValueScopeBlockAlternative3.fsx`:
A block may be created using parentheses.

```
1 (
2   let p = 3
3   let p = 4
4   do printfn "%A" p
5 )

-----

1 $ fsharp --nologo letValueScopeBlockAlternative3.fsx
2 $ mono letValueScopeBlockAlternative3.exe
3 4
```

In both cases, we used indentation, which is good practice, but not required here. Bindings inside are not available outside the nested scope as shown in Listing 6.13.

¹Jon: Drawings would be good to describe scope

Listing 6.13 `letValueScopeNestedScope.fsx`:
Bindings inside a scope are not available outside.

```

1  let p = 3
2  (
3      let q = 4
4      do printfn "%A" q
5  )
6  do printfn "%A %A" p q

```

```

1  $ fsharpc --nologo -a letValueScopeNestedScope.fsx
2
3  letValueScopeNestedScope.fsx(6,22): error FS0039: The value or
4  constructor 'q' is not defined. Maybe you want one of the following:
   p

```

Nesting is a natural part of structuring code, e.g., through function definitions to be discussed in Section 6.2 and flow control structures to be discussed in Chapter 8. Blocking code by nesting is a key concept for making robust code that is easy to use by others without the user necessarily needing to know the details of the inner workings of a block of code.

Defining blocks is used for controlling the extent of a lexical scope of bindings. For example, adding a second `printfn` statement as in Listing 6.14. will print the value 4 last bound to the identifier `p`, since F# interprets the above as `let p = 3 in let p = 4 in (printfn "%A" p; printfn "%A" p)`.

Listing 6.14 `letValueScopeBlockProblem.fsx`:
Overshadowing hides the first binding.

```

1  let p = 3 in let p = 4 in do printfn "%A" p; do printfn "%A" p

```

```

1  $ fsharpc --nologo letValueScopeBlockProblem.fsx
2  $ mono letValueScopeBlockProblem.exe
3  4
4  4

```

Had we intended to print the two different values of `p`, then we should have created a block as in Listing 6.15.

Listing 6.15 `letValueScopeBlock.fsx`:
Blocks allow for the return to the previous scope.

```

1  let p = 3 in (let p = 4 in do printfn "%A" p); do printfn "%A" p;

```

```

1  $ fsharpc --nologo letValueScopeBlock.fsx && mono letValueScopeBlock.exe
2  4
3  3

```

Here, the lexical scope of `let p = 4 in ...` is for the nested scope, which ends at “)”, returning to the lexical scope of `let p = 3 in ...`.

6.2 Function bindings

A function is a mapping between an input and output domain. A key advantage of using functions, when programming, is that they encapsulate code into smaller units, that are easier to debug and may be reused. F# is a functional first programming language and offers a number of alternative methods for specifying parameters, which will be discussed in this section. Binding identifiers to functions follows a syntax similar to value-binding,

Listing 6.16 Function binding expression

```
1 let <funcIdent> <arg> {<arg>} | () = <bodyExpr> [in <expr>]
```

Here `<funcIdent>` is an identifier and is the name of the function, `<arg>` is zero or more identifiers, that bind to the value used when calling the function, and which is to be used in the body of the function, the expression `<bodyExpr>`. The `|` notation denotes a choice, i.e., either that on the left-hand-side or that of the right-hand-side. Thus `let f x = x * x` and `let f () = 3` are valid function bindings, but `let f = 3` would be a value binding not a function binding. The arguments and the function may be annotated with a type, in which case for arguments we write

Listing 6.17 Function binding expression

```
1 let <funcIdent> (<arg> : <type>) {(<arg> : <type>)} : <type> | () : <type>
  = <bodyExpr> [in <expr>]
```

where `<type>` is a name of an existing type. The argument types are given in parentheses, and the return type is given last.

Functions are a key concept in F# and in this chapter, we will discuss the very basics. Recursive functions will be discussed in Chapter 8, and higher-order functions in Chapter 16.

An example of defining a function and using in interactive mode is shown in Listing 6.18.

Listing 6.18: An example of a binding of an identifier and a function.

```
1 > let sum (x : float) (y : float) : float = x + y in
2 - let c = sum 357.6 863.4 in
3 - do printfn "%A" c;;
4 1221.0
5 val sum : x:float -> y:float -> float
6 val c : float = 1221.0
7 val it : unit = ()
```

and we see that the function is interpreted to have the type `val sum : x:float -> y:float -> float`. The “->” lexeme means a mapping between sets, in this case, floats. The function is also a higher order function, to be discussed in detail below, and here it suffices to think of `sum` as a function that takes 2 floats as argument and returns a float.

Not all types need to be declared, just sufficient for F# to be able to infer the types for the full statement. In the example, one specification is sufficient, and we could just have specified the type of the result as in Listing 6.19.

Listing 6.19 letFunctionAlterative.fsx:
All types need most often not be specified.

```
1 let sum x y : float = x + y
```

Or even just one of the arguments as in Listing 6.20.

Listing 6.20 letFunctionAlterative2.fsx:
Just one type is often enough for F# to infer the rest.

```
1 let sum (x : float) y = x + y
```

In both cases, since the `+` operator is only defined for operands of the same type, then when the type of either the result, any or both operands are declared, then the type of the remaining follows directly. As for values, lightweight syntax automatically inserts the keyword `in` and the lexeme `;` as shown in Listing 6.21.

Listing 6.21 letFunctionLightWeight.fsx:
Lightweight syntax for function definitions.

```
1 let sum x y : float = x + y
2 let c = sum 357.6 863.4
3 do printfn "%A" c

-----

1 $ fsharpc --nologo letFunctionLightWeight.fsx
2 $ mono letFunctionLightWeight.exe
3 1221.0
```

Arguments need not always be inferred to types, but may be of the generic type, which F# prefers, when *type safety* is ensured as shown in Listing 6.22.

Listing 6.22: Type safety implies that a function will work for any type, and hence it is generic.

```
1 > let second x y = y
2 - let a = second 3 5
3 - do printfn "%A" a
4 - let b = second "horse" 5.0
5 - do printfn "%A" b;;
6 5
7 5.0
8 val second : x:'a -> y:'b -> 'b
9 val a : int = 5
10 val b : float = 5.0
11 val it : unit = ()
```

Here, the function `second` does not use the first argument `x`, which therefore can be of any type, and which F#, therefore, calls `'a`, and the type of the second element, `y`, can also be of any type and not necessarily the same as `x`, so it is called `'b`. Finally, the result is the same type as `y`, whatever it is. This is an example of a *generic function* since it will work on any type.

· operator
· operand

· type safety

· generic function

A function may contain a sequence of expressions but must return a value. E.g., the quadratic formula may be written as shown in Listing 6.23.

Listing 6.23 `identifiersExampleAdvance.fsx`:

A function may contain sequences of expressions.

```

1  let solution a b c sgn =
2      let discriminant a b c =
3          b ** 2.0 - 2.0 * a * c
4      let d = discriminant a b c
5      (-b + sgn * sqrt d) / (2.0 * a)
6
7  let a = 1.0
8  let b = 0.0
9  let c = -1.0
10 let xp = solution a b c +1.0
11 let xn = solution a b c -1.0
12 do printfn "0 = %A * x ** 2.0 + %A * x + %A" a b c
13 do printfn "  has solutions %A and %A" xn xp

```

```

1  $ fsharp --nologo identifiersExampleAdvance.fsx
2  $ mono identifiersExampleAdvance.exe
3  0 = 1.0 * x ** 2.0 + 0.0 * x + -1.0
4  has solutions -0.7071067812 and 0.7071067812

```

Here, we used the lightweight syntax, where the “=” identifies the start of a nested scope, and F# identifies the scope by indentation. The amount of space used for indentation does not matter, but all lines following the first must use the same. The scope ends before the first line with the previous indentation or none. Notice how the last expression is not bound to an identifier but is the result of the function, i.e., in contrast to many other languages, F# does not have an explicit keyword for returning values but requires a final expression, which will be returned to the caller of the function. Note also that since the function `discriminant` is defined in the nested scope of `solution` then `discriminant` cannot be called outside `solution` since the scope ends before `let a = 1.0`.

Lexical scope and function definitions can be a cause of confusion as the following example in Listing 6.24 shows.²

Listing 6.24 `lexicalScopeNFunction.fsx`:

Lexical scope means that $f(z) = 3x$ and not $4x$ at the time of calling.

```

1  let testScope x =
2      let a = 3.0
3      let f z = a * z
4      let a = 4.0
5      f x
6  do printfn "%A" (testScope 2.0)

```

```

1  $ fsharp --nologo lexicalScopeNFunction.fsx
2  $ mono lexicalScopeNFunction.exe
3  6.0

```

²Jon: Add a drawing or possibly a spell-out of lexical scope here.

Here, the value-binding for `a` is redefined, after it has been used to define a helper function `f`. So which value of `a` is used, when we later apply `f` to an argument? To resolve the confusion, remember that value-binding is lexically defined, i.e., the binding `let f z = a * z` uses the value of `a`, as it is defined by the ordering of the lines in the script, not dynamically by when `f` was called. Hence, **think of lexical scope as substitution of an identifier with its value or function immediately at the place of definition**. I.e., since `a` and `3.0` are synonymous in the first lines of the program, then the function `f` is really defined as, `let f z = 3.0 * z`. Advice

Functions do not need a name, but may be declared as an *anonymous function* using the `fun` keyword and the “->” lexeme, as shown in Listing 6.25. · anonymous function
· `fun`
· “->”

Listing 6.25 `functionDeclarationAnonymous.fsx`:
Anonymous functions are functions as values.

```
1 let first = fun x y -> x
2 do printfn "%d" (first 5 3)

-----

1 $ fsharp --nologo functionDeclarationAnonymous.fsx
2 $ mono functionDeclarationAnonymous.exe
3 5
```

Here, a name is bound to an anonymous function, which returns the first of two arguments. The difference to `let first x y = x` is that anonymous functions may be treated as values, meaning that they may be used as arguments to other functions, and new values may be reassigned to their identifiers, when mutable, as will be discussed in Section 6.7. A common use of anonymous functions is as arguments to other functions, as demonstrated in Listing 6.26.

Listing 6.26 `functionDeclarationAnonymousAdvanced.fsx`:
Anonymous functions are often used as arguments for other functions.

```
1 let apply f x y = f x y
2 let mul = fun a b -> a * b
3 do printfn "%d" (apply mul 3 6)

-----

1 $ fsharp --nologo functionDeclarationAnonymousAdvanced.fsx
2 $ mono functionDeclarationAnonymousAdvanced.exe
3 18
```

Note that here `apply` is given 3 arguments, the function `mul` and 2 integers. It is not given the result of `mul 3 6` since that would not match the definition of `apply`. **Anonymous functions and functions as arguments are powerful concepts but tend to make programs harder to read, and their use should be limited.** Advice

The result of one function is often used as an argument of another. This is function composition, and an example is shown in Listing 6.27.

Listing 6.27 functionComposition.fsx:
Composing functions using intermediate bindings.

```

1  let f x = x + 1
2  let g x = x * x
3
4  let a = f 2
5  let b = g a
6  let c = g (f 2)
7  do printfn "a = %A, b = %A, c = %A" a b c

```

```

1  $ fsharp --nologo functionComposition.fsx
2  $ mono functionComposition.exe
3  a = 3, b = 9, c = 9

```

In the example we combine two functions `f` and `g` by storing the result of `f 2` in `a` and using that as argument of `g`. This is the same as `g (f 2)`, and in the later case, the compiler creates a temporary value for `f 2`. Such compositions are so common in F# that a special operator has been invented called the *piping* operators, “`>|`” and “`<|`”. They are used as demonstrated in Listing 6.28.

· piping
· `lstineline|>`
· `<|`

Listing 6.28 functionPiping.fsx:
Composing functions by piping.

```

1  let f x = x + 1
2  let g x = x * x
3
4  let a = g (f 2)
5  let b = 2 |> f |> g
6  let c = g <| (f <| 2)
7  do printfn "a = %A, b = %A, c = %A" a b c

```

```

1  $ fsharp --nologo functionPiping.fsx && mono functionPiping.exe
2  a = 9, b = 9, c = 9

```

The example shows regular composition and left-to-right and right-to-left piping. The word piping is a picture of data flowing through pipes, where functions are places, where the pipes have been assembled and the data changes. The three expressions in Listing 6.28 performs the same calculation. The left-to-right piping in line 5 corresponds to the left-to-right reading direction of many human languages, i.e., the value 2 is used as argument to `f`, and the result is used as argument to `g`. However, this is opposite arithmetic composition in line 4. Right-to-left piping has the order of arithmetic composition. Unfortunately, since F# evaluates the expression from the left, without the parenthesis in line 6, `g <| f <| 2`, then F# would read the expression as `(g <| f) <| 2`, which is an error, since `g` takes an integer as argument not a function. F# can also define composite on a function level, further discussion on this is deferred to Chapter 16. The piping operator comes in four variants: “`|>`”, “`<|`”, “`||>`”, and “`<||`”. The allow for piping between pairs and triples to functions of 2 and 3 arguments, see Listing 6.29 for an example.

Listing 6.29 functionTuplePiping.fsx:

Tuples can be piped to functions of more than one argument.

```

1 let f x = printfn "%A" x
2 let g x y = printfn "%A %A" x y
3 let h x y z = printfn "%A %A %A" x y z
4
5 1 |> f
6 (1, 2) ||> g
7 (1, 2, 3) |||> h

```

```

1 $ fsharp -nologo functionTuplePiping.fsx
2 $ mono functionTuplePiping.exe
3 1
4 1 2
5 1 2 3

```

The example demonstrates left-to-right piping, right-to-left works analogously.³

A *procedure* is a generalization of the concept of functions, and in contrast to functions, procedures need not return values. This is demonstrated in Listing 6.30.

Listing 6.30 procedure.fsx:

A procedure is a function that has no return value, and in F# returns “()”.

```

1 let printIt a = printfn "This is '%A'" a
2 do printIt 3
3 do printIt 3.0

```

```

1 $ fsharp -nologo procedure.fsx && mono procedure.exe
2 This is '3'
3 This is '3.0'

```

In F# this is automatically given the unit type as the return value. Procedural thinking is useful for *encapsulation* of scripts, but is prone to *side-effects* and should be minimized by being replaced by functional thinking. More on side-effects in Section 6.7. **Procedural thinking is useful for encapsulation, but is prone to side-effects and should be minimized by being replaced by functional thinking.**

In F# functions (and procedures) are *first-class citizens*, which means that functions are values: They may be passed as arguments, returned from a function, and bound to a name. For first-class citizens, the name, it is bound to, does not carry significance to the language, as, e.g., illustrated with the use of anonymous functions. Technically, a function is stored as a *closures*. A closure is a description of the function, its arguments, its expression, and the environment, at the time it was created, i.e., the triple: (*args, exp, env*). Consider the listing in Listing 6.31.

³Jon: Tuples have not yet been introduced!

Listing 6.31 functionFirstClass.fsx:
The function timesTwo has a non-trivial closure.

```

1  let mul x y = x * y
2  let factor = 2.0
3  let applyFactor fct x =
4      let a = fct factor x
5      string a
6
7  do printfn "%g" (mul 5.0 3.0)
8  do printfn "%s" (applyFactor mul 3.0)

```

```

1  $ fsharp --nologo functionFirstClass.fsx && mono functionFirstClass.exe
2  15
3  6

```

It defines two functions `mul` and `applyFactor`, where the later is a higher order function taking another function as an argument and uses part of the environment to produce its result. The two closures are,

$$\text{mul} : (\text{args}, \text{exp}, \text{env}) = ((x, y), (x * y), ()) \quad (6.3)$$

$$\text{applyFactor} : (\text{args}, \text{exp}, \text{env}) = ((x, \text{fct}), (\text{let } a = \text{fct factor } x), (\text{factor} \rightarrow 2.0)) \quad (6.4)$$

The function `mul` does not use its environment, and everything needed to evaluate its expression are values for its arguments. The function `applyFactor` also takes two arguments, a function and a value. It uses `factor` from the environment, thus this is stored in its closure. When `mul` is given as an argument in Listing 6.31 line 8, then it is its closure, which is given to `applyFactor`, and the closure contains everything that `applyFactor` requires to use `mul`. Likewise, if `applyFactor` is given as argument to yet another function, then its closure includes the relevant part of its environment at time of definition, i.e., `factor`, such that when `applyFactor` is applied to two arguments, then its closure contains everything needed to evaluate its expression.

6.3 Operators

Operators are functions, and in F#, the infix multiplication operator `+` is equivalent to the function `(+)`, e.g.,

Listing 6.32 addOperatorNFunction.fsx:
Operators have function equivalents.

```

1  let a = 3.0
2  let b = 4.0
3  let c = a + b
4  let d = (+) a b
5  do printfn "%A plus %A is %A and %A" a b c d

```

```

1  $ fsharp --nologo addOperatorNFunction.fsx
2  $ mono addOperatorNFunction.exe
3  3.0 plus 4.0 is 7.0 and 7.0

```

All operators have this option, and you may redefine them and define your own operators, but in F#

names of user-defined operators are limited.

operator

- A *unary operator* name can be: “+”, “-”, “+.”, “-.”, “%”, “&”, “&&”, “~”, “~~”, “~~~”, ..., apostropheOp. Here apostropheOp is an operator name starting with “!” followed by one or more of either “!”, “%”, “&”, “*”, “+”, “-”, “.”, “/”, “<”, “=”, “>”, “@”, “^”, “|”, “~”, but apostropheOp cannot be “!=”. · unary operator
- An *binary operator* name can be: “+”, “-”, “+.”, “-.”, “%”, “&”, “&&”, “:=”, “:.”, “\$”, “?”, dotOp. Here dotOp is an operator name starting with “.” followed by “+”, “-”, “+.”, “-.”, “%”, “&”, “&&”, “-”, “+”, “|”, “<”, “>”, “=”, “|”, “&”, “^”, “*”, “/”, “%”, “!=”. Only “?” and “?<-” may start with “?”. · binary operator

The precedence rules and associativity of user-defined operators follow the rules for which they share prefixes with built-in rules, see Table E.5. E.g., `.*`, `+++`, and `<+` are valid operator names for infix operators, they have precedence as ordered, and their associativities are all left. Using `~` as the first character in the definition of an operator makes the operator unary and will not be part of the name. Examples of definitions and use of operators are,

Listing 6.33 operatorDefinitions.fsx:

Operators may be (re)defined by their function equivalent.

```
1 let (.*) x y = x * y + 1
2 printfn "%A" (3 .* 4)
3 let (+++) x y = x * y + y
4 printfn "%A" (3 +++ 4)
5 let (<+) x y = x < y + 2.0
6 printfn "%A" (3.0 <+ 4.0)
7 let (~+.) x = x+1
8 printfn "%A" (+.1)

1 $ fsharpc --nologo operatorDefinitions.fsx
2 $ mono operatorDefinitions.exe
3 13
4 16
5 true
6 2
```

Operators beginning with `*` must use a space in its definition, (`*` in order for it not to be confused with the beginning of a comment (`*`, see Chapter 7 for more on comments in the code).

Beware, redefining existing operators lexically redefines all future uses of the operators for all types, hence **it is not a good idea to redefine operators, but better to define new**. In Chapter 20 Advice we will discuss how to define type-specific operators including prefix operators.

6.4 Do bindings

Aside from `let` bindings that binds names with values or functions, sometimes we just need to execute code. This is called a `do` binding or alternatively a *statement*. The syntax is as follows.

- `do`
- do binding
- statement

Listing 6.34 Syntax for `do` bindings.

```
1  [do ]<expr>
```

The expression `<expr>` must return `unit`. The keyword `do` is optional in most cases, but using it emphasizes that the expression is not a function that returns a useful value. Procedures are examples of such expressions and a very useful family of procedures are the `printf` family described below. In the remainder of this book, we will refrain from using the `do` keyword.

6.5 The Printf function

A common way to output information to the console is to use one of the family of *printf* commands. These functions are special, since they take a variable number of arguments, and the number is decided by the first - the format string,

Listing 6.35 `printf` statement.

```
1  printf <format-string> {<ident>}
```

where a `formatString` is a string (simple or verbatim) with placeholders. The function `printf` prints `formatString` to the console, where all placeholder has been replaced by the value of the corresponding argument formatted as specified, e.g., in `printfn "1 2 %d" 3` the `formatString` is `"1 2 %d"`, and the placeholder is `%d`, and the `printf` replaced the placeholder with the value of the corresponding argument, and the result is printed to the console, in this case, `1 2 3`. There are specifiers for all the basic types and more as elaborated in Table 6.2. The placeholder can be given a specified with, either by setting a specific integer, or using the `*` character, indicating that the with is given as an argument prior to the replacement value. E.g., the placeholder string `%8s` will print a right-aligned string that is eight characters wide padded with spaces, if needed. For floating point numbers, `%8f` will print a number that is exactly seven digits and a decimal making eight characters in total. Zeros are added after the decimal as needed. Alternatively, we may specify the number of decimals, such that `%8.1f` will print a floating point number, aligned to the right with one digit after the decimal padded with spaces where needed. The default is for the value to be right justified in the field, but left justification can be specified by the `-` character. For number types, you can specify their format by `"0"` for padding the number with zeros to the left, when right justifying the number; `"+"` to explicitly show a plus sign for positive numbers; `SP` to enforce a space, where there otherwise would be a plus sign for positive numbers. Examples of some of these combinations are shown in Listing 6.36.

Specifier	Type	Description
%b	bool	Replaces with boolean value
%s	string	
%c	char	
%d, %i	basic integer	
%u	basic unsigned integers	
%x	basic integer	formatted as unsigned hexadecimal with lower case letters
%X	basic integer	formatted as unsigned hexadecimal with upper case letters
%o	basic integer	formatted as unsigned octal integer
%f, %F,	basic floats	formatted on decimal form
%e, %E,	basic floats	formatted on scientific form. Lower case uses "e" while upper case uses "E" in the formatting.
%g, %G,	basic floats	formatted on the shortest of the corresponding decimal or scientific form.
%M	decimal	
%O	Objects ToString method	
%A	any built-in types	Formatted as a literal type
%a	Printf.TextWriterFormat -> 'a -> ()	
%t	(Printf.TextWriterFormat -> ()	

Table 6.2: Printf placeholder string

Listing 6.36 printfExample.fsx:

Examples of printf and some of its formatting options.

```

1 let pi = 3.1415192
2 let hello = "hello"
3 printf "An integer: %d\n" (int pi)
4 printf "A float %f on decimal form and on %e scientific form\n" pi pi
5 printf "A char '%c' and a string \"%s\"\n" hello.[0] hello
6 printf "Float using width 8 and 1 number after the decimal:\n"
7 printf "  \"%8.1f\" \"%-8.1f\"\n" pi -pi
8 printf "  \"%08.1f\" \"%-08.1f\"\n" pi -pi
9 printf "  \"% 8.1f\" \"% 8.1f\"\n" pi -pi
10 printf "  \"%-8.1f\" \"%-8.1f\"\n" pi -pi
11 printf "  \"%+8.1f\" \"%+8.1f\"\n" pi -pi
12 printf "  \"%8s\" \"%-8s\"\n" "hello" "hello"

```

```

1 $ fsharp --nologo printfExample.fsx && mono printfExample.exe
2 An integer: 3
3 A float 3.141519 on decimal form and on 3.141519e+000 scientific form
4 A char 'h' and a string "hello"
5 Float using width 8 and 1 number after the decimal:
6     "    3.1" "    -3.1"
7     "000003.1" "-00003.1"
8     "    3.1" "    -3.1"
9     "3.1" "    -3.1"
10     "    +3.1" "    -3.1"
11     "    hello"
12 "hello"

```

Function	Example	Description
<code>printf</code> <code>printfn</code>	<code>printf "%d apples" 3</code>	Prints to the console, i.e., <code>stdout</code> as <code>printf</code> and adds a newline.
<code>fprintf</code> <code>fprintfn</code>	<code>fprintf stream "%d apples" 3</code>	Prints to a stream, e.g., <code>stderr</code> and <code>stdout</code> , which would be the same as <code>printf</code> and <code>fprintf</code> . as <code>fprintf</code> but with added newline.
<code>fprintf</code> <code>fprintfn</code>	<code>fprintf "%d apples" 3</code>	Print to <code>stderr</code> as <code>fprintf</code> but with added newline.
<code>sprintf</code>	<code>printf "%d apples" 3</code>	Return printed string
<code>failwithf</code>	<code>failwithf "%d failed apples" 3</code>	prints to a string and used for raising an exception.

Table 6.3: The family of `printf` functions.

Not all combinations of flags and identifier types are supported, e.g., strings cannot have number of integers after the decimal specified. The placeholder types `%A`, `%a`, and `%t` are special for `F#`, examples of their use are shown in Listing 6.37.

Listing 6.37 `printfExampleAdvance.fsx`:

Custom format functions may be used to specialise output.

```

1 let noArgument writer = printf "I will not print anything"
2 let customFormatter writer arg = printf "Custom formatter got: \"%A\""
  arg
3 printf "Print examples: %A, %A, %A\n" 3.0m 3uy "a string"
4 printf "Print function with no arguments: %t\n" noArgument
5 printf "Print function with 1 argument: %a\n" customFormatter 3.0

1 $ fsharpc --nologo printfExampleAdvance.fsx
2 $ mono printfExampleAdvance.exe
3 Print examples: 3.0M, 3uy, "a string"
4 Print function with no arguments: I will not print anything
5 Print function with 1 argument: Custom formatter got: "3.0"

```

The `%A` is special in that all built-in types including tuples, lists, and arrays to be discussed in Chapter 11 can be printed using this formatting string, but notice that the formatting performed includes the named literal string. The two formatting strings `%t` and `%a` are options for user-customizing the formatting, and will not be discussed further.

Beware, `formatString` is not a `string` but a `Printf.TextWriterFormat`, so to predefine a `formatString` as, e.g., `let str = "hello %s" in printf str "world"` will be a type error.

The family of `printf` is shown in Table 6.3. The function `fprintf` prints to a stream, e.g., `stderr` and `stdout`, of type `System.IO.TextWriter`. Streams will be discussed in further detail in Chapter 19. The function `failwithf` is used with exceptions, see Chapter 18 for more details. The function has a number of possible return value types, and for testing the `ignore` function ignores it all, e.g., `ignore (failwithf "%d failed apples" 3)`

6.6 Reading from the console

The `printf` and `printfn` functions allow us to write text on the screen, i.e., `stdout`. A program often needs to ask a user to input data, e.g., by typing text on a keyboard. Text typed on the keyboard is accessible through the `stdin` stream, and F# provides several library functions for capturing text typed on the keyboard. In the following, we will briefly discuss the `System.Console.ReadLine` function. For more details and other methods of input see Chapters 19 and 23.

The function `System.Console.ReadLine` takes a unit value as an argument and returns the string the user typed. The program will not advance until the user presses newline. An example of a program that multiplies two floating point numbers supplied by a user is given in Listing 6.38,

Listing 6.38 `userDialoguePrintf.fsx`:
Interacting with a user with `ReadLine`.

```
1 printfn "To perform the multiplication of a and b"
2 printf "Enter a: "
3 let a = float (System.Console.ReadLine ())
4 printf "Enter b: "
5 let b = float (System.Console.ReadLine ())
6 printfn "a * b = %A" (a * b)
```

and an example dialogue is shown in Listing 6.39.

Listing 6.39: An example dialogue of running Listing 6.38. What the user typed has been framed in red boxes.

```
1 $ fsharp --nologo userDialoguePrintf.fsx && mono userDialoguePrintf.exe
2 To perform the multiplication of a and b
3 Enter a: 3.5
4 Enter b: 7.4
5 a * b = 25.9
```

Note that the string is immediately cast to floats, such that we can multiply the input using the float multiplication operator. This also implies, that if the user inputs anything but floating point value, then `mono` will halt with an error message.

6.7 Variables

The *mutable* in `let` bindings means that the identifier may be rebound to a new value using the “<-” lexeme with the following syntax,⁴

Listing 6.40 Syntax for defining variables and their initial value.

```
1 let mutable <ident> = <expr> [in <expr>]
```

and values are changed using the assignment operator,

⁴Jon: Discussion on heap and stack should be added here.

Listing 6.41 Value reassignment for mutable variables.

```
1 <ident> <- <ident>
```

Mutable data is synonymous with the term *variable*. A variable is an area in the computers working memory associated with an identifier and a type, and this area may be read from and written to during program execution, see Listing 6.42 for an example.

Listing 6.42 mutableAssignReassingShort.fsx:
A variable is defined and later reassigned a new value.

```
1 let mutable x = 5
2 printfn "%d" x
3 x <- -3
4 printfn "%d" x

-----

1 $ fsharpc --nologo mutableAssignReassingShort.fsx
2 $ mono mutableAssignReassingShort.exe
3 5
4 -3
```

Here, an area in memory was denoted `x`, initially assigned the integer value 5, hence the type was inferred to be `int`. Later, this value of `x` was replaced with another integer using the “<-” lexeme. The “<-” lexeme is used to distinguish the assignment from the comparison operator, i.e., if we by mistake had written as shown in Listing 6.43, then we instead would have obtained the default assignment of the result of the comparison of the content of `a` with the integer 3, which is false.

Listing 6.43: It is a common error to mistake “=” and “<-” lexemes for mutable variables.

```
1 > let mutable a = 0
2 - a = 3;;
3 val mutable a : int = 0
4 val it : bool = false
```

However, it is important to note, that when the variable is initially defined, then the “=” operator must be used, while later reassignments must use the “<-” expression.

Assignment type mismatches will result in an error as demonstrated in Listing 6.44.

Listing 6.44 mutableAssignReassingTypeError.fsx:
Assignment type mismatching causes a compile-time error.

```

1  let mutable x = 5
2  printfn "%d" x
3  x <- -3.0
4  printfn "%d" x

```

```

1  $ fsharpc --nologo mutableAssignReassingTypeError.fsx
2
3  mutableAssignReassingTypeError.fsx(3,6): error FS0001: This expression
    was expected to have type
4      'int'
5  but here has type
6      'float'
7  $ mono mutableAssignReassingTypeError.exe
8  Cannot open assembly 'mutableAssignReassingTypeError.exe': No such file
    or directory.

```

I.e., once the type of an identifier has been declared or inferred, then it cannot be changed.

A typical variable is a counter of type integer, and a typical use of counters is to increment them, see Listing 6.45 for an example.

Listing 6.45 mutableAssignIncrement.fsx:
Variable increment is a common use of variables.

```

1  let mutable x = 5 // Declare a variable x and assign the value 5 to it
2  printfn "%d" x
3  x <- x + 1 // Assign a new value -3 to x
4  printfn "%d" x

```

```

1  $ fsharpc --nologo mutableAssignIncrement.fsx
2  $ mono mutableAssignIncrement.exe
3  5
4  6

```

Using variables in expressions, as opposed to the left-hand-side of an assignment operation, reads the value of the variable. Thus, when using a variable as the return value of a function, then the value is copied from the local scope of the function to the scope from which it is called. This is demonstrated in Listing 6.46.

Listing 6.46: Returning a mutable variable returns its value.

```

1  > let g () =
2  -   let mutable y = 0
3  -   y
4  -   printfn "%d" (g ());;
5  0
6  val g : unit -> int
7  val it : unit = ()

```

In the example, we see that the type is a value, and not mutable.

Variables implement dynamic scope, e.g., in comparison with the lexical scope, where the value of an identifier depends on *where* it is defined, dynamic scope depends on, *when* it is used. E.g., the script in Listing 6.24 defines a function using lexical scope and returns the number 6.0, however, if `a` is made `mutable`, then the behavior is different as shown in Listing 6.47.

Listing 6.47 `dynamicScopeNFunction.fsx`:

Mutual variables implement dynamics scope rules. Compare with Listing 6.24.

```
1 let testScope x =
2     let mutable a = 3.0
3     let f z = a * z
4     a <- 4.0
5     f x
6     printfn "%A" (testScope 2.0)

1 $ fsharp --nologo dynamicScopeNFunction.fsx
2 $ mono dynamicScopeNFunction.exe
3 8.0
```

Here, the response is 8.0 since the value of `a` changed before the function `f` was called.

6.8 Reference cells

F# has a variation of mutable variables called *reference cells*. Reference cells have built-in function `ref` and operators `!` and `:=`, where `ref` creates a reference variable, and the `!` and the `:=` operators reads and writes its value. An example of using reference cells is given in Listing 6.48.

Listing 6.48 `refCell.fsx`:

Reference cells are variants of mutable variables.

```
1 let x = ref 0
2 printfn "%d" !x
3 x := !x + 1
4 printfn "%d" !x

1 $ fsharp --nologo refCell.fsx && mono refCell.exe
2 0
3 1
```

Reference cells are different from mutable variables since their content is allocated on *The Heap*. The Heap is a global data storage that is not destroyed, when a function returns, which is in contrast to the *call stack* also known as *The Stack*. The Stack maintains all the local data for a specific instance of a function call, see Section 13.2 for more detail on the call stack. As a consequence, when a reference cell is returned from a function, then it is the reference to the location on The Heap, which is returned as a value. Since this points outside the local data area of the function, then this location is still valid after the function returns, and the variable stored there is accessible to the caller. This is illustrated in Figure 6.1

· reference cells
· `ref`
· `:=`

· The Heap
· call stack
· The Stack




Figure 6.1: A reference cell is a pointer to The Heap and the content is not destroyed when its reference falls out of scope.

Reference cells may cause *side-effects*, where variable changes are performed across independent scopes. Some side-effects are useful, e.g., the `printf` family changes the content of the screen, and the screen is outside the scope of the caller. Another example of a useful side-effect is a counter shown in Listing 6.49.

Listing 6.49 `refEncapsulation.fsx`:

An increment function with a local state using a reference cell.

```

1  let incr =
2      let counter = ref 0
3      fun () ->
4          counter := !counter + 1
5          !counter
6  printfn "%d" (incr ())
7  printfn "%d" (incr ())
8  printfn "%d" (incr ())

```

```

1  $ fsharpc --nologo refEncapsulation.fsx && mono refEncapsulation.exe
2  1
3  2
4  3

```

Here `incr` is an anonymous function with an internal state `counter`. At first glance, it may be surprising, that `incr ()` does not return the value 1 at every call. The reason is that the value of the `incr` is the closure of the anonymous function `fun () -> counter := ...`, which is

$$\text{incr} : (\text{args}, \text{exp}, \text{env}) = ((\text{counter} := !\text{counter} + 1), (\text{counter} \rightarrow \text{ref } 0)). \quad (6.5)$$

Thus, `counter` is only initiated once at the initial binding, while every call of `incr ()` updates its value on The Heap. Such a programming structure is called *encapsulation*, since the `counter` state has been encapsulated in the anonymous function, and the only way to access it is by calling the same anonymous function. **Encapsulation is good programming practice, but side-effects should be avoided wherever possible.**

The `incr` example in Listing 6.49 is an example of a useful side-effect. An example to be avoided is shown in Listing 6.50.

Listing 6.50 refSideEffect.fsx:

Intertwining independent scopes is typically a bad idea.

```

1  let updateFactor factor =
2      factor := 2
3
4  let multiplyWithFactor x =
5      let a = ref 1
6      updateFactor a
7      !a * x
8
9  printfn "%d" (multiplyWithFactor 3)

```

```

1  $ fsharpc --nologo refSideEffect.fsx && mono refSideEffect.exe
2  6

```

In the example, the function `updateFactor` changes a variable in the scope of `multiplyWithFactor`, which is prone to errors, since the computations are not local at the place of writing, i.e., in `multiplyWithFactor`, and if `updateFactor` were defined in a library, then the source code may not be available. Better style of programming is shown in Listing 6.51.

Listing 6.51 refWithoutSideEffect.fsx:

A solution of Listing 6.50 avoiding side-effects.

```

1  let updateFactor () =
2      2
3
4  let multiplyWithFactor x =
5      let a = ref 1
6      a := updateFactor ()
7      !a * x
8
9  printfn "%d" (multiplyWithFactor 3)

```

```

1  $ fsharpc --nologo refWithoutSideEffect.fsx
2  $ mono refWithoutSideEffect.exe
3  6

```

Here, there can be no doubt in `multiplyWithFactor` that the value of `a` is changing. Side-effects do have their use, but should, in general, be avoided at almost all costs, and in general, it is advised to **minimize the use of side effects**.

Advice

Reference cells give rise to an effect called *aliasing*, where two or more identifiers refer to the same data as illustrated in Listing 6.52. · aliasing

Listing 6.52 `refCellAliasing.fsx`:

Aliasing can cause surprising results and should be avoided.

```

1  let a = ref 1
2  let b = a
3  printfn "%d, %d" !a !b
4  b := 2
5  printfn "%d, %d" !a !b

```

```

1  $ fsharp --nologo refCellAliasing.fsx && mono refCellAliasing.exe
2  1, 1
3  2, 2

```

Here, `a` is defined as a reference cell, and by defining `b` to be equal to `a`, we have created an alias. This can be very confusing since as the example shows, changing the value of `b` causes `a` to change as well. Aliasing is a variant of side-effects, and **aliasing should be avoided at all costs**.

Advice

Since F# version 4.0, the compiler has automatically converted mutable variables to reference cells, when needed. E.g., Listing 6.49 can be rewritten using a mutable variable as shown in Listing 6.53.

Listing 6.53 `mutableEncapsulation.fsx`:

Local mutable content can be indirectly accessed outside its scope.

```

1  let incr =
2      let mutable counter = 0
3      fun () ->
4          counter <- counter + 1
5          counter
6  printfn "%d" (incr ())
7  printfn "%d" (incr ())
8  printfn "%d" (incr ())

```

```

1  $ fsharp --nologo mutableEncapsulation.fsx
2  $ mono mutableEncapsulation.exe
3  1
4  2
5  3

```

Reference cells are preferred over mutable variables for encapsulation to avoid confusion.

6.9 Tuples

Tuples are a direct extension of constants. They are immutable and do not have concatenations nor `·` tuple indexing operations. Tuples are unions of immutable types and have the following syntax,

Listing 6.54

```

1  <expr>{, <expr>}

```

Tuples are identified by the “,” lexeme and often enclosed in parentheses, but that is not required. An example is a triple, also known as a 3-tuple, `(2, true, "hello")`. In interactive mode, the type of tuples is shown as demonstrated in Listing 6.55.

Listing 6.55: Tuple types are products of sets.

```
1 > let tp = (2, true, "hello")
2 - printfn "%A" tp;;
3 (2, true, "hello")
4 val tp : int * bool * string = (2, true, "hello")
5 val it : unit = ()
```

The values `2`, `true`, and `"hello"` are *members*, and the number of elements of a tuple is its *length*. From the response of F#, we see that the tuple is inferred to have the type `int * bool * string`. The “*” denotes the Cartesian product between sets. Tuples can be products of any types and have the lexical scope like value and function bindings. Notice also that a tuple may be printed as a single entity by the `%A` placeholder. In the example, we bound `tp` to the tuple, the opposite is also possible, where an identifier is bound to a tuple as shown in Listing 6.56.

Listing 6.56: Definition of a tuple.

```
1 > let deconstructNPrint tp =
2 -   let (a, b, c) = tp
3 -   printfn "tp = (%A, %A, %A)" a b c
4 -
5 - deconstructNPrint (2, true, "hello")
6 - deconstructNPrint (3.14, "Pi", 'p');;
7 tp = (2, true, "hello")
8 tp = (3.14, "Pi", 'p')
9 val deconstructNPrint : 'a * 'b * 'c -> unit
10 val it : unit = ()
```

In this, a function is defined that takes 1 argument, a 3-tuple. If we wanted a function of 3 arguments, then the function binding should have been `let deconstructNPrint a b c = ...`. The value binding `let (a, b, c) = tp`, binds a tuple with 3 named members to a value, thus deconstructing it in terms of its members. This is called pattern matching and will be discussed in further details in Chapter 15. Since we used the `\%A` placeholder in the `printfn` function, then the function can be called with 3-tuples of different types. F# informs us that the tuple type is variable by writing `'a * 'b * 'c`. The “'” notation means that the type can be decided at run-time, see Section 14.6 for more on variable types.

Pairs or 2-tuples are so common that F# includes two built-in functions `fst` and `snd`, to extract the first and second element of a pair. This is demonstrated in Listing 6.57.

Listing 6.57 pair.fsx:
Deconstruction of pairs with the built-in functions `fst` and `snd`.

```
1 let pair = ("first", "second")
2 printfn "fst(pair) = %s, snd(pair) = %s" (fst pair) (snd pair)

-----

1 $ fsharp --nologo pair.fsx && mono pair.exe
2 fst(pair) = first, snd(pair) = second
```

Tuples comparison are defined similarly as strings. Tuples of different lengths are different. For tuples of equal length, then they are compared element by element. E.g., $(1,2) = (1,3)$ is false, while $(1,2) = (1,2)$ is true. The “<” operator is the boolean negation of the “=” operator. For the “<”, “<=”, “>”, and “>=” operators, the strings are ordered lexicographically, such that $(\text{'a'}, \text{'b'}, \text{'c'}) < (\text{'a'}, \text{'b'}, \text{'s'})$ && $(\text{'a'}, \text{'b'}, \text{'s'}) < (\text{'c'}, \text{'o'}, \text{'s'})$ is true, that is, the “<” operator on two tuples is true, if the left operand should come before the right, when sorting alphabetically, see Listing 6.58 for an example.

Listing 6.58 tupleCompare.fsx:

Tuples are compared as strings are compared alphabetically.

```

1 let lessThan (a, b, c) (d, e, f) =
2     if a <> d then a < d
3     elif b <> e then b < d
4     elif c <> f then c < f
5     else false
6
7 let printTest x y =
8     printfn "%A < %A is %b" x y (lessThan x y)
9
10 let a = ('a', 'b', 'c');
11 let b = ('d', 'e', 'f');
12 let c = ('a', 'b', 'b');
13 let d = ('a', 'b', 'd');
14 printTest a b
15 printTest a c
16 printTest a d

```

```

1 $ fsharpc --nologo tupleCompare.fsx && mono tupleCompare.exe
2 ('a', 'b', 'c') < ('d', 'e', 'f') is true
3 ('a', 'b', 'c') < ('a', 'b', 'b') is false
4 ('a', 'b', 'c') < ('a', 'b', 'd') is true

```

The algorithm for deciding the boolean value of $(a_1, a_2) < (b_1, b_2)$ is as follows: we start by examining the first elements, and if a_1 and b_1 are different, in which case the result of $(a_1, a_2) < (b_1, b_2)$ is equal to the result of $a_1 < b_1$. If a_1 and b_1 are equal, then we move onto the next letter and repeat the investigation. The “<=”, “>”, and “>=” operators are defined similarly.

Binding tuples to mutables does not make the tuple mutable. This is demonstrated in Listing 6.59.

Listing 6.59 tupleOfMutables.fsx:

A mutable change value, but the tuple defined by it does not refer to the new value.

```

1 let mutable a = 1
2 let mutable b = 2
3 let c = (a, b)
4 printfn "%A, %A, %A" a b c
5 a <- 3
6 printfn "%A, %A, %A" a b c

```

```

1 $ fsharpc --nologo tupleOfMutables.fsx && mono tupleOfMutables.exe
2 1, 2, (1, 2)
3 3, 2, (1, 2)

```


However, it is possible to define a mutable variable of type tuple such that new tuple values can be assigned to it as shown in Listing 6.60.

Listing 6.60 mutableTuple.fsx:

A mutable tuple can be assigned a new value.

```
1 let mutable pair = 1,2
2 printfn "%A" pair
3 pair <- (3,4)
4 printfn "%A" pair
```

```
1 $ fsharp --nologo mutableTuple.fsx && mono mutableTuple.exe
2 (1, 2)
3 (3, 4)
```

Mutable tuples are value types meaning that binding to new names makes copies not aliases, as demonstrated in Listing 6.61.

Listing 6.61 mutableTupleValue.fsx:

A mutable tuple is a value type.

```
1 let mutable pair = 1,2
2 let mutable aCopy = pair
3 pair <- (3,4)
4 printfn "%A %A" pair aCopy
```

```
1 $ fsharp --nologo mutableTupleValue.fsx && mono mutableTupleValue.exe
2 (3, 4) (1, 2)
```

The use of tuples shortens code and highlights semantic content at a higher level, e.g., instead of focussing on the elements, tuples focus on their union. While this may look elegant and short there is the risk of *obfuscation*, i.e., writing compact code that is difficult to read, where an unprepared reader of the code may not easily understand the computation nor appreciate its elegance without an accompanying explanation. Hence, **always keep an eye out for compact and concise ways to write code, but never at the expense of readability.** Advice

7 | In-code documentation

Documentation is a very important part of writing programs, since it is most unlikely, that you will be writing really obvious code. Moreover, what seems obvious at the point of writing may be mystifying months later to the author and to others. Documentation serves several purposes:

1. Communicate what the code should be doing
2. Highlight big insights essential for the code
3. Highlight possible conflicts and/or areas, where the code could be changed later

The essential point is that coding is a journey in problem-solving, and proper documentation is an aid in understanding the solution and the journey that lead to it. Documentation is most often a mixture of in-code documentation and accompanying documents. Here, we will focus on in-code documentation, which arguably causes problems in multi-language environments and run the risk of bloating code.

F# has two different syntaxes for comments. Comments are either block comments,

Listing 7.1 Block comments.

```
1 (<any text>)
```

The comment text (<any text>) can be any text and it is still parsed by F# as keywords and basic types implying that (* a comment (* in a comment *) *) and (* "*)" *) are valid comments, while (* " *) is invalid.

Alternatively, comments may be given as line comments,

Listing 7.2 Line comments.

```
1 //<any text>newline
```

where the comment text ends after the first newline.

The F# compiler has an option for generating *Extensible Markup Language (XML)* files from scripts using the C# documentation comments tags¹. The XML documentation starts with a triple-slash ///, i.e., a lineComment and a slash, which serves as comments for the code construct, that follows immediately after. XML consists of tags which always appears in pairs, e.g., the tag “tag” would look like <tag> ... </tag>. The F# accept any tags, but recommends those listed in Table 7.1. If no tags

- Extensible Markup Language
- XML

¹For specification of C# documentations comments see ECMA-334 3rd Edition, Annex E, Section 2: <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-334.pdf>

Tag	Description
<c>	Set text in a code-font.
<code>	Set one or more lines in code-font.
<example>	Set as an example.
<exception>	Describe the exceptions a function can throw.
<list>	Create a list or table.
<para>	Set text as a paragraph.
<param>	Describe a parameter for a function or constructor.
<paramref>	Identify that a word is a parameter name.
<permission>	Document the accessibility of a member.
<remarks>	Further describe a function.
<returns>	Describe the return value of a function.
<see>	Set as link to other functions.
<seealso>	Generate a See Also entry.
<summary>	Main description of a function or value.
<typeparam>	Describe a type parameter for a generic type or method.
<typeparamref>	Identify that a word is a type parameter name.
<value>	Describe a value.

Table 7.1: Recommended XML tags for documentation comments, from ECMA-334 3rd Edition, Annex E, Section 2.

are used, then it is automatically assumed to be a `<summary>`. An example of a documented script is,

Listing 7.3 commentExample.fsx:
Code with XML comments.

```

1  /// The discriminant of a quadratic equation with parameters a, b, and c
2  let discriminant a b c = b ** 2.0 - 4.0 * a * c
3
4  /// <summary>Find x when 0 = ax^2+bx+c.</summary>
5  /// <remarks>Negative discriminant are not checked.</remarks>
6  /// <example>
7  ///     The following code:
8  ///     <code>
9  ///         let a = 1.0
10 ///         let b = 0.0
11 ///         let c = -1.0
12 ///         let xp = (solution a b c +1.0)
13 ///         printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
14 ///     </code>
15 ///     prints <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> to the console.
16 /// </example>
17 /// <param name="a">Quadratic coefficient.</param>
18 /// <param name="b">Linear coefficient.</param>
19 /// <param name="c">Constant coefficient.</param>
20 /// <param name="sgn">+1 or -1 determines the solution.</param>
21 /// <returns>The solution to x.</returns>
22 let solution a b c sgn =
23     let d = discriminant a b c
24     (-b + sgn * sqrt d) / (2.0 * a)
25
26 let a = 1.0
27 let b = 0.0
28 let c = -1.0
29 let xp = (solution a b c +1.0)
30 printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp

```

```

1  $ fsharp --nologo commentExample.fsx && mono commentExample.exe
2  0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 1.0

```

Mono's `fsharpc` command may be used to extract the comments into an XML file as demonstrated in Listing 7.4.

Listing 7.4, Converting in-code comments to XML.

```

1  $ fsharp --doc:commentExample.xml commentExample.fsx
2  F# Compiler for F# 4.0 (Open Source Edition)
3  Freely distributed under the Apache 2.0 Open Source License

```

This results in an XML file with the following content as shown in Listing 7.5.

Listing 7.5, An XML file generated by fsharpc.

```

1  <?xml version="1.0" encoding="utf-8"?>
2  <doc>
3  <assembly><name>commentExample</name></assembly>
4  <members>
5  <member name="M:CommentExample.solution(System.Double, System.Double,
    System.Double, System.Double)">
6  <summary>Find x when 0 = ax^2+bx+c.</summary>
7  <remarks>Negative discriminant are not checked.</remarks>
8  <example>
9  The following code:
10 <code>
11     let a = 1.0
12     let b = 0.0
13     let c = -1.0
14     let xp = (solution a b c +1.0)
15     printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
16 </code>
17     prints <c>0 = 1.0x^2 + 0.0x + -1.0 => x_+ = 0.7</c> to the console.
18 </example>
19 <param name="a">Quadratic coefficient.</param>
20 <param name="b">Linear coefficient.</param>
21 <param name="c">Constant coefficient.</param>
22 <param name="sgn">+1 or -1 determines the solution.</param>
23 <returns>The solution to x.</returns>
24 </member>
25 <member name="M:CommentExample.discriminant(System.Double, System.Double,
    System.Double)">
26 <summary>
27 The discriminant of a quadratic equation with parameters a, b, and c
28 </summary>
29 </member>
30 </members>
31 </doc>

```

The extracted XML is written in C# type by convention, since F# is part of the Mono and .Net framework that may be used by any of the languages using Assemblies. Besides the XML inserted in the script, the XML has added `<?xml ...>` header, `<doc>`, `<assembly>`, `<members>`, and `<member>` tags. The header and the `<doc>` tag are standards for XML. The extracted XML is geared towards documenting big libraries of codes and thus highlights the structured programming organisation, see Chapters 9 and 20, and `<assembly>`, `<members>`, and `<member>` are indications for where the functions belong in the hierarchy. As an example, the prefix `M:CommentExample.` means that it is a method in the namespace `commentExample`, which in this case is the name of the file. Further, the function type `val solution : a:float -> b:float -> c:float -> sgn:float -> float` is in the XML documentation `M:CommentExample.solution(System.Double, System.Double, System.Double, System.Double)`, which is the C# equivalent.

An accompanying program in the Mono suite is `mdoc`, whose primary use is to perform a syntax analysis of an assembly and generate a scaffold XML structure for an accompanying document. With the `-i` flag, it is further possible to include the in-code comments as initial descriptions in the XML. The XML may be updated gracefully by `mdoc` as the code develops, without destroying manually entered documentation in the accompanying documentation. Finally, the XML may be exported to HTML.

The primary use of the `mdoc` command is to analyze compiled code and generate an empty XML structure with placeholders to describe functions, values, and variables. This structure can then be updated and edited as the program develops. The edited XML files can then be exported to *Hyper*

Text Markup Language (HTML) files, which can be viewed in any browser. Using the console, all of this is accomplished by the procedure shown in Listing 7.6, and the result is shown in Figure 7.1.

· Hyper Text Markup Language
· HTML

Listing 7.6, Converting an XML file to HTML.

```
1 $ mdoc update -o commentExample -i commentExample.xml commentExample.exe
2 New Type: CommentExample
3 Member Added: public static double determinant (double a, double b,
4   double c);
5 Member Added: public static double solution (double a, double b, double
6   c, double sgn);
7 Member Added: public static double a { get; }
8 Member Added: public static double b { get; }
9 Member Added: public static double c { get; }
10 Member Added: public static double xp { get; }
11 Namespace Directory Created:
12 New Namespace File:
13 Members Added: 6, Members Deleted: 0
14 $ mdoc export-html -out commentExampleHTML commentExample
15 .CommentExample
```

A full description of how to use mdoc is found [here](http://www.mono-project.com/docs/tools+libraries/tools/monodoc/generating-documentation/)².

²<http://www.mono-project.com/docs/tools+libraries/tools/monodoc/generating-documentation/>

solution Method

Find x when $0 = ax^2 + bx + c$.

Syntax

```
[Microsoft.FSharp.Core.CompilationArgumentCounts(Mono.Cecil.CustomAttributeArgument[])]  
public static double solution (double a, double b, double c, double sgn)
```

Parameters

a
Quadratic coefficient.
b
Linear coefficient.
c
Constant coefficient.
sgn
+1 or -1 determines the solution.

Returns

The solution to x .

Remarks

Negative discriminant are not checked.

Example

The following code:

```
Example  
let a = 1.0  
let b = 0.0  
let c = -1.0  
let xp = (solution a b c +1.0)  
printfn "0 = %.1fx^2 + %.1fx + %.1f => x_+ = %.1f" a b c xp
```

prints $0 = 1.0x^2 + 0.0x + -1.0 \Rightarrow x_+ = 0.7$ to the console.

Requirements

Namespace:
Assembly: commentExample (in commentExample.dll)
Assembly Versions: 0.0.0.0

Figure 7.1: Part of the HTML documentation as produce by `mdoc` and viewed in a browser.

8 | Controlling program flow

Non-recursive functions encapsulate code and allow for control of execution flow. That is, if a piece of code needs to be executed many times, then we can encapsulate it in the body of a function, and call the function several times. In this chapter, we will look at more general control of flow via loops and conditional execution. Recursion is another mechanism for controlling flow, but this is deferred to Chapter 13.

8.1 While and for loops

Many programming constructs need to be repeated, and F# contains many structures for repetition. A *while*-loop has the syntax,

· while

Listing 8.1 While loop.

```
1 while <condition> do <expr> [done]
```

The *condition* `<condition>` is an expression that evaluates to true or false. A while-loop repeats the `<expr>` expression as long as the condition is true. Using lightweight syntax the block following the `do` keyword up to and including the `done` keyword may be replaced by a newline and indentation. As an example, the program in Listing 8.5 counts from 1 to 10.

· condition

· do

· done

Listing 8.2 countWhile.fsx:

Count to 10 with a counter variable.

```
1 let mutable i = 1 in while i <= 10 do printf "%d " i; i <- i + 1 done;
2 printf "\n"
```

```
1 $ fsharp --nologo countWhile.fsx && mono countWhile.exe
2 1 2 3 4 5 6 7 8 9 10
```

We will call `i` for the counter variable. The counting is done by performing the following computation: In line 1, the counter variable is first given an initial value of 1. Then execution enters the while-loop and examines the condition. Since `1 <= 10` then the condition is true, and execution enters the body of the loop. The body prints the value of the counter to the screen and increases counter by 1. Then execution returns to the top of the while-loop. Now the condition is `2 <= 10` which is also true, and so execution enters the body and so on until the counter has reach value 11, in which case the condition `11 <= 10` is false, and execution continues in line 2.

In lightweight syntax, this would be as shown in Listing 8.3.

Listing 8.3 `countWhileLightweight.fsx`:

Count to 10 with a counter variable using lightweight syntax, see Listing 8.2.

```
1 let mutable i = 1
2 while i <= 10 do
3     printf "%d " i
4     i <- i + 1
5 printf "\n"

1 $ fsharp --nologo countWhileLightweight.fsx
2 $ mono countWhileLightweight.exe
3 1 2 3 4 5 6 7 8 9 10
```

Notice that although the expression following the condition is preceded with a `do` keyword, and `<expr>` is a do binding, the keyword `do` is mandatory.

Counters are so common that a special syntax has been reserved for loops using counters. These are called *for*-loops. For-loops comes in several variants, and here we will focus on the one using an explicit counter. Its syntax is,

Listing 8.4 For loop.

```
1 for <ident> = <firstExpr> to <lastExpr> do <bodyExpr> [done]
```

A for-loop initially binds the counter identifier `<ident>` to be the value `<firstExpr>`. Then execution enters the body and `<bodyExpr>` is evaluated. Once done, then the counter is increased and execution evaluates once again `<bodyExpr>`. This is repeated until and including the counter has the value `<lastExpr>`. As for while-loops, using lightweight syntax the block following the `do` keyword up to and including the `done` keyword may be replaced by a newline and indentation.

The counting example from Listing 8.2 using a *for*-loop is shown in Listing 8.5

Listing 8.5 `count.fsx`:

Counting from 1 to 10 using a *for*-loop.

```
1 for i = 1 to 10 do printf "%d " i done
2 printfn ""

1 $ fsharp --nologo count.fsx && mono count.exe
2 1 2 3 4 5 6 7 8 9 10
```

As this interactive script demonstrates, the identifier `i` takes all the values between 1 and 10, but in spite of its changing state, it is not mutable. Note also that the return value of the `for` expression is “()” like the `printf` functions. The lightweight equivalent is shown in Listing 8.6.




Figure 8.1: The Fibonacci spiral is an approximation of the golden spiral. Each square has side lengths of successive Fibonacci numbers, and the curve in each square is the circular arc with a radius of the square it is drawn in.

Listing 8.6 countLightweight.fsx:

Counting from 1 to 10 using a `for`-loop, see Listing 8.5.

```

1  for i = 1 to 10 do
2      printf "%d " i
3      printfn ""

1  $ fsharpc --nologo countLightweight.fsx && mono countLightweight.exe
2  1 2 3 4 5 6 7 8 9 10

```

To further compare `for`- and `while`-loops, consider the following problem.

Problem 8.1

Write a program that calculates the n 'th Fibonacci number.

The Fibonacci's numbers is a sequence of numbers starting with 1, 1, and where the next number is calculated as the sum of the previous two. Hence the first ten numbers are: 1, 1, 2, 3, 5, 8, 13, 21, 34, 55. Fibonacci's numbers are related to Golden spirals shown in Figure 8.1. Often the sequence is extended with a preceding number 0, to be 0, 1, 1, 2, 3, \dots , which we will do here as well.

We could solve this problem with a `for`-loop as follows,

Listing 8.7 fibFor.fsx:

The n 'th Fibonacci number is the sum of the previous 2.

```
1  let fib n =  
2      let mutable pair = (0, 1)  
3      for i = 2 to n do  
4          pair <- (snd pair, (fst pair) + (snd pair))  
5      snd pair  
6  
7  printfn "fib(1) = %d" (fib 1)  
8  printfn "fib(2) = %d" (fib 2)  
9  printfn "fib(3) = %d" (fib 3)  
10 printfn "fib(10) = %d" (fib 10)  
-----  
1  $ fsharp --nologo fibFor.fsx && mono fibFor.exe  
2  fib(1) = 1  
3  fib(2) = 1  
4  fib(3) = 2  
5  fib(10) = 55
```

The basic idea of the solution is that if we are given the $(n - 1)$ 'th and $(n - 2)$ 'th numbers, then the n 'th number is trivial to compute. And assume that `fib(1)` and `fib(2)` are given, then it is trivial to calculate the `fib(3)`. For the `fib(4)` we only need `fib(3)` and `fib(2)`, hence we may disregard `fib(1)`. Thus, we realize that we can cyclicly update the previous, current, and next values by shifting values until we have reached the desired `fib(n)`. This is implement in Listing 8.7 as function `fib`, which takes an integer `n` as argument and returns the n 'th Fibonacci number. The function does this iteratively using a `for`-loop, where `i` is the counter value, `pair` is the pair of the $i - 1$ 'th and i 'th Fibonacci numbers. In the body of the loop, the i 'th and $i + 1$ 'th numbers are assigned to `pair`. The `for`-loop automatically updates `i` for next iteration. When $n < 2$ then the body of the for-loop is not evaluated, and 1 is returned. This is of course wrong for $n < 1$, but we will ignore this for now.

The same program but using a while-loop is shown in Listing 8.8.

Listing 8.8 fibWhile.fsx:

Search for the largest Fibonacci number less than a specified number.

```

1  let fib (n : int) : int =
2      let mutable pair = (0, 1)
3      let mutable i = 1
4      while i < n do
5          pair <- (snd pair, fst pair + snd pair)
6          i <- i + 1
7      snd pair
8
9  printfn "fib(1) = %d" (fib 1)
10 printfn "fib(2) = %d" (fib 2)
11 printfn "fib(3) = %d" (fib 3)
12 printfn "fib(10) = %d" (fib 10)

```

```

1  $ fsharp --nologo fibWhile.fsx && mono fibWhile.exe
2  fib(1) = 1
3  fib(2) = 1
4  fib(3) = 2
5  fib(10) = 55

```

As can be seen, the program is almost identical. In this case, the `for`-loop is to be preferred, since more lines of code typically mean more chances of making a mistake. However, while-loops are still simpler and possibly easier to argue correctness about. To understand what is being calculated in code such as the while-loop in Listing 8.8, we can describe the loop in terms of its *loop invariant*. An *invariant* is a statement that is always true at a particular point in a program, and a loop invariant is a statement which is true at the beginning and end of a loop. In line 4 in Listing 8.8, we may state the invariant: The variable `pair` is the pair of the $i - 1$ 'th and i 'th Fibonacci numbers. This is provable by induction:

Base case: Before entering the while loop, i is 1, `pair` is (0, 1). Thus, the invariant is true.

Induction step: Assuming that `pair` is the $i - 1$ 'th and i 'th Fibonacci numbers, then the body first assigns a new value to `pair` as the i 'th and $i + 1$ 'th Fibonacci numbers, then increases i with one such that at the end of the loop the `pair` again contains the the $i - 1$ 'th and i 'th Fibonacci numbers. Thus, the invariant is true.

Thus when know that the second value in `pair` holds the value of the i 'th Fibonacci number, and since we further may prove that when line 7 is only reached, then $i = n$, and thus that `fib` returns the n 'th Fibonacci number.

While-loops also allows for other logical structures than for-loops, such as the case when the number of iteration cannot easily be decided when entering the loop. As an example, consider a slight variation of the above problem, where we wish to find the largest Fibonacci number less than some number. A solution to this problem is shown in Listing 8.9.

Listing 8.9 fibWhileLargest.fsx:

Search for the largest Fibonacci number less than a specified number.

```

1  let largestFibLeq n =
2      let mutable pair = (0, 1)
3      while snd pair <= n do
4          pair <- (snd pair, fst pair + snd pair)
5          snd pair
6
7  for i = 1 to 10 do
8      printfn "largestFibLeq(%d) = %d" i (largestFibLeq i)

```

```

1  $ fsharp --nologo fibWhileLargest.fsx && mono fibWhileLargest.exe
2  largestFibLeq(1) = 2
3  largestFibLeq(2) = 3
4  largestFibLeq(3) = 5
5  largestFibLeq(4) = 5
6  largestFibLeq(5) = 8
7  largestFibLeq(6) = 8
8  largestFibLeq(7) = 8
9  largestFibLeq(8) = 13
10 largestFibLeq(9) = 13
11 largestFibLeq(10) = 13

```

The strategy here is to iteratively calculate numbers in Fibonacci's sequence until we've found one larger than the argument `n`, and then return the previous. This could not be calculated with a for-loop.

8.2 Conditional expressions

Programs often contains code, which only should be executed under certain conditions. This can be expressed as `if`-expressions, whos syntax is as follows.

Listing 8.10 Conditional expressions.

```

1  if <cond> then <expr> {elif <cond> then <expr>} [else <expr>]

```

· `if`
 · `then`
 · `elif`
 · `else`

The condition `<con>` is an expression resulting in a Boolean value, and there can be zero or more `elif` conditions as indicated by `{}`. Each expression `<expr>` is called a *branch*, and all branches must have the identical type, such that regardless which branch is chosen, then the type of the result of the conditional expression is the same. The result of the conditional expression is the first branch, for which its condition was true and if all conditions are false then the `else`-branch is evaluated. If no `else` expression is present, then `"()`" will be returned. See Listing 8.11 for a simple example.

· branch

Listing 8.11 condition.fsx:

Conditions evaluates their branches depending on the value of the condition.

```

1  if true then printfn "hi" else printfn "bye"
2  if false then printfn "hi" else printfn "bye"
-----
1  $ fsharpc --nologo condition.fsx && mono condition.exe
2  hi
3  bye

```

The lightweight syntax allows for newlines entered everywhere, but indentation must be used to express scope.

To demonstrate conditional expressions, let us write a program, which writes the sentence, “I have n apple(s)”, where the plural ‘s’ is added appropriately for various *ns*. This is done in Listing 8.12 using the lightweight syntax.

Listing 8.12 conditionalLightweight.fsx:

Using conditional expression to generate different strings.

```

1  let applesIHave n =
2      if n < -1 then
3          "I owe " + (string -n) + " apples"
4      elif n < 0 then
5          "I owe " + (string -n) + " apple"
6      elif n < 1 then
7          "I have no apples"
8      elif n < 2 then
9          "I have 1 apple"
10     else
11         "I have " + (string n) + " apples"
12
13  printfn "%A" (applesIHave -3)
14  printfn "%A" (applesIHave -1)
15  printfn "%A" (applesIHave 0)
16  printfn "%A" (applesIHave 1)
17  printfn "%A" (applesIHave 2)
18  printfn "%A" (applesIHave 10)
-----
1  $ fsharpc --nologo conditionalLightWeight.fsx
2  $ mono conditionalLightWeight.exe
3  "I owe 3 apples"
4  "I owe 1 apple"
5  "I have no apples"
6  "I have 1 apple"
7  "I have 2 apples"
8  "I have 10 apples"

```

The sentence structure and its variants give rise to a more compact solution since the language to be returned to the user is a variant of “I have/or no/number apple(s)”, i.e., under certain conditions should the sentence use “have” and “owe” etc. So, we could instead make decisions on each of these sentence parts and then built the final sentence from its parts. This is accomplished in the following example:

Listing 8.13 conditionalLightweightAlt.fsx:
Using sentence parts to construct the final sentence.

```

1  let applesIHave n =
2      let haveOrOwe = if n < 0 then "owe" else "have"
3      let pluralS = if (n = 0) || (abs n) > 1 then "s" else ""
4      let number = if n = 0 then "no" else (string (abs n))
5
6      "I " + haveOrOwe + " " + number + " apple" + pluralS
7
8  printfn "%A" (applesIHave -3)
9  printfn "%A" (applesIHave -1)
10 printfn "%A" (applesIHave 0)
11 printfn "%A" (applesIHave 1)
12 printfn "%A" (applesIHave 2)
13 printfn "%A" (applesIHave 10)

```

```

1  $ fsharp --nologo conditionalLightWeightAlt.fsx
2  $ mono conditionalLightWeightAlt.exe
3  "I owe 3 apples"
4  "I owe 1 apple"
5  "I have no apples"
6  "I have 1 apple"
7  "I have 2 apples"
8  "I have 10 apples"

```

While arguably shorter, this solution is also denser, and for a small problem like this, it is most likely more difficult to debug and maintain.

Note that both `elif` and `else` branches are optional, which may cause problems. For example, both `let a = if true then 3` and `let a = if true then 3 elif false then 4` will be invalid, since F# is not smart enough to realize that the type of the expression is uniquely determined. Instead, F# looks for the `else` to ensure all cases have been covered, and that `a` always will be given a unique value of the same type regardless of the branch taken in the conditional statement, hence, `let a = if true then 3 else 4` is the only valid expression of the 3. In practice, F# assumes that the omitted branch returns `()`, and thus it is fine to say `let a = if true then ()` and `if true then printfn "hej"`. Nevertheless, it is good practice in F# always to include an `else` branch.

8.3 Programming intermezzo: Automatic conversion of decimal to binary numbers

Using loops and conditional expressions, we are now able to solve the following problem:

Problem 8.2

Given an integer on decimal form, write its equivalent value on the binary form.

To solve this problem, consider odd numbers: They all have the property, that the least significant bit is 1, e.g., $1_2 = 1$, $101_2 = 5$ in contrast to even numbers such as $110_2 = 6$. Division by 2 is equal to right-shifting by 1, e.g., $1_2/2 = 0.1_2 = 0.5$, $101_2/2 = 10.1_2 = 2.5$, $110_2/2 = 11_2 = 3$. Thus, by integer division by 2 and checking the remainder, we may sequentially read off the least significant bit. This

leads to the algorithm shown in Listing 8.14.

Listing 8.14 dec2bin.fsx:

Using integer division and remainder to write any positive integer on binary form.

```

1  let dec2bin n =
2      if n < 0 then
3          "Illegal value"
4      elif n = 0 then
5          "0b0"
6      else
7          let mutable v = n
8          let mutable str = ""
9          while v > 0 do
10             str <- (string (v % 2)) + str
11             v <- v / 2
12             "0b" + str
13
14
15  printfn "%4d -> %s" -1 (dec2bin -1)
16  printfn "%4d -> %s" 0 (dec2bin 0)
17  for i = 0 to 3 do
18      printfn "%4d -> %s" (pown 10 i) (dec2bin (pown 10 i))

```

```

1  $ fsharp --nologo dec2bin.fsx && mono dec2bin.exe
2      -1 -> Illegal value
3          0 -> 0b0
4          1 -> 0b1
5          10 -> 0b1010
6          100 -> 0b1100100
7          1000 -> 0b1111101000

```

In the code, the states `v` and `str` are iteratively updated until `str` finally contains the desired solution.

To prove that Listing 8.14 calculates the correct sequence we use induction. First we realize, that for $v < 1$ then the `while` loop is skipped, and the result is trivially true. We will concentrate on line 9 in Listing 8.14, and we will prove the following loop invariant: The string `str` contains all the bits of `n` to the right of the bit pattern remaining in variable `v`.

Base case $n = 000 \dots 000x$: If n only uses the lowest bit then $n = 0$ or $n = 1$. If $n = 0$ then it is trivially correct. Considering the case $n = 1$: Before entering into the loop, `v` is 1, `str` is the empty string, so the invariant is true. The condition of the while-loop is $1 > 0$ so execution enters the loop. Since integer division of 1 by 2 gives 0 with remainder 1, then `str` is set to `"1"` and `v` to 0. Now we reexamine the while-loop's condition, $0 > 0$, which is false, so we exit the loop. At this point, `v` is 0 and `str` is `"1"`, so all bits have been shifted from `n` to `str` and none are left in `v`. Thus the invariant is true. Finally, the program returns `"0b1"`.

Induction step: Consider the case of $n > 1$, and assume that the invariant is true when entering the loop, i.e., that m bits already have been shifted to `str` and that $n > 2^m$. In this case `v` contains the remaining bits of `n`, which is the integer division $v = n / 2^{*m}$. Since $n > 2^m$ then `v` is non-zero and the loop condition is true, so we enter the loop body. In the loop body we concatenate the rightmost bit of `v` to the left of `str` using `v % 2`, and right-shift `v` one bit to the right with `v <- v / 2`. Thus, when returning to the condition, the invariant is true, since the right-most bit in `v` has been shifted to `str`. This continues until all bits have been shifted to `str` and `v = 0`, in which case the loop terminates, and `"0b"+str` is returned.

Thus we have proven, that `dec2bin` correctly converts integers to strings representing binary numbers.

9 | Organising code in libraries and application programs

In this chapter, we will focus on a number of ways to make the code available as *library* functions in F#. A library is a collection of types, values, and functions that an application program can use. A library does not perform calculations on its own.

F# includes several programming structures to organize code in libraries: Modules, namespaces, and classes. In this chapter, we will describe modules and namespaces. Classes will be described in detail in Chapter 20.

9.1 Modules

An F# *module*, not to be confused with a Common Language Infrastructure module (see Appendix D), is a programming structure used to organize type declarations, values, functions, etc.

Every implementation and script file in F# implicitly defines a module, and the module name is given by the filename. Thus, creating a script file `Meta.fsx` as shown in Listing 9.1.¹

Listing 9.1 Meta.fsx:
A script file defining the apply function.

```
1 type floatFunction = float -> float -> float
2 let apply (f : floatFunction) (x : float) (y : float) : float = f x y
```

Here we’ve implicitly defined a module of name `Meta`. Another script file may now use this function, which is accessed using the “.” notation, i.e., `Meta.apply` will refer to this function in other programs. An application program could be as shown in Listing 9.3.

Listing 9.2 MetaApp.fsx:
Defining a script calling the module.

```
1 let add : Meta.floatFunction = fun x y -> x + y
2 let result = Meta.apply add 3.0 4.0
3 printfn "3.0 + 4.0 = %A" result
```

In the above, we have explicitly used the module’s type definition for illustration purposes. A shorter

¹Jon: Type definitions have not been introduced at this point!

and possibly simpler program would have been to define `add` as `let add x y = x + y`, since F#'s type system will infer the implied type. However, **explicit definitions of types is recommended for readability**. Hence, an alternative to the above's use of lambda functions is, `let add (x: float) (y: float) : float = x + y`. To compile the module and the application program, we write as demonstrated in Listing 9.3. Advice

Listing 9.3: Compiling both the module and the application code. Note that file order matters, when compiling several files.

```
1 $ fsharp --nologo Meta.fsx MetaApp.fsx && mono MetaApp.exe
2 3.0 + 4.0 = 7.0
```

Notice, since the F# compiler reads through the files once, the order of the filenames in the compile command is very important. Hence, the script containing the module and function definitions must be to the left of the script containing their use. Notice also that if not otherwise specified, then the F# compiler produces an `.exe` file derived from the last filename in the list of filenames.

We may also explicitly define the module name using the `module` using the syntax,

· `module`

Listing 9.4 Outer module.

```
1 module <ident>
2 <script>
```

Here, the identifier `<ident>` is a name not necessarily related to the filename, and the script `<script>` is expression. An example is given in Listing 9.20.

**Listing 9.5 MetaExplicit.fsx:
Explicit definition of the outermost module.**

```
1 module Meta
2 type floatFunction = float -> float -> float
3 let apply (f : floatFunction) (x : float) (y : float) : float = f x y
```

Since we have created a new file, where the module `Meta` is explicitly defined, we can use the same application program. This is demonstrated in Listing 9.6.

Listing 9.6: Changing the module definition to explicit naming has no effect on the application nor the compile command.

```
1 $ fsharp --nologo MetaExplicit.fsx MetaApp.fsx && mono MetaApp.exe
2 3.0 + 4.0 = 7.0
```

Notice that, since `MetaExplicit.fsx` explicitly defines the module name, `apply` is not available to an application program as `MetaExplicit.apply`. **It is recommended that module names are defined explicitly, since filenames may change due to external conditions**. I.e., filenames are typically set from the perspective of the filesystem. The user may choose to change names to suit a filesystem structure, or different platforms may impose different file naming convention. Thus, direct linking of filenames with the internal workings of a program is a needless complication of structure. Advice

The definitions inside a module may be accessed directly from an application program, omitting the

“.”-notation, by use of the `open` keyword,

· `open`

Listing 9.7 Open module.

```
1 open <ident>
```

I.e., we can modify `MetaApp.fsx` as shown in Listing 9.9.

Listing 9.8 MetaAppWOpen.fsx:
Avoiding the “.”-notation by the `open` keyword.

```
1 open Meta
2 let add : floatFunction = fun x y -> x + y
3 let result = apply add 3.0 4.0
4 printfn "3.0 + 4.0 = %A" result
```

In this case, the namespace of our previously defined module is included into the scope of the application functions, and its types, values, functions, etc. can be used directly, as shown in Listing 9.9.

Listing 9.9: How the application program opens the module has no effect on the module code nor compile command.

```
1 $ fsharpc --nologo MetaExplicit.fsx MetaAppWOpen.fsx && mono
   MetaAppWOpen.exe
2 3.0 + 4.0 = 7.0
```

The `open`-keyword should be used sparingly, since including a library’s definitions into the application scope can cause surprising naming conflicts, since the user of a library typically has no knowledge of the inner workings of the library. E.g., the user may accidentally use code defined in the library, but with different type and functionality than intended, which the type system will use to deduce types in the application program, and therefore will either give syntax or runtime errors that are difficult to understand. This problem is known as *namespace pollution*, and for clarity, **it is recommended to use the `open`-keyword sparingly**. Notice that for historical reasons, the work namespace pollution is used to cover pollution both due to modules and namespaces.

· namespace pollution
Advice

Modules may also be nested, in which case the nested definitions must use the “=”-sign and must be appropriately indented.

Listing 9.10 Nested modules.

```
1 module <ident> = <script>
```

In lightweight syntax, a newline may be entered before the script `<script>`, and the script must be indented. An example is shown in Listing 9.11.

Listing 9.11 nestedModules.fsx:
Modules may be nested.

```

1 module Utilities
2 let PI = 3.1415
3 module Meta =
4     type floatFunction = float -> float -> float
5     let apply (f : floatFunction) (x : float) (y : float) : float = f x y
6 module MathFcts =
7     let add : Meta.floatFunction = fun x y -> x + y

```

In this case, `Meta` and `MathFcts` are defined at the same level and said to be siblings, while `Utilities` is defined at a higher level. In this relation, the former two are said to be the children of the latter. Note that the nesting respects the lexical scope rules, such that the constant `PI` is directly accessible in both modules `Meta` and `MathFcts`, as is the module `Meta` in `MathFcts` but not `MathFcts` in `Meta`. The “.”-notation is reused to index deeper into the module hierarchy as the example in Listing 9.12 shows.

Listing 9.12 nestedModulesApp.fsx:
Applications using nested modules require additional usage of the “.” notation to navigate the nesting tree.

```

1
2 let add : Utilities.Meta.floatFunction = fun x y -> x + y
3 let result = Utilities.Meta.apply Utilities.MathFcts.add 3.0 Utilities.PI
4 printfn "3.0 + 4.0 = %A" result

```

Modules can be recursive using the `rec`-keyword, meaning that in our example we can make the outer module recursive as demonstrated in Listing 9.13.²

Listing 9.13 nestedRecModules.fsx:
Mutual dependence on nested modules requires the `rec` keyword in the module definition.

```

1 module rec Utilities
2     module Meta =
3         type floatFunction = float -> float -> float
4         let apply (f : floatFunction) (x : float) (y : float) : float = f x y
5     module MathFcts =
6         let add : Meta.floatFunction = fun x y -> x + y

```

The consequence is that the modules `Meta` and `MathFcts` are accessible in both modules, but compilation will now give a warning since soundness of the code will first be checked at runtime. In general, it is advised to **avoid programming constructions, whose validity cannot be checked at compile-time.** Advice

9.2 Namespaces

An alternative to structure code in modules is to use a *namespace*, which only can hold modules and `namespace`

²Jon: **Dependence on version 4.1 and higher.**

type declarations and only works in compiled mode. Namespaces are defined as explicitly defined outer modules using the `namespace` keyword,

· `namespace`

Listing 9.14 Namespace.

```
1 namespace <ident>
2 <script>
```

An example is given in Listing 9.15.

Listing 9.15 namespace.fsx:

Defining a namespace is similar to explicitly named modules.

```
1 namespace Utilities
2 type floatFunction = float -> float -> float
3 module Meta =
4     let apply (f : floatFunction) (x : float) (y : float) : float = f x y
```

Notice that when organizing code in a namespace, the first line of the file, other than comments and compiler directives, must be the one starting with `namespace`.

As for modules, the content of a namespace is accessed using the “.” notation, as demonstrated in Listing 9.17.

Listing 9.16 namespaceApp.fsx:

The “.”-notation lets the application program accessing functions and types in a namespace.

```
1 let add : Utilities.floatFunction = fun x y -> x + y
2 let result = Utilities.Meta.apply add 3.0 4.0
3 printfn "3.0 + 4.0 = %A" result
```

Likewise, the compilation is performed identically, see Listing 9.17.

Listing 9.17: Compilation of files including namespace definitions uses the same procedure as modules.

```
1 $ fsharp --nologo namespace.fsx namespaceApp.fsx && mono
   namespaceApp.exe
2 3.0 + 4.0 = 7.0
```

Hence, from an application point of view, it is not immediately possible to see, that `Utilities` is defined as a namespace and not a module. However, in contrast to modules, namespaces may span several files. E.g., we may add a third file extending the `Utilities` namespace with the `MathFcts` module as demonstrated in Listing 9.18.

Listing 9.18 namespaceExtension.fsx:

Namespaces may span several files. Here is shown an extra file, which extends the Utilities namespace.

```
1 namespace Utilities
2 module MathFcts =
3     let add : floatFunction = fun x y -> x + y
```

To compile we now need to include all three files in the right order. Likewise, the compilation is performed identically, see Listing 9.19.

Listing 9.19: Compilation of namespaces defined in several files requires careful consideration of order, since the compiler reads once and only once through the files in the order they are given.

```
1 $ fsharp --nologo namespace.fsx namespaceExtension.fsx namespaceApp.fsx
   && mono namespaceApp.exe
2 3.0 + 4.0 = 7.0
```

The order matters since `namespaceExtension.fsx` relies on the definition of `floatFunction` in the file `namespace.fsx`. You can use extensions to extend existing namespaces included with the F# compiler.³⁴

Namespaces may also be nested. In contrast to modules, nesting is defined using the “.” notation. That is, to create a child namespace `more of Utilities` we must use initially write `namespace Utilities.more`. Indentation is ignored in the `namespace` line, thus left-most indentation is almost always used. Namespaces follow lexical scope rules, and identically to modules, namespaces containing mutually dependent children can be declared using the `rec` keyword, e.g., `namespace rec Utilities`.

9.3 Compiled libraries

Libraries may be distributed in compiled form as `.dll` files. This saves the user from having to recompile a possibly large library every time library functions needs to be compiled with an application program. In order to produce a library file from `MetaExplicitModuleDefinition.fsx` and then compile an application program, we first use the compiler’s `-a` option to produce the `.dll`. A demonstration is given in Listing 9.20.

Listing 9.20: A stand-alone `.dll` file is created and used with special compile commands.

```
1 $ fsharp --nologo -a MetaExplicit.fsx
```

This produces the file `MetaExplicit.dll`, which may be linked to an application using the `-r` option during compilation, see Listing 9.21.⁵

³Jon: **Something about intrinsic and optional extension** <https://docs.microsoft.com/en-us/dotnet/fsharp/language-reference/type-extensions>.

⁴Jon: **Perhaps something about the global namespace** `global`.

⁵Jon: **This is the MacOS option standard, Windows is slightly different.**

Listing 9.21: The library is linked to an application during compilation to produce runnable code.

```
1 $ fsharpc --nologo -r MetaExplicit.dll MetaApp.fsx && mono MetaApp.exe
2 3.0 + 4.0 = 7.0
```

A library can be the result of compiling a number of files into a single .dll file. .dll-files may be loaded dynamically in script files (.fsx-files) using the *#r directive* as illustrated in Listing 9.23.

· #r directive

Listing 9.22 MetaHashApp.fsx:

The .dll file may be loaded dynamically in .fsx script files and in interactive mode. Nevertheless, this usage is not recommended.

```
1 #r "MetaExplicit.dll"
2 let add : Meta.floatFunction = fun x y -> x + y
3 let result = Meta.apply add 3.0 4.0
4 printfn "3.0 + 4.0 = %A" result
```

We may now omit the explicit mentioning of the library when compiling as shown in Listing 9.23.

Listing 9.23: When using the #r directive, then the .dll file need not be explicitly included in the list of files to be compiled.

```
1 $ fsharpc --nologo MetaHashApp.fsx && mono MetaHashApp.exe
2 3.0 + 4.0 = 7.0
```

The *#r* directive is also used to include a library in interactive mode. However, for the code to be compiled, the use of the *#r* directive requires that the filesystem path to the library is coded inside the script. As for module names, direct linking of filenames with the internal workings of a program is a needless complication of structure, and **it is recommended not to rely on the use of the #r directive**.

Advice

In the above, we have compiled *script files* into libraries. However, F# has reserved the .fs filename suffix for library files and such files are called *implementation files*. In contrast to script files, implementation files do not support the *#r* directive. When compiling a list of implementation and script files, all but the last file must explicitly define a module or a namespace.

· script files
· implementation files

Both script and implementation files may be augmented with *signature files*. A signature file contains no implementation but only type definitions. Signature files offer three distinct features:

· signature files

1. Signature files can be used as part of the documentation of code since type information is of paramount importance for an application programmer to use a library.
2. Signature files may be written before the implementation file. This allows for a higher-level programming design that focusses on *which* functions should be included and *how* they can be composed.
3. Signature files allow for access control. Most importantly, if a type definition is not available in the signature file, then it is not available to the application program. Such definitions are private and can only be used internally in the library code. More fine-grained control is available relating to classes and will be discussed in Chapter 20.

Signature files can be generated automatically using the `--sig:<filename>` compiler directive. To

demonstrate this feature, we will first move the definition of `add` to the implementation file, see Listing 9.28.

Listing 9.24 MetaWAdd.fs:
An implementation file including the `add` function.

```
1 module Meta
2 type floatFunction = float -> float -> float
3 let apply (f : floatFunction) (x : float) (y : float) : float = f x y
4 let add (x : float) (y : float) : float = x + y
```

A signature file may be automatically generated as shown in Listing 9.25.

Listing 9.25: Automatic generation of a signature file at compile time.

```
1 $ fsharpc --nologo --sig:MetaWAdd.fsi MetaWAdd.fs
2
3 MetaWAdd.fs(4,48): warning FS0988: Main module of program is empty:
   nothing will happen when it is run
```

The warning can safely be ignored since at this point it is not our intention to produce runnable code. The above has generated the signature file in Listing 9.28.

Listing 9.26 MetaWAdd.fsi:
An automatically generated signature file from `MetaWAdd.fs`.

```
1 module Meta
2 type floatFunction = float -> float -> float
3 val apply : f:floatFunction -> x:float -> y:float -> float
4 val add : x:float -> y:float -> float
```

We can generate a library using the automatically generated signature file using `fsharpc -a MetaWAdd.fsi MetaWAdd.fs`, which is identical to compiling the `.dll` file without the signature file. However, if we remove, e.g., the type definition for `add` in the signature file, then this function becomes private to the module, and cannot be accessed outside. Hence, using the signature file in Listing 9.30 and recompiling the `.dll` as Listing 9.28 generates no error.

Listing 9.27 MetaWAddRemoved.fsi:
Removing the type definition for `add` from `MetaWAdd.fsi`.

```
1 module Meta
2 type floatFunction = float -> float -> float
3 val apply : f:floatFunction -> x:float -> y:float -> float
```

Listing 9.28: Automatic generation of a signature file at compile time.

```
1 $ fsharpc --nologo -a MetaWAddRemoved.fsi MetaWAdd.fs
```

But, when using the newly created `MetaWAdd.dll` with a modified version of Listing 9.3, which does

not itself supply a definition of `add` as shown in Listing 9.29, we get a syntax error, since `add` now is inaccessible to the application program. This is demonstrated in Listing 9.30.

Listing 9.29 `MetaWOAddApp.fsx`:

A version of Listing 9.3 without a definition of `add`.

```
1 let result = Meta.apply add 3.0 4.0
2 printfn "3.0 + 4.0 = %A" result
```

Listing 9.30: Automatic generation of a signature file at compile time.

```
1 $ fsharpc --nologo -r MetaWAdd.dll MetaWOAddApp.fsx
2
3 MetaWOAddApp.fsx(1,25): error FS0039: The value or constructor 'add' is
   not defined.
```

10 | Testing programs

A software bug is an error in a computer program that causes it to produce an incorrect result or behave in an unintended manner. The term bug was used by Thomas Edison in 1878¹, but made popular in computer science by Grace Hopper, who found a moth interfering with the electronic circuits of the Harvard Mark II electromechanical computer and coined the term *bug* for errors in computer programs. The original bug is shown in Figure 10.1. Software is everywhere, and errors therein have a huge economic impact on our society and can threaten lives².

The ISO/IEC organizations have developed standards for software testing³. To illustrate basic concepts of software quality, consider a hypothetical route planning system. Essential factors of its quality are,

- functionality
 - reliability
 - usability
 - efficiency
- Functionality:** Does the software compile and run without internal errors. Does it solve the problem, it was intended to solve? E.g., does the route planning software find a suitable route from point a to b?
- Reliability:** Does the software work reliably over time? E.g., does the route planning software work in case of internet dropouts?
- Usability:** Is the software easy and intuitive to use by humans? E.g., is it easy to enter addresses and alternative routes in the software's interface?

¹https://en.wikipedia.org/wiki/Software_bug, possibly <http://edison.rutgers.edu/NamesSearch/DocImage.php3?DocId=LB003487>

²https://en.wikipedia.org/wiki/List_of_software_bugs

³ISO/IEC 9126, International standard for the evaluation of software quality, December 19, 1991, later replaced by ISO/IEC 25010:2011

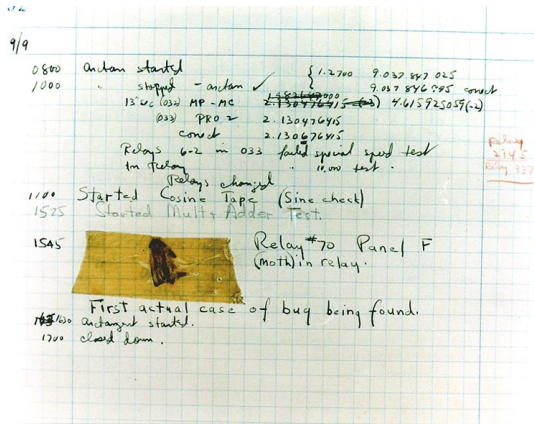


Figure 10.1: The first computer bug caught by Grace Hopper, U.S. Naval Historical Center Online Library Photograph NH 96566-KN.

Efficiency: How many computer and human resources does the software require? E.g., does it take milliseconds or hours to find a requested route? Can the software run on a mobile platform with limited computer speed and memory?

Maintainability: In case of the discovery of new bugs, is it easy to test and correct the software? Is it easy to extend the software with new functionality? E.g., is it easy to update the map with updated roadmaps and new information? Can the system be improved to work both for car drivers and bicyclists?

Portability: Is it easy to port the software to new systems such as new server architecture and screen sizes? E.g., if the routing software originally was written for IOS devices, will it be easy to port to Android systems?

· maintainability

· portability

The above-mentioned concepts are ordered based on the requirements of the system. Functionality and reliability are perhaps the most important concepts, since if the software does not solve the specified problem, then the software design process has failed. However, many times the problem definition will evolve along with the software development process. But as a bare minimum, the software should run without internal errors and not crash under a well-defined set of circumstances. Furthermore, it is often the case that software designed for the general public requires a lot of attention to the usability of the software since in many cases non-experts are expected to be able to use the software with little or no prior training. On the other hand, software used internally in companies will be used by a small number of people, who become experts in using the software, and it is often less important that the software is easy to understand by non-experts. An example is text processing software like Microsoft Word versus Gnu Emacs and LaTeX. Word is designed to be used by non-experts for small documents such as letters and notes and relies heavily on interfacing with the system using click-interaction. On the other hand, Emacs and LaTeX are for experts for longer and professionally typeset documents and relies heavily on keyboard shortcuts and text-codes for typesetting document entities.

The purpose of *software testing* is to find bugs. When errors are found, then we engage in *debugging*, which is the process of diagnosing and correcting bugs. Once we have a failed software test, i.e., one that does not find any bugs, then we have strengthened our belief in the software, but it is important to note that software testing and debugging rarely removes all bugs, and with each correction or change of software, there is a fair chance of introducing new bugs. It is not exceptional that the software testing the software is as large as the software itself.

· software testing

· debugging

In this chapter, we will focus on two approaches to software testing, which emphasizes functionality: *white-box* and *black-box testing*. An important concept in this context is *unit testing*, where the program is considered in smaller pieces, called units, and for which accompanying programs for testing can be made, which tests these units automatically. Black-box testing considers the problem formulation and the program interface, and can typically be written early in the software design phase. In contrast, white-box testing considers the program text, and thus requires the program to be available. Thus, there is a tendency for black-box test programs to be more stable, while white-box testing typically is developed incrementally alongside the software development.

· white-box testing

· black-box testing

· unit testing

To illustrate software testing we'll start with a problem:

Problem 10.1

Given any date in the Gregorian calendar, calculate the day of the week.

Facts about dates in the Gregorian calendar are:

- combinations of dates and weekdays repeat themselves every 400 years;
- the typical length of the months January, February, ... follow the knuckle rule, i.e., January belongs to the index knuckle, February to the space between the index and the middle finger,

and August restarts or starts on the other hand. All knuckle months have 31 days, all spacing months have 30 days except February, which has 29 days on leap years and 28 days all other years.

- A leap year is a multiple of 4, except if it is also a multiple of 100 but not of 400.

Many solutions to the problem have been discovered, and here we will base our program on Gauss' method, which is based on integer division and calculates the weekday of the 1st of January of a given year. For any other date, we will count our way through the weeks from the previous 1st of January. The algorithm relies on an enumeration of weekdays starting with Sunday = 0, Monday = 1, ..., and Saturday = 6. Our proposed solution is shown in Listing 10.1.⁴

Listing 10.1 date2Day.fsx:

A function that can calculate day-of-week from any date in the Gregorian calendar.

```

1  let januaryFirstDay (y : int) =
2      let a = (y - 1) % 4
3      let b = (y - 1) % 100
4      let c = (y - 1) % 400
5      (1 + 5 * a + 4 * b + 6 * c) % 7
6
7  let rec sum (lst : int list) j =
8      if 0 <= j && j < lst.Length then
9          lst.[0] + sum lst.[1..] (j - 1)
10     else
11         0
12
13  let date2Day d m y =
14      let dayPrefix =
15          ["Sun"; "Mon"; "Tues"; "Wednes"; "Thurs"; "Fri"; "Satur"]
16      let feb = if (y % 4 = 0) && ((y % 100 > 0) || (y % 400 = 0)) then 29
17              else 28
18      let daysInMonth = [31; feb; 31; 30; 31; 30; 31; 31; 30; 31; 30; 31]
19      let dayOne = januaryFirstDay y
20      let daysSince = (sum daysInMonth (m - 2)) + d - 1
21      let weekday = (dayOne + daysSince) % 7;
22      dayPrefix.[weekday] + "day"

```

10.1 White-box testing

White-box testing considers the text of a program. The degree to which the text of the program is covered in the test is called *coverage*. Since our program is small, we do have the opportunity to ensure that all functions are called at least once, which is called *function coverage*, we will also be able to test every branching in the program, which is called *branching coverage*, and in this case that implies *statement coverage*. The procedure is as follows:

- white-box testing
- coverage
- function coverage
- branching coverage
- statement coverage

1. Decide which are the units to test: The program, shown in Listing 10.1, has 3 functions, and we will consider these each as a unit, but we might as well just have chosen `date2Day` as a single unit. The important part is that the union of units must cover the whole program text, and since `date2Day` calls both `januaryFirstDay` and `sum`, designing test cases for the two later is superfluous. However, we may have to do this anyway, when debugging, and we may choose at

⁴Jon: This example relies on lists, which has not been introduced yet.

a later point to use these functions separately, and in both cases, we will be able to reuse the testing of the smaller units.

2. Identify branching points: The function `januaryFirstDay` has no branching function, `sum` has one, and depending on the input values two paths through the code may be used, and `date2Day` has one, where the number of days in February is decided. Note that in order to test this, our test-date must be March 1 or later. In this example, there are only examples of `if`-branch points, but they may as well be loops and pattern matching expressions. In the following code, the branch points have been given a comment and a number, as shown in Listing 10.2.

Listing 10.2 `date2DayAnnotated.fsx`:

In white-box testing, the branch points are identified.

```

1  // Unit: januaryFirstDay
2  let januaryFirstDay (y : int) =
3      let a = (y - 1) % 4
4      let b = (y - 1) % 100
5      let c = (y - 1) % 400
6      (1 + 5 * a + 4 * b + 6 * c) % 7
7
8  // Unit: sum
9  let rec sum (lst : int list) j =
10     (* WB: 1 *)
11     if 0 <= j && j < lst.Length then
12         lst.[0] + sum lst.[1..] (j - 1)
13     else
14         0
15
16 // Unit: date2Day
17 let date2Day d m y =
18     let dayPrefix =
19         ["Sun"; "Mon"; "Tues"; "Wednes"; "Thurs"; "Fri"; "Satur"]
20     (* WB: 1 *)
21     let feb = if (y % 4 = 0) && ((y % 100 <> 0) || (y % 400 = 0)) then
22         29 else 28
23     let daysInMonth = [31; feb; 31; 30; 31; 30; 31; 31; 30; 31; 30; 31]
24     let dayOne = januaryFirstDay y
25     let daysSince = (sum daysInMonth (m - 2)) + d - 1
26     let weekday = (dayOne + daysSince) % 7;
27     dayPrefix.[weekday] + "day"

```

3. For each unit, produce an input set that tests each branch: In our example, the branch points depend on a Boolean expression, and for good measure, we are going to test each term that can lead to branching. Thus,

Unit	Branch	Condition	Input	Expected output
januaryFirstDay	0	-	2016	5
sum	1	$0 \leq j \ \&\& \ j < \text{lst.Length}$		
	1a	<code>true && true</code>	[1; 2; 3] 1	3
	1b	<code>false && true</code>	[1; 2; 3] -1	0
	1c	<code>true && false</code>	[1; 2; 3] 10	0
	1d	<code>false && false</code>	-	-
date2Day	1	$(y \% 4 = 0) \ \&\& \ ((y \% 100 \neq 0) \ \ (y \% 400 = 0))$		
	-	<code>true && (true true)</code>	-	-
	1a	<code>true && (true false)</code>	8 9 2016	Thursday
	1b	<code>true && (false true)</code>	8 9 2000	Friday
	1c	<code>true && (false false)</code>	8 9 2100	Wednesday
	-	<code>false && (true true)</code>	-	-
	1d	<code>false && (true false)</code>	8 9 2015	Tuesday
	-	<code>false && (false true)</code>	-	-
	-	<code>false && (false false)</code>	-	-
	-			

The impossible cases have been intentionally blank, e.g., it is not possible for $j < 0$ and $j > n$ for some positive value n .

- Write a program that tests all these cases and checks the output, e.g.,

Listing 10.3 `date2DayWhiteTest.fsx`:

The tests identified by white-box analysis. The program from Listing 10.2 has been omitted for brevity.

```

1 printfn "White-box testing of date2Day.fsx"
2 printfn "  Unit: januaryFirstDay"
3 printfn "    Branch: 0 - %b" (januaryFirstDay 2016 = 5)
4
5 printfn "  Unit: sum"
6 printfn "    Branch: 1a - %b" (sum [1; 2; 3] 1 = 3)
7 printfn "    Branch: 1b - %b" (sum [1; 2; 3] -1 = 0)
8 printfn "    Branch: 1c - %b" (sum [1; 2; 3] 10 = 0)
9
10 printfn "  Unit: date2Day"
11 printfn "    Branch: 1a - %b" (date2Day 8 9 2016 = "Thursday")
12 printfn "    Branch: 1b - %b" (date2Day 8 9 2000 = "Friday")
13 printfn "    Branch: 1c - %b" (date2Day 8 9 2100 = "Wednesday")
14 printfn "    Branch: 1d - %b" (date2Day 8 9 2015 = "Tuesday")

```

```

1 $ fsharpc --nologo date2DayWhiteTest.fsx && mono
   date2DayWhiteTest.exe
2 White-box testing of date2Day.fsx
3   Unit: januaryFirstDay
4     Branch: 0 - true
5   Unit: sum
6     Branch: 1a - true
7     Branch: 1b - true
8     Branch: 1c - true
9   Unit: date2Day
10     Branch: 1a - true
11     Branch: 1b - true
12     Branch: 1c - true
13     Branch: 1d - true

```

Notice that the output of the tests is organized such that they are enumerated per unit, hence we can rearrange as we like and still uniquely refer to a unit's test. Also, the output of the test program produces a list of tests that should return true or success or a similar positively loaded word, but without further or only little detail, such that we at a glance can identify any test that produced unexpected results.

After the white-box testing has failed to find errors in the program, we have some confidence in the program, since we have run every line at least once. It is, however, in no way a guarantee that the program is error free, which is why white-box testing is often accompanied with black-box testing to be described next.

10.2 Black-box testing

In black-box testing, the program is considered a black box, and no knowledge is required about how a particular problem is solved, in fact, it is often useful not to have that knowledge at all. It is rarely possible to test all input to a program, so in black-box testing, the solution is tested for typical and extreme cases based on knowledge of the problem. The procedure is as follows:

Decide on the interface to use: It is useful to have an agreement with the software developers about what interface is to be used, e.g., in our case, the software developer has made a function `date2Day d m y` where `d`, `m`, and `y` are integers specifying the day, month, and year.

Make an overall description of the tests to be performed and their purpose:

- 1 a consecutive week, to ensure that all weekdays are properly returned
- 2 two set of consecutive days across boundaries that may cause problems: across a new year, across a regular month boundary.
- 3 a set of consecutive days across February-March boundaries for a leap and non-leap year
- 4 four dates after February in a non-multiple-of-100 leap year and in a non-leap year, a multiple-of-100-but-not-of-400 non-leap year, and a multiple-of-400 leap year.

Given no information about the program's text, there are other dates that one could consider as likely candidates for errors, but the above is judged to be a fair coverage.

Choose a specific set of input and expected output relations on the tabular form:

Test number	Input	Expected output
1a	1 1 2016	Friday
1b	2 1 2016	Saturday
1c	3 1 2016	Sunday
1d	4 1 2016	Monday
1e	5 1 2016	Tuesday
1f	6 1 2016	Wednesday
1g	7 1 2016	Thursday
2a	31 12 2014	Wednesday
2b	1 1 2015	Thursday
2c	30 9 2017	Saturday
2d	1 10 2017	Sunday
3a	28 2 2016	Sunday
3b	29 2 2016	Monday
3c	1 3 2016	Tuesday
3d	28 2 2017	Tuesday
3e	1 3 2017	Wednesday
4a	1 3 2015	Sunday
4b	1 3 2012	Thursday
4c	1 3 2000	Wednesday
4d	1 3 2100	Monday

Write a program executing the tests as shown in Listing 10.4 and 10.5.

Listing 10.4 date2DayBlackTest.fsx:

The tests identified by black-box analysis. The program from Listing 10.2 has been omitted for brevity.

```

28 let testCases = [
29     ("A complete week",
30      [(1, 1, 2016, "Friday");
31       (2, 1, 2016, "Saturday");
32       (3, 1, 2016, "Sunday");
33       (4, 1, 2016, "Monday");
34       (5, 1, 2016, "Tuesday");
35       (6, 1, 2016, "Wednesday");
36       (7, 1, 2016, "Thursday");]);
37     ("Across boundaries",
38      [(31, 12, 2014, "Wednesday");
39       (1, 1, 2015, "Thursday");
40       (30, 9, 2017, "Saturday");
41       (1, 10, 2017, "Sunday")]);
42     ("Across february boundary",
43      [(28, 2, 2016, "Sunday");
44       (29, 2, 2016, "Monday");
45       (1, 3, 2016, "Tuesday");
46       (28, 2, 2017, "Tuesday");
47       (1, 3, 2017, "Wednesday")]);
48     ("Leap years",
49      [(1, 3, 2015, "Sunday");
50       (1, 3, 2012, "Thursday");
51       (1, 3, 2000, "Wednesday");
52       (1, 3, 2100, "Monday")]);
53 ]
54
55 printfn "Black-box testing of date2Day.fsx"
56 for i = 0 to testCases.Length - 1 do
57     let (setName, testSet) = testCases.[i]
58     printfn "    %d. %s" (i+1) setName
59     for j = 0 to testSet.Length - 1 do
60         let (d, m, y, expected) = testSet.[j]
61         let day = date2Day d m y
62         printfn "        test %d - %b" (j+1) (day = expected)

```

Listing 10.5: Output from Listing 10.4.

```

1  $ fsharpc --nologo date2DayBlackTest.fsx && mono
    date2DayBlackTest.exe
2  Black-box testing of date2Day.fsx
3    1. A complete week
4      test 1 - true
5      test 2 - true
6      test 3 - true
7      test 4 - true
8      test 5 - true
9      test 6 - true
10     test 7 - true
11   2. Across boundaries
12     test 1 - true
13     test 2 - true
14     test 3 - true
15     test 4 - true
16   3. Across february boundary
17     test 1 - true
18     test 2 - true
19     test 3 - true
20     test 4 - true
21     test 5 - true
22   4. Leap years
23     test 1 - true
24     test 2 - true
25     test 3 - true
26     test 4 - true

```

Notice how the program has been made such that it is almost a direct copy of the table, produced in the previous step.

A black-box test is a statement of what a solution should fulfill for a given problem. Hence, **it is a good idea to make a black-box test early in the software design phase, in order to clarify the requirements for the code to be developed, and take an outside view of the code prior to developing it.** Advice

After the black-box testing has failed to find errors in the program, we have some confidence in the program, since, from a user's perspective, the program produces sensible output in many cases. It is, however, in no way a guarantee that the program is error free.

10.3 Debugging by tracing

Once an error has been found by testing, then the *debugging* phase starts. The cause of a bug can either be that the chosen algorithm is the wrong one for the job, or the implementation of it has an error. In the debugging process, we have to keep an open mind, and not rely on assumptions, since assumptions tend to blind the reader of a text. A frequent source of errors is that the state of a program is different, than expected, e.g., because the calculation performed is different than intended, or that the return of a library function is different than expected. The most important tool for debugging is *simplification*. This is similar to white-box testing, but where the units tested are very small. E.g., the suspected piece of code could be broken down into smaller functions or code snippets, which is given well-defined input, and, e.g., use `printfn` statements to obtain the output of the code snippet. Another related technique is to use *mockup code*, which replaces parts of the code with code that produces safe and · debugging
· mockup code

relevant results. If the bug is not obvious then more rigorous techniques must be used such as *tracing*.⁵ Some development interfaces have built-in tracing system, e.g., **fsharp**i will print inferred types and some binding values. However, often a source of a bug is due to a misunderstanding of the flow of data through a program execution, and we will in the following introduce *hand tracing* a technique to simulate the execution of a program by hand.

Consider the program in Listing 10.6.⁵

Listing 10.6 gcd.fsx:
gcd

```

1  let rec gcd a b =
2      if a < b then
3          gcd b a
4      elif b > 0 then
5          gcd b (a % b)
6      else
7          a
8
9  let a = 10
10 let b = 15
11 printfn "gcd %d %d = %d" a b (gcd a b)

```

```

1  $ fsharpc --nologo gcd.fsx && mono gcd.exe
2  gcd 10 15 = 5

```

The greatest common divisor of 2 integers. which includes a function for calculating the greatest common divisor of 2 integers, and calls this function with the numbers 10 and 15. Hand tracing this program means that we simulate its execution and as part of that keep track of the bindings, assignments and input and output of the program. To do this, we need to consider code snippet's *environment*. E.g., to hand trace the above program, we start by noting the outer environment, called E_0 for short. In line 1, then the `gcd` identifier is bound to a function, hence we write:

$$E_0 : \\ \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset)$$

Function bindings like this one is noted as a closure, which is the triplet (arguments, expression, environment). The closure is everything needed for the expression to be calculated. Here we wrote `gcd-body` to denote everything after the equal sign in the function binding. Next, F# executes line 9 and 10, and we update our environment to reflect the bindings as,

$$E_0 : \\ \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\ a \rightarrow 10 \\ b \rightarrow 15$$

⁵Jon: This program uses recursion, which has not been introduced yet.

In line 11 the function is evaluated. This initiates a new environment E_1 , and we update our trace as,

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow ? \\
 E_1 : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset)
 \end{aligned}$$

where the new environment is noted to have gotten its argument names **a** and **b** bound to the values 10 and 15 respectively, and where the return of the function to environment E_0 is yet unknown, so it is noted as a question mark. In line 2 the comparison **a** < **b** is checked, and since we are in environment E_1 then this is the same as checking $10 < 15$, which is true so the program executes line 3. Hence, we initiate a new environment E_2 and update our trace as,

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow ? \\
 E_1 : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset) \\
 & \text{line 3: gcd b a} \rightarrow ? \\
 E_2 : & ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset)
 \end{aligned}$$

where in the new environment **a** and **b** bound to the values 15 and 10 respectively. In E_2 , $10 < 15$ is false, so the program evaluates **b** > 0, which is true, hence line 5 is executed. This calls **gcd** once again, but with new arguments, and **a** % **b** is parenthesized, then it is evaluated before **gcd** is called.

Hence, we update our trace as,

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow ? \\
 E_1 : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset) \\
 & \text{line 3: gcd b a} \rightarrow ? \\
 E_2 : & ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 5 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow ? \\
 E_3 : & ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset)
 \end{aligned}$$

Again we fall through to line 5, evaluate the remainder operator and initiates a new environment,

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow ? \\
 E_1 : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset) \\
 & \text{line 3: gcd b a} \rightarrow ? \\
 E_2 : & ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 5 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow ? \\
 E_3 : & ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 0 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow ? \\
 E_4 : & ((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \emptyset)
 \end{aligned}$$

This time both $a < b$ and $b > 0$ are false, so we fall through to line 7, and `gcd` from E_4 returns its value of `a`, which is 5, so we scratch E_4 and change the question mark in E_3 to 5:

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow ? \\
 E_1 : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset) \\
 & \text{line 3: gcd b a} \rightarrow ? \\
 E_2 : & ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 5 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow ? \\
 E_3 : & ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 0 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow 5 \\
 \cancel{E_4} : & ((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \emptyset)
 \end{aligned}$$

Now line 5 in E_3 is also a return point of `gcd`, hence we scratch E_3 and change the question mark in

E_2 to 5,

$E_0 :$
 gcd $\rightarrow ((a, b), \text{gcd-body}, \emptyset)$
 $a \rightarrow 10$
 $b \rightarrow 15$
 line 11: gcd a b $\rightarrow ?$
 $E_1 : ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset)$
 line 3: gcd b a $\rightarrow ?$
 $E_2 : ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset)$
 line 5: a % b $\rightarrow 5$
 line 5: gcd b (a % b) $\rightarrow \text{? } 5$
 ~~$E_3 : ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset)$~~
 line 5: a % b $\rightarrow 0$
 line 5: gcd b (a % b) $\rightarrow \text{? } 5$
 ~~$E_4 : ((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \emptyset)$~~

and likewise, for E_2 and E_1 :

$E_0 :$
 gcd $\rightarrow ((a, b), \text{gcd-body}, \emptyset)$
 $a \rightarrow 10$
 $b \rightarrow 15$
 line 11: gcd a b $\rightarrow \text{? } 5$
 ~~$E_1 : ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset)$~~
 line 3: gcd b a $\rightarrow \text{? } 5$
 ~~$E_2 : ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset)$~~
 line 5: a % b $\rightarrow 5$
 line 5: gcd b (a % b) $\rightarrow \text{? } 5$
 ~~$E_3 : ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset)$~~
 line 5: a % b $\rightarrow 0$
 line 5: gcd b (a % b) $\rightarrow \text{? } 5$
 ~~$E_4 : ((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \emptyset)$~~

Now we are able to continue the program in environment E_0 with the `printfn` statement, and we

write:

$$\begin{aligned}
 E_0 : & \\
 & \text{gcd} \rightarrow ((a, b), \text{gcd-body}, \emptyset) \\
 & a \rightarrow 10 \\
 & b \rightarrow 15 \\
 & \text{line 11: gcd a b} \rightarrow \backslash 5 \\
 & \text{line 11: stdout} \rightarrow \text{"gcd 10 15 = 5"} \\
 \cancel{E_1} : & ((a \rightarrow 10, b \rightarrow 15), \text{gcd-body}, \emptyset) \\
 & \text{line 3: gcd b a} \rightarrow \backslash 5 \\
 \cancel{E_2} : & ((a \rightarrow 15, b \rightarrow 10), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 5 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow \backslash 5 \\
 \cancel{E_3} : & ((a \rightarrow 10, b \rightarrow 5), \text{gcd-body}, \emptyset) \\
 & \text{line 5: a \% b} \rightarrow 0 \\
 & \text{line 5: gcd b (a \% b)} \rightarrow \backslash 5 \\
 \cancel{E_4} : & ((a \rightarrow 5, b \rightarrow 0), \text{gcd-body}, \emptyset)
 \end{aligned}$$

which completes the hand tracing of `gcd.fsx`.

F# uses the lexical scope, which implies that besides function arguments, we also at times need to consider the environment at the place of writing. Consider the program in Listing 10.7.

Listing 10.7 lexicalScopeTracing.fsx: lexicalScopeTracing

```

1  let testScope x =
2      let a = 3.0
3      let f z = a * z
4      let a = 4.0
5      f x
6  printfn "%A" (testScope 2.0)

```

```

1  $ fsharp --nologo lexicalScopeTracing.fsx
2  $ mono lexicalScopeTracing.exe
3  6.0

```

Example of lexical scope and closure environment. To hand trace this, we start by creating the outer environment, define the closure for `testScope` and reach line 6,

$$\begin{aligned}
 E_0 : & \\
 & \text{testScope} \rightarrow (x, \text{testScope-body}, \emptyset) \\
 & \text{line 6: testScope 2.0} \rightarrow ?
 \end{aligned}$$

We create new environment for `testScope` and note the bindings,

$$\begin{aligned}
 E_0 : & \\
 & \text{testScope} \rightarrow (x, \text{testScope-body}, \emptyset) \\
 & \text{line 6: testScope 2.0} \rightarrow ? \\
 E_1 : & (x \rightarrow 2.0, \text{testScope-body}, \emptyset) \\
 & a \rightarrow 3.0 \\
 & f \rightarrow (z, a * z, (a \rightarrow 3.0, x \rightarrow 2.0)) \\
 & a \rightarrow 4.0
 \end{aligned}$$

Since we are working with lexical scope, then `a` is noted twice, and its interpretation is by lexical order. Hence, the environment for the closure of `f` is everything above in E_1 , so we add $a \rightarrow 3.0$ and $x \rightarrow 2.0$. In line 5 `f` is called, so we create an environment based on its closure,

$$\begin{aligned}
 E_0 : & \\
 & \text{testScope} \rightarrow (x, \text{testScope-body}, \emptyset) \\
 & \text{line 6: testScope 2.0} \rightarrow ? \\
 E_1 : & (x \rightarrow 2.0, \text{testScope-body}, \emptyset) \\
 & a \rightarrow 3.0 \\
 & f \rightarrow (z, a * z, (a \rightarrow 3.0, x \rightarrow 2.0)) \\
 & a \rightarrow 4.0 \\
 & \text{line 5: f x} \rightarrow ? \\
 E_2 : & (z \rightarrow 2.0, a * z, (a \rightarrow 3.0, x \rightarrow 2.0))
 \end{aligned}$$

The expression in the environment E_2 evaluates to `6.0`, and unraveling the scopes we get,

$$\begin{aligned}
 E_0 : & \\
 & \text{testScope} \rightarrow (x, \text{testScope-body}, \emptyset) \\
 & \text{line 6: testScope 2.0} \rightarrow \text{\textbackslash} 6.0 \\
 & \text{line 6: stdout} \rightarrow "6.0" \\
 E_1 : & (x \rightarrow 2.0, \text{testScope-body}, \emptyset) \\
 & a \rightarrow 3.0 \\
 & f \rightarrow (z, a * z, (a \rightarrow 3.0, x \rightarrow 2.0)) \\
 & a \rightarrow 4.0 \\
 & \text{line 5: f x} \rightarrow \text{\textbackslash} 6.0 \\
 E_2 : & (z \rightarrow 2.0, a * z, (a \rightarrow 3.0, x \rightarrow 2.0))
 \end{aligned}$$

For mutable bindings, i.e., variables, the scope is dynamic. For this, we need the concept of storage. Consider the program in Listing 10.8.

**Listing 10.8 dynamicScopeTracing.fsx:
dynamicScopeTracing**

```

1  let testScope x =
2      let mutable a = 3.0
3      let f z = a * z
4      a <- 4.0
5      f x
6  printfn "%A" (testScope 2.0)

```

```

1  $ fsharp --nologo dynamicScopeTracing.fsx
2  $ mono dynamicScopeTracing.exe
3  8.0

```

Example of lexical scope and closure environment. We add a storage area to our hand tracing, e.g., line 6,

Store :

E_0 :

testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$

line 6: testScope 2.0 $\rightarrow ?$

So when we generate environment E_1 , the mutable binding is to a storage location,

Store :

$\alpha_1 \rightarrow 3.0$

E_0 :

testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$

line 6: testScope 2.0 $\rightarrow ?$

$E_1 : (x \rightarrow 2.0, \text{testScope-body}, \emptyset)$

$a \rightarrow \alpha_1$

which is assigned the value 3.0 at the definition of **a**. Now the definition of **f** is using the storage location

Store :

$\alpha_1 \rightarrow 3.0$

E_0 :

testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$

line 6: testScope 2.0 $\rightarrow ?$

$E_1 : (x \rightarrow 2.0, \text{testScope-body}, \emptyset)$

$a \rightarrow \alpha_1$

$f \rightarrow (z, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$

and in line 4 it is the value in the storage, which is updated,

Store :
 $\alpha_1 \rightarrow \text{3.0 4.0}$
 $E_0 :$
 testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$
 line 6: testScope 2.0 $\rightarrow ?$
 $E_1 : (x \rightarrow 2.0, \text{testScope-body}, \emptyset)$
 $a \rightarrow \alpha_1$
 $f \rightarrow (z, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$

Hence,

Store :
 $\alpha_1 \rightarrow \text{3.0 4.0}$
 $E_0 :$
 testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$
 line 6: testScope 2.0 $\rightarrow ?$
 $E_1 : (x \rightarrow 2.0, \text{testScope-body}, \emptyset)$
 $a \rightarrow \alpha_1$
 $f \rightarrow (z, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$
 line 5: f x $\rightarrow ?$
 $E_2 : (z \rightarrow 2.0, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$

and the return value from **f** evaluated in environment E_2 now reads the value 4.0 for **a** and returns 8.0. Hence,

Store :
 $\alpha_1 \rightarrow \text{3.0 4.0}$
 $E_0 :$
 testScope $\rightarrow (x, \text{testScope-body}, \emptyset)$
 line 6: testScope 2.0 $\rightarrow \text{8.0}$
 line 6: stdout $\rightarrow \text{"8.0"}$
 $E_1 : (x \rightarrow 2.0, \text{testScope-body}, \emptyset)$
 $a \rightarrow \alpha_1$
 $f \rightarrow (z, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$
 line 5: f x $\rightarrow \text{8.0}$
 $E_2 : (z \rightarrow 2.0, a * z, (a \rightarrow \alpha_1, x \rightarrow 2.0))$

By the above examples, it is seen that hand tracing can be used to in detail study the flow of data through a program. It may seem tedious in the beginning, but the care illustrated above is useful at the start to ensure rigor in the analysis. Most will find that once accustomed to the method, the analysis can be performed rigorously but with less paperwork, and in conjunction with strategically placed debugging **printfn** statements, it is a very valuable tool for debugging.

11 | Collections of data

F# is tuned to work with collections of data, and there are several built-in types of collections with various properties making them useful for different tasks. Examples include strings, lists, arrays, and sequences. Strings were discussed in Chapter 5 and will be revisited here in more details. Sequences will not be discussed,¹ and we will concentrate on lists and one- and two-dimensional arrays.

11.1 Strings

Strings have been discussed in Chapter 5, the content of which will be briefly revisited here followed by a description of some of the many supporting built-in functions in F# on strings.

A *string* is a sequence of characters. Each character is represented using UTF-16, see Appendix C for further details on the unicode standard. The type `string` is an alias for `System.string`. String literals are delimited by quotation marks “” and inside the delimiters, character escape sequences are allowed (see Table 5.2), which are replaced by the corresponding character code. Examples are `"This is a string"`, `"\tTabulated string"`, `"A \"quoted\" string"`, and `""`. Strings may span several lines, and new lines inside strings are part of the string unless the line is ended with a backslash. Strings may be *verbatim* by preceding the string with “@”, in which case escape sequences are not replaced, but a double quotation marks is an escape sequence which is replaced by a single, e.g., `@"This is a string"`, `@"\tNon-tabulated string"`, `@"A \"quoted\" string"`, and `@""`. Alternatively, a verbatim string may be delimited by tripple quotes, e.g., `"""This is a string"""`, `"""\tNon-tabulated string"""`, `"""A "quoted" string"""`, and `""""`. Strings may be indexed using the `.[]` notation, as demonstrated in Listing 5.27.

11.1.1 String properties

Strings have a few properties, which are values attached to each string and access using the “.” notation. The only to be mentioned here is:

Length: Returns the length of the string. Compare with `String.length` method.

Listing 11.1: Length

```
1 > "abcd".Length;;  
2 val it : int = 4
```

¹Jon: Should we discuss sequences?

11.1.2 String module

In the `String` module the following functions are available.

`String.collect: (char -> string) -> string -> string`. Creates a new string whose characters are the results of applying a specified function to each of the characters of the input string and concatenating the resulting strings.

Listing 11.2: `String.collect`

```
1 > String.collect (fun c -> (string c) + ", ") "abc";;  
2 val it : string = "a, b, c, "
```

`String.concat: string -> seq<string> -> string`. Returns a new string made by concatenating the given strings with a separator. Here `seq<string>` is a sequence but can also be a list or an array.

Listing 11.3: `String.concat`

```
1 > String.concat ", " ["abc"; "def"; "ghi"];;  
2 val it : string = "abc, def, ghi"
```

`String.exists: (char -> bool) -> string -> bool`. Tests if any character of the string satisfies the given predicate.

Listing 11.4: `String.exists`

```
1 > String.exists (fun c -> c = 'd') "abc";;  
2 val it : bool = false
```

`String.forall: (char -> bool) -> string -> bool`. Tests if all characters in the string satisfy the given predicate.

Listing 11.5: `String.forall`

```
1 > String.forall (fun c -> c < 'd') "abc";;  
2 val it : bool = true
```

`String.init: int -> (int -> string) -> string`. Creates a new string whose characters are the results of applying a specified function to each index and concatenating the resulting strings.

Listing 11.6: `String.init`

```
1 > String.init 5 (fun i -> (string i) + ", ");;  
2 val it : string = "0, 1, 2, 3, 4, "
```

`String.iter: (char -> unit) -> string -> unit`. Applies a specified function to each character in the string.

Listing 11.7: String.iter

```
1 > String.iter (fun c -> printfn "%c" c) "abc";;  
2 a  
3 b  
4 c  
5 val it : unit = ()
```

String.iteri: (int -> char -> unit) -> string -> unit. Applies a specified function to the index of each character in the string and the character itself.

Listing 11.8: String.iteri

```
1 > String.iteri (fun i c -> printfn "%d: %c" i c) "abc";;  
2 0: a  
3 1: b  
4 2: c  
5 val it : unit = ()
```

String.length: string -> int. Returns the length of the string.

Listing 11.9: String.length

```
1 > String.length "abcd";;  
2 val it : int = 4
```

String.map: (char -> char) -> string -> string. Creates a new string whose characters are the results of applying a specified function to each of the characters of the input string.

Listing 11.10: String.map

```
1 > let dc = int 'A' - int 'a'  
2 - String.map (fun c -> c + char dc) "abcd";;  
3 val dc : int = -32  
4 val it : string = "ABCD"
```

String.mapi: (int -> char -> char) -> string -> string. Creates a new string whose characters are the results of applying a specified function to each character and index of the input string.

Listing 11.11: String.mapi

```
1 > String.mapi (fun i c -> char (int c + i)) "aaaa";;  
2 val it : string = "abcd"
```

String.replicate: int -> string -> string. Returns a string by concatenating a specified number of instances of a string.

Listing 11.12: String.replicate

```
1 > String.replicate 4 "abc, ";;  
2 val it : string = "abc, abc, abc, abc, "
```

11.2 Lists

Lists are unions of immutable values of the same type and have a more flexible structure than tuples. Lists can be expressed as a *sequence expression*,

· list
· sequence expression

Listing 11.13 Lists with a sequence expression.

```
1 [[<expr>; <expr>]]
```

Examples are `[1; 2; 3; 4; 5]`, which represents a list of integers, `["This"; "is"; "a"; "list"]`, which represents a list of strings, `[(fun x -> x); (fun x -> x*x)]`, which represents a list of anonymous functions, and `[]`, which is an empty list. Lists may also be given as ranges,

Listing 11.14 Lists with a range expressions.

```
1 [<expr> .. <expr> [... <expr>]]
```

where `<expr>` in *range expressions* must be of integers, floats, or characters. Examples are `[1 .. 5]`, `[-3.0 .. 2.0]`, and `['a' .. 'z']`. Range expressions may include a step size, thus, `[1 .. 2 .. 10]` evaluates to `[1; 3; 5; 7; 9]`.

· range expressions

A list type is identified with the `list` keyword, such that a list of integers has the type `int list`. Like strings, lists may be indexed using the `.[]` notation, the lengths of lists is retrieved using the `Length` property, and we may test whether a list is empty using the `isEmpty` property. These features are demonstrated in Listing 11.15.

· list
· .[]
· Length
· isEmpty

Listing 11.15 listIndexing.fsx:

Lists are indexed as strings and has a `Length` property.

```
1 let printList (lst : int list) : unit =
2     for i = 0 to lst.Length - 1 do
3         printf "%A " lst.[i]
4     printfn ""
5
6 let lst = [3; 4; 5]
7 printfn "lst = %A, lst.[1] = %A" lst lst.[1]
8 printfn "lst.Length = %A, lst.isEmpty = %A" lst.Length lst.IsEmpty
9 printList lst

```

```
1 $ fsharpc --nologo listIndexing.fsx && mono listIndexing.exe
2 lst = [3; 4; 5], lst.[1] = 4
3 lst.Length = 3, lst.isEmpty = false
4 3 4 5
```

F# implements lists as linked lists, see Figure 11.1, which is why indexing element i has computational complexity $\mathcal{O}(i)$, since the list has to be traversed from the beginning until element i is located. Thus, **indexing lists is slow and should be avoided**.

Advice

Notice especially that lists are zero-indexed, and thus, the last element in a list `lst` is `lst.Length - 1`. This is a very common source of error! Therefore, indexing in lists using *for*-loops is supported using a special notation with the `in` keyword,

· for
· in




Figure 11.1: A list is a linked list: Here is illustrated the linked list of ["This"; "is"; "a"; "list"].

Listing 11.16 For-in loop with in expression.

```
1 for <ident> in <list> do <bodyExpr> [done]
```

In `for-in` loops, the loop runs through each element of the `<list>`, and assigns it to the identifier `<ident>`. This is demonstrated in Listing 11.17.

Listing 11.17 listFor.fsx:

The `for-in` loops are preferred over `for-to`.

```

1 let printList (lst : int list) : unit =
2   for elm in lst do
3     printfn "%A" elm
4     printfn ""
5
6   printList [3; 4; 5]

-----

1 $ fsharpc --nologo listFor.fsx && mono listFor.exe
2 3 4 5

```

Using `for-in`-expressions remove the risk of off-by-one indexing errors, and thus, **`for-in` is to be preferred over `for-to`**. Advice

Lists support slicing identically to strings as demonstrated in Listing 11.18.

Listing 11.18 listSlicing.fsx:

Examples of list slicing. Compare with Listing 5.27.

```

1 let lst = ['a' .. 'g']
2 printfn "%A" lst.[0]
3 printfn "%A" lst.[3]
4 printfn "%A" lst.[3..]
5 printfn "%A" lst[..3]
6 printfn "%A" lst.[1..3]
7 printfn "%A" lst.[*]

-----

1 $ fsharpc --nologo listSlicing.fsx && mono listSlicing.exe
2 'a'
3 'd'
4 ['d'; 'e'; 'f'; 'g']
5 ['a'; 'b'; 'c'; 'd']
6 ['b'; 'c'; 'd']
7 ['a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g']

```





Figure 11.2: A list is a ragged linked list: Here is illustrated the linked list of `[[1;2];[3;4;5]]`.

Lists may be concatenated using either the “@”² *concatenation* operator or the “::” *cons* operators. The differences is that “@” concatenates two lists of identical types, while “::” concatenates an element and a list of identical types. This is demonstrated in Listing 11.19.

· @
· list concatenation
· ::
· list cons

Listing 11.19 listCon.fsx:
Examples of list concatenation.

```
1 printfn "[1] @ [2; 3] = %A" ([1] @ [2; 3])
2 printfn "[1; 2] @ [3; 4] = %A" ([1; 2] @ [3; 4])
3 printfn "1 :: [2; 3] = %A" (1 :: [2; 3])

1 $ fsharp -nologo listCon.fsx && mono listCon.exe
2 [1] @ [2; 3] = [1; 2; 3]
3 [1; 2] @ [3; 4] = [1; 2; 3; 4]
4 1 :: [2; 3] = [1; 2; 3]
```

Since lists are represented as linked lists, then the cons operator is very efficient and has computational complexity $\mathcal{O}(1)$, while concatenation has computational complexity $\mathcal{O}(n)$, where n is the length of the first list.

It is possible to make multidimensional lists as lists of lists as shown in Listing 11.20.

Listing 11.20 listMultidimensional.fsx:
A ragged multidimensional list, built as lists of lists, and its indexing.

```
1 let a = [[1;2];[3;4;5]]
2 let row = a.Item 0 in printfn "%A" row
3 let elm = row.Item 1 in printfn "%A" elm
4 let elm = (a.Item 0).Item 1 in printfn "%A" elm

1 $ fsharp -nologo listMultidimensional.fsx
2 $ mono listMultidimensional.exe
3 [1; 2]
4 2
5 2
```

The example shows a *ragged multidimensional list* since each row has a different number of elements. This is also illustrated in Figure 11.2.

· ragged
· multidimensional list

The indexing of a particular element is slow due to the linked list implementation of lists, which is

²Jon: why does the at-symbol not appear in the index?

why arrays are often preferred for two- and higher-dimensional data structures, see Section 11.3.

11.2.1 List properties

Lists support a number of properties, i.e., values that are attached to each list and access using the “.” notation, some of which are:

Head: Returns the first element of a list.

Listing 11.21: Head

```
1 > [1; 2; 3].Head;;  
2 val it : int = 1
```

IsEmpty: Returns true if the list is empty.

Listing 11.22: Head

```
1 > [1; 2; 3].IsEmpty;;  
2 val it : bool = false
```

Length: Returns the number of elements in the list.

Listing 11.23: Length

```
1 > [1; 2; 3].Length;;  
2 val it : int = 3
```

Tail: Returns the list except its first element.

Listing 11.24: Tail

```
1 > [1; 2; 3].Tail;;  
2 val it : int list = [2; 3]
```

11.2.2 List module

The built-in `List` module contains a wealth of functions for lists, some of which are briefly summarized below:

List.collect: ('T -> 'U list) -> 'T list -> 'U list. Apply the supplied function to each element in a list and return a concatenated list of the results.

Listing 11.25: List.collect

```
1 > List.collect (fun elm -> [elm; elm; elm]) [1; 2; 3];;  
2 val it : int list = [1; 1; 1; 2; 2; 2; 3; 3; 3]
```

List.contains: 'T -> 'T list -> bool. Returns true or false depending on whether an element is contained in the list.

Listing 11.26: List.contains

```
1 > List.contains 3 [1; 2; 3];;
2 val it : bool = true
```

List.empty: 'T list. An empty list of inferred type.

Listing 11.27: List.empty

```
1 > let a : int list = List.empty;;
2 val a : int list = []
```

List.exists: ('T -> bool) -> 'T list -> bool. Returns true or false depending on whether any element is true for a given function.

Listing 11.28: List.exists

```
1 > let odd x = (x % 2 = 1) in List.exists odd [0 .. 2 .. 4];;
2 val it : bool = false
```

List.filter: ('T -> bool) -> 'T list -> 'T list. Returns a new list, of all the elements of the original list for which the supplied function evaluates to true.

Listing 11.29: List.filter

```
1 > let odd x = (x % 2 = 1) in List.filter odd [0 .. 9];;
2 val it : int list = [1; 3; 5; 7; 9]
```

List.find: ('T -> bool) -> 'T list -> 'T. Return the first element for which the given function is true.

Listing 11.30: List.find

```
1 > let odd x = (x % 2 = 1) in List.find odd [0 .. 9];;
2 val it : int = 1
```

List.findIndex: ('T -> bool) -> 'T list -> int. Return the index of the first element for which the given function is true.

Listing 11.31: List.findIndex

```
1 > let isK x = (x = 'k') in List.findIndex isK ['a' .. 'z'];;
2 val it : int = 10
```

List.fold: ('State -> 'T -> 'State) -> 'State -> 'T list -> 'State. Update an accumulator iteratively by applying the supplied function to each element in a list, e.g. for a list consisting of $x_0, x_1, x_2, \dots, x_n$, a supplied function f , and an initial value for the accumulator s , calculate $f(\dots f(f(f(s, x_0), x_1), x_2), \dots, x_n)$.

Listing 11.32: List.fold

```

1 > let addSquares acc elm = acc + elm*elm
2 - List.fold addSquares 0 [0 .. 9];;
3 val addSquares : acc:int -> elm:int -> int
4 val it : int = 285

```

List.foldBack: ('T -> 'State -> 'State) -> 'T list -> 'State -> 'State. Update an accumulator iteratively by applying function to each element in a list, e.g. for a list consisting of $x_0, x_1, x_2, \dots, x_n$, a supplied function f , and an initial value for the accumulator s , calculate $f(x_0, f(x_1, f(x_2, \dots, f(x_n, s))))$.

Listing 11.33: List.foldBack

```

1 > let addSquares elm acc = acc + elm*elm
2 - List.foldBack addSquares [0 .. 9] 0;;
3 val addSquares : elm:int -> acc:int -> int
4 val it : int = 285

```

List.forall: ('T -> bool) -> 'T list -> bool. Apply a function to all element and logically and the result.

Listing 11.34: List.forall

```

1 > let odd x = (x % 2 = 1) in List.forall odd [0 .. 9];;
2 val it : bool = false

```

List.head: 'T list -> int. The first element in the list. Exception if empty.

Listing 11.35: List.head

```

1 > let a = [1; -2; 0] in List.head a;;
2 val it : int = 1

```

List.isEmpty: 'T list -> bool. Compare with the empty list

Listing 11.36: List.isEmpty

```

1 > List.isEmpty [1; 2; 3];;
2 val it : bool = false
3
4 > let a = [1; 2; 3] in List.isEmpty a;;
5 val it : bool = false

```

List.item: 'T list -> int. Retrieve an element of a list by its index.

Listing 11.37: List.item

```

1 > let a = [1; -3; 0] in List.item 1 a;;
2 val it : int = -3

```

`List.iter: ('T -> unit) -> 'T list -> unit`. Apply a procedure to every element in the list.

Listing 11.38: List.iter

```
1 > let prt x = printfn "%A " x in List.iter prt [0; 1; 2];;
2 0
3 1
4 2
5 val it : unit = ()
```

`List.Length: 'T list -> int`. The number of elements in a list

Listing 11.39: List.Length

```
1 > List.length [1; 2; 3];;
2 val it : int = 3
```

`List.map: ('T -> 'U) -> 'T list -> 'U list`. Return a list, where the supplied function has been applied to every element.

Listing 11.40: List.map

```
1 > let square x = x*x in List.map square [0 .. 9];;
2 val it : int list = [0; 1; 4; 9; 16; 25; 36; 49; 64; 81]
```

`List.ofArray: 'T list -> int`. Return a list whose elements are the same as the supplied array.

Listing 11.41: List.ofArray

```
1 > List.ofArray [|1; 2; 3|];;
2 val it : int list = [1; 2; 3]
```

`List.rev: 'T list -> 'T list`. Return a list whose elements have been reversed.

Listing 11.42: List.rev

```
1 > List.rev [1; 2; 3];;
2 val it : int list = [3; 2; 1]
```

`List.sort: 'T list -> 'T list`. Return a list whose elements have been sorted.

Listing 11.43: List.sort

```
1 > List.sort [3; 1; 2];;
2 val it : int list = [1; 2; 3]
```

`List.tail: 'T list -> 'T list`. The list except its first element. Exception if empty.

Listing 11.44: List.tail

```

1 > List.tail [1; 2; 3];;
2 val it : int list = [2; 3]
3
4 > let a = [1; 2; 3] in List.tail a;;
5 val it : int list = [2; 3]

```

List.toArray: 'T list -> 'T []. Return an array whos elements are the same as the supplied list.

Listing 11.45: List.toArray

```

1 > List.toArray [1; 2; 3];;
2 val it : int [] = [|1; 2; 3|]

```

List.unzip: ('T1 * 'T2) list -> 'T1 list * 'T2 list. Return a pair of lists, whos elements are take from pairs of a list.

Listing 11.46: List.unzip

```

1 > List.unzip [(1, 'a'); (2, 'b'); (3, 'c')];;
2 val it : int list * char list = ([1; 2; 3], ['a'; 'b'; 'c'])
3
4 >

```

List.zip: 'T1 list -> 'T2 list -> ('T1 * 'T2) list. Return a list of pairs, whos elements are take iteratively from two lists.

Listing 11.47: List.zip

```

1 > List.zip [1; 2; 3] ['a'; 'b'; 'c'];;
2 val it : (int * char) list = [(1, 'a'); (2, 'b'); (3, 'c')]

```

11.3 Arrays

One dimensional *arrays* or just arrays for short are mutable lists of the same type and follow a similar syntax as lists. Arrays can be stated as *sequence expressions*,

· arrays
· sequence expressions

Listing 11.48 Arrays with a sequence expression.

```

1 [|<expr>; <expr>|]

```

and examples are [|1; 2; 3; 4; 5|], which is an array of integers, [|"This"; "is"; "an"; "array"|], which is an array of strings, [|fun x -> x; (fun x -> x*x)|], which is an array of anonymous functions, and [|]|, which is an empty array. Arrays may also be given as ranges,

Listing 11.49 Arrays with a range expressions.

```
1 [|<expr> .. <expr> [|.. <expr>]|]
```

but arrays of *range expressions* must be of `<expr>` integers, floats, or characters. Examples are `[|1 .. 5|]`, `[|-3.0 .. 2.0|]`, and `[|'a' .. 'z'|]`. Range expressions may include a step size, thus, `[|1 .. 2 .. 10|]` evaluates to `[|1; 3; 5; 7; 9|]`.

The array type is defined using the `array` keyword or alternatively the “`[]`” lexeme. Like strings and lists, arrays may be indexed using the “`.[]`” notation. Arrays cannot be resized, but are mutable as shown in Listing 11.50.

Listing 11.50 arrayReassign.fsx:

Arrays are mutable in spite the missing `mutable` keyword.

```
1 let square (a : int array) =
2     for i = 0 to a.Length - 1 do
3         a.[i] <- a.[i] * a.[i]
4
5 let A = [| 1; 2; 3; 4; 5 |]
6 printfn "%A" A
7 square A
8 printfn "%A" A

1 $ fsharp -nologo arrayReassign.fsx && mono arrayReassign.exe
2 [|1; 2; 3; 4; 5|]
3 [|1; 4; 9; 16; 25|]
```

Notice that in spite the missing `mutable` keyword, the function `square` still had the *side-effect* of squaring all entries in `A`. F# implements arrays as chunks of memory and indexes arrays via address arithmetic. I.e., element i in an array, whose first element is in memory address α and whose elements fill β addresses each is found at address $\alpha + i\beta$.³ Hence, indexing has computational complexity of $\mathcal{O}(1)$, but appending and prepending values to arrays and array concatenation requires copying the new and existing values to a fresh area in memory and thus has computational complexity $\mathcal{O}(n)$, where n is the total number of elements. Thus, **indexing arrays is fast, but cons and concatenation is slow and should be avoided.** Advice

Arrays support *slicing*, that is, indexing an array with a range results in a copy of the array with values corresponding to the range. This is demonstrated in Listing 11.51.

³Jon: Add a figure illustrating address indexing.

Listing 11.51 arraySlicing.fsx:

Examples of array slicing. Compare with Listing 11.18 and Listing 5.27.

```

1 let arr = [|'a' .. 'g'|]
2 printfn "%A" arr.[0]
3 printfn "%A" arr.[3]
4 printfn "%A" arr.[3..]
5 printfn "%A" arr[..3]
6 printfn "%A" arr.[1..3]
7 printfn "%A" arr.[*]

-----

1 $ fsharpc --nologo arraySlicing.fsx && mono arraySlicing.exe
2 'a'
3 'd'
4 [|'d'; 'e'; 'f'; 'g'|]
5 [|'a'; 'b'; 'c'; 'd'|]
6 [|'b'; 'c'; 'd'|]
7 [|'a'; 'b'; 'c'; 'd'; 'e'; 'f'; 'g'|]

```

As illustrated, the missing start or end index implies from the first or to the last element.

Arrays have explicit operator support for appending and concatenation, instead the `Array` namespace includes an `Array.append` function, as shown in Listing 11.52.

Listing 11.52 arrayAppend.fsx:Two arrays are appended with `Array.append`.

```

1 let a = [|1; 2;|]
2 let b = [|3; 4; 5|]
3 let c = Array.append a b
4 printfn "%A, %A, %A" a b c

-----

1 $ fsharpc --nologo arrayAppend.fsx && mono arrayAppend.exe
2 [|1; 2|], [|3; 4; 5|], [|1; 2; 3; 4; 5|]

```

Arrays are *reference types*, meaning that identifiers are references and thus suffers from aliasing, as reference types illustrated in Listing 11.53.

Listing 11.53 arrayAliasing.fsx:

Arrays are reference types and suffer from aliasing.

```

1 let a = [|1; 2; 3|];
2 let b = a
3 a.[0] <- 0
4 printfn "a = %A, b = %A" a b;;

-----

1 $ fsharpc --nologo arrayAliasing.fsx && mono arrayAliasing.exe
2 a = [|0; 2; 3|], b = [|0; 2; 3|]

```


11.3.1 Array properties and methods

Arrays support a number of properties and methods, i.e., values and functions that are attached to each array and access using the “.” notation, some of which are:

Clone(): Returns a copy of the array.

Listing 11.54: Clone

```
1 > let a = [|1; 2; 3|];
2 - let b = a.Clone()
3 - a.[0] <- 0
4 - printfn "a = %A, b = %A" a b;;
5 a = [|0; 2; 3|], b = [|1; 2; 3|]
6 val a : int [] = [|0; 2; 3|]
7 val b : obj = [|1; 2; 3|]
8 val it : unit = ()
```

Length: Returns the number of elements in the array.

Listing 11.55: Length

```
1 > [|1; 2; 3|].Length;;
2 val it : int = 3
```

11.3.2 Array module

There are quite a number of built-in procedures for arrays in the **Array** module, some of which are summarized below.⁴

Array.append: 'T [] -> 'T [] -> 'T []. Creates an array that contains the elements of one array followed by the elements of another array.

Listing 11.56: Array.append

```
1 > Array.append [|1; 2;|] [|3; 4; 5|];;
2 val it : int [] = [|1; 2; 3; 4; 5|]
```

Array.compareWith: ('T -> 'T -> int) -> 'T [] -> 'T [] -> int. Compares two arrays using the given comparison function, element by element.

Listing 11.57: Array.compareWith

```
1 > let compArr elm1 elm2 =
2 -   if elm1 > elm2 then 1
3 -   elif elm1 < elm2 then -1
4 -   else 0
5 - Array.compareWith compArr [|1; 2; 4|] [|1; 2; 3|];;
6 val compArr : elm1:'a -> elm2:'a -> int when 'a : comparison
7 val it : int = 1
```

⁴Jon: rewrite description

`Array.concat: seq<'T []> -> 'T []`. Creates an array that contains the elements of each of the supplied sequence of arrays.

Listing 11.58: Array.concat

```
1 > Array.concat [|1; 2; 3|]; [|4; 5|]; [|6; 7; 8|];;  
2 val it : int [] = [|1; 2; 3; 4; 5; 6; 7; 8|]
```

`Array.contains: .` Evaluates to true if the given element is in the input array.

Listing 11.59: Array.contains

```
1 > Array.contains 3 [|1; 2; 3|];;  
2 val it : bool = true
```

`Array.copy: 'T [] -> 'T []`. Creates an array that contains the elements of the supplied array.

Listing 11.60: Array.copy

```
1 > let a = [|1; 2; 3|]  
2 - let b = Array.copy a;;  
3 val a : int [] = [|1; 2; 3|]  
4 val b : int [] = [|1; 2; 3|]
```

`Array.create: int -> 'T -> 'T []`. Creates an array whose elements are initialized the supplied value.

Listing 11.61: Array.create

```
1 > Array.create 4 3.14;;  
2 val it : float [] = [|3.14; 3.14; 3.14; 3.14|]
```

`Array.empty: 'T []`. Returns an empty array of the given type.

Listing 11.62: Array.empty

```
1 > let a : int [] = Array.empty;;  
2 val a : int [] = [|]|
```

`Array.exists: ('T -> bool) -> 'T [] -> bool`. Tests whether any element of an array satisfies the supplied predicate.

Listing 11.63: Array.exists

```
1 > let odd x = (x % 2 = 1) in Array.exists odd [|0 .. 2 .. 4|];;  
2 val it : bool = false
```

`Array.fill: 'T [] -> int -> int -> 'T -> unit`. Fills a range of elements of an array with the supplied value.

Listing 11.64: Array.fill

```

1 > let arr = Array.zeroCreate 10;
2 - Array.fill arr 2 5 2;;
3 val arr : int [] = [|0; 0; 2; 2; 2; 2; 2; 0; 0; 0|]
4 val it : unit = ()

```

Array.filter: ('T -> bool) -> 'T [] -> 'T []. Returns a collection that contains only the elements of the supplied array for which the supplied condition returns true.

Listing 11.65: Array.filter

```

1 > let odd x = (x % 2 = 1) in Array.filter odd [|0 .. 9|];;
2 val it : int [] = [|1; 3; 5; 7; 9|]

```

Array.find: ('T -> bool) -> 'T [] -> 'T. Returns the first element for which the supplied function returns true. Raises `System.Collections.Generic.KeyNotFoundException`

Listing 11.66: Array.find

```

1 > let odd x = (x % 2 = 1) in Array.find odd [|0 .. 9|];;
2 val it : int = 1

```

Array.findIndex: ('T -> bool) -> 'T [] -> int. Returns the index of the first element in an array that satisfies the supplied condition. Raises `System.Collections.Generic.KeyNotFoundException` if none of the elements satisfy the condition.

Listing 11.67: Array.findIndex

```

1 > let isK x = (x = 'k') in Array.findIndex isK [|'a' .. 'z'|];;
2 val it : int = 10

```

Array.fold: ('State -> 'T -> 'State) -> 'State -> 'T [] -> 'State. Applies a function to each element of an array, threading an accumulator argument through the computation. If the input function is `f` and the array elements are `i0...iN`, this function computes `f (...(f s i0)...) iN`.

Listing 11.68: Array.fold

```

1 > let addSquares acc elm = acc + elm*elm
2 - Array.fold addSquares 0 [|0 .. 9|];;
3 val addSquares : acc:int -> elm:int -> int
4 val it : int = 285

```

Array.foldBack: ('T -> 'State -> 'State) -> 'T [] -> 'State -> 'State. Applies a function to each element of an array, threading an accumulator argument through the computation. If the input function is `f` and the array elements are `i0...iN`, this function computes `f i0 (...(f iN s))`.

Listing 11.69: Array.foldBack

```

1 > let addSquares elm acc = acc + elm*elm
2 - Array.foldBack addSquares [|0 .. 9|] 0;;
3 val addSquares : elm:int -> acc:int -> int
4 val it : int = 285

```

Array.forall: ('T -> bool) -> 'T [] -> bool. Tests whether all elements of an array satisfy the supplied condition.

Listing 11.70: Array.forall

```

1 > let odd x = (x % 2 = 1) in Array.forall odd [|0 .. 9|];;
2 val it : bool = false

```

Array.get: 'T [] -> int -> 'T. Gets an element from an array.

Listing 11.71: Array.get

```

1 > Array.get [|1; 2; 3|] 2;;
2 val it : int = 3

```

Array.init: int -> (int -> 'T) -> 'T []. Uses a supplied function to create an array of the supplied dimension.

Listing 11.72: Array.init

```

1 > let squareIndex ind = ind*ind
2 - Array.init 10 squareIndex;;
3 val squareIndex : ind:int -> int
4 val it : int [] = [|0; 1; 4; 9; 16; 25; 36; 49; 64; 81|]

```

Array.isEmpty: 'T [] -> bool. Tests whether an array has any elements.

Listing 11.73: Array.isEmpty

```

1 > Array.isEmpty [|];;
2 val it : bool = true

```

Array.iter: ('T -> unit) -> 'T [] -> unit. Applies the supplied function to each element of an array.

Listing 11.74: Array.iter

```

1 > let prt x = printfn "%A " x in Array.iter prt [|0; 1; 2|];;
2 0
3 1
4 2
5 val it : unit = ()

```

Array.length: 'T [] -> int. Returns the length of an array. The `System.Array.Length` property does the same thing.

Listing 11.75: Array.length

```
1 > let a = [|1; 2; 3|] in a.Length;;
2 val it : int = 3
```

Array.map: ('T -> 'U) -> 'T [] -> 'U []. Creates an array whose elements are the results of applying the supplied function to each of the elements of a supplied array.

Listing 11.76: Array.map

```
1 > let square x = x*x in Array.map square [|0 .. 9|];;
2 val it : int [] = [|0; 1; 4; 9; 16; 25; 36; 49; 64; 81|]
```

Array.ofList: 'T list -> 'T []. Creates an array from the supplied list.

Listing 11.77: Array.ofList

```
1 > Array.ofList [|1; 2; 3|];;
2 val it : int [] = [|1; 2; 3|]
```

Array.rev: 'T [] -> 'T []. Reverses the order of the elements in a supplied array.

Listing 11.78: Array.rev

```
1 > Array.rev [|1; 2; 3|];;
2 val it : int [] = [|3; 2; 1|]
```

Array.set: 'T [] -> int -> 'T -> unit. Sets an element of an array.

Listing 11.79: Array.set

```
1 > let arr = [|1; 2; 3|]
2 - Array.set arr 2 10
3 - printfn "%A" arr;;
4 [|1; 2; 10|]
5 val arr : int [] = [|1; 2; 10|]
6 val it : unit = ()
```

Array.sort: 'T[] -> 'T []. Sorts the elements of an array and returns a new array. `Operators.compare` is used to compare the elements.

Listing 11.80: Array.sort

```
1 > Array.sort [|3; 1; 2|];;
2 val it : int [] = [|1; 2; 3|]
```

Array.sub: 'T [] -> int -> int -> 'T []. Creates an array that contains the supplied subrange, which is specified by starting index and length.

Listing 11.81: Array.sub

```

1 > Array.sub [|0..9|] 2 5;;
2 val it : int [] = [|2; 3; 4; 5; 6|]

```

`Array.toList: 'T [] -> 'T list`. Converts the supplied array to a list.

Listing 11.82: Array.toList

```

1 > Array.toList [|1; 2; 3|];;
2 val it : int list = [1; 2; 3]

```

`Array.unzip: ('T1 * 'T2) [] -> 'T1 [] * 'T2 []`. Splits an array of tuple pairs into a tuple of two arrays.

Listing 11.83: Array.unzip

```

1 > Array.unzip [| (1, 'a'); (2, 'b'); (3, 'c') |];;
2 val it : int [] * char [] = ([|1; 2; 3|], [|'a'; 'b'; 'c'|])

```

`Array.zip: 'T1 [] -> 'T2 [] -> ('T1 * 'T2) []`. Combines three arrays into an array of tuples that have three elements. The three arrays must have equal lengths; otherwise, `System.ArgumentException` is raised.

Listing 11.84: Array.zip

```

1 > Array.zip [|1; 2; 3|] [|'a'; 'b'; 'c'|];;
2 val it : (int * char) [] = [| (1, 'a'); (2, 'b'); (3, 'c') |]

```

11.4 Multidimensional arrays

Multidimensional arrays can be created as arrays of arrays (of arrays ...). These are known as *jagged arrays* since there is no inherent guarantee that all sub-arrays are of the same size. E.g., the example in Listing 11.85 is a jagged array of increasing width.

- multidimensional arrays
- jagged arrays

Listing 11.85 arrayJagged.fsx:

An array of arrays. When row lengths are of non-equal elements, then it is a Jagged array.

```

1 let arr = [| [|1|]; [|1; 2|]; [|1; 2; 3|] |]
2
3 for row in arr do
4     for elm in row do
5         printf "%A " elm
6     printf "\n"

```

```

1 $ fsharp -nologo arrayJagged.fsx && mono arrayJagged.exe
2 1
3 1 2
4 1 2 3

```

Indexing arrays of arrays is done sequentially, in the sense that in the above example, the number of outer arrays is `a.Length`, `a.[i]` is the *i*'th array, the length of the *i*'th array is `a.[i].Length`, and the *j*'th element of the *i*'th array is thus `a.[i].[j]`. Often 2-dimensional rectangular arrays are used, which can be implemented as a jagged array as shown in Listing 11.86.

Listing 11.86 arrayJaggedSquare.fsx:

A rectangular array.

```

1 let pownArray (arr : int array array) p =
2     for i = 1 to arr.Length - 1 do
3         for j = 1 to arr.[i].Length - 1 do
4             arr.[i].[j] <- pown arr.[i].[j] p
5
6 let printArrayOfArrays (arr : int array array) =
7     for row in arr do
8         for elm in row do
9             printf "%3d " elm
10        printf "\n"
11
12 let A = [| [|1 .. 4|]; [|1 .. 2 .. 7|]; [|1 .. 3 .. 10|] |]
13 pownArray A 2
14 printArrayOfArrays A

```

```

1 $ fsharp -nologo arrayJaggedSquare.fsx && mono arrayJaggedSquare.exe
2 1 2 3 4
3 1 9 25 49
4 1 16 49 100

```

Notice, the `for-in` cannot be used in `pownArray`, e.g.,

```
for row in arr do for elm in row do elm <- pown elm p done done,
```

since the iterator value `elm` is not mutable even though `arr` is an array. In fact, square arrays of dimensions 2 to 4 are so common that F# has built-in modules for their support. In the following, we describe `Array2D`. The workings of `Array3D` and `Array4D` are very similar. An example of creating the same 2-dimensional array as above but as an `Array2D` is shown in Listing 11.87.

· `Array2D`
 · `Array3D`
 · `Array4D`

Listing 11.87 array2D.fsx:
Creating a 3 by 4 rectangular arrays of integers.

```

1 let arr = Array2D.create 3 4 0
2 for i = 0 to (Array2D.length1 arr) - 1 do
3     for j = 0 to (Array2D.length2 arr) - 1 do
4         arr.[i,j] <- j * Array2D.length1 arr + i
5 printfn "%A" arr

```

```

1 $ fsharpc --nologo array2D.fsx && mono array2D.exe
2 [[0; 3; 6; 9]
3  [1; 4; 7; 10]
4  [2; 5; 8; 11]]

```

Notice that the indexing uses a slightly different notation `[,]` and the length functions are also slightly different. The statement `A.Length` would return the total number of elements in the array, in this case, 12. As can be seen, the `printfn` supports direct printing of the 2-dimensional array. Higher dimensional arrays support slicing as shown in Listing 11.88.

Listing 11.88 array2DSlicing.fsx:
Examples of Array2D slicing. Compare with Listing 11.87.

```

1 let arr = Array2D.create 3 4 0
2 for i = 0 to (Array2D.length1 arr) - 1 do
3     for j = 0 to (Array2D.length2 arr) - 1 do
4         arr.[i,j] <- j * Array2D.length1 arr + i
5 printfn "%A" arr.[2,3]
6 printfn "%A" arr.[1..,3..]
7 printfn "%A" arr[..1,*]
8 printfn "%A" arr.[1,*]
9 printfn "%A" arr.[1..1,*]

```

```

1 $ fsharpc --nologo array2DSlicing.fsx && mono array2DSlicing.exe
2 11
3 [[10]
4  [11]]
5 [[0; 3; 6; 9]
6  [1; 4; 7; 10]]
7 [[1; 4; 7; 10]]
8 [[1; 4; 7; 10]]

```

Note that in almost all cases, slicing produces a sub rectangular 2 dimensional array except for `arr.[1,*]`, which is an array, as can be seen by the single “`[`”. In contrast, `A.[1..1,*]` is an `Array2D`. Note also, that `printfn` typesets 2 dimensional arrays as `[[...]]` and not `[|[| ... |]|]`, which can cause confusion with lists of lists.⁵

Multidimensional arrays have the same properties and methods as arrays, see Section 11.3.1.

⁵Jon: `Array2D.ToString` produces `[[...]]` and not `[|[| ... |]|]`, which can cause confusion.

11.4.1 Array2D module

There are quite a number of built-in procedures for arrays in the `Array2D` namespace, some of which are summarized below.⁶

`copy: 'T [,] -> 'T [,]`. Creates a new array whose elements are the same as the input array.

Listing 11.89: `Array2D.copy`

```
1 > let a = Array2D.init 3 4 (fun i j -> i + 10 * j)
2 - let b = Array2D.copy a;;
3 val a : int [,] = [[0; 10; 20; 30]
4                   [1; 11; 21; 31]
5                   [2; 12; 22; 32]]
6 val b : int [,] = [[0; 10; 20; 30]
7                   [1; 11; 21; 31]
8                   [2; 12; 22; 32]]
```

`create: int -> int -> 'T -> 'T [,]`. Creates an array whose elements are all initially the given value.

Listing 11.90: `Array2D.create`

```
1 > Array2D.create 2 3 3.14;;
2 val it : float [,] = [[3.14; 3.14; 3.14]
3                      [3.14; 3.14; 3.14]]
```

`get: 'T [,] -> int -> int -> 'T`. Fetches an element from a 2D array. You can also use the syntax `array.[index1,index2]`.

Listing 11.91: `Array2D.get`

```
1 > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
2 - Array2D.get arr 1 2;;
3 val arr : int [,] = [[0; 10; 20; 30]
4                   [1; 11; 21; 31]
5                   [2; 12; 22; 32]]
6 val it : int = 21
```

`init: int -> int -> (int -> int -> 'T) -> 'T [,]`. Creates an array given the dimensions and a generator function to compute the elements.

Listing 11.92: `Array2D.init`

```
1 > let idxFct i j = i + 10 * j
2 - Array2D.init 3 4 idxFct;;
3 val idxFct : i:int -> j:int -> int
4 val it : int [,] = [[0; 10; 20; 30]
5                   [1; 11; 21; 31]
6                   [2; 12; 22; 32]]
```

`iter: ('T -> unit) -> 'T [,] -> unit`. Applies the given function to each element of the array.

⁶Jon: rewrite description

Listing 11.93: Array2D.iter

```

1 > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
2 - Array2D.iter (fun elm -> printf "%A " elm) arr
3 - printfn ";;;
4 0 10 20 30 1 11 21 31 2 12 22 32
5 val arr : int [,] = [[0; 10; 20; 30]
6                        [1; 11; 21; 31]
7                        [2; 12; 22; 32]]
8 val it : unit = ()

```

length1: 'T [,] -> int. Returns the length of an array in the first dimension.

Listing 11.94: Array2D.length1

```

1 > let arr = Array2D.create 2 3 0.0 in Array2D.length1 arr;;
2 val it : int = 2

```

length2: 'T [,] -> int. Returns the length of an array in the second dimension.

Listing 11.95: Array2D.forall length2

```

1 > let arr = Array2D.create 2 3 0.0 in Array2D.length2 arr;;
2 val it : int = 3

```

map: ('T -> 'U) -> 'T [,] -> 'U [,]. Creates a new array whose elements are the results of applying the given function to each of the elements of the array.

Listing 11.96: Array2D.map

```

1 > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
2 - let square x = x*x in Array2D.map square arr;;
3 val arr : int [,] = [[0; 10; 20; 30]
4                        [1; 11; 21; 31]
5                        [2; 12; 22; 32]]
6 val it : int [,] = [[0; 100; 400; 900]
7                        [1; 121; 441; 961]
8                        [4; 144; 484; 1024]]

```

set: 'T [,] -> int -> int -> 'T -> unit. Sets the value of an element in an array. You can also use the syntax `array.[index1,index2] <- value`.

Listing 11.97: Array2D.set

```

1 > let arr = Array2D.init 3 4 (fun i j -> i + 10 * j)
2 - Array2D.set arr 1 2 100
3 - printfn "%A" arr;;
4 [[0; 10; 20; 30]
5  [1; 11; 100; 31]
6  [2; 12; 22; 32]]
7 val arr : int [,] = [[0; 10; 20; 30]
8                        [1; 11; 100; 31]
9                        [2; 12; 22; 32]]
10 val it : unit = ()

```

12 | The imperative programming paradigm

Imperative programming is a paradigm for programming states. In imperative programming, the focus is on how a problem is to be solved as a list of *statements* that affects *states*. In F# states are mutable and immutable values, and they are affected by functions and procedures. An imperative program is typically identified as using

- Imperative programming
- statements
- states
- mutable values

mutable values

Mutable values are holders of state, they may change over time, and thus have a dynamic scope.

Procedures

Procedures are functions that returns “()”, instead of functions that transform data. They are the embodiment of side-effects.

- procedures
- side-effects

Side-effects

Side-effects are changes of state that are not reflected in the arguments and return values of a function. The `printf` is an example of a procedure that uses side-effects to communicate with the terminal.

Loops

The `for`- and `while`-loops typically uses an iteration value to update some state, e.g., `for`-loops are often used to iterate through a list and summarize its content.

- `for`
- `while`

In contrast, mono state or stateless programs as *functional programming* can be seen as a subset of imperative programming and is discussed in Chapter 17. *Object oriented programming* is an extension of imperative programming, where statements and states are grouped into classes. For a discussion on object-oriented programming, see Chapter 22.

- functional programming
- Object oriented programming

Imperative programs are like Turing machines, a theoretical machine introduced by Alan Turing in 1936 [10]. Almost all computer hardware is designed for *machine code*, which is a common term used for many low-level computer programming languages, and almost all machine languages follow the imperative programming paradigm.

- machine code

A prototypical example is a baking recipe, e.g., to make a loaf of bread do the following:

1. Mix yeast with water
2. Stir in salt, oil, and flour
3. Knead until the dough has a smooth surface
4. Let the dough rise until it has double size
5. Shape dough into a loaf
6. Let the loaf rise until double size

7. Bake in the oven until the bread is golden brown

Each line in this example consists of one or more statements that are to be executed, and while executing them, states such as the size of the dough and the color of the bread changes. Some execution will halt execution until certain conditions of these states are fulfilled, e.g., the bread will not be put into the oven for baking before it has risen sufficiently.

12.1 Imperative design

Programming is the act of solving a problem by writing a program to be executed on a computer. And imperative programming focusses on states. To solve a problem, you could work through the following list of actions

1. Understand the problem. As Pólya described it, see Chapter 2, the first step in any solution is to understand the problem. A good trick to check, whether you understand the problem is to briefly describe it in your own words.
2. Identify the main values, variables, functions, and procedures needed. If the list of procedures is large, then you most likely should organize them in modules. It is also useful to start in a coarse to fine manner.
3. For each function and procedure, write a precise description of what it should do. This can conveniently be performed as an in-code comment for the procedure using the F# XML documentation standard.
4. Make mockup functions and procedures using the intended types, but don't necessarily compute anything sensible. Run through examples in your mind, using this mockup program to identify any obvious oversights.
5. Write a suite of unit-tests that tests the basic requirements for your code. The unit tests should be runnable with your mockup code. Writing unit-tests will also allow you to evaluate the usefulness of the code pieces as seen from an application point of view.
6. Replace the mockup functions in a prioritized order, i.e., write the must-have code before you write the nice-to-have code, while regularly running your unit-tests to keep track on your progress.
7. Evaluate the code in relation to the desired goal and reiterate earlier actions as needed until the task has been sufficiently completed.
8. Complete your documentation both in-code and outside to ensure that the intended user has sufficient knowledge to effectively use your program and to ensure that you or a fellow programmer will be able to maintain and extend the program in the future.

13 | Recursion

Recursion is a central concept in F# and is used to control flow in loops without the `for` and `while` constructions. Figure 13.1 illustrates the concept of an infinite loop with recursion.




Figure 13.1: An infinitely long string of "forever forever forever...", conceptually calculated by `let rec forever () = "fsharp " + (forever ())`.

13.1 Recursive functions

A *recursive function* is a function, which calls itself, and the syntax for defining recursive functions is an extension of that for regular functions:

Listing 13.1 Syntax for defining one or more mutually dependent recursive functions.

```
1 let rec <ident> = <expr> {and <ident> = <expr>} [in] <expr>
```

From a compiler point of view, the `rec` is necessary, since the function is used before the compiler has completed its analysis. If two functions are mutually recursive, then they must be defined jointly using the `and` keyword.

An example of a recursive function that counts from 1 to 10 similarly to Listing 8.5 is given in Listing 13.2.

Listing 13.2 `countRecursive.fsx`:
Counting to 10 using recursion.

```

1  let rec prt a b =
2      if a > b then
3          printf "\n"
4      else
5          printf "%d " a
6          prt (a + 1) b
7
8  prt 1 10

```

```

1  $ fsharp --nologo countRecursive.fsx && mono countRecursive.exe
2  1 2 3 4 5 6 7 8 9 10

```

Here the `prt` function calls itself repeatedly, such that the first call is `prt 1 10`, which calls `prt 2 10`, and so on until the last call `prt 11 10`. Each time `prt` is called, new bindings named `a` and `b` are made to new values. This is illustrated in Figure 13.2. The old values are no longer accessible as indicated by subscript in the figure. E.g., in `prt3` the scope has access to `a3` but not `a2` and `a1`. Thus, in this program, process is similar to a `for` loop, where the counter is `a` and in each loop its value is reduced.

The structure of the function is typical for recursive functions. They very often follow the following pattern.

Listing 13.3 Recursive functions consists of a stopping criterium, a stopping expression, and a recursive step.

```

1  let rec f a =
2      if <stopping condition>
3      then <stopping step>
4      else <recursion step>

```

The `match` – `with` are also very common conditional structures. In Listing 13.2 `a > b` is the *stopping condition*, `printfn "\n"` is *stopping step*, and `printfn "%d " a; prt (a + 1) b` is the *recursion step*.

- `match`
- `with`
- stopping condition
- stopping step
- recursion step

13.2 The call stack and tail recursion

Fibonacci's sequence of numbers is a recursive sequence of numbers with relations to the Golden ratio and structures in biology. Fibonacci's sequence is the sequence of numbers 1, 1, 2, 3, 5, 8, 13, ... The sequence starts with 1, 1 and the next number is recursively given as the sum of the two previous. A direct implementation of this is given in Listing 8.7.

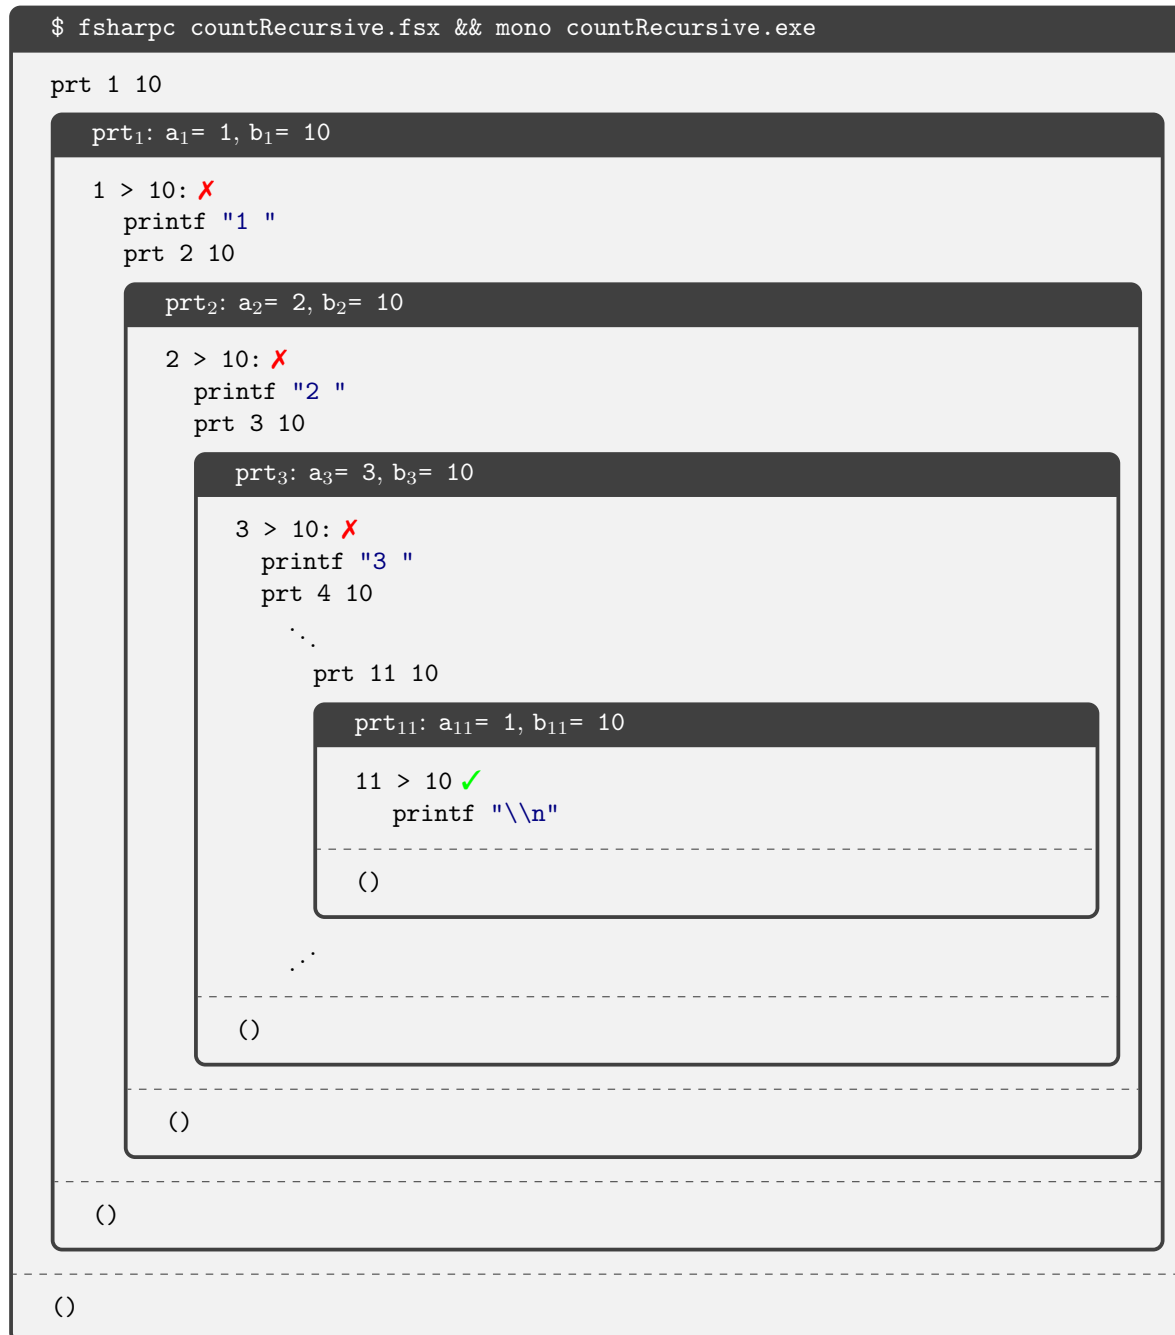


Figure 13.2: Illustration of the recursion used to write the sequence “1 2 3 ... 10” in line 8 in Listing 13.2. Each frame corresponds to a call to `prt`, where new values overshadow old. All return `unit`.

Listing 13.4 fibRecursive.fsx:
The n 'th Fibonacci number using recursive.

```

1  let rec fib n =
2      if n < 1 then
3          0
4      elif n = 1 then
5          1
6      else
7          fib (n - 1) + fib (n - 2)
8
9  for i = 0 to 10 do
10     printfn "fib(%d) = %d" i (fib i)

```

```

1  $ fsharp --nologo fibRecursive.fsx && mono fibRecursive.exe
2  fib(0) = 0
3  fib(1) = 1
4  fib(2) = 1
5  fib(3) = 2
6  fib(4) = 3
7  fib(5) = 5
8  fib(6) = 8
9  fib(7) = 13
10 fib(8) = 21
11 fib(9) = 34
12 fib(10) = 55

```

Here we extended the sequence to 0, 1, 1, 2, 3, 5, ... and starting sequence 0, 1 allowing us to define all $\text{fib}(n) = 0, n < 1$. Thus, our function is defined for all integers, and for the irrelevant negative arguments, it fails gracefully by returning 0. This is a general advice: **make functions that fail gracefully**. Advice

A visualization of the calls and the scopes created by `fibRecursive` is shown in Figure 13.3. The figure illustrates that each recursive step results in two calls to the function, thus creating two new scopes. And it gets worse. Figure 13.4 illustrates the tree of calls for `fib 5`. Thus a call to the function `fib` generates a tree of calls that is five levels deep and has `fib(5)` number of nodes. In general for the program in Listing 13.4, a call to `fib(n)` produces a tree with $\text{fib}(n) \leq c\alpha^n$ calls to the function for some positive constant c and $\alpha \geq \frac{1+\sqrt{5}}{2} \sim 1.6^1$. Each call takes time and requires memory, and we have thus created a slow and somewhat memory intensive function. This is a hugely ineffective implementation of calculating entries into Fibonacci's sequence, since many of the calls are identical. E.g., in Figure 13.4 `fib 1` is called five times. Before we examine a faster algorithm, we first need to discuss how F# executes function calls.

When a function is called, then memory is dynamically allocated internally for the function on what is known as the *call stack*. Stacks are used for many things in programming, but typically the call stack is considered special since it is almost always implicitly part of any program execution. Hence, it is often just referred to as *The Stack*. When a function is called, a new *stack frame* is stacked (pushed) on the call stack including its arguments, local storage such as mutable values, and where execution should return to when the function is finished. When the function finishes, the stack frame is unstacked (popped) and in its stead, the return value of the function is stacked. This return value is then unstacked and used by the caller. After unstacking the return value, the call stack is identical to its state prior to the call. Figure 13.5 shows snapshots of the call stack, when calling `fib 5` in Listing 13.4. The call first stacks a frame onto the call stack with everything needed to execute the

- call stack
- The Stack
- stack frame

¹Jon: <https://math.stackexchange.com/questions/674533/prove-upper-bound-big-o-for-fibonnaccis-sequence>

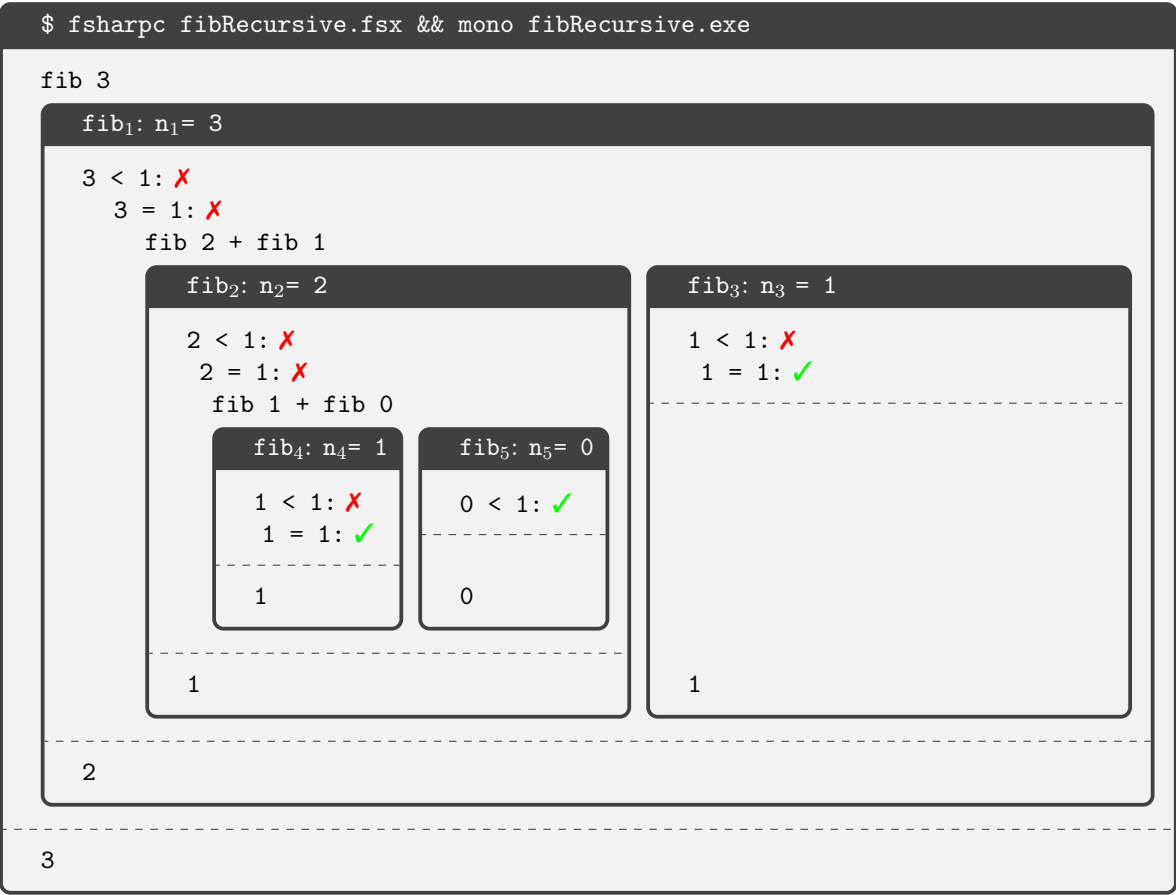


Figure 13.3: Illustration of the recursion used to write the sequence “1 2 3 ... 10” in line 8 in Listing 13.2. Each frame corresponds to a call to `fib`, where new values overshadow old.




Figure 13.4: The function calls involved in calling `fib 5`.



Figure 13.5: A call to `fib 5` in Listing 13.4 starts a sequence of function calls and stack frames on the call stack.

function body plus a reference to where the return to, when the execution is finished. Then the body of `fib` is executed, which includes calling `fib 4` and `fib 3` in turn. The call to `fib 4` stacks a frame onto the call stack, and its body is executed. Once execution is returned from the call to `fib 4`, the result of the function is on top of the stack. It is unstacked, saved and the call to `fib 3` is treated equally. When the end of `fib 5` is reached, its frame is unstacked, and its result is stacked. In this way, the call stack is returned to its original state except for the result of the function, and execution is returned to the point right after the original call to `fib 5`. Thus, for Listing 13.4 $\mathcal{O}(\alpha^n)$, $\alpha = \frac{1+\sqrt{5}}{2}$ stacking operations are performed for a call to `fib n`. The $\mathcal{O}(f(n))$ is the *Landau symbol* used to denote the order of a function, such that if $g(n) = \mathcal{O}(f(n))$ then there exists two real numbers $M > 0$ and a n_0 such that for all $n \geq n_0$, $|g(n)| \leq M|f(n)|$.² As indicated by the tree in Figure 13.4, the call tree is at most n high, which corresponds to a maximum of n additional stack frames as compared to the starting point.

The implementation of Fibonacci's sequence in Listing 13.4 can be improved to run faster and use less memory. One such algorithm is given in Listing 13.5

Listing 13.5 `fibRecursiveAlt.fsx`:

A fast, recursive implementation of Fibonacci's numbers. Compare with Listing 13.4.

```

1 let fib n =
2     let rec fibPair n pair =
3         if n < 2 then pair
4         else fibPair (n - 1) (snd pair, fst pair + snd pair)
5     if n < 1 then 0
6     elif n = 1 then 1
7     else fibPair n (0, 1) |> snd
8
9 printfn "fib(10) = %d" (fib 10)

```

```

1 $ fsharp --nologo fibRecursiveAlt.fsx && mono fibRecursiveAlt.exe
2 fib(10) = 55

```

Calculating the 45th Fibonacci number a MacBook Pro, with a 2.9 GHz Intel Core i5 using Listing 13.4 takes about 11.2s while using Listing 13.5 is about 224 times faster and only takes 0.050s. The reason is that `fib` in Listing 13.5 calculates every number in the sequence once and only once by processing the list recursively while maintaining the previous two values needed to calculate the next in the sequence. I.e., the function `helper` transforms the pair `(a,b)` to `(b,a+b)` such that, e.g., the 4th and 5th pair `(3,5)` is transformed into the 5th and the 6th pair `(5,8)` in the sequence. What complicates the algorithm is that besides the transformation, we must keep track of when to stop, which here is done using a counter variable, that is recursively reduced by 1 until our stopping criterium.

²Jon: Introduction of Landau notation needs to be moved earlier, since it used in Collections chapter.

Listing 13.5 also uses much less memory than Listing 13.4, since its recursive call is the last expression in the function, and since the return value of two recursive calls to `helper` is the same as the return value of the last. In fact, the return value of any number of recursive calls to `helper` is the return value of the last. This structure is called *tail-recursion*. Compilers can easily optimize the call stack usage for tail recursion, since when in this example `helper` calls itself, then its frame is no longer needed, and may be replaced by the new `helper` with the slight modification, that the return point should be to `fib` and not the end of the previous `helper`. Once the recursion reaches the stopping criteria, then instead of popping a long list of calls of `helper` frames, then there is only one, and the return value is equal to the return value of the last call and the return point is to `fib`. Thus, many stack frames in tail recursion are replaced by one. Hence, **prefer tail-recursion whenever possible**. · tail-recursion
Advice

13.3 Mutual recursive functions

Functions that recursively call each other are called *mutually recursive* functions. F# offers the `let - rec - and` notation for co-defining mutually recursive functions. As an example, consider the function `even : int -> bool`, which returns true if its argument is even and false otherwise, and the opposite function `odd : int -> bool`. A mutually recursive implementation of these functions can be developed from the following relations: `even 0 = true`, `odd 0 = false`, and for $n > 0$, `even n = odd (n-1)`, which implies that for $n > 0$, `odd n = even (n-1)`: · mutually recursive
· let
· rec
· and

Listing 13.6 mutuallyRecursive.fsx:
Using mutual recursion to implement even and odd functions.

```

1  let rec even x =
2      if x = 0 then true
3      else odd (x - 1)
4  and odd x =
5      if x = 0 then false
6      else even (x - 1);;
7
8  let w = 5;
9  printfn "%*s %*s %*s" w "i" w "even" w "odd"
10 for i = 1 to w do
11     printfn "%*d %*b %*b" w i w (even i) w (odd i)

```

```

1  $ fsharp --nologo mutuallyRecursive.fsx && mono mutuallyRecursive.exe
2      i  even  odd
3      1 false  true
4      2  true  false
5      3 false  true
6      4  true  false
7      5 false  true

```

Notice that in the lightweight notation the `and` must be on the same indentation level as the original `let`.

Without the `and` keyword, F# will issue a compile error at the definition of `even`. However, it is possible to implement mutual recursion by using functions as an argument, e.g.,

Listing 13.7 mutuallyRecursiveAlt.fsx:

Mutual recursion without the `and` keyword needs a helper function.

```

1  let rec evenHelper (notEven: int -> bool) x =
2      if x = 0 then true
3      else notEven (x - 1)
4
5  let rec odd x =
6      if x = 0 then false
7      else evenHelper odd (x - 1);;
8
9  let even x = evenHelper odd x
10
11 let w = 5;
12 printfn "%s %s %s" w "i" w "Even" w "Odd"
13 for i = 1 to w do
14     printfn "%d %b %b" w i w (even i) w (odd i)

```

```

1  $ fsharp --nologo mutuallyRecursiveAlt.fsx
2  $ mono mutuallyRecursiveAlt.exe
3      i  Even  Odd
4      1 false  true
5      2  true false
6      3 false  true
7      4  true false
8      5 false  true

```

But, Listing 13.6 is clearly to be preferred over Listing 13.7.

In the above, we used the `even` and `odd` function problems to demonstrate mutual recursion. There is, of course, a much simpler solution, which does not use recursion at all:

Listing 13.8 parity.fsx:

A better way to test for parity without recursion.

```

1  let even x = (x % 2 = 0)
2  let odd x = not (even x)

```

This is to be preferred anytime as the solution to the problem. ³

³Jon: Here it would be nice to have an *intermezzo*, giving examples of how to write a recursive program by thinking the problem has been solved.

14 | Programming with types

F# is a strongly typed language, meaning that types are known or inferred at compile time. In the previous chapters, we have used *primitive types* such as `float` and `bool`, function types, and compound types implicitly defined by tuples. These types are used for simple programming tasks, and everything that can be programmed can be accomplished using these types. However, larger programs are often easier to read and write when using more complicated type structures. In this chapter, we will discuss type abbreviations, enumerated types, discriminated unions, records, and structs. Class types are discussed in depth in Chapter 20.

· primitive types

14.1 Type abbreviations

F# allows for renaming of types, which is called *type abbreviation* or *type aliasing*. The syntax is

· type abbreviation

· type aliasing

Listing 14.1 Syntax for type abbreviation.

```
1 type <ident> = <aType>
```

where the identifier is a new name, and the type-name is an existing or a compound of existing types. E.g., in Listing 14.2 several type abbreviations are defined.

Listing 14.2 typeAbbreviation.fsx:

Defining 3 type abbreviations, two of which are compound types.

```
1 type size = int
2 type position = float * float
3 type person = string * int
4 type intToFloat = int -> float
5
6 let sz : size = 3
7 let pos : position = (2.5, -3.2)
8 let pers : person = ("Jon", 50)
9 let conv : intToFloat = fun a -> float a
10 printfn "%A, %A, %A, %A" sz pos pers (conv 2)
-----
1 $ fsharpc --nologo typeAbbreviation.fsx && mono typeAbbreviation.exe
2 3, (2.5, -3.2), ("Jon", 50), 2.0
```

Here we define the abbreviations `size`, `position`, `person`, and `intToFloat`, and later make bindings

enforcing the usage of these abbreviations.

Type abbreviations are used as short abbreviations of longer types, and they add semantic content to the program text, thus making programs shorter and easier to read. Type abbreviations allow the programmer to focus on the intended structure at a higher level by, e.g., programming in terms of a type `position` rather than `float * float`. Thus, they often result in programs with fewer errors. Type abbreviations also make maintenance easier. For instance, if we at a later stage decide that positions only can have integer values, then we only need to change the definition of the type abbreviation, not every place, a value of type `position` is used.

14.2 Enumerations

Enumerations or *enums* for short are types with named values. Names in enums are assigned to a subset of integer or char values. Their syntax is as follows:

- enumerations
- enums

Listing 14.3 Syntax for enumerations.

```
1 type <ident> =
2   [ | ] <ident> = <integerOrChar>
3   | <ident> = <integerOrChar>
4   | <ident> = <integerOrChar>
5   ...
```

An example of using enumerations is given in Listing 14.4.

Listing 14.4 enum.fsx:

An enum type acts as a typed alias to a set of integers or chars.

```
1 type medal =
2   Gold = 0
3   | Silver = 1
4   | Bronze = 2
5
6 let aMedal = medal.Gold
7 printfn "%A has value %d" aMedal (int aMedal)
```

```
1 $ fsharp --nologo enum.fsx && mono enum.exe
2 Gold has value 0
```

In the example, we define an enumerated type for medals, which allows us to work with the names rather than the values. Since the values most often are arbitrary, we can program using semantically meaningful names instead. Being able to refer to an underlying integer type allows us to interface with other – typically low-level – programs that require integers, and to perform arithmetic. E.g., for the medal example, we can typecast the enumerated types to integers and calculate an average medal harvest.

14.3 Discriminated Unions

A discriminated union is a union of a set of named cases. These cases can further be of specified types. The syntax for defining a discriminated union is as follows:

Listing 14.5 Syntax for type abbreviation.

```

1  [<attributes>]
2  type <ident> =
3      [| ]<ident> [of [<ident> :] <aType> [* [<ident> :] <aType> ...]]
4      | <ident> [of [<ident> :] <aType> [* [<ident> :] <aType> ...]]
5      ...

```

Discriminated unions are reference types, i.e., their content is stored on *The Heap*, see Section 6.8 for a discussion on reference types. Since they are immutable, there is no risk of side-effects. As reference types, when used as arguments to and returned from a function, then only a reference is passed. This is in contrast to value types, which transport a complete copy of the data structure. Discriminated unions are thus effective for large data structures. However, there is a slight overhead, since working with the content of reference types is indirect through their reference. Discriminated unions can also be represented as structures using the [<Struct>] attribute, in which case they are value types. See Section 14.5 for a discussion on structs.

An example just using the named cases but no further specification of types is given in Listing 14.6.

Listing 14.6 discriminatedUnions.fsx:

A discriminated union of medals. Compare with Listing 14.4.

```

1  type medal =
2      Gold
3      | Silver
4      | Bronze
5
6  let aMedal = medal.Gold
7  printfn "%A" aMedal

```

```

1  $ fsharp --nologo discriminatedUnions.fsx
2  $ mono discriminatedUnions.exe
3  Gold

```

Here we define a discriminated union as three named cases signifying three different types of medals. Comparing with the enumerated type in Listing 14.4, we see that the only difference is that the cases of the discriminated unions have no value. A commonly used discriminated union is the *option type*, see Section 18.2 for more detail.

Discriminated unions may also be used to store data. Where the names in enumerated types are aliases of single values, the names used in discriminated unions can hold any value specified at the time of creation. An example is given in Listing 14.7.




Figure 14.1: The tree with 3 leaves.

Listing 14.7 discriminatedUnionsOf.fsx:
A discriminated union using explicit subtypes.

```

1 type vector =
2     Vec2D of float * float
3     | Vec3D of x : float * y : float * z : float
4
5 let v2 = Vec2D (1.0, -1.2)
6 let v3 = Vec3D (x = 1.0, z = -1.2, y = 0.9)
7 printfn "%A and %A" v2 v3

```

```

1 $ fsharp -nologo discriminatedUnionsOf.fsx
2 $ mono discriminatedUnionsOf.exe
3 Vec2D (1.0,-1.2) and Vec3D (1.0,0.9,-1.2)

```

In this case, we define a discriminated union of two and three-dimensional vectors. Values of these types are created using their names followed by a tuple of their arguments. As can be seen, the arguments may be given field names, and if they are, then the names may be used when creating values of this type. As also demonstrated, the field names can be used to specify the field values in arbitrary order. However, values for all fields must be given.

Discriminated unions can be defined recursively. This feature is demonstrated in Listing 14.8.

Listing 14.8 discriminatedUnionTree.fsx:
A discriminated union modelling binary trees.

```

1 type Tree =
2     Leaf of int
3     | Node of Tree * Tree
4
5 let one = Leaf 1
6 let two = Leaf 2
7 let three = Leaf 3
8 let tree = Node (Node (one, two), three)
9 printfn "%A" tree

```

```

1 $ fsharp -nologo discriminatedUnionTree.fsx
2 $ mono discriminatedUnionTree.exe
3 Node (Node (Leaf 1,Leaf 2),Leaf 3)

```

In this example we define a tree as depicted in Figure 14.1.

Pattern matching must be used in order to define functions on values of a discriminated union. E.g.,

in Listing 14.9 we define a function that traverses a tree and prints the content of the nodes.¹

Listing 14.9 `discriminatedUnionPatternMatching.fsx`:
A discriminated union modelling binary trees.

```

1  type Tree = Leaf of int | Node of Tree * Tree
2  let rec traverse (t : Tree) : string =
3      match t with
4          Leaf(v) -> string v
5          | Node(left, right) -> (traverse left) + ", " + (traverse right)
6
7  let tree = Node (Node (Leaf 1, Leaf 2), Leaf 3)
8  printfn "%A: %s" tree (traverse tree)

```

```

1  $ fsharp --nologo discriminatedUnionPatternMatching.fsx
2  $ mono discriminatedUnionPatternMatching.exe
3  Node (Node (Leaf 1,Leaf 2),Leaf 3): 1, 2, 3

```

Discriminated unions are very powerful and can often be used instead of class hierarchies. Class hierarchies are discussed in Section 21.1.

14.4 Records

A record is a compound of named values, and a record type is defined as follows:

Listing 14.10 Syntax for defining record types.

```

1  [ <attributes> ]
2  type <ident> = {
3      [ mutable ] <label1> : <type1>
4      [ mutable ] <label2> : <type2>
5      ...
6  }

```

Records are collections of named variables and values of possibly different types. They are reference types, and thus their content is stored on *The Heap*, see Section 6.8 for a discussion on reference types. Records can also be *struct records* using the [`<Struct>`] attribute, in which case they are value types. See Section 14.5 for a discussion on structs. An example of using records is given in Listing 14.11. The values of individual members of a record are obtained using the “.” notation

· The Heap
· struct records
· .

¹Jon: Example uses pattern matching, which has yet to be introduced.

Listing 14.11 records.fsx:

A record is defined for holding information about a person.

```

1  type person = {
2      name : string
3      age : int
4      height : float
5  }
6
7  let author = {name = "Jon"; age = 50; height = 1.75}
8  printfn "%A\nname = %s" author author.name

```

```

1  $ fsharp --nologo records.fsx && mono records.exe
2  {name = "Jon";
3   age = 50;
4   height = 1.75;}
5  name = Jon

```

The examples illustrate a how record type is defined to store varied data about a person, and how a value is created by a record expression defining its field values.

If two record types are defined with the same label set, then the latter dominates the former, and the compiler will at a binding infer that later. This is demonstrated in Listing 14.12.

Listing 14.12 recordsDominance.fsx:

Redefined types dominates old record types, but earlier definitions are still accessible using explicit or implicit specification for bindings.

```

1  type person = { name : string; age : int; height : float }
2  type teacher = { name : string; age : int; height : float }
3
4  let lecturer = {name = "Jon"; age = 50; height = 1.75}
5  printfn "%A : %A" lecturer (lecturer.GetType())
6  let author : person = {name = "Jon"; age = 50; height = 1.75}
7  printfn "%A : %A" author (author.GetType())
8  let father = {person.name = "Jon"; age = 50; height = 1.75}
9  printfn "%A : %A" author (author.GetType())

```

```

1  $ fsharp --nologo recordsDominance.fsx && mono recordsDominance.exe
2  {name = "Jon";
3   age = 50;
4   height = 1.75;} : RecordsDominance+teacher
5  {name = "Jon";
6   age = 50;
7   height = 1.75;} : RecordsDominance+person
8  {name = "Jon";
9   age = 50;
10  height = 1.75;} : RecordsDominance+person

```

In the example, two identical record types are defined, and we use the built-in `GetType()` method to inspect the type of bindings. We see that `lecturer` is of `RecordsDominance+teacher` type, since `teacher` dominates the identical `author` type definition. However, we may enforce the `person` type by either specifying it for the name as in `let author : person = ...` or by fully or partially

specifying it in the record expression following the “=” sign. In both cases, they are therefore of `RecordsDominance+author` type. The built-in `GetType()` method is inherited from the base class for all types, see Chapter 20 for a discussion on classes and inheritance.

Note that when creating a record, you must supply a value to all fields, and you cannot refer to other fields of the same record, e.g., `{name = "Jon"; age = height * 3; height = 1.75}` is illegal.

Since records are per default reference types, binding creates aliases not copies. This matters for mutable members, in which case when copying, we must explicitly create a new record with the old data. Copying can be done either by using referencing to the individual members of the source or using the short-hand *with* notation. This is demonstrated in Listing 14.13.

· *with*

Listing 14.13 recordCopy.fsx:

Bindings are references. To copy and not make an alias, explicit copying must be performed.

```

1  type person = {
2      name : string;
3      mutable age : int;
4  }
5
6  let author = {name = "Jon"; age = 50}
7  let authorAlias = author
8  let authorCopy = {name = author.name; age = author.age}
9  let authorCopyAlt = {author with name = "Noj"}
10 author.age <- 51
11 printfn "author : %A" author
12 printfn "authorAlias : %A" authorAlias
13 printfn "authorCopy : %A" authorCopy
14 printfn "authorCopyAlt : %A" authorCopyAlt

```

```

1  $ fsharpc --nologo recordCopy.fsx && mono recordCopy.exe
2  author : {name = "Jon";
3      age = 51;}
4  authorAlias : {name = "Jon";
5      age = 51;}
6  authorCopy : {name = "Jon";
7      age = 50;}
8  authorCopyAlt : {name = "Noj";
9      age = 50;}

```

Here `age` is defined as a mutable value, and can be changed using the usual “<-” assignment operator. The example demonstrates two different ways to create records. Note that when the mutable value `author.age` is changed in line 10, then `authorAlias` also changes, since it is an alias of `author`, but neither `authorCopy` nor `authorCopyAlt` changes, since they are copies. As illustrated, copying using *with* allows for easy copying and partial updates of another record value.

14.5 Structures

Structures or *structs* for short have much in common with records. They specify a compound type with named fields, but they are value types, and they allow for some customization of what is to happen when a value of its type is created. Since they are value types, then they are best used for small amount of data. The syntax for defining struct types are:

· structures
· structs

Listing 14.14 Syntax for type abbreviation.

```

1  [ <attributes> ]
2  [<Struct>]
3  type <ident> =
4      val [ mutable ] <label1> : <type1>
5      val [ mutable ] <label2> : <type2>
6      ...
7      [new (<arg1>, <arg2>, ...) = {<label1> = <arg1>; <label1> = <arg2>; ...}]
8      [new (<arg1>, <arg2>, ...) = {<label1> = <arg1>; <label1> = <arg2>; ...}]
9      ...

```

The syntax makes use of the *val* and *new* keywords. Keyword *val* like *let* binds a name to a value, but unlike *let* the value is always the type's default value. The *new* keyword denotes the function used to fill values into the fields at time of creation. This function is called the *constructor*. No *let* nor *do* bindings are allowed in structure definitions. Fields are accessed using the “.” notation. An example is given in Listing 14.15.

· *val*
·
· *new*
· constructor
·

Listing 14.15 struct.fsx:
Defining a struct type and creating a value of it.

```

1  [<Struct>]
2  type position =
3      val x : float
4      val y : float
5      new (a : float, b : float) = {x = a; y = b}
6
7  let p = position (3.0, 4.2)
8  printfn "%A: x = %A, y = %A" p p.x p.y

```

```

1  $ fsharp --nologo struct.fsx && mono struct.exe
2  Struct+position: x = 3.0, y = 4.2

```

Structs are small versions of classes and allows, e.g., for overloading of the *new* constructor and for overriding of the inherited *ToString()* function. This is demonstrated in Listing 14.16.

· overload
· override
· ToString()

Listing 14.16 structOverloadNOverride.fsx:

Overloading the `new` constructor and overriding the default `ToString()` function.

```

1  [<Struct>]
2  type position =
3      val x : float
4      val y : float
5      new (a : float, b : float) = {x = a; y = b}
6      new (a : int, b : int) = {x = float a; y = float b}
7      override this.ToString() =
8          "(" + (string this.x) + ", " + (string this.y) + ")"
9
10 let pFloat = position (3.0, 4.2)
11 let pInt = position (3, 4)
12 printfn "%A and %A" pFloat pInt

```

```

1  $ fsharpc --nologo structOverloadNOverride.fsx
2  $ mono structOverloadNOverride.exe
3  (3, 4.2) and (3, 4)

```

We defer further discussion of these concepts to Chapter 20.

The use of structs are generally discouraged, and instead, it is recommended to use enums, records, and discriminated unions possibly with the `[<Struct>]` attribute for the last two in order to make them value types.

14.6 Variable types

An advanced topic in F# is *variable types*. There are three different versions of variable types in F#: *runtime resolved*, which has the syntax `'<ident>`, *anonymous*, which are written as `"_"`, and *statically resolved*, which have the syntax `~<ident>`. Variable types are particularly useful for functions that work for many types. An example of a generic function and its use is given in Listing 14.17.

Listing 14.17 variableType.fsx:

A function apply with runtime resolved types.

```

1  let apply (f : 'a -> 'a -> 'a) (x : 'a) (y : 'a) : 'a = f x y
2  let intPlus (x : int) (y : int) : int = x + y
3  let floatPlus (x : float) (y : float) : float = x + y
4
5  printfn "%A %A" (apply intPlus 1 2) (apply floatPlus 1.0 2.0)

```

```

1  $ fsharpc --nologo variableType.fsx && mono variableType.exe
2  3 3.0

```

- variable types
- runtime resolved variable type
- anonymous variable type
- `-`
- statically resolved variable type

In this example, the function `apply` has runtime resolved variable type, and it accepts three parameters `f`, `x`, and `y`. The function will work as long as the parameters for `f` is a function of two parameters of identical type, and `x` and `y` are values of the same type. Thus, in the `printfn` statement, we are able to use `apply` for both an integer and a float variant.

The example in Listing 14.17 illustrates a very complicated way to add two numbers. And the “+” operator works for both types out of the box, so why not something simpler like relying on the F# type inference system by not explicitly specifying types as attempted in Listing 14.18.

Listing 14.18 `variableTypeError.fsx`:

Even though the “+” operator is defined for both integers and floats, the type inference is static and infers `plus : int -> int`.

```

1  let plus x y = x + y
2
3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)

```

```

1  $ fsharpc --nologo variableTypeError.fsx && mono variableTypeError.exe
2
3  variableTypeError.fsx(3,34): error FS0001: This expression was expected
   to have type
4      'int'
5  but here has type
6      'float'
7
8  variableTypeError.fsx(3,38): error FS0001: This expression was expected
   to have type
9      'int'
10 but here has type
11      'float'

```

Unfortunately, the example fails to compile, since the type inference is performed at compile time, and by `plus 1 2`, it is inferred that `plus : int -> int`. Hence, calling `plus 1.0 2.0` is a type error. Function bindings allow for the use of the `inline` keyword, and adding this successfully reuses the `inline` definition of `plus` for both types as shown in Listing 14.19.

Listing 14.19 `variableTypeInline.fsx`:

The keyword `inline` forces static and independent inference each place the function is used. Compare to the error case in Listing 14.18.

```

1  let inline plus x y = x + y
2
3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)

```

```

1  $ fsharpc --nologo variableTypeInline.fsx && mono variableTypeInline.exe
2  3 3.0

```

In the example, adding the `inline` does two things: Firstly, it copies the code to be performed to each place, the function is used, and secondly, it forces statically resolved variable type checking independently in each place. The type annotations inferred as a result of the `inline`-keyword may be written explicitly as shown in Listing 14.20.

Listing 14.20 `compiletimeVariableType.fsx`:

Explicitly spelling out of the statically resolved type variables from Listing 14.18.

```

1  let inline plus (x : ^a) (y : ^a) : ^a when ^a : (static member ( + ) :
    ^a * ^a -> ^a) = x + y
2
3  printfn "%A %A" (plus 1 2) (plus 1.0 2.0)
-----
1  $ fsharp -nologo compiletimeVariableType.fsx
2  $ mono compiletimeVariableType.exe
3  3 3.0

```

The example in Listing 14.20 demonstrates the statically resolved variable type syntax, `^<ident>`, as well as the use of *type constraints* using the keyword `when`. Type constraints have a rich syntax, but will not be discussed further in this book.² In the example, the type constraint `when ^a : (static member (+) : ^a * ^a -> ^a)` is given using the object oriented properties of the type variable `^a`, meaning that the only acceptable type values are those, which have a member function `(+)` taking a tuple and giving a value all of identical type, and which where the type can be inferred at compile time. See Chapter 20 for details on member functions.

The `inline` construction is useful when generating generic functions and still profiting from static type checking. However, explicit copying of functions is often something better left to the compiler to optimize over. An alternative seems to be using runtime resolved variable types with the `^<ident>` syntax. Unfortunately, this is not possible in case of most operators, since they have been defined in the `FSharp.Core` namespace to be statically resolved variable type. E.g., the “+” operator has type `(+) : ^T1 -> ^T2 -> ^T3` (requires `^T1 with static member (+)` and `^T2 with static member (+)`).

²Jon: Should I extend on type constraints? Perhaps it is better left for a specialize chapter on generic functions.

15 | Pattern matching

Pattern matching is used to transform values and variables into a syntactical structure. The simplest example is value-bindings. The `let`-keyword was introduced in Section 6.1, its extension with pattern matching is given as,

Listing 15.1 Syntax for `let`-expressions with pattern matching.

```
1  [[<Literal>]]
2  let [mutable] <pat> [: <returnType>] = <bodyExpr> [in <expr>]
```

A typical use of this is to extract elements of tuples as demonstrated in Listing 15.2.

Listing 15.2 `letPattern.fsx`:

Patterns in `let` expressions may be used to extract elements of tuples.

```
1  let a = (3,4)
2  let (x,y) = a
3  let (alsoX,_) = a
4  printfn "%A: %d %d %d" a x y alsoX

-----

1  $ fsharp --nologo letPattern.fsx && mono letPattern.exe
2  (3, 4): 3 4 3
```

Here we extract the elements of a pair twice. First by binding to `x` and `y`, and second by binding to `alsoX` while using the wildcard pattern to ignore the second element. Thus, again the wildcard pattern in value-bindings is used to underline a disregarded value.

Another common use of patterns is as alternative to `if – then – else` expressions particularly when parsing input for a function. Consider the example in Listing 15.3.

Listing 15.3 switch.fsx:

Using `if – then – else` to print discriminated unions.

```

1 type Medal = Gold | Silver | Bronze
2 let statement (m : Medal) : string =
3     if m = Gold then "You won"
4     elif m = Silver then "You almost won"
5     else "Maybe you can win next time"
6
7 let m = Silver
8 printfn "%A : %s" m (statement m)

```

```

1 $ fsharp --nologo switch.fsx && mono switch.exe
2 Silver : You almost won

```

In the example, is a discriminated union and a function defined. The function converts each case to a supporting statement using an `if`-expression. The same can be done with the `match – with` expression and patterns as is demonstrated in Listing 15.4.

Listing 15.4 switchPattern.fsx:

Using `match – with` to print discriminated unions.

```

1 type Medal = Gold | Silver | Bronze
2 let statement (m : Medal) : string =
3     match m with
4         Gold -> "You won"
5         | Silver -> "You almost won"
6         | _ -> "Maybe you can win next time"
7
8 let m = Silver
9 printfn "%A : %s" m (statement m)

```

```

1 $ fsharp --nologo switchPattern.fsx && mono switchPattern.exe
2 Silver : You almost won

```

Here we used a pattern for the discriminated union cases and a wildcard pattern as default. The lightweight syntax for `match`-expressions is,

Listing 15.5 Syntax for `match`-expressions.

```

1 match <inputExpr> with
2     [| ]<pat> [when <guardExpr>] -> <caseExpr>
3     | <pat> [when <guardExpr>] -> <caseExpr>
4     | <pat> [when <guardExpr>] -> <caseExpr>
5     ...

```

where `<inputExpr>` is the *input pattern* to find matches of, `<pat>` is a pattern to match with, `<guardExpr>` is an optional guard expression, and `<caseExpr>` is the resulting expression. Each set starting with `<pat>` is called a case. In lightweight syntax, the indentation must be equal to or higher than the indentation of `match`. All cases must return a value of the same type, and F# report an error, when not the complete domain of the input pattern is covered by cases in `match`-expressions.

Patterns are also used in a version of *for*-loop expressions, and its lightweight syntax is given as, *· for*

Listing 15.6 Syntax for *for*-expressions with pattern matching.

```
1 for <pat> in <sourceExpr> do
2   <bodyExpr>
```

Typically, *<sourceExpr>* is a list or an array. An example is given in Listing 15.7.

Listing 15.7 forPattern.fsx:
Patterns may be used in *for*-loops.

```
1 for (_,y) in [(1,3); (2,1)] do
2   printfn "%d" y
```

```
1 $ fsharpc --nologo forPattern.fsx && mono forPattern.exe
2 3
3 1
```

The wildcard pattern is used to disregard the first element in a pair while iterating over the complete list. It is good practice to **use wildcard patterns to emphasize unused values**.

Advice

The final expression involving patterns to be discussed is *anonymous functions*. Patterns for anonymous functions have the syntax, *· anonymous functions*

Listing 15.8 Syntax for anonymous functions with pattern matching.

```
1 fun <pat> [<pat> ...] -> <bodyExpr>
```

This is an extension of the syntax discussed in Section 6.2. A typical use for patterns in *fun*-expressions is shown in Listing 15.9. *· fun*

Listing 15.9 funPattern.fsx:
Patterns may be used in *fun*-expressions.

```
1 let f = fun _ -> "hello"
2 printfn "%s" (f 3)
```

```
1 $ fsharpc --nologo funPattern.fsx && mono funPattern.exe
2 hello
```

Here we use an anonymous function expression and bind it to *f*. The expression has one argument of any type, which it ignores using the wildcard pattern. Some limitations apply to the patterns allowed in *fun*-expressions.¹ The wildcard pattern in *fun*-expressions are often used for *mockup functions*, where the code requires the said function, but its content has yet to be decided. Thus, mockup functions can be used as loose place-holders, while experimenting with program design. *· mockup functions*

Patterns are also used in exceptions to be discussed in Section 18.1, and in conjunction with the

¹Jon: Remove or elaborate.

`function`-keyword, a keyword we discourage in this book. We will now demonstrate a list of important patterns in F#.

15.1 Wildcard pattern

A *wildcard pattern* is denoted “`_`” and matches anything, see e.g., Listing 15.10.

· wildcard pattern
· `_`

Listing 15.10 `wildcardPattern.fsx`:
Constant patterns matches to constants.

```
1 let whatever (x : int) : string =
2     match x with
3         _ -> "If you say so"
4
5 printfn "%s" (whatever 42)
```

```
1 $ fsharp --nologo wildcardPattern.fsx && mono wildcardPattern.exe
2 If you say so
```

In this example, anything matches the wildcard pattern, so all cases are covered and the function always returns the same sentence. This is rarely a useful structure on its own since this could be replaced by a value binding or by a function ignoring its input. However, wildcard patterns are extremely useful, since they act as the final `else` in `if`-expressions.

15.2 Constant and literal patterns

A *constant pattern* matches any input pattern with constants, see e.g., Listing 15.11.

· constant pattern

Listing 15.11 `constPattern.fsx`:
Constant patterns matches to constants.

```
1 type Medal = Gold | Silver | Bronze
2 let intToMedal (x : int) : Medal =
3     match x with
4         0 -> Gold
5         | 1 -> Silver
6         | _ -> Bronze
7
8 printfn "%A" (intToMedal 0)
```

```
1 $ fsharp --nologo constPattern.fsx && mono constPattern.exe
2 Gold
```

In this example, the input pattern is queried for a match with 0 and 1 or the wildcard pattern. Any simple literal type constants may be used in the constant pattern such as 8, 23y, 1010u, 1.2, "hello world", 'c', and , false. Here we also use the wildcard pattern. Notice, matching is performed in a lazy manner and stops for the first matching case from the top. Thus, although the wildcard pattern

matches everything, its case expression is only executed if none of the previous patterns matches the input.

Constants can also be pre-bound by the [`<Literal>`] attribute for value-bindings. This is demonstrated in Listing 15.12.

Listing 15.12 `literalPattern.fsx`:

A variant of constant patterns are literal patterns.

```

1  [<Literal>]
2  let TheAnswer = 42
3  let whatIsTheQuestion (x : int) : string =
4      match x with
5          TheAnswer -> "We will need to build a bigger machine..."
6          | _ -> "Don't know that either"
7
8  printfn "%A" (whatIsTheQuestion 42)

```

```

1  $ fsharpc --nologo literalPattern.fsx && mono literalPattern.exe
2  "We will need to build a bigger machine..."

```

The attributed is used to identify the value-binding `TheAnswer` to be used as if it were a simple literal type. Literal patterns must be either uppercase or module prefixed identifiers.

15.3 Variable patterns

A *variable pattern* is a single lower-case letter identifier. Variable pattern identifiers are assigned the value and type of the input pattern. Combinations of constant and variable patterns are also allowed together with records and arrays. This is demonstrated in Listing 15.13.

Listing 15.13 `variablePattern.fsx`:

Variable patterns are useful for extracting and naming fields etc.

```

1  let (name, age) = ("Jon", 50)
2  let getAgeString (age : int) : string =
3      match age with
4          0 -> "newborn"
5          | 1 -> "1 year old"
6          | n -> (string n) + " years old"
7
8  printfn "%s is %s" name (getAgeString age)

```

```

1  $ fsharpc --nologo variablePattern.fsx && mono variablePattern.exe
2  Jon is 50 years old

```

In this example, the use of the value identifier `n` has the function of a named wildcard pattern. Hence, the case could as well have been `| _ -> (string age) + "years old"`, since `age` is already defined in this scope. However, variable patterns syntactically act as an argument to an anonymous function and thus act to isolate the dependencies. They are also very useful together with guards, see Section 15.4.

15.4 Guards

A *guard* is a pattern used together with `match`-expressions including the `when`-keyword, as shown in Listing 15.5.

Listing 15.14 guardPattern.fsx:

Guard expressions can be used with other patterns to restrict matches.

```
1 let getAgeString (age : int) : string =
2     match age with
3         n when n < 1 -> "infant"
4         | n when n < 13 -> "child"
5         | n when n < 20 -> "teen"
6         | _ -> "adult"
7
8 printfn "A person aged %d is a/an %s" 50 (getAgeString 50)
```

```
1 $ fsharp --nologo guardPattern.fsx && mono guardPattern.exe
2 A person aged 50 is a/an adult
```

Here guards are used to iteratively carve out subset of integers to assign different strings to each set. The guard expression in `<pat> when <guardExpr> -> <caseExpr>` is any expression evaluating to a Boolean, and the case expression is only executed for the matching case.

15.5 List patterns

Lists have a concatenation pattern associated with them. The `“::”` cons-operator is used to match the head and the rest of a list, and `“[]”` is used to match an empty list, which is also sometimes called the nil-case. This is very useful, when recursively processing lists as shown in Listing 15.15

Listing 15.15 listPattern.fsx:

Recursively parsing a list using list patterns.

```
1 let rec sumList (lst : int list) : int =
2     match lst with
3         n :: rest -> n + (sumList rest)
4         | [] -> 0
5
6 let rec sumThree (lst : int list) : int =
7     match lst with
8         [a; b; c] -> a + b + c
9         | _ -> sumList lst
10
11 let aList = [1; 2; 3]
12 printfn "The sum of %A is %d, %d" aList (sumList aList) (sumThree aList)
```

```
1 $ fsharp --nologo listPattern.fsx && mono listPattern.exe
2 The sum of [1; 2; 3] is 6, 6
```

In the example, the function `sumList` uses the `cons` operator to match the head of the list with `n` and the tail with `rest`. The pattern `n :: tail` also matches `3 :: []`, and in that case `tail` would be assigned the value `[]`. When `lst` is empty, then it matches with `[]`. List patterns can also be matched explicitly named elements as demonstrated in the `sumThree` function. The elements to be matched can be any mix of constants and variables.

15.6 Array, record, and discriminated union patterns

Array, *record*, and *discriminated union* patterns are direct extensions on constant, variable, and wildcard patterns. Listing 15.16 gives examples of array patterns.

- array pattern
- record pattern
- discriminated union patterns

Listing 15.16 `arrayPattern.fsx`:

Using variable patterns to match on size and content of arrays.

```

1  let arrayToString (x : int []) : string =
2      match x with
3          [|1;_:_|] -> "3 elements, first of is a one"
4          | [|x;1;_]| -> "3 elements, first is " + (string x) + "Second is one"
5          | x -> "A general array"
6
7  printfn "%s" (arrayToString [|1; 1; 1|])
8  printfn "%s" (arrayToString [|3; 1; 1|])
9  printfn "%s" (arrayToString [|1|])

```

```

1  $ fsharp --nologo arrayPattern.fsx && mono arrayPattern.exe
2  3 elements, first of is a one
3  3 elements, first is 3Second is one
4  A general array

```

In the function `arrayToString`, the first case matches arrays of 3 elements, where the first is the integer 1, the second case matches arrays of 3 elements, where the second is a 1 and names the first `x`, and the final case matches all arrays and works as a default match case. As demonstrated, the cases are treated from first to last, and only the expression of the first case that matches is executed.

For record pattern, we use the field names to specify matching criteria. This is demonstrated in Listing 15.17.

Listing 15.17 recordPattern.fsx:

Variable patterns for records to match on field values.

```

1  type Address = {street : string; zip : int; country : string}
2  let contact : Address = {
3      street = "Universitetsparken 1";
4      zip = 2100;
5      country = "Denmark"}
6  let getZip (adr : Address) : int =
7      match adr with
8      {street = _; zip = z; country = _} -> z
9
10 printfn "The zip-code is: %d" (getZip contact)

```

```

1  $ fsharp --nologo recordPattern.fsx && mono recordPattern.exe
2  The zip-code is: 2100

```

Here, the record type `Address` is created, and in the function `getZip`, a variable pattern `z` is created for naming zip values, and the remaining fields are ignored. Since the fields are named, the pattern match needs not mention the ignored fields, and the example match is equivalent to `{zip = z} -> z`. The curly brackets are required for record patterns.

Discriminated union patterns are similar. For discriminated unions with arguments, the arguments can be matched as constants, variables, or wildcards. A demonstration is given in Listing 15.18.

Listing 15.18 unionPattern.fsx:

Matching on discriminated union types.

```

1  type vector =
2      Vec2D of float * float
3      | Vec3D of float * float * float
4
5  let project (vec : vector) : vector =
6      match vec with
7      Vec3D (a, b, _) -> Vec2D (a, b)
8      | v -> v
9
10 let v = Vec3D (1.0, -1.2, 0.9)
11 printfn "%A -> %A" v (project v)

```

```

1  $ fsharp --nologo unionPattern.fsx && mono unionPattern.exe
2  Vec3D (1.0,-1.2,0.9) -> Vec2D (1.0,-1.2)

```

In the `project`-function, three-dimensional vectors are projected to two dimensions by removing the third element. Two-dimensional vectors are unchanged. The example uses the wildcard pattern to emphasize, that the third element of three-dimensional vectors is ignored. Named arguments can also be matched, in which case “;” is used to delimit the fields in the match instead of “,”.

15.7 Disjunctive and conjunctive patterns

Patterns may be combined disjunctively using the “/” lexeme and conjunctively using the “&” lexeme. *Disjunctive patterns* combine as fall-through as illustrated in Listing 15.19.

· |
· &
· disjunctive pattern

Listing 15.19 `disjunctivePattern.fsx`:

Patterns can be combined logically as ‘or’ syntax structures.

```
1 let vowel (c : char) : bool =
2     match c with
3         'a' | 'e' | 'i' | 'o' | 'u' | 'y' -> true
4         | _ -> false
5
6 String.iter (fun c -> printf "%A " (vowel c)) "abcdefg"
-----
1 $ fsharp --nologo disjunctivePattern.fsx && mono disjunctivePattern.exe
2 true false false false true false false
```

Here one or more cases must match for the final case expression, and thus, any vowel results in the same case expression `true`. All else is matched with the wildcard pattern.

For *conjunctive patterns* all patterns must match, which is illustrated in Listing 15.20.

· conjunctive patterns

Listing 15.20 `conjunctivePattern.fsx`:

Patterns can be combined logically as ‘or’ syntax structures.

```
1 let is11 (v : int * int) : bool =
2     match v with
3         (1,_) & (_,1) -> true
4         | _ -> false
5
6 printfn "%A" (List.map is11 [(0,0); (0,1); (1,0); (1,1)])
-----
1 $ fsharp --nologo conjunctivePattern.fsx && mono conjunctivePattern.exe
2 [false; false; false; true]
```

In this case, we separately check the elements of a pair for the constant value 1 and return true only when both elements are 1. In many cases, conjunctive patterns can be replaced by more elegant matches, e.g., using tuples, and in the above example a single case `(1,1) -> true` would have been simpler. Nevertheless, conjunctive patterns are used together with active patterns, to be discussed below.

15.8 Active Pattern

The concept of patterns is extendable to functions. Such functions are called *active patterns*, and active patterns comes in two flavors: regular and option types. The active pattern cases are constructed as function bindings, but using a special notation. They all take the pattern input as last argument, and may take further preceding arguments. The syntax for active patterns is one of,

· active patterns

Listing 15.21 Syntax for binding active patterns to expressions.

```

1  let (|<caseName>|[_| ]) [ <arg> [<arg> ... ] ] <inputArgument> = <expr>
2  let (|<caseName>|<caseName>|...|<caseName>|) <inputArgumentn> = <expr>

```

When using the (`|<caseName>|[_|]`) variants, then the active pattern function must return an option type. The multi-case variant (`|<caseName>|<caseName>|...|<caseName>|`) must return a `Fsharp.Core.Choice` type. All other variants can return any type. There are no restrictions on arguments `<arg>`, and `<inputArgumentn>` is the input pattern to be matched. Notice in particular that the multi-case variant only takes one argument and cannot be combined with the option-type syntax. Below we will demonstrate how the various patterns are used by example.

The single case, (`|<caseName>|`) matches all, and is useful for extracting information from complex types, as demonstrated in Listing 15.22.

Listing 15.22 activePattern.fsx:
Single case active pattern for deconstructing complex types.

```

1  type vec = {x : float; y : float}
2  let (|Cartesian|) (v : vec) = (v.x, v.y)
3  let (|Polar|) (v : vec) = (sqrt(v.x*v.x + v.y * v.y), atan2 v.y v.x)
4  let printCartesian (p : vec) : unit =
5      match p with
6          Cartesian (x, y) -> printfn "%A:\n Cartesian (%A, %A)" p x y
7  let printPolar (p : vec) : unit =
8      match p with
9          Polar (a, d) -> printfn "%A:\n Cartesian (%A, %A)" p a d
10
11  let v = {x = 2.0; y = 3.0}
12  printCartesian v
13  printPolar v

```

```

1  $ fsharpc --nologo activePattern.fsx && mono activePattern.exe
2  {x = 2.0;
3    y = 3.0;}:
4    Cartesian (2.0, 3.0)
5  {x = 2.0;
6    y = 3.0;}:
7    Cartesian (3.605551275, 0.9827937232)

```

Here we define a record to represent two-dimensional vectors and two different single case active patterns. Note that in the binding of the active pattern functions in line 2 and 3, the argument is the input expression `match <inputExpr> with ...`, see Listing 15.5. However, the argument for the cases in line 6 and 9 are names bound to the output of the active pattern function.

Both `Cartesian` and `Polar` matches a vector record, but they dismantle the contents differently. For an alternative solution using Class types, see Section 20.1.

More complicated behavior is obtainable by supplying additional arguments to the single case. This is demonstrated in Listing 15.23.

Listing 15.23 activeArgumentsPattern.fsx:

All but the multi-case active pattern may take additional arguments.

```

1 type vec = {x : float; y : float}
2 let (|Polar|) (o : vec) (v : vec) =
3     let x = v.x - o.x
4     let y = v.y - o.y
5     (sqrt(x*x + y * y), atan2 y x)
6 let printPolar (o : vec) (p : vec) : unit =
7     match p with
8     | Polar o (a, d) -> printfn "%A:\n Cartesian (%A, %A)" p a d
9
10 let v = {x = 2.0; y = 3.0}
11 let offset = {x = 1.0; y = 1.0}
12 printPolar offset v

```

```

1 $ fsharp -nologo activeArgumentsPattern.fsx
2 $ mono activeArgumentsPattern.exe
3 {x = 2.0;
4  y = 3.0;}:
5 Cartesian (2.236067977, 1.107148718)

```

Here we supply an offset, which should be subtracted prior to calculating lengths and angles. Notice in line 8, that the argument is given prior to the result binding.

Active pattern functions return option types are called *partial pattern functions*. The option type allows for specifying mismatches as illustrated in Listing 15.24. · partial pattern functions

Listing 15.24 activeOptionPattern.fsx:

Option type active patterns mismatches on None results.

```

1 let (|Div|_|) (e,d) = if d <> 0.0 then Some (e/d) else None
2
3 let safeDiv (p : float * float) =
4     match p with
5     | (0.0, 0.0) -> printfn "Div %A = undefined" p
6     | Div res -> printfn "Div %A = %A" p res
7     | _ -> printfn "Div %A = infinity" p
8
9 List.iter safeDiv [(1.0,1.0); (0.0,1.0); (1.0,0.0); (0.0,0.0)]

```

```

1 $ fsharp -nologo activeOptionPattern.fsx
2 $ mono activeOptionPattern.exe
3 Div (1.0, 1.0) = 1.0
4 Div (0.0, 1.0) = 0.0
5 Div (1.0, 0.0) = infinity
6 Div (0.0, 0.0) = undefined

```

In the example, we use the (`|<caseName>|_|`) variant to indicate, that the active pattern returns an option type. Nevertheless, the result binding `res` in line 6 uses the underlying value of `Some`. And in contrast to the two previous examples of single case patterns, the value `None` results in a mismatch. Thus in this case, if the denominator is 0.0, then `Div res` does not match but the wildcard pattern does.

Multicase active patterns work similarly to discriminated unions without arguments.² An example is given in Listing 15.25. · multicase active patterns

Listing 15.25 `activeMultiCasePattern.fsx`:

Multi-case active patterns have a syntactical structure similar to discriminated unions.

```

1  let (|Gold|Silver|Bronze|) inp =
2      if inp = 0 then Gold
3      elif inp = 1 then Silver
4      else Bronze
5
6  let intToMedal (i : int) =
7      match i with
8      | Gold -> printfn "%d: Its gold!" i
9      | Silver -> printfn "%d: Its silver." i
10     | Bronze -> printfn "%d: Its no more than bronze." i
11
12  List.iter intToMedal [0..3]

```

```

1  $ fsharpc --nologo activeMultiCasePattern.fsx
2  $ mono activeMultiCasePattern.exe
3  0: Its gold!
4  1: Its silver.
5  2: Its no more than bronze.
6  3: Its no more than bronze.

```

In this example, we define three cases in line 1. The result of the active pattern function must be one of these cases. For the `match`-expression, the match is based on the output of the active pattern function, hence in line 8, the case expression is executed, when the result of applying the active pattern function to the input expression `i` is `Gold`. In this case, a solution based on discriminated unions would probably be clearer.

15.9 Static and dynamic type pattern

Input patterns can also be matched on type. For *static type matching* the matching is performed at compile time, and indicated using the “`:`” lexeme followed by the type name to be matched. Static type · static type pattern
· :
matching is further used as input to the type inference performed at compile time to infer non-specified types as illustrated in Listing 15.26.

²Jon: This maybe too advanced for this book.

Listing 15.26 staticTypePattern.fsx:

Static matching on type binds the type of other values by type inference.

```

1  let rec sum lst =
2      match lst with
3          (n : int) :: rest -> n + (sum rest)
4          | [] -> 0
5
6  printfn "The sum is %d" (sum [0..3])

```

```

1  $ fsharpc --nologo staticTypePattern.fsx && mono staticTypePattern.exe
2  The sum is 6

```

Here the head of the list `n` in the list pattern is explicitly matched as an integer, and the type inference system thus concludes that `lst` must be a list of integers.

In contrast to static type matching, *dynamic type matching* is performed at runtime, and indicated using the “`:?`” lexeme followed by a type name. Dynamic type patterns allow for matching generic values at runtime. This is an advanced topic, which is included here for completeness. An example is given in Listing 15.27.

· dynamic type pattern
· `:?`

Listing 15.27 dynamicTypePattern.fsx:

Static matching on type binds the type of other values by type inference.

```

1  let isString (x : obj) : bool =
2      match x with
3          :? string -> true
4          | _ -> false
5
6  let a = "hej"
7  printfn "Is %A a string? %b" a (isString a)
8  let b = 3
9  printfn "Is %A a string? %b" b (isString b)

```

```

1  $ fsharpc --nologo dynamicTypePattern.fsx && mono dynamicTypePattern.exe
2  Is "hej" a string? true
3  Is 3 a string? false

```

In F# all types are also objects, whose type is denoted `obj`. Thus, the example uses the generic type when defining the argument to `isString`, and then dynamic type pattern matching for further processing. See Chapter 20 for more on objects. Dynamic type patterns are often used for analyzing exceptions, which is discussed in Section 18.1. While dynamic type patterns are useful, they imply runtime checking, and **it is almost always better to prefer compile time than runtime type checking**. Advice

16 | Higher order functions

A *higher order function* is a function that takes a function as an argument and/or returns a function. Higher order functions are also sometimes called functionals or functors. F# is a functions-first programming language with strong support for working with functions as values: Functions evaluate as *closures*, see Section 6.2, which can be passed to and from functions as any other value. An example of a higher order function is `List.map`, which takes a function and a list and produces a list, demonstrated in Listing 16.1.

- higher order function
- closures

Listing 16.1 higherOrderMap.fsx:

`List.map` is a higher order function, since it takes a function as argument.

```
1 let inc x = x + 1
2 let newList = List.map inc [2; 3; 5]
3 printfn "%A" newList

1 $ fsharp -nologo higherOrderMap.fsx && mono higherOrderMap.exe
2 [3; 4; 6]
```

Here `List.map` applies the function `inc` to every element of the list. *Anonymous functions* are expressions that results in a function, which need not be named. A typical example is to use anonymous functions as arguments to higher order functions such as `List.map` shown in Listing 16.2.

- anonymous functions

Listing 16.2 higherOrderAnonymous.fsx:

An anonymous function is a higher order function used here as an unnamed argument. Compare with Listing 16.1.

```
1 let newList = List.map (fun x -> x + 1) [2; 3; 5]
2 printfn "%A" newList

1 $ fsharp -nologo higherOrderAnonymous.fsx
2 $ mono higherOrderAnonymous.exe
3 [3; 4; 6]
```

A very brief version of Listing 16.2 is shown in Listing 16.3.

Listing 16.3 `higherOrderAnonymousBrief.fsx`:
A compact version of Listing 16.1.

```
1 printfn "%A" (List.map (fun x -> x + 1) [2; 3; 5])

-----

1 $ fsharp --nologo higherOrderAnonymousBrief.fsx
2 $ mono higherOrderAnonymousBrief.exe
3 [3; 4; 6]
```

What was originally three lines in Listing 16.1 as been reduced to a single line in Listing 16.3. While the result is exactly the same, the later is more difficult to read: First the reader has to understand what the anonymous function does on a single argument, then how this applies to the list through `list.Map`, and finally what `printfn` does with the result.

Anonymous functions are also useful as return values of functions as shown in Listing 16.4

Listing 16.4 `higherOrderReturn.fsx`:
The procedure `inc` returns an increment function. Compare with Listing 16.1.

```
1 let inc n =
2     fun x -> x + n
3 printfn "%A" (List.map (inc 1) [2; 3; 5])

-----

1 $ fsharp --nologo higherOrderReturn.fsx && mono higherOrderReturn.exe
2 [3; 4; 6]
```

Here the `inc` function produces a customized incrementation function as argument to `List.map`: It adds a prespecified number to an integer argument. Note that the closure of this customized function is only produced once, when the arguments for `List.map` is prepared, and not every time `List.map` maps the function to the elements of the list. Compare with Listing 16.1.

Piping is another example of a set of higher order function: `(<|)`, `(|>)`, `(<||)`, `(||>)`, `(<|||)`, `(|||>)`.¹ E.g., the functional equivalent of the right-to-left piping operator takes a value and a function and applies the function to the value as demonstrated in Listing 16.5.

Listing 16.5 `higherOrderPiping.fsx`:
The functional equivalent of the right-to-left piping operator is a higher order function.

```
1 let inc x = x + 1
2 let aValue = 2
3 let anotherValue = (|>) aValue inc
4 printfn "%d -> %d" aValue anotherValue

-----

1 $ fsharp --nologo higherOrderPiping.fsx && mono higherOrderPiping.exe
2 2 -> 3
```

Here the piping operator is used to apply the `inc` function to `aValue`. A more elegant way to write

¹Jon: Make piping operators go into index.

this would be `aValue |> inc`, or even just `inc aValue`.

16.1 Function composition

Piping is a useful shorthand for composing functions where the focus is on the transformation of arguments and results. Using higher order functions, we can forgo the arguments and compose functions as functions directly. This is done with the “>>” and “<<” operators. An example is given in Listing 16.6.

· function composition
· >>
· <<

Listing 16.6 `higherOrderComposition.fsx`:
Functions defined as composition of other functions.

```
1 let f x = x + 1
2 let g x = x * x
3 let h = f >> g
4 let k = f << g
5 printfn "%d" (g (f 2))
6 printfn "%d" (h 2)
7 printfn "%d" (f (g 2))
8 printfn "%d" (k 2)

1 $ fsharp -nologo higherOrderComposition.fsx
2 $ mono higherOrderComposition.exe
3 9
4 9
5 5
6 5
```

In the example we see that `(f >> g) x` gives the same result as `g (f x)`, while `(f << g) x` gives the same result as `f (g x)`. A memo technique for remembering the order of the application, when using the function composition operators, is that `(f >> g) x` is the same as `x |> f |> g`, i.e., the result of applying `f` to `x` is the argument to `g`. However, there is a clear distinction between the piping and composition operators. The type of the piping operator is

`(|>) : ('a, 'a -> 'b) -> 'b`

i.e., the piping operator takes a value of type `'a` and a function of type `'a -> 'b` and applies the function to the value and produces the value `'b`. In contrast, the composition operator has type

`(>>) : ('a -> 'b, 'b -> 'c) -> ('a -> 'c)`

i.e., it takes two functions of type `'a -> 'b` and `'b -> 'c` respectively, and produces a new function of type `'a -> 'c`.

16.2 Currying

Consider a function `f` of two generic arguments. It's type in F# will be `f : 'a -> 'b -> 'c`, meaning that `f` takes an argument of type `'a` and returns a function of type `'b -> 'c`. That is, if just one argument is given, then the result is a function, not a value. This is called *partial specification* or

· partial specification

currying in tribute of Haskell Curry². An example is given in Listing 16.7.

Listing 16.7 `higherOrderCurrying.fsx`:
Currying: defining a function as a partial specification of another.

```

1  let mul x y = x*y
2  let timesTwo = mul 2.0
3  printfn "%g" (mul 5.0 3.0)
4  printfn "%g" (timesTwo 3.0)

```

```

1  $ fsharpc --nologo higherOrderCurrying.fsx
2  $ mono higherOrderCurrying.exe
3  15
4  6

```

Here, `mul 2.0` is a partial application of the function `mul x y`, where the first argument is fixed, and hence, `timesTwo` is a function of 1 argument being the second argument of `mul`. The same can be achieved using tuple arguments, as shown in Listing 16.8.

Listing 16.8 `higherOrderTuples.fsx`:
Partial specification of functions using tuples is less elegant. Compare with Listing 16.7.

```

1  let mul (x, y) = x*y
2  let timesTwo y = mul (2.0, y)
3  printfn "%g" (mul (5.0, 3.0))
4  printfn "%g" (timesTwo 3.0)

```

```

1  $ fsharpc --nologo higherOrderTuples.fsx && mono higherOrderTuples.exe
2  15
3  6

```

Conversion between multiple and tuple arguments is easily done with higher order functions as demonstrated in Listing 16.9.

Listing 16.9: Two functions to convert between two and 2-tuple arguments.

```

1  > let curry f x y = f (x,y)
2  - let uncurry f (x,y) = f x y;;
3  val curry : f:(('a * 'b -> 'c) -> x:'a -> y:'b -> 'c)
4  val uncurry : f:(('a -> 'b -> 'c) -> x:'a * y:'b -> 'c)

```

Conversion between multiple and tuple arguments are useful, when working with higher order functions such as `List.map`. E.g., if `let mul (x, y) = x * y` as in Listing 16.8, then `curry mul` has the type `x:'a -> y:'b -> 'c` as can be seen in Listing 16.9, and thus, is equal to the anonymous function `fun x y -> x * y`. Hence, `curry mul 2.0` is equal to `fun y -> 2.0 * y`, since the precedence of function calls is `(curry mul) 2.0`.

Currying makes elegant programs and is often used in functional programming. Nevertheless, currying

²Haskell Curry (1900–1982) was an American mathematician and logician, who also has a programming language named after him: Haskell.

may lead to obfuscation, and in general **currying should be used with care and be well documented for proper readability of code.** Advice

17 | The functional programming paradigm

Functional programming is a style of programming which performs computations by evaluating functions. Functional programming avoids mutable values and side-effects. It is declarative in nature, e.g., by the use of value- and function-bindings – **let**-bindings – and avoids statements – **do**-bindings. Thus, the result of a function in functional programming depends only on its arguments, and therefore functions have no side-effect and are deterministic, such that repeated call to a function with the same arguments always gives the same result. In functional programming, data and functions are clearly separated, and hence data structures are dum as compared to objects in object-oriented programming paradigm, see Chapter 22. Functional programs clearly separate behavior from data and subscribes to the view that *it is better to have 100 functions operate on one data structure than 10 functions on 10 data structures*. Simplifying the data structure has the advantage that it is much easier to communicate data than functions and procedures between programs and environments. The .Net, mono, and java’s virtual machine are all examples of an attempt to rectify this, however, the argument still holds.

The functional programming paradigm can trace its roots to lambda calculus introduced by Alonzo Church in 1936 [1]. Church designed lambda calculus to discuss computability. Some of the forces of the functional programming paradigm are that it is often easier to prove the correctness of code, and since no states are involved, then functional programs are often also much easier to parallelize than other paradigms.

Functional programming has a number of features:

Pure functions

Functional programming is performed with pure functions. A pure function always returns the same value, when given the same arguments, and it has no side-effects. A function in F# is an example of a pure function. Pure functions can be replaced by their result without changing the meaning of the program. This is known as *referential transparency*.

· pure function

· referential
transparency

Higher order functions

Functional programming makes use of higher order functions, where functions may be given as arguments and returned as results of a function application. Higher order functions and *first-class citizenship* are related concepts, where higher-order functions are the mathematical description of functions that operator on functions, while a first-class citizen is the computer science term for functions as values. F# implements higher-order functions.

· higher order function

· first-class citizenship

Recursion

Functional programs use recursion instead of **for**- and **while**-loops. Recursion can make programs ineffective, but compilers are often designed to optimize tail-recursion calls. Common recursive programming structures are often available as optimized higher-order functions such as *iter*, *map*, *reduce*, *fold*, and *foldback*. F# has good support for all of these features.

· recursion

· iter

· map

· reduce

· fold

· foldback

Immutable states

Functional programs operate on values not on variables. This implies lexicographical scope in contrast to mutable values, which implies dynamic scope.

· immutable state

· immutable state

· strongly typed

Strongly typed

Functional programs are often strongly typed, meaning that types are set no later than at compile-time. F# does have the ability to perform runtime type assertion, but for most parts it relies on explicit type annotations and type inference at compile-time. The implication is that type errors are caught at compile time instead of at runtime.

Lazy evaluation

Due to referential transparency, values can be compute any time up until the point, when needed. Hence, they need not be computed at compilation time, which allows for infinite data structures. F# has support for lazy evaluations using the *lazy*-keyword, sequences using the *seq*-type, and computation expressions, all of which are advanced topics and not treated in this book.

· lazy evaluation

· lazy
· seq

Immutable states imply that data structures in functional programming are different than in imperative programming. E.g., in F# lists are immutable, so if an element of a list is to be changed, a new list must be created by copying all old values except that which is to be changed. Such an operation is therefore linear in computational complexity. In contrast, arrays are mutable values, and changing a value is done by reference to the value's position and change the value at that location. This has constant computational complexity. While fast, mutable values give dynamic scope, makes reasoning about the correctness of a program harder, since mutable states do not have referential transparency.

Functional programming may be considered a subset of *imperative programming*, in the sense that functional programming does not include the concept of a state, or one may think of functional programming as only having one unchanging state. Functional programming has also a bigger focus on declaring rules for *what* should be solved, and not explicitly listing statements describing *how* these rules should be combined and executed in order to solve a given problem. Functional programming is often found to be less error-prone at runtime making more stable, safer programs, and less open for, e.g., hacking.

· imperative
programming

17.1 Functional design

A key to all good programming designs is encapsulating code into modules. For functional programs, the essence is to consider data and functions as transformations of data. I.e., the basic pattern is a piping sequence,

```
x |> f |> g |> h,
```

where *x* is the input data, and *f*, *g*, and *h* are functions that transform the data. Of course, most long programs include lots of control structure implying that we would need junctions in the pipe system, however, piping is a useful memo technique.

In functional programming there are some pitfalls, that you should avoid:

- Creating large data structures, such as a single record containing all data. Since data is immutable, changing a single field in a monstrous record would mean a lot of copying in many parts of your program. Better is to use a range of data structures that express isolated semantic units of your problem.
- Non-tail recursion. Relying on the built-in functions `map`, `fold`, etc., is a good start for efficiency.
- Single character identifiers. Since functional programming tends to produce small, well-defined functions, there is a tendency to the use of single character identifiers, e.g., `let f x = ...`. In the very small, this can be defended, but the names used as identifiers can be used to increase the readability of code to yourself or to others. Typically, identifiers are long and informative in the outer most scope, while decreasing in size as you move in.

- Few comments. Since functional programming is very concise, there is a tendency, that we as programmers forget to add sufficient comments to the code, since at the time of writing, the meaning may be very clear and well thought through, but experience shows that this clarity deteriorates fast with time.
- Identifiers that are meaningless clones of each other. Since identifiers cannot be reused except by overshadowing in deeper scopes, there is often a tendency to have a family of identifiers like `a`, `a2`, `newA` etc. Better is to use names, that more clearly states the semantic meaning of the values, or, if only used as temporary storage, to discard them completely in lieu of piping and function composition. However, the latter most often requires comments describing the transformation being performed.

Thus, a design pattern for functional programs must focus on,

- What is the input data to be processed
- How is the data to be transformed

For large programs, the design principle is often similar to other programming paradigms, which are often visualized graphically as components that take input, interact, and produces results often together with a user. The effect of functional programming is mostly seen in the small, i.e., where a subtask is to be structure functionally.

18 | Handling Errors and Exceptions

18.1 Exceptions

Exceptions are runtime errors, such as division by zero. E.g., attempting integer division by zero halts execution and a long somewhat cryptic error message is written to screen as illustrated in Listing 18.1.

Listing 18.1: Division by zero halts execution with an error message.

```
1 > 3 / 0;;
2 System.DivideByZeroException: Attempted to divide by zero.
3   at <StartupCode$FSI_0002>.$FSI_0002.main@ () [0x00001] in
4     <0e5b9fd12a6649c598d7fa8c09a58dd3>:0
5     at (wrapper managed-to-native)
6       System.Reflection.MonoMethod:InternalInvoke
7       (System.Reflection.MonoMethod,object,object[],System.Exception&)
8     at System.Reflection.MonoMethod.Invoke (System.Object obj,
9       System.Reflection.BindingFlags invokeAttr, System.Reflection.Binder
10        binder, System.Object[] parameters, System.Globalization.CultureInfo
11        culture) [0x00032] in <c9f8153c41de4f8cbafd0e32f9bf6b28>:0
12 Stopped due to error
```

The error message starts by `System.DivideByZeroException: Attempted to divide by zero`, followed by a description of which libraries were involved when the error occurred, and finally F# informs us that it `Stopped due to error`. The type `System.DivideByZeroException` is a built-in exception type, and the built-in integer division operator chooses to raise the exception when the undefined division by zero is attempted. Many times such errors can be avoided by clever program design. However, this is not always possible or desirable, which is why F# implements exception handling for graceful control.

Exceptions are a basic-type called *exn*, and F# has a number of built-in, a few of which are listed in Table 18.1.

Exceptions are handled by the `try`-keyword expressions. We say that an expression may *raise* or *cast* an exception, the `try`-expression may *catch* and *handle* the exception by another expression.

Exceptions like in Listing 18.1 may be handled by `try-with` expressions as demonstrated in Listing 18.2.

- *exn*
- raising exception
- casting exceptions
- catching exception
- handling exception

Attribute	Description
ArgumentException	Arguments provided are invalid.
DivideByZeroException	Division by zero.
NotFiniteNumberException	floating point value is plus or minus infinity, or Not-a-Number (NaN).
OverflowException	Arithmetic or casting caused an overflow.
IndexOutOfRangeException	Attempting to access an element of an array using an index which is less than zero or equal or greater than the length of the array.

Table 18.1: Some built-in exceptions. The prefix `System.` has been omitted for brevity.

Listing 18.2 `exceptionDivByZero.fsx`:
A division by zero is caught and a default value is returned.

```
1 let div enum denom =
2     try
3         enum / denom
4     with
5         | :? System.DivideByZeroException -> System.Int32.MaxValue
6
7 printfn "3 / 1 = %d" (div 3 1)
8 printfn "3 / 0 = %d" (div 3 0)
```

```
1 $ fsharpc --nologo exceptionDivByZero.fsx && mono exceptionDivByZero.exe
2 3 / 1 = 3
3 3 / 0 = 2147483647
```

In the example, when the division operator raises the `System.DivideByZeroException` exception, then `try-with` catches it and returns the value `System.Int32.MaxValue`. Division by zero is still an undefined operation, but with the exception system, the program is able to receive a message about this undefined situation and choose an appropriate action.

The `try` expressions comes in two flavors: `try-with` and `try-finally` expressions.

The `try-with` expression has the following syntax,

· `try-with`

Listing 18.3 Syntax for the `try-with` exception handling.

```
1 try
2     <testExpr>
3 with
4     [ | ] <pat1> -> <exprHndl1>
5     | <pa2> -> <exprHndl2>
6     | <pat3> -> <exprHndl3>
7     ...
```

where `<testExpr>` is an expression, which might raise an exception, `<patn>` is a pattern, and `<exprHndl n>` is the corresponding exception handler. The value of the `try`-expression is either the value of `<testExpr>`, if it does not raise an exception, or the value of the exception handler `<exprHndl n>` of the first matching pattern `<patn>`. The above is lightweight syntax. Regular syntax omits newlines.

In Listing 18.2 *dynamic type matching* is used (see Section 15.9) using the “`?:`” lexeme, i.e., the `?:` dynamic type pattern matches exceptions at runtime which has the `System.DivideByZeroException` type. The

exception value may contain further information and can be accessed if named using the `as`-keyword as demonstrated in Listing 18.4.

Listing 18.4 `exceptionDivByZeroNamed.fsx`:
Exception value is bound to a name. Compare to Listing 18.2.

```

1 let div enum denom =
2     try
3         enum / denom
4     with
5         | :? System.DivideByZeroException as ex ->
6             printfn "Error: %s" ex.Message
7             System.Int32.MaxValue
8
9 printfn "3 / 1 = %d" (div 3 1)
10 printfn "3 / 0 = %d" (div 3 0)

```

```

1 $ fsharp --nologo exceptionDivByZeroNamed.fsx
2 $ mono exceptionDivByZeroNamed.exe
3 3 / 1 = 3
4 Error: Attempted to divide by zero.
5 3 / 0 = 2147483647

```

Here the exception value is bound to the name `ex`.

All exceptions may be caught as the dynamic type `System.Exception`, and F# implements a short-hand for catching exceptions and binding its value to a name as demonstrated in

Listing 18.5 `exceptionDivByZeroShortHand.fsx`:
An exception of type `System.Exception` is bound to a name. Compare to Listing 18.4.

```

1 let div enum denom =
2     try
3         enum / denom
4     with
5         | ex -> printfn "Error: %s" ex.Message; System.Int32.MaxValue
6
7 printfn "3 / 1 = %d" (div 3 1)
8 printfn "3 / 0 = %d" (div 3 0)

```

```

1 $ fsharp --nologo exceptionDivByZeroShortHand.fsx
2 $ mono exceptionDivByZeroShortHand.exe
3 3 / 1 = 3
4 Error: Attempted to divide by zero.
5 3 / 0 = 2147483647

```

Finally, the short-hand may be guarded with a `when`-guard as demonstrated in Listing 18.6.

· `when`

Listing 18.6 exceptionDivByZeroGuard.fsx:

An exception of type `System.Exception` is bound to a name and guarded. Compare to Listing 18.5.

```
1 let div enum denom =
2     try
3         enum / denom
4     with
5         | ex when enum = 0 -> 0
6         | ex -> System.Int32.MaxValue
7
8 printfn "3 / 1 = %d" (div 3 1)
9 printfn "3 / 0 = %d" (div 3 0)
10 printfn "0 / 0 = %d" (div 0 0)
-----
1 $ fsharp --nologo exceptionDivByZeroGuard.fsx
2 $ mono exceptionDivByZeroGuard.exe
3 3 / 1 = 3
4 3 / 0 = 2147483647
5 0 / 0 = 0
```

The first pattern only matches the `System.Exception` exception when `enum` is 0, in which case the exception handler returns 0.

Thus, if you don't care about the type of exception, then you need only use the short-hand pattern matching and name binding demonstrated in Listing 18.5 and Listing 18.6, but if you would like to distinguish between types of exceptions, then you must use explicit type matching and possibly value binding demonstrated in Listing 18.2 and Listing 18.4

The `try-finally` expression has the following syntax,

· `try-finally`

Listing 18.7 Syntax for the `try-finally` exception handling.

```
1 try
2     <testExpr>
3 finally
4     <cleanupExpr>
```

The `try-finally` expression evaluates the `<cleanupExpr>` expression following evaluation of the `<testExpr>` regardless of whether an exception is raised or not as illustrated in Listing 18.8.

Listing 18.8 exceptionDivByZeroFinally.fsx:

The `finally` branch is executed regardless of an exception.

```

1  let div enum denom =
2      printf "Doing division:"
3      try
4          printf " %d %d." enum denom
5          enum / denom
6      finally
7          printfn " Division finished."
8
9  printfn "3 / 1 = %d" (try div 3 1 with ex -> 0)
10 printfn "3 / 0 = %d" (try div 3 0 with ex -> 0)

```

```

1  $ fsharpc --nologo exceptionDivByZeroFinally.fsx
2  $ mono exceptionDivByZeroFinally.exe
3  Doing division: 3 1. Division finished.
4  3 / 1 = 3
5  Doing division: 3 0. Division finished.
6  3 / 0 = 0

```

Here, the `finally` branch is evaluated following the evaluation of the test expression regardless of whether the test expression raises an exception or not. However, if an exception is raised in a `try-finally` expression and there is no outer `try-with` expression, then execution stops without having evaluated the `finally` branch.

Exceptions can be raised using the *raise*-function

· `raise`Listing 18.9 Syntax for the `raise` function that raises exceptions.

```

1  raise (<expr>)

```

An example of raising the `System.ArgumentException` is shown in Listing 18.10

Listing 18.10 `raiseArgumentException.fsx`:

Raising the division by zero with customized message.

```

1  let div enum denom =
2      if denom = 0 then
3          raise (System.ArgumentException "Error: \"division by 0\"")
4      else
5          enum / denom
6
7  printfn "3 / 0 = %s" (try (div 3 0 |> string) with ex -> ex.Message)

```

```

1  $ fsharpc --nologo raiseArgumentException.fsx
2  $ mono raiseArgumentException.exe
3  3 / 0 = Error: "division by 0"

```

In this example, division by zero is never attempted, instead an exception is raised, which must be handled by the caller. Note that the type of `div` is `int -> int -> int` because `denom` is compared with

an integer in the conditional statement. This contradicts the typical requirements for `if` statements, where every branch has to return the same type. However, any code that explicitly raises exceptions are ignored, and the type is inferred by the remaining branches.

Programs may define new exceptions using the syntax,

Listing 18.11 Syntax for defining new exceptions.

```
1  exception <ident> of <typeId> {* <typeId>}
```

An example of defining a new exception and raising it is given in Listing 18.12.

Listing 18.12 exceptionDefinition.fsx:

A user-defined exception is raised but not caught by outer construct.

```
1  exception DontLikeFive of string
2
3  let picky a =
4      if a = 5 then
5          raise (DontLikeFive "5 sucks")
6      else
7          a
8
9  printfn "picky %A = %A" 3 (try picky 3 |> string with ex -> ex.Message)
10 printfn "picky %A = %A" 5 (try picky 5 |> string with ex -> ex.Message)
-----
1  $ fsharpc --nologo exceptionDefinition.fsx
2  $ mono exceptionDefinition.exe
3  picky 3 = "3"
4  picky 5 = "Exception of type 'ExceptionDefinition+DontLikeFive' was
   thrown."
```

Here an exception called `DontLikeFive` is defined, and it is raised in the function `picky`. The example demonstrates that catching the exception as a `System.Exception` as in Listing 18.5 the `Message` property includes information about the exception name but not its argument. To retrieve the argument `"5 sucks"`, we must match the exception with correct exception name as demonstrated in Listing 18.13.

Listing 18.13 exceptionDefinitionNCatch.fsx:
Catching a user-defined exception.

```

1  exception DontLikeFive of string
2
3  let picky a =
4      if a = 5 then
5          raise (DontLikeFive "5 sucks")
6      else
7          a
8
9  try
10     printfn "picky %A = %A" 3 (picky 3)
11     printfn "picky %A = %A" 5 (picky 5)
12 with
13 | DontLikeFive msg -> printfn "Exception caught with message: %s" msg

```

```

1  $ fsharpc --nologo exceptionDefinitionNCatch.fsx
2  $ mono exceptionDefinitionNCatch.exe
3  picky 3 = 3
4  Exception caught with message: 5 sucks

```

F# includes the *failwith* function to simplify the most common use of exceptions. It is defined as `failwith : string -> exn` and takes a string and raises the built-in `System.Exception` exception. An example of its use is shown in Listing 18.14.

Listing 18.14 exceptionFailwith.fsx:
An exception raised by failwith.

```

1  if true then failwith "hej"

```

```

1  $ fsharpc --nologo exceptionFailwith.fsx && mono exceptionFailwith.exe
2
3  Unhandled Exception:
4  System.Exception: hej
5      at <StartupCode$exceptionFailwith>.$ExceptionFailwith$fsx.main@ ()
6      [0x0000b] in <599574c21515099da7450383c2749559>:0
7  [ERROR] FATAL UNHANDLED EXCEPTION: System.Exception: hej
8      at <StartupCode$exceptionFailwith>.$ExceptionFailwith$fsx.main@ ()
9      [0x0000b] in <599574c21515099da7450383c2749559>:0

```

To catch the *failwith* exception, there are two choices, either use the `:?` or the `Failure` pattern. the `:?` pattern matches types, and we can match with the type of `System.Exception` as shown in Listing 18.15.

Listing 18.15 exceptionSystemException.fsx:
Catching a failwith exception using type matching pattern.

```

1 let _ =
2     try
3         failwith "Arrrrrg"
4     with
5         :? System.Exception -> printfn "So failed"

```

```

1 $ fsharpc --nologo exceptionSystemException.fsx
2
3 exceptionSystemException.fsx(5,5): warning FS0067: This type test or
  downcast will always hold
4
5 exceptionSystemException.fsx(5,5): warning FS0067: This type test or
  downcast will always hold
6 $ mono exceptionSystemException.exe
7 So failed

```

However, this gives annoying warnings, since F# internally is built such that all exception matches the type of `System.Exception`. Instead it is better to either match using the wildcard pattern as in Listing 18.16,

Listing 18.16 exceptionMatchWildcard.fsx:
Catching a failwith exception using the wildcard pattern.

```

1 let _ =
2     try
3         failwith "Arrrrrg"
4     with
5         _ -> printfn "So failed"

```

```

1 $ fsharpc --nologo exceptionMatchWildcard.fsx
2 $ mono exceptionMatchWildcard.exe
3 So failed

```

or use the built-in `Failure` pattern as in Listing 18.17.

Listing 18.17 exceptionFailure.fsx:
Catching a failwith exception using the `Failure` pattern.

```

1 let _ =
2     try
3         failwith "Arrrrrg"
4     with
5         Failure msg ->
6             printfn "The castle of %A" msg

```

```

1 $ fsharpc --nologo exceptionFailure.fsx && mono exceptionFailure.exe
2 The castle of "Arrrrrg"

```

Notice how only the `Failure` pattern allows for the parsing of the message given to `failwith` as argument.

Invalid arguments is such a common reason for failures that a built-in function has been supplied in F#. The `invalidArg` takes 2 strings and raises the built-in `ArgumentException`

· `invalidArg`

Listing 18.18 exceptionInvalidArg.fsx:

An exception raised by `invalidArg`. Compare with Listing 18.10.

```
1  if true then invalidArg "a" "is too much 'a'"

-----

1  $ fsharp --nologo exceptionInvalidArg.fsx
2  $ mono exceptionInvalidArg.exe
3
4  Unhandled Exception:
5  System.ArgumentException: is too much 'a'
6  Parameter name: a
7      at <StartupCode$exceptionInvalidArg>.$ExceptionInvalidArg$fsx.main@ ()
8      [0x0000b] in <599574c911642f55a7450383c9749559>:0
9  [ERROR] FATAL UNHANDLED EXCEPTION: System.ArgumentException: is too much
10     'a'
11     Parameter name: a
12     at <StartupCode$exceptionInvalidArg>.$ExceptionInvalidArg$fsx.main@ ()
13     [0x0000b] in <599574c911642f55a7450383c9749559>:0
```

The `invalidArg` function raises an `System.ArgumentException` as shown in Listing 18.19.

Listing 18.19 exceptionInvalidArgNCAatch.fsx:

Catching the exception raised by `invalidArg`.

```
1  let _ =
2      try
3          invalidArg "a" "is too much 'a'"
4      with
5          :? System.ArgumentException -> printfn "Argument is no good!"

-----

1  $ fsharp --nologo exceptionInvalidArgNCAatch.fsx
2  $ mono exceptionInvalidArgNCAatch.exe
3  Argument is no good!
```

The `try` construction is typically used to gracefully handle exceptions, but there are times, where you may want to pass on the bucket, so to speak, and re-raise the exception. This can be done with the `reraise` as shown in Listing 18.20.

· `reraise`

Listing 18.20 exceptionReraise.fsx:
Reraising an exception.

```

1  let _ =
2      try
3          failwith "Arrrrrg"
4      with
5          Failure msg ->
6              printfn "The castle of %A" msg
7              reraise()

```

```

1  $ fsharp --nologo exceptionReraise.fsx && mono exceptionReraise.exe
2  The castle of "Arrrrrg"
3
4  Unhandled Exception:
5  System.Exception: Arrrrrg
6      at <StartupCode$exceptionReraise>.$ExceptionReraise$fsx.main@ ()
7      [0x00041] in <599574d491e0c9eea7450383d4749559>:0
8  [ERROR] FATAL UNHANDLED EXCEPTION: System.Exception: Arrrrrg
9      at <StartupCode$exceptionReraise>.$ExceptionReraise$fsx.main@ ()
10     [0x00041] in <599574d491e0c9eea7450383d4749559>:0

```

The `reraise` function is only allowed to be the final call in the expression of a `with` rule.

18.2 Option types

At exceptions, it is not always obvious what should be returned. E.g., in the Listing 18.2, the exception is handled gracefully, but the return value is somewhat arbitrarily chosen to be the largest possible integer, which is still far from infinity, which is the correct result. Instead we could use the *option type*. The option type is a wrapper, that can be put around any type, and which extends the type with the special value *None*. All other values are preceded by the *Some* identifier. An example of rewriting Listing 18.2 to correctly represent the non-computable value is shown in Listing 18.21.

· option type
· None
· Some

Listing 18.21: Option types can be used, when the value in case of exceptions is unclear.

```

1  > let div enum denom =
2      - try
3      -     Some (enum / denom)
4      - with
5      -     | :? System.DivideByZeroException -> None;;
6  val div : enum:int -> denom:int -> int option
7
8  >
9  - let a = div 3 1;;
10 val a : int option = Some 3
11
12 > let b = div 3 0;;
13 val b : int option = None

```

The value of an option type can be extracted by and tested for by its member function, *IsNone*, *IsSome*, and *Value* as illustrated in Listing 18.22.

· IsNone
· IsSome
· Value

Listing 18.22 option.fsx:
Simple operations on option types.

```

1  let a = Some 3;
2  let b = None;
3  printfn "%A %A" a b
4  printfn "%A %b %b" a.Value b.IsSome b.IsNone

$ fsharpc --nologo option.fsx && mono option.exe
Some 3 <null>
3 false true

```

The `Value` member is not defined for `None`, thus **explicit pattern matching for extracting values from an option type is preferred**, e.g., `let get (opt : 'a option) (def : 'a) = match opt with Some x -> x | _ -> def`. Note also that `printf` prints the value `None` as `<null>`. This author hopes, that future versions of the option type will have better visual representations of the `None` value. Advice

Functions of option types are defined using the *option*-keyword. E.g., to define a function with explicit option type annotation that always returns `None`, write `let f (x : 'a option) = None`. option

F# includes an extensive `Option` module. It defines among many other functions *Option.bind* `Option.bind` which implements `let bind f opt = match opt with None -> None | Some x -> f x`. The function `Option.bind` is demonstrated in Listing 18.23.

Listing 18.23: `Option.bind` is useful for cascading calculations on option types.

```

1  > Option.bind (fun x -> Some (2*x)) (Some 3);;
2  val it : int option = Some 6

```

The `Option.bind` is a useful tool for cascading functions that evaluates to option types.

18.3 Programming intermezzo: Sequential division of floats

The following problem illustrates cascading error handling:

Problem 18.1

Given a list of floats such as `[1.0; 2.0; 3.0]`, calculate the sequential division `1.0/2.0/3.0`.

A sequential division is safe if the list does not contain zero values. However, if any element in the list is zero, then error handling must be performed. An example using `failwith` is given in Listing 18.24.

Listing 18.24 seqDiv.fsx:
Sequentially dividing a list of numbers.

```

1 let rec seqDiv acc lst =
2     match lst with
3     | [] -> acc
4     | elm::rest when elm <> 0.0 -> seqDiv (acc/elm) rest
5     | _ -> failwith "Division by zero"
6
7 try
8     printfn "%A" (seqDiv 1.0 [1.0; 2.0; 3.0])
9     printfn "%A" (seqDiv 1.0 [1.0; 0.0; 3.0])
10 with
11     Failure msg -> printfn "%s" msg

```

```

1 $ fsharp -nologo seqDiv.fsx && mono seqDiv.exe
2 0.1666666667
3 Division by zero

```

In this example, a recursive function is defined which updates an accumulator element, initially set to the neutral value 1.0. Division by zero results in a `failwith` exception, wherefore we must wrap its use in a `try-with` expression.

Instead of using exceptions, we may use `Option.bind`. In order to use `Option.bind` for a sequence of non-option floats, we will define a division operator, that reverses the order of operands. This is shown in Listing 18.25.

Listing 18.25 seqDivOption.fsx:
Sequentially dividing a sequence of numbers using `Option.bind`. Compare with Listing 18.24.

```

1 let divideBy denom enum =
2     if denom = 0.0 then
3         None
4     else
5         Some (enum/denom)
6
7 let success =
8     Some 1.0
9     |> Option.bind (divideBy 2.0)
10    |> Option.bind (divideBy 3.0)
11 printfn "%A" success
12
13 let fail =
14     Some 1.0
15     |> Option.bind (divideBy 0.0)
16     |> Option.bind (divideBy 3.0)
17 printfn "%A; isNone: %b" fail fail.IsNone

```

```

1 $ fsharp -nologo seqDivOption.fsx && mono seqDivOption.exe
2 Some 0.1666666667
3 <null>; isNone: true

```


Here the function `divideBy` takes two non-option arguments and returns an option type. Thus, `Option.bind (divideBy 2.0) (Some 1.0)` is equal to `Some 0.5`, since `divideBy 2.0` is a function that divides any float argument by 2.0. Iterating `Option.bind (divideBy 3.0) (Some 0.5)` we calculate `Some 0.1666666667` or `Some (1.0/6.0)` as expected. In Listing 18.25 this is written as a single `let` binding using piping. And since `Option.bind` correctly handles the distinction between `Some` and `None` values, such piping sequences correctly handles possible errors as shown in Listing 18.25.

The sequential application can be extended to lists using `List.foldBack` as demonstrated in Listing 18.26.

Listing 18.26 seqDivOptionAdv.fsx:

Sequentially dividing a list of numbers using `Option.bind` and `List.foldBack`. Compare with Listing 18.25.

```

1  let divideBy denom enum =
2      if denom = 0.0 then
3          None
4      else
5          Some (enum/denom)
6  let divideByOption x acc =
7      Option.bind (divideBy x) acc
8
9  let success = List.foldBack divideByOption [3.0; 2.0; 1.0] (Some 1.0)
10 printfn "%A" success
11
12 let fail = List.foldBack divideByOption [3.0; 0.0; 1.0] (Some 1.0)
13 printfn "%A; isNone: %A" fail fail.IsNone

```

```

1  $ fsharpc --nologo seqDivOptionAdv.fsx && mono seqDivOptionAdv.exe
2  Some 0.16666666667
3  <null>; isNone: true

```

Since `List.foldBack` processes the list from the right, the list of integers have been reversed. Notice how `divideByOption` is the function spelled out in each piping step of Listing 18.25.

Exceptions and option type are systems to communicate errors up through a hierarchy of function calls. While exceptions favor imperative style programming, option types are functional style programming. Exceptions allow for a detailed report of the type of error to the caller, whereas option types only allow for flagging that an error has occurred.

19 | Working with files

An important part of programming is handling data. A typical source of data is hard-coded bindings and expressions from libraries or the program itself, and the result is often shown on a screen either as text output on the console. This is a good starting point when learning to program, and one which we have relied heavily upon in this book until now. However, many programs require more: We often need to ask a user to input data via, e.g., typing text on a keyboard, clicking with a mouse, striking a pose in front of a camera. We also often need to load and save data to files, retrieve and deposit information from the internet, and visualize data as graphically, as sounds, or by controlling electrical appliances. Graphical user interfaces will be discussed in Chapter 23, and here we will concentrate on working with the console, with files, and with the general concept of streams.

File and stream input and output are supported via built-in namespaces and classes. For example, the `printf` family of functions discussed in Section 6.5 is defined in the `Printf` module of the `Fsharp.Core` namespace, and it is used to put characters on the `stdout` stream, i.e., to print on the screen. Likewise, `ReadLine` discussed in Section 6.6 is defined in the `System.Console` class, and it fetches characters from the `stdin` stream, that is, reads the characters the user types on the keyboard until the newline is pressed.

A *file* on a computer is a resource used to store data in and retrieve data from. Files are often associated with a physical device, such as a hard disk, but can also be a virtual representation in memory. Files are durable, such that other programs can access them independently, given certain rules for access. A file has a name, a size, and a type, where the type is related to the basic unit of storage such as characters, bytes, and words, (`char`, `byte`, and `int32`). Often data requires a conversion between the internal format to and from the format stored in the file. E.g., floating point numbers are sometimes converted to a UTF8 string using `fprintf` in order to store them in a file in a human-readable form and interpreted from UTF8 when retrieving them at a later point from the file. Files have a low-level structure and representation, which varies from device to device, and the low-level details are less relevant for the use of the file and most often hidden for the user. Basic operations on files are *creation*, *opening*, *reading from*, *writing to*, *closing*, and *deleting* files.

- file
- create file
- open file
- read file
- write file
- close file
- delete file
- stream

A *stream* is similar to files in that they are used to store data in and retrieve data from, but streams only allow for handling of data one element at a time like the readout of a thermometer: we can make temperature readings as often as we like, making notes and thus saving a history of temperatures, but we cannot access the future. Hence, streams are in principle without an end, and thus have infinite size, and data from streams are programmed locally by considering the present and previous elements. In contrast, files are finite in size and allow for global operations on all the file's data. Files may be considered a stream, but the opposite is not true.¹

¹Jon: Maybe add a figure illustrating the difference between files and streams.

19.1 Command line arguments

Compiled programs may be started from the console with one or more arguments. E.g., if we have made a program called `prog`, then arguments may be passed as `mono prog arg1 arg2`. To read the arguments in the program must define a function with the *EntryPoint* attribute, and this function must be of type `string array -> int`.

Listing 19.1 Defining an entry point function with arguments from the console.

```
1  [<EntryPoint>]
2  let <funcIdent> <arg> =
3      <bodyExpr>
```

`<funcIdent>` is the function's name, `<arg>` is the name of an array of strings, and `<bodyExpr>` is the function body. Return value 0 implies a successful execution of the program, while non-zero means failure. The entry point function must be in the last file compiled. An example is given in Listing 19.2.

Listing 19.2 `commandLineArgs.fsx`:
Interacting with a user with `ReadLine` and `WriteLine`.

```
1  [<EntryPoint>]
2  let main args =
3      printfn "Arguments passed to function : %A" args
4      0 // Signals that program terminated successfully
```

An example execution with arguments is shown in Listing 19.3.

Listing 19.3: An example dialogue of running Listing 19.2.

```
1  $ fsharp --nologo commandLineArgs.fsx
2  $ mono commandLineArgs.exe Hello World
3  Arguments passed to function : ["Hello"; "World"]
```

In Bash, the return value is called the *exit status* and can be tested using Bash's `if` statements as demonstrated in Listing 19.4.

Listing 19.4: Testing return values in Bash when running Listing 19.2.

```
1  $ fsharp --nologo commandLineArgs.fsx
2  $ if mono commandLineArgs.exe Hello World; then echo "success"; else
3      echo "failure"; fi
4  Arguments passed to function : ["Hello"; "World"]
5  success
```

Also in Bash, the exit status of the last executed program can be accessed using the `$?` built-in environment variable. In Windows this same variable is called `%errorlevel%`.

Stream	Description
<code>stdout</code>	Standard output stream used by <code>printf</code> and <code>printfn</code> .
<code>stderr</code>	Standard error stream used to display warnings and errors by Mono.
<code>stdin</code>	Standard input stream used to read keyboard input.

Table 19.1: Three built-in streams in `System.Console`.

Function	Description
<code>Write: string -> unit</code>	Write to the console. E.g., <code>System.Console.Write "Hello world"</code> . Similar to <code>printf</code> .
<code>WriteLine: string -> unit</code>	As <code>Write</code> but followed by a newline character, e.g., <code>WriteLine "Hello world"</code> . Similar to <code>printfn</code> .
<code>Read: unit -> int</code>	Read the next key from the keyboard blocking execution as long, e.g., <code>Read ()</code> . The key pressed is echoed to the screen.
<code>ReadKey: bool -> System.ConsoleKeyInfo</code>	As <code>Read</code> but returns more information about the key pressed. When given the value <code>true</code> as argument, then the key pressed is not echoed to the screen. E.g., <code>ReadKey true</code> .
<code>ReadLine unit -> string</code>	Read the next sequence of characters until newline from the keyboard, e.g., <code>ReadLine ()</code> .

Table 19.2: Some functions for interacting with the user through the console in the `System.Console` class. Prefix “`System.Console.`” is omitted for brevity.

19.2 Interacting with the console

From a programming perspective, the console is a stream: A program may send new data to the console, but cannot return to previously sent data and make changes. Likewise, the program may retrieve input from the user, but cannot go back and ask the user to have inputted something else, nor can we peak into the future and retrieve what the user will input in the future. The console uses three built-in streams in `System.Console` listed in Table 19.1. On the console, the standard output and error streams are displayed as text, and it is typically not possible to see a distinction between them. However, command-line interpreters such as Bash can, and it is possible from the command-line to filter output from programs according to these streams. However, a further discussion on this is outside the scope of this text. In `System.Console` there are many functions supporting interaction with the console, and the most important ones are shown in Table 19.2. Note that you must supply the empty argument “`()`” to the `Read` functions, in order to run most of the functions instead of referring to them as values. A demonstration of the use of `Write`, `WriteLine`, and `ReadLine` is given in Listing 19.5.

Listing 19.5 `userDialogue.fsx`:Interacting with a user with `ReadLine` and `WriteLine`. The user typed “3.5” and “7.4”.

```

1 System.Console.WriteLine "To perform the multiplication of a and b"
2 System.Console.Write "Enter a: "
3 let a = float (System.Console.ReadLine ())
4 System.Console.Write "Enter b: "
5 let b = float (System.Console.ReadLine ())
6 System.Console.WriteLine ("a * b = " + string (a * b))

```

```

1 $ fsharp --nologo userDialogue.fsx && mono userDialogue.exe
2 To perform the multiplication of a and b
3 Enter a: 3.5
4 Enter b: 7.4
5 a * b = 25.9

```

The functions `Write` and `WriteLine` acts as `printfn` but without a formatting string. These functions have many overloaded definitions the description of which is outside the scope of this book. **For Advice writing to the console, `printf` is to be preferred.**

Often `ReadKey` is preferred over `Read`, since the former returns a value of type `System.ConsoleKeyInfo`, which is a structure with three properties:

Key A `System.ConsoleKey` enumeration of the key pressed. E.g., the character 'a' is `ConsoleKey.A`.

KeyChar A unicode representation of the key.

Modifiers A `System.ConsoleModifiers` enumeration of modifier keys shift, ctrl, and alt.

An example of a dialogue is shown in Listing 19.6.

Listing 19.6 `readKey.fsx`:

Reading keys and modifiers. The user press 'a', 'shift-a', and 'ctrl-a', and the program was terminated by pressing 'ctrl-c'. The 'alt-a' combination does not work on MacOS.

```

1 open System
2
3 printfn "Start typing"
4 while true do
5     let key = Console.ReadKey true
6     let shift =
7         if key.Modifiers = ConsoleModifiers.Shift then "SHIFT+" else ""
8     let alt =
9         if key.Modifiers = ConsoleModifiers.Alt then "ALT+" else ""
10    let ctrl =
11        if key.Modifiers = ConsoleModifiers.Control then "CTRL+" else ""
12    printfn "You pressed: %s%s%s" shift alt ctrl (key.Key.ToString ())

```

```

1 $ fsharp --nologo readKey.fsx && mono readKey.exe
2 Start typing
3 You pressed: A
4 You pressed: SHIFT+A
5 You pressed: CTRL+A

```

System.IO.File	Description
<code>Open:</code> <code>(path : string) * (mode : FileMode)</code> <code>-> FileStream</code>	Request the opening of a file on <code>path</code> for reading and writing with access mode <code>FileMode</code> , see Table 19.4. Other programs are not allowed to access the file, before this program closes it.
<code>OpenRead: (path : string)</code> <code>-> FileStream</code>	Request the opening of a file on <code>path</code> for reading. Other programs may read the file regardless of this opening.
<code>OpenText: (path : string)</code> <code>-> StreamReader</code>	Request the opening of an existing UTF8 file on <code>path</code> for reading. Other programs may read the file regardless of this opening.
<code>OpenWrite: (path : string)</code> <code>-> FileStream</code>	Request the opening of a file on <code>path</code> for writing with <code>FileMode.OpenOrCreate</code> . Other programs may not access the file, before this program closes it.
<code>Create: (path : string)</code> <code>-> FileStream</code>	Request the creation of a file on <code>path</code> for reading and writing, overwriting any existing file. Other programs may not access the file, before this program closes it.
<code>CreateText: (path : string)</code> <code>-> StreamWriter</code>	Request the creation of an UTF8 file on <code>path</code> for reading and writing, overwriting any existing file. Other programs may not access the file, before this program closes it.

Table 19.3: The family of `System.IO.File.Open` functions. See Table 19.4 for a description of `FileMode`, Tables 19.5 and 19.6 for a description of `FileStream`, Table 19.7 for a description of `StreamReader`, and Table 19.8 for a description of `StreamWriter`.

19.3 Storing and retrieving data from a file

A file stored on the filesystem has a name, and it must be opened before it can be accessed and closed when finished. Opening files informs the operating system that your program is now going to use the file. While a file is open, the operating system will protect it depending on how the file is opened. E.g., if you are going to write to the file, then this typically implies that no one else may write to the file at the same time, since simultaneous writing to a file may leave the resulting file in an uncertain state. Sometimes the operating system will realize that a file, that was opened by a program, is no longer being used, e.g., since the program is no longer running, but **it is good practice always to release reserved files, e.g., by closing them as soon as possible, such that other programs may have access to it.** On the other hand, it is typically safe for several programs to read the same file at the same time, but it is still important to close files after their use, such that the operating system can effectively manage the computer's resources. Reserved files are just one of the possible obstacles that you may meet when attempting to open a file. Other points of failure may be that the file may not exist, your program may not have sufficient rights for accessing it, or the device, where the file is stored, may have unreliable access. Thus, **never assume that accessing files always works, but program defensively, e.g., by checking the return status of the file accessing functions and by `try` constructions.**

Advice

Advice

Data in files may have been stored in various ways, e.g., it may contain UTF8 encoded characters or sequences of floating point numbers stored as raw bits in chunks of 64 bits, or it may be a sequence of bytes that are later going to be interpreted as an image in jpeg or tiff format. To aid in retrieving the data, F# has a family of open functions, all residing in the `System.IO.File` class. These are described in Table 19.3.

For the general `Open` function, you must also specify how the file is to be opened. This is done with

FileMode	Description
Append	Open a file and seek to its end, if it exists, or create a new file. Can only be used together with FileAccess.Write. May throw <code>IOException</code> and <code>NotSupportedException</code> exceptions.
Create	Create a new file, and delete an already existing file. May throw the <code>UnauthorizedAccessException</code> exception.
CreateNew	Create a new file, but throw the <code>IOException</code> exception, if the file already exists.
Open	Open an existing file, and <code>System.IO.FileNotFoundException</code> exception is thrown if the file does not exist.
OpenOrCreate	Open a file, if exists, or create a new file.
Truncate	Open an existing file and truncate its length to zero. Cannot be used together with FileAccess.Read.

Table 19.4: File mode values for the `System.IO.Open` function.

a special set of values described in Table 19.4. An example of how a file is opened and later closed is shown in Listing 19.7.

Listing 19.7 `openFile.fsx`:

Opening and closing a file, in this case, the source code of this same file.

```

1  let filename = "openFile.fsx"
2
3  let reader =
4      try
5          Some (System.IO.File.Open (filename, System.IO.FileMode.Open))
6      with
7          _ -> None
8
9  if reader.IsSome then
10     printfn "The file %A was successfully opened." filename
11     reader.Value.Close ()

```

```

1  $ fsharpc --nologo openFile.fsx && mono openFile.exe
2  The file "openFile.fsx" was successfully opened.

```

Notice how the example uses the defensive programming style, where the `try`-expression is used to return the optional datatype, and further processing is made dependent on the success of the opening operation.

In F# the distinction between files and streams are not very clear. F# offers built-in support for accessing files as bytes through the `System.IO.FileStream` class, and for characters in a particular encoding through the `System.IO.TextReader` and `System.IO.TextWriter`.

A successfully opened `System.IO.FileStream` file, e.g., using `System.IO.File.OpenRead` from Table 19.3, will result in an `FileStream` object. From this object we can extract information about the file, such as the permitted operations and more listed in Table 19.5. This information is important in order to restrict the operation that we will perform on the file. Some typical operations are listed in and 19.6. E.g., we may `Seek` a particular position in the file, but only within the range of legal positions from 0 until the length of the file. Most operating systems do not necessarily write information to files immediately after one of the `Write` functions, but will often for optimization purposes will often collect information in a buffer that is to be written to a device in batches. However, sometimes it is useful to be able to force the operating system to empty its buffer to the device. This is called *flushing* and can be forced using the `Flush` function.

Property	Description
CanRead	Gets a value indicating whether the current stream supports reading. (Overrides <code>Stream.CanRead</code> .)
CanSeek	Gets a value indicating whether the current stream supports seeking. (Overrides <code>Stream.CanSeek</code> .)
CanWrite	Gets a value indicating whether the current stream supports writing. (Overrides <code>Stream.CanWrite</code> .)
Length	Gets the length in bytes of the stream. (Overrides <code>Stream.Length</code> .)
Name	Gets the name of the <code>FileStream</code> that was passed to the constructor.
Position	Gets or sets the current position of this stream. (Overrides <code>Stream.Position</code> .)

Table 19.5: Some properties of the `System.IO.FileStream` class.

Method	Description
Close ()	Closes the stream.
Flush ()	Causes any buffered data to be written to the file.
Read byte[] * int * int	Reads a block of bytes from the stream and writes the data in a given buffer.
ReadByte ()	Read a byte from the file and advances the read position to the next byte.
Seek int * SeekOrigin	Sets the current position of this stream to the given value.
Write byte[] * int * int	Writes a block of bytes to the file stream.
WriteByte byte	Writes a byte to the current position in the file stream.

Table 19.6: Some methods of the `System.IO.FileStream` class.

Text is typically streamed through the `StreamReader` and `StreamWriter`. These may be considered higher order stream processing, since they include an added interpretation of the bits to strings. A `StreamReader` has methods similar to a `FileStream` object and a few new properties and methods, such as the `EndOfStream` property and `ReadToEnd` method, see Table 19.7. Likewise, a `StreamWriter` has an added method for automatically flushing following every writing operation. A simple example of opening a text-file and processing it is given in Listing 19.8.

Property/Method	Description
EndOfStream	Check whether the stream is at its end.
Close ()	Closes the stream.
Flush ()	Causes any buffered data to be written to the file.
Peek ()	Reads the next character, but does not advance the position.
Read ()	Reads the next character.
Read char[] * int * int	Reads a block of bytes from the stream and writes the data in a given buffer.
ReadLine ()	Reads the next line of characters until a newline. Newline is discarded.
ReadToEnd ()	Reads the remaining characters till end-of-file.

Table 19.7: Some methods of the `System.IO.StreamReader` class.

Property/Method	Description
AutoFlush : bool	Get or set the auto-flush. If set, then every call to Write will flush the stream.
Close ()	Closes the stream.
Flush ()	Causes any buffered data to be written to the file.
Write 'a	Write a basic type to the file.
WriteLine string	As Write but followed by newline.

Table 19.8: Some methods of the `System.IO.StreamWriter` class.

Function	Description
Copy (src : string, dest : string)	Copy a file from src to dest possibly overwriting any existing file.
Delete string	Delete a file
Exists string	Check whether the file exists
Move (from : string, to : string)	Move a file from src to to possibly overwriting any existing file.

Table 19.9: Some methods of the `System.IO.File` class.**Listing 19.8** `readFile.fsx`:An example of opening a text file, and using the `StreamReader` properties and methods.

```

1 let printFile (reader : System.IO.StreamReader) =
2     while not(reader.EndOfStream) do
3         let line = reader.ReadLine ()
4         printfn "%s" line
5
6 let filename = "readFile.fsx"
7 let reader = System.IO.File.OpenText filename
8 printFile reader

```

```

1 $ fsharpc --nologo readFile.fsx && mono readFile.exe
2 let printFile (reader : System.IO.StreamReader) =
3     while not(reader.EndOfStream) do
4         let line = reader.ReadLine ()
5         printfn "%s" line
6
7 let filename = "readFile.fsx"
8 let reader = System.IO.File.OpenText filename
9 printFile reader

```

Here the program reads the source code of itself, and prints it to the console.

19.4 Working with files and directories.

F# has support for managing files summarized in the `System.IO.File` class and summarized in Table 19.9.

In the `System.IO.Directory` class there are a number of other frequently used functions summarized in Table 19.10.

Function	Description
CreateDirectory string	Create the directory and all implied sub-directories.
Delete string	Delete a directory
Exists string	Check whether the directory exists
GetCurrentDirectory ()	Get working directory of the program
GetDirectories (path : string)	Get directories in path
GetFiles (path : string)	Get files in path
Move (from : string, to : string)	Move a directory and its content from src to to.
SetCurrentDirectory : (path : string) -> unit	Set the current working directory of the program to path.

Table 19.10: Some methods of the `System.IO.Directory` class.

Function	Description
Combine string * string	Combine 2 paths into a new path.
GetDirectoryName (path: string)	Extract the directory name from path.
GetExtension (path: string)	Extract the extension from path.
GetFileName (path: string)	Extract the name and extension from path.
GetFileNameWithoutExtension (path : string)	Extract the name without the extension from path.
GetFullPath (path : string)	Extract the absolute path from path.
GetTempFileName ()	Create a uniquely named and empty file on disk and return its full path.

Table 19.11: Some methods of the `System.IO.Path` class.

In the `System.IO.Path` class there are a number of other frequently used functions summarized in Table 19.11.

19.5 Reading from the internet

The internet is a global collection of computers that are connected in a network using the internet protocol suite TCP/IP. The internet is commonly used for transport of data such as emails and for offering services such as web pages on the World Wide Web. Web resources are identified by a *Uniform Resource Locator (URL)* popularly known as a web page, and an URL contains information about where and how data from the web page is to be obtained. E.g., the URL `https://en.wikipedia.org/wiki/F_Sharp_(programming_language)`, contains 3 pieces of information: `https` is the protocol to be used to interact with the resource, `en.wikipedia.org` is the host's name, and `wiki/F_Sharp_(programming_language)` is the filename.

- Uniform Resource Locator
- URL

F#'s `System` namespace contains functions for accessing web pages as stream as illustrated in Listing 19.9.

Listing 19.9 webRequest.fsx:

Downloading a web page and printing the first few characters.

```

1  /// Set a url up as a stream
2  let url2Stream url =
3      let uri = System.Uri url
4      let request = System.Net.WebRequest.Create uri
5      let response = request.GetResponse ()
6      response.GetResponseStream ()
7
8  /// Read all contents of a web page as a string
9  let readUrl url =
10     let stream = url2Stream url
11     let reader = new System.IO.StreamReader(stream)
12     reader.ReadToEnd ()
13
14  let url = "http://fsharp.org"
15  let a = 40
16
17  let html = readUrl url
18  printfn "Downloaded %A. First %d characters are: %A" url a html.[0..a]

```

```

1  $ fsharpc --nologo webRequest.fsx && mono webRequest.exe
2  Downloaded "http://fsharp.org". First 40 characters are: "<!DOCTYPE html>
3  <html lang="en">
4  <head>"

```

To connect to a URL as a stream, we first need first format the URL string as a *Uniform Resource Identifiers (URI)*, which is a generalization of the URL concept, using the `System.Uri` function. Then we must initialize the request by the `System.Net.WebRequest` function, and the response from the host is obtained by the `GetResponse` method. Finally, we can access the response as a stream by the `GetResponseStream` method. In the end, we convert the stream to a `StreamReader`, such that we can use the methods from Table 19.7 to access the web page. ²

- Uniform Resource
- Identifiers
- URI

19.6 Resource Management

Streams and files are examples of computer resources, that may be shared by several applications. Most operating systems allow for several applications to be running in parallel, and to avoid unnecessarily blocking and hogging of resources, all responsible applications must release resources as soon as they are done using them. F# has language constructions for automatic releasing of resources: the `use` binding and the `using` function. These automatically dispose of resources, when the resource's name binding falls out of scope. Technically this is done by calling the `Dispose` method on objects that implement the `System.IDisposable` interface. See Section 21.3 for more on interfaces.

- Dispose
- System.IDisposable

The `use` keyword is similar to `let`:

- use

Listing 19.10 Use binding expression.

```

1  use <valueIdent> = <bodyExpr> [in <expr>]

```

²Jon: This section could be extended...

A `use` binding provides a binding between the `<bodyExpr>` expression to the name `<valueIdent>` in the following expression(s), and adds a call to `Dispose()` on `<valueIdent>` if it implements `System.IDisposable`. See for example Listing 19.11.

Listing 19.11 useBinding.fsx:

Using `use` instead of `let` releases disposable resources at end of scope.

```

1  open System.IO
2
3  let writeToFile (filename : string) (str : string) : unit =
4      use file = File.CreateText filename
5      file.Write str
6      // file.Dispose() is implicitly called here,
7      // implying that the file is closed.
8
9  writeToFile "use.txt" "Using 'use' closes the file, when out of scope."
```

Here, `file` is an `System.IDisposable` object, and before `writeToFile` returns, `file.Dispose()` is called automatically, which in this case implies that the file is closed. Had we used `let` instead, then the file would first be closed, when the program terminates.

The higher order function `using` takes a disposable object and a function, executes the function on the disposable objects and then calls `Dispose()` on the disposable object. This is illustrated in Listing 19.12

Listing 19.12 using.fsx:

The `using` function executes a function on an object and releases its disposable resources. Compare with Listing 19.11.

```

1  open System.IO
2
3  let writeToFile (str : string) (file : StreamWriter) : unit =
4      file.Write str
5
6  using (File.CreateText "use.txt") (writeToFile "Disposed after call.")
7  // Dispose() is implicitly called on the anonymous file handle implying
8  // that the file is automatically closed.
```

The main difference between `use` and `using` is that resources allocated using `use` are disposed at the end of its scope, while `using` disposes the resources after the execution of the function in its argument. In spite of the added control of `using`, we **prefer `use` over `using` due to its simpler structure.** Advice

19.7 Programming intermezzo: Ask user for existing file

A typical problem, when working with files, is

Problem 19.1

Ask the user for the name of an existing file.

Such a dialogue most often requires the program to aid the user, e.g., by telling the user, which files are available, and to check that the filename entered is an existing file. A solution could be,

Listing 19.13 filename dialogue.fsx:

...

```
1  let getAFileName () =  
2      let mutable filename = Unchecked.defaultof<string>  
3      let mutable fileExists = false  
4      while not(fileExists) do  
5          System.Console.Write("Enter Filename: ")  
6          filename <- System.Console.ReadLine()  
7          fileExists <- System.IO.File.Exists filename  
8          filename  
9  
10     let listOfFiles = System.IO.Directory.GetFiles "."  
11     printfn "Directory contains: %A" listOfFiles  
12     let filename = getAFileName ()  
13     printfn "You typed: %s" filename
```

20 | Classes and objects

Object-oriented programming is a programming paradigm that focusses on *objects* such as a person, place, thing, event, and concept relevant for the problem.

- object-oriented programming
- objects

Object-oriented programming has a rich language for describing objects and their relations, which can seem overwhelming at first, and they will be explained in detail in this and following chapters. Here is a brief overview: Objects may contain data and code, which may be either *public* or *private*. An object's *members* are the object's public values and functions. Public values are called *properties* or *attributes*, and public functions are called *methods*, and these can be accessed using the “.” notation similarly to modules and namespaces. Private values are called *fields* and private functions are just called *functions*, and these can only be used by code inside the object. The type definition of an object is called a *class*, while values of the class are called *objects*. When objects are created, a special function called the *constructor* is executed. Creating objects is also often referred to as *instantiating* objects.

- public
- private
- members
- properties
- attributes
- methods
- fields
- functions

Object-oriented programming is an extension of data types, in the sense that objects contain both data and functions in a similar manner as a module, but object-oriented programming emphasizes the semantic unity of the data and functions. Thus, objects are often *models* of real-world entities, and object-oriented programming leads to a particular style of programming analysis and design called *object-oriented analysis and design* to be discussed in Chapter 22.

- class
- objects
- constructor
- instantiation
- models
- object-oriented analysis
- object-oriented design

20.1 Constructors and members

An *object* is a variable of a *class* type. A class is defined using the **type** keyword, and there are *always* parentheses after the class name to distinguish it from a regular type definition.

- object
- class

Listing 20.1 A class definition.

```
1 type <classIdent> ({<arg>}) [as <selfIdent>]
2   [class]
3   [inherit <baseClassIdent>({<arg>})]
4   {[let <binding>] | [do <statement>]}
5   {(member | abstract member | default | override) <memberDef>}
6   [end]
```

The **<classIdent>** is the name of the class, **<arg>** are its optional arguments, **<selfIdent>** is an optional *self identifier*, **<baseClassIdent>** is the name of another class that this class optionally builds upon using the **inherit** keyword (see Section 21.1), the optional **let**-bindings and **do**-statements define *fields* and *functions*, and **<memberDef>** are public members, i.e., *properties* and *methods*. Members may be regular members using the **member** keyword or abstract members using either the **abstract**

- self identifier
- fields
- functions
- properties
- methods

`member`, `default`, or `override` keywords (see Section 21.2). The *primary constructor* is everything until the first member. Mutably recursive class definition can be defined using the `and` keyword, e.g., `type aClass () = ... and bClass () =` Do statements must use the `do`-keyword.

An example of a simple program defining a class and creating objects of their type is given in Listing 20.2.

Listing 20.2 class.fsx:

A class definition and an object of this class.

```

1 type aClass (anArgument : int) =
2     // Constructor body section
3     let objectValue = anArgument
4     do printfn "A class has been created (%A)" objectValue
5     // Member section
6     member this.value = anArgument
7     member this.scale (factor : int) = factor * objectValue
8
9 let a = aClass (2)
10 let b = aClass (1)
11 printfn "%d %d %d" a.value (a.scale 3) b.value

```

```

1 $ fsharp --nologo class.fsx && mono class.exe
2 A class has been created (2)
3 A class has been created (1)
4 2 6 1

```

In the example, the class `aClass` is defined in the header in line 1, and it includes one integer argument, `anArgument`. Classes can also be defined without arguments, but the parentheses cannot be omitted. Together with the header, line 2–4 is the primary constructor. In the member section line 5–7 is the value `value : int` and function `scale : int -> int` defined using the name `this` as a *self identifier*. If not declared using the `as` keyword in the header, then the self identifier can be any valid identifier. In line 9 and 10 two objects `a` and `b` of type `aClass` are created, which implies that memory is reserved on *The Heap* (see Section 6.8) and the constructor is run for each of them. Thus, for `a`, `this.value` refers to the memory set aside for `a`, and for `b`, `this.value` refers to the memory set aside for `b`. In line 11 are shown examples of their use. Notice, that members are accessed outside the object using the `."` notation in the same manner as an application program would access elements of a module. In many languages, objects are instantiated using the `new` keyword, but in F# this is optional. I.e., `let a = aClass (2)` is identical to `let a = new aClass (2)`.

Class types allow for defining code, which is executed when values of its type are created, i.e., when objects are instantiated. This initialization code is called the *constructor*, and in contrast to many other languages, the constructor is always stated as an integral part of the class header in F# as described above. It is called the *primary constructor*, its arguments are specified in the header, and the primary constructor's body is the `let` and `do`-bindings following the header. The values and variables in the constructor are called *fields*, while functions are just called *functions*. Note that members are not available in the constructor unless the self identifier has been declared in the header using the keyword `as`, e.g., `type classMutable(name : string) as this =`

Members are declared using the `member`-keyword, which defines values and functions that are accessible from outside the class using the `."`-notation. In this manner, the members define the *interface* between the internal bindings in the constructor and an application program. Member values are called *properties* or *attributes*. Note that the concept of attributes as member values is different from the concept of functions and `let`-binding attributes, which is specified with the `[<>]` notation. For this reason, this author prefers the name member property in F#. Member functions are called *methods*.

Note that members are immutable. In the example in Listing 20.2, line 6 and 7 defines a property and a method. Properties and methods belong to objects, and the implication is the example `value` and `scale` 'resides' on or 'belongs' to each object. The body of a member has access to arguments, the primary constructor's bindings, and to all class' members, regardless of the member's lexicographical order.

In the class definition in Listing 20.2 we bind the primary constructor's arguments to the property `this.value`. The prefix `this.` is a *self identifier* used in the definition of the class such that, e.g., `this.value` is the name of the `objectValue` value for the particular object being constructed. As a quirk, F# is very flexible regarding what name can be used for the self-identifier, and the member section could as well have been `self.value`, `__.value`, or anything else, and it need not be the same in every member definition, however, **consistency in the name used as self-identifier is strongly encouraged, preferably using a name that reflects the nature of the reference, such as `this` or `me`.** Advice

As an aside, if we wanted to use a tuple argument for the class, then this must be explicitly annotated since the call to the constructor looks identical. This is demonstrated in Listing 20.3.

Listing 20.3 classTuple.fsx:

Beware: Creating objects from classes with several arguments and tuples have the same syntax.

```
1 type vectorWTupleArgs (x : float * float) =
2     member this.cartesian = x
3 type vectorWTwoArgs (x : float, y : float) =
4     member this.cartesian = (x,y)
5 let v = vectorWTupleArgs (1.0, 2.0)
6 let w = vectorWTwoArgs (1.0, 2.0)
```

Whether the full list of arguments should be transported from the caller to the object as a tuple or not is a matter of taste that mainly influences the header of the class. The same cannot be said about how the elements of the vector are stored inside the object and made accessible outside the object. In Listing 20.3, the difference between storing the vector's elements in individual members `member this.x = x` and `member this.y = y` or as a tuple `member this.cartesian = (x, y)`, is that in order to access the first element in a vector `v`, an application program in the first case must write `v.x`, while in the second case the application program must first retrieve the tuple and then extract the first element, e.g., as `fst v.cartesian`. Which is to be preferred depends very much on the application: Is it the individual elements or the complete tuple of elements that is to have focus, when using the objects. Said differently, which choice will make the easiest to read application program with the lowest risk of programming errors. Hence, when designing classes, **consider carefully how application programs will use the class and aim for simplicity and versatility while minimizing the risk of error in the application program.** Advice

20.2 Accessors

Methods are most often used as an interface between the fields of an object and the application program. Consider the example in Listing 20.4.

Listing 20.4 classAccessor.fsx:

Accessor methods interface with internal bindings.

```

1  type aClass () =
2      let mutable v = 1
3      member this.setValue (newValue : int) : unit =
4          v <- newValue
5      member this.getValue () : int = v
6
7  let a = aClass ()
8  printfn "%d" (a.getValue ())
9  a.setValue (2)
10 printfn "%d" (a.getValue ())

```

```

1  $ fsharp --nologo classAccessor.fsx && mono classAccessor.exe
2  1
3  2

```

In the example, the data contained in objects of type `aClass` is stored in the mutable field `v`. Since only members can be accessed from an application, it is not possible to retrieve or change the data of these object of class `aClass` directly. We could have programmed `v` as a member instead, i.e., `member this.v = 1`, however, often we are in the situation, where there is a range of possible choices of data representation, details of which we do wish to share with an application program. E.g., implementation details of arrays are not important for our ability to use them in applications. What matters is that the members that work on the array elements are well defined and efficient. Thus, the example demonstrates how we can build two simple methods `setValue` and `getValue` to set and get the data stored `v`. By making a distinction between the internal representation, and how members give access to the data, we retain the possibility to change the internal representation without having to reprogram all the application programs. Analogously, we can change the engine in a car from one type to another without having to change the car's interaction with the driver and the road: steering wheel, pedals, wheels etc.

Such functions are called *accessors*. Internal states with setters and getters are a typical construction, since it allows for complicated computations, when states are read to and written from, and gives the designer of the class the freedom to change the internal representation while keeping the interface the same. Accessors are so common that F# includes a special syntax for them: Classes can be made to act like variables using `member...with...and` keywords and the special function bindings `get()` and `set()` as demonstrated in Listing 20.5. · accessors

Listing 20.5 classGetSet.fsx:

Members can act as variables with the built-in getter and setter functions.

```

1  type aClass () =
2      let mutable v = 0
3      member this.value
4          with get () = v
5          and set (a) = v <- a
6
7  let a = aClass ()
8  printfn "%d" a.value
9  a.value <- 2
10 printfn "%d" a.value

```

```

1  $ fsharp -nologo classGetSet.fsx && mono classGetSet.exe
2  0
3  2

```

The expression defining `get: () -> 'a` and `set: 'a -> ()`, where `'a` is any type, can be any usual expression. The application calls the `get` and `set` as if the property were a mutable value. If `set` is omitted, then the property act as a value rather than a variable, and values cannot be assigned to it in the application program.

Setters and getters are so common that F# has a short-hand for this using `member val value = 0 with get, set`, which creates the internal mutable value `value`, but this is discouraged in this text.

Defining an *Item* property with extended `get` and `set` makes objects act as indexed variables as demonstrated in Listing 20.6.

Listing 20.6 classGetSetIndexed.fsx:

Properties can act as index variables with the built-in getter and setter functions.

```

1  type aClass (size : int) =
2      let arr = Array.create<int> size 0
3      member this.Item
4          with get (ind : int) = arr.[ind]
5          and set (ind : int) (p : int) = arr.[ind] <- p
6
7  let a = aClass (3)
8  printfn "%A" a
9  printfn "%d %d %d" a.[0] a.[1] a.[2]
10 a.[1] <- 3
11 printfn "%d %d %d" a.[0] a.[1] a.[2]

```

```

1  $ fsharp -nologo classGetSetIndexed.fsx && mono classGetSetIndexed.exe
2  ClassGetSetIndexed+aClass
3  0 0 0
4  0 3 0

```

Higher dimensional indexed properties are defined by adding more indexing arguments to the definition of `get` and `set` such as demonstrated in Listing 20.7.

Listing 20.7 classGetSetHigherIndexed.fsx:

Properties can act as index variables with the built-in getter and setter functions.

```

1 type aClass (rows : int, cols : int) =
2     let arr = Array2D.create<int> rows cols 0
3     member this.Item
4         with get (i : int, j : int) = arr.[i,j]
5         and set (i : int, j : int) (p : int) = arr.[i,j] <- p
6
7 let a = aClass (3, 3)
8 printfn "%A" a
9 printfn "%d %d %d" a.[0,0] a.[0,1] a.[2,1]
10 a.[0,1] <- 3
11 printfn "%d %d %d" a.[0,0] a.[0,1] a.[2,1]

```

```

1 $ fsharp -nologo classGetSetHigherIndexed.fsx
2 $ mono classGetSetHigherIndexed.exe
3 ClassGetSetHigherIndexed+aClass
4 0 0 0
5 0 3 0

```

20.3 Objects are reference types

Objects are reference type values, implying that copying objects copies their references not their values, and their content is stored on *The Heap*, see also Section 6.8. Consider the example in Listing 20.8.

· The Heap

Listing 20.8 classReference.fsx:

Objects are reference types means assignment is aliasing.

```

1 type aClass () =
2     let mutable v = 0
3     member this.value with get () = v and set (a) = v <- a
4
5 let a = aClass ()
6 let b = a
7 a.value <- 2
8 printfn "%d %d" a.value b.value

```

```

1 $ fsharp -nologo classReference.fsx && mono classReference.exe
2 2 2

```

Thus, the binding to `b` in line 6 is an alias to `a`, not a copy, and changing object `a` also changes `b`! This is a common cause of error, and you should **think of objects as arrays**. For this reason, it is often seen that classes implement a copy function, returning a new object with copied values, e.g., Listing 20.9.

Advice

Listing 20.9 classCopy.fsx:

A copy method is often needed. Compare with Listing 20.8.

```

1  type aClass () =
2      let mutable v = 0
3      member this.value with get () = v and set (a) = v <- a
4      member this.copy () =
5          let o = aClass ()
6          o.value <- v
7          o
8  let a = aClass ()
9  let b = a.copy ()
10 a.value <- 2
11 printfn "%d %d" a.value b.value

```

```

1  $ fsharp --nologo classCopy.fsx && mono classCopy.exe
2  2 0

```

In the example, we see that since `b` now is a copy, we do not change it by changing `a`. This is called a *copy constructor*.

· copy constructor

20.4 Static classes

Classes can act as modules and hold data, which is identical for all objects of its type. These are defined using the `static`-keyword. And since they do not belong to a single object, but are shared between all objects, they are defined without the self-identifier and accessed using the class name, and they cannot refer to the arguments of the constructor. For an example, consider a class whose objects each should hold a unique identification number (id): When an object is instantiated, the object must be given the next available identification number. The next available id could be given as an argument to the constructor, however, this delegates the task of maintaining the uniqueness of ids to the application program. Better is to use a static field and delegate the administration of ids completely to the class' and object's constructors as demonstrated in Listing 20.10.

· static

Listing 20.10 classStatic.fsx:

Static fields and members are identical to all objects of the type.

```

1  type student (name : string) =
2      static let mutable nextAvailableID = 0 // A global id for all objects
3      let studentID = nextAvailableID // A per object id
4      do nextAvailableID <- nextAvailableID + 1
5      member this.id with get () = studentID
6      member this.name = name
7      static member nextID = nextAvailableID // A global member
8  let a = student ("Jon") // Students will get unique ids, when
    instantiated
9  let b = student ("Hans")
10 printfn "%s: %d, %s: %d" a.name a.id b.name b.id
11 printfn "Next id: %d" student.nextID // Accessing the class's member

```

```

1  $ fsharp --nologo classStatic.fsx && mono classStatic.exe
2  Jon: 0, Hans: 1
3  Next id: 2

```

Notice in the example line 2, a static field `nextAvailableID` is created for the value to be shared by all objects. The initialization of its value is only performed once, at the beginning of program execution. However, every time an object is instantiated, then the value of `nextAvailableID` is copied to the object's field `studentID` in line 3, and `nextAvailableID` is updated. The static field can be accessed with a static accessor as demonstrated in line 7. Notice how this definition does not include a self-identifier, and that the member is accessible from the application in line 11 using the class' name, in both cases since it is not a member of any particular object.

20.5 Recursive members and classes

The members of a class are inherently recursive: static and non-static methods may recurse using the self identifier and other members regardless of their lexicographical scope. This is demonstrated in Listing 20.11.

Listing 20.11 classRecursion.fsx:

Members can recurse without the `rec` keyword and refer to other members regardless of their lexicographical scope.

```

1  type twice (v : int) =
2      static member fac n = if n > 1 then n * (twice.fac (n-1)) else 1 // No
    rec
3      member this.copy = this.twice // No lexicographical scope
4      member this.twice = 2*v
5
6  let a = twice (2)
7  let b = twice.fac 3
8  printfn "%A %A %A" a.copy a.twice b

```

```

1  $ fsharp --nologo classRecursion.fsx && mono classRecursion.exe
2  4 4 6

```

For mutually recursive classes, the keyword `and` must be used as shown in Listing 20.12.

· `and`

Listing 20.12 `classAssymetry.fsx`:

Mutually recursive classes are defined using the `and` keyword.

```

1  type anInt (v : int) =
2      member this.value = v
3      member this.add (w : aFloat) : aFloat = aFloat ((float this.value) +
4          w.value)
5  and aFloat (w : float) =
6      member this.value = w
7      member this.add (v : anInt) : aFloat = aFloat ((float v.value) +
8          this.value)
9  let a = anInt (2)
10 let b = aFloat (3.2)
11 let c = a.add b
12 let d = b.add a
13 printfn "%A %A %A %A" a.value b.value c.value d.value

```

```

1  $ fsharp -nologo classAssymetry.fsx && mono classAssymetry.exe
2  2 3.2 5.2 5.2

```

Here `anInt` and `aFloat` hold an integer and a floating point value respectively, and they both implement an addition of `anInt` and `aFloat` that returns an `aFloat`. Thus, they are mutually dependent and must be defined in the same `type` definition using `and`.

20.6 Function and operator overloading

It is often convenient to define different methods with the same name, but whose functionality depends on the number and type of arguments given. This is called *overloading* and F# supports method overloading. An example is shown in Listing 20.13.

Listing 20.13 classOverload.fsx:

Overloading methods `set : int -> ()` and `set : int * int -> ()` is permitted since they differ in argument number or type.

```

1  type Greetings () =
2      let mutable greetings = "Hi"
3      let mutable name = "Programmer"
4      member this.str = greetings + " " + name
5      member this.setName (newName : string) : unit =
6          name <- newName
7      member this.setName (newName : string, newGreetings : string) : unit =
8          greetings <- newGreetings
9          name <- newName
10  let a = Greetings ()
11  printfn "%s" a.str
12  a.setName ("F# programmer")
13  printfn "%s" a.str
14  a.setName ("Expert", "Hello")
15  printfn "%s" a.str

1  $ fsharp -nologo classOverload.fsx && mono classOverload.exe
2  Hi Programmer
3  Hi F# programmer
4  Hello Expert

```

In the example we define an object, which can produce greetings strings on the form `<greeting> <name>` using the `str` member. It has a default greeting “Hi” and name “Programmer”, but the name can be changed by calling the `setName` accessor with one argument, and both greeting and name can be changed by calling the overloaded `setName` with two arguments. Overloading in class definition is allowed as long as the arguments differ in number or type.

In Listing 20.12 the notation for addition is less than elegant. For such situations, F# supports *operator overloading*. All usual operators may be overloaded (see Section 6.3), and in contrast to regular operator overloading, the compiler uses type inference to decide which function is to be called. All operators have a functional equivalence, and to overload the binary “+” and unary “-” operators we overload their functional equivalence (+) and (~-) as static members. This is demonstrated in Listing 20.14.

Listing 20.14 classOverloadOperator.fsx:
Operators can be overloaded using.

```

1  type anInt (v : int) =
2      member this.value = v
3      static member (+) (v : anInt, w : anInt) = anInt (v.value + w.value)
4      static member (~-) (v : anInt) = anInt (-v.value)
5  and aFloat (w : float) =
6      member this.value = w
7      static member (+) (v : aFloat, w : aFloat) = aFloat (v.value + w.value)
8      static member (+) (v : anInt, w : aFloat) =
9          aFloat ((float v.value) + w.value)
10     static member (+) (w : aFloat, v : anInt) = v + w // reuse def. above
11     static member (~-) (v : aFloat) = aFloat (-v.value)
12
13 let a = anInt (2)
14 let b = anInt (3)
15 let c = aFloat (3.2)
16 let d = a + b // anInt + anInt
17 let e = c + a // aFloat + anInt
18 let f = a + c // anInt + aFloat
19 let g = -a // unitary minus anInt
20 let h = a + -b // anInt + unitary minus anInt
21 printf "a=%A, b=%A, c=%A, d=%A" a.value b.value c.value d.value
22 printf ", e=%A, f=%A, g=%A, h=%A" e.value f.value g.value h.value

```

```

1  $ fsharp --nologo classOverloadOperator.fsx
2  $ mono classOverloadOperator.exe
3  a=2, b=3, c=3.2, d=5, e=5.2, f=5.2, g=-2, h=-1

```

Thus, writing `v + w` is equivalent to writing `anInt.(+) (v, w)`. Presently the former is to be preferred, but at times, e.g., when using functions as arguments, it is useful to be able to refer to an operator by its function-equivalent. Note that the functional equivalence of the multiplication operator `(*)` shares a prefix with the begin block comment lexeme “`(*`”, which is why the multiplication function is written as `(*)`. Note also that unitary operators have a special notation using the “`~-`”-lexeme as illustrated in the above example for unitary minus. With the unitary minus, we are able to subtract objects of `anInt` by first negating the right-hand operand and then adding the result to the left-hand operand, thus demonstrating the difference between unary and binary minus operators, where the binary minus would have been defined as `static member (-) (v : anInt, w : aFloat) = anInt ((float v.value) - w.value)`.

In Listing 20.14, notice how the second `(+)` operator overloads the first by calling the first with the proper order of arguments. This is a general principle, **avoid duplication of code, reuse of existing code is almost always preferred**. Here it is to be preferred for two reasons. Firstly, if we discover a mistake in the multiplication code, then we need only correct it once, which implies that both multiplication methods are corrected once and reducing the chance of introducing new mistakes by attempting to correct old. Secondly, if we later decide to change the internal representation of the vector, then we only need to update one version of the multiplication function, hence we reduce programming time and risk of errors as well. Advice

Beware that operator overloading outside class definitions overwrites *all* definitions of the operator. E.g., overloading `(+)` `(v, w)` outside a class will influence integer, real, string, etc. Thus, **operator overloading should only be done inside class definitions**. Advice

20.7 Additional constructors

Like methods, constructors can also be overloaded using the `new` keyword. E.g., the example in Listing 20.13 may be modified, such that the name and possibly greeting is set at object instantiation rather than by using the accessor. This is illustrated in Listing 20.15.

Listing 20.15 `classExtraConstructor.fsx`:
Extra constructors can be added using `new`.

```

1  type classExtraConstructor (name : string, greetings : string) =
2      static let defaultGreetings = "Hello"
3      // Additional constructor are defined by new ()
4      new (name : string) =
5          classExtraConstructor (name, defaultGreetings)
6      member this.name = name
7      member this.str = greetings + " " + name
8
9  let s = classExtraConstructor ("F#") // Calling additional constructor
10 let t = classExtraConstructor ("F#", "Hi") // Calling primary
11      constructor
12 printfn "%A, %A" s.str t.str

```

```

1  $ fsharp -nologo classExtraConstructor.fsx
2  $ mono classExtraConstructor.exe
3  "Hello F#", "Hi F#"

```

The body of the additional constructor must call the primary constructor, and the body cannot extend the primary constructor's fields and functions. It is useful to **think of the primary constructor as a superset of arguments and the additional as subsets or specialisations**. As regular scope rules dictate, the additional constructor has access to the primary constructor's bindings. However, in order to access the object's members, the self identifier has to be explicitly declared using the `as`-keyword in the header. E.g., writing `new(x : float, y : float) as alsoThis = ...`. However beware, even though the body of the additional constructor now may access the property `alsoThis.x`, this value has first been created once the primary constructor has been called. E.g., calling the primary constructor in the additional constructor as `new(x : float, y : float) as alsoThis = classExtraConstructor(fst alsoThis.x, y, defaultSeparator)` will cause an exception at run-time. Code may be executed in additional constructors: Before the call to the primary constructor, `let` and `do` statements are allowed. If code is to be executed after the primary constructor has been called, then it must be preceded by the `then` keyword as shown in Listing 20.16.

Advice

Listing 20.16 classDoThen.fsx:

The optional `do`- and `then`-keywords allows for computations before and after the primary constructor is called.

```

1 type classDoThen (aValue : float) =
2     // "do" is mandatory to execute code in the primary constructor
3     do printfn "    Primary constructor called"
4     // Some calculations
5     do printfn "    Primary done" (* *)
6     new () =
7         // "do" is optional in additional constructors
8         printfn "    Additional constructor called"
9         classDoThen (0.0)
10        // Use "then" to execute code after construction
11        then
12            printfn "    Additional done"
13    member this.value = aValue
14
15 printfn "Calling additional constructor"
16 let v = classDoThen ()
17 printfn "Calling primary constructor"
18 let w = classDoThen (1.0)

```

```

1 $ fsharpc --nologo classDoThen.fsx && mono classDoThen.exe
2 Calling additional constructor
3     Additional constructor called
4     Primary constructor called
5     Primary done
6     Additional done
7 Calling primary constructor
8     Primary constructor called
9     Primary done

```

The `do`-keyword is often understood to be implied by F#, e.g., in front of all `printf`-statements, but in the above examples they are required where used. This may change in future releases of F#. F# allows for many additional constructors, but they must be distinguishable by type.

20.8 Interfacing with printf family

In previous examples, we accessed the property in order to print the content of the objects. Luckily, a more elegant solution is available. Objects can be printed directly, but the result is most often not very useful as can be seen in Listing 20.17.

Listing 20.17 classPrintf.fsx:

Printing classes yields low-level information about the class.

```

1 type vectorDefaultToString (x : float, y : float) =
2     member this.x = (x,y)
3
4 let v = vectorDefaultToString (1.0, 2.0)
5 printfn "%A" v // Printing objects gives lowlevel information

```

```

1 $ fsharp --nologo classPrintf.fsx && mono classPrintf.exe
2 ClassPrintf+vectorDefaultToString

```

All classes are given default members through a process called *inheritance*, to be discussed below in Section 21.1. One example is the `ToString() : () -> string` function, which is useful in conjunction with, e.g., `printf`. When an object is given as argument to a `printf` function, then `printf` calls the object's `ToString()` function. The default implementation returns low-level information about the object, as can be seen above, but we may *override* this member using the *override*-keyword as demonstrated in Listing 20.18.¹

Listing 20.18 classToString.fsx:

Overriding `ToString()` function for better interaction with members of the `printf` family of procedures. Compare with Listing 20.17.

```

1 type vectorWToString (x : float, y : float) =
2     member this.x = (x,y)
3     // Custom printing of objects by overriding this.ToString()
4     override this.ToString() =
5         sprintf "(%A, %A)" (fst this.x) (snd this.x)
6
7 let v = vectorWToString(1.0, 2.0)
8 printfn "%A" v // No change in application but result is better

```

```

1 $ fsharp --nologo classToString.fsx && mono classToString.exe
2 (1.0, 2.0)

```

We see that as a consequence, the `printf` statement is much simpler. However beware, an application program may require other formatting choices than selected at the time of designing the class, e.g., in our example the application program may prefer square brackets as delimiters for vector tuples. So in general **when designing an override to `ToString()`, choose simple, generic formatting for the widest possible use.**

Advice

The most generic formatting is not always obvious, and in the vector case some candidates for the formatting string of `ToString()` are `"%A %A"`, `"%A, %A"`, `"(%A, %A)"`, or `"[%A, %A]"`. Considering each carefully it seems that arguments can be made against all. A common choice is to let the formatting be controlled by static members that can be changed by the application program by accessors.

¹Jon: something about `ToString` not working with 's' format string in `printf`.




Figure 20.1: Illustration of vector addition in two dimensions.

20.9 Programming intermezzo

Consider the following problem.

Problem 20.1

A Euclidean vector is a geometric object that has a direction and a length and two operations: vector addition and scalar multiplication, see Figure 20.1. Define a class for a vector in two dimensions.

An essential part in designing a solution for the above problem is to decide, which representation to use internally for vectors. The Cartesian representation of a vector is as a tuple of real values (x, y) , where x and y are real values, and where we can imagine that the tail of the vector is in the origin, and its tip is at the coordinate (x, y) . For vectors on Cartesian form,

$$\vec{v} = (x, y), \quad (20.1)$$

the basic operations are defined as

$$\vec{v}_1 + \vec{v}_2 = (x_1 + x_2, y_1 + y_2), \quad (20.2)$$

$$a\vec{v} = (ax, ay), \quad (20.3)$$

$$\text{dir}(\vec{v}) = \tan \frac{y}{x}, \quad x \neq 0, \quad (20.4)$$

$$\text{len}(\vec{v}) = \sqrt{x^2 + y^2}, \quad (20.5)$$

where x_i and y_i are the elements of vector \vec{v}_i , a is a scalar, and dir and len are the direction and length functions. The polar representation of vectors is also a tuple of real values (θ, l) , where θ and l are the vector's direction and length. This representation is closely tied to the definition of a vector, and with the constraint that $0 \leq \theta < 2\pi$ and $0 \leq l$. This representation reminds us that vectors do not have a position. For vectors on polar form,

$$\vec{v} = (\theta, l), \quad (20.6)$$

their basic operations are defined as

$$x(\theta, l) = l \cos(\theta), \quad (20.7)$$

$$y(\theta, l) = l \sin(\theta), \quad (20.8)$$

$$\vec{v}_1 + \vec{v}_2 = (x(\theta_1, l_1) + x(\theta_2, l_2), y(\theta_1, l_1) + y(\theta_2, l_2)) \quad (20.9)$$

$$a\vec{v} = (\theta, al), \quad (20.10)$$

where θ_i and l_i are the elements of vector \vec{v}_i , a is a scalar, and x and y are the Cartesian coordinate functions.

So far in our analysis, we have realized that:

- both the Cartesian and polar representation uses a pair of reals to represent the vector,
- both require functions to calculate the elements of the other representation,
- the polar representation is invalid for negative lengths, and
- the addition operator under the polar representation is also more complicated and essentially requires access to the Cartesian representation.

The first step in shaping our solution is to decide on file structure: For conceptual separation, we choose to use a library and an application file. F# wants files to define namespaces or modules, so we choose the library to be a `Geometry` module, which implements the vector class to be called `vector`. Further, when creating vector objects, we would like to give the application program the ability to choose either Cartesian or polar form. This can be done using *discriminated unions*. Discriminated unions allow us to tag values of possibly identical form, but they also implied longer programs. Thus, we will also provide an additional constructor on implicit Cartesian form, since this is the most common representation.

· discriminated unions

A key point, when defining libraries, is to consider their interface with the application program. Hence, our second step is to write an application using the yet to be written library in order to get a feel for how such an interface could be. This is demonstrated in the application program Listing 20.19.

Listing 20.19 vectorApp.fsx:

An application using the library in Listing 20.20.

```
1 open Geometry
2 let v = vector(Cartesian (1.0,2.0))
3 let w = vector(Polar (3.2,1.8))
4 let p = vector()
5 let q = vector(1.2, -0.9)
6 let a = 1.5
7 printfn "%A * %A = %A" a v (a * v)
8 printfn "%A + %A = %A" v w (v + w)
9 printfn "vector() = %A" p
10 printfn "vector(1.2, -0.9) = %A" q
11 printfn "v.dir = %A" v.dir
12 printfn "v.len = %A" v.len
```

The application of the vector class seems natural, makes use of the optional discriminated unions, and uses the infix operators “+” and “*” in a manner close to standard arithmetic, and interacts smoothly with the `printf` family. Thus, we have further sketched requirements to the library with the emphasis on application.

After a couple of trials, our library implementation has ended up as shown in Listing 20.20.

Listing 20.20 vector.fs:
A library serving the application in Listing 20.21.

```

1 module Geometry
2 type Coordinate =
3   Cartesian of float * float // (x, y)
4   | Polar of float * float // (dir, len)
5 type vector(c : Coordinate) =
6   let (_x, _y, _dir, _len) =
7     match c with
8     | Cartesian (x, y) ->
9       (x, y, atan2 y x, sqrt (x * x + y * y))
10    | Polar (dir, len) when len >= 0.0 ->
11      (len * cos dir, len * sin dir, dir, len)
12    | Polar (dir, _) ->
13      failwith "Negative length in polar representation."
14 new(x : float, y : float) =
15   vector(Cartesian (x, y))
16 new() =
17   vector(Cartesian (0.0, 0.0))
18 member this.x = _x
19 member this.y = _y
20 member this.len = _len
21 member this.dir = _dir
22 static member val left = "(" with get, set
23 static member val right = ")" with get, set
24 static member val sep = ", " with get, set
25 static member ( * ) (a : float, v : vector) : vector =
26   vector(Polar (v.dir, a * v.len))
27 static member ( * ) (v : vector, a : float) : vector =
28   a * v
29 static member (+) (v : vector, w : vector) : vector =
30   vector(Cartesian (v.x + w.x, v.y + w.y))
31 override this.ToString() =
32   sprintf "%s%s%s%s" vector.left this.x vector.sep this.y
   vector.right

```

Realizations achieved during writing this code are: Firstly, in order to implement a vector class using discriminated unions, we had to introduce a constructor with helper variables `_x`, `_y`, etc. The consequence is that the Cartesian and polar representation is evaluated once and only once every time an object is created. Unfortunately, discriminated unions do not implement guards on subsets, so we still have to cast an exception, when the application attempts to create an object with a negative length. Secondly, for the `ToString` override we have implemented static members for typesetting vectors since it seems more appropriate that all vectors should be typeset identically. Changing typesetting thus respect dynamic scope.

The output of our combined library and application is shown in Listing 20.21.

Listing 20.21: Compiling and running the code from Listing 20.20 and 20.19.

```
1 $ fsharpc --nologo vector.fs vectorApp.fsx && mono vectorApp.exe
2 1.5 * (1.0, 2.0) = (1.5, 3.0)
3 (1.0, 2.0) + (-1.796930596, -0.1050734582) = (-0.7969305964, 1.894926542)
4 vector() = (0.0, 0.0)
5 vector(1.2, -0.9) = (1.2, -0.9)
6 v.dir = 1.107148718
7 v.len = 2.236067977
```

The output is as expected and for the vector class, our solution seems to be a good compromise between versatility and syntactical bloating.

21 | Derived classes

21.1 Inheritance

Sometimes it is useful to derive new classes from old in order to reuse code or to emphasize a program structure. For example, consider the concepts of a *car* and *bicycle*. They are both *vehicles* that can move forward and turn, but a car can move in reverse, has 4 wheels uses gasoline or electricity, while a bicycle has 2 wheels and needs to be pedaled. Structurally we can say that “a car is a vehicle” and “a bicycle is a vehicle”. Such a relation is sometimes drawn as a tree as shown in Figure 21.1 and is called an *is-a relation*. Is-a relations can be implemented using class *inheritance*, where vehicle is called the *base class* and car and bicycle are each a *derived class*. The advantage is that a derived class can inherit the members of the base class, *override* and add possibly new members. Inheritance is indicated using the `inherit` keyword. Listing 20.1 shows the syntax for class definitions using inheritance, and an example of defining base and derived classes for vehicles is shown In Listing 21.1.




Figure 21.1: Both a car and a bicycle is a (type of) vehicle.

Listing 21.1 vehicle.fsx:

New classes can be derived from old.

```

1  /// A general vehicle, which moves on a plane and has a heading
2  type vehicle () =
3      let mutable p = (0.0, 0.0) // coordinate on a plane
4      let mutable d = 0.0 // heading direction in radians
5      member this.dir with get() = d
6      member this.pos with get() = p and set aPos = p <- aPos
7      member this.turn angle = // turn heading
8          d <- d + angle
9      member this.forward step = // move forward (abs step)
10         let s = abs step
11         let vec = (s * (cos d), s * (sin d))
12         p <- (fst p + fst vec, snd p + snd vec)
13  /// A car is a vehicle, has wheels and can move in reverse
14  type car (name) =
15      inherit vehicle () // inherit dir, pos, turn, and forward
16      member this.wheels = 4 // A car has 4 wheels
17      member this.reverse step = // move backwards (abs step)
18         let s = - abs step
19         let vec = (s * (cos this.dir), s * (sin this.dir))
20         this.pos <- (fst this.pos + fst vec, snd this.pos + snd vec)
21  /// A bicycle is a vehicle and has wheels
22  type bicycle () =
23      inherit vehicle () // inherit dir, pos, turn, and forward
24      member this.wheels = 2 // A bike has 4 wheels
25
26  let aVehicle = vehicle () // has dir, pos, turn, forward
27  let aCar = car () // has dir, pos, turn, forward, wheels, reverse
28  let aBike = bicycle () // has dir, pos, turn, forward, wheels
29  printfn "The car aCar has %d wheels" aCar.wheels
30  printfn "The bicycle aBike has %d wheels" aBike.wheels

```

```

1  $ fsharp --nologo vehicle.fsx && mono vehicle.exe
2  The car aCar has 4 wheels
3  The bicycle aBike has 2 wheels

```

In the example, a base class `vehicle` is defined with members `dir`, `pos`, `turn`, and `forward`. The derived classes inherit all the members of the base class, but do not have access to any non-members of the base's constructor. I.e., `car` and `bicycle` automatically have methods `turn` and `forward`, and properties `dir` and `pos` with their accessors, but they do not have direct access to the fields `p` and `d`. Both derived classes additionally define a property `wheels` and `car` also define a method `reverse`. Note that inheritance is one-way, and in spite that both derived classes define a member `wheels`, the base class does not have a `wheels` member.

Derived classes can replace base class members by defining new members *overshadow* the base' member. · overshadow
The base' members are still available using the `base`-keyword. Consider the example in the Listing 21.2. · base

Listing 21.2 memberOvershadowing.fsx:

Inherited members can be overshadowed, but we can still access the base' member.

```

1  /// A counter has an internal state initialized at instantiation and
2  /// is incremented in steps of 1
3  type counter (init : int) =
4      let mutable i = init
5      member this.value with get () = i and set (v) = i <- v
6      member this.inc () = i <- i + 1
7  /// A counter2 is a counter which increments in steps of 2.
8  type counter2 (init : int) =
9      inherit counter (init)
10     member this.inc () = this.value <- this.value + 2
11     member this.incByOne () = base.inc () // inc by 1 implemented in base
12
13 let c1 = counter (0) // A counter by 1 starting with 0
14 printf "c1: %d" c1.value
15 c1.inc() // inc by 1
16 printfn " %d" c1.value
17 let c2 = counter2 (1) // A counter by 2 starting with 1
18 printf "c2: %d" c2.value
19 c2.inc() // inc by 2
20 printf " %d" c2.value
21 c2.incByOne() // inc by 1
22 printfn " %d" c2.value

```

```

1  $ fsharp --nologo memberOvershadowing.fsx
2  $ mono memberOvershadowing.exe
3  c1: 0 1
4  c2: 1 3 4

```

In this case, we have defined two counters, each with an internal field `i` and with members `value` and `inc`. The `inc` method in `counter` increments `i` with 1, and in `counter2` the field `i` is incremented with 2. Note how `counter2` inherits both members `value` and `inc`, but overshadows `inc` by defining its own. Note also how `counter2` defines another method `incByOne` by accessing the inherited `inc` method using the `base` keyword.

Even though derived classes are different from their base, the derived class includes the base class, which can be recalled using *upcasting* by the upcast operator “`:>`”. At compile-time this operator removes the additions and overshadowing of the derived class, as illustrated in Listing 21.3.

· upcast
· :>

Listing 21.3 upCasting.fsx:

Objects can be upcasted resulting in an object as if it were its base. Implementations from the derived class are ignored.

```

1  /// hello holds property str
2  type hello () =
3      member this.str = "hello"
4  /// howdy is a hello class and has property altStr
5  type howdy () =
6      inherit hello ()
7      member this.str = "howdy"
8      member this.altStr = "hi"
9
10 let a = hello ()
11 let b = howdy ()
12 let c = b :> hello // a howdy object as if it were a hello object
13 printfn "%s %s %s %s" a.str b.str b.altStr c.str

```

```

1  $ fsharp --nologo upCasting.fsx && mono upCasting.exe
2  hello howdy hi hello

```

Here `howdy` is derived from `hello`, overshadows `str`, and adds property `altStr`. By upcasting object `b`, we create object `c` as a copy of `b` with all its fields, functions, and members as if it had been of type `hello`. I.e., `c` contains the base class version of `str` and does not have property `altStr`. Objects `a` and `c` are now of same type and can be put into, e.g., an array as `let arr = [|a, c|]`. Previously upcasted objects can also be downcasted again using the *downcast* operator `:?>`, but the validity of the operation is checked at runtime. Thus, **avoid downcasting when possible**.

· downcast
· :?>
Advice

In the above, inheritance is used to modify and extend any class. I.e., the definition of the base classes were independent on the definitions of inherited classes. In that sense, the base classes were oblivious to any future derivation of them. Sometimes it is useful to define base classes, which are not independent on derived classes, and which impose design constraints on derived classes. Two such dependencies in F# are abstract classes and interfaces to be described in the following sections.

21.2 Abstract class

An *abstract class* contains members defined using the *abstract member* and optionally the *default* keywords. An *abstract member* in the base class is a type definition, and derived classes must provide an implementation using the *override* keyword. Optionally, the base class may provide a default implementation using the *default* keyword, in which case overriding is not required in derived classes. Objects of classes containing abstract members without default implementations cannot be instantiated, but derived classes that provide the missing implementations can be. Note that abstract classes must be given the `[<AbstractClass>]` attribute. Note also that in contrast to overshadowing, upcasting keeps the implementations of the derived classes. Examples of this are shown in Listing 21.4.

· abstract class
· *abstract member*
· *default*
· *override*
· [`<AbstractClass>`]

Listing 21.4 abstractClass.fsx:

In contrast to regular objects, upcasted derived object use the derived implementation of abstract methods.

```

1  /// An abstract class for general greeting classes with property str
2  [<AbstractClass>]
3  type greeting () =
4      abstract member str : string
5  /// hello is a greeting
6  type hello () =
7      inherit greeting ()
8      override this.str = "hello"
9  /// howdy is a greeting
10 type howdy () =
11     inherit greeting ()
12     override this.str = "howdy"
13
14 let a = hello ()
15 let b = howdy ()
16 let c = [| a :> greeting; b :> greeting |] // arrays of greetings
17 Array.iter (fun (elm : greeting) -> printfn "%s" elm.str) c

```

```

1  $ fsharp --nologo abstractClass.fsx && mono abstractClass.exe
2  hello
3  howdy

```

In the example, we define a base class and two derived classes. Note how the abstract member is defined in the base class using the “:”-operator as a type declaration rather than a name binding. Note also that since the base class does not provide a default implementation, the derived classes supply an implementation using the `override`-keyword. In the example, objects of `baseClass` cannot be created, since such objects would have no implementation for `this.hello`. Finally, the two different derived and upcasted objects can be put in the same array, and when calling their implementation of `this.hello` we still get the derived implementations, which is in contrast to overshadowing.

Abstract classes may also specify a default implementation, such that derived classes have the option of implementing an overriding member, but are not forced to. In spite that implementations are available in the abstract class, the abstract class still cannot be used to instantiate objects. Such a variant is shown in Listing 21.5.

Listing 21.5 abstractDefaultClass.fsx:

Default implementations in abstract classes makes implementations in derived classes optional. Compare with Listing 21.4.

```

1  /// An abstract class for general greeting classes with property str
2  [<AbstractClass>]
3  type greeting () =
4      abstract member str : string
5      default this.str = "hello" // Provide default implementation
6  /// hello is a greeting
7  type hello () =
8      inherit greeting ()
9  /// howdy is a greeting
10 type howdy () =
11     inherit greeting ()
12     override this.str = "howdy"
13
14 let a = hello ()
15 let b = howdy ()
16 let c = [| a :> greeting; b :> greeting |] // arrays of greetings
17 Array.iter (fun (elm : greeting) -> printfn "%s" elm.str) c

```

```

1  $ fsharp -nologo abstractDefaultClass.fsx
2  $ mono abstractDefaultClass.exe
3  hello
4  howdy

```

In the example, the program in Listing 21.4 has been modified such that `greeting` is given a default implementation for `str`, in which case, `hello` does not need to supply one. However, in order for `howdy` to provide a different greeting, it still needs provide an override member.

Note that even if all abstract members in an abstract class has defaults, objects of its type can still not be created, but must be derived as, e.g., shown with `hello` above.

As a side note, every class implicitly derives from a base class `System.Object` which, which is an `System.Object` abstract class defining among other members the `ToString` method with default implementation.

21.3 Interfaces

Inheritance of an abstract base class allows an application to rely on the definition of the base regardless of any future derived classes. This gives great flexibility, but at times even less knowledge is needed about objects in order to write useful applications. This is what *interfaces* offer. An interface specifies which members must exist but nothing more. Interfaces are defined as an abstract class *without arguments* and *only with abstract members*. Classes implementing interfaces must specify implementations for the abstract members using the `interface with` keywords. Objects of classes implementing interfaces can be upcasted as if they had an abstract base class of the interface' name. Consider the example in Listing 21.6.

Listing 21.6 classInterface.fsx:

Interfaces specifies which members classes contain, and with upcasting gives more flexibility than abstract classes.

```

1  /// An interface for classes that have method fct and member value
2  type IValue =
3      abstract member fct : float -> float
4      abstract member value : int
5  /// A house implements the IValue interface
6  type house (floors: int, baseArea: float) =
7      interface IValue with
8          // calculate total price based on per area average
9          member this.fct (pricePerArea : float) =
10             pricePerArea * (float floors) * baseArea
11          // return number of floors
12          member this.value = floors
13  /// A person implements the IValue interface
14  type person(name : string, height: float, age : int) =
15      interface IValue with
16          // calculate body mass index (kg/(m*m)) using hypothetic mass
17          member this.fct (mass : float) = mass / (height * height)
18          // return the length of name
19          member this.value = name.Length
20          member this.data = (name, height, age)
21
22  let a = house(2, 70.0) // a two storage house with 70 m*m base area
23  let b = person("Donald", 1.8, 50) // a 50 year old person 1.8 m high
24  let lst = [a :> IValue; b :> IValue]
25  let printInterfacePart (o : IValue) =
26      printfn "value = %d, fct(80.0) = %g" o.value (o.fct 80.0)
27  List.iter printInterfacePart lst

```

```

1  $ fsharpc --nologo classInterface.fsx && mono classInterface.exe
2  value = 2, fct(80.0) = 11200
3  value = 6, fct(80.0) = 24.6914

```

Here, two distinctly different classes are defined: `house` and `person`. These are not related by inheritance since no sensible common structure seems available. However, they share structures in the sense that they both have an integer property and a `float -> float` method. For each of the derived classes, these members have different meanings. Still, some treatment of these members by an application will only rely on their type and not their meaning. E.g., in Listing 21.6 the `printfn` function only needs to know the member's type not their meaning. As a consequence, the application can upcast them both to the implicit abstract base class `IValue`, put them in an array, and apply a function using the member definition of `IValue` with the higher-order `List.iter` function. Another example could be a higher order function calculating average values: For average values of the number of floors and average value of the length of people's names, the higher order function would only need to know that both of these classes implements the `IValue` interfaces in order to calculate the average of list of either objects types.

As a final note, inheritance ties classes together in a class hierarchy. Abstract members enforce inheritance and impose constraints on the derived classes. Like abstract classes, interfaces impose constraints on derived classes, but without requiring a hierarchical structure.




Figure 21.2: Starting position for the game of chess.

21.4 Programming intermezzo: Chess

To demonstrate the use of hierarchies, consider the following problem

Problem 21.1

The game of chess is a turn-based game for two, which consists of a board of 8×8 squares and a set of 16 black and 16 white pieces. A piece can be either a king, queen, rook, bishop, knight or pawn and each piece has a specific movement pattern on the board. Pieces are added to, moved on, and removed from the board during the game, and there can be at most one piece per square. A piece strikes another piece of opposing color by moving to its square and the piece of opposing color is removed from the game. The game starts with the configuration shown in Figure 21.2.

Make a program that allows two humans to play simple chess using only kings and rooks. The king must be able to move to all neighboring squares not occupied by a piece of the same color and cannot move onto a square, where it can be struck in the next turn. The rook must be able to move in horizontal and vertical lines until a piece of the same color or up to and including a piece of opposing color. Make a program that allows two humans to play simple chess.

Since we expect that the solution to the above problem is going to be a relatively long program, we have decided to split the code into a library and an application program. Before writing a library, it is often useful to start thinking about how the library should be used. Thus we start by sketching the application program, and in the process consider options for the main methods and properties to be used.

We also foresee future extensions to include more pieces, but also that these pieces will obey the same game mechanics that we design for the present problem. Thus, we will put the main part of the library in a file defining the module called **Chess** and the derived pieces in another file defining the module **Pieces**.

Every game needs a board, and we will define a class **Board**. A board is like an array, so it seems useful to be able to move pieces by index notation. Thus, the board must have a two-dimensional **Item** property. We also decide that each position will hold an option type such when a square is empty it holds **None** otherwise it holds piece **p** as **Some p**. Although chess notation would be neat, for ease of programming we will let index (0,0) correspond to position a1 in chess notation etc. The most common operation will probably be to move pieces around, so we will give the board a **move** method. We will most likely also like to print the board with pieces in their right locations. For simplicity we choose

to override the `ToString` method in `Board`, and that this method also prints information about each individual piece such as where it is, where it can move to, and which pieces it can either protect or hit. The pieces that a piece can protect or hit we will call the piece's neighbor pieces.

A piece can be one of several types, so this gives a natural hierarchical structure, which is well suited for inheritance. Each piece must be given a color, which may conveniently be given as argument at instantiation. Thus, we have decided to make a base class called `chessPiece` with argument `Color`, and derived classes `king` and `rook`. The color may conveniently define as a discriminated union type of either `White` or `Black`. Each piece will also override the `ToString` method for ease of printing. The override will be used in conjunction with the board's override, so it should only give information about the piece's type and color. For compact printing, we will use a single letter for the type of piece, upper case if white, and lower case if black. We expect the pieces also to need to know something about the relation to board, so we will make a `position` property, which holds the coordinates of the piece, and we will make a `availableMoves` method that lists the possible moves, a piece can make. Thus, we produce the application in Listing 21.7, and an illustration of what the program should do is shown in Figure 21.3.

Listing 21.7 chessApp.fsx:
A chess application.

```

1  open Chess
2  open Pieces
3  /// Print various information about a piece
4  let printPiece (board : Board) (p : chessPiece) : unit =
5      printfn "%A: %A %A" p p.position (p.availableMoves board)
6
7  // Create a game
8  let board = Chess.Board () // Create a board
9  // Pieces are kept in an array for easy testing
10 let pieces = [
11     king (White) :> chessPiece;
12     rook (White) :> chessPiece;
13     king (Black) :> chessPiece ]
14 // Place pieces on the board
15 board.[0,0] <- Some pieces.[0]
16 board.[1,1] <- Some pieces.[1]
17 board.[4,1] <- Some pieces.[2]
18 printfn "%A" board
19 Array.iter (printPiece board) pieces
20
21 // Make moves
22 board.move (1,1) (3,1) // Moves a piece from (1,1) to (3,1)
23 printfn "%A" board
24 Array.iter (printPiece board) pieces

```

At this point, we are fairly happy with the way the application is written. The double bookkeeping of pieces in an array and on the board seems a bit excessive, but for testing, it seems useful to be able to easily access all pieces both those in play and struck. Although the `position` property of a `chessPiece` could be replaced by a function searching for a specific piece on the board, we have a hunch that we will need to retrieve a piece's position often, and that this double will most likely save execution time later.

Continuing our outer to inner approach, as a second step, we consider the specific pieces: They will inherit a base piece and implement the details that are special for that piece. Each piece is signified by its color and its type, and each type has a specific motion pattern. Since we have already decided to use discriminated unions for the color, it seems natural to let the color be part of the constructor of




Figure 21.3: Starting at the left and moving white rook to b4.

the base class. As in the example application in Listing 21.7, pieces are upcasted to `chessPiece`, then the base class must know how to print the piece type. For this, we will define an abstract property, such that everything needed for overriding `ToString` is available to the base class, but also such that the name of the type of the piece is set in the derived class.

For a piece on the board, its available moves depend on its type and the other pieces. The application program will need to make a decision on whether to move the piece depending on which vacant squares, it can move to, and its relation to its neighbors, i.e., is the piece protecting one of its own color, or does it have the opportunity to hit an opponent. Thus given the board with all the pieces, it seems useful that `availableMoves` returns two lists: a list of vacant squares and a list of neighboring pieces of either color. Each piece has certain movement pattern, which we will specify regardless of the piece's position on the board and relation to other pieces. Thus, this will be an abstract member called `candidateRelativeMoves` implemented in the derived pieces. These candidate relative moves are then to be sifted for legal moves, and the process will be the same for all pieces, which thus can be implemented in the base class as the `availableMoves`.

Many pieces move in runs, e.g., the rook can move horizontally and vertically until there is another piece. Vacant squares behind the blocking piece are unavailable. For a rook, we thus must analyze four runs: northward, eastward, southward, and westward. For each run, we must consult the board to see, how many vacant fields there are in that direction, and which is the piece blocking if any. Thus, we decide that the board must have a function that can analyze a list of runs and that the result is concatenated into a single list of vacant squares and a single list of neighboring pieces if any. This function we call `getVacentNNeighbours`. And so we arrive at Listing 21.8.

Listing 21.8 pieces.fs:
An extension of chess base.

```

1 module Pieces
2 open Chess
3 /// A king is a chessPiece which moves 1 square in any direction
4 type king(col : Color) =
5   inherit chessPiece(col)
6   override this.nameOfType = "king"
7   // king has runs of 1 in 8 directions: (N, NE, E, SE, S, SW, W, NW)
8   override this.candidateRelativeMoves =
9     [[(-1,0)];[(-1,1)];[(0,1)];[(1,1)];
10      [(1,0)];[(1,-1)];[(0,-1)];[(-1,-1)]]
11 /// A rook is a chessPiece which moves horizontally and vertically
12 type rook(col : Color) =
13   inherit chessPiece(col)
14   // rook can move horizontally and vertically
15   // Make a list of relative coordinate lists. We consider the
16   // current position and try all combinations of relative moves
17   // (1,0); (2,0) ... (7,0); (-1,0); (-2,0); ...; (0,-7).
18   // Some will be out of board, but will be assumed removed as
19   // illegal moves.
20   // A list of functions for relative moves
21   let indToRel = [
22     fun elm -> (elm,0); // South by elm
23     fun elm -> (-elm,0); // North by elm
24     fun elm -> (0,elm); // West by elm
25     fun elm -> (0,-elm) // East by elm
26   ]
27   // For each function in indToRel, we calculate List.map f [1..7].
28   // swap converts (List.map fct indices) to (List.map indices fct).
29   let swap f a b = f b a
30   override this.candidateRelativeMoves =
31     List.map (swap List.map [1..7]) indToRel
32   override this.nameOfType = "rook"

```

The king has the simplest relative movement candidates being the hypothetical eight neighboring squares. For rooks, the relative movement candidates are somewhat more complicated. For rooks, we would like to use `List.map` to convert a list of single indices into double indices to calculate each run. And we have gathered all the elemental functions for this in `indToRel`. E.g., function at index 0, we may write `List.map indToRel.[0] indices`. However, we would also like to use `List.map` to perform this operation for all elemental functions in `indToRel`. Direct joining such two applications of `List.map` does not work, since `List.map` takes a function and a list as its arguments, and for the second application, these two arguments should switch order. I.e., the first time it is `indices` that takes the role of the list, while the second it is `indToRel` that takes the role of the list. A standard solution in functional programming is to use currying and the `swap` function as illustrated in line 31: `· swap`. The function is equivalent to the anonymous function `fun elm -> swap List.map indices elm`, and since `swap` swaps the arguments of a function, this reduces to `fun elm -> List.map elm indices`, which is exactly what is needed.

The final step will be to design the `Board` and `chessPiece` classes. The `Chess` module implements discriminated unions for color and an integer tuple for a position. These are shown in Listing 21.9.

Listing 21.9 chess.fs:

A chess base: Module header and discriminated union types.

```

1 module Chess
2 type Color = White | Black
3 type Position = int * int

```

The `chessPiece` will need to know what a board is, so we must define it as a mutually recursive class with `Board`. Further, since all pieces must supply an implementation of `availableMoves`, we set it to be abstract by the abstract class attribute and with an abstract member. The board will need to be able to ask for a string describing each piece and to keep the board on the screen we include an abbreviated description of the piece's properties color and piece type. The result is shown in Listing 21.10.

Listing 21.10 chess.fs:

A chess base. Abstract type `chessPiece`.

```

4 /// An abstract chess piece
5 [<AbstractClass>]
6 type chessPiece(color : Color) =
7     let mutable _position : Position option = None
8     abstract member nameOfType : string // "king", "rook", ...
9     member this.color = color // White, Black
10    member this.position // E.g., (0,0), (3,4), etc.
11        with get() = _position
12        and set(pos) = _position <- pos
13    override this.ToString () = // E.g. "K" for white king
14        match color with
15        | White -> (string this.nameOfType.[0]).ToUpper ()
16        | Black -> (string this.nameOfType.[0]).ToLower ()
17    /// A list of runs, which is a list of relative movements, e.g.,
18    /// [(1,0); (2,0);...]; [(-1,0); (-2,0)]...]. Runs must be
19    /// ordered such that the first in a list is closest to the piece
20    /// at hand.
21    abstract member candiateRelativeMoves : Position list list
22    /// Available moves and neighbours [(1,0); (2,0);...], [p1; p2])
23    member this.availableMoves (board : Board) : (Position list *
24        chessPiece list) =
25        board.getVacantNNeighbours this

```

Our `Board` class is by far the largest and will be discussed by Listing 21.11–21.13. The constructor is shown in Listing 21.11.

Listing 21.11 chess.fs:
A chess base: the constructor

```

25  /// A board
26  and Board () =
27    let _array = Collections.Array2D.create<chessPiece option> 8 8 None
28    /// Wrap a position as option type
29    let validPositionWrap (pos : Position) : Position option =
30      let (rank, file) = pos // square coordinate
31      if rank < 0 || rank > 7 || file < 0 || file > 7
32      then None
33      else Some (rank, file)
34    /// Convert relative coordinates to absolute and remove out
35    /// of board coordinates.
36    let relativeToAbsolute (pos : Position) (lst : Position list) :
      Position list =
37      let addPair (a : int, b : int) (c : int, d : int) : Position =
38        (a+c,b+d)
39      // Add origin and delta positions
40      List.map (addPair pos) lst
41      // Choose absolute positions that are on the board
42      |> List.choose validPositionWrap

```

For memory efficiency, the board has been implemented using a `Array2D`, since pieces will move around often. For later use in the members shown in Listing 21.13 we define two functions that convert relative coordinates into absolute coordinates on the board, and removes those that fall outside the board. These are called `validPositionWrap` and `relativeToAbsolute`.

For ease of use in an application, `Board` implements `Item`, such that the board can be read and writing to using array notation. And `ToString` is overridden, such that an application may print the board anytime using a `printf` function. This is shown in Listing 21.12.

Listing 21.12 chess.fs:

A chess base: Board header, constructor, and non-static members.

```

43  /// Board is indexed using .[,] notation
44  member this.Item
45      with get(a : int, b : int) = _array.[a, b]
46      and set(a : int, b : int) (p : chessPiece option) =
47          if p.IsSome then p.Value.position <- Some (a,b)
48          _array.[a, b] <- p
49  /// Produce string of board for, e.g., the printfn function.
50  override this.ToString() =
51      let rec boardStr (i : int) (j : int) : string =
52          match (i,j) with
53              (8,0) -> ""
54              | _ ->
55                  let stripOption (p : chessPiece option) : string =
56                      match p with
57                          None -> ""
58                          | Some p -> p.ToString()
59                  // print top to bottom row
60                  let pieceStr = stripOption _array.[7-i,j]
61                  //let pieceStr = sprintf "(%d, %d)" i j
62                  let lineSep = " " + String.replicate (8*4-1) "-"
63                  match (i,j) with
64                      (0,0) ->
65                          let str = sprintf "%s\n| %1s " lineSep pieceStr
66                          str + boardStr 0 1
67                      | (i,7) ->
68                          let str = sprintf "| %1s |\n%s\n" pieceStr lineSep
69                          str + boardStr (i+1) 0
70                      | (i,j) ->
71                          let str = sprintf "| %1s " pieceStr
72                          str + boardStr i (j+1)
73      boardStr 0 0

```

Note that for efficiency, location is also stored in each piece, so `set` also needs to update the particular piece' position as done in line 47. Note also that the board is printed with the first coordinate of the board being rows and second columns and such that element (0,0) is at the bottom right complying with standard chess notation.

The main computations are done in the static methods of the board as shown in Listing 21.13.

Listing 21.13 chess.fs:

A chess base: Board static members.

```

74  /// Move piece by specifying source and target coordinates
75  member this.move (source : Position) (target : Position) : unit =
76      this.[fst target, snd target] <- this.[fst source, snd source]
77      this.[fst source, snd source] <- None
78  /// Find the tuple of empty squares and first neighbour if any.
79  member this.getVacantNOccupied (run : Position list) : (Position list
80      * (chessPiece option)) =
81      try
82          // Find index of first non-vacant square of a run
83          let idx = List.findIndex (fun (i, j) -> this.[i,j].IsSome) run
84          let (i,j) = run.[idx]
85          let piece = this.[i, j] // The first non-vacant neighbour
86          if idx = 0
87          then ([], piece)
88          else (run[..(idx-1)], piece)
89      with
90      _ -> (run, None) // outside the board
91  /// find the list of all empty squares and list of neighbours
92  member this.getVacantNNeighbours (piece : chessPiece) : (Position list
93      * chessPiece list) =
94      match piece.position with
95      | None ->
96          ([], [])
97      | Some p ->
98          let convertNWrap =
99              (relativeToAbsolute p) >> this.getVacantNOccupied
100          let vacantPieceLists = List.map convertNWrap
101          piece.candidateRelativeMoves
102          // Extract and merge lists of vacant squares
103          let vacant = List.collect fst vacantPieceLists
104          // Extract and merge lists of first obstruction pieces and
105          filter out own pieces
106          let opponent =
107              vacantPieceLists
108              |> List.choose snd
109          (vacant, opponent)

```

A chess piece must implement `candidateRelativeMoves`, and we decided in Listing 21.10 that moves should be specified relative to the piece's position. Since the piece does not know, which other pieces are on the board, it can only specify all potential positions. For convenience, we will allow pieces to also specify positions outside the board, such that, e.g., the rook can specify the 7 nearest neighboring squares up, down, left, and right regardless that some may be outside the board. Thus `getVacantNNeighbours` must first convert the relative positions to absolute and clip any outside the board. This is done by `relativeToAbsolute`. Then for each run, the first occupied square must be identified. Since `availableMoves` must return two lists, vacant squares, and immediate neighbors, this structure is imposed on the output of `convertNWrap` as well. This is computed in `getVacantNOccupied` by use of the built-in `List.findIndex` function. This function returns the index of the first element in a list for which the supplied function is true and otherwise throws an exception. Exceptions are always somewhat inelegant, but in this case, it is harmless, since the exception signifies a valid situation where no pieces exist on the run. After having analyzed all runs independently, then all the vacant lists are merged and all the neighboring pieces are merged and both are returned to the caller.

Compiling the library files with the application and executing gives the result shown in Listing 21.14.

Listing 21.14: Running the program. Compare with Figure 21.3.

```

1  $ fsharp --nologo chess.fs pieces.fs chessApp.fsx && mono chessApp.exe
2  -----
3  |   |   |   |   |   |   |   |   |
4  -----
5  |   |   |   |   |   |   |   |   |
6  -----
7  |   |   |   |   |   |   |   |   |
8  -----
9  |   | k |   |   |   |   |   |   |
10 -----
11 |   |   |   |   |   |   |   |   |
12 -----
13 |   |   |   |   |   |   |   |   |
14 -----
15 |   | R |   |   |   |   |   |   |
16 -----
17 | K |   |   |   |   |   |   |   |
18 -----
19
20 K: Some (0, 0) ([ (0, 1); (1, 0) ], [R])
21 R: Some (1, 1) ([ (2, 1); (3, 1); (0, 1); (1, 2); (1, 3); (1, 4); (1, 5);
22   (1, 6); (1, 7); (1, 0) ],
23   [k])
24 k: Some (4, 1) ([ (3, 1); (3, 2); (4, 2); (5, 2); (5, 1); (5, 0); (4, 0);
25   (3, 0) ], [])
26 -----
27 |   |   |   |   |   |   |   |   |
28 -----
29 |   |   |   |   |   |   |   |   |
30 -----
31 |   | k |   |   |   |   |   |   |
32 -----
33 |   | R |   |   |   |   |   |   |
34 -----
35 |   |   |   |   |   |   |   |   |
36 -----
37 |   |   |   |   |   |   |   |   |
38 -----
39 | K |   |   |   |   |   |   |   |
40 -----
41
42 K: Some (0, 0) ([ (0, 1); (1, 1); (1, 0) ], [])
43 R: Some (3, 1) ([ (2, 1); (1, 1); (0, 1); (3, 2); (3, 3); (3, 4); (3, 5);
44   (3, 6); (3, 7); (3, 0) ],
45   [k])
46 k: Some (4, 1) ([ (3, 2); (4, 2); (5, 2); (5, 1); (5, 0); (4, 0); (3,
47   0) ], [R])

```

We see that the program has correctly determined that initially, the white king has the white rook as its neighbor and due to its location in the corner only has two free positions to move to. The white rook has many and the black king as its neighbor. The black king is free to move to all its eight neighboring fields. After moving the white rook to (3,1) or b4 in regular chess notation, then the white king has no neighbors, the white rook and the black king are now neighbors with an appropriate restriction on their respective vacant squares. These simple use-tests are in no way a thorough test of

the quality of the code, but they give us a good indication that our library offers a tolerable interface for the application, and that at least major parts of the code function as expected. Thus, we conclude this intermezzo.

22 | The object-oriented programming paradigm

Object-oriented programming is a paradigm for encapsulating data and methods into cohesive units. · Object-oriented programming
Key features of object-oriented programming are:

Encapsulation

Data and methods are collected into a cohesive unit, and an application program need only focus on how to use the object, not on its implementation details.

Inheritance

Objects are organized in a hierarchy of gradually increased specialty. This promotes a design of code that is of general use, and code reuse by specializing the general to the specific.

Polymorphism

By overriding methods from a base class, derived classes define new data types while their methods still produce results compatible with the base class definitions.

Object-oriented programming has a well-developed methodology for analysis and design. The analysis serves as input to the design phase, where the analysis reveals *what* a program is supposed to do, · what
and the design *how* it is supposed to be doing it. The analysis should be expressed in general terms · how
irrespective of the technologic constraints, while the design should include technological constraints such as defined by the targeted language and hardware.

The primary steps for *object oriented analysis and design* are: · object oriented analysis and design

1. identify objects,
2. describe object behavior,
3. describe object interactions,
4. describe some details of the object's inner workings,
5. write a precise description for classes, properties and methods using, e.g., F#'s XML documentation standard,
6. write mockup code,
7. write unit-tests and test the basic framework using the mockup code,
8. replace the mockup with real code while testing to keep track of your progress. Extend the unit-test as needed,
9. evaluate code in relation to the desired goal

10. complete your documentation both in-code and outside.

Step 1–4 are the analysis phase and gradually stops in step 4, while the design phase gradually starts step 4 and gradually stops when actual code is written in step 7. Notice that the last steps are identical with imperative programming Chapter 12. Programming is never a linear experience, and you will often need to go back to previous steps to update or change decisions. You should not refrain from improving your program design and implementation, but you should always be mindful of the goal. Often the perfect solution is much less than needed to complete a task, often less will suffice.

An object-oriented analysis can be a daunting process. A good starting point is a *use case*, *problem statement*, or a *user story*, which in human language describes of a number of possibly hypothetical interactions between a user and a system performs in order to solve some task. Two useful methodologies for performing an object-oriented analysis is the method of nouns-and-verbs and the unified modeling language described in the following sections.

- use case
- problem statement
- user story

22.1 Identification of objects, behaviors, and interactions by nouns-and-verbs

A key point in object-oriented programming is that objects should to a large extent be independent and reusable. As an example, the type `int` models the concept of integer numbers. It can hold integer values from -2,147,483,648 to 2,147,483,647, and a number of standard operations and functions are defined for it. We may use integers in many different programs, and it is certain that the original designers did not foresee our use, but strived to make a general type applicable for many uses. Such a design is a useful goal, when designing objects, that is, our objects should model the general concepts and be applicable in future uses.

Analyzing a specific use-case, good candidates for objects are persons, places, things, events, concept etc., which are almost always characterized by being *nouns* in the text. Interactions between objects are actions that bind objects together, and actions are often associated with *verbs*. When choosing methods, it is important to maintain an object centered perspective, i.e., for a general-purpose object, we should limit the need for including information about other objects. E.g., a value of type `int` need not know anything about the program, in which it is being used.

- nouns
- verbs

Said briefly, the *nouns-and-verbs method* is:

- nouns-and-verbs method

Nouns are object candidates, verbs are candidate methods, that describe interactions between objects.

22.2 Class diagrams in the Unified Modelling Language

Having found an initial list of candidate objects and interactions, it is often useful to make a drawing of these relations and with an increased focus on the object's inner workings. A *class diagram* is a schematic drawing of the program highlighting its object-oriented structure, and we will use the *Unified Modelling Language 2 (UML)* [5] standard. The standard is very broad, and here we will discuss structure diagrams for use of describing objects.

- class diagram
- Unified Modelling Language 2
- UML

A class is drawn as a shown in Figure 22.1. In UML, classes are represented as boxes with their class name. Depending on the desired level of details zero or more properties and methods are described. These describe the basic interface to the class and objects of its type. Abstract members that require an implementation, are shown in cursive. Here we have used F# syntax, to conform with this book theme,




Figure 22.1: A UML diagram for a class, consists of it’s name, zero or more attributes, and zero or more methods.




Figure 22.2: An interface is a class that requires an implementation.

but typically C# syntax is used. Interfaces are a special type of class that require an implementation. To highlight this, UML uses the notation shown in Figure 22.2.¹

Relations between classes and objects are indicated by lines and arrows. The most common ones are summarized in Figure 22.3. Their meaning will be described in detail in the following.

Classes may inherit other classes, where the parent is called the base class and the children its derived classes. Such a relation is often called an *is-a* relation, since the derived class *is a* kind of base class. An illustration of inheritance in UML is shown in Figure 22.4. Here two classes inherit the base class. The syntax is analogous for interfaces, except a stippled line is used to indicate that a derived class implements an interface, as shown in Figure 22.5.

Other relations between classes are association, aggregation, and composition:

Association

In associated relations, one class knows about the other, e.g., uses it as arguments of a function or similar.




Figure 22.3: Arrows used in class diagrams to show relations between objects.




Figure 22.4: Inheritance is shown by a closed arrow head pointing to the base.




Figure 22.5: Implementations of interfaces is shown with stippled line and closed arrow head pointing to the base.

Aggregation

Aggregated relationships is a specialization of associations. In aggregated relations, the host object has a local copy of a guest object, but the host did not create the guest. E.g., the guest object is given as an argument to a function of the host, and the host makes a local alias for later use. When the host is deleted, then the guest is not.

· composition

· has-a

· package



Figure 22.6: Bidirectional association is shown as a line with optional annotation.

n	exactly n instances
*	zero or more instances
n..m	n to m instances
n..*	from n to infinite instances

Table 22.1: Notation for association multiplicities is similar to F#’s slicing notation.



Figure 22.7: Unidirectional association shows a one-side *has-a* relations.

Composition

A composed relationship is a specialization of aggregations. In composed relations, the host creates the guest, and when the host is deleted so is the guest.

Aggregational and compositional relations are often called *has-a* relations since host objects have one or more guests either as aliases or as owner.

Bidirectional association means that classes know about each other. The UML notation is shown in Figure 22.6. Association may be annotated by an identifier and a multiplicity. In the figure, HostA has 0 or more variables of type HostB named objectsBinA, while HostB has 0 or 1 variables of HostA named objectsAinB. The multiplicity notation is very similar to F#’s slicing notation. Typical values are shown in Table 22.1. If the association is unidirectional, then an arrow is added for emphasis as shown in Figure 22.7. In this example, Host knows about Guest and has one instance of it, and Guest is oblivious about Host.

Aggregation is illustrated using a diamond tail and an open arrow as shown in Figure 22.8. Here the Host class has stored aliases to 4 different Guest objects. A stronger relation is composition. This is shown like aggregation but with a filled diamond as illustrated in Figure 22.9. In this example, Owner has created 1 or more objects of type Dependent, and when Owner is deleted so are these objects.

Finally, for visual flair, modules and namespaces are often visualized as a *package* as shown in Figure 22.10. A package is like a module in F#.

22.3 Programming intermezzo: designing a racing game

An example is the following *problem statement*:

· problem statement

¹Jon: Add programming examples for each of these UML structures




Figure 22.8: Aggregation relations are a subset of associations, where local aliases are stored for later use.




Figure 22.9: Composition relations are a subset of aggregation, where the host controls the lifetime of the guest objects.

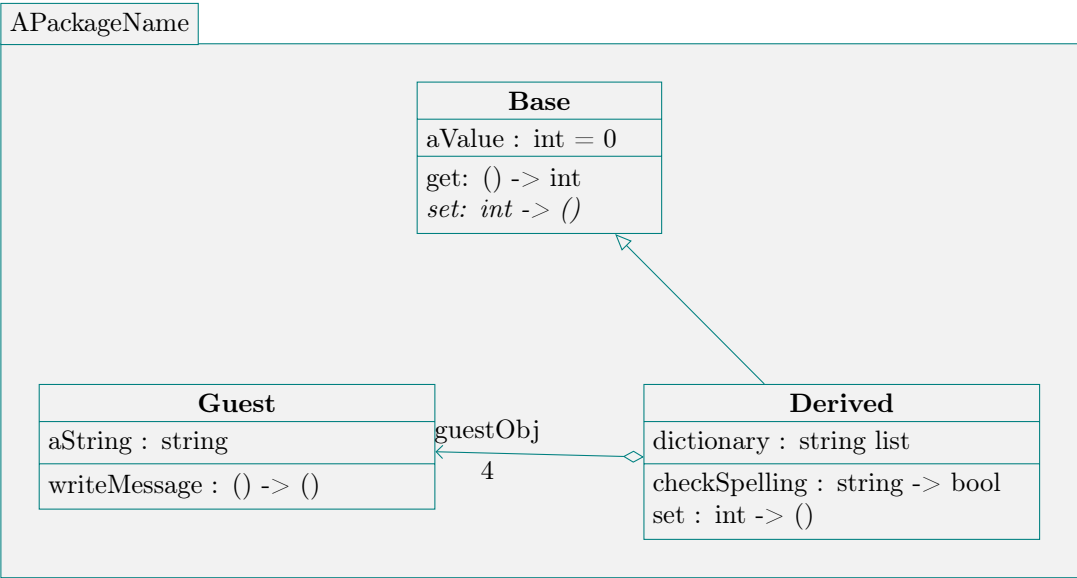


Figure 22.10: Packages are a visualizations of modules and namespaces.

Problem 22.1

Write a racing game, where each player controls his or her vehicle on a track. Each vehicle must have individual features such as top acceleration, speed, and handling. The player must be able to turn the vehicle left and right, and to accelerate up and down. At the beginning of the game, each vehicle is placed behind the starting line. Once the start signal is given, then the players may start to operate their vehicles. The player who first completes 3 rounds wins.

To seek a solution, we will use the *nouns-and-verbs method*. Below the problem statement is repeated with **nouns** and **verbs** highlighted.

Write a racing game, where each player controls his or her vehicle on a track. Each vehicle must have individual features such as top acceleration, speed, and handling. The player must be able to turn the vehicle left and right, and to accelerate up and down. At the beginning of the game, each vehicle is placed behind the starting line. Once the start signal is given, then the players may start to operate their vehicles. The player who first completes 3 rounds wins.

The above nouns and verbs are candidates for objects, their behaviour and interaction. A deeper analysis is:

Identification of objects by nouns (Step 1):

Identified unique nouns are: racing game (game), player, vehicle, track, feature, top acceleration, speed, handling, beginning, starting line, start signal, rounds. From this list we seek cohesive units that are independent and reusable. The nouns

game, player, vehicle, and track

seems to fulfill these requirements, while all the rests seems to be features of the former and thus not independent concepts. E.g., **top acceleration** is a feature of a **vehicle**, and **starting line** is a feature of a **track**.

Object behavior and interactions by verbs (Step 2 and 3):

To continue our object oriented analysis, we will consider the object candidate identified above, and verbalize how they would act as models of general concepts useful in our game.

player The **player** is associated with the following verbs:

- A **player** **controls/operates** a **vehicle**.
- A **player** **turns** and **accelerates** a **vehicle**.
- A **player** **completes** a **rounds**.
- A **player** **wins**.

Verbalizing a **player**, we say that a **player** in general must be able to control the **vehicle**. In order to do this, the **player** must receive information about the **track** and all **vehicles** or at least some information about the nearby **vehicles** and **track**. And the **player** must receive information about the state of the **game**, i.e., when does the race start and stop.

vehicle A **vehicle** is controlled by a **player** and further associated with the following verbs:

- A **vehicle** **has** features **top acceleration**, **speed**, and **handling**.
- A **vehicle** **is placed** and on the **track**.

To further describe a **vehicle** we say that a **vehicle** is a model of a physical object, which moves around on the **track** under the influence of a **player**. A **vehicle** must have a number of attributes such as top acceleration, speed, and handling, and must be able to receive information about when to turn and accelerate. A **vehicle** must be able to determine its location in particular if it is on or off **track** and, and it must be able to determine if it has crashed into an obstacle such as another **vehicle**.

track A **track** is the place where vehicles operate and are further associated with the following verbs:

- A **track** **has** a **starting line**.
- A **track** **has** **rounds**.

Thus, a **track** is a fixed entity on which the **vehicles** race. It has a size and a shape, a starting and a finishing line, which may be the same, and **vehicles** may be placed on the **track** and can move on and possibly off the **track**.

game Finally, a **game** is associated with the following verbs:

- A **game** **has** a **beginning** and a **start signal**.
- A **game** **can be won**.

A **game** is the total sum of all the **players**, the **vehicles**, the **tracks**, and their interactions. A **game** controls the flow of a particular **game** including inviting **players** to race, sending the **start signal**, and monitoring when a **game** is finished and who **won**.

From the above we see that the object candidates **feature** seems to be a natural part of the description of the **vehicle**'s attributes, and similarly, **starting line** may be an intricate part of a **track**. Also, many of the *verbs* used in the problem statement and in our extended verbalization of the general concepts indicate methods that are used to interact with the object. The object centered perspective tells us that for a general-purpose **vehicle** object, we need not include information about the **player**, analogous to a value of type **int** need not know anything the program, in which it is being used. In contrast, the candidate **game** is not as easily dismissed and could be used as a class which contains all the above, i.e.,

With this description, we see that 'start signal' can be included as a natural part of the game object. Being confident that a good working hypothesis of the essential objects for the solution, we continue our investigating into further details about the objects and their interactions.

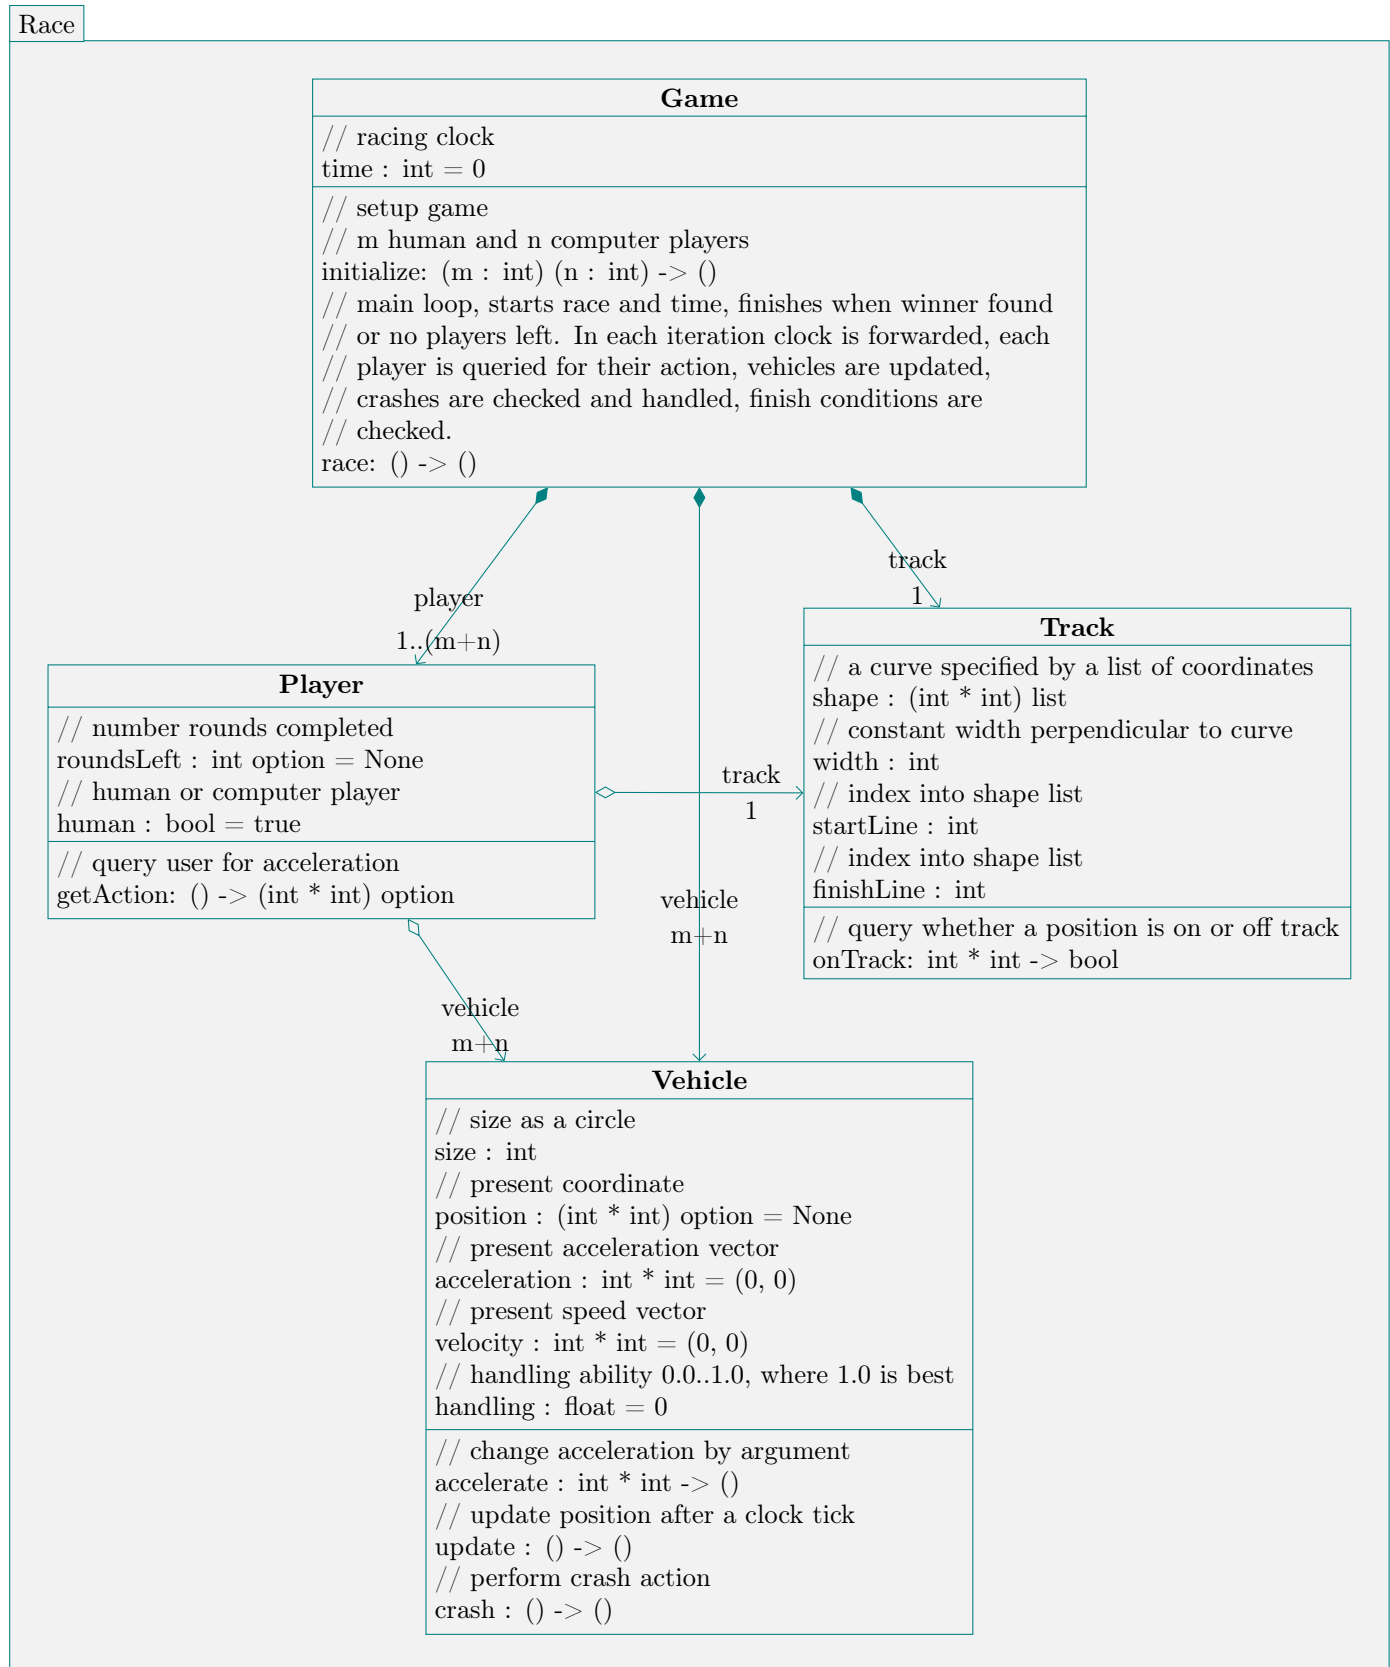


Figure 22.11: A class diagram for a racing game.

Analysis details (Step 4):

A class diagram of our design for the proposed classes and their relations is shown in Figure 22.11.

In the present description, there will be a single Game object, that initializes the other objects, and execute a loop updating the clock, query the players for actions, and informs the vehicles that they should move and under what circumstances. The track has been chosen to be dumb and does not participate much in the action. Player's method `getAction` will be an input from a user by keyboard, joystick or similar, but the complexity of the code for a computer player will be large since it needs to take a sensible decision based on the track and the location of the other vehicles. What at present is less clear, is whether it is the responsibility of Game or Vehicle to detect an off track or a crash event. If a vehicle is to do this, then each vehicle must have aggregated association to all other vehicles and obstacles. So, on the one hand, it would seem an elegant delegation of responsibilities that a vehicle knows, whether it has crashed into an obstacle or not, but on the other hand, it seems wasteful of memory resources to have duplicated references of all obstacles in every vehicle. The final choice is thus one of elegance versus resource management, and in the above, we have favored resource management. Thus, the main loop in Game must check all vehicles for a crash event, after the vehicle's positions have been updated, and in case inform the relevant vehicles.

Having created a design for a racing game, we are now ready to write start coding (Step 6–). It is not uncommon, that transforming our design into code will reveal new structures and problems, that possibly require our design to be updated. Nevertheless, a good design phase is almost always a sure course to avoid many problems once coding, since the design phase allows the programmer to think about the problem from a helicopter perspective before tackling details of specific sub-problems.

23 | Graphical User Interfaces

A *graphical user interface* (*GUI*) uses graphical elements such as windows, icons, and sound to communicate with the user, and a typical way to activate these elements is through a pointing device such as the mouse or by touch. Some of these elements may themselves be textual, and thus most operating systems offer access to a *command-line interface* (*CLI*) in a window alongside other interface types.

- graphical user interface
- GUI
- command-line interface
- CLI

An example of a graphical user interface is a web-browser, as shown in Figure 23.1. The program presents information to the user in terms of text and images and has active areas that may be activated by clicking and which allows the user to go to other web-pages by typing a URL, to follow hyperlinks, and to generate new pages by entering search queries.

F# includes a number of implementations of graphical user interfaces, and at time of writing both *GTK+* and *WinForms 2.0* are supported on both the Microsoft .Net and the Mono platform. WinForms can be used without extra libraries during compilation, and therefore will be the subject of the following chapter.

- GTK+
- WinForms 2.0

WinForms is a set of libraries that simplify many common tasks for applications, and in this chapter, we will focus on the graphical user interface part of WinForms. A *form* is a visual interface used to communicate information with the user, typically a window. Communication is done through *controls*, which are elements that displays information or accepts input. Examples of controls are a box with text, a button, and a menu. When the user gives input to a control element, this generates an *event*, which you can write code to react to. Winforms is designed for *event-driven programming*, meaning that at runtime, most time is spent on waiting for the user to give input. See Chapter 24 for more on event-driven programming.

- form
- control
- event
- event-driven programming

Designing easy to use graphical user interfaces is a challenging task. This chapter will focus on examples of basic graphical elements and how to program these in WinForms.

23.1 Opening a window

Programming graphical user interface with WinForms uses the namespaces: *System.Windows.Forms* and *System.Drawing*. *System.Windows.Forms* includes code for generating forms, controls, and handling events. *System.Drawing* is used for low-level drawing, and it gives access to the *Windows Graphics Device Interface* (*GDI+*), which allows you to create and manipulate graphics objects targeting several platforms such as screens and paper. All controls in *System.Windows.Forms* in Mono are drawn using *System.Drawing*.

- *System.Windows.Forms*
- *System.Drawing*
- Windows Graphics Device Interface
- GDI+

To display a graphical user interface on the screen, the first thing to do is open a window, which acts as a reserved screen-space for our output. In WinForms, windows are called forms. Code for opening a window is shown in Listing 23.1, and the result is shown in Figure 23.2. Note, the present version of WinForms on MacOS only works with the 32-bit implementation of mono, *mono32*, as demonstrated in the example.

- *mono32*

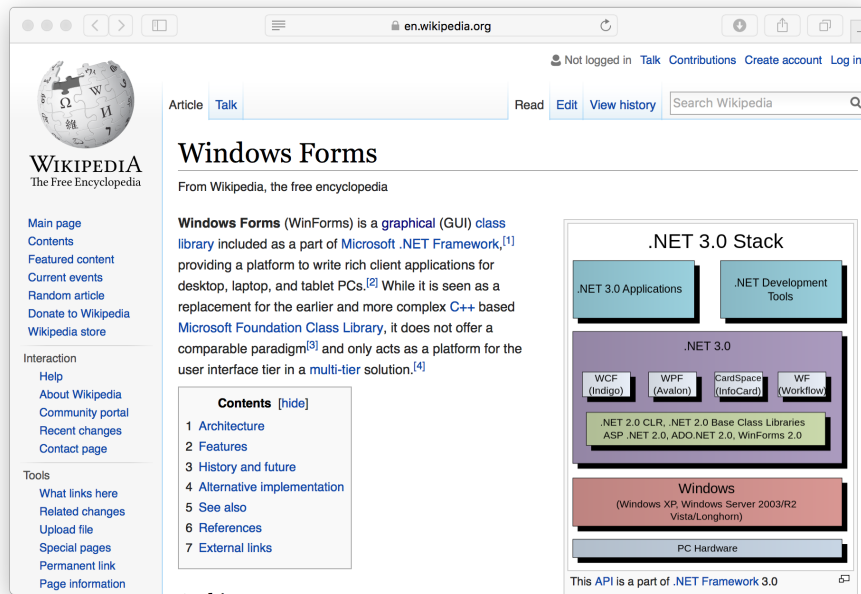


Figure 23.1: A web-browser is a graphical user interface for accessing a web-server and interacting with its services. Here the browser is showing the page `https://en.wikipedia.org/wiki/Windows_Forms` at time of writing.

Listing 23.1 winforms/openWindow.fsx:

Create the window and turn over control to the operating system. See Figure 23.2.

```
1 // Create a window
2 let win = new System.Windows.Forms.Form ()
3 // Start the event-loop.
4 System.Windows.Forms.Application.Run win

1 $ fsharpc --nologo openWindow.fsx && mono32 openWindow.exe
```

The `new System.Windows.Forms.Form ()` creates an object (See Chapter 20), but does not display the window on the screen. We use the optional `new` keyword since the form is an `IDisposable` object and may be implicitly disposed of. I.e., it is recommended to **instantiate `IDisposable` objects using `new` to contrast them with other object types**. Executing `System.Windows.Forms.Application.Run` is applied to the object, then the control is handed over to the WinForms' *event-loop*, which continues until the window is closed by, e.g., pressing the icon designated by the operating system. On the Mac OSX that is the red button in the top left corner of the window frame, and on Windows it is the cross on the top right corner of the window frame.

Advice
· event-loop

The window form has a long list of *methods* and *properties*. E.g., the background color may be set by `BackColor`, the title of the window may be set by `Text`, and you may get and set the size of the window with the `Size`. This is demonstrated in Listing 23.2.

· methods
· properties




Figure 23.2: A window opened by Listing 23.1.

Listing 23.2 winforms/windowProperty.fsx:

Create the window and changing its properties. See Figure 23.3

```

1 // Prepare window form
2 let win = new System.Windows.Forms.Form ()
3
4 // Set some properties
5 win.BackColor <- System.Drawing.Color.White
6 win.Size <- System.Drawing.Size (600, 200)
7 win.Text <- sprintf "Color '%A' and Size '%A'" win.BackColor win.Size
8
9 // Start the event-loop.
10 System.Windows.Forms.Application.Run win

```

These properties are *accessors* implying that they act as mutable variables.

· accessors

23.2 Drawing geometric primitives

The *System.Drawing.Color* is a structure for specifying colors as 4 channels: alpha, red, green, and blue. Some methods and properties for the Color structure is shown in Table 23.1. Each channel is an 8-bit unsigned integer, but often referred as the 32 bit unsigned integer by concatenating the channels. The alpha channel specifies the transparency of a color, where values 0–255 denotes the range of fully transparent to fully opaque, and the remaining channels denote the amount of red, green, and blue, where 0 is none and 255 is full intensity. Any color may be created using the `FromArgb` method, e.g., an opaque red is given by `System.Drawing.Color.FromArgb (255, 255, 0, 0)`. There are also many build-in colors, e.g., the same red color is also a known color and may be obtained as `System.Drawing.Color.Red`. For a given color, then the 4 alpha, red, green, and blue channel's values may be obtained as the `A`, `R`, `G`, and `B` members, see Listing 23.3

· Color




Figure 23.3: A window with user-specified size and background color, see Listing 23.2.

Method/Property	Description
A	Get the value of the alpha channel of a color.
B	Get the value of the blue channel of a color.
Black	Get a predefined color with ARGB value of 0xFF000000.
Blue	Get a predefined color with ARGB value of 0xFF0000FF.
Brown	Get a predefined color with ARGB value of 0xFFA52A2A.
FromArgb : int -> Color FromArgb : int*int*int -> Color FromArgb : int*int*int*int -> Color	Create a color structure.
G	Get the value of the green channel of a color.
Gray	Get a predefined color with ARGB value of 0xFF808080.
Green	Get a predefined color with ARGB value of 0xFF00FF00.
Orange	Get a predefined color with ARGB value of 0xFFFFFA500.
Purple	Get a predefined color with ARGB value of 0xFF800080.
R	Get the value of the red channel of a color.
Red	Get a predefined color with ARGB value of 0xFFFFF0000.
ToArgb : Color -> int	Get the 32-bit integer representation of a color.
White	Get a predefined color with ARGB value of 0xFFFFFFFF.
Yellow	Get a predefined color with ARGB value of 0xFFFFFFF00.

Table 23.1: Some methods and properties of the `System.Drawing.Color` structure.

Constructor	Description
Bitmap(int, int)	Create a new empty Image of specified size.
Bitmap(Stream) Bitmap(string)	Create a Image from a System.IO.Stream or from a file specified by a filename.
Font(string, single)	Create a new font from the font's name and em-size.
Pen(Brush) Pen(Brush), single) Pen(Color) Pen(Color, single)	Create a pen to paint either with a brush or solid color and possibly with specified width.
Point(int, int) Point(Size) PointF(single, single)	Create an ordered pair of integers or singles specifying x- and y-coordinates in the plane.
Size(int, int) Size(Point) SizeF(single, single) SizeF(PointF)	Create an ordered pair of integers or singles specifying height and width in the plane.
SolidBrush(Color) TextureBrush(Image)	Create a Brush as a solid color or from an image to fill the interior of a geometric shapes.

Table 23.2: Basic geometrical structures in WinForms. Brush and Image are abstract classes.

Listing 23.3 drawingColors.fsx:
Defining colors and accessing their values.

```
1 // open namespace for brevity
2 open System.Drawing
3 // Define a color from ARGB
4 let c = Color.FromArgb (0xFF, 0x7F, 0xFF, 0xD4) //Aquamarine
5 printfn "The color %A is (%x, %x, %x, %x)" c c.A c.R c.G c.B
6 // Define a list of named colors
7 let colors =
8     [Color.Red; Color.Green; Color.Blue;
9      Color.Black; Color.Gray; Color.White]
10 for col in colors do
11     printfn "The color %A is (%x, %x, %x, %x)" col col.A col.R col.G col.B
```

```
1 $ fsharpc --nologo drawingColors.fsx && mono drawingColors.exe
2 The color Color [A=255, R=127, G=255, B=212] is (ff, 7f, ff, d4)
3 The color Color [Red] is (ff, ff, 0, 0)
4 The color Color [Green] is (ff, 0, 80, 0)
5 The color Color [Blue] is (ff, 0, 0, ff)
6 The color Color [Black] is (ff, 0, 0, 0)
7 The color Color [Gray] is (ff, 80, 80, 80)
8 The color Color [White] is (ff, ff, ff, ff)
```

System.Drawing contains many often used structures. Listing 23.2 used System.Drawing.Size to specify a size by a pair of integers. Other important structures in WinForms are *System.Drawing.Point*, which specifies a coordinate as a pair of points; *System.Drawing.Pen*, which specifies how to draw lines and curves; *System.Drawing.Font*, which specifies the font of a string; *System.Drawing.SolidBrush* and *System.Drawing.TextureBrush*, used for filling geometric primitives, and *System.Drawing.Bitmap*, which is a type of *System.Drawing.Image*. These are summarized in Table 23.2. ¹

- Point
- Pen
- Font
- SolidBrush
- TextureBrush
- Bitmap
- Image

Constructor	Description
<code>DrawImage : Image * (Point []) -> unit</code> <code>DrawImage : Image * (PointF []) -> unit</code>	Draw an image at a specific point and size.
<code>DrawImage : Image * Point -> unit</code> <code>DrawImage : Image * PointF -> unit</code>	Draw an image at a specific point.
<code>DrawLines : Pen * (Point []) -> unit</code> <code>DrawLines : Pen * (PointF []) -> unit</code>	Draw a series of lines between the n 'th and the $n + 1$ 'th points.
<code>DrawString :</code> <code>string * Font * Brush * PointF -> unit</code>	Draw a string at the specified point.

Table 23.3: Basic geometrical structures in WinForms.




Figure 23.4: Coordinate systems in Winforms have the y axis pointing down.

The `System.Drawing.Graphics` is a class for drawing geometric primitives to a display device, and some of its methods are summarized in Table 23.3.

The location and shape of geometrical primitives are specified in a coordinate system, and WinForms operates with 2 coordinate systems: *screen coordinates* and *client coordinates*. Both coordinate systems have their origin in the top-left corner, with the first coordinate, x , increases to the right, and the second, y , increases down as illustrated in Figure 23.4. The Screen coordinate system has its origin in the top-left corner of the screen, while the client coordinate system has its origin in the top-left corner of the drawable area of a form or a control, i.e., for a window, this will be the area without the window borders, scroll and title bars. A control is a graphical object such as a clickable button, which will be discussed later. Conversion between client and screen coordinates is done with `System.Drawing.PointToClient` and `System.Drawing.PointToScreen`. To draw geometric primitives, we must also specify the pen using for a line like primitives and the brush for filled regions.

Displaying graphics in WinForms is performed as the reaction to an event. E.g., windows are created by the program, moved, minimized, occluded by other windows, resized, etc., by the user or the program, and each action may require that the content of the window is refreshed. Thus, we must create a function that WinForms can call any time. This is known as a *call-back function*, and it is added to an existing form using the form's `Paint.Add` method. Due to the event-driven nature of WinForms, functions for drawing graphics primitives are only available when responding to an event, e.g., `System.Drawing.Graphics.DrawLine` draws a line in a window, and *it is only possible to call this function, as part of an event handling*.

As an example, consider the problem of drawing a triangle in a window. For this we need to make a function that can draw a triangle not once, but any time. An example of such a program is shown in Listing 23.4.

¹Jon: Do something about the vertical alignment of minpage.




Figure 23.5: Drawing a triangle using Listing 23.4.

Listing 23.4 winforms/triangle.fsx:
Adding line graphics to a window. See Figure 23.5

```

1 // Open often used libraries, be ware of namespace pollution!
2 open System.Windows.Forms
3 open System.Drawing
4
5 // Prepare window form
6 let win = new Form ()
7 win.Size <- Size (320, 170)
8
9 // Set paint call-back function
10 let paint (e : PaintEventArgs) : unit =
11     let pen = new Pen (Color.Black)
12     let points =
13         [|Point (0,0); Point (10,170); Point (320,20); Point (0,0)|]
14     e.Graphics.DrawLines (pen, points)
15 win.Paint.Add paint
16
17 // Start the event-loop.
18 Application.Run win // Start the event-loop.

```

A walk-through of the code is as follows: First, we open the two libraries that we will use heavily. This will save us some typing but also pollute our namespace. E.g., now `Point` and `Color` are existing types, and we cannot define our own identifier with these names. Then we create the form with size 320×170 , we add a paint call-back function and we start the event-loop. The event-loop will call the paint function, whenever the system determines that the window's content needs to be refreshed. This function is to be called as a response to a paint event and takes a `System.Windows.Forms.PaintEventArgs` object, which includes the `System.Drawing.Graphics` object. The function `paint` chooses a pen and a set of points and draws a set of lines connecting the points.

The code in Listing 23.4 is not optimal. In spite that the triangle spans the rectangle $(0,0)$ to $(320,170)$ and the window's size is set to $(320,170)$, then our window is too small and the triangle is clipped at the window border. The error is that we set the window's `Size` property which determines the size of the window including top bar and borders. Alternatively we may set the `ClientSize`, which determines the size of the drawable area, and this is demonstrated in Listing 23.5 and Figure 23.6.




Figure 23.6: Setting the `ClientSize` property gives a predictable drawing area, see Listing 23.5 for code.

Listing 23.5 `winforms/triangleClientSize.fsx`:
Adding line graphics to a window. See Figure 23.6.

```

1 // Open often used libraries, be ware of namespace pollution!
2 open System.Windows.Forms
3 open System.Drawing
4
5 // Prepare window form
6 let win = new Form ()
7 win.ClientSize <- Size (320, 170)
8
9 // Set paint call-back function
10 let paint (e : PaintEventArgs) : unit =
11     let pen = new Pen (Color.Black)
12     let points =
13         [|Point (0,0); Point (10,170); Point (320,20); Point (0,0)|]
14     e.Graphics.DrawLines (pen, points)
15 win.Paint.Add paint
16
17 // Start the event-loop.
18 Application.Run win // Start the event-loop.

```

Thus, **prefer the `ClientSize` over the `Size` property for internal consistency.**

Advice

Considering the program in Listing 23.4, we may identify a part that concerns the specification of the triangle, or more generally the graphical model, and some which concerns system specific details. For future maintenance, it is often a good idea to **separate the model from how it is viewed on a specific system**. E.g., it may be that at some point, you decide that you would rather use a different library than WinForms. In this case, the general graphical model will be the same but the specific details on initialization and event handling will be different. We think of the model and the viewing part of the code as top and bottom layers, and these are often connected with a connection layer. This *Model-View paradigm* is shown in Figure 23.7. While it is not easy to completely separate the general

Advice

· Model-View paradigm

In Listing 23.6, the program has been redesigned to follow the Model-View paradigm, to make use of a Winforms and a model specific function called `view` and `model`.




Figure 23.7: Separating model from view gives flexibility later.

Listing 23.6 winforms/triangleOrganized.fsx:

Improved organization of code for drawing a triangle. See Figure 23.8.

```

1 // Open often used libraries, be ware of namespace polution!
2 open System.Windows.Forms
3 open System.Drawing
4
5 ////////////// WinForm specifics ///////////////////
6 // Setup a window form and return function to activate
7 let view (sz : Size) (pen : Pen) (pts : Point []) : (unit -> unit) =
8     let win = new System.Windows.Forms.Form ()
9     win.ClientSize <- sz
10    win.Paint.Add (fun e -> e.Graphics.DrawLines (pen, pts))
11    fun () -> Application.Run win // function as return value
12
13 ////////////// Model ///////////////////
14 // A black triangle, using winform primitives for brevity
15 let model () : Size * Pen * (Point []) =
16     let size = Size (320, 170)
17     let pen = new Pen (Color.FromArgb (0, 0, 0))
18     let lines =
19         [|Point (0,0); Point (10,170); Point (320,20); Point (0,0)|]
20     (size, pen, lines)
21
22 ////////////// Connection ///////////////////
23 // Tie view and model together and enter main event loop
24 let (size, pen, lines) = model ()
25 let run = view size pen lines
26 run ()
  
```

This program is longer, but there is a much better separation of *what* is to be displayed (model) from the *how* it is to be done (view).

To further our development of a general program for displaying graphics, consider the case, where we are to draw another two triangles, that are a translation and a rotations of the original, and where we would like to specify the color of each triangle individually. A simple extension of model in Listing 23.6 for generating many shapes of different colors is `model : unit -> Size * ((Point []) * Pen) list`, i.e., semantically augment each point array with a pen and return a list of such pairs. For this example we also program translation and rotation transformations. See Listing 23.7 for the result.




Figure 23.8: Better organization of the code for drawing a triangle, see Listing 23.6.

Listing 23.7 winforms/transformWindows.fsx:

Model of a triangle and simple transformations of it. See also Listing 23.8 and 23.9, and Figure 23.9.

```

15 /////////////// Model ///////////////////
16 // A black triangle, using winform primitives for brevity
17 let model () : Size * ((Pen * (Point [])) list) =
18     /// Translate a primitive
19     let translate (d : Point) (arr : Point []) : Point [] =
20         let add (d : Point) (p : Point) : Point =
21             Point (d.X + p.X, d.Y + p.Y)
22         Array.map (add d) arr
23
24     /// Rotate a primitive
25     let rotate (theta : float) (arr : Point []) : Point [] =
26         let toInt = int << round
27         let rot (t : float) (p : Point) : Point =
28             let (x, y) = (float p.X, float p.Y)
29             let (a, b) = (x * cos t - y * sin t, x * sin t + y * cos t)
30             Point (toInt a, toInt b)
31         Array.map (rot theta) arr
32
33     let size = Size (400, 200)
34     let lines =
35         [|Point (0,0); Point (10,170); Point (320,20); Point (0,0)|]
36     let black = new Pen (Color.FromArgb (0, 0, 0))
37     let red = new Pen (Color.FromArgb (255, 0, 0))
38     let green = new Pen (Color.FromArgb (0, 255, 0))
39     let shapes =
40         [(black, lines);
41          (red, translate (Point (40, 30)) lines);
42          (green, rotate (1.0 * System.Math.PI / 180.0) lines)]
43     (size, shapes)

```

We update `view` accordingly to iterate through this list as shown in Listing 23.8.




Figure 23.9: Transformed versions of the same triangle resulting from running the code in Listing 23.7–23.9.

Listing 23.8 winforms/transformWindows.fsx:

View of lists of pairs of pen and point arrays. See also Listing 23.7 and 23.9, and Figure 23.9.

```

1 // Open often used libraries, be ware of namespace pollution!
2 open System.Windows.Forms
3 open System.Drawing
4
5 ////////////// WinForm specifics ///////////////////
6 /// Setup a window form and return function to activate
7 let view (sz : Size) (shapes : (Pen * (Point [])) list) : (unit -> unit)
8 =
9     let win = new System.Windows.Forms.Form ()
10    win.ClientSize <- sz
11    let paint (e : PaintEventArgs) ((p, pts) : (Pen * (Point []))) : unit =
12        e.Graphics.DrawLine (p, pts)
13    win.Paint.Add (fun e -> List.iter (paint e) shapes)
14    fun () -> Application.Run win // function as return value

```

Since we are using Winforms primitives in the model, then the connection layer is trivial as shown in Listing 23.9.

Listing 23.9 winforms/transformWindows.fsx:

Model of a triangle and simple transformations of it. See also Listing 23.7 and 23.8, and Figure 23.9.

```

45 ////////////// Connection ///////////////////
46 // Tie view and model together and enter main event loop
47 let (size, shapes) = model ()
48 let run = view size shapes
49 run ()

```

23.3 Programming intermezzo: Hilbert Curve

A curve in 2 dimensions has a length but no width, and we can only visualize it by giving it a width. Thus, it came as a surprise to many when Giuseppe Peano in 1890 demonstrated that there exist curves, which fill every point in a square. The method he used to achieve this was recursion:

Problem 23.1

Consider a curve consisting of piecewise straight lines all with the same length but with varying angles 0° , 90° , 180° , or 270° w.r.t. the horizontal axis. To draw this curve we need 3 basic operations: Move forward (*F*), turn right (*R*), and turn left (*L*). The turning is w.r.t. the present direction. A Hilbert Curve is a space-filling curve, which can be expressed recursively as:

$$A \rightarrow LBFRAFARFBL \quad (23.1)$$

$$B \rightarrow RAFLBFBLFAR \quad (23.2)$$

starting with *A*. For practical illustrations, we typically only draw space-filling curves to a specified depth of recursion, which is called the order of the curve. Hence, to draw a first order curve, we don't recurse at all, i.e., ignore all occurrences of the symbols *A* and *B* on the right-hand-side of (23.1), and get

$$A \rightarrow LFRFRFL.$$

For the second order curve, we recurse once, i.e.,

$$\begin{aligned} A &\rightarrow LBFRAFARFBL \\ &\rightarrow L(RAFLBFBLFAR)F \\ &\quad R(LBFRAFARFBL)F(LBFRAFARFBL) \\ &\quad RF(RAFLBFBLFAR)L \\ &\rightarrow LRFLFLFRFRFLFRFRFLFLFRFRFLRFRFLFLFL. \end{aligned}$$

Since $LR = RL = \emptyset$ then the above simplifies to $FLFLFRFRFRFLFLFRFRFRFLFLFL$. Make a program, that given an order produces an image of the Hilbert curve.

Our strategy to solve this problem will be first to define the curves in terms of movement commands $LRFL\dots$. For this, we will define a discriminated union `type Command = F | L | R`. The movement commands can then be defined as a `Command list` type. The list for a specific order is a simple set of recursive functions in `F#`, which we will call *A* and *B*.

To produce a graphical drawing of a command list, we must transform it into coordinates, and during the conversion, we need to keep track of both the present position and the present heading, since not all commands draw. This is a concept similar to Turtle Graphics, which is often associated with the Logo programming language from the 1960's. In Turtle graphics, we command a little robot - a turtle - which moves in 2 dimensions and can turn on the spot or move forward, and its track is the line being drawn. Thus we introduce a `type Turtle = {x : float; y : float; d : float}` record. Conversion of command lists to turtle lists is a fold programming structure, where the command list is read from left-to-right, building up an accumulator by adding each new element. For efficiency, we choose to prepend the new element to the accumulator. This we have implemented as the `addRev` function. Once the full list of turtles has been produced, then it is reversed.

Finally, the turtle list is converted to Winforms `Point` array, and a window of appropriate size is chosen. The resulting model part is shown in Listing 23.10. The view and connection parts are identical to Listing 23.8 and 23.9, and Figure 23.10 shows the result of using the program to draw Hilbert curves of orders 1, 2, 3, and 5.

Listing 23.10 winforms/hilbert.fsx:

Using simple turtle graphics to produce a list of points on a polygon. The view and connection parts are identical to Listing 23.8 and 23.9.

```

15 // Turtle commands, type definitions must be in outer most scope
16 type Command = F | L | R
17 type Turtle = {x : float; y : float; d : float}
18
19 // A black Hilbert curve using winform primitives for brevity
20 let model () : Size * ((Pen * (Point [])) list) =
21     /// Hilbert recursion production rules
22     let rec A n : Command list =
23         if n > 0 then
24             [L]@B (n-1)@[F; R]@A (n-1)@[F]@A (n-1)@[R; F]@B (n-1)@[L]
25         else
26             []
27     and B n : Command list =
28         if n > 0 then
29             [R]@A (n-1)@[F; L]@B (n-1)@[F]@B (n-1)@[L; F]@A (n-1)@[R]
30         else
31             []
32
33     /// Convert a command to a turtle record and prepend to a list
34     let addRev (lst : Turtle list) (cmd : Command) (len : float) : Turtle
35         list =
36         let toInt = int << round
37         match lst with
38         | t::rest ->
39             match cmd with
40             | L -> {t with d = t.d + 3.141592/2.0}::rest // turn left
41             | R -> {t with d = t.d - 3.141592/2.0}::rest // turn right
42             | F -> {t with x = t.x + len * cos t.d;
43                     y = t.y + len * sin t.d}::rest // forward
44             | _ -> failwith "Turtle list must be non-empty."
45
46     let maxPoint (p1 : Point) (p2 : Point) : Point =
47         Point (max p1.X p2.X, max p1.Y p2.Y)
48
49     // Calculate commands for a specific order
50     let curve = A 5
51     // Convert commands to point array
52     let initTrtl = {x = 0.0; y = 0.0; d = 0.0}
53     let len = 20.0
54     let line =
55         // Convert command list to reverse turtle list
56         List.fold (fun acc elm -> addRev acc elm len) [initTrtl] curve
57         // Reverse list
58         |> List.rev
59         // Convert turtle list to point list
60         |> List.map (fun t -> Point (int (round t.x), int (round t.y)))
61         // Convert point list to point array
62         |> List.toArray
63     let black = new Pen (Color.FromArgb (0, 0, 0))
64     // Set size to as large as shape
65     let minVal = System.Int32.MinValue
66     let maxPoint = Array.fold maxPoint (Point (minVal, minVal)) line
67     let size = Size (maxPoint.X + 1, maxPoint.Y + 1)
68     // return shapes as singleton list
69     (size, [(black, line)])

```




Figure 23.10: Hilbert curves of order 1, 2, 3, and 5 by code in Listing 23.10.

23.4 Handling events

In the previous section, we have looked at how to draw graphics using the `Paint` method of an existing form object. Forms have many other event handlers, that we may use to interact with the user. Listing 23.11 demonstrates event handlers for moving and resizing a window, for clicking in a window, and for typing on the keyboard.

Listing 23.11 `winforms/windowEvents.fsx`:
Catching window, mouse, and keyboard events.

```
1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form () // create a form
6
7  // Window event
8  let windowMove (e : EventArgs) =
9      printfn "Move: %A" win.Location
10     win.Move.Add windowMove
11
12  let windowResize (e : EventArgs) =
13      printfn "Resize: %A" win.DisplayRectangle
14     win.Resize.Add windowResize
15
16  // Mouse event
17  let mutable record = false; // records when button down
18  let mouseMove (e : MouseEventArgs) =
19      if record then printfn "MouseMove: %A" e.Location
20     win.MouseMove.Add mouseMove
21
22  let mouseDown (e : MouseEventArgs) =
23      printfn "MouseDown: %A" e.Location; (record <- true)
24     win.MouseDown.Add mouseDown
25
26  let mouseUp (e : MouseEventArgs) =
27      printfn "MouseUp: %A" e.Location; (record <- false)
28     win.MouseUp.Add mouseUp
29
30  let mouseClicked (e : MouseEventArgs) =
31      printfn "MouseClicked: %A" e.Location
32     win.MouseClick.Add mouseClicked
33
34  // Keyboard event
35  win.KeyPreview <- true
36  let keyPress (e : KeyPressEventArgs) =
37      printfn "KeyPress: %A" (e.KeyChar.ToString ())
38     win.KeyPress.Add keyPress
39
40  Application.Run win // Start the event-loop.
```


Listing 23.12: Output from an interaction with the program in Listing 23.11.

```

1  Move: {X=22,Y=22}
2  Move: {X=22,Y=22}
3  Move: {X=50,Y=71}
4  Resize: {X=0,Y=0,Width=307,Height=290}
5  MouseDown: {X=144,Y=118}
6  MouseClick: {X=144,Y=118}
7  MouseUp: {X=144,Y=118}
8  MouseDown: {X=144,Y=118}
9  MouseUp: {X=144,Y=118}
10 MouseDown: {X=96,Y=66}
11 MouseMove: {X=96,Y=67}
12 MouseMove: {X=97,Y=69}
13 MouseMove: {X=99,Y=71}
14 MouseMove: {X=103,Y=74}
15 MouseMove: {X=107,Y=77}
16 MouseMove: {X=109,Y=79}
17 MouseMove: {X=112,Y=81}
18 MouseMove: {X=114,Y=82}
19 MouseMove: {X=116,Y=84}
20 MouseMove: {X=117,Y=85}
21 MouseMove: {X=118,Y=85}
22 MouseClick: {X=118,Y=85}
23 MouseUp: {X=118,Y=85}
24 KeyPress: "a"
25 KeyPress: "b"
26 KeyPress: "c"

```

In Listing 23.11 is shown the output from an interaction with the program which is the result of the following actions: moving the window, resizing the window, clicking the left mouse key, pressing and holding the down the left mouse key while moving the mouse, and releasing the left mouse key, and type “abc”. As demonstrated, some actions like moving the mouse results in a lot of events, and some like the initial window moves results are surprising. Thus, event-driven programming should take care to interpret the events robustly and carefully.

Common for all event-handlers is that they listen for an event, and when the event occurs, then the functions that have been added using the `Add` method are called. This is also known as sending a message. Thus, a single event can give rise to calling zero or more functions.

Graphical user interfaces and other systems often need to perform actions that depends on specific lengths of time or at a certain point in time. To time events, F# has the `System.Timer` class, which has been optimized for graphical user interfaces as `System.Windows.Forms.Timer`, and which can be set to create an event after a specified duration of time. F# also has the `System.DateTime` class to specify points in time. An often used property is `System.DateTime.Now`, which returns a `DateTime` object for the date and time, when the property is accessed. The use of these two classes is demonstrated in Listing 23.13 and Figure 23.11.

· `Timer`
· `DateTime`




Figure 23.11: See Listing 23.13.

Listing 23.13 winforms/clock.fsx:

Using `System.Windows.Forms.Timer` and `System.DateTime.Now` to update the display of the present date and time. See Figure 23.11 for the result.

```

1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form () // make a window form
6  win.ClientSize <- Size (200, 50)
7
8  // make a label to show time
9  let label = new Label()
10 win.Controls.Add label
11 label.Width <- 200
12 label.Text <- string System.DateTime.Now // get present time and date
13
14 // make a timer and link to label
15 let timer = new Timer()
16 timer.Interval <- 1000 // create an event every 1000 millisecond
17 timer.Enabled <- true // activate the timer
18 timer.Tick.Add (fun e -> label.Text <- string System.DateTime.Now)
19
20 Application.Run win // start event-loop

```

In the code, a label has been created to show the present date and time. The label is a type of control, and it is displayed using the default font, which is rather small. How to change this and other details on controls will be discussed in the next section.

23.5 Labels, buttons, and pop-up windows

In WinForms buttons, menus and other interactive elements are called *Controls*. A form is a type of control, and thus, programming controls are very similar to programming windows. In Listing 23.14 is shown a small program that displays a label and a button in a window, and when the button is pressed, then the label is updated.




Figure 23.12: After pressing the button 3 times. See Listing 23.14.

Listing 23.14 winforms/buttonControl.fsx:
Create the button and an event, see also Figure 23.12.

```

1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form () // make a window form
6  win.ClientSize <- Size (140, 120)
7
8  // Create a label
9  let label = new Label()
10 win.Controls.Add label
11 label.Location <- new Point (20, 20)
12 label.Width <- 120
13 let mutable clicked = 0
14 let setLabel clicked =
15     label.Text <- sprintf "Clicked %d times" clicked
16     setLabel clicked
17
18 // Create a button
19 let button = new Button ()
20 win.Controls.Add button
21 button.Size <- new Size (100, 40)
22 button.Location <- new Point (20, 60)
23 button.Text <- "Click me"
24 button.Click.Add (fun e -> clicked <- clicked + 1; setLabel clicked)
25
26 Application.Run win // Start the event-loop.

```

As the listing demonstrates, the button is created using the *System.Windows.Forms.Button* constructor, and it is added to the window's form's control list. The *Location* property controls its position w.r.t. the enclosing form. Other accessors are *Width*, *Text*, and *Size*.

System.Windows.Forms includes a long list of controls, some of which are summarized in Table 23.4. Examples are given in controls is shown in Listing 23.15 and Figure 23.13.

- Button
- CheckBox
- DateTimePicker
- Label
- ProgressBar
- RadioButton
- TextBox

Method/Property	Description
Button	A clickable button.
CheckBox	A clickable check box.
DateTimePicker	A box showing a date with a drop-down menu for choosing another.
Label	A displayable text.
ProgressBar	A box showing a progress bar.
RadioButton	A single clickable radio button. Can be paired with other radio buttons.
TextBox	A text area, which can accept input from the user.

Table 23.4: Some types of `System.Windows.Forms.Control`.**Listing 23.15** `winforms/controls.fsx`:

Examples of control elements added to a window form, see also Figure 23.13.

```

1  open System.Windows.Forms
2  open System.Drawing
3
4  let win = new Form ()
5  win.ClientSize <- Size (300, 300)
6
7  let button = new Button ()
8  win.Controls.Add button
9  button.Location <- new Point (20, 20)
10 button.Text <- "Click me"
11
12 let lbl = new Label ()
13 win.Controls.Add lbl
14 lbl.Location <- new Point (20, 60)
15 lbl.Text <- "A text label"
16
17 let chkbox = new CheckBox ()
18 win.Controls.Add chkbox
19 chkbox.Location <- new Point (20, 100)
20
21 let pick = new DateTimePicker ()
22 win.Controls.Add pick
23 pick.Location <- new Point (20, 140)
24
25 let prgrss = new ProgressBar ()
26 win.Controls.Add prgrss
27 prgrss.Location <- new Point (20, 180)
28 prgrss.Minimum <- 0
29 prgrss.Maximum <- 9
30 prgrss.Value <- 3
31
32 let rdbtn = new RadioButton ()
33 win.Controls.Add rdbtn
34 rdbtn.Location <- new Point (20, 220)
35
36 let txtbox = new TextBox ()
37 win.Controls.Add txtbox
38 txtbox.Location <- new Point (20, 260)
39 txtbox.Text <- "Type something"
40
41 Application.Run win

```

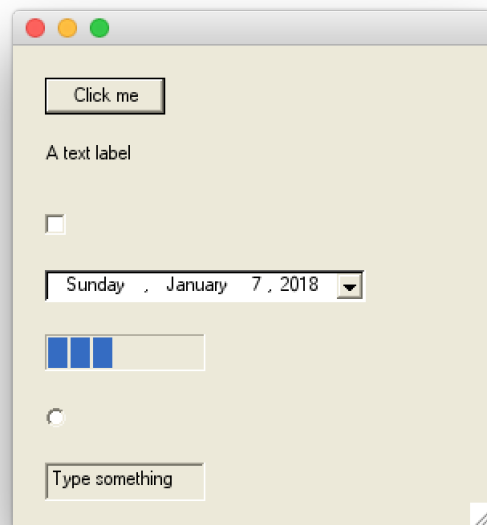


Figure 23.13: Examples of control elements. See Listing 23.15.

Some controls opens separate windows for more involved dialogue with the user. Some examples are `MessageBox`, `OpenFileDialog`, and `SaveFileDialog`.

`System.Windows.Forms.MessageBox` is used to have a simple but restrictive dialogue with the user is `MessageBox` needed, which is demonstrated in Listing 23.16 and Figure 23.14.




Figure 23.14: After pressing the button 3 times. See Listing 23.16.

Listing 23.16 winforms/messageBox.fsx:
Create the MessageBox, see also Figure 23.14.

```

1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form ()
6  win.ClientSize <- Size (140, 80)
7
8  let button = new Button ()
9  win.Controls.Add button
10 button.Size <- new Size (100, 40)
11 button.Location <- new Point (20, 20)
12 button.Text <- "Click me"
13 // Open a message box, when button clicked
14 let buttonClicked (e : EventArgs) =
15     let question = "Is this statement false?"
16     let caption = "Window caption"
17     let boxType = MessageBoxButtons.YesNo
18     let response = MessageBox.Show (question, caption, boxType)
19     printfn "The user pressed %A" response
20 button.Click.Add buttonClicked
21
22 Application.Run win

```

As alternative to YesNo response button, message box also offers `AbortRetryIgnore`, `OK`, `OKCancel`, `RetryCancel`, and `YesNoCancel`. Opens a new window for a simple dialogue with the user. All other windows of the process are blocked until the user closes the dialogue window

With `System.Windows.Forms.OpenFileDialog` you can ask the user to select an existing filename · `OpenFileDialog` as demonstrated in Listing 23.17 and Figure 23.15.

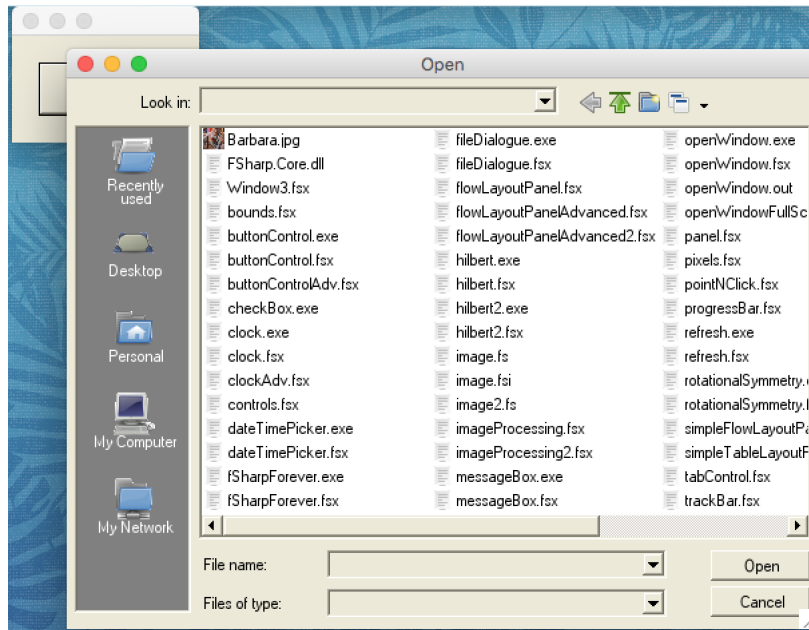


Figure 23.15: Ask the user for a filename to read from. See Listing 23.17.

Listing 23.17 winforms/openFileDialog.fsx:
Create the OpenFileDialog, see also Figure 23.15.

```

1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form ()
6  win.ClientSize <- Size (140, 80)
7
8  let button = new Button ()
9  win.Controls.Add button
10 button.Size <- new Size (100, 40)
11 button.Location <- new Point (20, 20)
12 button.Text <- "Click me"
13 // Open a message box, when button clicked
14 let buttonClicked (e : EventArgs) =
15     let opendlg = new OpenFileDialog()
16     let okOrCancel = opendlg.ShowDialog()
17     printfn "The user pressed %A and selected %A" okOrCancel
18     opendlg.FileName
19 button.Click.Add buttonClicked
20 Application.Run win

```

Similarly to `OpenFileDialog`, `System.Windows.Forms.SaveFileDialog` asks for a file name, but if `SaveFileDialog` an existing file is selected, then the user will be asked to confirm the choice.




Figure 23.16: A panel including a label and a text input field, see Listing 23.18.

23.6 Organising controls

It is often useful to organise the controls in groups, and such groups are called *Panels* in WinForms. An example of creating a `System.Windows.Forms.Panel` that includes a `System.Windows.Forms.TextBox` and `System.Windows.Forms.Label` for getting user input is shown in Listing 23.18 and Figure 23.16.

Listing 23.18 winforms/panel.fsx:
Create a panel, label, text input controls.

```

1  open System.Drawing
2  open System.Windows.Forms
3
4  let win = new Form () // Create a window form
5  win.ClientSize <- new Size (200, 100)
6
7  // Customize the Panel control
8  let panel = new Panel ()
9  panel.ClientSize <- new Size (160, 60)
10 panel.Location <- new Point (20,20)
11 panel.BorderStyle <- BorderStyle.Fixed3D
12 win.Controls.Add panel // Add panel to window
13
14 // Customize the Label and TextBox controls
15 let label = new Label ()
16 label.ClientSize <- new Size (120, 20)
17 label.Location <- new Point (15,5)
18 label.Text <- "Input"
19 panel.Controls.Add label // add label to panel
20
21 let textBox = new TextBox ()
22 textBox.ClientSize <- new Size (120, 20)
23 textBox.Location <- new Point (20,25)
24 textBox.Text <- "Initial text"
25 panel.Controls.Add textBox // add textbox to panel
26
27 Application.Run win // Start the event-loop

```

The label and textbox are children of the panel, and the main advantage of using panels is that the coordinates of the children are relative to the top left corner of the panel. I.e., moving the panel will move the label and the textbox at the same time.

A very flexible panel is the *System.Windows.Forms.FlowLayoutPanel*, which arranges its objects according to the space available. This is useful for graphical user interfaces targeting varying device sizes, such as a computer monitor and a tablet, and it also allows the program to gracefully adapt, when the user changes window size. A demonstration of *System.Windows.Forms.FlowLayoutPanel* together with *System.Windows.Forms.CheckBox* and *System.Windows.Forms.RadioButton* is given in Listing 23.20 and in Figure 23.17. The program illustrates how the button elements flow under four possible *System.Windows.FlowDirections*, and how *System.Windows.WrapContents* influences, what happens to content that flows outside the panels region.

Listing 23.19 winforms/flowLayoutPanel.fsx:
Create a FlowLayoutPanel, with checkbox and radiobuttons.

```

1  open System.Windows.Forms
2  open System.Drawing
3
4  let flowLayoutPanel = new FlowLayoutPanel ()
5  let buttonLst =
6      [(new Button (), "Button0");
7       (new Button (), "Button1");
8       (new Button (), "Button2");
9       (new Button (), "Button3")]
10 let panel = new Panel ()
11 let wrapContentsCheckBox = new CheckBox ()
12 let initiallyWrapped = true
13 let radioButtonLst =
14     [(new RadioButton (), (3, 34), "TopDown", FlowDirection.TopDown);
15      (new RadioButton (), (3, 62), "BottomUp", FlowDirection.BottomUp);
16      (new RadioButton (), (111, 34), "LeftToRight",
17       FlowDirection.LeftToRight);
18      (new RadioButton (), (111, 62), "RightToLeft",
19       FlowDirection.RightToLeft)]
20
21 // customize buttons
22 for (btn, txt) in buttonLst do
23     btn.Text <- txt
24
25 // customize wrapContentsCheckBox
26 wrapContentsCheckBox.Location <- new Point (3, 3)
27 wrapContentsCheckBox.Text <- "Wrap Contents"
28 wrapContentsCheckBox.Checked <- initiallyWrapped
29 wrapContentsCheckBox.CheckedChanged.Add (fun _ ->
30     flowLayoutPanel.WrapContents <- wrapContentsCheckBox.Checked)
31
32 // customize radio buttons
33 for (btn, loc, txt, dir) in radioButtonLst do
34     btn.Location <- new Point (fst loc, snd loc)
35     btn.Text <- txt
36     btn.Checked <- flowLayoutPanel.FlowDirection = dir
37     btn.CheckedChanged.Add (fun _ -> flowLayoutPanel.FlowDirection <- dir)

```

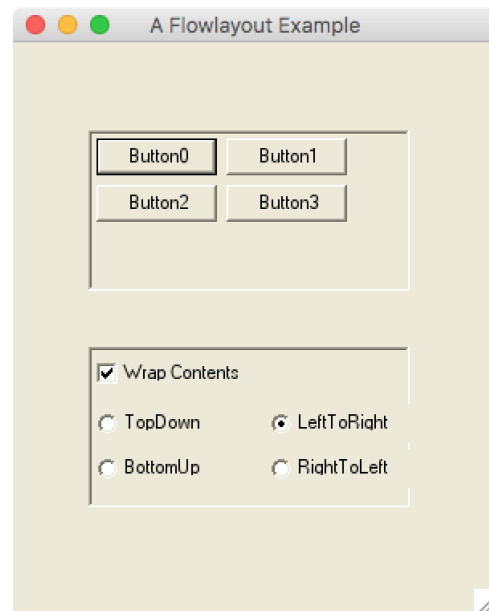


Figure 23.17: Demonstration of the FlowLayoutPanel panel, CheckBox, and RadioButton controls, see Listing 23.20.

Listing 23.20 winforms/flowLayoutPanel.fsx:

Create a FlowLayoutPanel, with checkbox and radiobuttons. Continued from Listing 23.19.

```

36 // customize flowLayoutPanel
37 for (btn, txt) in buttonLst do
38     flowLayoutPanel.Controls.Add btn
39 flowLayoutPanel.Location <- new Point (47, 55)
40 flowLayoutPanel.BorderStyle <- BorderStyle.Fixed3D
41 flowLayoutPanel.WrapContents <- initiallyWrapped
42
43 // customize panel
44 panel.Controls.Add (wrapContentsCheckBox)
45 for (btn, loc, txt, dir) in radioButtonLst do
46     panel.Controls.Add (btn)
47 panel.Location <- new Point (47, 190)
48 panel.BorderStyle <- BorderStyle.Fixed3D
49
50 // Create a window, add controls, and start event-loop
51 let win = new Form ()
52 win.ClientSize <- new Size (302, 356)
53 win.Controls.Add flowLayoutPanel
54 win.Controls.Add panel
55 win.Text <- "A Flowlayout Example"
56 Application.Run win

```

A walkthrough of the program is as follows. The goal is to make 2 areas, one giving the user control over display parameters, and another displaying the result of the user's choices. These are FlowLayoutPanel and Panel. In the FlowLayoutPanel there are four Buttons, to be displayed in a region, that is not tall enough for the buttons to be shown in vertical sequence and not wide enough to be shown in horizontal sequence. Thus the FlowDirection rules come into play, i.e., the buttons are added in sequence as they are named, and the default FlowDirection.LeftToRight arranges places the buttonLst.[0] in

the top left corner, and `buttonLst.[1]` to its right. Other flow directions do it differently, and the reader is encouraged to experiment with the program.

The program in Listing 23.20 has not completely separated the semantic blocks of the interface and relies on explicitly setting of coordinates of controls. This can be avoided by using nested panels. E.g., in Listing 23.22, the program has been rewritten as a nested set of `FloatLayoutPanel` in three groups: The button panel, the checkbox, and the radio button panel. Adding a `Resize` event handler for the window to resize the outermost panel according to the outer window, allows for the three groups to change position relative to each other, which results in three different views all shown in Figure 23.18.

Listing 23.21 `winforms/flowLayoutPanelAdvanced.fsx:`
Create nested `FlowLayoutPanel`.

```

1  open System.Windows.Forms
2  open System.Drawing
3  open System
4
5  let win = new Form ()
6  let mainPanel = new FlowLayoutPanel ()
7  let mainPanelBorder = 5
8  let flowLayoutPanel = new FlowLayoutPanel ()
9  let buttonLst =
10     [(new Button (), "Button0");
11      (new Button (), "Button1");
12      (new Button (), "Button2");
13      (new Button (), "Button3")]
14  let wrapContentsCheckBox = new CheckBox ()
15  let panel = new FlowLayoutPanel ()
16  let initiallyWrapped = true
17  let radioButtonLst =
18     [(new RadioButton (), "TopDown", FlowDirection.TopDown);
19      (new RadioButton (), "BottomUp", FlowDirection.BottomUp);
20      (new RadioButton (), "LeftToRight", FlowDirection.LeftToRight);
21      (new RadioButton (), "RightToLeft", FlowDirection.RightToLeft)]
22
23  // customize buttons
24  for (btn, txt) in buttonLst do
25     btn.Text <- txt
26
27  // customize radio buttons
28  for (btn, txt, dir) in radioButtonLst do
29     btn.Text <- txt
30     btn.Checked <- flowLayoutPanel.FlowDirection = dir
31     btn.CheckedChanged.Add (fun _ -> flowLayoutPanel.FlowDirection <- dir)
32
33  // customize flowLayoutPanel
34  for (btn, txt) in buttonLst do
35     flowLayoutPanel.Controls.Add btn
36  flowLayoutPanel.BorderStyle <- BorderStyle.Fixed3D
37  flowLayoutPanel.WrapContents <- initiallyWrapped
38
39  // customize wrapContentsCheckBox
40  wrapContentsCheckBox.Text <- "Wrap Contents"
41  wrapContentsCheckBox.Checked <- initiallyWrapped
42  wrapContentsCheckBox.CheckedChanged.Add (fun _ ->
43     flowLayoutPanel.WrapContents <- wrapContentsCheckBox.Checked)

```

Listing 23.22 winforms/flowLayoutPanelAdvanced.fsx:
Create nested FlowLayoutPanel. Continued from Listing 23.21.

```
44 // customize panel
45 // changing border style changes ClientSize
46 panel.BorderStyle <- BorderStyle.Fixed3D
47 let width = panel.ClientSize.Width / 2 - panel.Margin.Left -
    panel.Margin.Right
48 for (btn, txt, dir) in radioButtonLst do
49     btn.Width <- width
50     panel.Controls.Add (btn)
51
52 mainPanel.Location <- new Point (mainPanelBorder, mainPanelBorder)
53 mainPanel.BorderStyle <- BorderStyle.Fixed3D
54 mainPanel.Controls.Add flowLayoutPanel
55 mainPanel.Controls.Add wrapContentsCheckBox
56 mainPanel.Controls.Add panel
57
58 // customize window, add controls, and start event-loop
59 win.ClientSize <- new Size (220, 256)
60 let windowResize _ =
61     let size = win.DisplayRectangle.Size
62     mainPanel.Size <- new Size (size.Width - 2 * mainPanelBorder,
        size.Height - 2 * mainPanelBorder)
63 windowResize ()
64 win.Resize.Add windowResize
65 win.Controls.Add mainPanel
66 win.Text <- "Advanced Flowlayout"
67 Application.Run win
```

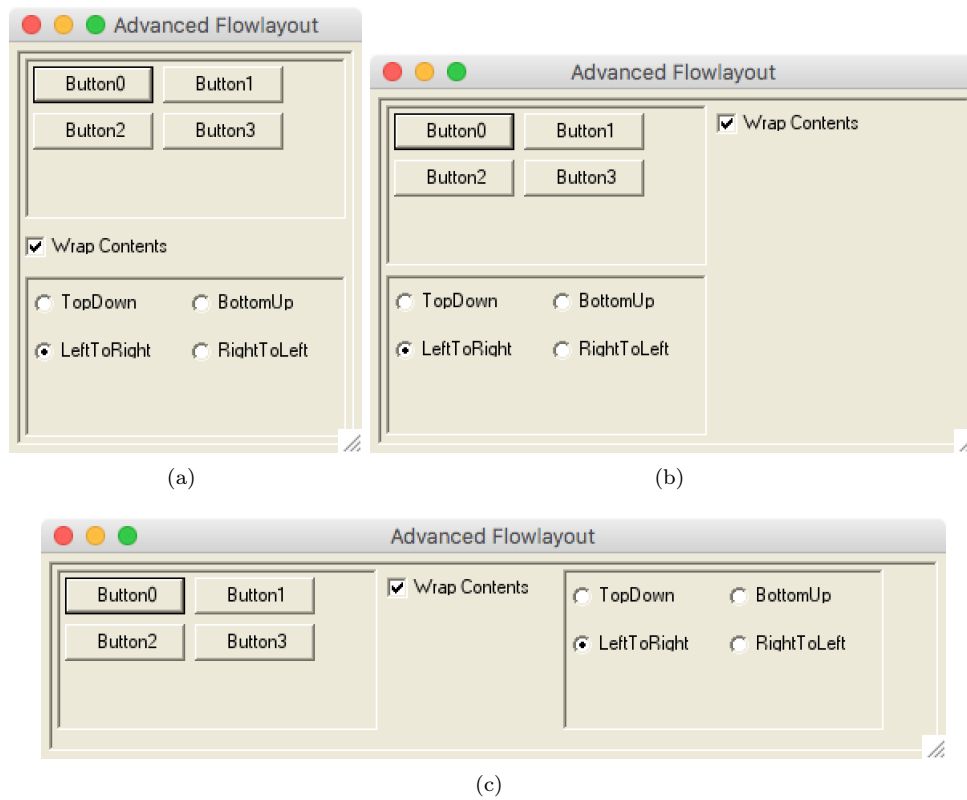


Figure 23.18: Nested `FlowLayoutPanel`, see Listing 23.22, allows for dynamic arrangement of content. Content flows, when the window is resized.

24 | The Event-driven programming paradigm

In *event-driven programming*, the flow of the program is determined by *events* such as the user moving the mouse, an alarm going off, a message arrives from another program, or an exception is thrown, and is very common for programs with extensive interaction with a user such as a graphical user-interface. The events are monitored by *listeners*, and the programmer can *handlers*, which are *call-back* functions to be executed when an event occurs. In event-driven programs, there is almost always a main loop, to which the program relinquishes control to when all handlers have been set up. Event-driven programs can be difficult to test since they often rely on difficult to automate mechanisms for triggering events, e.g., testing a graphical user-interface often requires users to point-and-click, which is very slow compared to automatic unit-testing.

- event-driven programming
- events
- listeners
- handlers
- call-back

25 | Where to go from here

You have now learned to program in a number of important paradigms and mastered the basics of F#, so where are good places to go now? I will highlight a number of options:

Program, program, program

You are at this stage no longer a novice programmer, so it is time for you to use your skills and create programs that solve problems. I have always found great inspiration in interacting with other domains and seek solutions by programming. And experience is a must if you want to become a good programmer, since your newly acquired skills need to settle in your mind, and you need to be exposed to new problems that require you to adapt and develop your skills.

Learn to use an Integrated Development Environment effectively

An Integrated Development Environment (IDE) is a tool that may increase your coding efficiency. IDEs can help you get started in different environments, such as on a laptop or a phone, it can quickly give you an overview of available options when you are programming. E.g., all IDEs will show you available members for identifiers as you type reducing time to search members and reducing the risk of spelling errors. Many IDEs will also help you to quickly refactor your code, e.g., by highlighting all occurrences of a name in a scope and letting you change it once in all places.

In this book, we have emphasized the console, compiling and running from the console is the basis of which all IDEs build, and many of the problems with using IDEs efficiently are related to understanding how it can best help you compiling and running programs.

Learn other cool features of F#

F# is a large language with many features. Some have been presented in this book, but more advanced topics have been neglected. Examples are:

- regular expressions: Much computations concerns processing of text. Regular expressions is a simple but powerful language for search and replace in strings. F# has built-in support for regular expressions as `System.Text.RegularExpressions`.
- sequence `seq`: All list type data structures in F# are built on sequences. Sequences are, however, more than lists and arrays. A key feature is that sequences can effectively contain large even infinite ordered lists, which you do not necessarily need or use, i.e., they are lazy and only compute its elements as needed. Sequences are programmed using computation expressions.
- computation expressions: Sequential expressions is an example of a computation expressions, e.g., the sequence of squares $i^2, i = 0..9$ can be written as `seq {for i in 0 .. 9 -> i * i}`
- asynchronous computations `async`: F# has a native implementation of asynchronous computation, which means that you can very easily set up computations that run independently on others such that they do not block for each other. This is extremely convenient if you, e.g., need to process a list of homepages, where each homepage may be slow to read, such that reading them in sequence will be slow, but with asynchronous computations, then they

can easily be read in parallel with a huge speedup for the total task as a result. Asynchronous workflows rely on computation expressions.

Learn another programming language

F# is just one of a great number of programming languages, and you should not limit yourself. Languages are often designed to be used for particular tasks, and when looking to solve a problem, you would do well in selecting the language that best fits the task. C# is an example from the Mono family, which emphasizes object-oriented programming, and many of the built-in libraries in F# are written in C#. C++ and C are ancestors of C# and are very popular since they allow for great control over the computer at the expense of programming convenience. Python is a popular prototyping language which emphasizes interactive programming like `fsharp`, and it is seeing a growing usage in web-backends and machine learning. And the list goes on. To get an idea of the wealth of languages, browse to <http://www.99-bottles-of-beer.net> which has examples of solutions to the simple problem: Write a program that types the lyrics of song “99 bottles of beer on the wall” and stops. At present many solutions in more than 1500 different languages have been submitted.

A | The console in Windows, MacOS X, and Linux

Almost all popular operating systems are accessed through a user-friendly *graphical user interface* (*GUI*) that is designed to make typical tasks easy to learn to solve. As a computer programmer, you often need to access some of the functionalities of the computer, which, unfortunately, are sometimes complicated by this particular graphical user interface. The *console*, also called the *terminal* and the *Windows command line*, is the right hand of a programmer. The console is a simple program that allows you to complete text commands. Almost all the tasks that can be done with the graphical user interface can be done in the console and vice versa. Using the console, you will benefit from its direct control of the programs we write, and in your education, you will benefit from the fast and raw information you get through the console.

- graphical user interface
- GUI
- console
- terminal
- Windows command line

A.1 The basics

When you open a *directory* or *folder* in your preferred operating system, the directory will have a location in the file system, whether from the console or through the operating system's graphical user interface. The console will almost always be associated with a particular directory or folder in the file system, and it is said that it is the directory that the console is in. The exact structure of file systems varies between Linux, MacOS X, and Windows, but common is that it is a hierarchical structure. This is illustrated in Figure A.1.

- directory
- folder

There are many predefined console commands, available in the console, and you can make your own. In the following, we will review the most important commands in the three different operating systems. These are summarized in Table A.1.

A.2 Windows

In this section we will discuss the commands summarized in Table A.1. Windows 7 and earlier versions: To open the console, press **Start->Run** in the lower left corner, and then type **cmd** in the box. In Windows 8 and 10, you right-click on the windows icon, choose **Run** or equivalent in your local language, and type **cmd**. Alternatively you can type **Windows-key + R**. Now you should open a console window with a prompt showing something like Listing A.1.

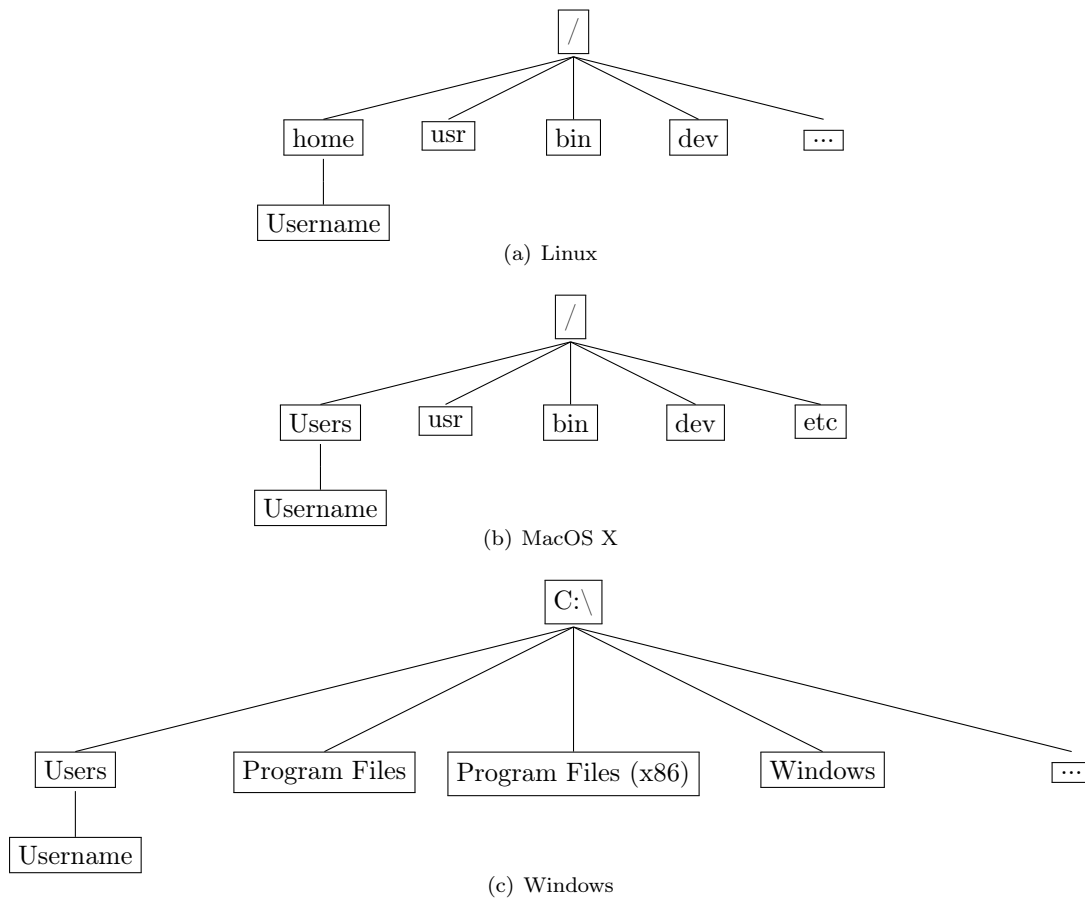


Figure A.1: The top file hierarchy levels of common operating systems.

Windows	MacOS X/Linux	Description
<code>dir</code>	<code>ls</code>	Show content of present directory.
<code>cd <dir></code>	<code>cd <dir></code>	Change present directory to <code><dir></code> .
<code>mkdir <dir></code>	<code>mkdir <dir></code>	create directory <code><dir></code> .
<code>rmdir <dir></code>	<code>rmdir <dir></code>	delete <code><dir></code> (Warning: cannot be reverted).
<code>move <file> <file or dir></code>	<code>mv <file> <file or dir></code>	Move <code><fil></code> to <code><file or dir></code> .
<code>copy <file> <file></code>	<code>cp <file> <file></code>	Create a new file called <code><file></code> as a copy of <code><file></code> .
<code>del <file></code>	<code>rm <file></code>	delete <code><file></code> (Warning: cannot be reverted).
<code>echo <string or variable></code>	<code>echo <string or variable></code>	Write a string or content of a variable to screen.

Table A.1: The most important console commands for Windows, MacOS X, and Linux.

Listing A.1: The Windows console.

```

1 Microsoft Windows [Version 6.1.7601]
2 Copyright (c) 2009 Microsoft Corporation. All rights reserved.
3
4 C:\Users\sporrington>

```

To see which files are in the directory, use *dir* as shown in Listing A.2.

· *dir*

Listing A.2: Directory listing with dir.

```

1 C:\Users\sporrington>dir
2 Volume in drive C has no label.
3 Volume Serial Number is 94F0-31BD
4
5 Directory of C:\Users\sporrington
6
7 30-07-2015 15:23 <DIR> .
8 30-07-2015 15:23 <DIR> ..
9 30-07-2015 14:27 <DIR> Contacts
10 30-07-2015 14:27 <DIR> Desktop
11 30-07-2015 17:40 <DIR> Documents
12 30-07-2015 15:11 <DIR> Downloads
13 30-07-2015 14:28 <DIR> Favorites
14 30-07-2015 14:27 <DIR> Links
15 30-07-2015 14:27 <DIR> Music
16 30-07-2015 14:27 <DIR> Pictures
17 30-07-2015 14:27 <DIR> Saved Games
18 30-07-2015 17:27 <DIR> Searches
19 30-07-2015 14:27 <DIR> Videos
20 0 File(s) 0 bytes
21 13 Dir(s) 95.004.622.848 bytes free
22
23 C:\Users\sporrington>

```

We see that there are no files and thirteen directories (DIR). The columns tell from left to right, the date and time of their creation, the file size or if it is a folder, and the name file or directory name. The first two folders “.” and “..” are found in each folder and refers to this folder as well as the one above in the hierarchy. In this case, the folder “.” is an alias for C:\Users\tracking and “..” for C:\Users.

Use *cd* to change directory, e.g., to Documents as in Listing A.3.

· *cd*

Listing A.3: Change directory with cd.

```

1 C:\Users\sporrington>cd Documents
2
3 C:\Users\sporrington\Documents>

```

The directory Documents is that Windows Explorer selects, when you press the Documents short-cut in Explorer. Note that some systems translate default filenames so their names may be given different names in different languages in the graphical user interface as compared to the console.

You can use *mkdir* to create a new directory called, e.g., myFolder as illustrated in Listing A.4.

· *mkdir*

Listing A.4: Creating a directory with mkdir.

```

1 C:\Users\sporrington\Documents>mkdir myFolder
2
3 C:\Users\sporrington\Documents>dir
4 Volume in drive C has no label.
5 Volume Serial Number is 94F0-31BD
6
7 Directory of C:\Users\sporrington\Documents
8
9 30-07-2015  19:17    <DIR>          .
10 30-07-2015  19:17    <DIR>          ..
11 30-07-2015  19:17    <DIR>          myFolder
12                0 File(s)                0 bytes
13                3 Dir(s)  94.656.638.976 bytes free
14
15 C:\Users\sporrington\Documents>

```

By using `dir` we inspect the result.

Files can be created by, e.g., *echo* and *redirection* as demonstrated in Listing A.5.

- `echo`
- *redirection*

Listing A.5: Creating a file with echo and redirection.

```

1 C:\Users\sporrington\Documents>echo "Hi" > hi.txt
2
3 C:\Users\sporrington\Documents>dir
4 Volume in drive C has no label.
5 Volume Serial Number is 94F0-31BD
6
7 Directory of C:\Users\sporrington\Documents
8
9 30-07-2015  19:18    <DIR>          .
10 30-07-2015  19:18    <DIR>          ..
11 30-07-2015  19:17    <DIR>          myFolder
12 30-07-2015  19:18                8 hi.txt
13                1 File(s)                8 bytes
14                3 Dir(s)  94.656.634.880 bytes free
15
16 C:\Users\sporrington\Documents>

```

To move the file `hi.txt` to the directory `myFolder` use *move* as shown in Listing A.6.

- `move`

Listing A.6: Move a file with move.

```

1 C:\Users\sporrington\Documents>move hi.txt myFolder
2 1 file(s) moved.
3
4 C:\Users\sporrington\Documents>

```

Finally, use *del* to delete a file and *rmdir* to delete a directory as shown in Listing A.7.

- `del`
- `rmdir`

Listing A.7: Delete files and directories with `del` and `rmdir`.

```

1 C:\Users\sporrington\Documents>cd myFolder
2
3 C:\Users\sporrington\Documents\myFolder>del hi.txt
4
5 C:\Users\sporrington\Documents\myFolder>cd ..
6
7 C:\Users\sporrington\Documents>rmdir myFolder
8
9 C:\Users\sporrington\Documents>dir
10 Volume in drive C has no label.
11 Volume Serial Number is 94F0-31BD
12
13 Directory of C:\Users\sporrington\Documents
14
15 30-07-2015  19:20    <DIR>          .
16 30-07-2015  19:20    <DIR>          ..
17                0 File(s)                0 bytes
18                2 Dir(s)  94.651.142.144 bytes free
19
20 C:\Users\sporrington\Documents>

```

The commands available from the console must be in its *search path*. The search path can be seen using `echo` as shown in Listing A.8.

Listing A.8: Displaying the search path.

```

1 C:\Users\sporrington\Documents>echo %Path%
2 C:\Windows\system32;C:\Windows;C:\Windows\System32\Wbem;
   C:\Windows\System32\WindowsPowerShell\v1.0\;"\Program
   Files\emacs-24.5\bin\"
3
4 C:\Users\sporrington\Documents>

```

The path can be changed using the Control panel in the graphical user interface. In Windows 7, choose the Control panel, choose System and Security → System → Advanced system settings → Environment Variables. In Windows 10 you can find this window by searching for “Environment” in the Control panel. In the window’s System variables box, double-click on Path and add or remove a path from the list. The search path is a list of paths separated by “;”. Beware, Windows uses the search path for many different tasks, so remove only paths that you are certain are not used for anything.

A useful feature of the console is that you can use the `tab`-key to cycle through filenames. E.g., if you write `cd` followed by a space and `tab` a couple of times, then the console will suggest you the available directories.

A.3 MacOS X and Linux

MacOS X (OSX) and Linux are very similar, and both have the option of using *bash* as console. It is in the standard console on MacOS X and on many Linux distributions. A summary of the most important *bash* commands are shown in Table A.1. In MacOS X, you find the console by opening

Finder and navigating to **Applications** → **Utilities** → **Terminal**. In Linux, the console can be started by typing **Ctrl + Alt + T**. Some Linux distributions have other key-combinations such as **Super + T**.

Once opened, the console is shown in a window with content as shown in Listing A.9.

Listing A.9: The Windows console.

```
1 Last login: Thu Jul 30 11:52:07 on ttys000
2 FN11194:~ sporring$
```

“FN11194” is the name of the computer, the character `~` is used as an alias for the user’s home directory, and “sporring” is the username for the user presently logged onto the system. Use `ls` to see which files are present, as shown in Listing A.10.

Listing A.10: Display a directory content with `ls`.

```
1 FN11194:~ sporring$ ls
2 Applications      Documents      Library       Music         Public
3 Desktop           Downloads     Movies        Pictures
4 FN11194:~ sporring$
```

More details about the files are available by using flags to `ls` as demonstrated in Listing A.11.

Listing A.11: Display extra information about files using flags to `ls`.

```
1 FN11194:~ sporring$ ls -l
2 drwx----- 6 sporring staff 204 Jul 30 14:07 Applications
3 drwx-----+ 32 sporring staff 1088 Jul 30 14:34 Desktop
4 drwx-----+ 76 sporring staff 2584 Jul 2 15:53 Documents
5 drwx-----+ 4 sporring staff 136 Jul 30 14:35 Downloads
6 drwx-----@ 63 sporring staff 2142 Jul 30 14:07 Library
7 drwx-----+ 3 sporring staff 102 Jun 29 21:48 Movies
8 drwx-----+ 4 sporring staff 136 Jul 4 17:40 Music
9 drwx-----+ 3 sporring staff 102 Jun 29 21:48 Pictures
10 drwxr-xr-x+ 5 sporring staff 170 Jun 29 21:48 Public
11 FN11194:~ sporring$
```

The flag `-l` means long, and many other flags can be found by querying the built-in manual with `man ls`. The output is divided into columns, where the left column shows a number of codes: “d” means they are directory followed by three sets of optional “rwx” denoting whether the owner, the associated group of users and anyone can “r” - read, “w” - write, or “x” - execute the files. In this case, only the owner can do any of the three. For directories, “x” means permission to enter. The second column can often be ignored, but shows how many links there are to the file or directory. Then follows the username of the owner, which in this case is `sporring`. The files are also associated with a group of users, and in this case, they all are associated with the group called `staff`. Then follows the file or directory size, the date of last change, and the file or directory name. There are always two hidden directories: “.” and “..”, where “.” is an alias for the present directory, and “..” for the directory above. Hidden files will be shown with the `-a` flag.

Use `cd` to change to the directory `Documents` as shown in Listing A.12.

· `cd`

Listing A.12: Change directory with `cd`.

```
1 FN11194:~ sporring$ cd Documents/  
2 FN11194:Documents sporring$
```

Note that some graphical user interfaces translate standard filenames and directories to the local language, such that navigating using the graphical user interface will reveal other files and directories, which, however, are aliases.

You can yourself create a new directory using *`mkdir`* as demonstrated in Listing A.13.

· `mkdir`

Listing A.13: Creating a directory using `mkdir`.

```
1 FN11194:Documents sporring$ mkdir myFolder  
2 FN11194:Documents sporring$ ls  
3 myFolder  
4 FN11194:tmp sporring$
```

A file can be created using *`echo`* and with *redirection* as shown in Listing A.14.

· `echo`
· redirection

Listing A.14: Creating a file with `echo` and redirection.

```
1 FN11194:Documents sporring$ echo "hi" > hi.txt  
2 FN11194:Documents sporring$ ls  
3 hi.txt          myFolder
```

To move the file `hi.txt` into `myFolder`, use *`mv`*. This is demonstrated in Listing A.15.

· `mv`

Listing A.15: Moving files with `mv`.

```
1 FN11194:Documents sporring$ echo mv hi.txt myFolder/  
2 FN11194:Documents sporring$
```

To delete the file and the directory, use *`rm`* and *`rmdir`* as shown in Listing A.16.

· `rm`
· `rmdir`

Listing A.16: The Windows console.

```
1 FN11194:Documents sporring$ cd myFolder/  
2 FN11194:myFolder sporring$ rm hi.txt  
3 FN11194:myFolder sporring$ cd ..  
4 FN11194:Documents sporring$ rmdir myFolder/  
5 FN11194:Documents sporring$ ls  
6 FN11194:Documents sporring$
```

Only commands found on the *search-path* are available in the console. The content of the search-path is seen using the `echo` command as demonstrated in Listing A.17.

· search-path

Listing A.17: The content of the search-path.

```
1 FN11194:Documents sporring$ echo $PATH
2 /Applications/Maple 17:/Applications/PackageMaker.app/Contents/MacOS:/
  /Applications/MATLAB_R2014b.app/bin:/opt/local/bin:/opt/local/sbin:
  /usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/opt/X11/bin:
  /Library/TeX/texbin
3 FN11194:Documents sporring$
```

The search-path can be changed by editing the setup file for Bash. On MacOS X it is called `~/.profile`, and on Linux it is either `~/.bash_profile` or `~/.bashrc`. Here new paths can be added by adding the following line: `export PATH=<new path>:<another new path>:$PATH`.

A useful feature of Bash is that the console can help you write commands. E.g., if you write `fs` followed by pressing the `tab`-key, and if `Mono` is in the search-path, then Bash will typically respond by completing the line as `fsharp`, and by further pressing the `tab`-key some times, Bash will show the list of options, typically `fshpari` and `fsharpc`. Also, most commands have an extensive manual, which can be accessed using the `man` command. E.g., the manual for `rm` is retrieved by `man rm`.

B | Number systems on the computer

B.1 Binary numbers

Humans like to use the *decimal number* system for representing numbers. Decimal numbers are *base 10* means that for a number consisting of a sequence of digits separated by a *decimal point*, where each *digit* can have values $d \in \{0, 1, 2, \dots, 9\}$ and the weight of each digit is proportional to its place in the sequence of digits with respect to the decimal point, i.e., the number $357.6 = 3 \cdot 10^2 + 5 \cdot 10^1 + 7 \cdot 10^0 + 6 \cdot 10^{-1}$ or in general:

$$v = \sum_{i=-m}^n d_i 10^i \quad (\text{B.1})$$

The basic unit of information in almost all computers is the binary digit or *bit* for short. A *binary number* consists of a sequence of binary digits separated by a decimal point, where each digit can have values $b \in \{0, 1\}$, and the base is 2. The general equation is,

$$v = \sum_{i=-m}^n b_i 2^i \quad (\text{B.2})$$

and examples are $1011.1_2 = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 + 1 \cdot 2^{-1} = 11.5$. Notice that we use subscript 2 to denote a binary number, while no subscript is used for decimal numbers. The left-most bit is called the *most significant bit*, and the right-most bit is called the *least significant bit*. Due to typical organisation of computer memory, 8 binary digits is called a *byte*, and 32 digits a *word*.

Other number systems are often used, e.g., *octal numbers*, which are base 8 numbers, where each digit is $o \in \{0, 1, \dots, 7\}$. Octals are useful short-hand for binary, since 3 binary digits maps to the set of octal digits. Likewise, *hexadecimal numbers* are base 16 with digits $h \in \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, b, c, d, e, f\}$, such that $a_{16} = 10$, $b_{16} = 11$ and so on. Hexadecimals are convenient since 4 binary digits map directly to the set of octal digits. Thus $367 = 101101111_2 = 557_8 = 16f_{16}$. A list of the integers 0–63 in various bases is given in Table B.1.

B.2 IEEE 754 floating point standard

The set of real numbers also called *reals* includes all fractions and irrational numbers. It is infinite in size both in the sense that there is no largest nor smallest number, and between any 2 given numbers there are infinitely many numbers. Reals are widely used for calculation, but since any computer only has finite memory, it is impossible to represent all possible reals on one. Hence, any computation performed on a computer with reals must rely on approximations. *IEEE 754 double precision floating-point format (binary64)*, known as a *double*, is a standard for representing an approximation of reals

Dec	Bin	Oct	Hex	Dec	Bin	Oct	Hex
0	0	0	0	32	100000	40	20
1	1	1	1	33	100001	41	21
2	10	2	2	34	100010	42	22
3	11	3	3	35	100011	43	23
4	100	4	4	36	100100	44	24
5	101	5	5	37	100101	45	25
6	110	6	6	38	100110	46	26
7	111	7	7	39	100111	47	27
8	1000	10	8	40	101000	50	28
9	1001	11	9	41	101001	51	29
10	1010	12	a	42	101010	52	2a
11	1011	13	b	43	101011	53	2b
12	1100	14	c	44	101100	54	2c
13	1101	15	d	45	101101	55	2d
14	1110	16	e	46	101110	56	2e
15	1111	17	f	47	101111	57	2f
16	10000	20	10	48	110000	60	30
17	10001	21	11	49	110001	61	31
18	10010	22	12	50	110010	62	32
19	10011	23	13	51	110011	63	33
20	10100	24	14	52	110100	64	34
21	10101	25	15	53	110101	65	35
22	10110	26	16	54	110110	66	36
23	10111	27	17	55	110111	67	37
24	11000	30	18	56	111000	70	38
25	11001	31	19	57	111001	71	39
26	11010	32	1a	58	111010	72	3a
27	11011	33	1b	59	111011	73	3b
28	11100	34	1c	60	111100	74	3c
29	11101	35	1d	61	111101	75	3d
30	11110	36	1e	62	111110	76	3e
31	11111	37	1f	63	111111	77	3f

Table B.1: A list of the integers 0–63 in decimal, binary, octal, and hexadecimal.

using 64 bits. These bits are divided into 3 parts: sign, exponent and fraction,

$$s e_1 e_2 \dots e_{11} m_1 m_2 \dots m_{52},$$

where s , e_i , and m_j are binary digits. The bits are converted to a number using the equation by first calculating the exponent e and the mantissa m ,

$$e = \sum_{i=1}^{11} e_i 2^{11-i}, \quad (\text{B.3})$$

$$m = \sum_{j=1}^{52} m_j 2^{-j}. \quad (\text{B.4})$$

I.e., the exponent is an integer, where $0 \leq e < 2^{11}$, and the mantissa is a rational, where $0 \leq m < 1$. For most combinations of e and m the real number v is calculated as,

$$v = (-1)^s (1 + m) 2^{e-1023} \quad (\text{B.5})$$

with the exception that

	$m = 0$	$m \neq 0$
$e = 0$	$v = (-1)^s 0$ (signed zero)	$v = (-1)^s m 2^{1-1023}$ (subnormals)
$e = 2^{11} - 1$	$v = (-1)^s \infty$	$v = (-1)^s \text{NaN}$ (not a number)

· subnormals
· NaN
· not a number

where $e = 2^{11} - 1 = 1111111111_2 = 2047$. The largest and smallest number that is not infinity is thus

$$e = 2^{11} - 2 = 2046 \quad (\text{B.6})$$

$$m = \sum_{j=1}^{52} 2^{-j} = 1 - 2^{-52} \simeq 1. \quad (\text{B.7})$$

$$v_{\max} = \pm (2 - 2^{-52}) 2^{1023} \simeq \pm 2^{1024} \simeq \pm 10^{308} \quad (\text{B.8})$$

The density of numbers varies in such a way that when $e - 1023 = 52$, then

$$v = (-1)^s \left(1 + \sum_{j=1}^{52} m_j 2^{-j} \right) 2^{52} \quad (\text{B.9})$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{-j} 2^{52} \right) \quad (\text{B.10})$$

$$= \pm \left(2^{52} + \sum_{j=1}^{52} m_j 2^{52-j} \right) \quad (\text{B.11})$$

$$\stackrel{k=52-j}{=} \pm \left(2^{52} + \sum_{k=51}^0 m_{52-k} 2^k \right) \quad (\text{B.12})$$

which are all integers in the range $2^{52} \leq |v| < 2^{53}$. When $e - 1023 = 53$, then the same calculation gives

$$v \stackrel{k=53-j}{=} \pm \left(2^{53} + \sum_{k=52}^1 m_{53-k} 2^k \right) \quad (\text{B.13})$$

which are every second integer in the range $2^{53} \leq |v| < 2^{54}$, and so on for larger values of e . When $e - 1023 = 51$, the same calculation gives,

$$v \stackrel{k=51-j}{=} \pm \left(2^{51} + \sum_{k=50}^{-1} m_{51-k} 2^k \right) \quad (\text{B.14})$$

which is a distance between numbers of $1/2$ in the range $2^{51} \leq |v| < 2^{52}$, and so on for smaller values of e . Thus we may conclude that the distance between numbers in the interval $2^n \leq |v| < 2^{n+1}$ is 2^{n-52} , for $-1022 = 1 - 1023 \leq n < 2046 - 1023 = 1023$. For subnormals, the distance between numbers are

$$v = (-1)^s \left(\sum_{j=1}^{52} m_j 2^{-j} \right) 2^{-1022} \quad (\text{B.15})$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j} 2^{-1022} \right) \quad (\text{B.16})$$

$$= \pm \left(\sum_{j=1}^{52} m_j 2^{-j-1022} \right) \quad (\text{B.17})$$

$$\stackrel{k=-j-1022}{=} \pm \left(\sum_{j=-1023}^{-1074} m_{-k-1022} 2^k \right) \quad (\text{B.18})$$

which gives a distance between numbers of $2^{-1074} \simeq 10^{-323}$ in the range $0 < |v| < 2^{-1022} \simeq 10^{-308}$.

C | Commonly used character sets

Letters, digits, symbols, and space are the core of how we store data, write programs, and communicate with computers and each other. These symbols are in short called characters and represent a mapping between numbers, also known as codes, and a pictorial representation of the character. E.g., the ASCII code for the letter 'A' is 65. These mappings are for short called character sets, and due to differences in natural languages and symbols used across the globe, many different character sets are in use. E.g., the English alphabet contains the letters 'a' to 'z', which is shared by many other European languages, but which have other symbols and accents. For example, Danish has further the letters 'æ', 'ø', and 'å'. Many non-European languages have completely different symbols, where the Chinese character set is probably the most extreme, and some definitions contain 106,230 different characters albeit only 2,600 are included in the official Chinese language test at the highest level.

Presently, the most common character set used is Unicode Transformation Format (UTF), whose most popular encoding schemes are 8-bit (UTF-8) and 16-bit (UTF-16). Many other character sets exist, and many of the later build on the American Standard Code for Information Interchange (ASCII). The ISO-8859 codes were an intermediate set of character sets that are still in use, but which is greatly inferior to UTF. Here we will briefly give an overview of ASCII, ISO-8859-1 (Latin1), and UTF.

C.1 ASCII

The *American Standard Code for Information Interchange* (ASCII) [8], is a 7 bit code tuned for the letters of the English language, numbers, punctuation symbols, control codes and space, see Tables C.1 and C.2. The first 32 codes are reserved for non-printable control characters to control printers and similar devices or to provide meta-information. The meaning of each control character is not universally agreed upon.

- American Standard Code for Information Interchange
- ASCII

The code order is known as *ASCIIbetical order*, and it is sometimes used to perform arithmetic on codes, e.g., an uppercase letter with code c may be converted to lower case by adding 32 to its code. The ASCIIbetical order also has a consequence for sorting, i.e., when sorting characters according to their ASCII code, 'A' comes before 'a', which comes before the symbol '{'.

- ASCIIbetical order

C.2 ISO/IEC 8859

The ISO/IEC 8859 report http://www.iso.org/iso/catalogue_detail?csnumber=28245 defines 10 sets of codes specifying up to 191 codes and graphics characters using 8 bits. Set 1, also known as ISO/IEC 8859-1, Latin alphabet No. 1 or *Latin1*, covers many European languages and is designed to be compatible with ASCII, such that code for the printable characters in ASCII is the same in ISO 8859-1. Table C.3 shows the characters above 7e. Codes 00-1f and 7f-9f are undefined in ISO 8859-1.

- Latin1

x0+0x	00	10	20	30	40	50	60	70
00	NUL	DLE	SP	0	@	P	'	p
01	SOH	DC1	!	1	A	Q	a	q
02	STX	DC2	"	2	B	R	b	r
03	ETX	DC3	#	3	C	S	c	s
04	EOT	DC4	\$	4	D	T	d	t
05	ENQ	NAK	%	5	E	U	e	u
06	ACK	SYN	&	6	F	V	f	v
07	BEL	ETB	'	7	G	W	g	w
08	BS	CAN	(8	H	X	h	x
09	HT	EM)	9	I	Y	i	y
0A	LF	SUB	*	:	J	Z	j	z
0B	VT	ESC	+	;	K	[k	{
0C	FF	FS	,	<	L	\	l	
0D	CR	GS	—	=	M]	m	}
0E	SO	RS	.	>	N	^	n	~
0F	SI	US	/	?	O	_	o	DEL

Table C.1: ASCII

C.3 Unicode

Unicode is a character standard defined by the Unicode Consortium, <http://unicode.org> as the *Unicode Standard*. Unicode allows for 1,114,112 different codes. Each code is called a *code point*, which represents an abstract character. However, not all abstract characters require a unit of several code points to be specified. Code points are divided into 17 planes each with $2^{16} = 65,536$ code points. Planes are further subdivided into named *blocks*. The first plane is called the *Basic Multilingual plane* and its block of the first 128 code points is called the *Basic Latin block* and are identical to ASCII, see Table C.1, and code points 128-255 is called the *Latin-1 Supplement block*, and are identical to the upper range of ISO 8859-1, see Table C.3. Each code-point has a number of attributes such as the *Unicode general category*. Presently more than 128,000 code points are defined as covering 135 modern and historical writing systems, and obtained at <http://www.unicode.org/Public/UNIDATA/UnicodeData.txt>, which includes the code point, name, and general category.

- Unicode Standard
- code point
- blocks
- Basic Multilingual plane
- Basic Latin block
- Latin-1 Supplement block
- Unicode general category

A Unicode code point is an abstraction from the encoding and the graphical representation of a character. A code point is written as “U+” followed by its hexadecimal number, and for the Basic Multilingual plane 4 digits are used, e.g., the code point with the unique name LATIN CAPITAL LETTER A has the Unicode code point “U+0041”, and in this text it is visualized as ‘A’. More digits are used for code points of the remaining planes.

The general category is used to specify valid characters that not necessarily have a visual representation but possibly transforms text. Some categories and their letters in the first 256 code points are shown in Table C.5.

To store and retrieve code points, they must be encoded and decoded. A common encoding is *UTF-8*, which encodes code points as 1 to 4 bytes, and which is backward-compatible with ASCII and ISO 8859-1. Hence, in all 3 coding systems, the character with code 65 represents the character ‘A’. Another popular encoding scheme is *UTF-16*, which encodes characters as 2 or 4 bytes, but which is not backward-compatible with ASCII or ISO 8859-1. UTF-16 is used internally in many compilers, interpreters, and operating systems.

- UTF-8
- UTF-16

Code	Description
NUL	Null
SOH	Start of heading
STX	Start of text
ETX	End of text
EOT	End of transmission
ENQ	Enquiry
ACK	Acknowledge
BEL	Bell
BS	Backspace
HT	Horizontal tabulation
LF	Line feed
VT	Vertical tabulation
FF	Form feed
CR	Carriage return
SO	Shift out
SI	Shift in
DLE	Data link escape
DC1	Device control one
DC2	Device control two
DC3	Device control three
DC4	Device control four
NAK	Negative acknowledge
SYN	Synchronous idle
ETB	End of transmission block
CAN	Cancel
EM	End of medium
SUB	Substitute
ESC	Escape
FS	File separator
GS	Group separator
RS	Record separator
US	Unit separator
SP	Space
DEL	Delete

Table C.2: ASCII symbols.

x0+0x	80	90	A0	B0	C0	D0	E0	F0
00			NBSP	°	À	Ð	à	ð
01			¡	±	Á	Ñ	á	ñ
02			¢	²	Â	Ò	â	ò
03			£	³	Ã	Ó	ã	ó
04			¤	´	Ä	Ö	ä	ö
05			¥	µ	Å	Õ	å	õ
06			¦	¶	Æ	Ö	æ	ö
07			§	·	Ç	×	ç	÷
08			¨	¸	È	Ø	è	ø
09			©	¹	É	Ù	é	ù
0a			ª	º	Ê	Ú	ê	ú
0b			«	»	Ë	Û	ë	û
0c			¬	$\frac{1}{4}$	Ì	Ü	ì	ü
0d			SHY	$\frac{1}{2}$	Í	Ý	í	ý
0e			®	$\frac{3}{4}$	Î	Þ	î	þ
0f			¯	¸	Ï	ß	ï	ÿ

Table C.3: ISO-8859-1 (latin1) non-ASCII part. Note that the codes 7f – 9f are undefined.

Code	Description
NBSP	Non-breakable space
SHY	Soft hyphen

Table C.4: ISO-8859-1 special symbols.

General category	Code points	Name
Lu	U+0041–U+005A, U+00C0–U+00D6, U+00D8–U+00DE	Upper case letters
Ll	U+0061–U+007A, U+00B5, U+00DF–U+00F6, U+00F8–U+00FF	Lower case letter
Lt	None	Digraphic letter, with first part uppercase
Lm	None	Modifier letter
Lo	U+00AA, U+00BA	Gender ordinal indicator
Nl	None	Letterlike numeric character
Pc	U+005F	Low line
Mn	None	Nonspacing combining mark
Mc	None	Spacing combining mark
Cf	U+00AD	Soft Hyphen

Table C.5: Some general categories for the first 256 code points.

D | Common Language Infrastructure

The *Common Language Infrastructure (CLI)*, not to be confused with *Command Line Interface* with the same acronym, is a technical standard developed by Microsoft [4, 3]. The standard specifies a language, its format, and a runtime environment that can execute the code. The main feature is that it provides a common interface between many languages and many platforms, such that programs can collaborate in a language-agnostic manner and can be executed on different platforms without having to be recompiled. Main features of the standard are:

- Common Language Infrastructure
- CLI
- Command Line Interface

Common Type System (CTS) which defines a common set of types that can be used across different languages as if it were their own.

- Common Type System
- CTS

Metadata which defines a common method for referencing programming structures such as values and functions in a language-independent manner.

- Metadata

Common Intermediate Language (CIL) which is a platform-independent, stack-based, object-oriented assembly language that can be executed by the Virtual Execution System.

- Common Intermediate Language
- CIL

Virtual Execution System (VES) which is a platform dependent, virtual machine, which combines the above into code that can be executed at runtime. Microsoft's implementation of VES is called *Common Language Runtime (CLR)* and uses *just-in-time* compilation.

- Virtual Execution System
- VES
- Common Language Runtime

The process of running an F# program is shown in Figure D.1¹: First the F# code is compiled or interpreted to CIL. This code possibly combined with other CIL code is then converted to a machine-readable code, and the result is then executed on the platform.

- CLR
- just-in-time

CLI defines a *module* as a single file containing executable code by VES. Hence, CLI's notion of a module is somewhat related to F#'s notion of module, but the two should not be confused. A collection of modules, a *manifest*, and possibly other resources, which jointly define a complete program is called an *assembly*. The manifest is the description of which files are included in the assembly together with its version, name, security information, and other bookkeeping information.

- module
- manifest
- assembly

¹Jon: update figure

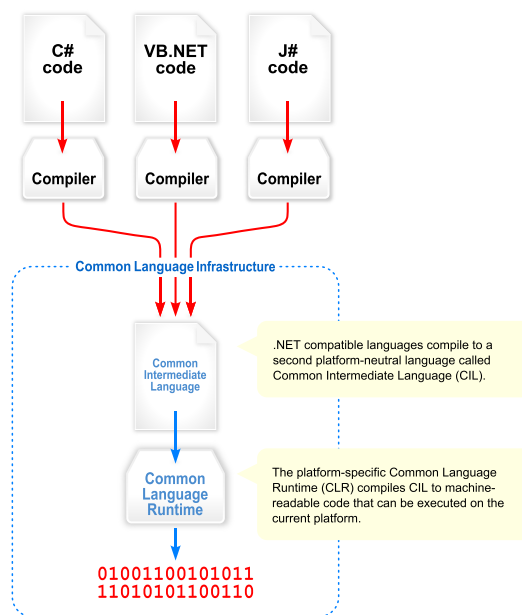


Figure D.1: Visual overview of CLI. Figure by Jarkko Piironen - Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=3602584>.

E | Language Details

This appendix lists various language details.

E.1 Arithmetic operators on basic types

Operator	leftOp	rightOp	Expression	Result	Description
leftOp + rightOp	ints	ints	5 + 2	7	Addition
	floats	floats	5.0 + 2.0	7.0	
	chars	chars	'a' + 'b'	'\195'	Addition of codes
	strings	strings	"ab" + "cd"	"abcd"	Concatenation
leftOp - rightOp	ints	ints	5 - 2	3	Subtraction
	floats	floats	5.0 - 2.0	3.0	
leftOp * rightOp	ints	ints	5 * 2	10	Multiplication
	floats	floats	5.0 * 2.0	10.0	
leftOp / rightOp	ints	ints	5 / 2	2	Integer division
	floats	floats	5.0 / 2.0	2.5	Division
leftOp % rightOp	ints	ints	5 % 2	1	Remainder
	floats	floats	5.0 % 2.0	1.0	
leftOp ** rightOp	floats	floats	5.0 ** 2.0	25.0	Exponentiation
leftOp && rightOp	bool	bool	true && false	false	boolean and
leftOp rightOp	bool	bool	true false	false	boolean or
leftOp &&& rightOp	ints	ints	0b1010 &&& 0b1100	0b1000	bitwise bool and
leftOp rightOp	ints	ints	0b1010 0b1100	0b1110	bitwise boolean or
leftOp ^^^ rightOp	ints	ints	0b1010 ^^^ 0b1101	0b0111	bitwise boolean exclusive or
leftOp <<< rightOp	ints	ints	0b00001100uy <<< 2	0b00110000uy	bitwise shift left
leftOp >>> rightOp	ints	ints	0b00001100uy >>> 2	0b00000011uy	bitwise and
+op	ints		+3	3	identity
	floats		+3.0	3.0	
-op	ints		-3	-3	negation
	floats		-3.0	-3.0	
not op	bool		not true	false	boolean negation
~~~op	ints		~~~0b00001100uy	0b11110011uy	bitwise boolean negation

Table E.1: Arithmetic operators on basic types. Ints, floats, chars, and strings means all built-in integer types etc. Note that for the bitwise operations, digits 0 and 1 are taken to be `true` and `false`.

Operator	leftOp	rightOp	Expression	Result	Description
leftOp < rightOp	bool	bool	true < false	false	Less than
	ints	ints	5 < 2	false	
	floats	floats	5.0 < 2.0	false	
	chars	chars	'a' < 'b'	true	
	strings	strings	"ab" < "cd"	true	
leftOp > rightOp	bool	bool	true > false	true	Greater than
	ints	ints	5 > 2	true	
	floats	floats	5.0 > 2.0	true	
	chars	chars	'a' > 'b'	false	
	strings	strings	"ab" > "cd"	false	
leftOp = rightOp	bool	bool	true = false	false	Equal
	ints	ints	5 = 2	false	
	floats	floats	5.0 = 2.0	false	
	chars	chars	'a' = 'b'	false	
	strings	strings	"ab" = "cd"	false	
leftOp <= rightOp	bool	bool	true <= false	false	Less than or equal
	ints	ints	5 <= 2	false	
	floats	floats	5.0 <= 2.0	false	
	chars	chars	'a' <= 'b'	true	
	strings	strings	"ab" <= "cd"	true	
leftOp >= rightOp	bool	bool	true >= false	true	Greater than or equal
	ints	ints	5 >= 2	true	
	floats	floats	5.0 >= 2.0	true	
	chars	chars	'a' >= 'b'	false	
	strings	strings	"ab" >= "cd"	false	
leftOp <> rightOp	bool	bool	true <> false	true	Not Equal
	ints	ints	5 <> 2	true	
	floats	floats	5.0 <> 2.0	true	
	chars	chars	'a' <> 'b'	true	
	strings	strings	"ab" <> "cd"	true	

Table E.2: Comparison operators on basic types. Ints, floats, chars, and strings means all built-in integer types etc..

## E.2 Basic arithmetic functions

Type	Function name	Example	Result	Description
Ints and floats	<code>abs</code>	<code>abs -3</code>	3	Absolute value
Floats	<code>acos</code>	<code>acos 0.8</code>	0.6435011088	Inverse cosine
Floats	<code>asin</code>	<code>asin 0.8</code>	0.927295218	Inverse sinus
Floats	<code>atan</code>	<code>atan 0.8</code>	0.6747409422	Inverse tangent
Floats	<code>atan2</code>	<code>atan2 0.8 2.3</code>	0.3347368373	Inverse tangentvariant
Floats	<code>ceil</code>	<code>ceil 0.8</code>	1.0	Ceiling
Floats	<code>cos</code>	<code>cos 0.8</code>	0.6967067093	Cosine
Floats	<code>cosh</code>	<code>cosh 0.8</code>	1.337434946	Hyperbolic cosine
Floats	<code>exp</code>	<code>exp 0.8</code>	2.225540928	Natural exponent
Floats	<code>floor</code>	<code>floor 0.8</code>	0.0	Floor
Floats	<code>log</code>	<code>log 0.8</code>	-0.2231435513	Natural logarithm
Floats	<code>log10</code>	<code>log10 0.8</code>	-0.09691001301	Base-10 logarithm
Ints, floats, chars, and strings	<code>max</code>	<code>max 3.0 4.0</code>	4.0	Maximum
Ints, floats, chars, and strings	<code>min</code>	<code>min 3.0 4.0</code>	3.0	Minimum
Ints	<code>pown</code>	<code>pown 3 2</code>	9	Integer exponent
Floats	<code>round</code>	<code>round 0.8</code>	1.0	Rounding
Ints and floats	<code>sign</code>	<code>sign -3</code>	-1	Sign
Floats	<code>sin</code>	<code>sin 0.8</code>	0.7173560909	Sinus
Floats	<code>sinh</code>	<code>sinh 0.8</code>	0.8881059822	Hyperbolic sinus
Floats	<code>sqrt</code>	<code>sqrt 0.8</code>	0.894427191	Square root
Floats	<code>tan</code>	<code>tan 0.8</code>	1.029638557	Tangent
Floats	<code>tanh</code>	<code>tanh 0.8</code>	0.6640367703	Hyperbolic tangent

Table E.3: Predefined functions for arithmetic operations

Name	Example	Description
<code>fst</code>	<code>fst (1, 2)</code>	
<code>snd</code>	<code>snd (1, 2)</code>	
<code>failwith</code>	<code>failwith</code>	
<code>invalidArg</code>	<code>invalidArg</code>	
<code>raise</code>	<code>raise</code>	
<code>reraise</code>	<code>reraise</code>	
<code>ref</code>	<code>ref</code>	
<code>ceil</code>	<code>ceil</code>	

Table E.4: Built-in functions.

### E.3 Precedence and associativity

· boolean or  
· boolean and

Operator	Associativity	Description
ident "<" types ">"	Left	High-precedence type application
ident "(" expr ")"	Left	High-precedence application
"."	Left	
prefixOp	Left	All prefix operators
" rule	Left	Pattern matching rule
ident expr, "lazy'' expr, "assert'' epxr	Left	
"**" opChar	Right	Exponent like
"*" opChar, "/" opChar, "%" opChar	Left	Infix multiplication like
"-" opChar, "+" opChar	Left	Infix addition like
":?'"	None	
"::'"	Right	
"^'" opChar	Right	
"!=" opChar, "<" opChar, ">" opChar, "=", " " opChar, "&" opChar, "\$" opChar	Left	Infix addition like
":>", "?:>"	Right	
"&", "&&"	Left	Boolean and like
"or", "  "	Left	Boolean or like
","	None	
":="	Right	
"->"	Right	
"if"	None	
"function", "fun", "match", "try"	None	
"let"	None	
";"	Right	
" "	Left	
"when"	Right	
"as"	Right	

Table E.5: Precedence and associativity of operators. Operators in the same row has same precedence. See ?? for the definition of `prefixOp`

# Bibliography

- [1] Alonzo Church. An unsolvable problem of elementary number theory. *American Journal of Mathematics*, 58:345—363, 1936.
- [2] Ole-Johan Dahl and Kristen Nygaard. SIMULA a language for programming and description of discrete event systems. introduction and user’s manual. Technical report, Norwegian Computing Center, 1967.
- [3] European Computer Manufacturers Association (ECMA). Standard ecma-335, common language infrastructure (cli). <http://www.ecma-international.org/publications/standards/Ecma-335.htm>.
- [4] International Organization for Standardization. Iso/iec 23271:2012, common language infrastructure (cli). <https://www.iso.org/standard/58046.html>.
- [5] Object Management Group. Uml version 2.0. <http://www.omg.org/spec/UML/2.0/>.
- [6] Programming Research Group. Specifications for the ibm mathematical formula translating system, fortran. Technical report, Applied Science Division, International Business Machines Corporation, 1954.
- [7] John McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, 1960.
- [8] X3: ASA Sectional Committee on Computers and Information Processing. American standard code for information interchange. Technical Report ASA X3.4-1963, American Standards Association (ASA), 1963. <http://worldpowersystems.com/projects/codes/X3.4-1963/>.
- [9] George Pólya. *How to solve it*. Princeton University Press, 1945.
- [10] Alan M. Turing. On computable numbers, with an application to the entscheidungsproblem. *Proceedings of the London Mathematical Society*, s2-42(1):230–265, 1936.

# Index

(**), 65  
->, 47  
., 144, 147  
//, 65  
:, 39, 162  
::, 112, 156  
:>, 217  
:?, 163  
:?>, 218  
;, 39  
<<, 166  
>>, 166  
_, 39  
#r directive, 87  
Item, 201  
System.Console.ReadKey, 187  
System.Console.ReadLine, 187  
System.Console.Read, 187  
System.Console.WriteLine, 187  
System.Console.Write, 187  
System.ConsoleKeyInfo.KeyChar, 188  
System.ConsoleKeyInfo.Key, 188  
System.ConsoleKeyInfo.Modifiers, 188  
System.IO.Directory.CreateDirectory, 192  
System.IO.Directory.Delete, 192  
System.IO.Directory.Exists, 192  
System.IO.Directory.GetCurrentDirectory, 192  
System.IO.Directory.GetDirectories, 192  
System.IO.Directory.GetFiles, 192  
System.IO.Directory.Move, 192  
System.IO.Directory.SetCurrentDirectory, 192  
System.IO.File.Copy, 192  
System.IO.File.CreateText, 189  
System.IO.File.Create, 189  
System.IO.File.Delete, 192  
System.IO.File.Exists, 192  
System.IO.File.Move, 192  
System.IO.File.OpenRead, 189  
System.IO.File.OpenText, 189  
System.IO.File.OpenWrite, 189  
System.IO.File.Open, 189  
System.IO.FileMode.Append, 189  
System.IO.FileMode.CreateNew, 189  
System.IO.FileMode.Create, 189  
System.IO.FileMode.OpenOrCreate, 189  
System.IO.FileMode.Open, 189  
System.IO.FileMode.Truncate, 189  
System.IO.FileStream.CanRead, 190  
System.IO.FileStream.CanSeek, 190  
System.IO.FileStream.CanWrite, 190  
System.IO.FileStream.Close, 190  
System.IO.FileStream.Flush, 190  
System.IO.FileStream.Length, 190  
System.IO.FileStream.Name, 190  
System.IO.FileStream.Position, 190  
System.IO.FileStream.ReadByte, 190  
System.IO.FileStream.Read, 190  
System.IO.FileStream.Seek, 190  
System.IO.FileStream.WriteByte, 190  
System.IO.FileStream.Write, 190  
System.IO.Path.Combine, 193  
System.IO.Path.GetDirectoryName, 193  
System.IO.Path.GetExtension, 193  
System.IO.Path.GetFileNameWithoutExtension, 193  
System.IO.Path.GetFileName, 193  
System.IO.Path.GetFullPath, 193  
System.IO.Path.GetTempFileName, 193  
System.IO.StreamReader.Close, 191  
System.IO.StreamReader.EndOfStream, 191  
System.IO.StreamReader.Flush, 191  
System.IO.StreamReader.Peek, 191  
System.IO.StreamReader.ReadLine, 191  
System.IO.StreamReader.ReadToEnd, 191  
System.IO.StreamReader.Read, 191  
System.IO.StreamWriter.AutoFlush, 191  
System.IO.StreamWriter.Close, 191  
System.IO.StreamWriter.Flush, 191  
System.IO.StreamWriter.WriteLine, 191  
System.IO.StreamWriter.Write, 191  
abs, 292  
acos, 292  
asin, 292  
atan2, 292  
atan, 292  
ceil, 292  
cosh, 292  
cos, 292  
exp, 292  
floor, 292



- log10, 292
- log, 292
- max, 292
- min, 292
- pown, 292
- round, 292
- sign, 292
- sinh, 292
- sin, 292
- sqrt, 292
- stderr, 187
- stdin, 187
- stdout, 187
- swap, 225
- tanh, 292
- tan, 292
- _, 148, 154
- &, 159
- (), 16
- [], 33, 110, 118
- :=, 58
- ;;, 18
- <-, 55
- [], 156
- abstract class, 218
- [abstract member](#), 218
- [<AbstractClass>], 218
- accessors, 200, 243
- active patterns, 159
- aggregation, 235
- aliasing, 60
- American Standard Code for Information Inter-  
change, 283
- and, 28
- [and](#), 132, 138
- anonymous function, 47
- anonymous functions, 153, 164
- anonymous variable type, 148
- ArgumentException, 172
- array pattern, 157
- Array.append, 120
- Array.compareWith, 120
- Array.concat, 121
- Array.contains, 121
- Array.copy, 121
- Array.create, 121
- Array.empty, 121
- Array.exists, 121
- Array.fill, 122
- Array.filter, 122
- Array.find, 122
- Array.findIndex, 122
- Array.fold, 122
- Array.foldBack, 123
- Array.forall, 123
- Array.get, 123
- Array.init, 123
- Array.isEmpty, 123
- Array.iter, 123
- Array.length, 124
- Array.map, 124
- Array.ofList, 124
- Array.rev, 124
- Array.set, 124
- Array.sort, 124
- Array.sub, 125
- Array.toList, 125
- Array.unzip, 125
- Array.zip, 125
- Array2D, 126
- Array2D.copy, 128
- Array2D.create, 128
- Array2D.get, 128
- Array2D.init, 128
- Array2D.iter, 129
- Array2D.length1, 129
- Array2D.length2, 129
- Array2D.map, 129
- Array2D.set, 129
- Array3D, 126
- Array4D, 126
- arrays, 117
- [as](#), 174
- ASCII, 283
- ASCIIbetical order, 32, 283
- assembly, 287
- association, 234
- attributes, 197, 198
- base, 21, 280
- [base](#), 216
- base class, 215
- bash, 276
- Basic Latin block, 284
- Basic Multilingual plane, 284
- basic types, 20
- binary number, 21, 280
- binary operator, 25, 51
- binary64, 280
- binding, 15
- bit, 21, 280
- Bitmap, 245
- black-box testing, 91
- block, 42
- blocks, 284
- bool, 20
- boolean and, 293
- boolean or, 293
- branch, 76
- branching coverage, 92
- bug, 90

- Button, 258
- byte, 280
- byte[], 23
- byte, 23
- call stack, 58, 135
- call-back, 269
- call-back function, 246
- call-back functions, 10
- casting exceptions, 172
- catching exception, 172
- cd, 274, 277
- char, 20, 22
- character, 22
- CheckBox, 258
- CIL, 287
- class, 25, 34, 197
- class, 197
- class diagram, 233
- CLI, 241, 287
- client coordinates, 246
- Clone, 120
- close file, 185
- closures, 49, 164
- CLR, 287
- code point, 22, 284
- Color, 243
- Command Line Interface, 287
- command-line interface, 241
- Common Intermediate Language, 287
- Common Language Infrastructure, 287
- Common Language Runtime, 287
- Common Type System, 287
- compile mode, 12
- composition, 236
- condition, 71
- conjunctive patterns, 159
- console, 12, 272
- constant pattern, 154
- constructor, 147, 197, 198
- control, 241, 257
- copy constructor, 203
- coverage, 92
- create file, 185
- CTS, 287
- currying, 167
- DateTime, 256
- DateTimePicker, 258
- debugging, 14, 17, 91, 98
- decimal, 23
- decimal number, 21, 280
- decimal point, 16, 21, 280
- declarative programming, 10
- default, 218
- del, 275
- delete file, 185
- derived class, 215
- digit, 21, 280
- dir, 274
- directory, 272
- discriminated union patterns, 157
- discriminated unions, 212
- disjunctive pattern, 159
- Dispose, 194
- DivideByZeroException, 172
- do, 15, 51, 71, 72
- do binding, 51
- do-binding, 15
- done, 71, 72
- dot notation, 34
- double, 280
- double, 23
- downcast, 218
- downcasting, 25
- dynamic type pattern, 163, 173
- echo, 275, 278
- efficiency, 90
- elif, 76
- else, 76
- encapsulate, 18
- encapsulation, 44, 49, 59
- end, 197
- EntryPoint, 186
- enumerations, 141
- enums, 141
- environment, 99
- eprintf, 54
- eprintfn, 54
- error message, 17
- escape sequences, 22
- event, 241
- event-driven programming, 10, 241, 269
- event-loop, 242
- events, 269
- exception, 30
- exclusive or, 31
- executable file, 12
- exit status, 186
- exn, 20, 172
- expression, 10, 15, 25, 38
- Extensible Markup Language, 65
- failwithf, 54
- failwith, 178
- fields, 197, 198
- file, 185
- first-class citizens, 49
- first-class citizenship, 169
- float, 20
- float32, 23

- floating point number, 21
- floating point numbers, 16
- `FlowLayoutPanel`, 264
- flushing, 190
- fold, 169
- foldback, 169
- folder, 272
- `Font`, 245
- for, 110
- `for`, 72, 153
- form, 241
- format string, 16
- `fprintf`, 54
- `fprintfn`, 54
- fractional part, 21, 25
- `fst`, 62
- `fun`, 47, 153
- function, 10, 15, 18
- function composition, 166
- function coverage, 92
- functional programming, 10, 130
- functionality, 90
- functions, 197, 198
  
- GDI+, 241
- generic function, 45
- graphical user interface, 241, 272
- `Graphics`, 246
- `Graphics.DrawLine`, 246
- GTK+, 241
- guard, 156
- GUI, 241, 272
  
- hand tracing, 99
- handlers, 269
- handling exception, 172
- has-a, 236
- `Head`, 113
- `Tail`, 113
- hexadecimal number, 21, 280
- higher order function, 164, 169
- how, 232
- HTML, 69
- Hyper Text Markup Language, 69
  
- identifier, 37
- IEEE 754 double precision floating-point format, 280
- `if`, 76
- `ignore`, 54
- `Image`, 245
- immutable state, 169
- Imperative programming, 130
- imperative programming, 9, 170
- implementation file, 12
- implementation files, 87
- `in`, 39, 110
- `IndexOutOfRangeException`, 172
- infix notation, 25
- inheritance, 210, 215, 234
- `inline`, 149
- input pattern, 152
- instantiation, 197
- `int`, 20
- `int16`, 23
- `int32`, 23
- `int64`, 23
- `int8`, 23
- integer, 21
- integer division, 30
- integer number, 15
- interactive mode, 12
- interface, 198, 220, 234
- `interface with`, 220
- `invalidArg`, 180
- invariant, 75
- is-a, 234
- is-a relation, 215
- `IsEmpty`, 113
- `isEmpty`, 110
- `IsNone`, 181
- `IsSome`, 181
- `it`, 16, 20
- `iter`, 169
  
- jagged arrays, 125
- just-in-time, 287
  
- keyword, 15, 38
  
- `Label`, 258
- Landau symbol, 137
- Latin-1 Supplement block, 284
- `Latin1`, 283
- `lazy`, 170
- lazy evaluation, 170
- least significant bit, 280
- `Length`, 107, 110, 113, 120
- length, 62
- `let`, 15, 39, 138, 151
- let-binding, 15
- lexeme, 18
- lexical scope, 17, 46
- lexically, 39
- library, 81
- library file, 12
- lightweight syntax, 39, 40
- list, 110
- list concatenation, 112
- list cons, 112
- list pattern, 156
- `List.collect`, 113

- List.contains, 114
- List.empty, 114
- List.exists, 114
- List.filter, 114
- List.find, 114
- List.findIndex, 114
- List.fold, 115
- List.foldBack, 115
- List.forall, 115
- List.head, 115
- List.isEmpty, 115
- List.item, 115
- List.iter, 116
- List.Length, 116
- List.map, 116
- List.ofArray, 116
- List.rev, 116
- List.sort, 116
- List.tail, 117
- List.toArray, 117
- List.unzip, 117
- List.zip, 117
- listeners, 269
- literal, 20
- literal type, 23
- loop invariant, 75
- lower camel case, 38
- ls, 277
- machine code, 130
- maintainability, 91
- manifest, 287
- map, 169
- match, 133, 152
- member, 25, 62
- members, 197
- MessageBox, 260
- Metadata, 287
- method, 34
- methods, 197, 198, 242
- mixed case, 38
- mkdir, 274, 278
- mockup code, 98
- mockup functions, 153
- Model-View paradigm, 248
- models, 197
- module, 81, 287
- module, 82
- mono32, 241
- most significant bit, 280
- move, 275
- multicase active patterns, 162
- multidimensional arrays, 125
- mutable, 55
- mutable data, 56
- mutable values, 130
- mutually recursive, 138
- mv, 278
- namespace, 25, 84
- namespace, 85
- namespace pollution, 83
- NaN, 282
- nested scope, 42
- new, 147, 198, 208
- newline, 22
- None, 181
- not, 28
- not a number, 282
- NotFiniteNumberException, 172
- nouns, 233
- nouns-and-verbs method, 233
- obfuscation, 64
- obj, 20
- object, 10, 34, 197
- object oriented analysis and design, 232
- Object oriented programming, 130
- object-oriented analysis, 197
- object-oriented design, 197
- Object-oriented programming, 10, 232
- object-oriented programming, 197
- objects, 197
- octal number, 21, 280
- open, 83
- open file, 185
- OpenFileDialog, 261
- operand, 45
- operands, 25
- operator, 26, 45, 51
- operator overloading, 206
- option type, 142, 181
- Option.bind, 182
- or, 28
- overflow, 29
- OverflowException, 172
- overload, 147
- overloading, 205
- override, 147, 210, 215
- override, 218
- overshadow, 216
- package, 236
- Paint.Add, 246
- Panels, 263
- partial pattern functions, 161
- partial specification, 166
- pascal case, 38
- Pen, 245
- pipng, 48, 165
- Point, 245
- PointToClient, 246

- PointToScreen, 246
- portability, 91
- precedence, 26
- prefix operator, 26
- primary constructor, 198
- primitive types, 140
- printf, 52, 54
- printfn, 15, 54
- private, 197
- problem statement, 233, 236
- procedure, 49
- procedures, 130
- ProgressBar, 258
- properties, 197, 198, 242
- public, 197
- pure function, 169
  
- RadioButton, 258
- ragged multidimensional list, 112
- raise, 176
- raising exception, 172
- range expressions, 110, 118
- read file, 185
- reals, 280
- rec, 84, 132, 138
- record pattern, 157
- recursion, 26, 169
- recursion step, 133
- recursive function, 132
- redirection, 275, 278
- reduce, 169
- ref, 58
- reference cells, 58
- reference types, 119
- referential transparency, 169
- reliability, 90
- remainder, 30
- rm, 278
- rmdir, 275, 278
- rounding, 25
- runtime error, 30
- runtime resolved variable type, 148
  
- SaveFileDialog, 262
- sbyte, 23
- scientific notation, 21
- scope, 41
- screen coordinates, 246
- script file, 12
- script files, 87
- script-fragment, 12, 18
- scripts, 12
- search path, 276
- search-path, 278
- self identifier, 197–199
- seq, 170
- sequence expression, 110
- sequence expressions, 117
- side-effect, 118
- side-effects, 49, 59, 130
- signature file, 12
- signature files, 87
- single, 23
- slicing, 118
- snd, 62
- software testing, 91
- SolidBrush, 245
- Some, 181
- source code, 12
- sprintf, 54
- stack frame, 135
- state, 10
- statement, 10, 51
- statement coverage, 92
- statements, 15, 130
- states, 130
- static type pattern, 162
- statically resolved variable type, 148
- stderr, 54
- stdout, 54
- stopping condition, 133
- stopping step, 133
- stream, 185
- string, 16, 22, 107
- string, 20
- String.collect, 108
- String.concat, 108
- String.exists, 108
- String.forall, 108
- String.init, 108
- String.iter, 109
- String.iteri, 109
- String.length, 109
- String.map, 109
- String.mapi, 109
- String.replicate, 109
- strongly typed, 170
- struct records, 144
- structs, 146
- structured programming, 10
- structures, 146
- subnormals, 282
- System.Drawing, 241
- System.IDisposable, 194
- System.Object, 220
- System.string, 107
- System.Windows.Forms, 241
  
- tail-recursion, 138
- terminal, 272
- TextBox, 258
- TextureBrush, 245

- The Heap, 58, 142, 144, 198, 202
- The Stack, 58, 135
- [then](#), 76
- Timer, 256
- `ToString()`, 147
- tracing, 99
- truth table, 28
- tuple, 61
- type, 16, 20
- type abbreviation, 140
- type aliasing, 140
- type constraints, 150
- type declaration, 16
- type inference, 14, 16
- type safety, 45
- typecasting, 24
  
- `uint16`, 23
- `uint32`, 23
- `uint64`, 23
- `uint8`, 23
- UML, 233
- unary operator, 51
- underflow, 29
- Unicode, 22
- Unicode general category, 284
- Unicode Standard, 284
- Unified Modelling Language 2, 233
- Uniform Resource Identifiers, 194
- Uniform Resource Locator, 193
- `unit`, 16, 20
- unit testing, 91
- unit-testing, 14
- upcast, 217
- upcasting, 25
- upper camel case, 38
- URI, 194
- URL, 193
- usability, 90
- [use](#), 194
- use case, 233
- user story, 233
- `using`, 195
- UTF-16, 284
- UTF-8, 284
  
- [val](#), 16, 147
- Value, 181
- value-binding, 39
- variable, 56
- variable pattern, 155
- variable types, 148
- verbatim, 24
- verbatim string, 107
- verbose syntax, 39
- verbs, 233, 238
  
- VES, 287
- Virtual Execution System, 287
  
- what, 232
- [when](#), 150, 156
- [while](#), 71
- white-box testing, 91, 92
- whitespace, 22
- whole part, 21, 25
- wildcard, 39
- wildcard pattern, 154
- Windows command line, 272
- Windows Graphics Device Interface, 241
- WinForms 2.0, 241
- [with](#), 133, 146, 152
- word, 280
- write file, 185
  
- XML, 65
- xor, 31