



Chapter 19

Object-oriented programming

also known as
or attributes

Object-oriented programming is a programming paradigm that focusses on *objects* such as a person, place, thing, event, and concept relevant for the problem. Objects may contain data and code, which in the object-oriented paradigm are called *properties* and *methods*. Object-oriented programming is an extension of data types, in the sense that objects contains both data and functions in a similar manner as a module, but object-oriented programming emphasizes the semantic unity of the data and functions. Thus, objects are often *models* of real world entities, and object-oriented programming leads to a particular style of programming analysis and design called *object-oriented analysis and design*.

- Object-oriented programming
- objects
- properties
- methods
- models
- object-oriented analysis
- object-oriented design

19.1 Constructors and members

An object is a variable of a class type. A class is defined using the “type” keyword, and there are *always* parentheses after the class name.

Listing 19.1, class.fsx:

A class definition and an object of this class.

```

1 type aClass(anArgument : int) =
2     // Constructor section of aClass
3     do printfn "A class has been created"
4     let objectValue = anArgument
5     // Member section of aClass
6     member this.value = anArgument
7     member this.scale (factor : int) = factor * objectValue
8
9 let a = aClass(2)
10 printfn "%d %d" a.value (a.scale(3))

```

```

1 $ fsharp --nologo class.fsx && mono class.exe
2 A class has been created
3 2 6

```

In the example, the class `aClass` is defined in the header in line 1, and it includes one integer argument, `anArgument`. Classes can also be defined without arguments, but the parentheses cannot be omitted. The body of the class line 2–7 is indicated by the indentation. The arguments are immutable. The body consists of two parts: The *constructor section* in line 2–4 and *member section* in line 5–7. In

- constructor section
- member section

but the constructor is implied in line 1.
primary

line 9 is an object `a` of `aClass` type created, which implies that memory is reserved and initialisation code is run, and in line 10 is the object used.

In object-oriented lingo, the initialization code is called the *constructor*, and in contrast to many other languages, the constructor is implicitly stated. It is called the *primary constructor*, the arguments given in the header are the primary constructor's arguments, and the primary constructor's body are the "let" and "do" statements following the header. Members are not available in the constructor unless the self-identifier has been declared in the header using the keyword "as", e.g., type `classMutable(` `name : string`) `as this = ...` → elaborate here. why do this? Example (concrete).

Class may have members declared using the keyword "member", which must come after the constructor, and which defines values and functions that are accessible from outside the class using the "." notation. In this manner, the members define the interface between the internal bindings in the constructor and an application program. In object-oriented lingo the value and variable members are called *properties* and functions are called *methods*. Members are values, i.e., immutable. In the example, line 6 and 7 defined a property and a method.

In the class definition in Listing 19.1 we bind the primary constructor's arguments to the member values `this.value`. The prefix `this.` is a *self-identifier* used in the definition of the class such that this.value is the name of the objectValue value for the particular object being constructed. E.g., when `a` is created in line 9, `this.value` refers to `a.value`. As a quirk, F# is very flexible regarding what name can be used for the self-identifier, and the member section could as well have been `self.value`, `_.value`, or anything else, and it need not be the same in every member definition. In Listing 19.1 we also declare a member function, `this.scale : int -> int`.

The body of member has access to arguments, the primary constructor's bindings, and to all class members, regardless of the member's lexicographical order.

Member values, member functions belong to objects, and the implication is the example value and `scale` 'resides' on or 'belongs' to each object.

→ I find this very confusing. I will discuss it with you.

As an aside, if we wanted to use a tuple argument for the class, then this must be explicitly annotated, since the call to the constructor looks identical. This is demonstrated in Listing 19.2.

Listing 19.2 classTuple.fsx:

Beware: Creating objects from classes with several arguments and tuples have the same syntax.

```
1 type vectorWTupleArgs(x : float * float) =  
2   member this.cartesian = x  
3 type vectorWTwoArgs(x : float, y : float) =  
4   member this.cartesian = (x,y)  
5 let v = vectorWTupleArgs(1.0, 2.0)  
6 let w = vectorWTwoArgs(1.0, 2.0)
```

Whether the full list of arguments should be transported from the caller to the object as a tuple or not is a matter of taste that mainly influences the header of the class. The same cannot be said about how the elements of the vector are stored inside the object and made accessible outside the object. In Listing 19.2, the difference between storing the vector's elements in individual members `member this.x = x` and `member this.y = y`, or as a tuple member `this.cartesian = (x, y)`, is that in order to access the first element in a vector `v` an application program would in the first case must write `v.x`, while in the second case the application program must first retrieve the tuple and then extract the first element, e.g., as `fst v.cartesian`. Which is to be preferred depends very much on the application: Is it the individual elements or the complete tuple of elements that is to have focus, when using the objects.

¹Jon: define "do" statement somewhere used in loops!

★ Advice about consistent self-identifiers?

177

class is like a library?
caller → application

Said ~~in a~~ differently, which choice will make the easiest to read application program with the lowest risk of programming errors. Hence, when designing classes, consider carefully how application programs will use the class, and aim for simplicity and versatility while minimizing the risk of error in the application program. Advice

19.2 Accessors

Methods are most often used as an interface between the local bindings of an object and the application program. Consider the example in Listing 19.3.

Listing 19.3, classAccessor.fsx:
Accessor methods interface with internal bindings.

```
1 type aClass() =  
2   let mutable v = 1  
3   member this.setValue (newValue : int) : unit =  
4     v <- newValue  
5   member this.getValue () : int = v  
6 let a = aClass()  
7 printfn "%d" (a.getValue ())  
8 a.setValue(2)  
9 printfn "%d" (a.getValue ())  
  
1 $ fsharp --nologo classAccessor.fsx && mono classAccessor.exe  
2 1  
3 2
```

In the above example, the methods `setValue` and `getValue` set and get the state of the objects. Such functions are called *accessors*. Internal states with setters and getters are a typical construction, since it allows for complicated computations, when states are read to and written from, and gives the designer of the class the freedom to change the internal representation while keeping the interface the same. Accessors are so common that F# includes a special syntax for them: Classes can be made to act like variables using “member”...“with”...“and” keywords and the special function bindings “get()” and “set()” as demonstrated in Listing 19.4.

So far, I would not know what set and get actually do.

I would be confused about:

- what type of complicated computations can be done (when I do not even understand the simple computations that they are meant to do)
- what ‘states’ exactly means
- what exactly is ‘change internal representation’

All this should come after an explanation of what set & get do, and why, and they should be expanded with examples.

Listing 19.4, classGetSet.fsx:

Members can act as variables with the builtin getter and setter functions.

```

1 type aClass() =
2     let mutable v = 0
3     member this.value
4         with get() = v
5           and set(a) = v <- a
6
7 let a = aClass()
8 printfn "%d" a.value
9 a.value <- 2
10 printfn "%d" a.value

```

```

1 $ fsharp --nologo classGetSet.fsx && mono classGetSet.exe
2 0
3 2

```

The expression defining `get: () -> 'a` and `set: 'a -> ()`, where 'a is any type, can be any usual expression. The application calls the `get` and `set` as if the member was a mutable value. If `set` is omitted, then the member act as a value rather than a variable, and values cannot be assigned to it in the application program.

Setters and getters are so common that F# has a short-hand for this using “`member val value = 0 with get, set`”, which creates the internal mutable value `value`, but this is discouraged in this text.

Defining a *Item* member with extended `get` and `set` makes objects act as indexed variables as demonstrated in Listing 19.5.

Listing 19.5, classGetSetIndexed.fsx:

Members can act as index variables with the builtin getter and setter functions.

```

1 type aClass(size : int) =
2     let arr = Array.create<int> size 0
3     member this.Item
4         with get(ind : int) = arr.[ind]
5           and set(ind : int) (p : int) = arr.[ind] <- p
6
7 let a = aClass(3)
8 printfn "%A" a
9 printfn "%d %d %d" a.[0] a.[1] a.[2]
10 a.[1] <- 3
11 printfn "%d %d %d" a.[0] a.[1] a.[2]

```

```

1 $ fsharp --nologo classGetSetIndexed.fsx && mono classGetSetIndexed.exe
2 ClassGetSetIndexed+aClass
3 0 0 0
4 0 3 0

```

Higher order indexed variables are defined by adding more indexing arguments to the definition of `get` and `set`.

Combinations of non-static member definitions is shown in Listing 19.6.

Listing 19.6 classMemberDefinition.fsx:

A large variation of class member definitions. This program intentionally does not compile, but demonstrates variation that will, and problematic lines are indicated by the in-code comments.

```

1 type Test() =
2     let letV = 0
3     [<DefaultValue>] val mutable valMutableV : int // Discouraged
4     let mutable letMutableV = 0
5     member val memberValV = 0 // Discouraged
6     member this.memberThisV = 0
7     member mutable this.memberMutableThisV = 0 // Error, mutable members
      are illegal
8     member val memberValGetV = 0 with get
9     member val memberValSetV = 0 with set // Error, must have get
10    member val memberValGetSetV = 0 with get, set
11    member this.memberThisGetV
12        with get() = letMutableV // Read from object variable
13    member this.memberThisSetV
14        with set(value) = letMutableV <- value // Save to object variable
15    member this.memberThisGetSetV
16        with get() = letMutableV // Read from object variable
17            and set(value) = letMutableV <- value // Save to object variable
18
19    let t = Test()
20    printfn "%A" t.letV // Error: internal
21    t.letV <- 1 // Error: internal
22    printfn "%A" t.valMutableV
23    t.valMutableV <- 1
24    printfn "%A" t.letMutableV // Error: internal
25    t.letMutableV <- 1 // Error: internal
26    printfn "%A" t.memberValV
27    t.memberValV <- 1 // Error: read-only
28    printfn "%A" t.memberThisV
29    t.memberThisV <- 1 // Error: read-only
30    printfn "%A" t.memberValGetV
31    t.memberValGetV <- 1 // Error: read-only
32    printfn "%A" t.memberValGetSetV
33    t.memberValGetSetV <- 1
34    printfn "%A" t.memberThisGetV
35    t.memberThisGetV <- 1 // Error: read-only
36    printfn "%A" t.memberThisSetV // Error: write-only
37    t.memberThisSetV <- 1
38    printfn "%A" t.memberThisGetSetV
39    t.memberThisGetSetV <- 1

```

(confusing)

The val-keyword in this context has not been discussed previously.² staticMemberV and staticMemberValV have the same interface. The [<DefaultValue>] val mutable valMutableV : int has not been discussed and is discouraged, but gives a mutable member value that is initialized to the type's default value, e.g., Unchecked.defaultof<int> in this case. [<DefaultValue>] is called an *attribute*, but will not be discussed further.³ Defining mutable member variables is illegal, but allowed using explicit attribute

²Jon: maybe mention "Explicit Field", <https://docs.microsoft.com/en-us/dotnet/fsharp/language-reference/members/explicit-fields-the-val-keyword>.

³Jon: Should attributes be included <https://docs.microsoft.com/en-us/dotnet/fsharp/language-reference/attributes>?

get and set functions. The definitions for `valMutableV`, `memberValGetSetV` and `memberThisGetSetV` gives the same interface to a mutable variable, but with slight variation in how their initial value is set, and how get and set actions can be programmed. In general, **minimize the use of constructions using the “val”-keyword in class definitions for brevity.** All the above holds for static definitions except [`<DefaultValue>`] `static val mutable staticValMutableV : int` is illegal. Advice

→ This last sentence is ungrammatical.

19.3 Objects are reference types

Objects are reference type values, implying that copying objects copies their references not their values. Consider the example in Listing 19.7.

Listing 19.7, `classReference.fsx`:

Objects are reference types means assignment is aliasing.

```

1 type aClass() =
2     let mutable v = 0
3     member this.value with get() = v and set(a) = v <- a
4
5 let a = aClass()
6 let b = a
7 a.value <- 2
8 printfn "%d %d" a.value b.value

```

```

1 $ fsharpc --nologo classReference.fsx && mono classReference.exe
2 2 2

```

Thus, the binding to `b` in line 6 is an alias to `a`, not a copy, and changing object `a` also changes `b`! This is a common cause of error, and you should **think of objects as arrays.** For this reason, it is often seen that classes implement a copy function, returning a new object with copied values, e.g., Listing 19.8. Advice

Listing 19.8, `classCopy.fsx`:

A copy method is often needed. Compare with Listing 19.7.

```

1 type aClass() =
2     let mutable v = 0
3     member this.value with get() = v and set(a) = v <- a
4     member this.copy() =
5         let o = aClass()
6         o.value <- v
7         o
8 let a = aClass()
9 let b = a.copy()
10 a.value <- 2
11 printfn "%d %d" a.value b.value

```

```

1 $ fsharpc --nologo classCopy.fsx && mono classCopy.exe
2 2 0

```

In the example we see that since `b` now is a copy, we do not change it by changing `a`.

19.4 Static classes

Classes can act as modules and hold properties and methods that are identical for all objects of its type. These are defined as *static members* using the “*static*”-keyword. They are accessed using the class name and not the object names. An example is given in Listing 19.9.

Listing 19.9, classStatic.fsx:

Static local variables and members are identical to all objects of the type.

```
1 type aClass(v : int) =
2     static let mutable c = 1
3     member this.value = c * v
4     static member factor with get() = c and set(a) = c <- a
5 let a = aClass(2)
6 let b = aClass(3)
7 printfn "%d %d" a.value b.value
8 aClass.factor <- 2 // Change value for all objects
9 printfn "%d %d" a.value b.value

-----

1 $ fsharpc --nologo classStatic.fsx && mono classStatic.exe
2 2 3
3 4 6
```

In line 8 the static member function is access changing the local static mutable variable, and since it is static, it is the same variable for both object a and b. Hence next time a.value and b.value are accessed, they both have changed.

19.5 Mutual recursive classes

Classes are inherently recursive and the body of the class can refer to static members of itself. For mutually recursive classes, use the keyword “*and*” as shown in Listing 19.10.

Listing 19.10, classAssymetry.fsx:

Mutually recursive classes are defined using the “*and*” keyword.

```
1 type anInt(v : int) =
2     member this.value = v
3     member this.add (w : float) : aFloat = aFloat((float this.value) + w)
4 and aFloat(w : float) =
5     member this.value = w
6     member this.add (v : int) : aFloat = aFloat((float v) + this.value)
7 let a = anInt(2)
8 let b = aFloat(3.2)
9 let c = a.add(b.value)
10 let d = b.add(a.value)
11 printfn "%A %A %A %A" a.value b.value c.value d.value

-----

1 $ fsharpc --nologo classAssymetry.fsx && mono classAssymetry.exe
2 2 3.2 5.2 5.2
```

Here `anInt` and `aFloat` are mutually dependent, and thus must be defined in the same “type” definition using “and”.

19.6 Function and operator overloading

Member functions may be overloaded, as shown in Listing 19.11.

Listing 19.11, `classOverload.fsx`:

Overloading methods. Members `set : int -> ()` and `set : int * int -> ()` are permitted since they differ in argument number or type.

```
1 type Greetings() =
2     let mutable greetings = "Hi"
3     let mutable name = "Programmer"
4     member this.str = greetings + " " + name
5     member this.setName(newName : string) : unit =
6         name <- newName
7     member this.setName(newName : string, newGreetings : string) : unit =
8         greetings <- newGreetings
9         name <- newName
10 let a = Greetings()
11 printfn "%s" a.str
12 a.setName("F# programmer")
13 printfn "%s" a.str
14 a.setName("Expert", "Hello")
15 printfn "%s" a.str

-----
1 $ fsharpc --nologo classOverload.fsx && mono classOverload.exe
2 Hi Programmer
3 Hi F# programmer
4 Hello Expert
```

In the example two methods are defined both named `set`, but with different number of arguments. This is called *overloading* and is allowed as long as the arguments differ in number or type.

· overloading

In Listing 19.10 the notation for addition is less than elegant. For such situations, F# supports operator overloading. To overload the “+” operator we overload its functional equivalence (+) as a static member as shown in Listing 19.12.

Listing 19.12, classOverloadOperator.fsx:
Operators can be overloaded using.

```
1 type anInt(v : int) =
2     member this.value = v
3     static member (+) (v : anInt, w : float) = (float v.value) + w
4     static member (+) (w : float, v : anInt) = v + w
5 let a = anInt(2)
6 let b = a + 3.2
7 let c = 3.2 + a
8 printfn "%A %A %A" a.value b c

-----

1 $ fsharp --nologo classOverloadOperator.fsx && mono
   classOverloadOperator.exe
2 2 5.2 5.2
```

refer to
earlier
chapter where
they have
been defined

All usual operators may be overloaded and the compiler uses type inference to decide which function is to be called.⁴ All operators have a functional equivalence, and writing $v + w$ is equivalent to writing `vector.(+) (v, w)`. Presently the former is to be preferred, but at times, e.g., when using functions as arguments, it is useful to be able to refer to an operator by its function equivalent. Note that the functional equivalence of the multiplication operator (*) shares a prefix with the begin block comment lexeme "(*", which is why the multiplication function is written as `(*)`. Note also that unitary operators have a special notation, and in the above case unitary minus should be defined as `static member (-) (v : anInt) = - v`. *I don't get this*

refer to
earlier
chapter
where this
was
introduced

In Listing 19.12, notice how the second (+) operator overloads the first by calling the first with the proper order of arguments. This is a general principle, **avoid duplication of code, reuse of existing code is almost always preferred**. Here it is to be preferred for two reasons. Firstly, if we discover a mistake in the multiplication code, then we need only correct it once, which implies that both multiplication methods are corrected once and reducing the chance of introducing new mistakes by attempting to correct old. Secondly, if we later decide to change the internal representation of the vector, then we only need to update one version of the multiplication function, hence we reduce programming time and risk of errors as well.

Can you give an example where this is useful?

Beware that operator overloading outside class definitions overwrites *all* definitions of the operator. E.g., overloading `vector.(+) (v, w)` outside a class will influence integer, real, string, etc. Thus, **operator overloading should only be done inside class definitions**.⁵

Advice

Overloading is a programming structure that does not work well in functional-first programming languages such as F#, since types are not easily inferred. Therefore, overloading has restricted usage. E.g., even in the above case, if the application program attempts to define the following function `let f = anInt.(+)`, the functional equivalent of the "+" operator, then the compiler will not be able to identify which of the two candidate functions is to be bound to `f`, and will throw an error. Although irrelevant until `f` is applied to arguments, an effective type inference system has yet to be implemented.

⁴Jon: Something about which the syntax for new operators allowed.

⁵Jon: Something about which operators can be overloaded <https://docs.microsoft.com/en-us/dotnet/fsharp/language-reference/operator-overloading>.

yes!
Also, why would we want to
overload operators? Motivation

19.7 Additional constructors

→ expand, explain, motivate

The constructor itself can be overloaded, to allow for varying parameters at time of object creation. Additional constructors may be created using the *new*-keyword as illustrated in Listing 19.13.

· new

Listing 19.13, classExtraConstructor.fsx:

Extra constructors can be added using “ ”.

```
1 type classExtraConstructor(name : string, greetings : string) =
2     static let defaultGreetings = "Hello"
3     // Additional constructors are defined by new()
4     new(name : string) =
5         classExtraConstructor(name, defaultGreetings)
6     member this.name = name
7     member this.greeting = greetings + " " + name
8
9 let s = classExtraConstructor("F#") // Calling additional constructor
10 let t = classExtraConstructor("F#", "Hi") // Calling primary constructor
11 printfn "\"%A\", \"%A\"\" s.greeting t.greeting"

-----

1 $ fsharp --nologo classExtraConstructor.fsx && mono
   classExtraConstructor.exe
2 "Hello F#", "Hi F#"
```

The body of the additional constructor must call the primary constructor, and the body cannot extend the primary constructor's values and functions. It is useful to think the primary constructor as a superset of arguments and the additional as subsets or specialisations. As regular scope rules dictate, the additional constructor has access to the primary constructor's bindings. However, in order to access the object's members, the self-identifier has to be explicitly declared using the “as”-keyword in the header. E.g., writing `new(x : float, y : float) as alsoThis = ...`. However beware, even though the body of the additional constructor now may access the member value `alsoThis.x`, this value has first been created once the primary constructor has been called. E.g., calling the primary constructor in the additional constructor as `new(x : float, y : float) as alsoThis = classExtraConstructor(fst alsoThis.x, y, defaultSeparator)` will cause an exception at runtime. Code may be executed in additional constructors: Before the call to the primary constructor, “let” and “do” statements are allowed. If code is to be executed after the primary constructor has been called, then it must be preceded by the “then” keyword as shown in Listing 19.14.

Advice